

# **Neural Radiance Fields**

CS194-26/294-26: Intro to Computer Vision and Computational  
Photography  
Angjoo Kanazawa  
UC Berkeley Fall 2022

**Lots of content from ECCV 2022 Tutorial on Neural  
Volumetric Rendering for Computer Vision**

# Capturing Reality



Earliest cave painting (45,500 years old) in Sulawesi, Indonesia

# Capturing Reality



Monet's Cathedral series: study of light 1893-1894

# Capturing Reality



First self-portrait Cornelius 1839




First Movie - Muybridge 1878

# Capturing Reality – in 3D



Building Rome in a Day, Agarwal et al. ICCV 2009

# Capturing Reality – in 3D



Google Earth 2016~

# **What is next?**

# 2020: Neural Radiance Field (NeRF)



Mildenhall\*, Srinivasan\*, Tancik\*, Barron, Ramamoorthi, Ng, ECCV 2020

# It has been two years

- Original NeRF paper: 1598 citations in 2 years



# Handling Appearance Changes



Nerf-W [Martin-Brualla et al. CVPR 2021]

# Real-time Rendering



Video from PlenOctrees [Yu et al. CVPR 2021]


# Real-time Inference

## INSTANT NEURAL GRAPHICS PRIMITIVES WITH A MULTIRESOLUTION HASH ENCODING

Thomas Müller Alex Evans Christoph Schied Alexander Keller

<https://nvlabs.github.io/instant-ngp>





@karenxcheng, with  
InstantNGP [Müller et  
al., SIGGRAPH 2022]

[GETTING STARTED](#)[GITHUB](#)[DOCUMENTATION](#)

VIEWPORT RENDER VIEW

0

1

2

▶ RESUME TRAINING

[Show Scene](#)[Show Images](#)[Refresh Page](#)

Resolution: 640x1024px

Time Allocation: 100% spent on viewer

Server Connected | Render Connected



CONTROLS



RENDER



SCENE

[LOAD PATH](#)[EXPORT PATH](#)

Height

1080

Width

1920

FOV

50

Seconds

4

FPS

24

[ADD CAMERA](#)

Smoothness

△



0.00

0

1

2

3

◀

◀

▶

▶

▶

CAMERA 0

 nerfstudio


# Dynamic NeRFs




[Xian et al., CVPR 2021]



HyperNeRF [Park et al., SigAsia 2021]  
Nerfies [Park et al., ICCV 2021]





# Generative 3D Faces

EG3D: Efficient Geometry-aware 3D Generative Adversarial Networks, Chan et al. CVPR 2022

Wang et al. SIGGRAPH 2022



© Disney

# City-Scale NeRFs




BlockNeRF  
[Tancik et al.  
CVPR 2022]

RawNeRF  
[Mildenhall et al.  
CVPR 2022]




RawNeRF  
[Mildenhall et al.  
CVPR 2022]


# Robotics



Dex-NeRF: Using a Neural Radiance field to Grasp Transparent Objects, [Ichnowski and Avigal et al. CoRL 2021]



NeRF-Supervision: Learning Dense Object Descriptors from Neural Radiance Fields, [Yen-Chen et al. ICRA 2022]



Vision-Only Robot Navigation in a Neural Radiance World  
[Adamkiewicz and Chen et al. ICRA 2022]

# Generating 3D scenes with diffusion models



DreamFusion  
[Poole et al.  
arXiv 2022]

# Goals of these lectures

- In 2 years, 1840 citations (as of November 28<sup>th</sup>) will not cover all these papers
- Visit the fundamentals in Neural Volumetric Rendering by abstracting away recent developments
- Provide first principles + background for you to go and read these papers & play around with the tools

# Menu

1. Birds Eye View & Background
2. Volumetric Rendering Function
3. Encoding and Representing 3D Volumes
4. Signal Processing Considerations
5. Challenges & Pointers



 nerfstudio

Capture of UC Berkeley redwoods with

# Birds Eye View & Background

# Birds Eye View

- What is NeRF?
- How is it different or similar to existing approaches?
- What is its historical context?

# Problem Statement

**Input:**


A set of calibrated Images



**Output:**


A 3D scene representation that  
renders novel views





# Three Key Components


$(x, y, z, \theta, \phi) \rightarrow F_\Omega \rightarrow (r, g, b, \sigma)$



Neural Volumetric 3D  
Scene Representation


Differentiable Volumetric  
Rendering Function

Objective: Synthesize  
all training views





Optimization via  
Analysis-by-Synthesis

## Representing a 3D scene as a continuous 5D function



What kind of a 3D representation is this?

# It is not a Mesh




Not a point cloud  
either



## It is volumetric

It's *continuous* voxels made of shiny transparent cubes

# What is the problem that is being solved?



# Plenoptic Function



Figure by Leonard McMillan


Q: What is the set of all things that we can ever see?

A: The Plenoptic Function (Adelson & Bergen '91)

Let's start with a stationary person and try to parameterize everything that they can see...

Slide credit:  
Alyosha Efros

# Grayscale Snapshot



$$P(\theta, \phi)$$

- **is intensity of light**
  - Seen from a single position (viewpoint)
  - At a single time
  - Averaged over the wavelengths of the visible spectrum

Slides from Alyosha Efros

# Color snapshot




$$P(\theta, \phi, \lambda)$$

- is intensity of light
  - Seen from a single position (viewpoint)
  - At a single time
  - As a function of wavelength

Slides from Alyosha Efros

# A movie




$$P(\theta, \phi, \lambda, t)$$

- is intensity of light
  - Seen from a single position (viewpoint)
  - Over time
  - As a function of wavelength

Slides from Alyosha Efros

# A holographic movie




$$P(\theta, \phi, \lambda, t, V_x, V_y, V_z)$$

- is intensity of light
  - Seen from ANY position and direction
  - Over time
  - As a function of wavelength

Slides from Alyosha Efros

# The plenoptic function




$$P(\theta, \phi, \lambda, t, V_x, V_y, V_z)$$

7D function, that can reconstruct every position & direction,  
at every moment, at every wavelength  
= it recreates the entirety of our visual reality!


Slides from Alyosha Efros

# Goal: Plenoptic Function from a set of images



- Objective: Recreate the visual reality
  - All about recovering photorealistic pixels, not about recording 3D point or surfaces
    - Image Based Rendering

# Goal: Plenoptic Function from a set of images




It is a conceptual device

Adelson & Bergen do not discuss how to solve this

# Plenoptic Function

7D function:  
2 – direction  
1 – wavelength  
1 – time  
3 – location



Look familiar




$$P(\theta, \phi, \cancel{\lambda}, \cancel{t}, V_x, V_y, V_z) \longrightarrow P(\theta, \phi, V_x, V_y, V_z)$$

Let's simplify:

1. Remove the time
2. Remove the wavelength & let the function output RGB colors

# An example of a sparse plenoptic function



If street view was super dense  
(360 view from any view point)  
then it is the Plenoptic Function

Levoy and Hanrahan, SIGGRAPH 1996  
Gortler et al. SIGGRAPH 1996

# Lightfield / Lumigraph

- An approach for modeling the Plenoptic Function
- Take a lot of pictures from many views
- Interpolate the rays to render a novel view

**Stanford Gantry**  
**128 cameras**



Levoy and Hanrahan, SIGGRAPH 1996  
Gortler et al. SIGGRAPH 1996

# Lightfield / Lumigraph

- An approach for modeling the Plenoptic Function
- Take a lot of pictures from many views
- Interpolate the rays to render a novel view






Figure from Marc Levoy

Levoy and Hanrahan, SIGGRAPH 1996  
Gortler et al. SIGGRAPH 1996

# Lightfield / Lumigraph

- An approach for modeling the Plenoptic Function
- Take a lot of pictures from many views
- Interpolate the rays to render a novel view



**Stanford Gantry  
128 cameras**



**Lytro camera**





Figure from Marc Levoy

Levoy and Hanrahan, SIGGRAPH 1996  
Gortler et al. SIGGRAPH 1996

# Lightfield / Lumigraph

Lightfields assume that the ray shooting out from a pixel is never occluded.



Because of this it only models the plenoptic surface:




Figure 1: The surface of a cube holds all the radiance information due to the enclosed object.


# How NeRF models the Plenoptic Function

$$P(\theta, \phi, V_x, V_y, V_z)$$


Look familiar  
😊?

NeRF takes the same input as the Plenoptic Function!

A subtle difference:



Plenoptic Function




NeRF

So NeRF requires the integration along the viewing ray to compute the Plenoptic Function

**Bottom line: it models the full plenoptic function!**

# 5D function





- For every location (3D), all possible views (2D) 
- NeRF models this space with a continuous view-dependent volume with opacity
- The color emitted by every point is composited to render a pixel
- Unlike a light field, the entire 5D plenoptic function can be modeled (you can fly through the world)

# Visualizing the 2D function on the sphere



Outgoing radiance distribution  
for point on side of ship



Outgoing radiance distribution  
for point on water's surface

# Baking in Light




- NeRF can capture non-Lambertian (specular, shiny surfaces) because it models the color in a view-dependent manner
- This is hard to do with meshes unless you model the physical materials & lighting interactions
- But, with Image Based Rendering — All lighting effects are baked in

# NeRF in a Slide


Objective: Reconstruct  
all training views



Volumetric 3D Scene  
Representation



Differentiable Volumetric  
Rendering Function



Optimization via  
Analysis-by-Synthesis

# Unmentioned caveat so far

- Training a NeRF requires a **calibrated** camera!!!!
- Need to know the camera parameters: extrinsic (viewpoint) & intrinsics (focal length, distortion, etc)



**How do we get this from images?**

# Structure from Motion

Or Photogrammetry (1850~)  
Long history in Computer Vision

*Proc. R. Soc. Lond. B.* 203, 405–426 (1979)

*Printed in Great Britain*

The interpretation of structure from motion

BY S. ULLMAN

*Artificial Intelligence Laboratory, Massachusetts Institute of Technology,  
545 Technology Square (Room 808), Cambridge, Massachusetts 02139 U.S.A.*

# NeRF is AFTER Structure from Motion


- In order to train NeRF you need to run SfM/SLAM on the images to estimate the camera parameters
- In this sense, the problem category is same as that of **Multi-view Stereo**



Colmap: Schönberger et al. 2016

# Multi-view Stereo


- Problem: Given calibrated cameras, recover highly detailed 3D **surface** model
- Dense photogrammetry, often the output is textured meshes



Figures by Carlos Hernandez, Yasutaka Furukawa

# Multi-View Stereo

Solutions to MVS is what you see for any existing 3D scanning system, ie sketchfab, or what's in your video game



# Multi-View Stereo


Because they often model surfaces, struggles on Thin / Amorphous / Shiny objects





 nerfstudio

# Where NeRF stands







# Analysis-by-Synthesis




Larry Roberts  
“Father of Computer Vision”



Input image




$2 \times 2$  gradient operator




computed 3D model  
rendered from new viewpoint

- History goes way back to the **first** Computer Vision paper!  
Roberts: Machine Perception of Three-Dimensional Solids, MIT, 1963

# Power of Analysis-by-Synthesis



- Space Carving: A MVS method that used Colored voxels
- But the optimization method was bottom up then.
- Key is optimization via Analysis-by-Synthesis [Plenoxels, Yu et al. 2022]



Input Image (1 of 45)



Reconstruction



Reconstruction



Reconstruction



Input Image  
(1 of 100)




Views of Reconstruction



Koltun and Seitz, A Theory of Shape by Space Carving IJCV 2000

# Analysis-by-Synthesis



Blanz & Vetter 1999

- With custom differentiable renders

# **Analysis by Synthesis Requires Differentiable Renderers**

Next: Deep dive into Volumetric Rendering Function

# Where we are

1. Birds Eye View & Background
- 2. Volumetric Rendering Function**
3. Encoding and Representing 3D Volumes
4. Signal Processing Considerations
5. Challenges & Pointers

# Volume Rendering

*"... in 10 years, all rendering will be volume rendering."*

Jim Kajiya at SIGGRAPH '91



# Neural Volumetric Rendering

# Neural Volumetric **Rendering**

computing color along rays  
through 3D space




*What color is this pixel?*

# Cameras and rays

# Cameras and rays

- We need the mathematical mapping from  $(camera, pixel) \rightarrow ray$
- Then can abstract underlying problem as learning the function  $ray \rightarrow color$  (the “plenoptic function”)



## Recap Coordinate frames: World-to-Camera Transforms




Figure credit: Peter Hedman

## Recap Coordinate frames: Camera-to-World Transforms







Figure credit: Peter Hedman

# Camera pose - pixel to camera


- Mapping from *(camera, pixel)* to ray in camera coordinate frame
- This coordinate system has camera situated at origin, with right/up/backwards aligned to x/y/z axes
  - Axis convention varies in different codebases :(
  - “Inverse intrinsic matrix” in a computer vision sense



# Camera pose - pixel to camera


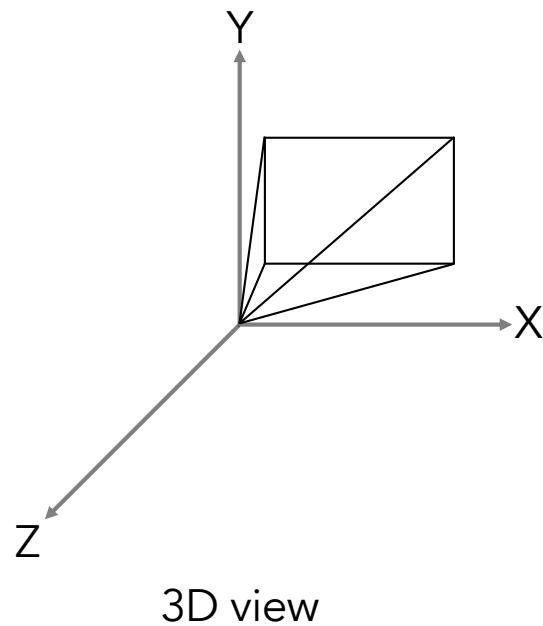


Top view  
(looking along Y)




Side view  
(looking along X)


# Camera pose - pixel to camera




# Camera pose - pixel to camera



Coordinates are  $(i, j)$  in pixel space




Top view  
(looking along Y)




Side view  
(looking along X)


# Camera pose - pixel to camera



Recenter using pixel coordinates of image center




Top view  
(looking along Y)





Side view  
(looking along X)

# Camera pose - pixel to camera




Rescale frustum by focal length  $f$  so that image plane is at distance 1



# Camera pose - pixel to camera

Full mapping is  $(i, j) \rightarrow \left( \frac{i-w/2}{f}, \frac{j-h/2}{f}, -1 \right)$  to get 3D coordinates for a point on the image plane.

Camera space ray points from origin toward this point.




# Camera pose - pixel to camera

- Omitted details
  - Half-pixel offset — add 0.5 to  $i$  and  $j$  so ray precisely hits pixel center
  - This is a *perfect* pinhole model — typically need to add a distortion model to correct for error found in real cameras


# Camera pose - camera to world

- Simply apply rigid rotation and translation to origin and image plane points (six degrees of freedom).
- This positions the camera in “world space”.



# Calculating points along a ray

In the world coordinate frame:




# **Neural Volumetric Rendering**

# Neural Volumetric Rendering

continuous, differentiable  
rendering model without  
concrete ray/surface intersections




# Surface vs. volume rendering




Want to know how ray interacts with scene

# Surface vs. volume rendering




Surface rendering — loop over geometry, check for ray hits

# Surface vs. volume rendering



Volume rendering — loop over ray points, query geometry

# History of volume rendering



# Early computer graphics

- ▶ Theory of volume rendering co-opted from physics in the 1980s: absorption, emission, out-scattering/in-scattering



Ray tracing simulated cumulus cloud [Kajiya]

Chandrasekhar 1950, Radiative Transfer  
Kajiya 1984, Ray Tracing Volume Densities

# Alpha compositing


- ▶ Theory of volume rendering co-opted from physics in the 1980s: absorption, emission, out-scattering/in-scattering
- ▶ Alpha rendering developed for digital compositing in VFX movie production



Pt. Reyes = Foreground over Hillside over Background.  
Alpha compositing [Porter and Duff]

Porter and Duff 1984, Compositing Digital Images

# Volume rendering for visualization






Medical data visualisation [Levoy]

- ▶ Theory of volume rendering co-opted from physics in the 1980s: absorption, emission, out-scattering/in-scattering
- ▶ Alpha rendering developed for digital compositing in VFX movie production
- ▶ Volume rendering applied to visualise 3D medical scan data in 1990s

Chandrasekhar 1950, *Radiative Transfer*  
Kajiya 1984, *Ray Tracing Volume Densities*  
Porter and Duff 1984, *Compositing Digital Images*  
Levoy 1988, *Display of Surfaces from Volume Data*  
Max 1995, *Optical Models for Direct Volume Rendering*

# Volume rendering for surfaces



Geometry and materials can be stored per-voxel and used with standard surface rendering methods

- Sparse voxel octrees
- Voxel hashing
- Anisotropic radiative transfer


# Volume rendering derivations




Absorption



<http://commons.wikimedia.org>



Scattering



Emission



<http://wikipedia.org>

# Simplify

Absorption



Scattering




Emission



<http://commons.wikimedia.org>


<http://wikipedia.org>

# Volumetric formulation for NeRF




Scene is a cloud of tiny colored particles

# Volumetric formulation for NeRF




## What does it mean for a ray to “hit” the volume?



This notion is *probabilistic*: chance that ray hits a particle in a small interval around  $t$  is  $\sigma(t)dt$ .  $\sigma$  is called the “volume density”


# Probabilistic interpretation



To determine if  $t$  is the *first* hit along the ray, need to know  $T(t)$ : the probability that the ray makes it through the volume up to  $t$ .

$T(t)$  is called “transmittance”


# Probabilistic interpretation




The product of these probabilities tells us how much you see the particles at  $t$ :

$$P[\text{first hit at } t] = P[\text{no hit before } t] \times P[\text{hit at } t] = T(t)\sigma(t)dt$$

# Calculating $T$ given $\sigma$



# Calculating $T$ given $\sigma$



# Calculating transmittance $T$

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

$$\begin{array}{ccc} T(t + dt) & = & T(t) \\ \nearrow & & \nearrow \\ & & (1 - \sigma(t)dt) \end{array}$$

# Solve for $T$

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

# Solve for $T$

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Taylor expansion for  $T \Rightarrow T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$

# Solve for $T$

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Taylor expansion for  $T \Rightarrow T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$

Rearrange  $\Rightarrow \frac{T'(t)}{T(t)} dt = -\sigma(t)dt$

# Solve for $T$

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Taylor expansion for  $T \Rightarrow T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$

Rearrange  $\Rightarrow \frac{T'(t)}{T(t)} dt = -\sigma(t)dt$

Integrate  $\Rightarrow \log T(t) = - \int_{t_0}^t \sigma(s)ds$

# Solve for $T$

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Taylor expansion for  $T \Rightarrow T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$


Rearrange  $\Rightarrow \frac{T'(t)}{T(t)} dt = -\sigma(t)dt$

Integrate  $\Rightarrow \log T(t) = -\int_{t_0}^t \sigma(s)ds$

Exponentiate  $\Rightarrow T(t) = \exp\left(-\int_{t_0}^t \sigma(s)ds\right)$

# PDF for ray termination

$$P[\text{no hits before } t] = T(t)$$



Finally, we can write the probability that a ray terminates at  $t$  as a function of only sigma

$$P[\text{first hit at } t] = P[\text{no hit before } t] \times P[\text{hit at } t]$$

$$= T(t)\sigma(t)dt$$

$$= \exp\left(-\int_{t_0}^t \sigma(s)ds\right)\sigma(t)dt$$


# Expected value of color along ray

This means the expected color returned by the ray will be

$$\int_{t_0}^{t_1} T(t) \sigma(t) \mathbf{c}(t) dt$$


Note the nested integral!

# Approximating the nested integral




We use quadrature to approximate the nested integral,

# Approximating the nested integral




We use quadrature to approximate the nested integral,  
splitting the ray up into  $n$  segments with endpoints  
 $\{t_1, t_2, \dots, t_{n+1}\}$

# Approximating the nested integral



We use quadrature to approximate the nested integral,  
splitting the ray up into  $n$  segments with endpoints  
 $\{t_1, t_2, \dots, t_{n+1}\}$   
with lengths  $\delta_i = t_{i+1} - t_i$

# Approximating the nested integral



We assume volume density and color are roughly constant within each interval

# Deriving quadrature estimate

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx$$

This allows us to break the outer integral

# Deriving quadrature estimate

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

This allows us to break the outer integral into a sum of analytically tractable integrals

# Deriving quadrature estimate

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

Caveat: piecewise constant density and color  
**do not** imply constant transmittance!

# Deriving quadrature estimate

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$


Caveat: piecewise constant density and color  
**do not** imply constant transmittance!

Important to account for how early part of a  
segment blocks later part when  $\sigma_i$  is high

# Evaluating $T$ for piecewise constant density

$$\text{For } t \in [t_i, t_{i+1}], T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$$


We need to evaluate at continuous  $t$  values  
that can lie *partway through* an interval



# Evaluating $T$ for piecewise constant density

For  $t \in [t_i, t_{i+1}]$ ,  $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$


$$\exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right) = T_i \quad \text{"How much light is blocked by all previous segments?"}$$



# Evaluating $T$ for piecewise constant density

$$\text{For } t \in [t_i, t_{i+1}], T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$$

"How much light is blocked partway through the current segment?"  $\exp(-\sigma_i(t - t_i))$



# Deriving quadrature estimate

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

# Deriving quadrature estimate

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

Substitute  $= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt$

# Deriving quadrature estimate

$$\begin{aligned}\int T(t)\sigma(t)\mathbf{c}(t)dt &\approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt \\ &= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt \\ \text{Integrate} &= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i}\end{aligned}$$

# Deriving quadrature estimate

$$\begin{aligned}\int T(t)\sigma(t)\mathbf{c}(t)dt &\approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt \\ &= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt \\ &= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i}\end{aligned}$$

Cancel  $\sigma_i$  =  $\sum_{i=1}^n T_i \mathbf{c}_i (1 - \exp(-\sigma_i \delta_i))$

# Connection to alpha compositing

$$= \sum_{i=1}^n T_i \mathbf{c}_i \underbrace{(1 - \exp(-\sigma_i \delta_i))}_{\text{segment opacity } \alpha_i}$$

# Connection to alpha compositing

$$= \sum_{i=1}^n T_i \mathbf{c}_i \underbrace{(1 - \exp(-\sigma_i \delta_i))}_{\text{segment opacity } \alpha_i}$$

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

$$color = \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

# Summary: volume rendering integral estimate

Rendering model for ray  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ :

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$


↑  
weights      colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment  $i$ :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$



# Volume rendering is trivially differentiable

Rendering model for ray  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ :

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

weights


colors

**differentiable w.r.t.  $\mathbf{c}, \sigma$**

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$


How much light is contributed by ray segment  $i$ :




$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$

# Further points on volume rendering


# Alpha mattes and compositing



# Alpha mattes and compositing



# Alpha mattes and compositing



# Alpha mattes and compositing



Mildenhall\*, Srinivasan\*, Tancik\* et al 2020, NeRF

Poole et al 2022, DreamFusion

Tang et al 2022, Compressible-composable NeRF via Rank-residual Decomposition

# Rendering weight PDF is important


Remember, expected color is equal to

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_i T_i \alpha_i \mathbf{c}_i$$

$T(t)\sigma(t)$  and  $T_i \alpha_i$  are “rendering weights” — probability distribution along the ray  
(continuous and discrete, respectively)

## Visual intuition – rendering weights not just 3D function


$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$



Rendering weights are not a 3D function — depends on ray, because of transmittance!

## Visual intuition – rendering weights not just 3D function

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$



Rendering weights are not a 3D function — depends on ray, because of transmittance!

# Rendering weight PDF is important — depth


We can use this distribution to compute expectations for other quantities, e.g. “expected depth”:

$$\bar{t} = \sum_i T_i \alpha_i t_i$$


This is often how people visualise NeRF depth maps.

Alternatively, other statistics like mode or median can be used.

# Rendering weight PDF is important — depth




Mean depth



Median depth

# Rendering weight PDF is important — depth



# Volume rendering other quantities

This idea can be used for any quantity we want to “volume render” into a 2D image.  
If  $\mathbf{v}$  lives in 3D space (semantic features, normal vectors, etc.)

$$\sum_i T_i \alpha_i \mathbf{v}_i$$


can be taken per-ray to produce 2D output images.

# Volume rendering other quantities



Various recent works have used this idea to render higher-level semantic feature maps (e.g., *Feature Field Distillation* and *Neural Feature Fusion Fields*).


# Density as geometry



Normal vectors (from analytic gradient of density)

# Applications/optimizing differentiable volume rendering

# Alpha compositing model in ML/computer vision



*Differentiable ray consistency* work used a forward model with “probabilistic occupancy” to supervise 3D-from-single-image prediction. Same rendering model as alpha compositing!

$$p(z_r = i) = \begin{cases} (1 - x_i^r) \prod_{j=1}^{i-1} x_j^r, & \text{if } i \leq N_r \\ \prod_{j=1}^{N_r} x_j^r, & \text{if } i = N_r + 1 \end{cases}$$

# Volume rendering for view synthesis

## Multiplane image methods

Stereo Magnification (Zhou et al. 2018)


Pushing the Boundaries... (Srinivasan et al. 2019)

Local Light Field Fusion (Mildenhall et al. 2019)

DeepView (Flynn et al. 2019)

Single-View... (Tucker & Snavely 2020)

Typical deep learning pipelines - images go into a 3D CNN, big RGBA 3D volume comes out



## Neural Volumes

(Lombardi et al. 2019)

Direct gradient descent to optimize an RGBA volume, regularized by a 3D CNN



