

# **Neural Radiance Fields 2**

CS194-26/294-26: Intro to Computer Vision and Computational  
Photography

Angjoo Kanazawa  
UC Berkeley Fall 2022

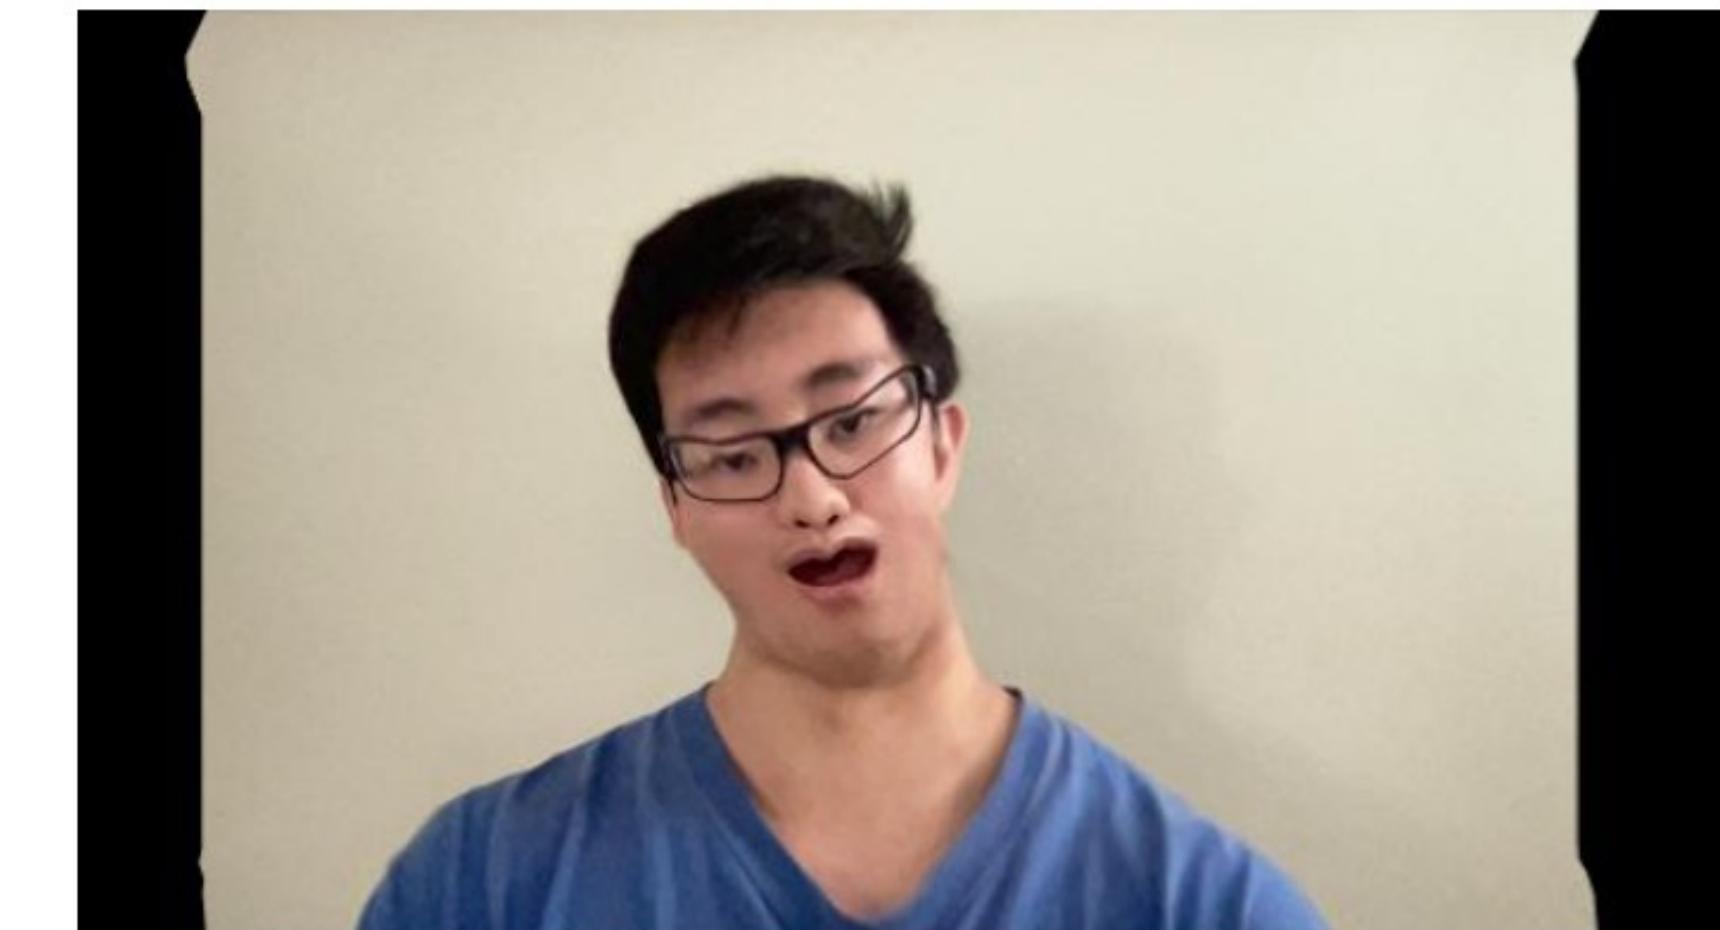
**Lots of content from ECCV 2022 Tutorial on Neural  
Volumetric Rendering for Computer Vision**

# Project 3 Winner!!

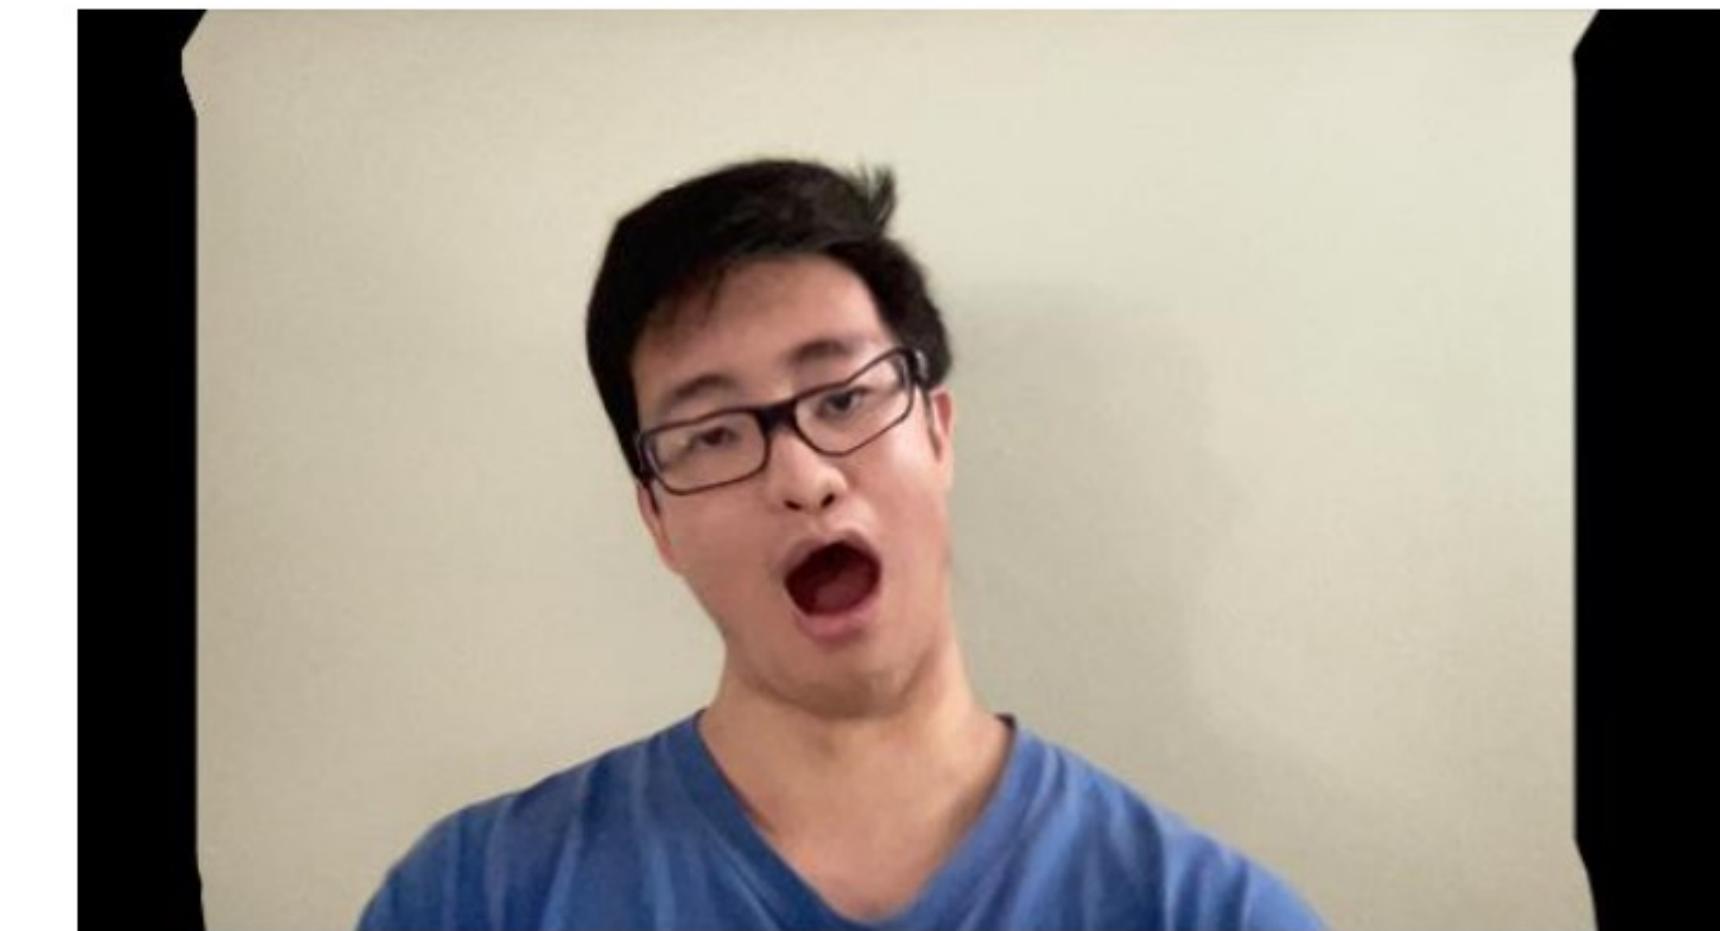
Original keyframe



My face morphed into keyframe shape



**Joshua Chen**



# Project 4 Highlight

Shinwoo Choi



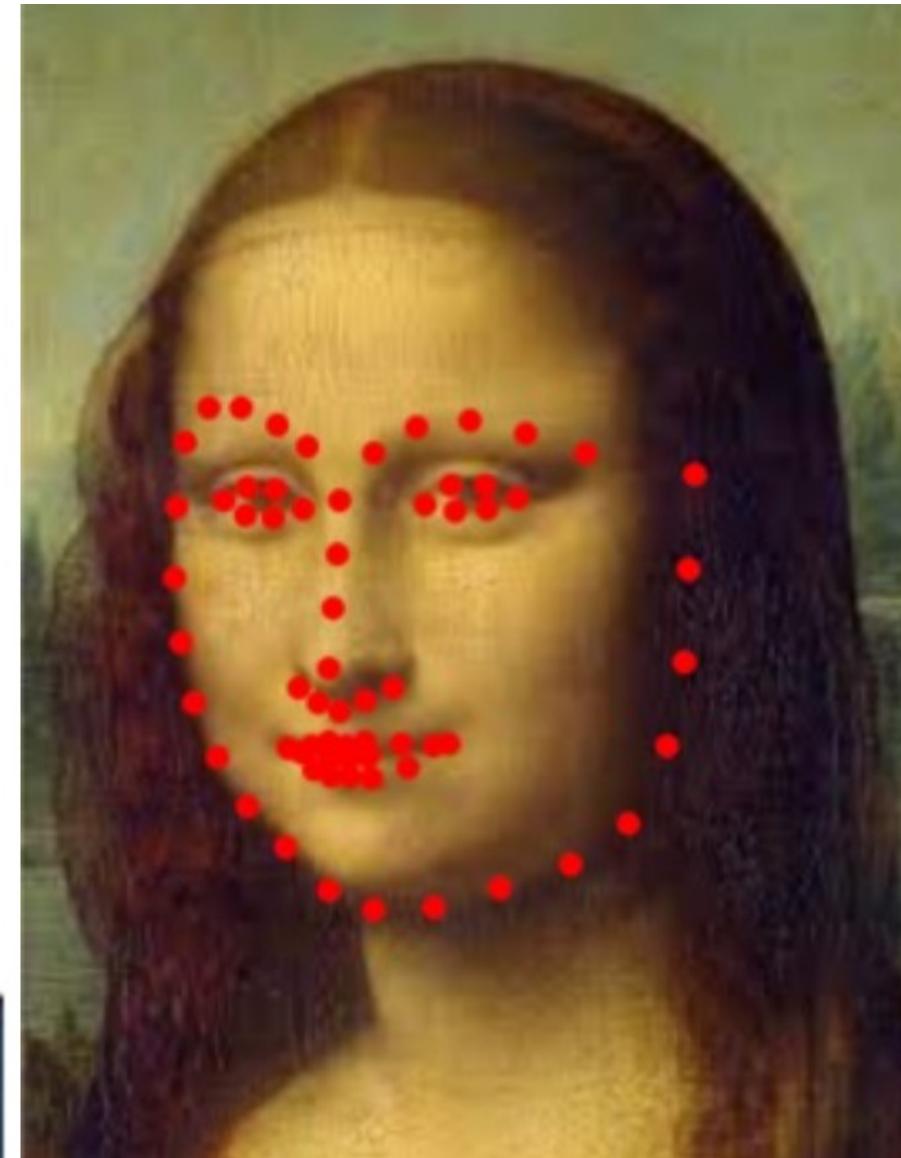
<https://inst.eecs.berkeley.edu/~cs194-26/fa22/upload/files/proj4B/cs194-26-afm/>

Class vote is happening now, do vote!

# Project 5 Winner!!

## Jules Dedieu

#	△	Team	Members	Score	Entries	Last	Code
1	—	Jules		5.06767	7	11d	
2	—	tna		5.47330	6	15d	
3	—	Val R		5.59510	7	14d	



# Where we are

1. Birds Eye View & Background
2. **Volumetric Rendering Function**
3. Encoding and Representing 3D Volumes
4. Signal Processing Considerations
5. Challenges & Pointers

# Simplify

Absorption



<http://commons.wikimedia.org>

Scattering



Emission



<http://wikipedia.org>

# Summary: volume rendering integral estimate

Rendering model for ray  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ :

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

differentiable w.r.t.  $\mathbf{c}, \sigma$

colors

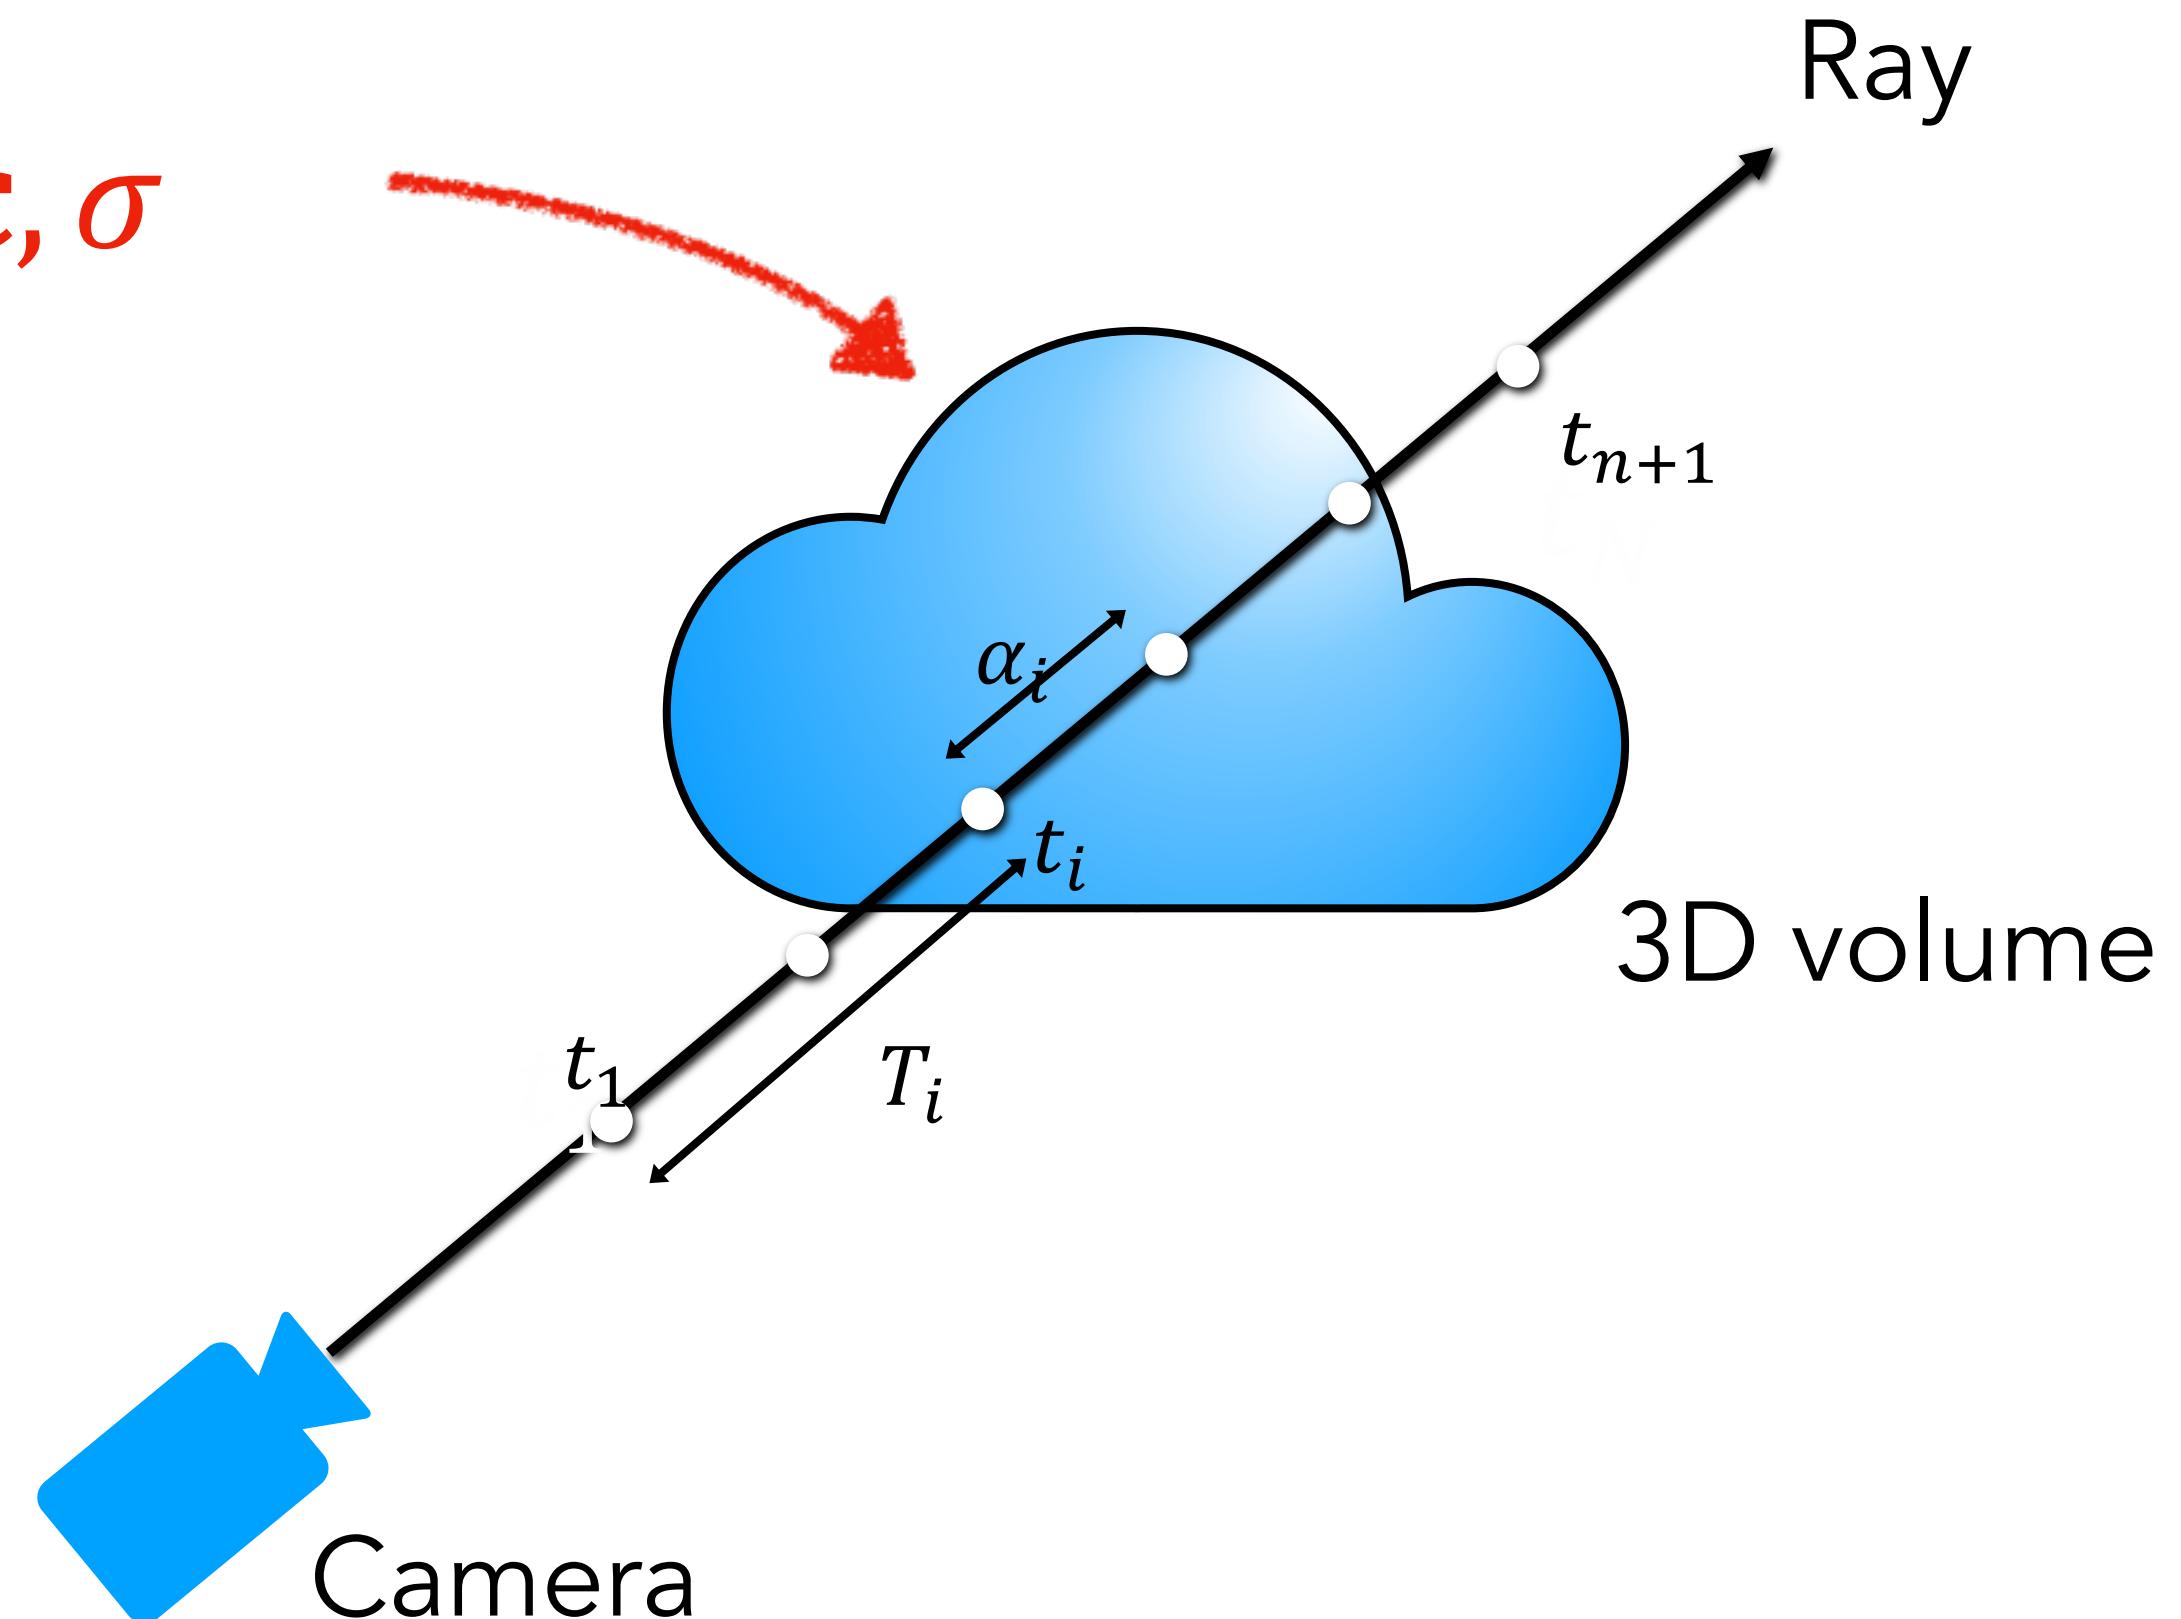
weights

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

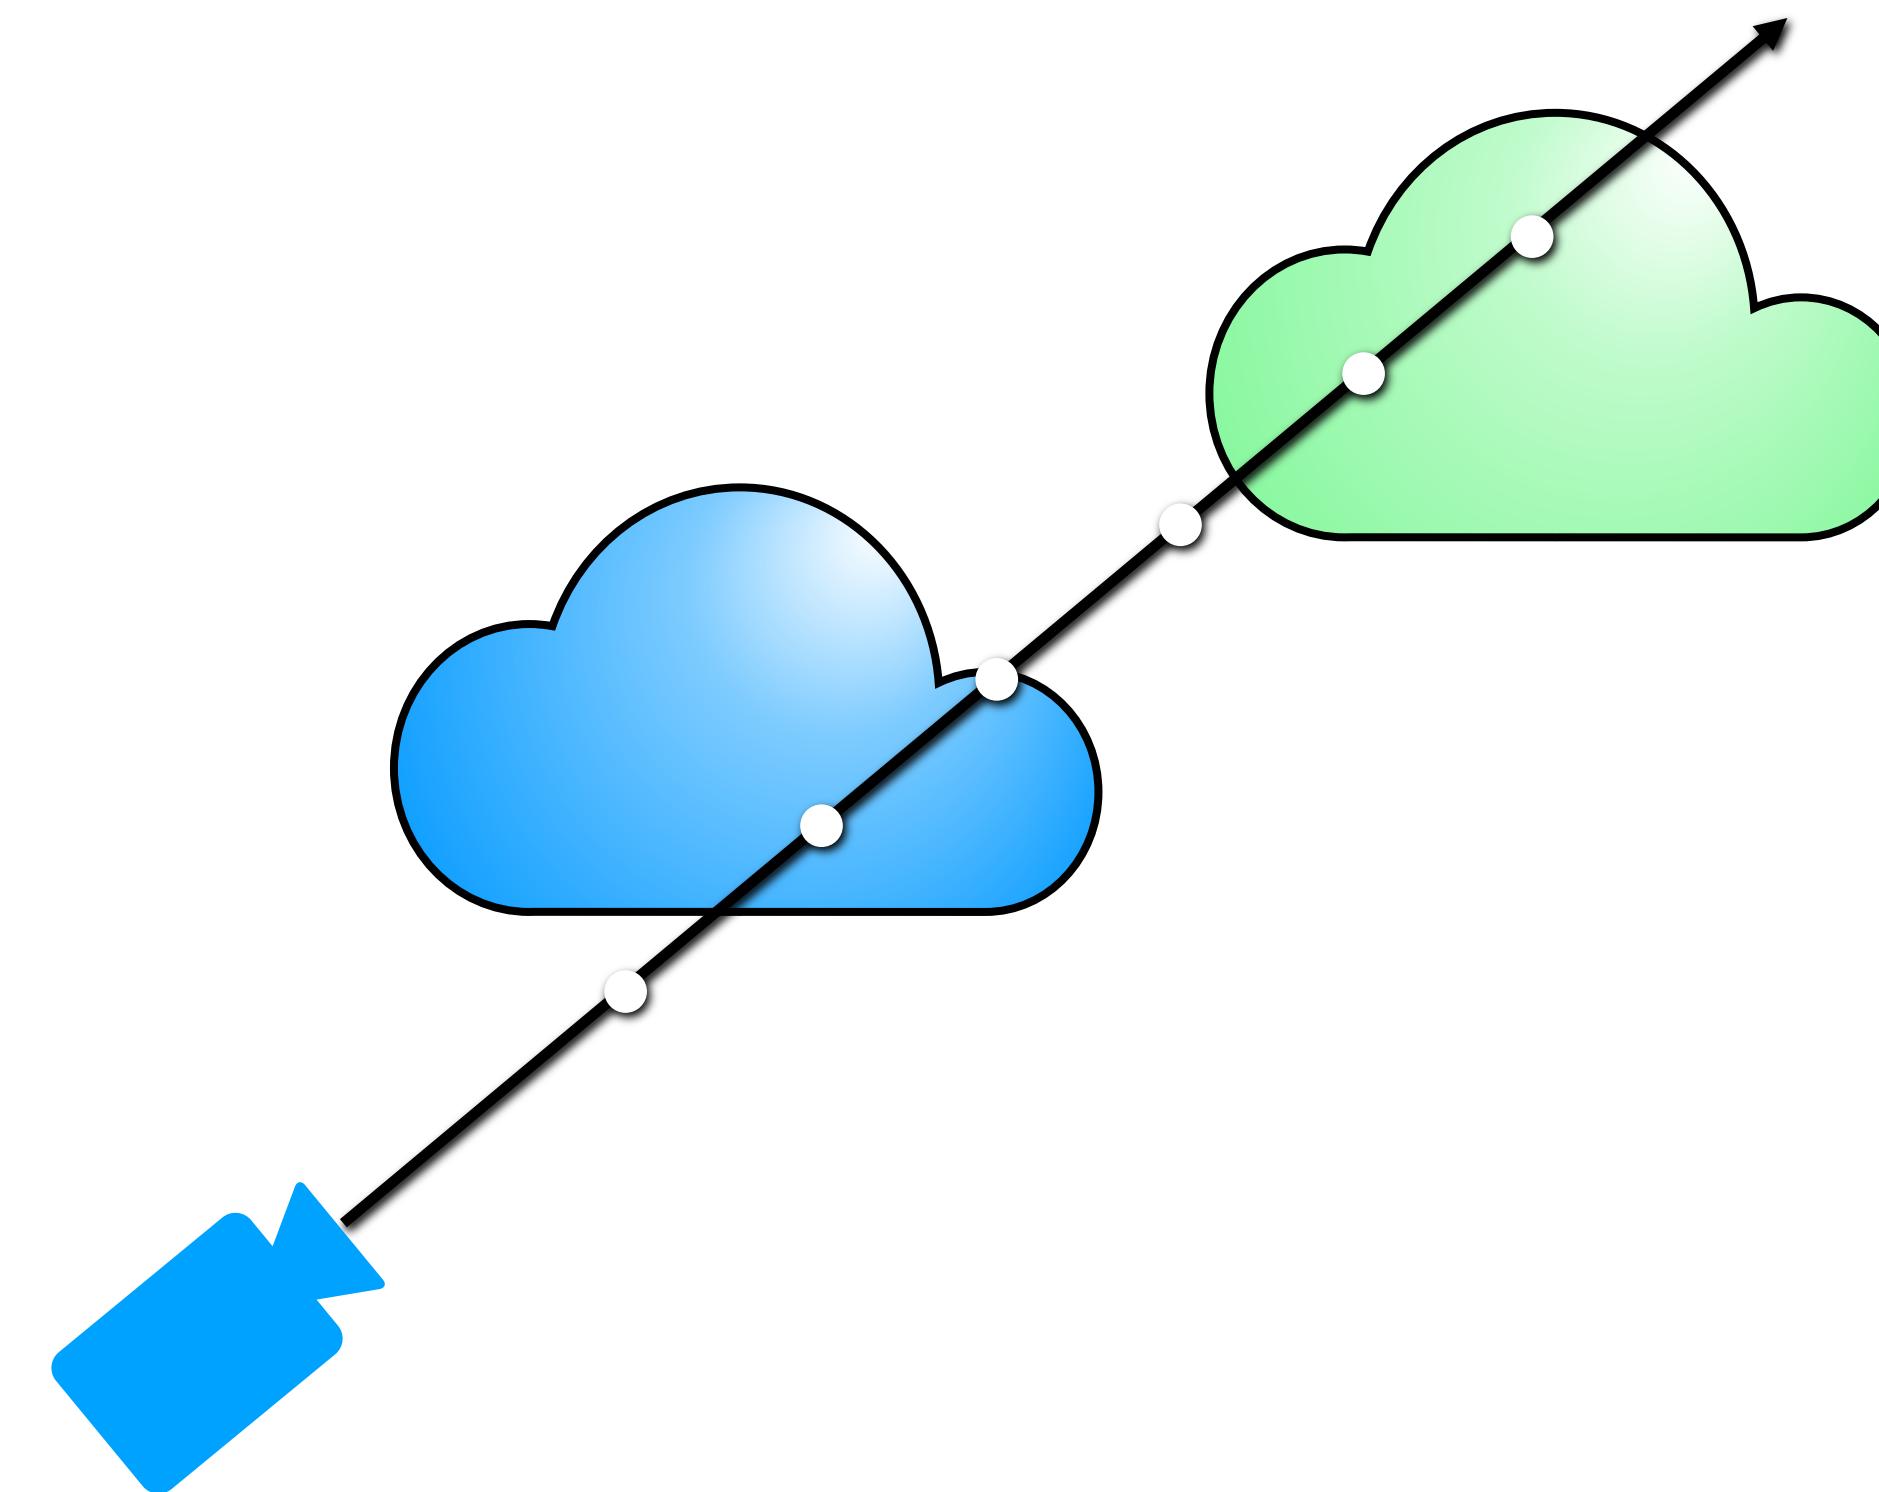
How much light is contributed by ray segment  $i$ :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$

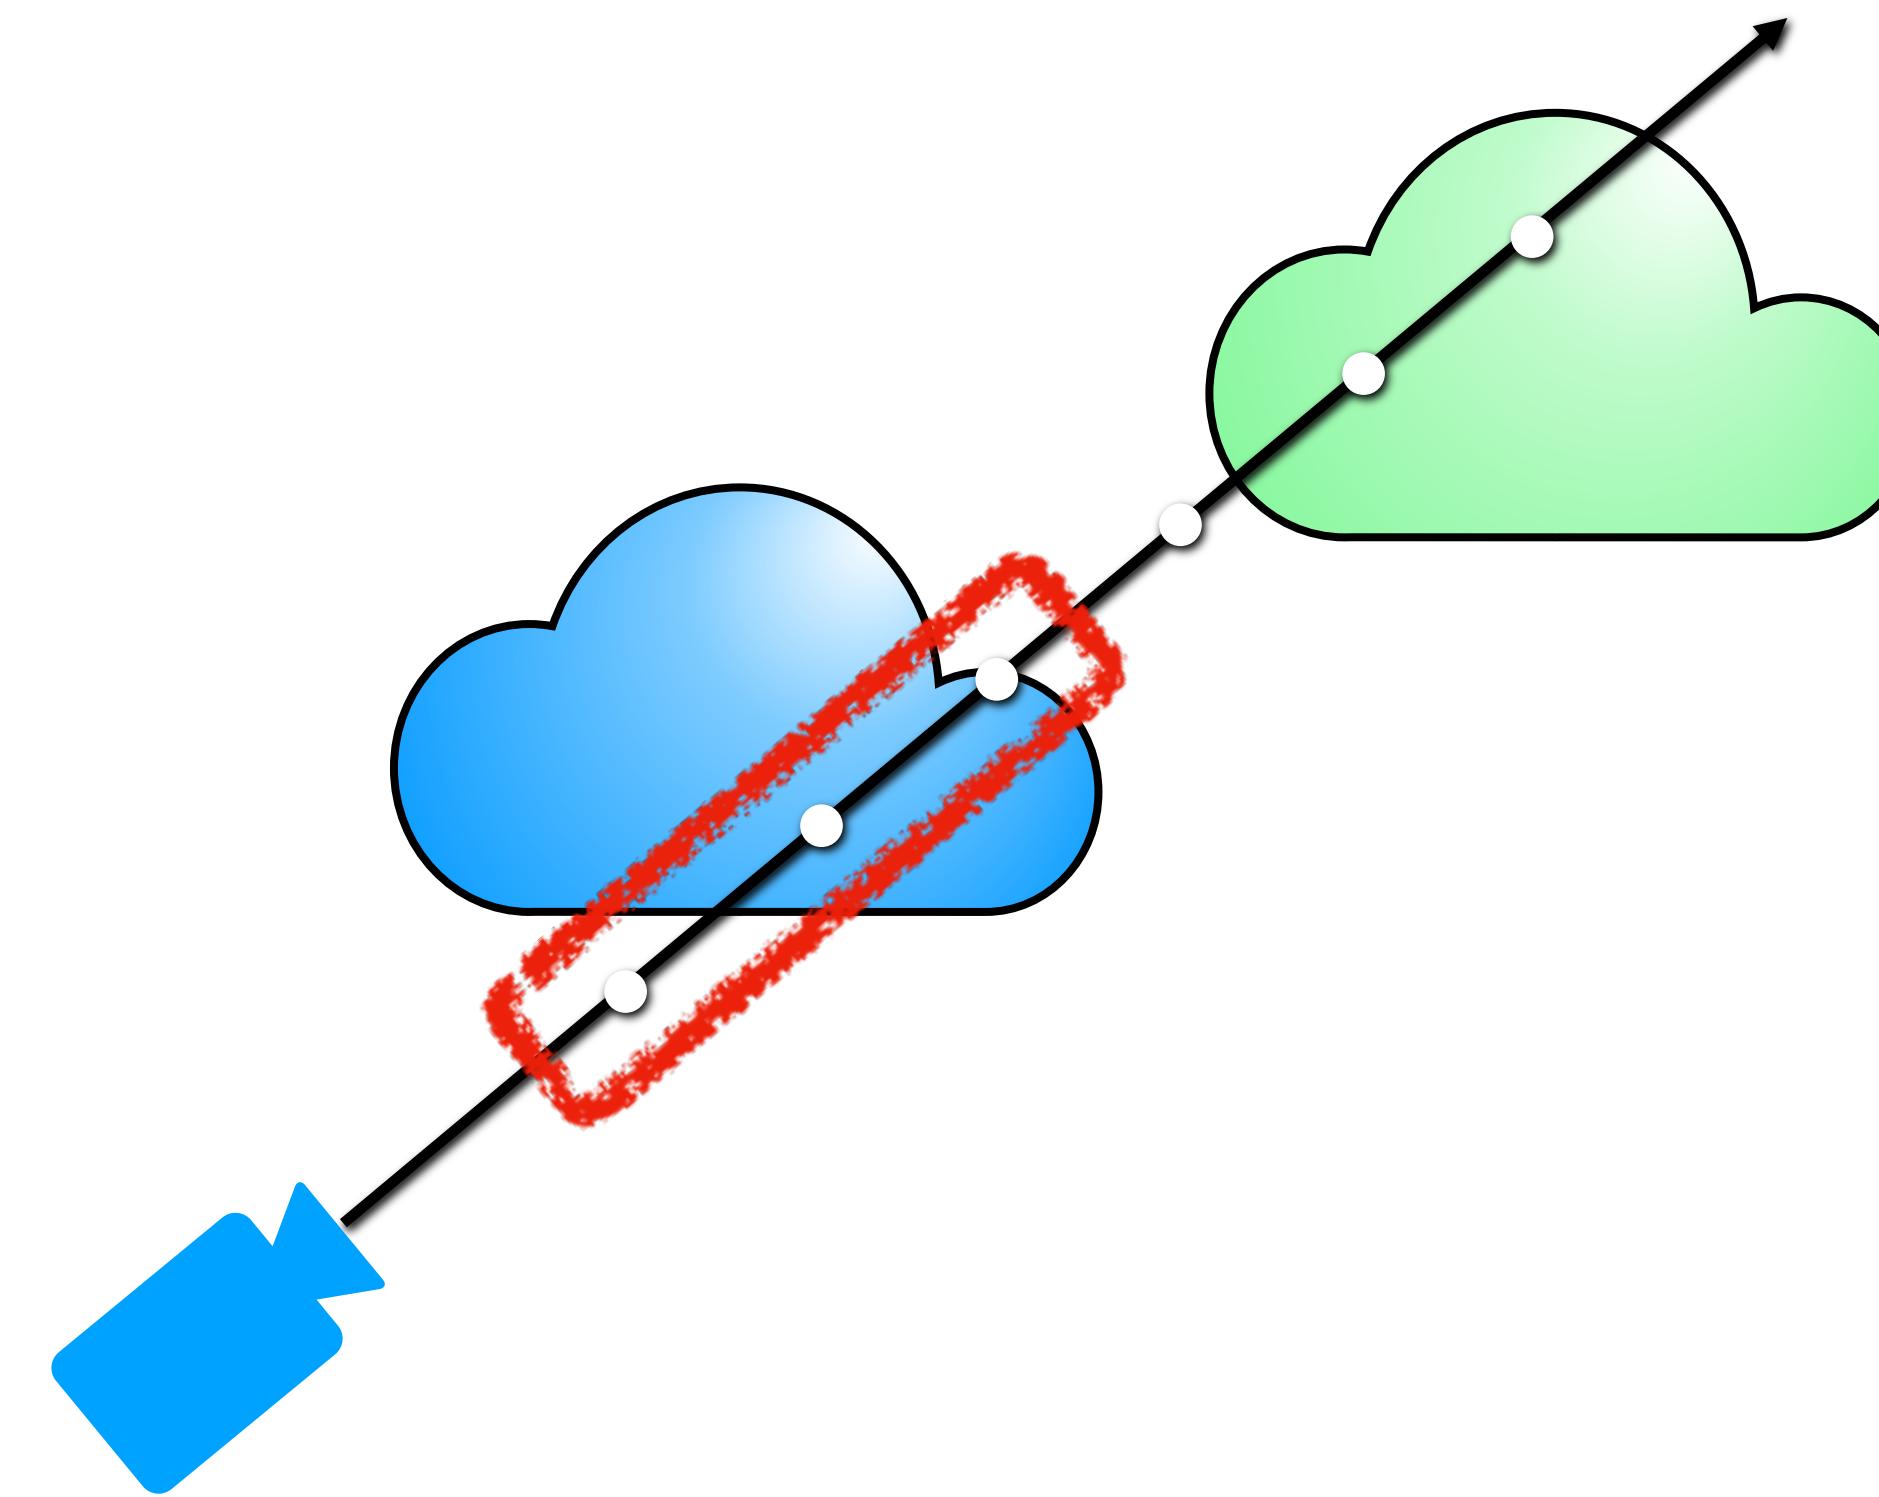


# **Further points on volume rendering**

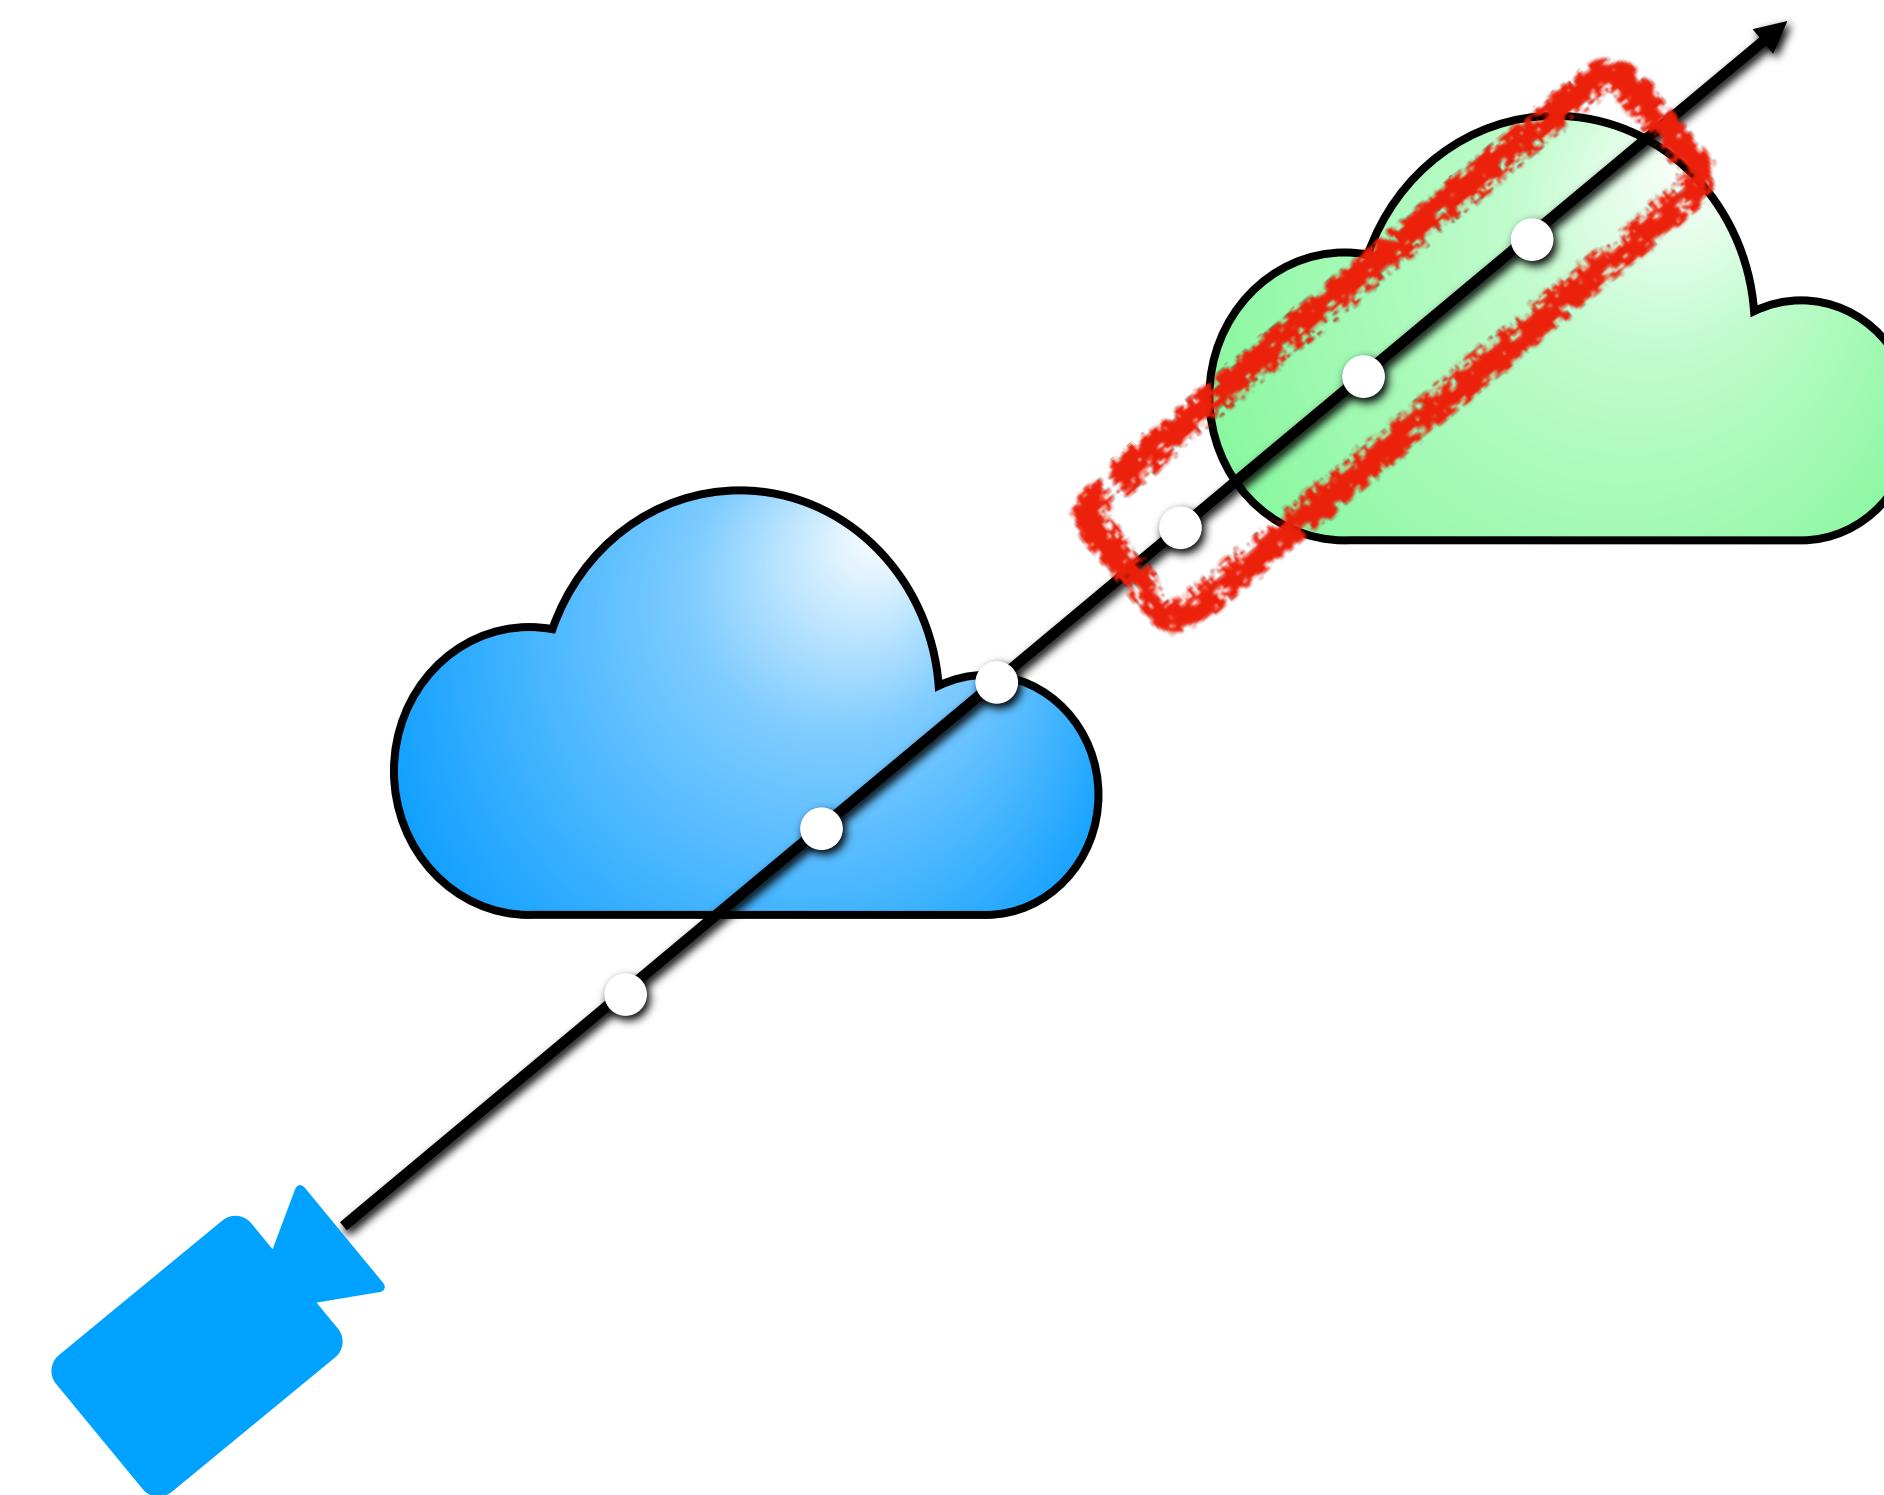
# Alpha mattes and compositing



# Alpha mattes and compositing



# Alpha mattes and compositing



# Alpha mattes and compositing



Mildenhall\*, Srinivasan\*, Tancik\* et al 2020, NeRF

Poole et al 2022, DreamFusion

Tang et al 2022, Compressible-composable NeRF via Rank-residual Decomposition

# Rendering weight PDF is important

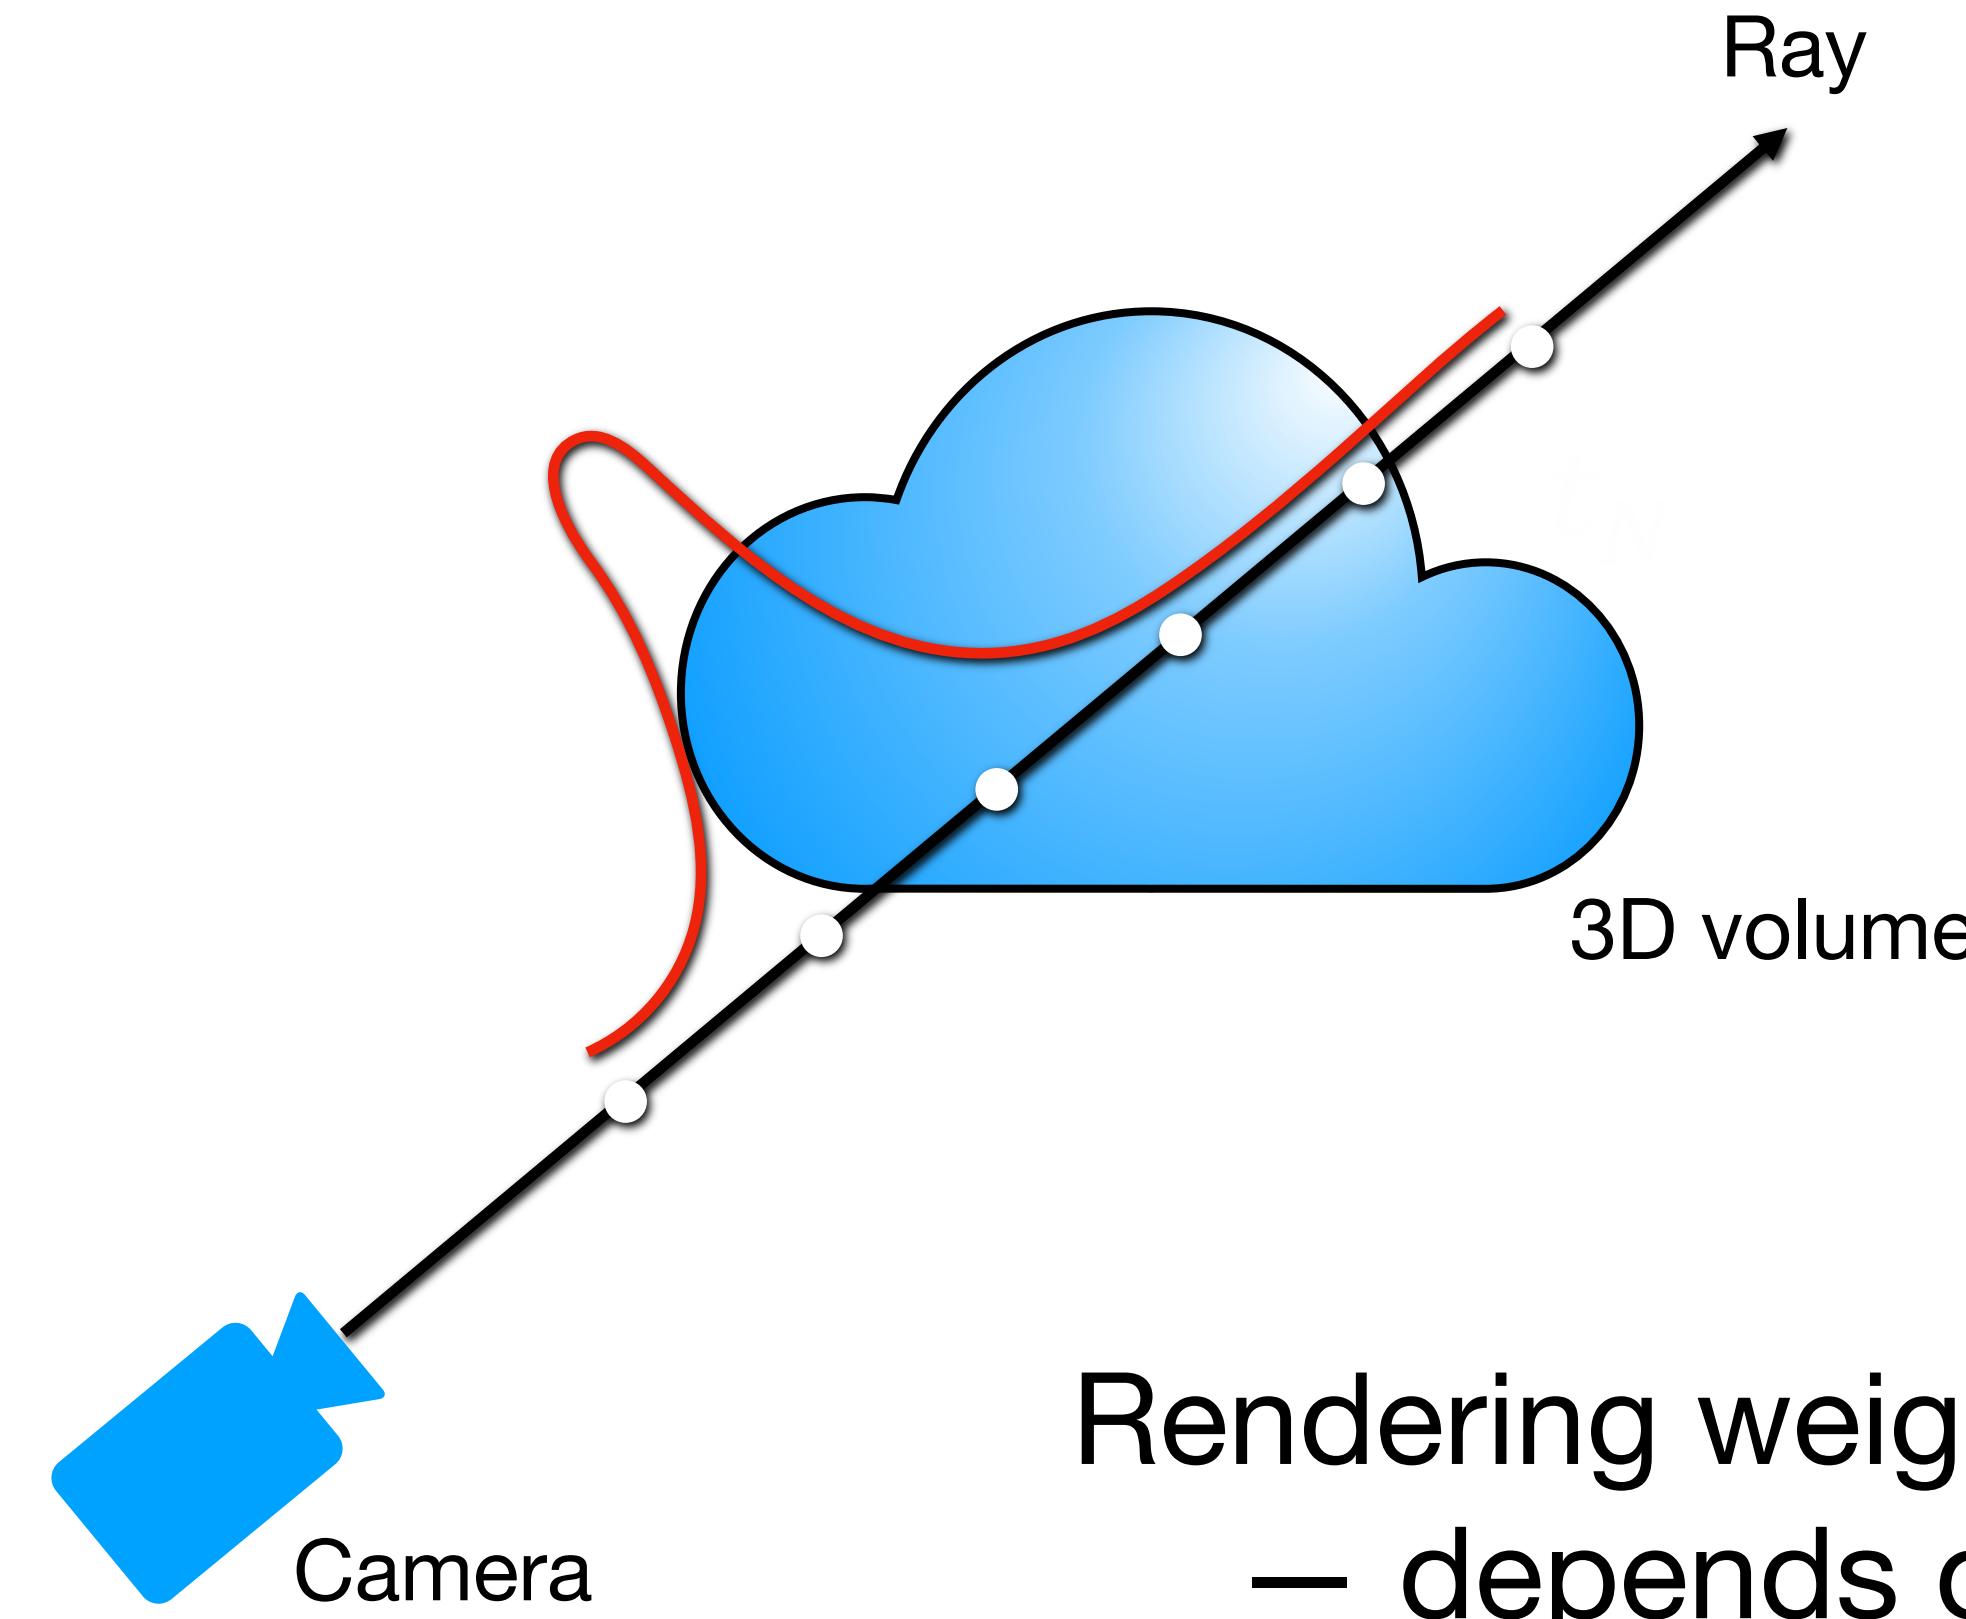
Remember, expected color is equal to

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_i T_i \underbrace{\alpha_i}_{w_i} \mathbf{c}_i$$

$T(t)\sigma(t)$  and  $T_i\alpha_i$  are “rendering weights”— probability distribution along the ray  
(continuous and discrete, respectively)

# Visual intuition – rendering weights depend on the ray!

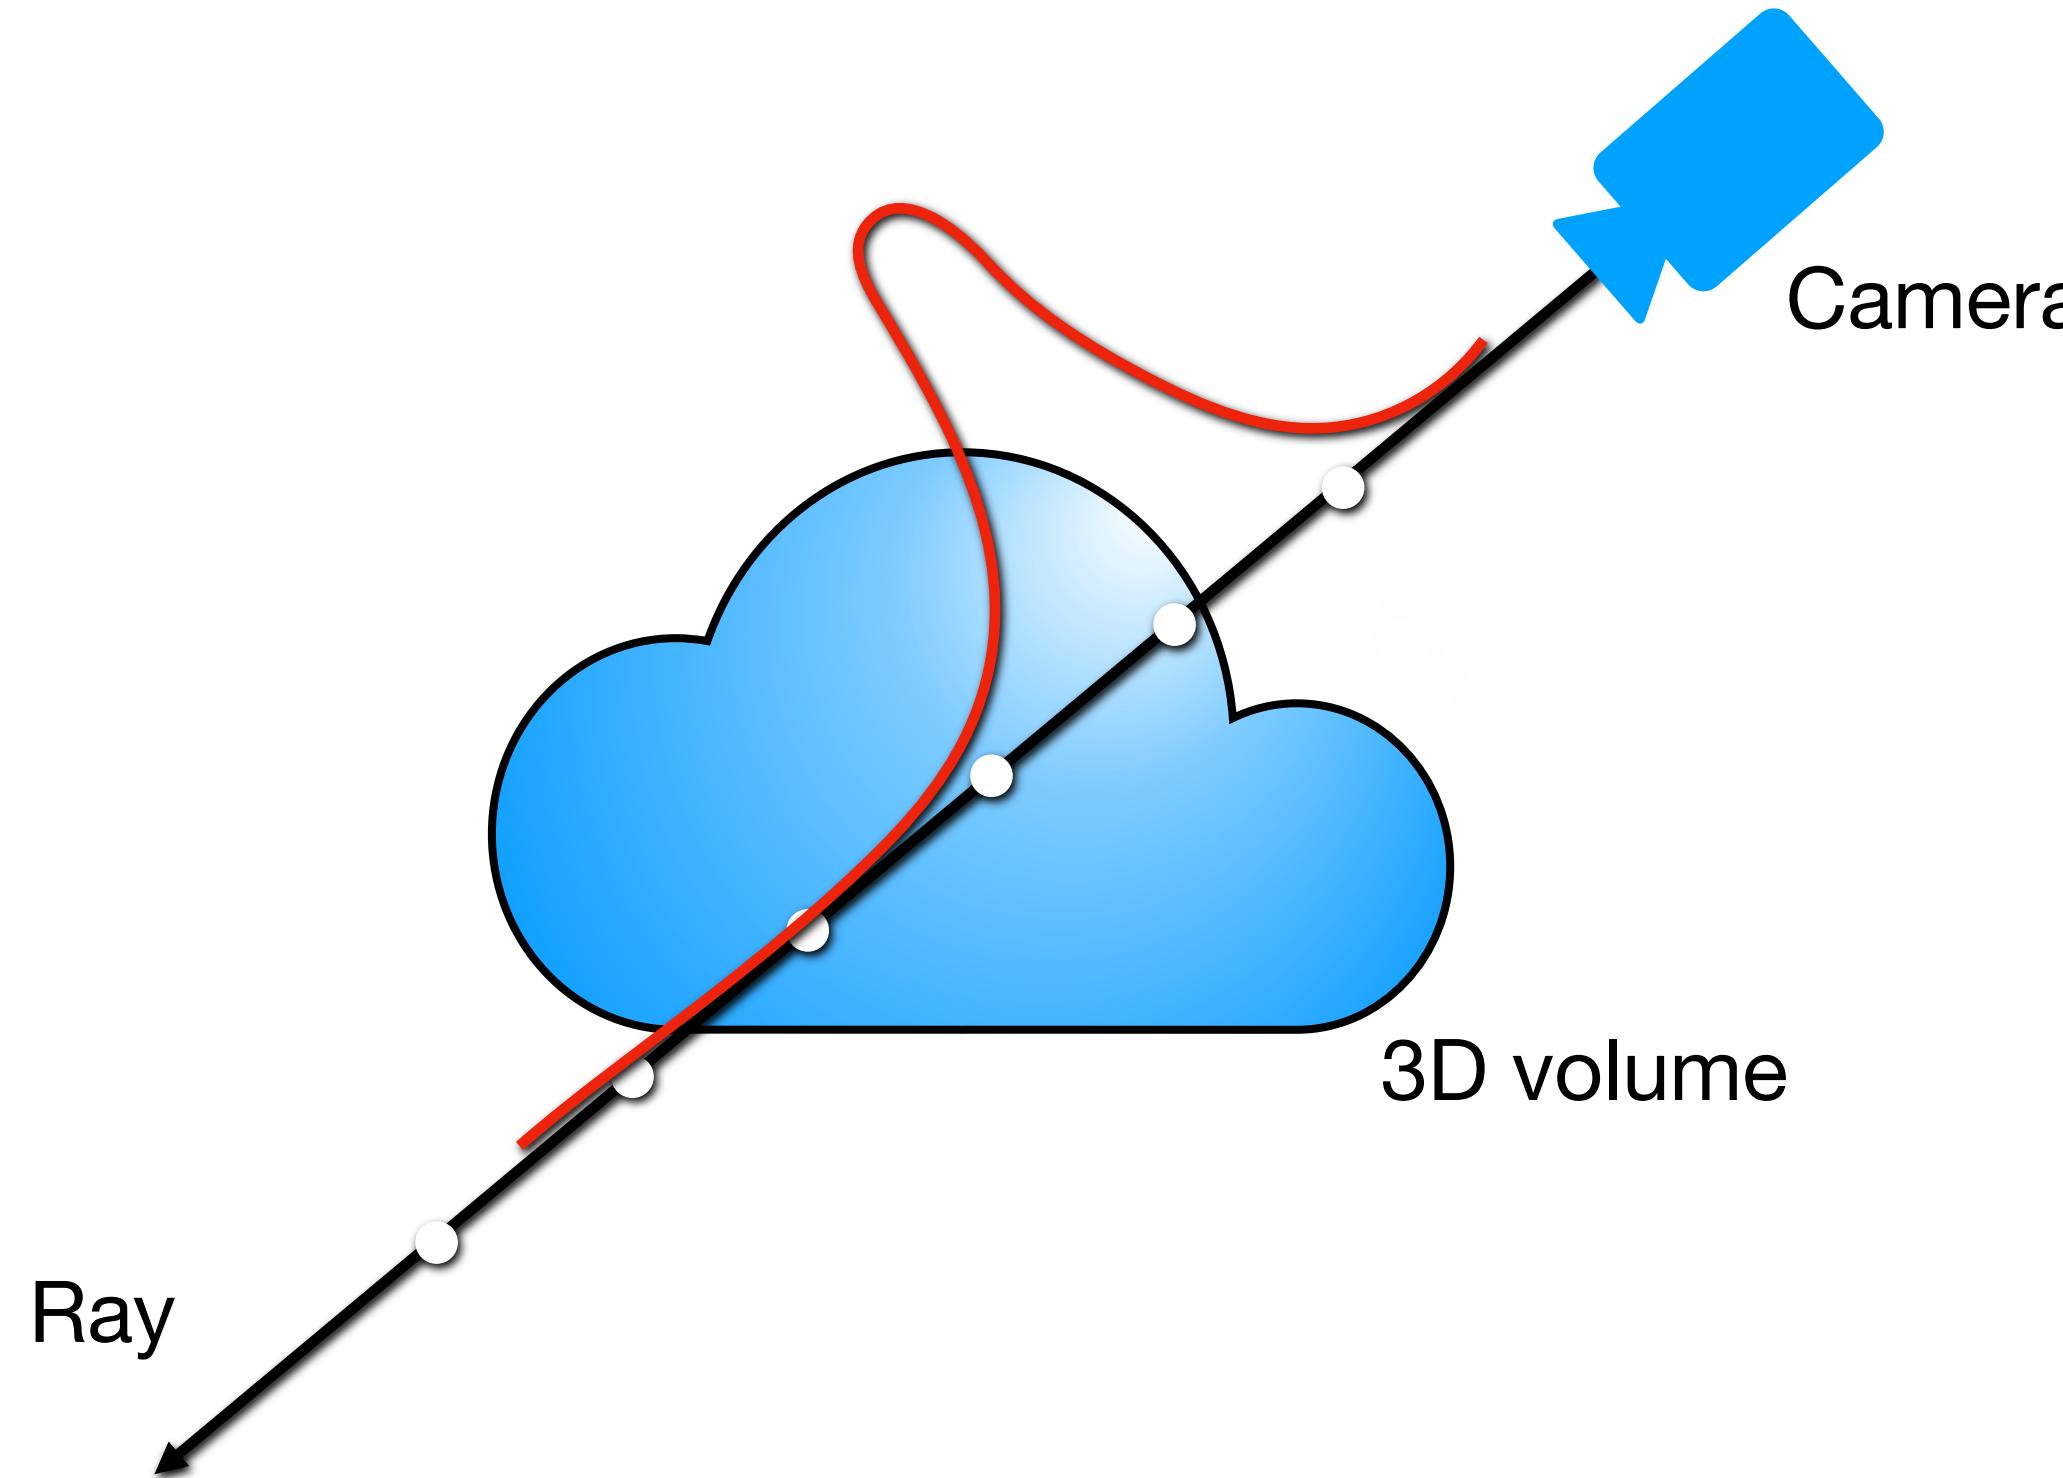
$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$



Rendering weights are not a 3D function  
– depends on the ray, because of  
transmittance!

# Visual intuition – rendering weights depend on the ray!

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$



Rendering weights are not a 3D function  
– depends on the ray, because of  
transmittance!

# Use the rendering weight (PDF) to render depth

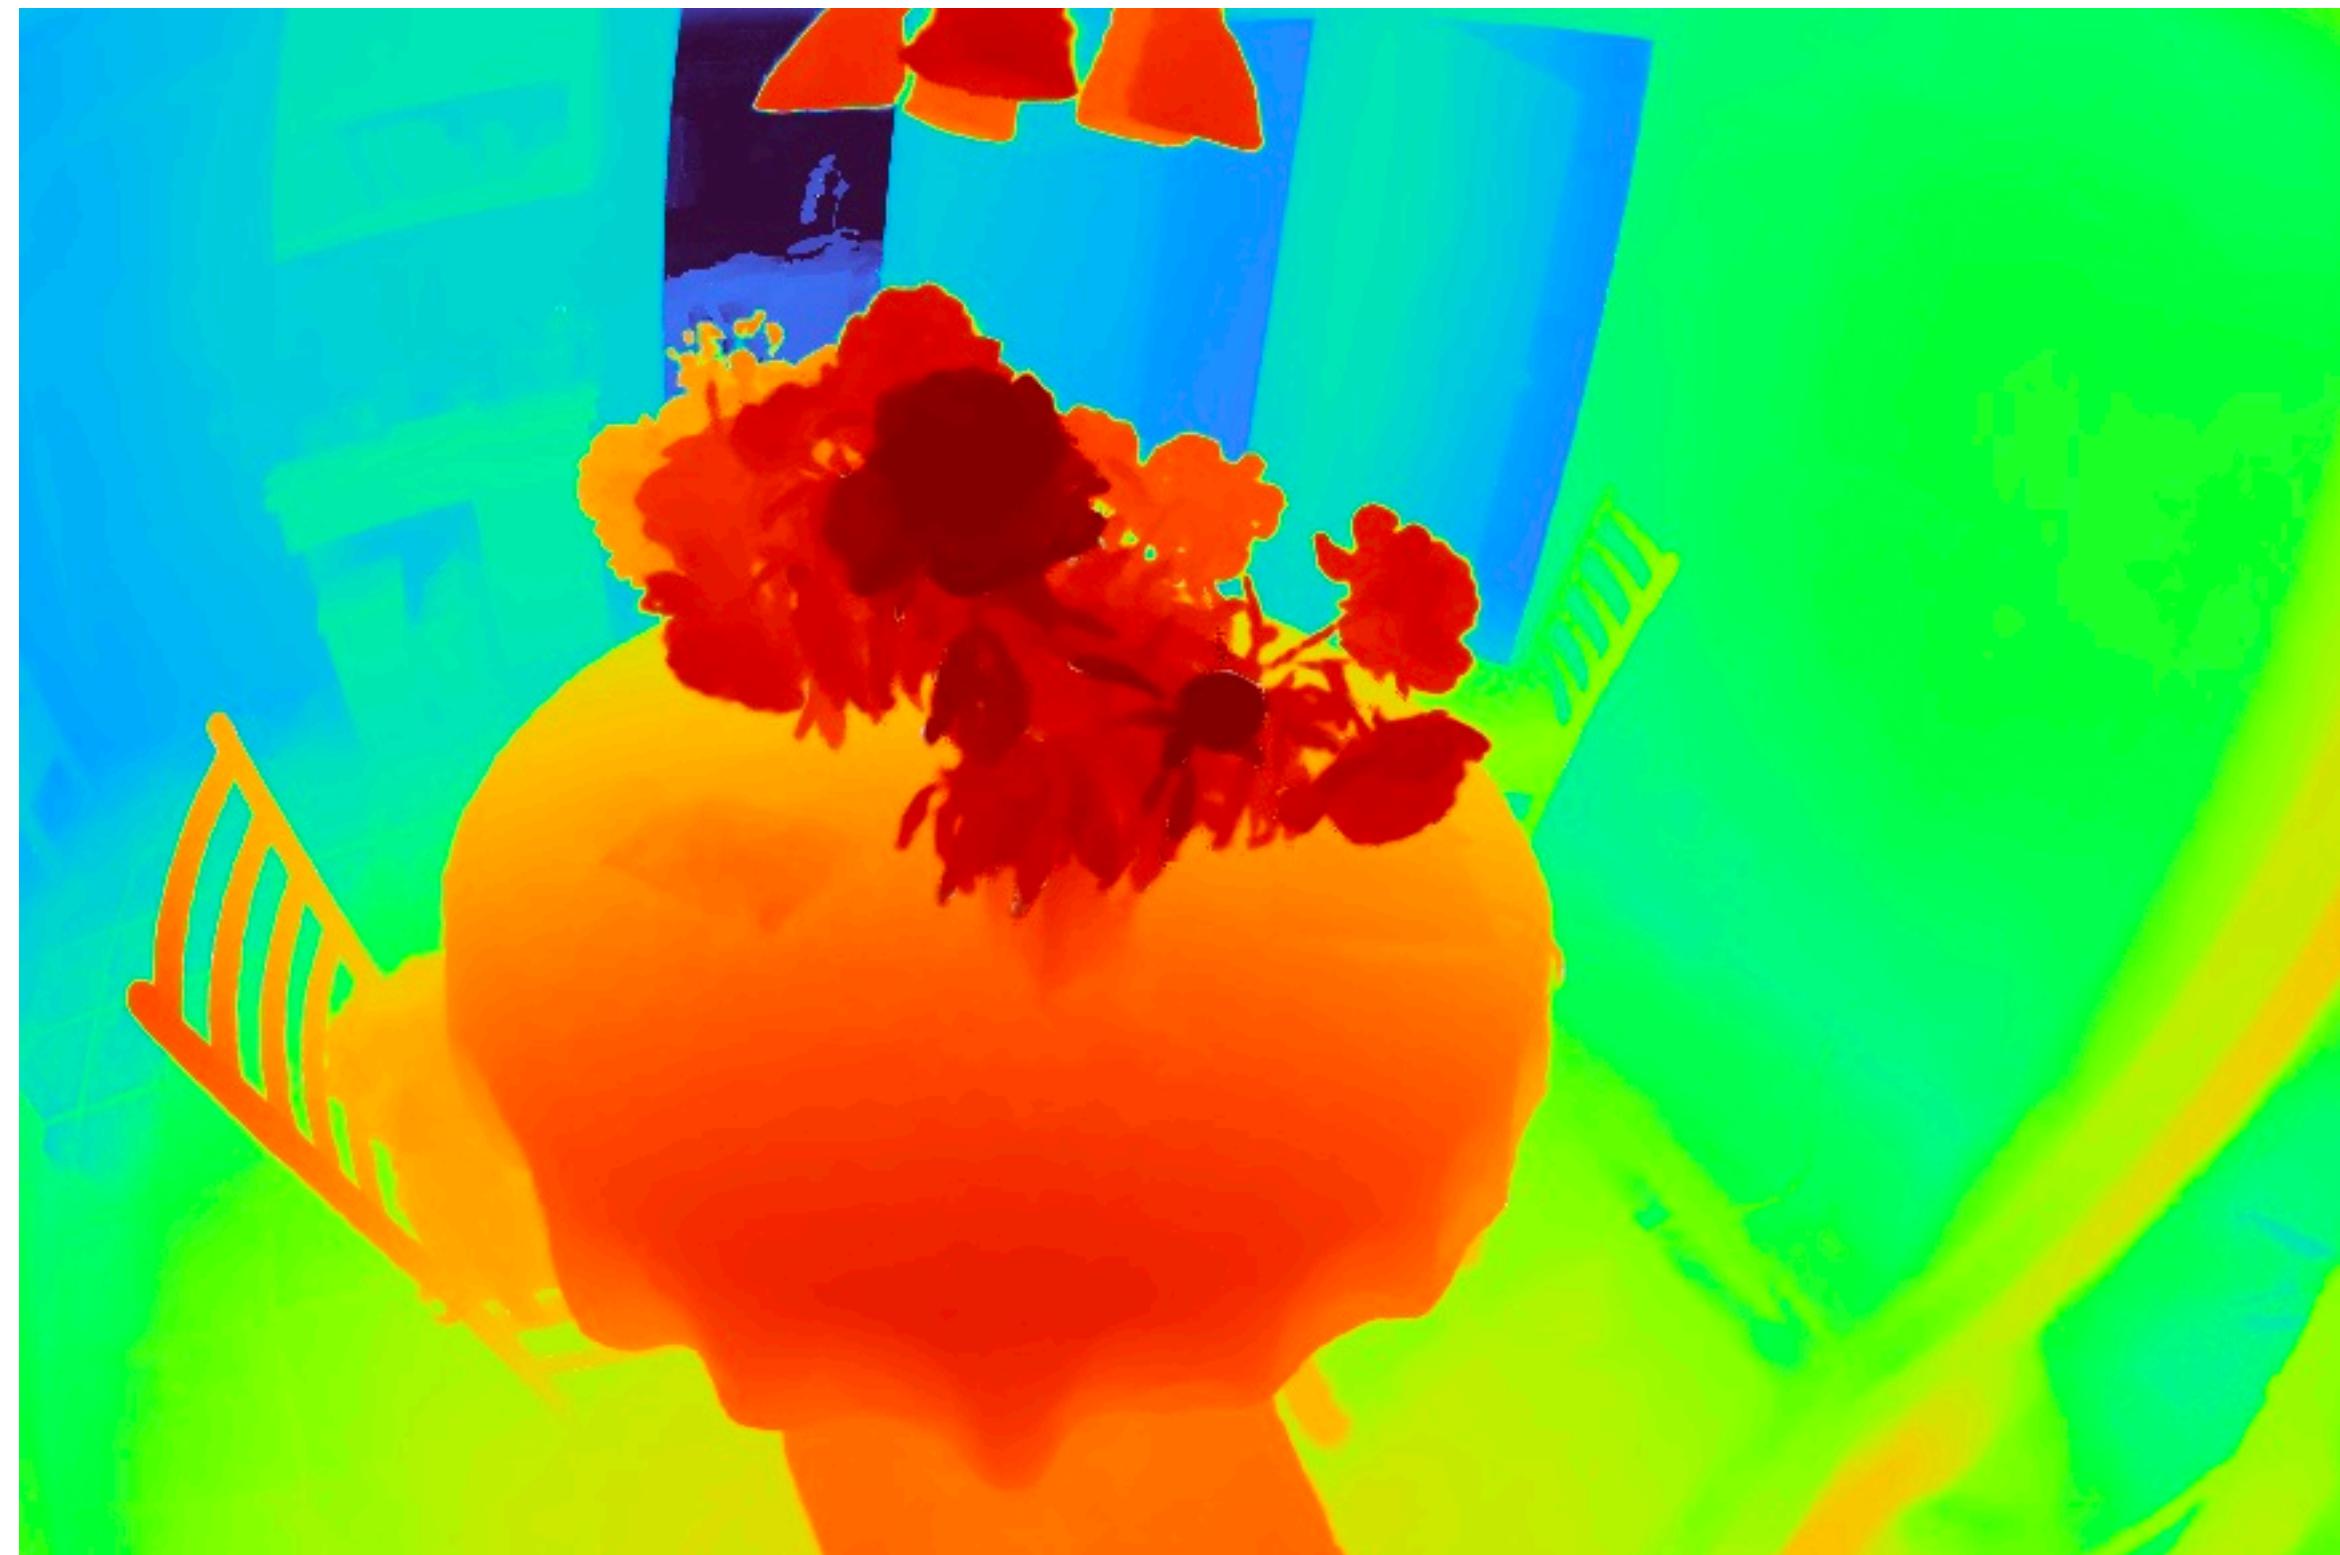
We can use this distribution to compute expectations for other quantities, e.g. “expected depth”:

$$\bar{t} = \sum_i T_i \alpha_i t_i$$

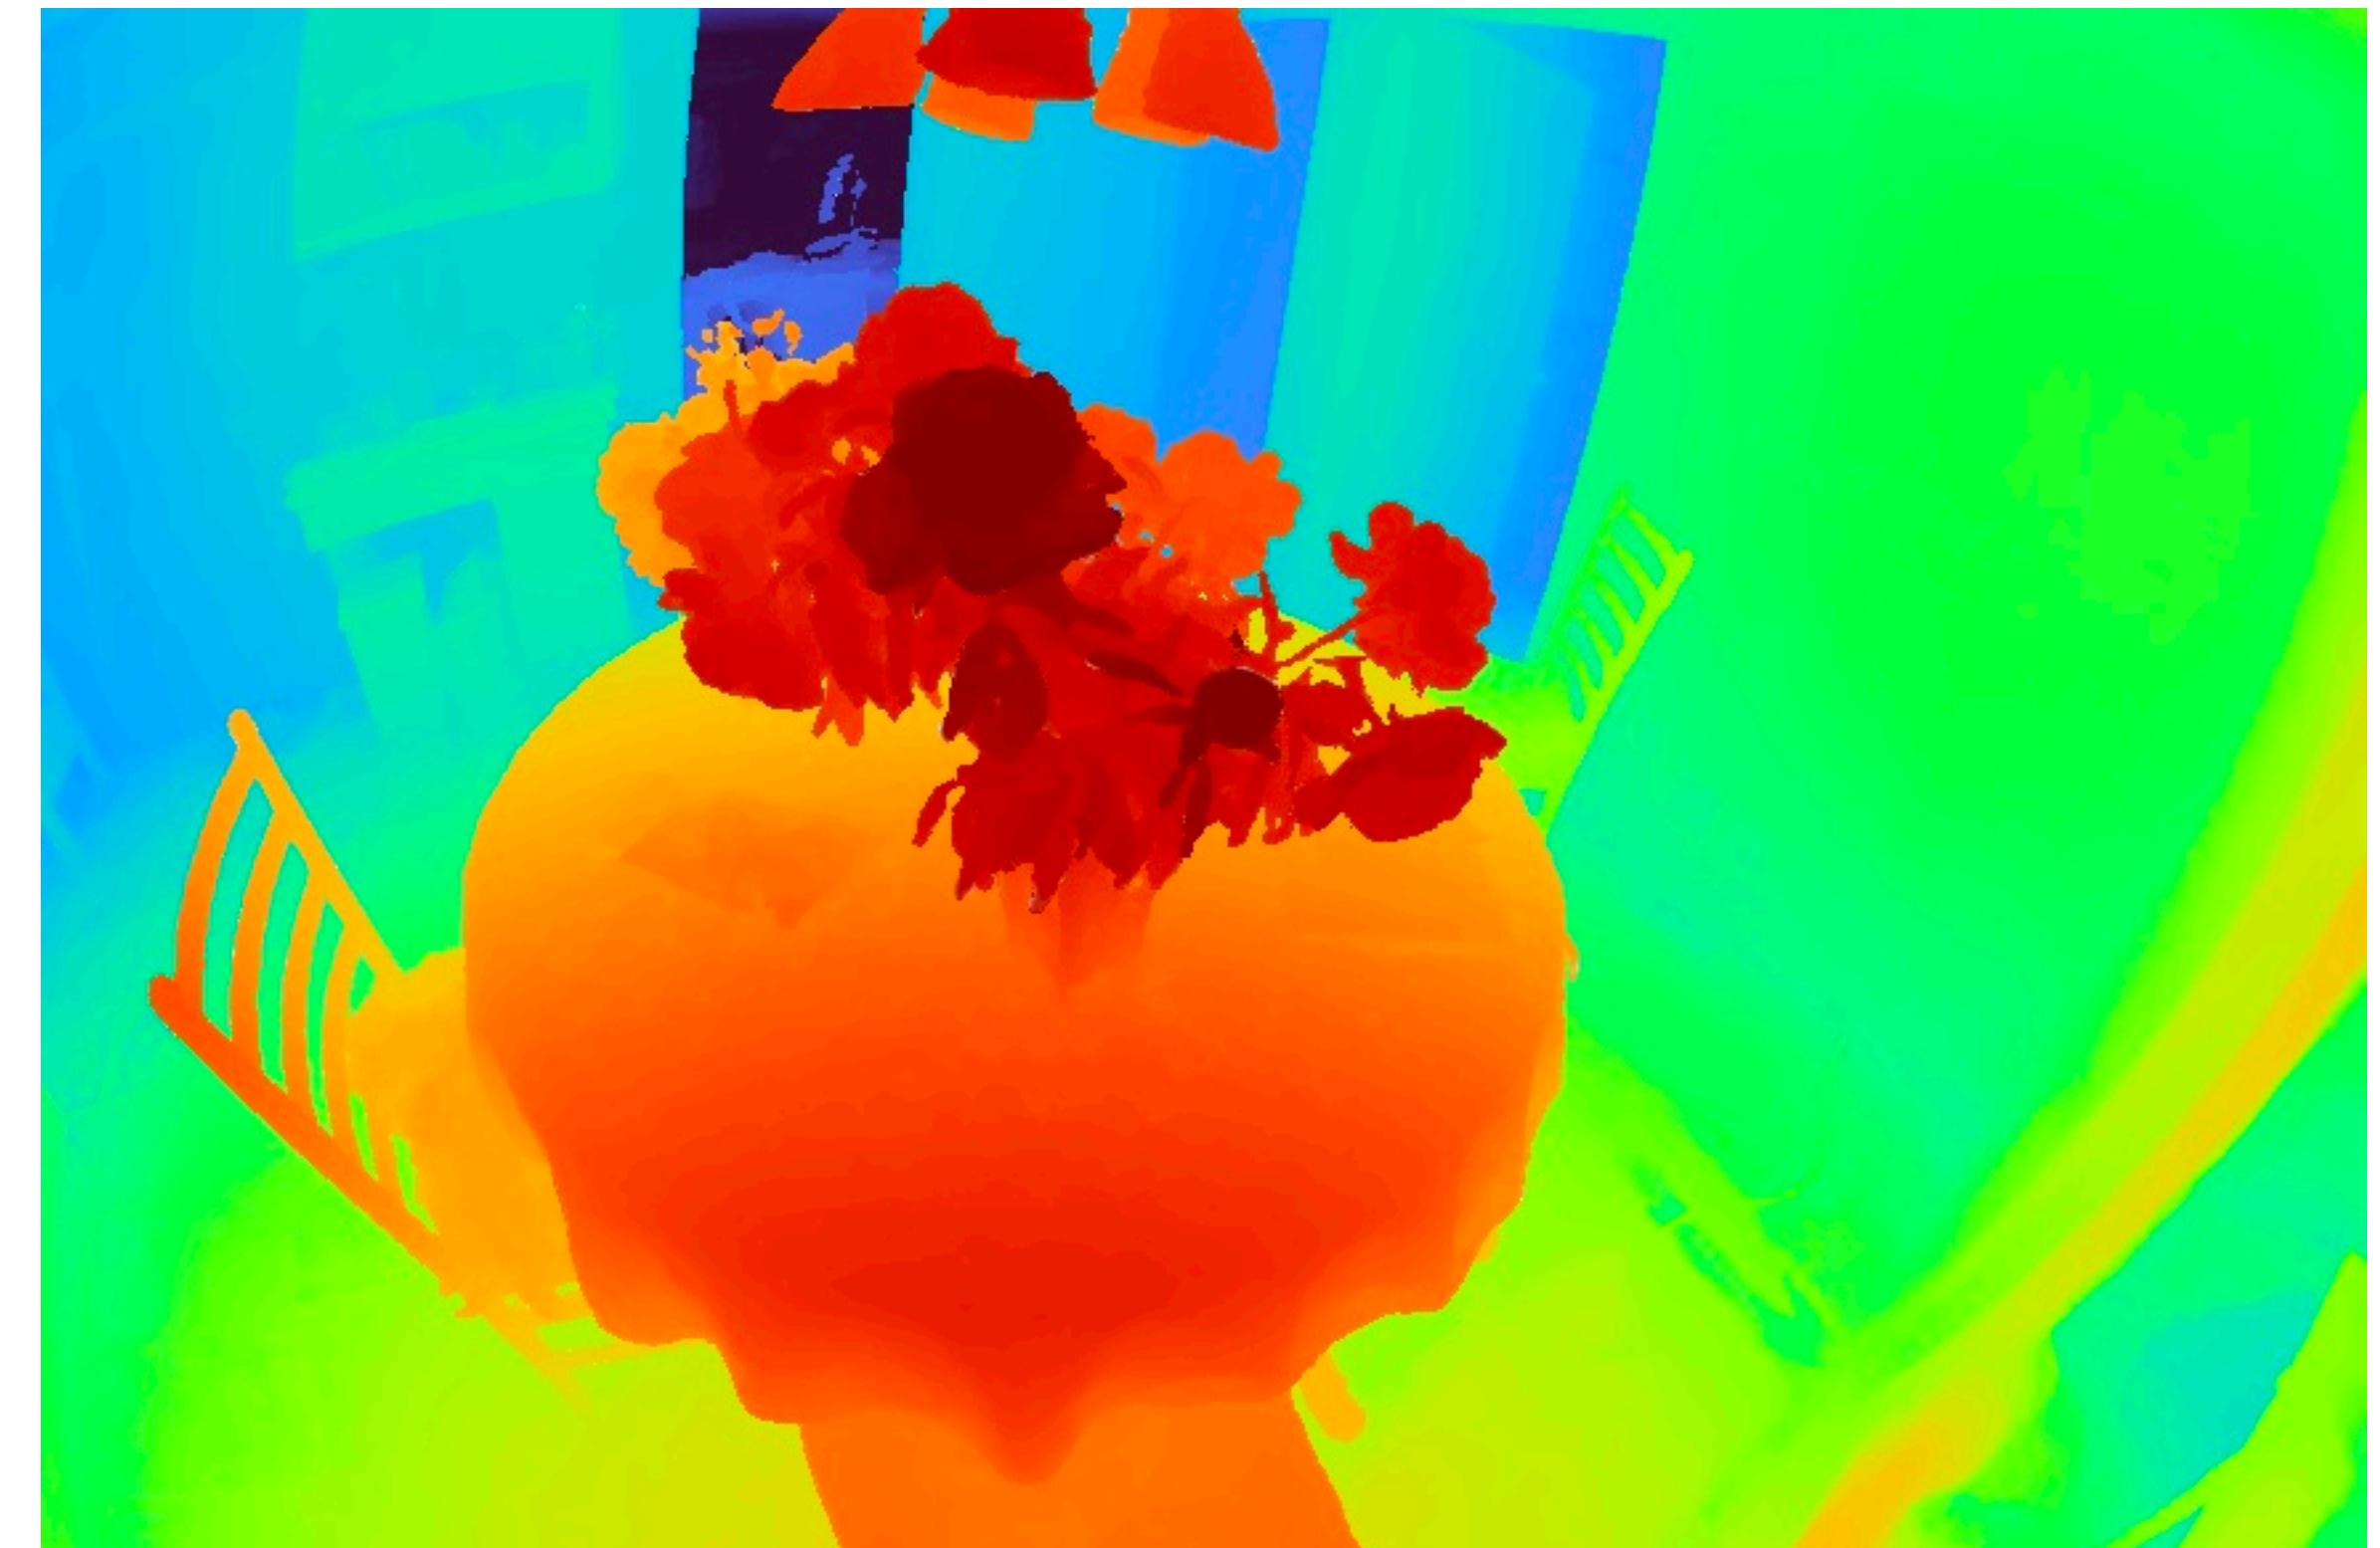
This is often how people visualise NeRF depth maps.

Alternatively, other statistics like mode or median can be used.

# Rendering depth value with the PDF

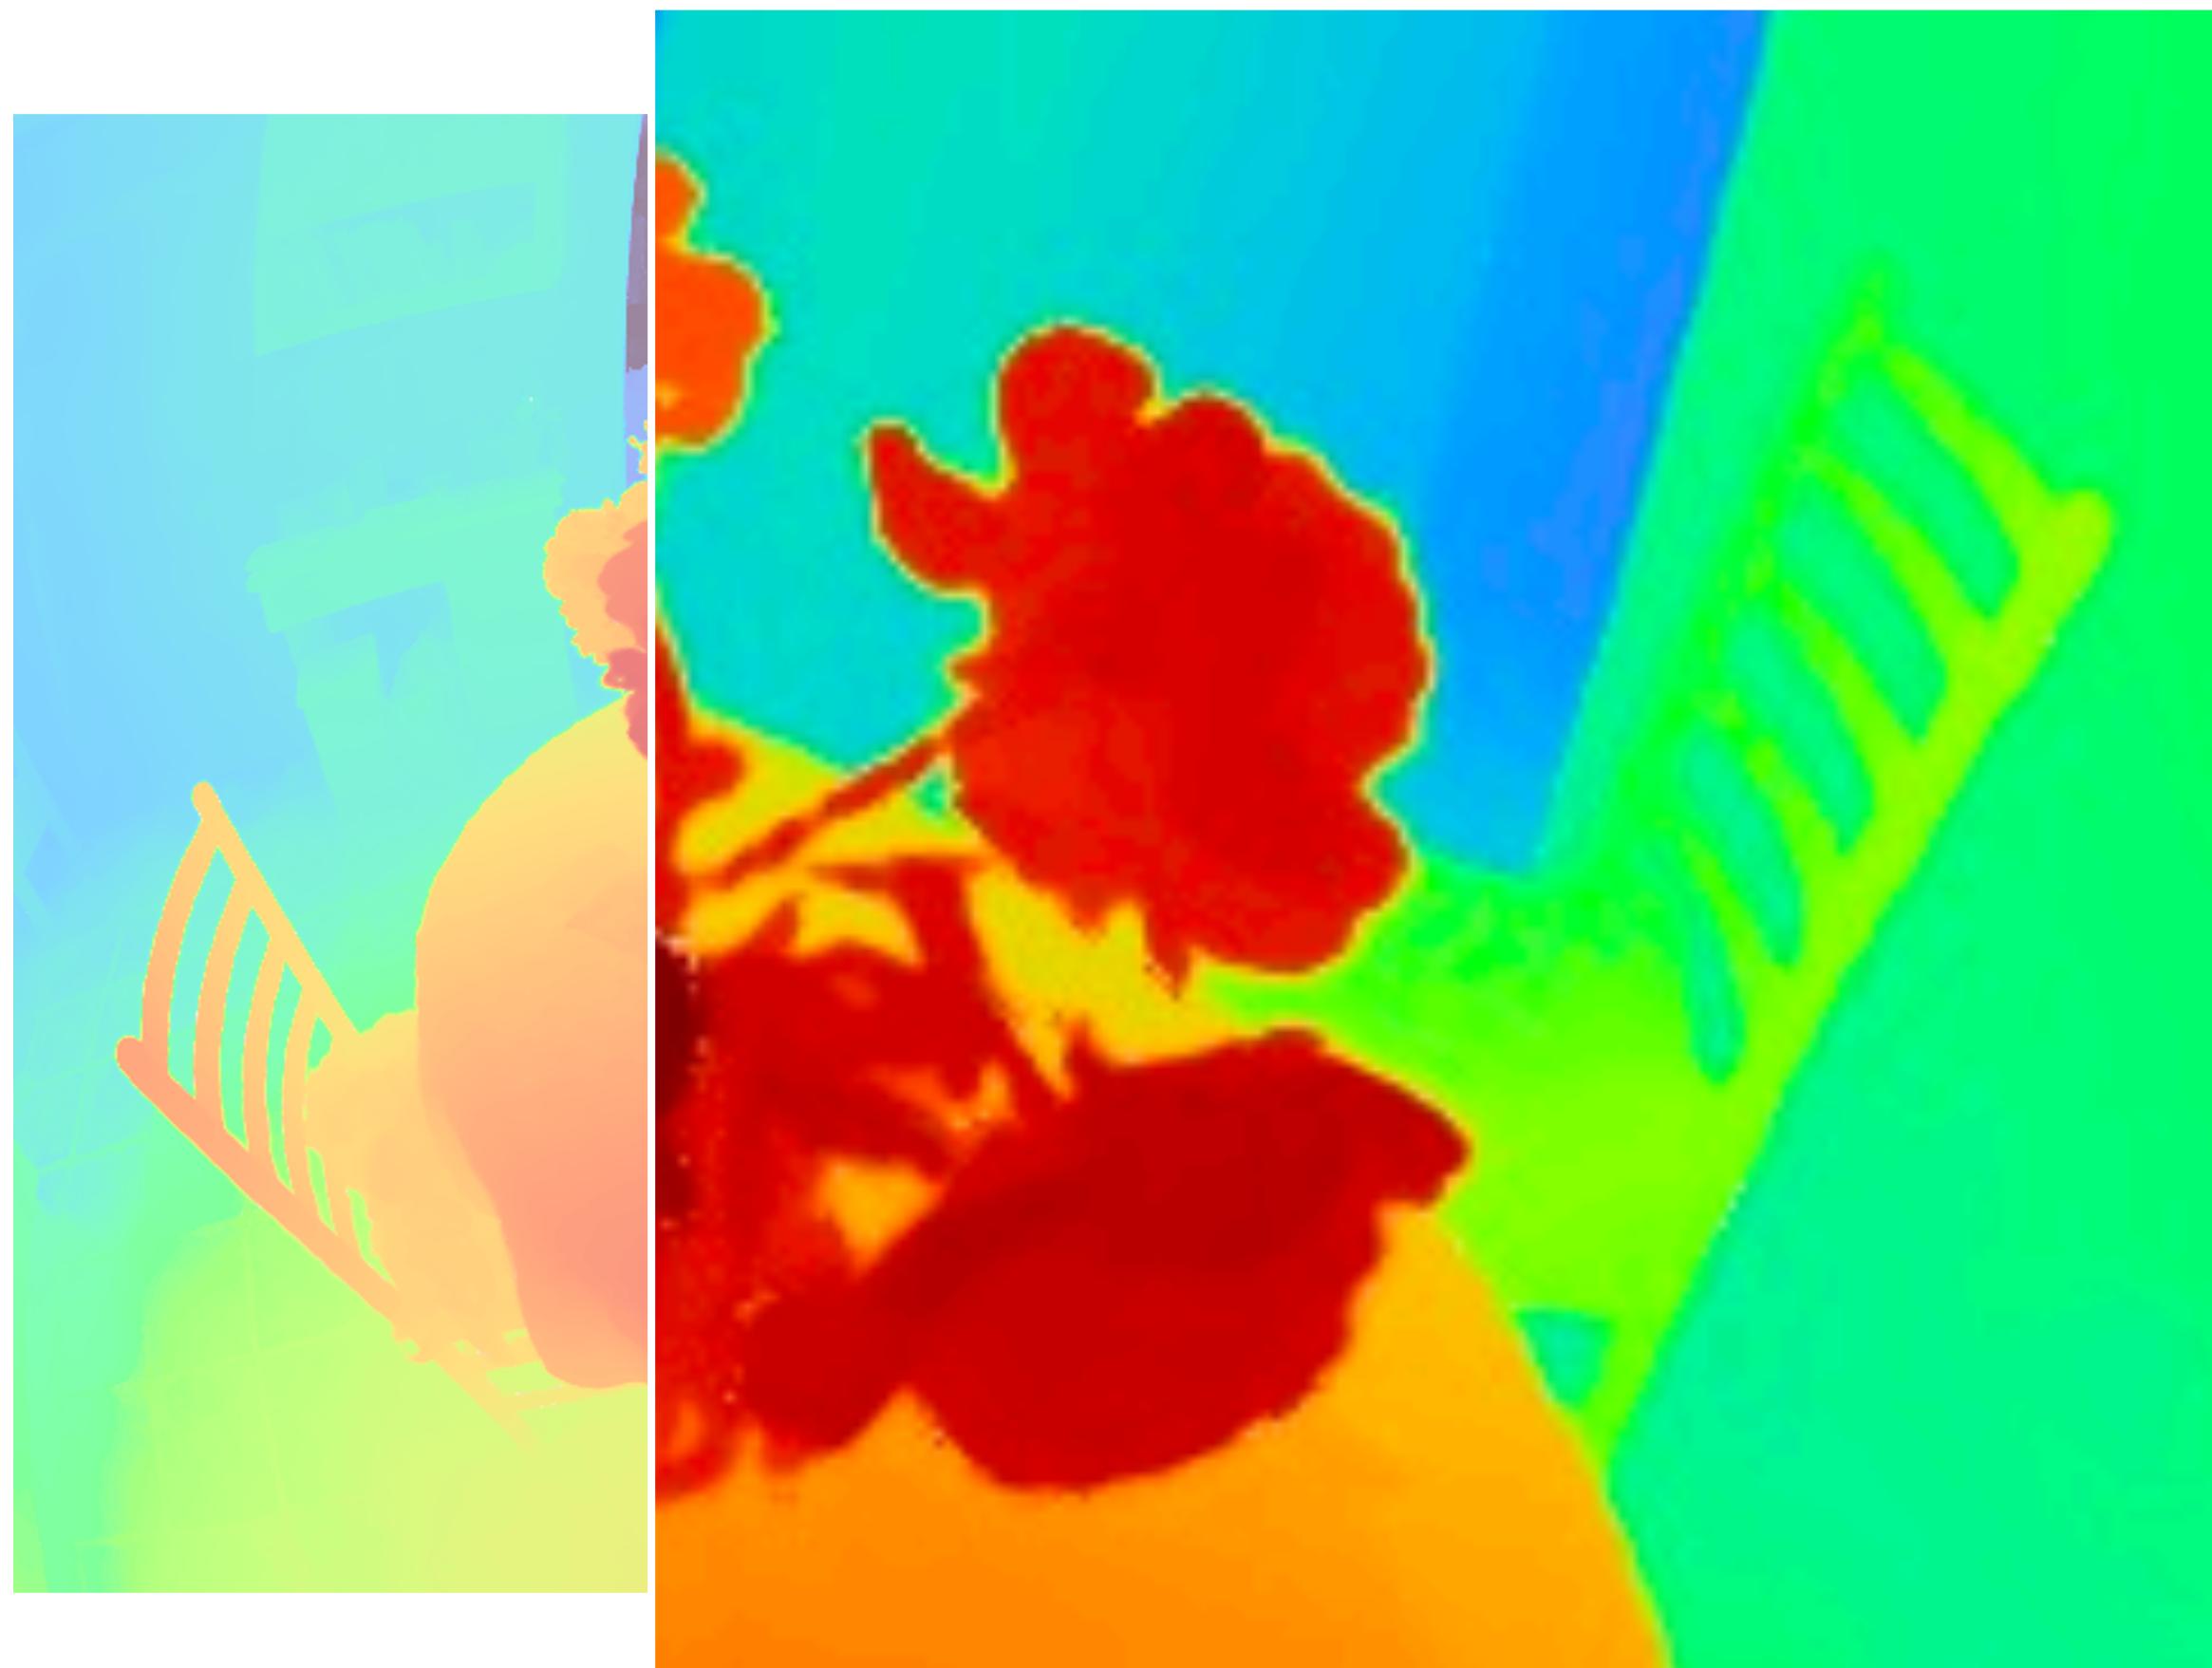


Mean depth

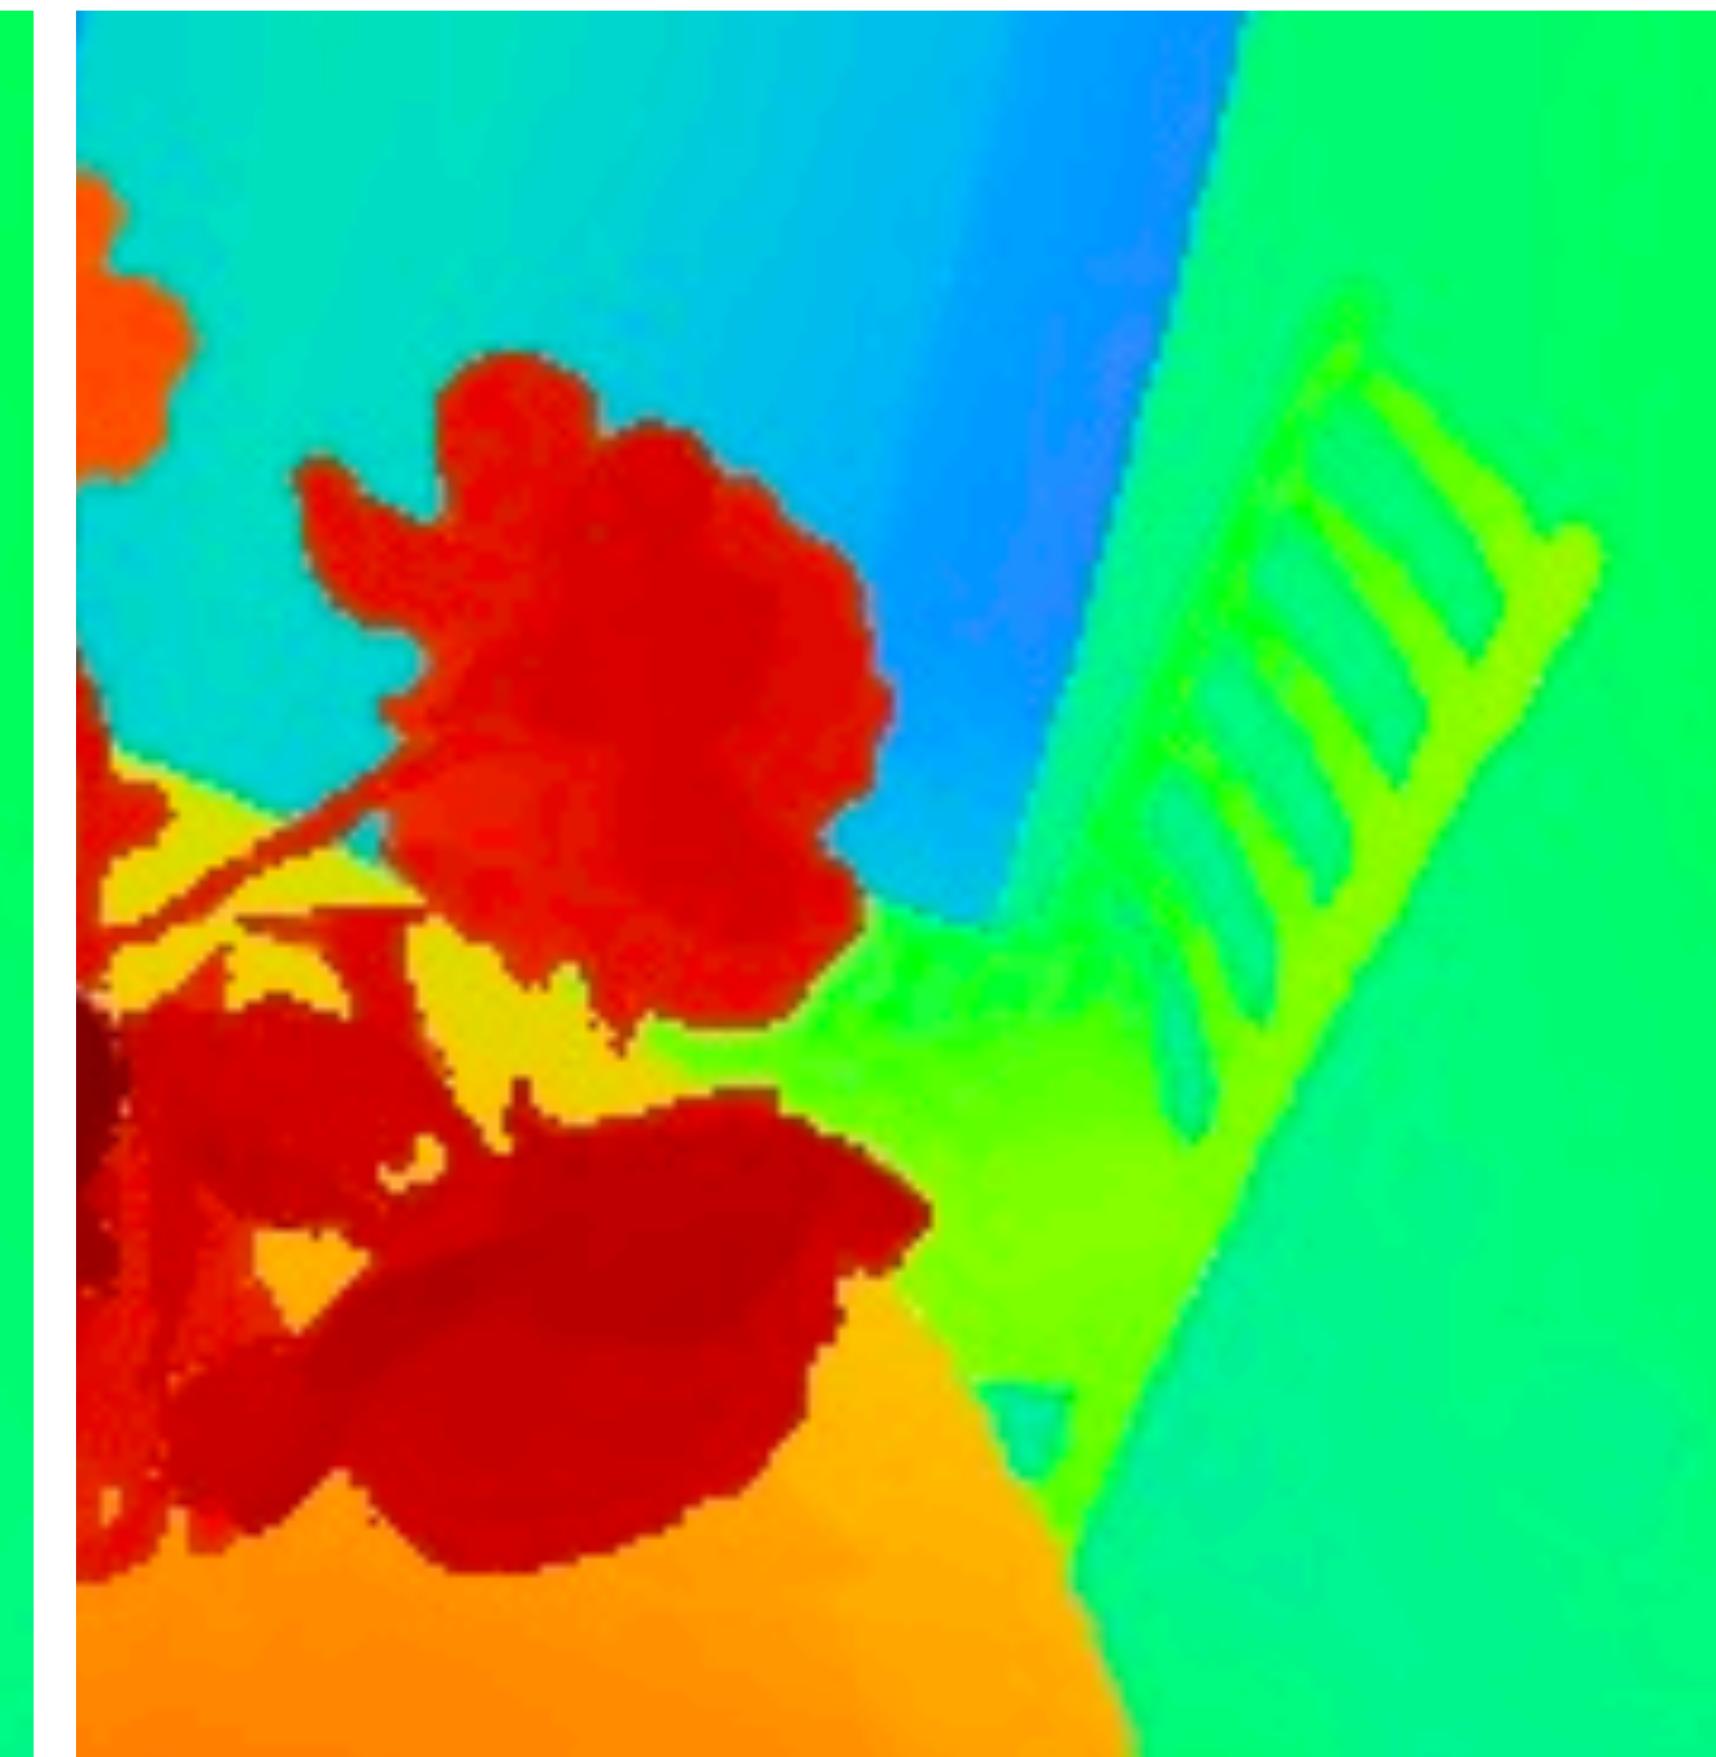


Median depth

# Rendering depth value with the PDF



Mean depth



Median depth

# Volume rendering other quantities

This idea can be used for any quantity we want to “volume render” into a 2D image. If  $\mathbf{v}$  lives in 3D space (semantic features, normal vectors, etc.)

$$\sum_i T_i \alpha_i \mathbf{v}_i$$

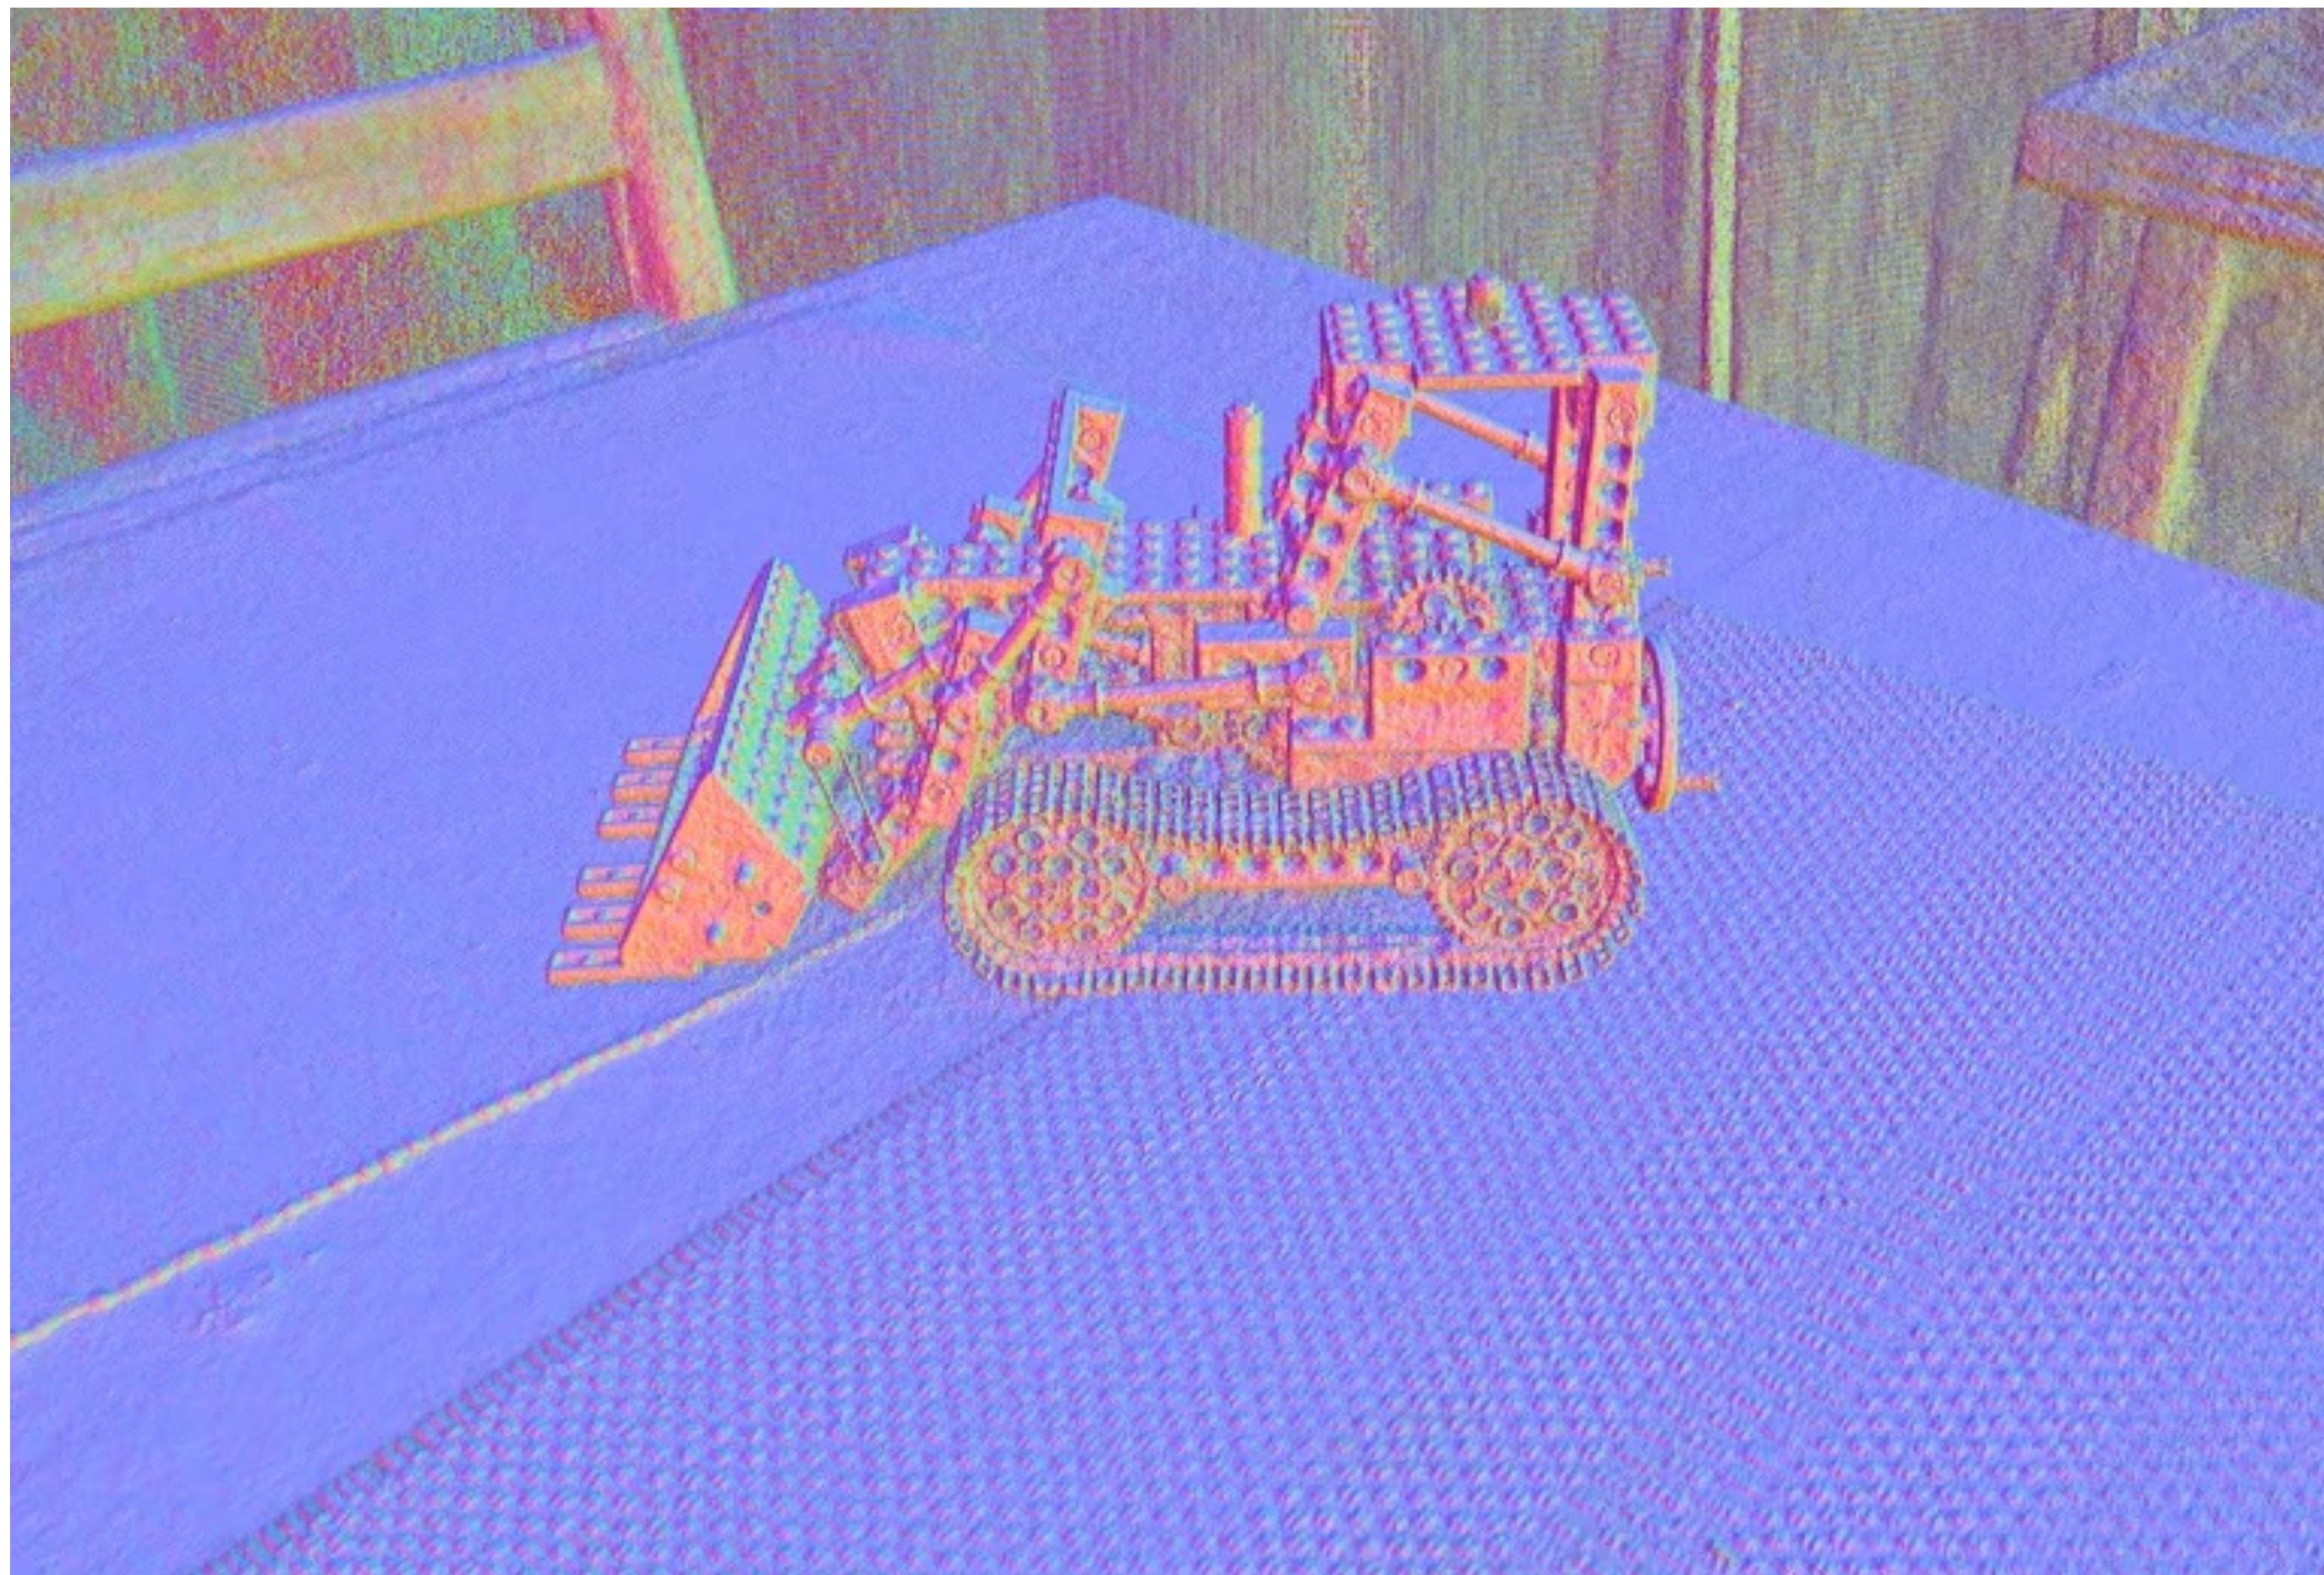
can be taken per-ray to produce 2D output images.

# Volume rendering other quantities



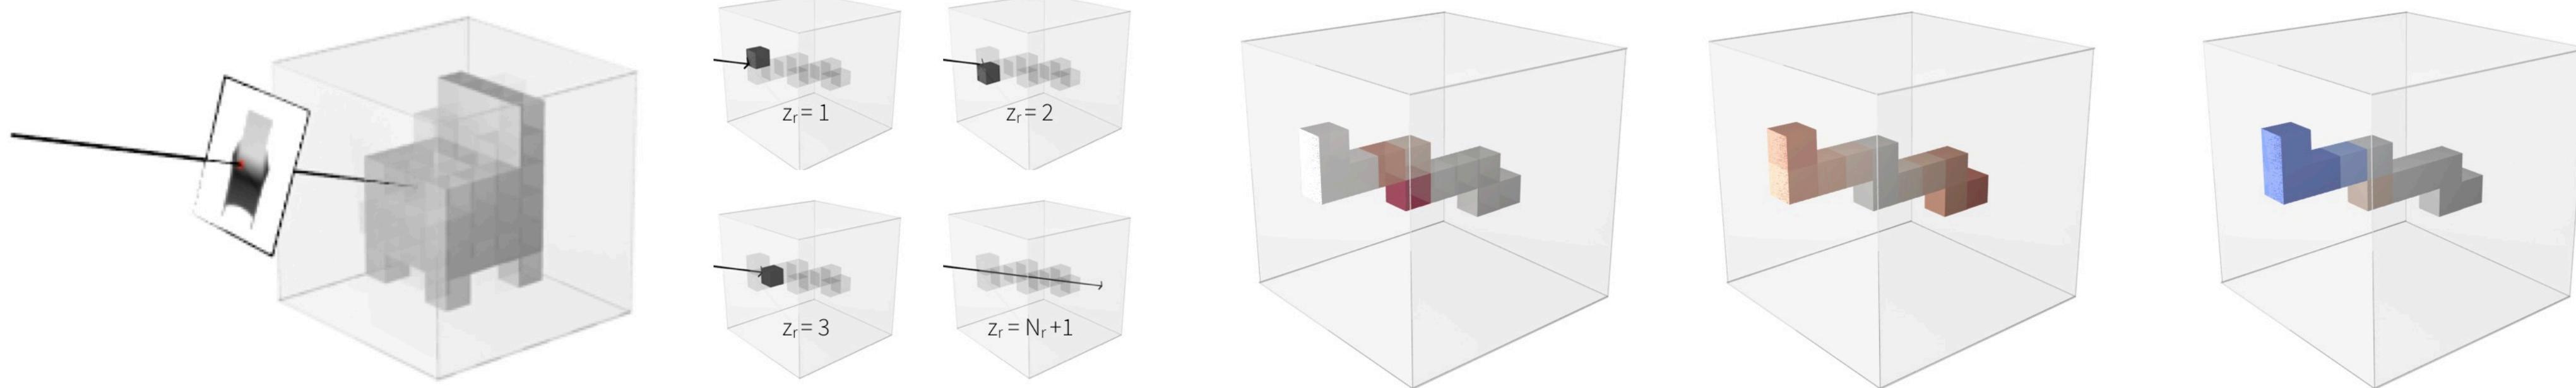
Various recent works have used this idea to render higher-level semantic feature maps (e.g., *Feature Field Distillation* and *Neural Feature Fusion Fields*).

# Density as geometry



Normal vectors (from analytic gradient of density)

# Alpha compositing model in ML/computer vision



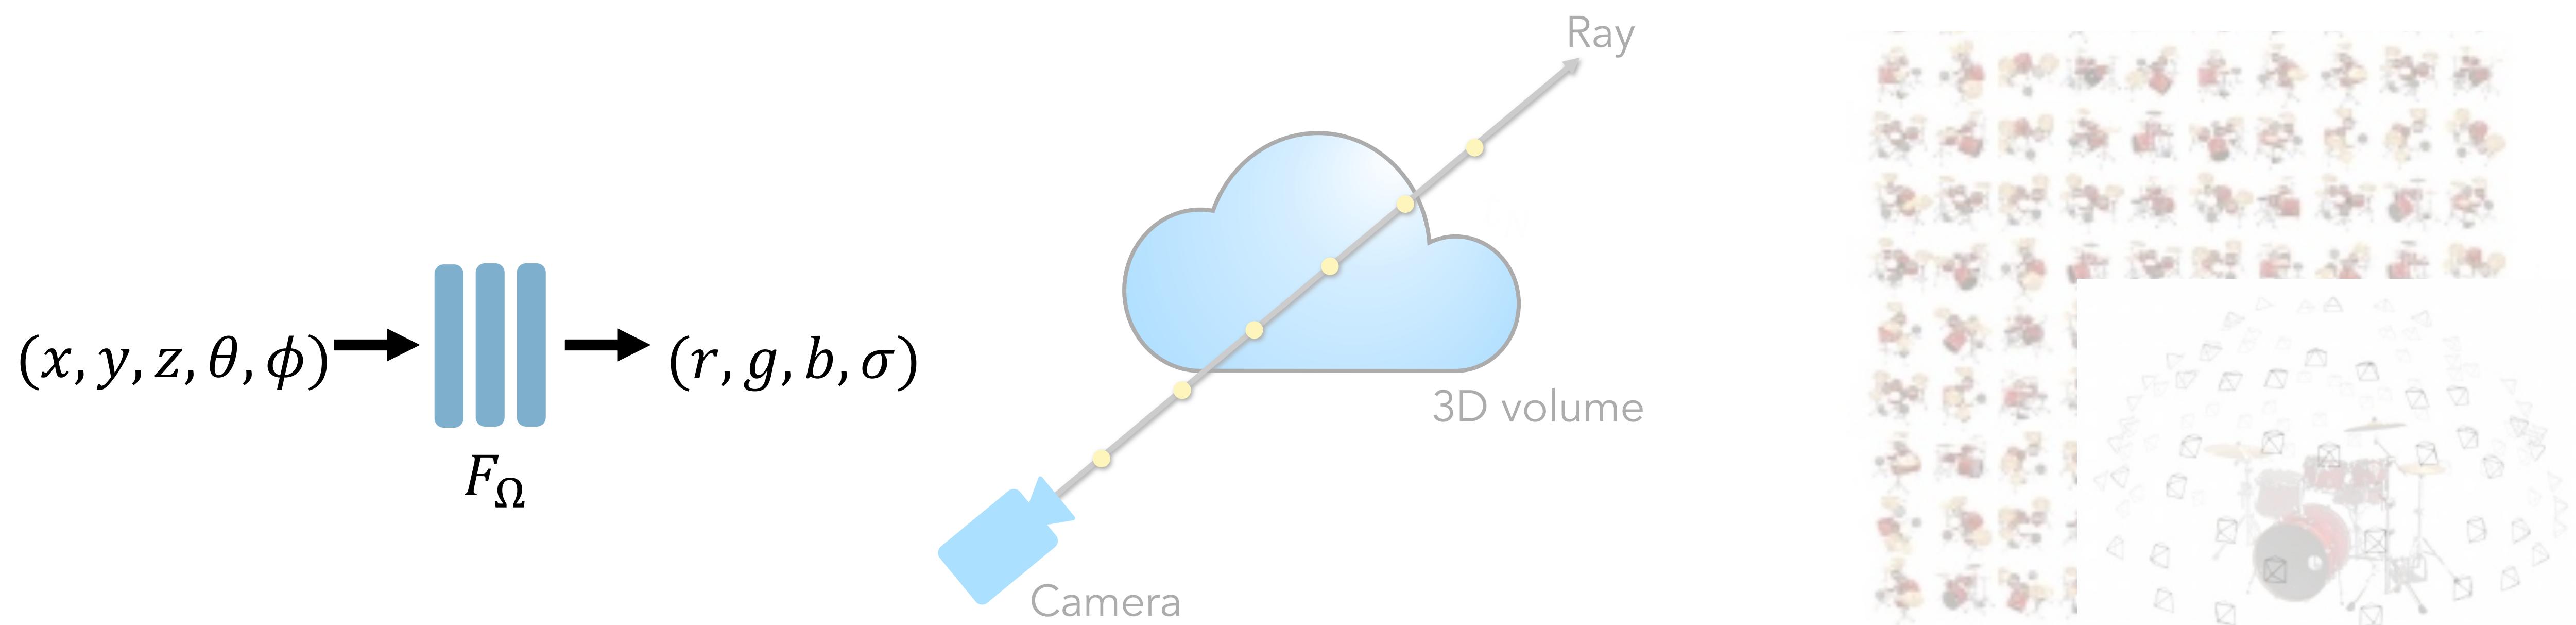
*Differentiable ray consistency* work used a forward model with “probabilistic occupancy” to supervise 3D-from-single-image prediction. Same rendering model as alpha compositing!

$$p(z_r = i) = \begin{cases} (1 - x_i^r) \prod_{j=1}^{i-1} x_j^r, & \text{if } i \leq N_r \\ \prod_{j=1}^{N_r} x_j^r, & \text{if } i = N_r + 1 \end{cases}$$

# Where we are

1. Birds Eye View & Background
2. Volumetric Rendering Function
- 3. Encoding and Representing 3D Volumes**
4. Signal Processing Considerations
5. Challenges & Pointers

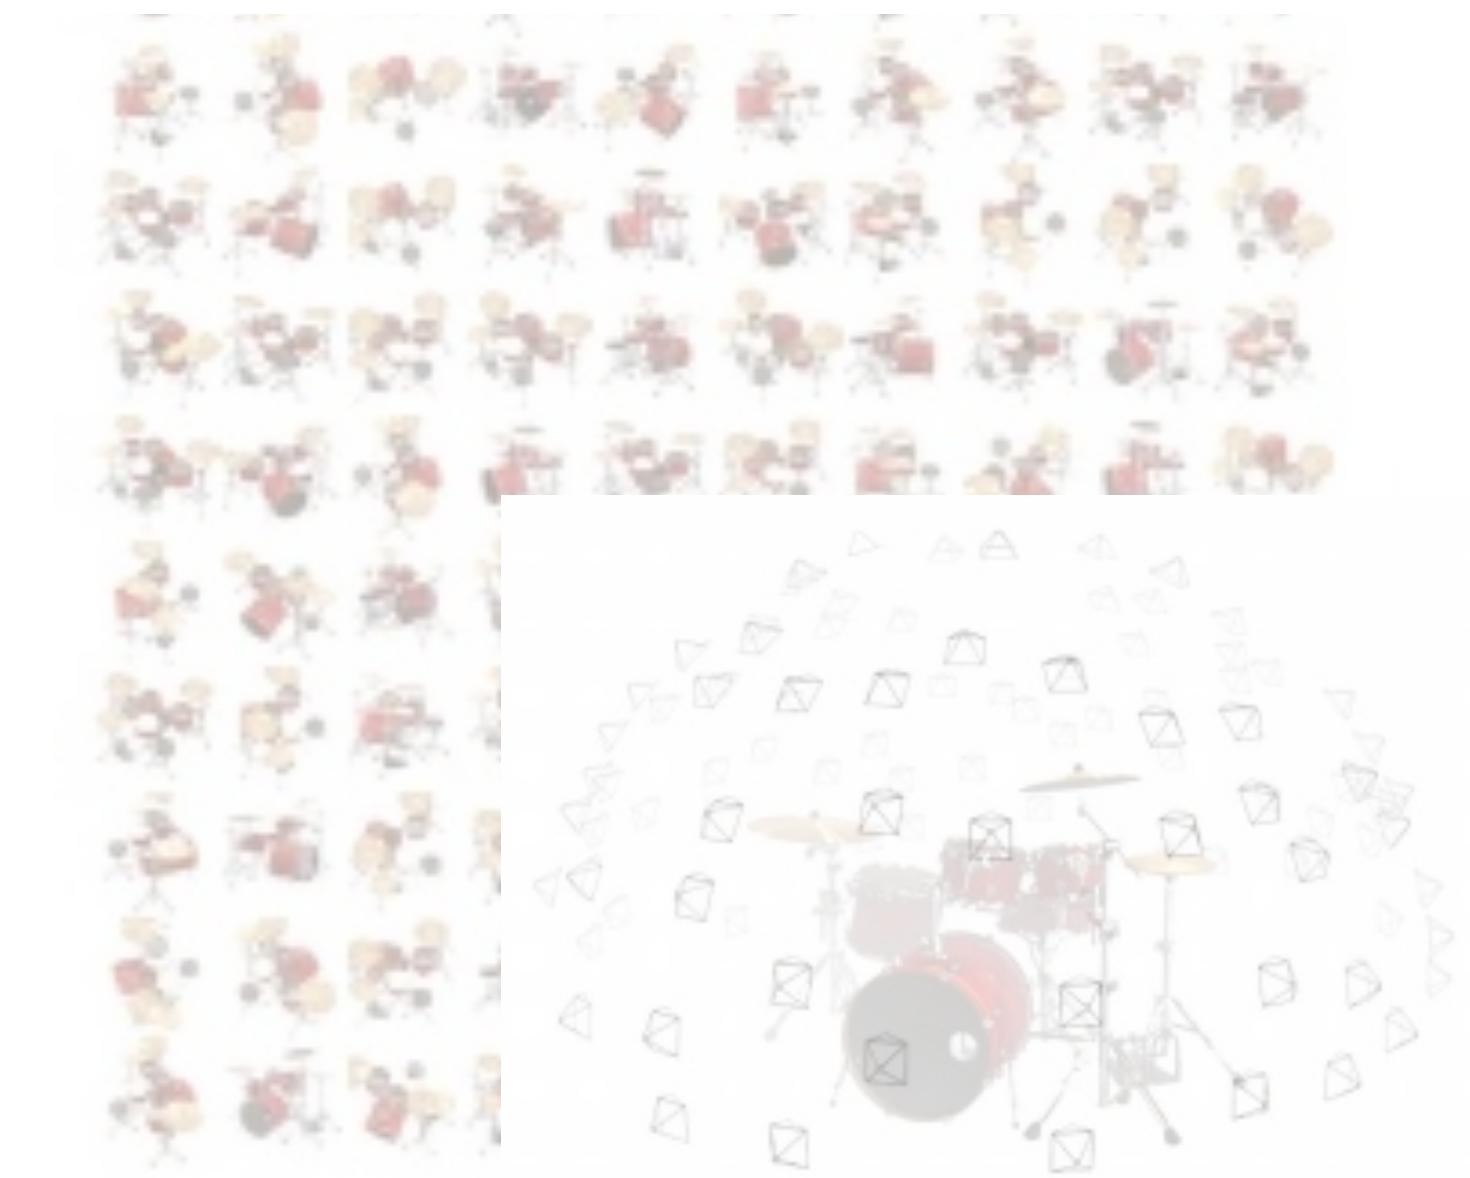
# Three Key Components



Neural Volumetric 3D  
Scene Representation

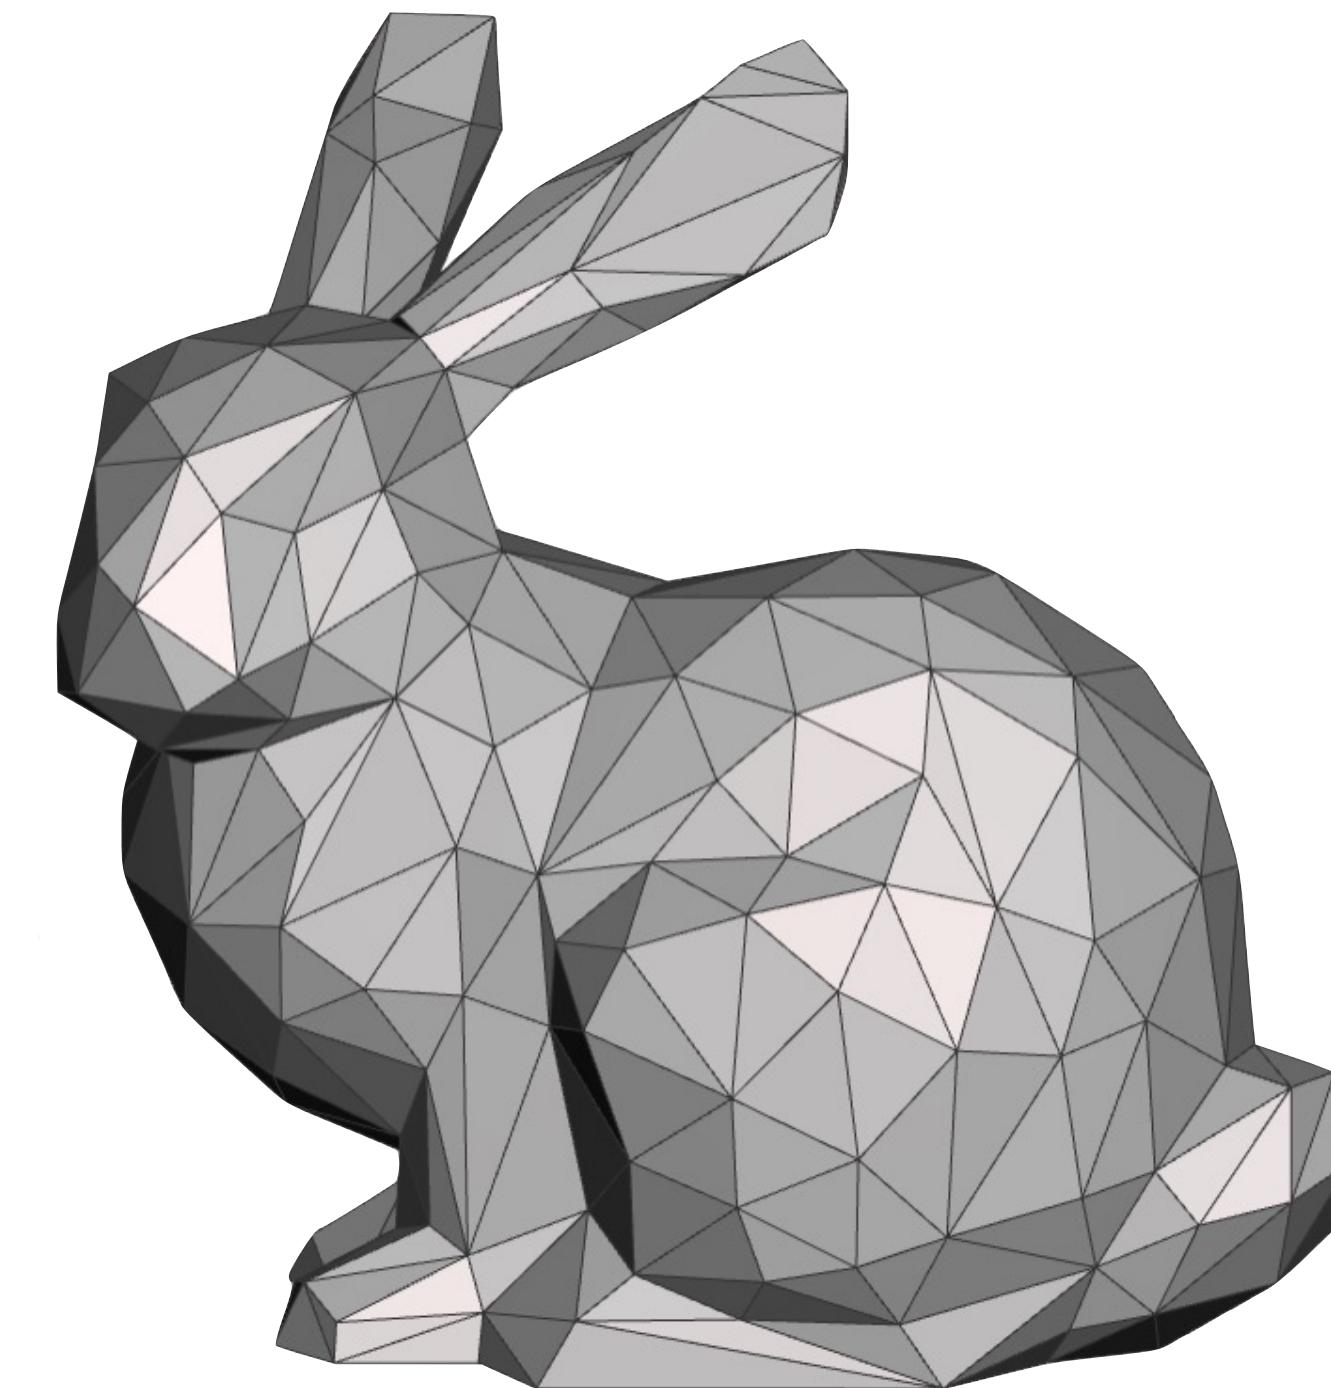
Differentiable Volumetric  
Rendering Function

Objective: Synthesize  
all training views

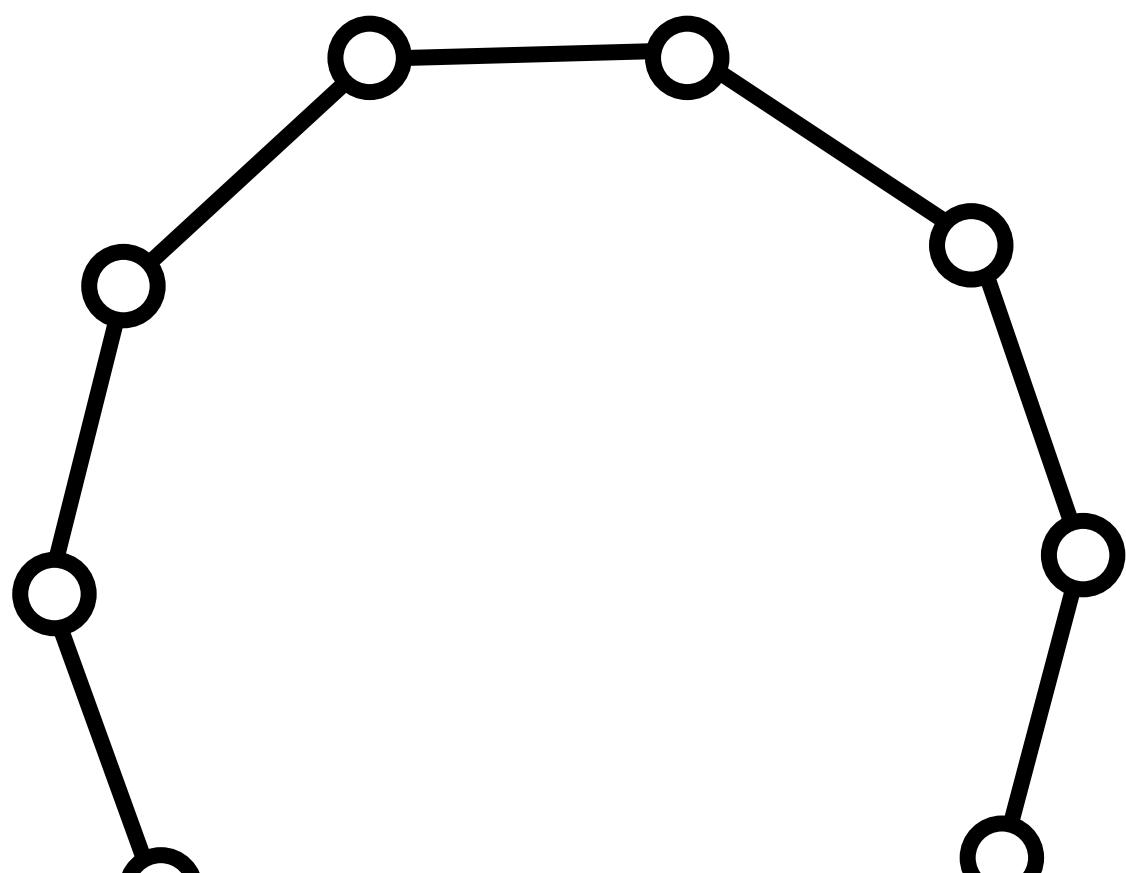


Optimization via  
Analysis-by-Synthesis

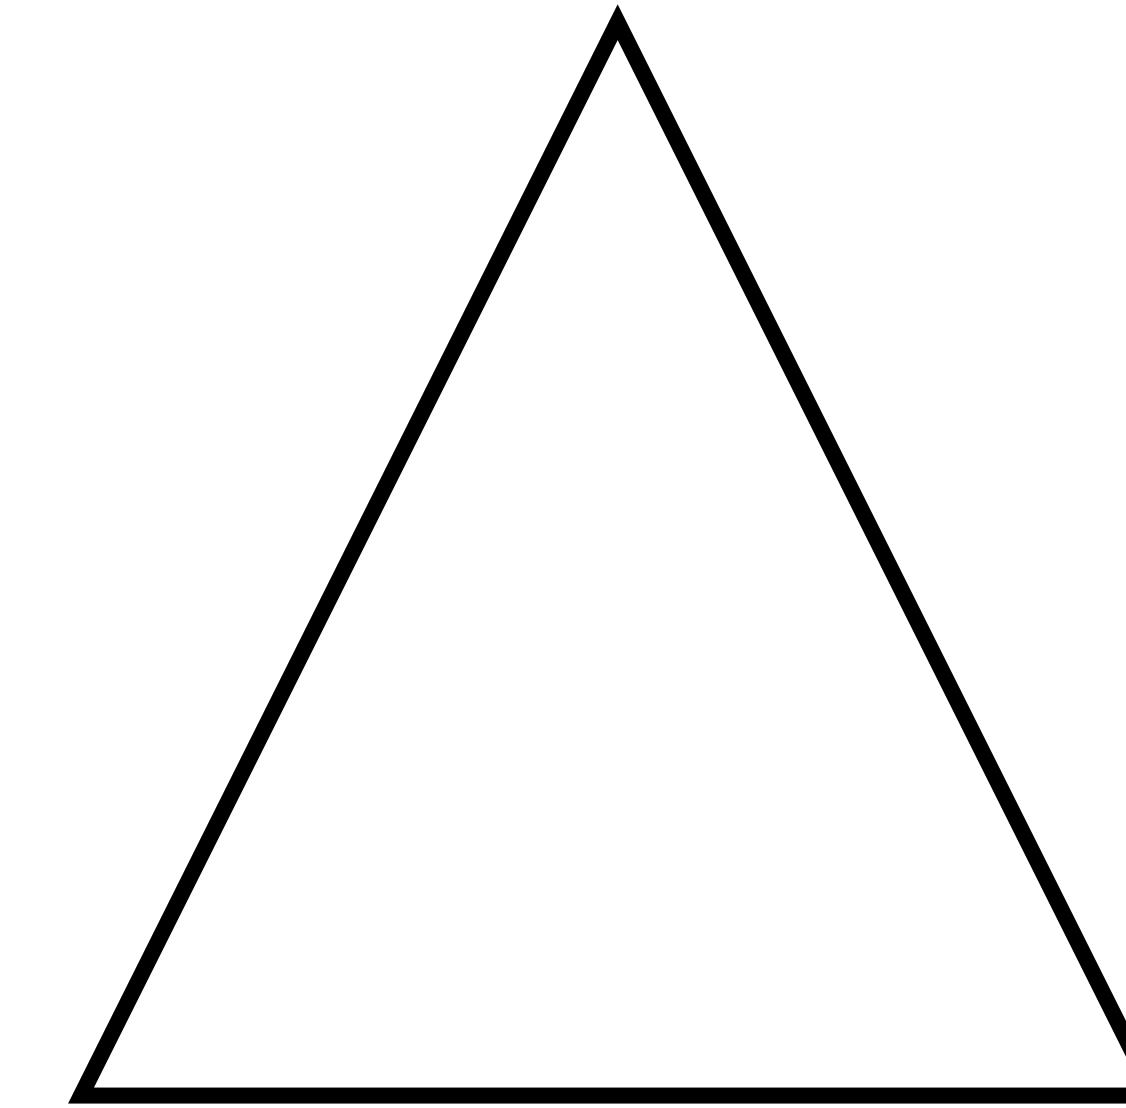
# Mesh Representation



# Gradient Based Optimization

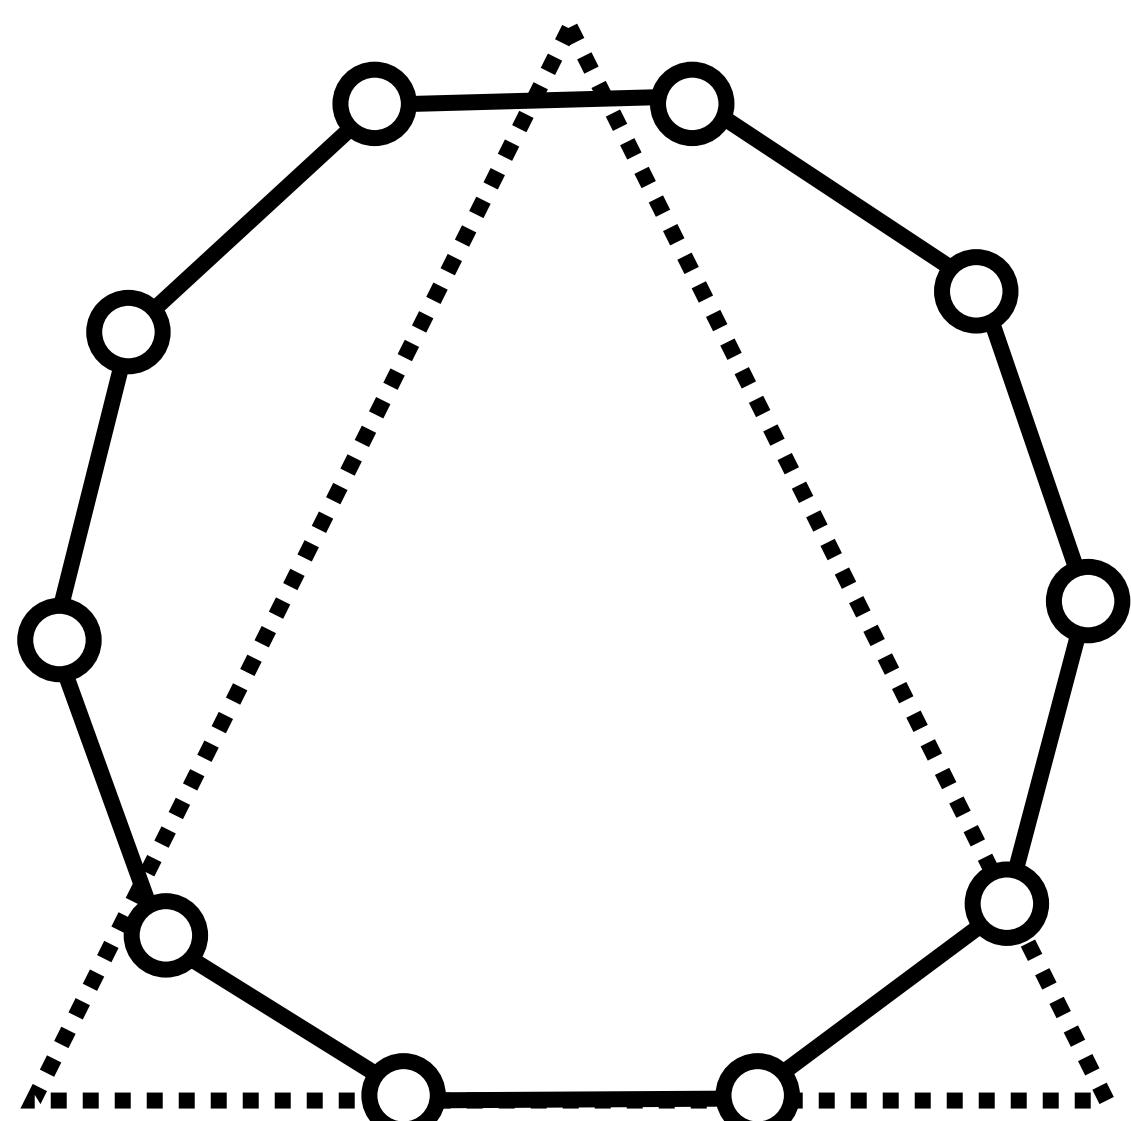


Initial Geometry

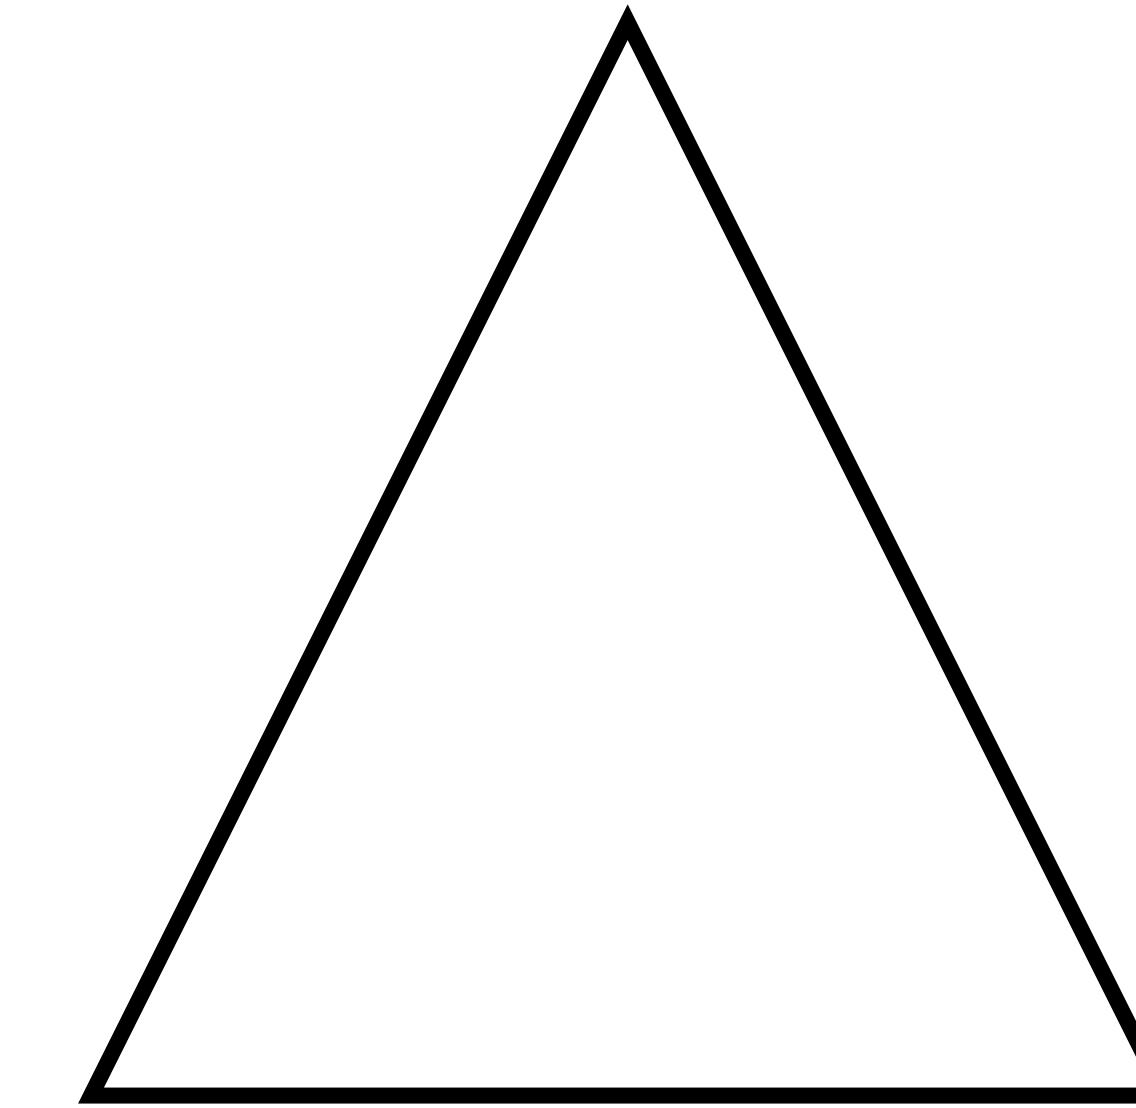


Target Geometry

# Gradient Based Optimization

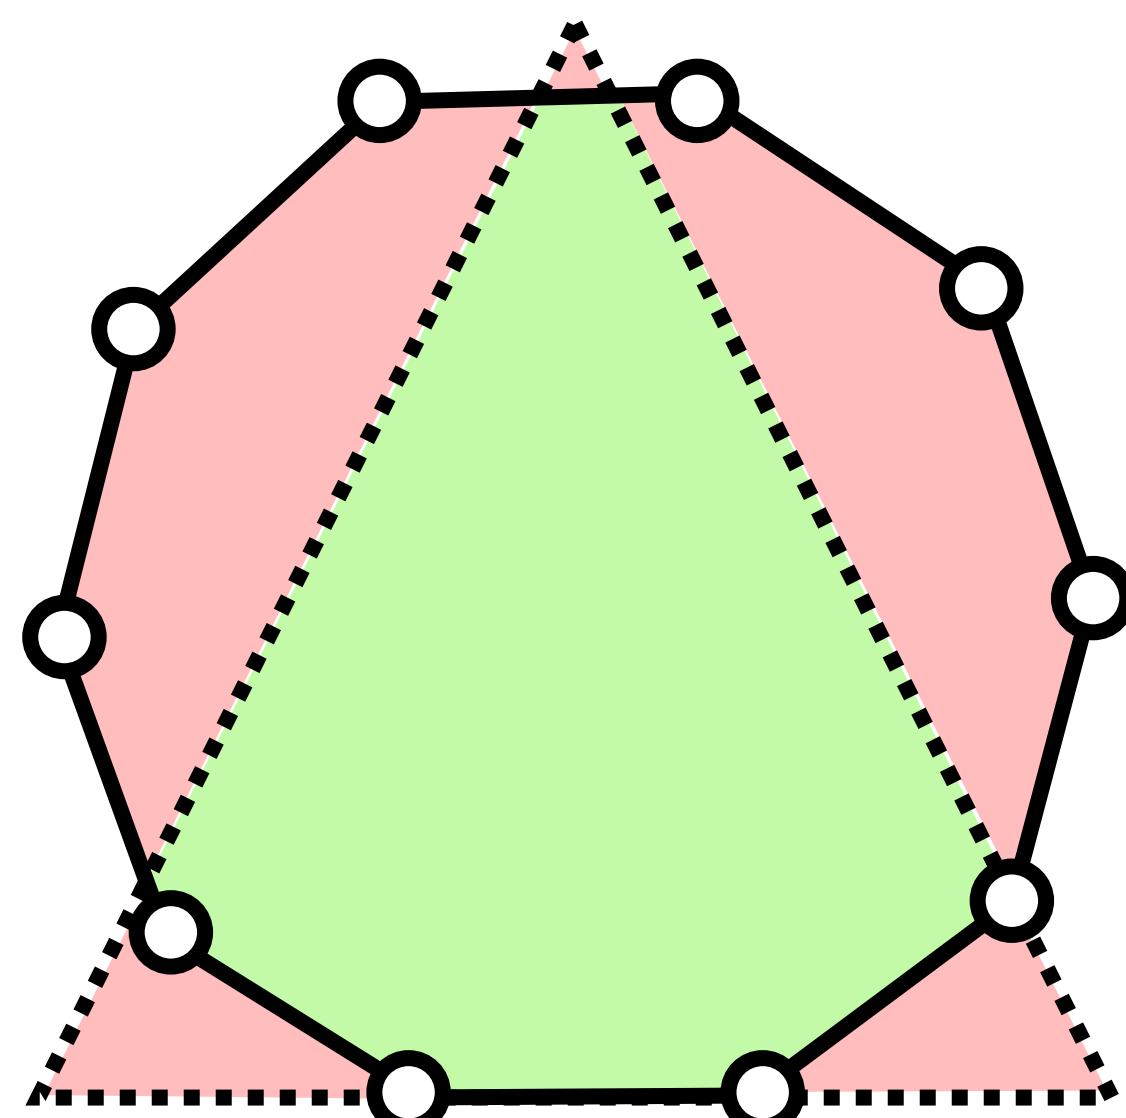


Initial Geometry

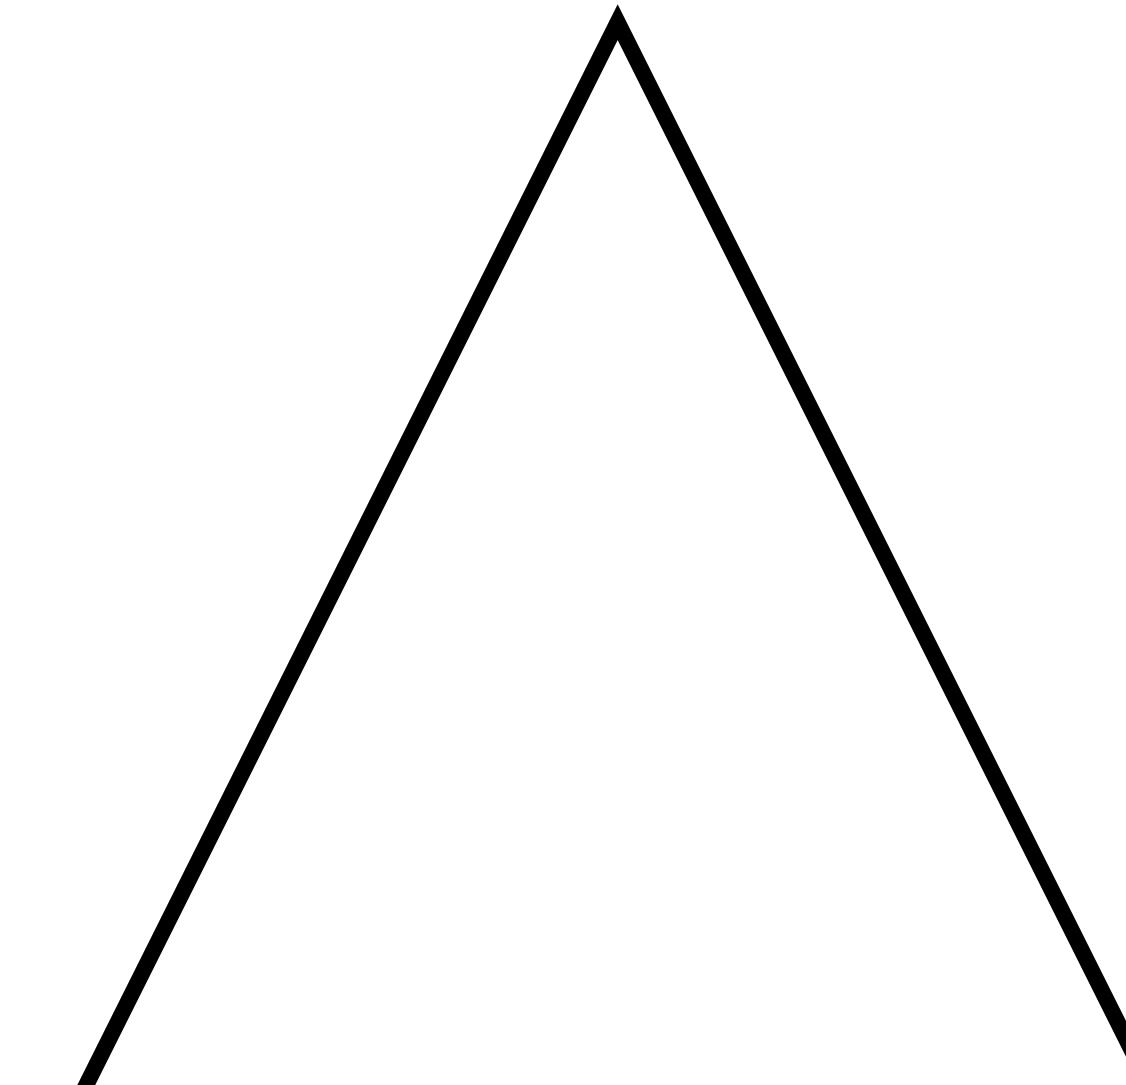


Target Geometry

# Gradient Based Optimization

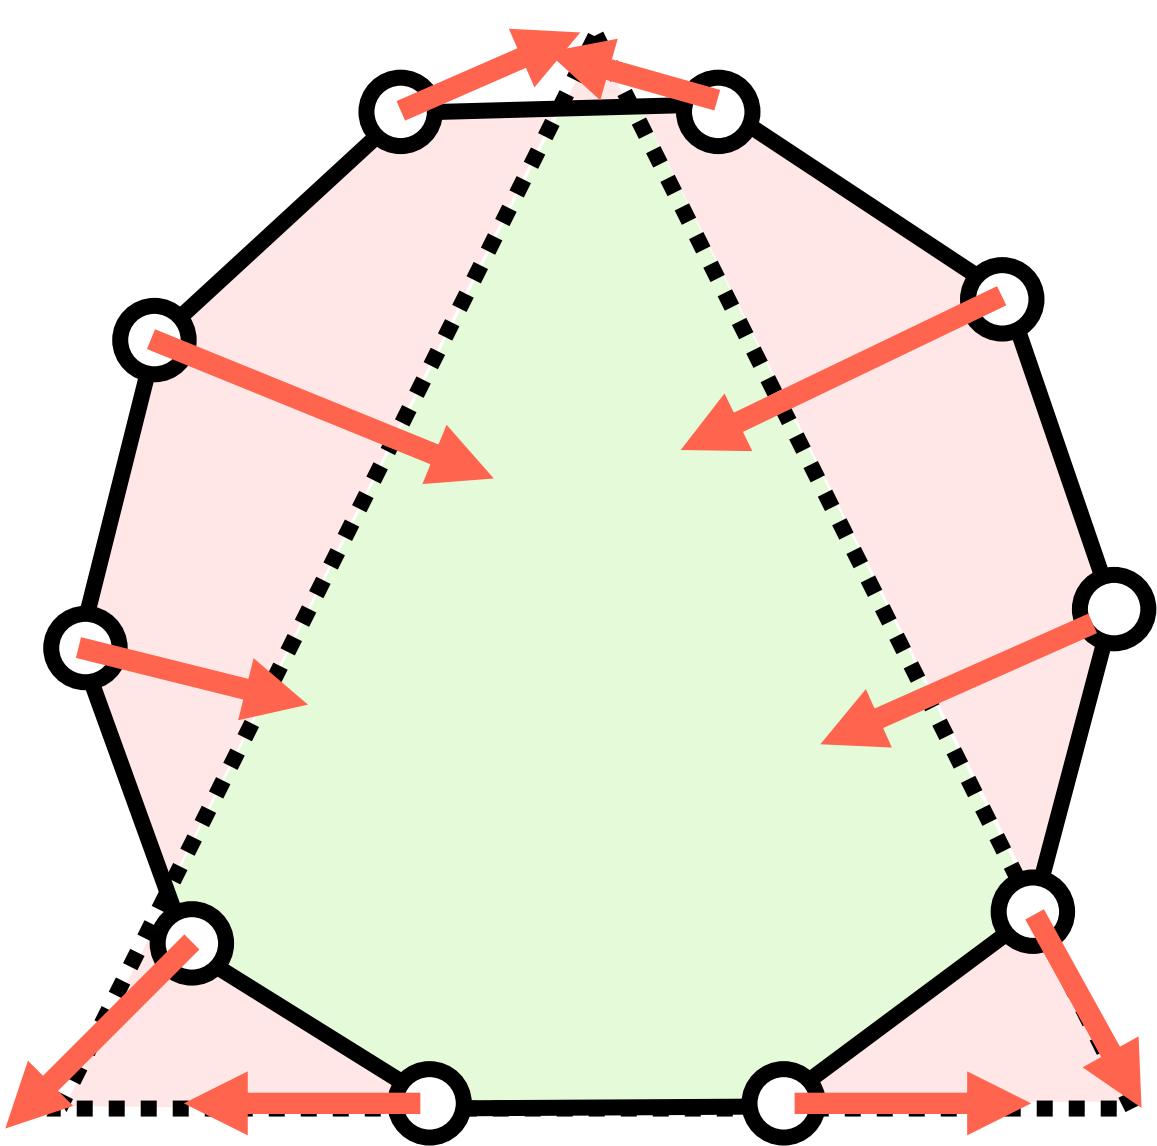


Compute Gradients

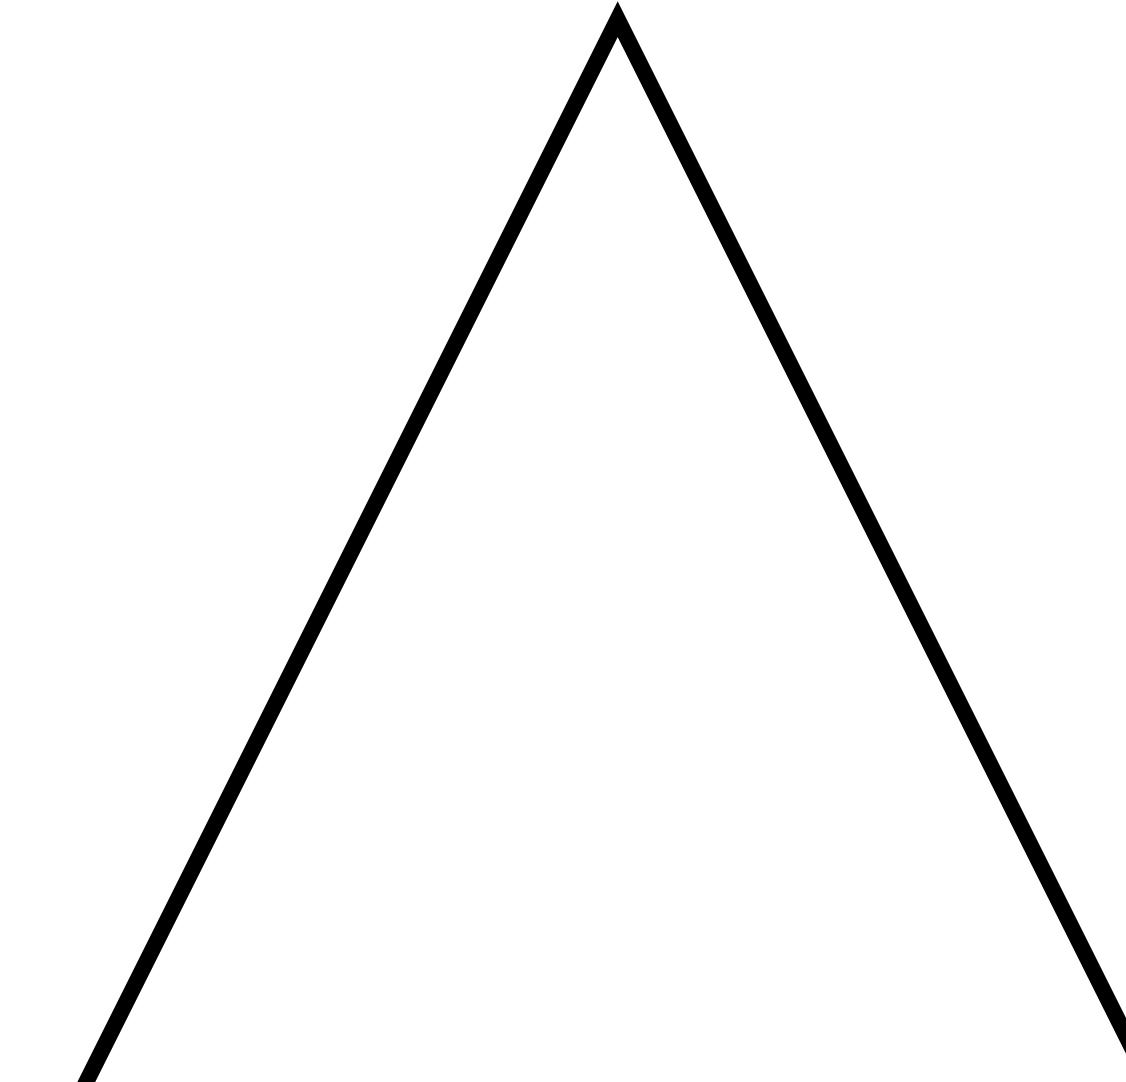


Target Geometry

# Gradient Based Optimization

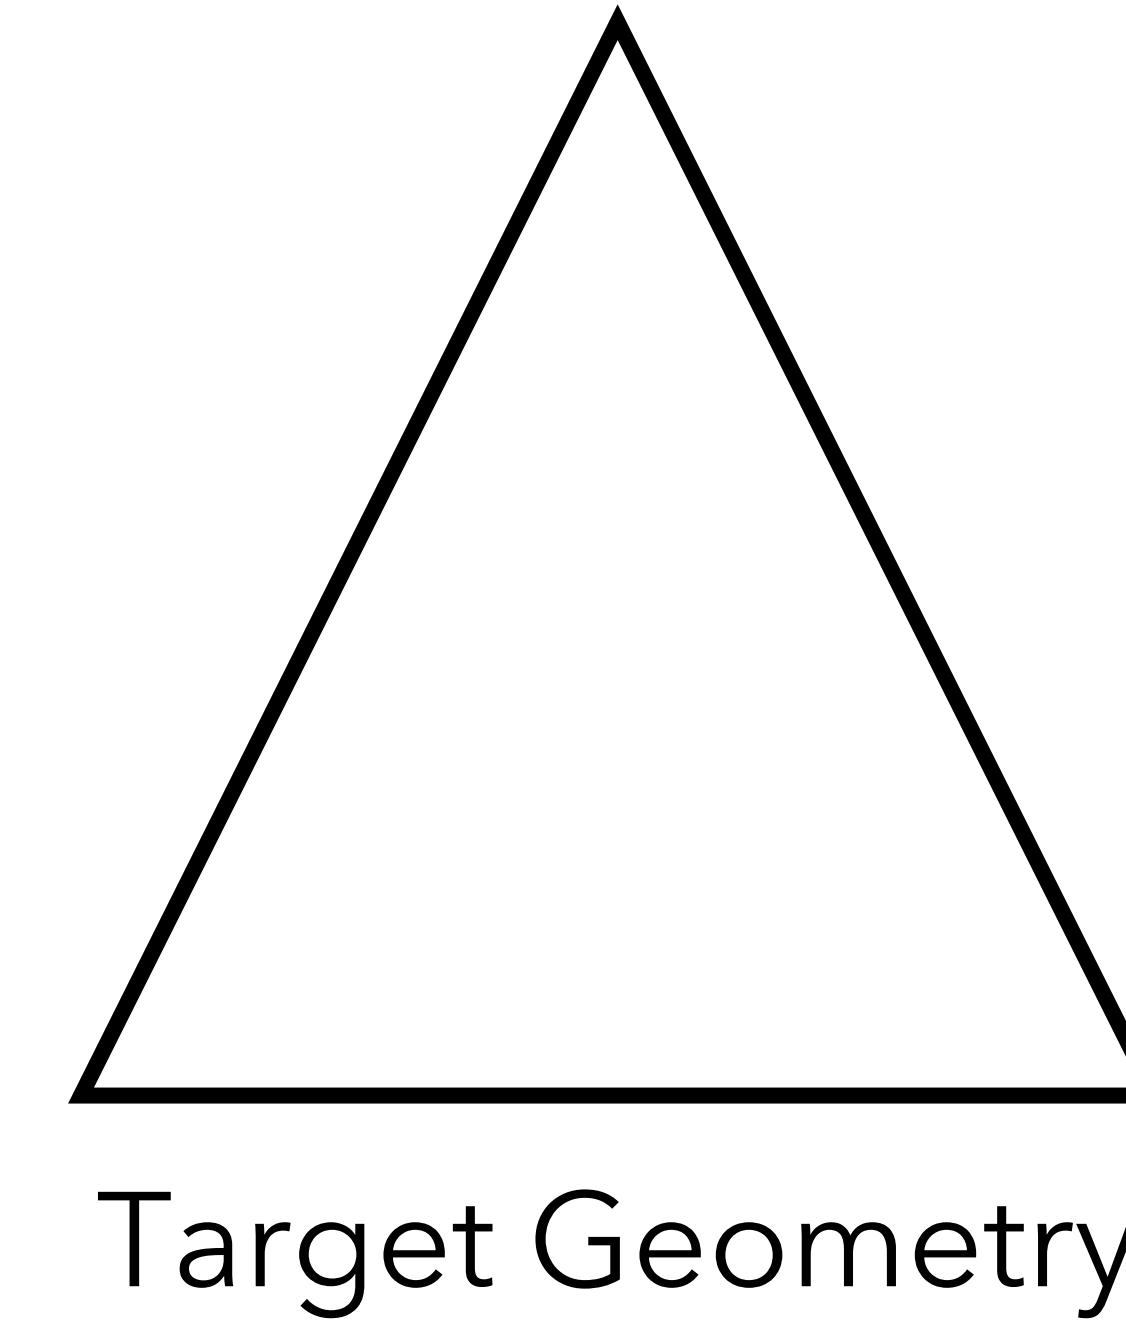
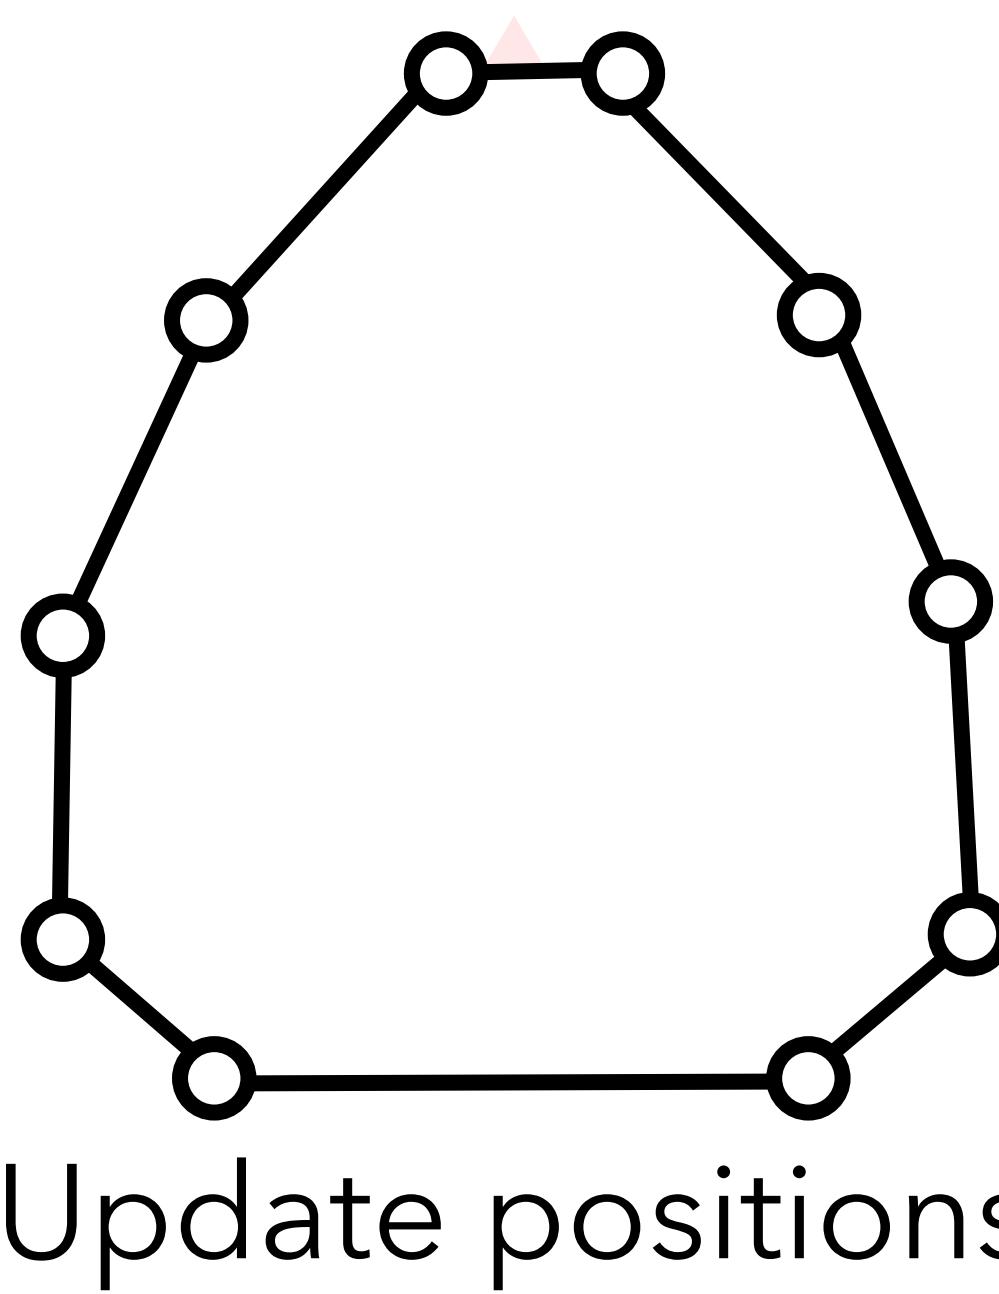


Compute Gradients

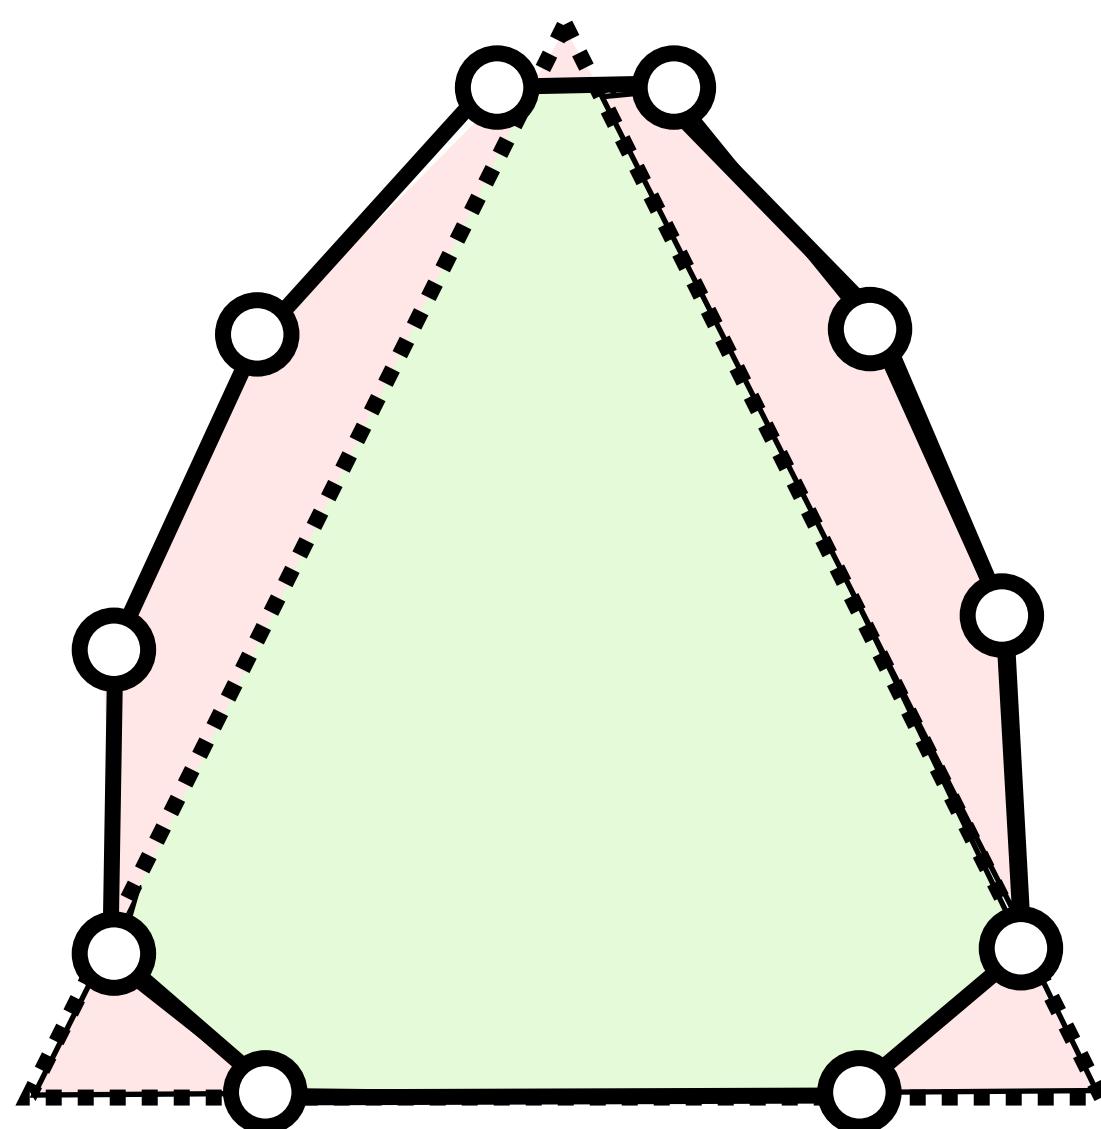


Target Geometry

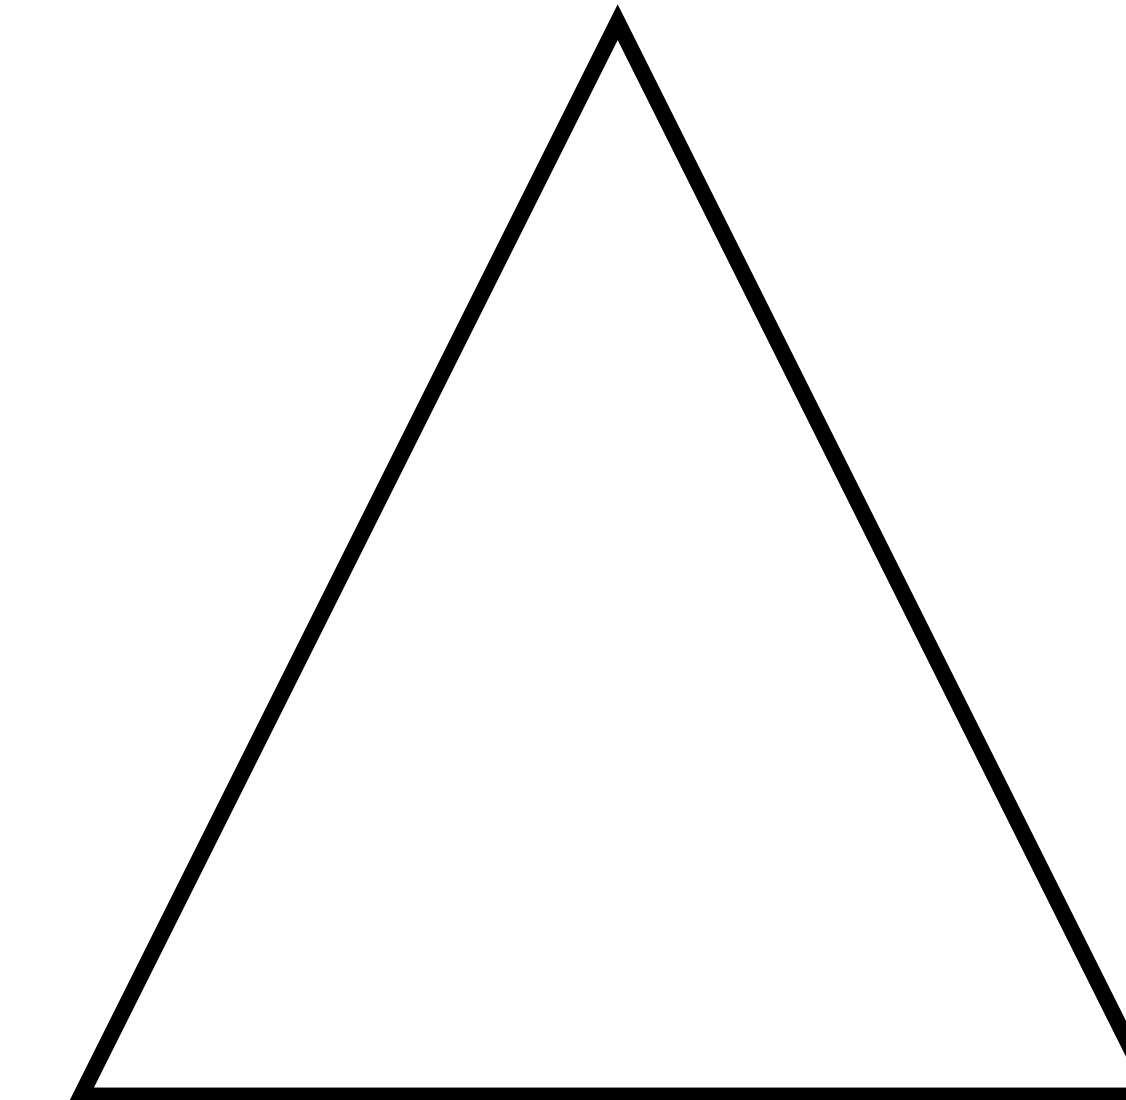
# Gradient Based Optimization



# Gradient Based Optimization

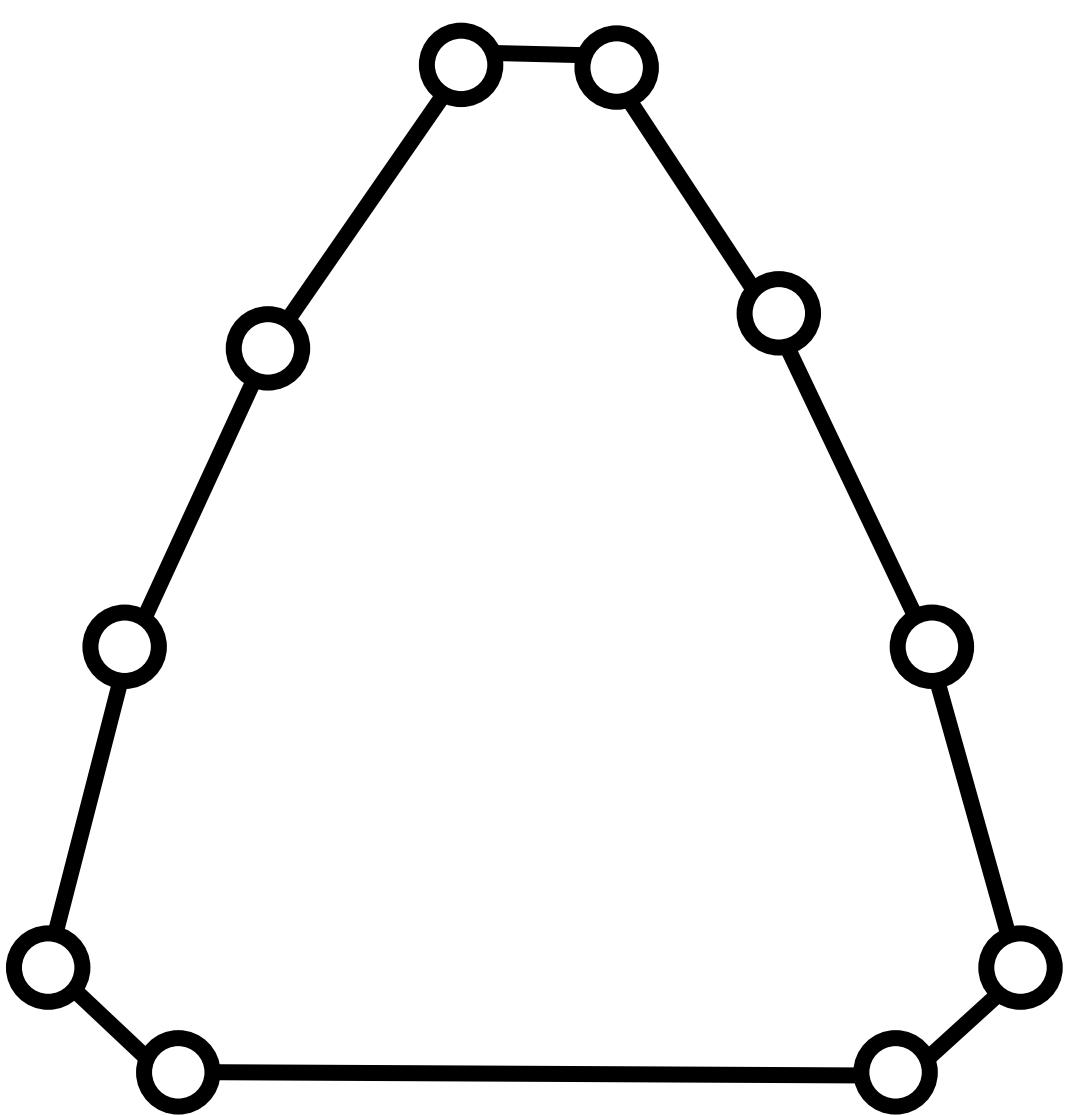


Compute New Error

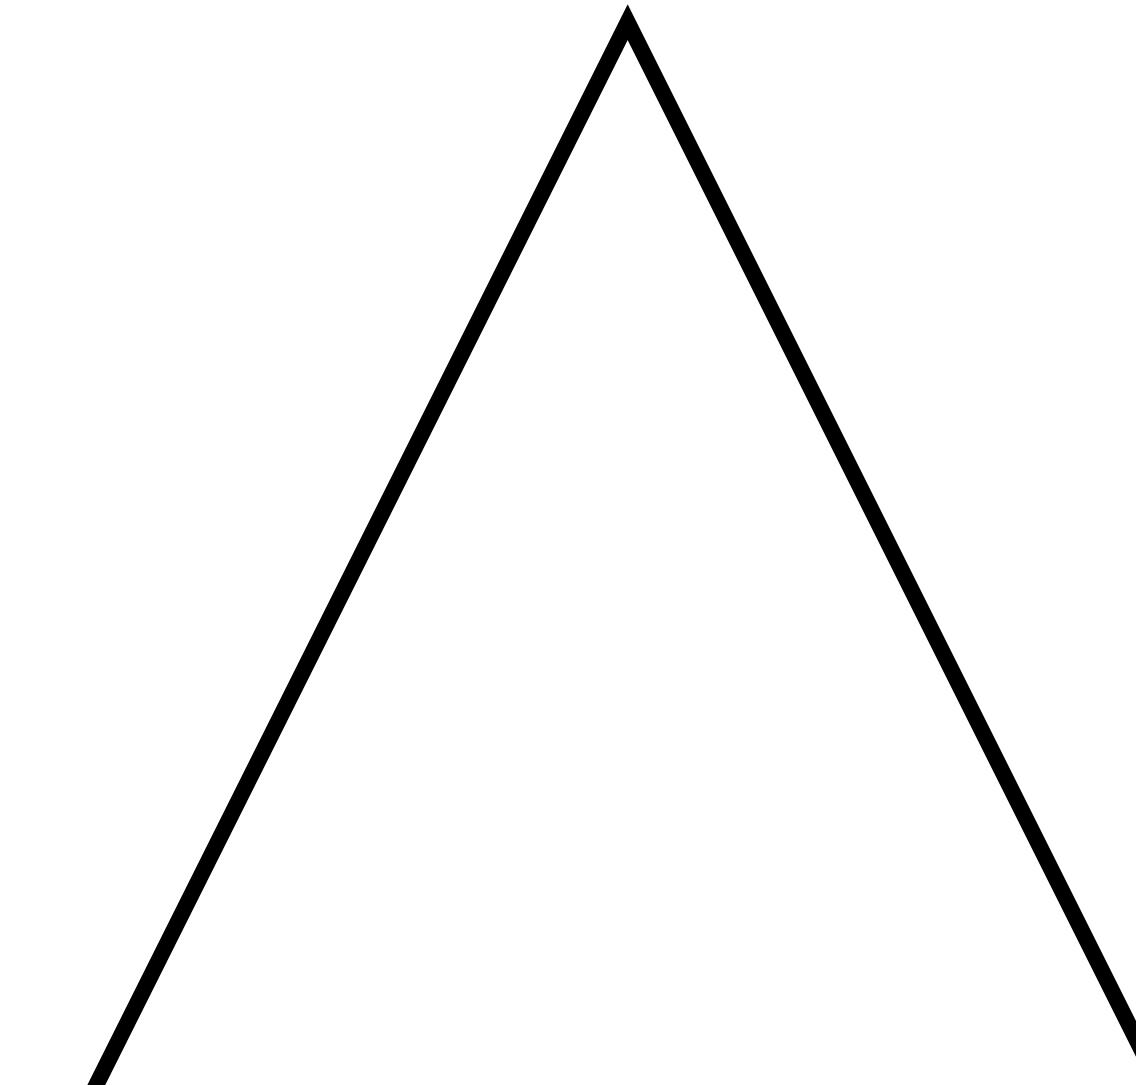


Target Geometry

# Gradient Based Optimization

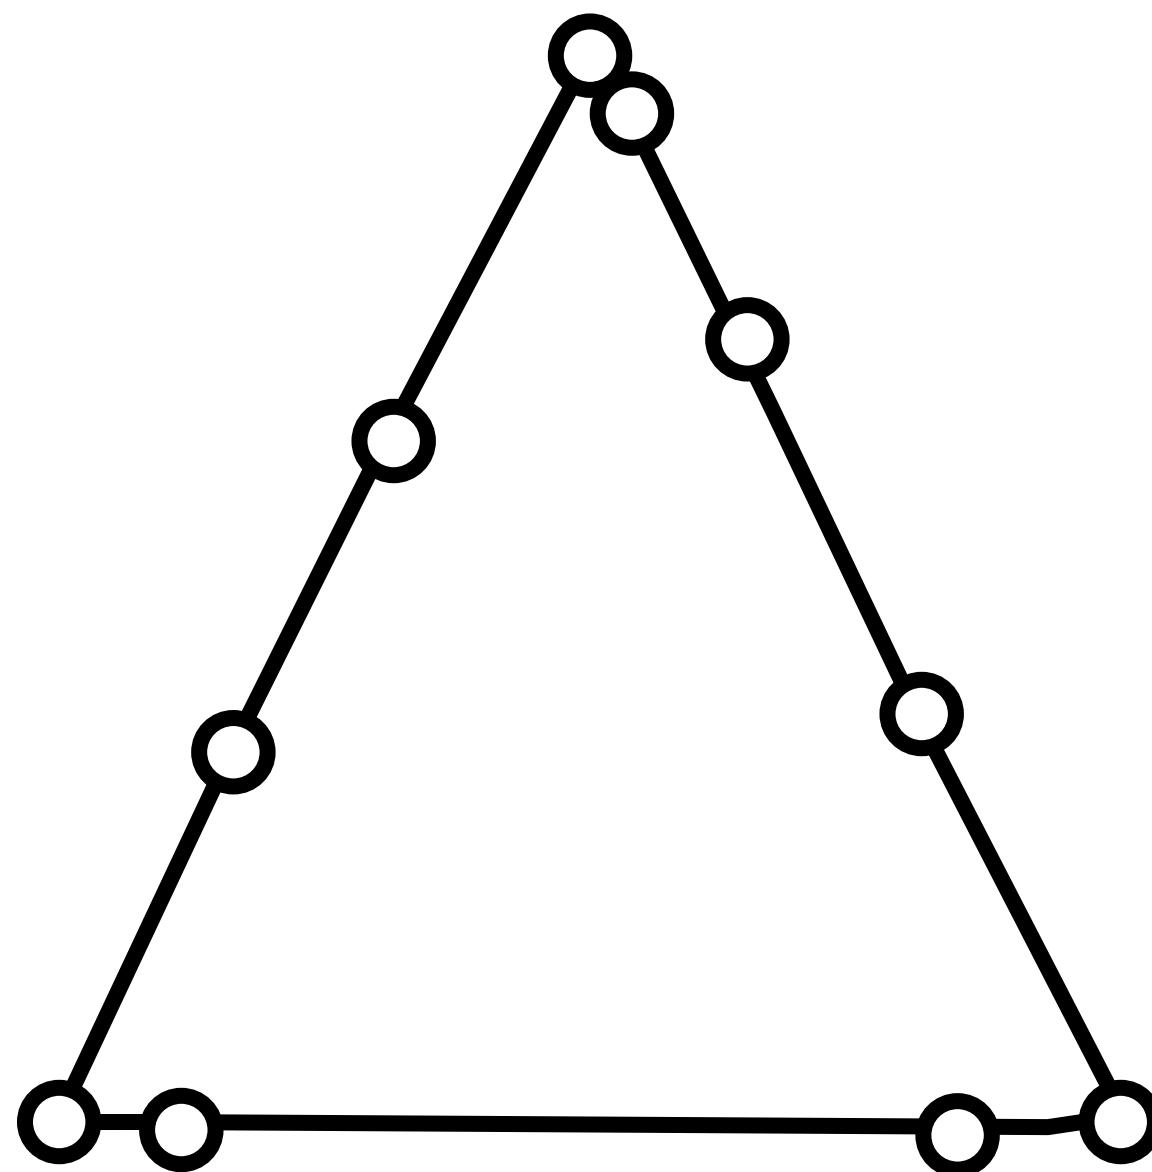


Repeat

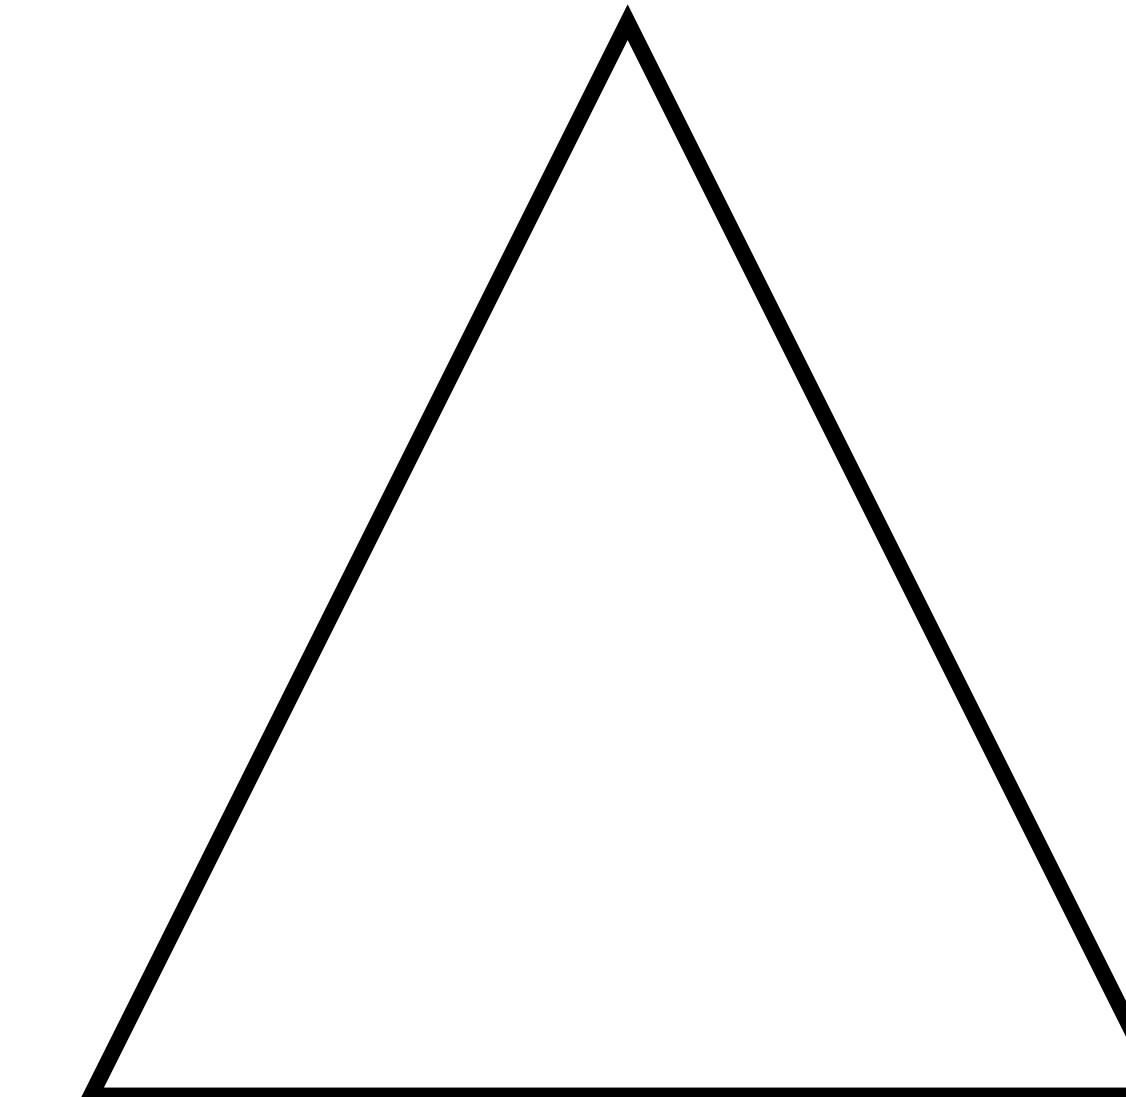


Target Geometry

# Gradient Based Optimization

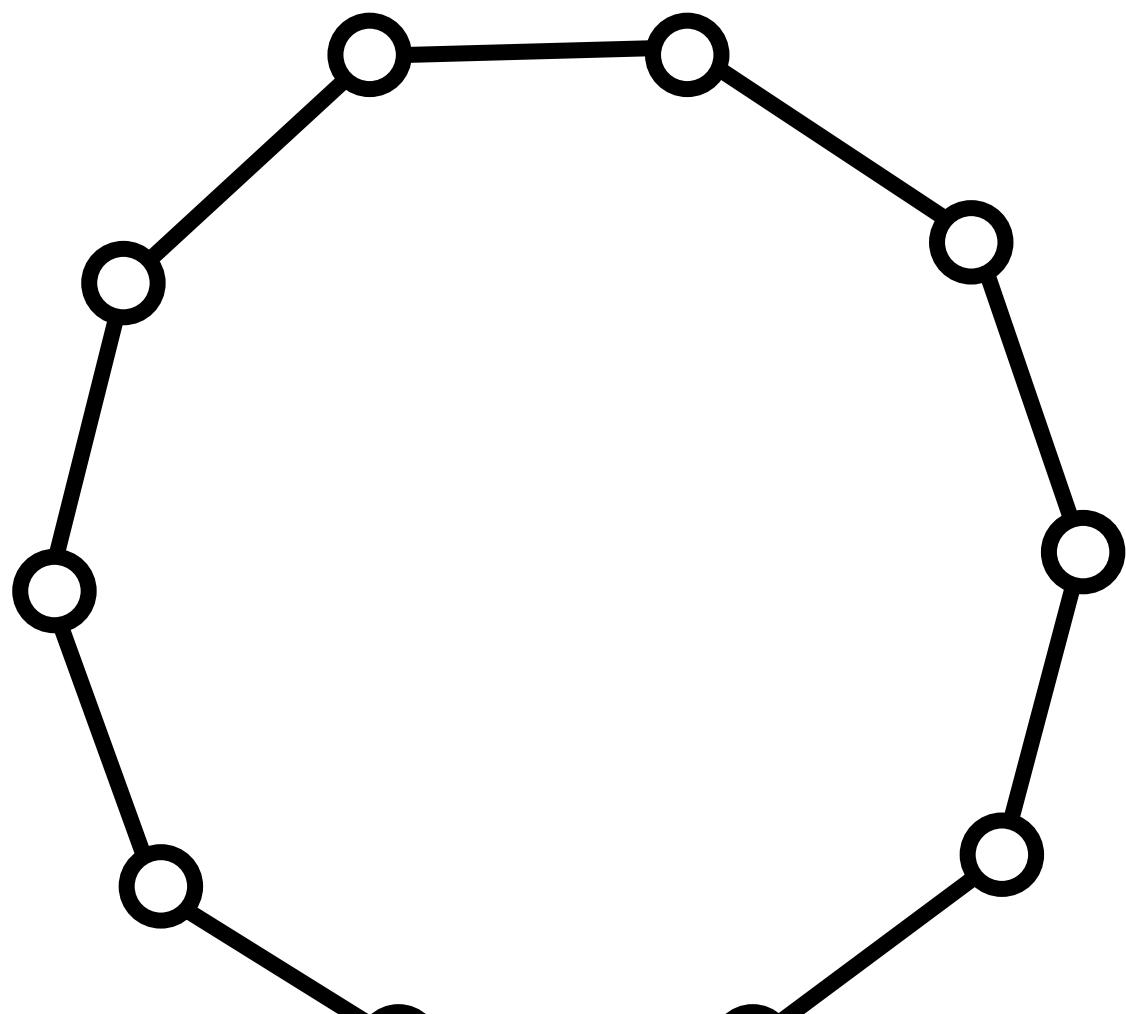


Repeat

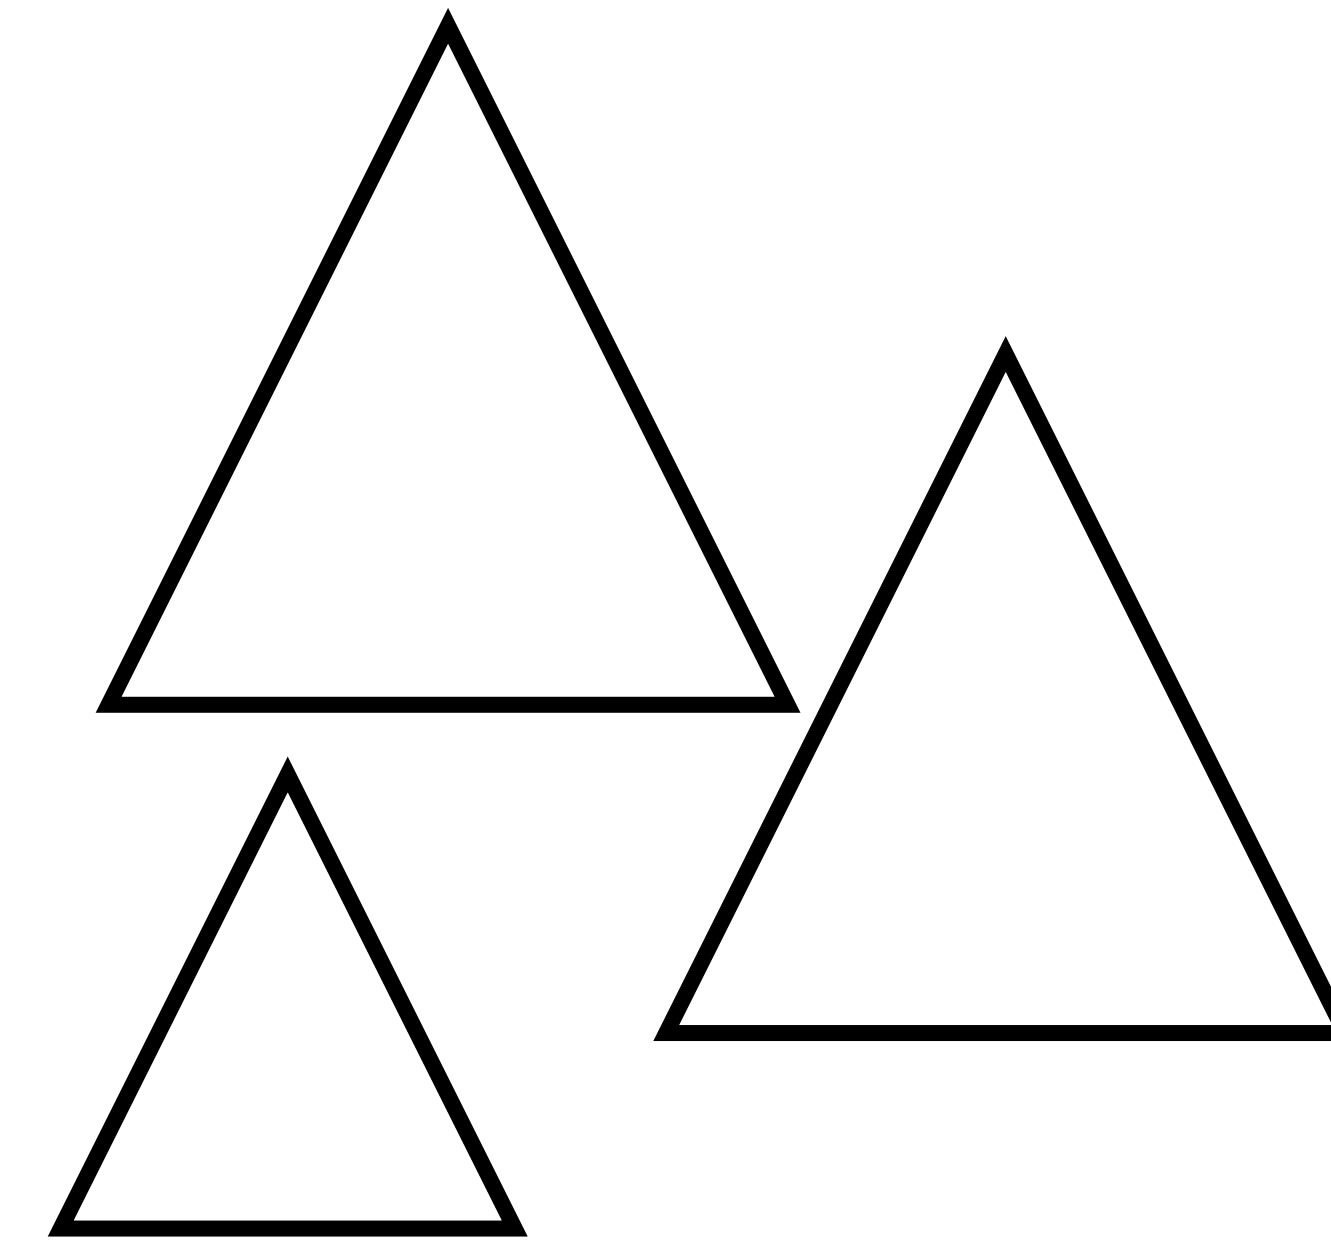


Target Geometry

# Gradient Based Optimization

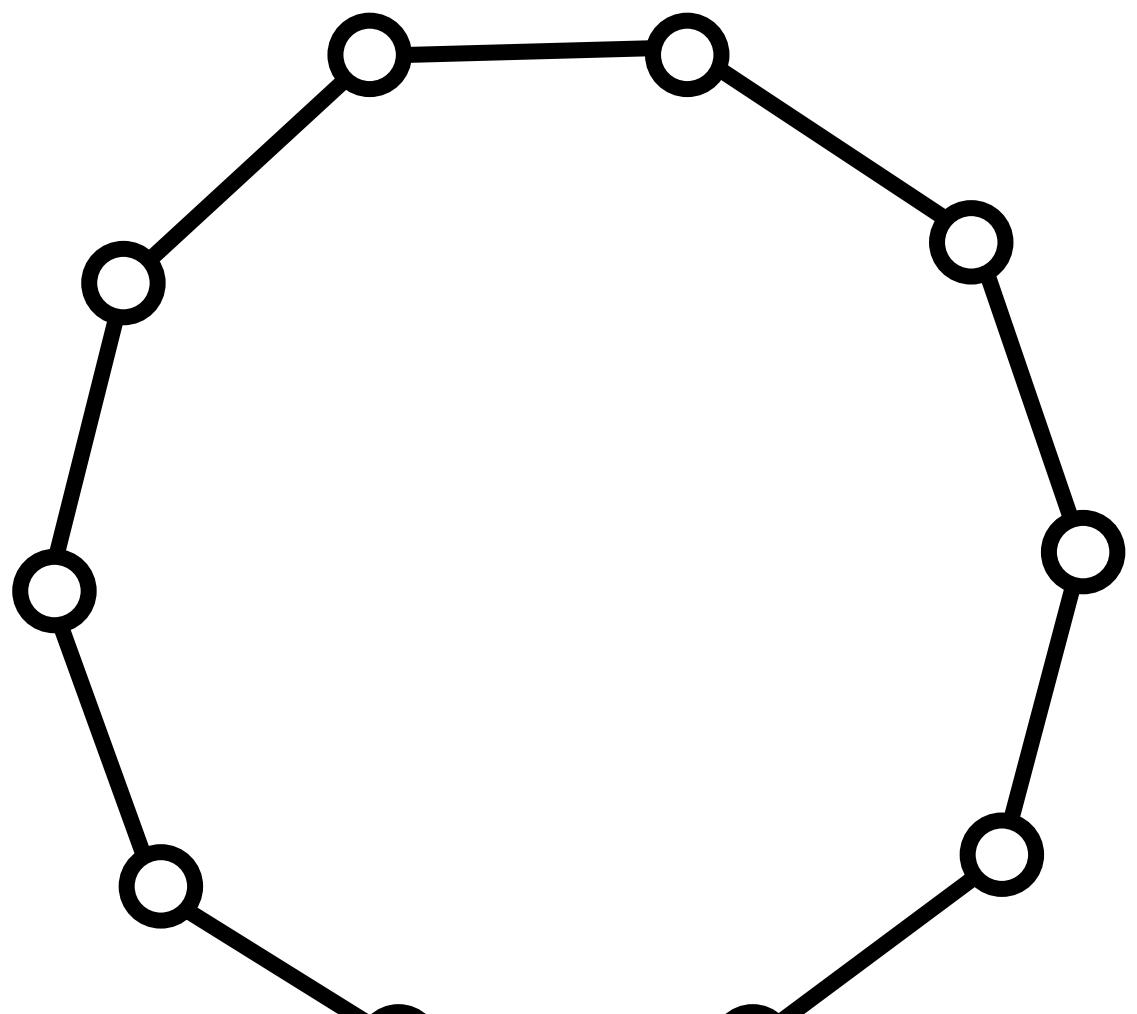


Initial Geometry

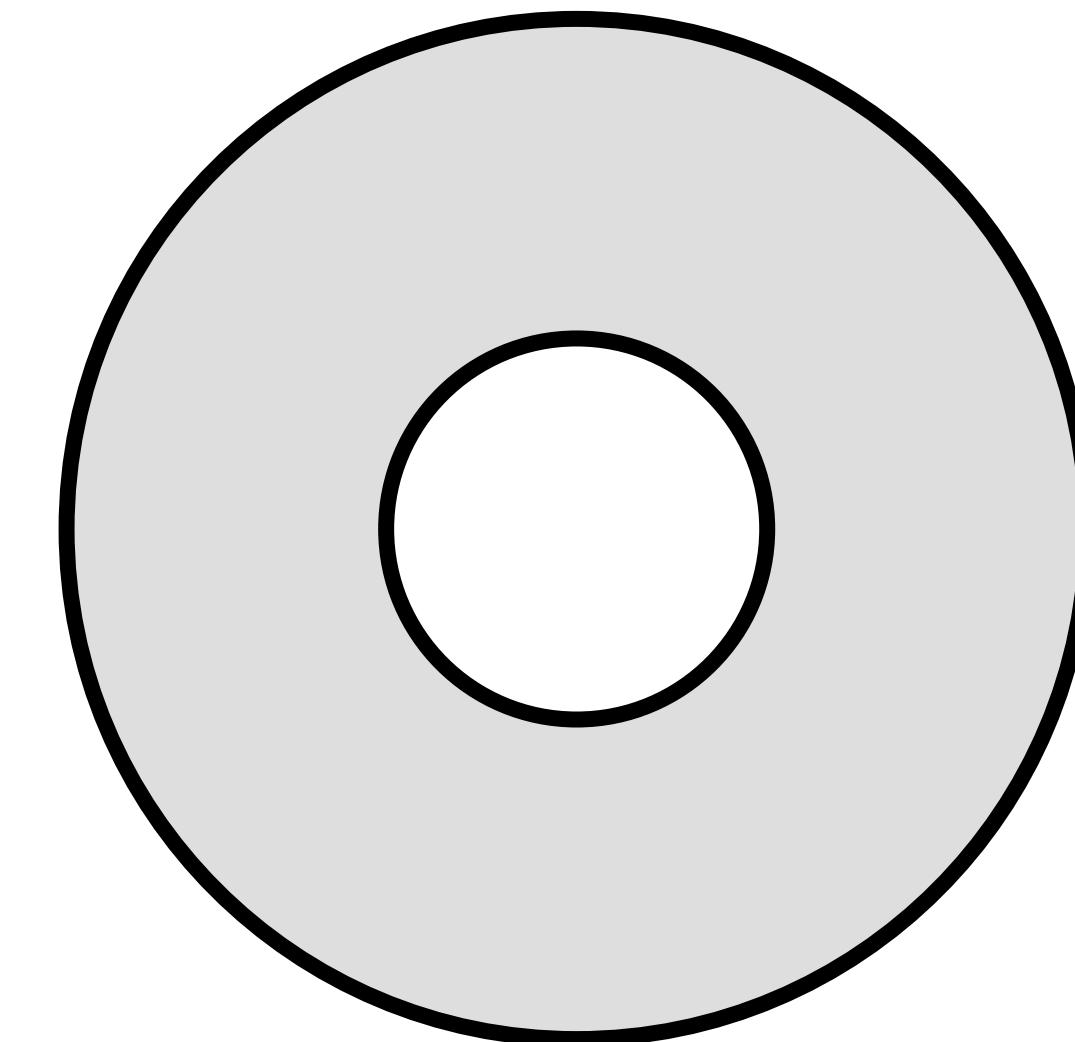


Target Geometry

# Gradient Based Optimization

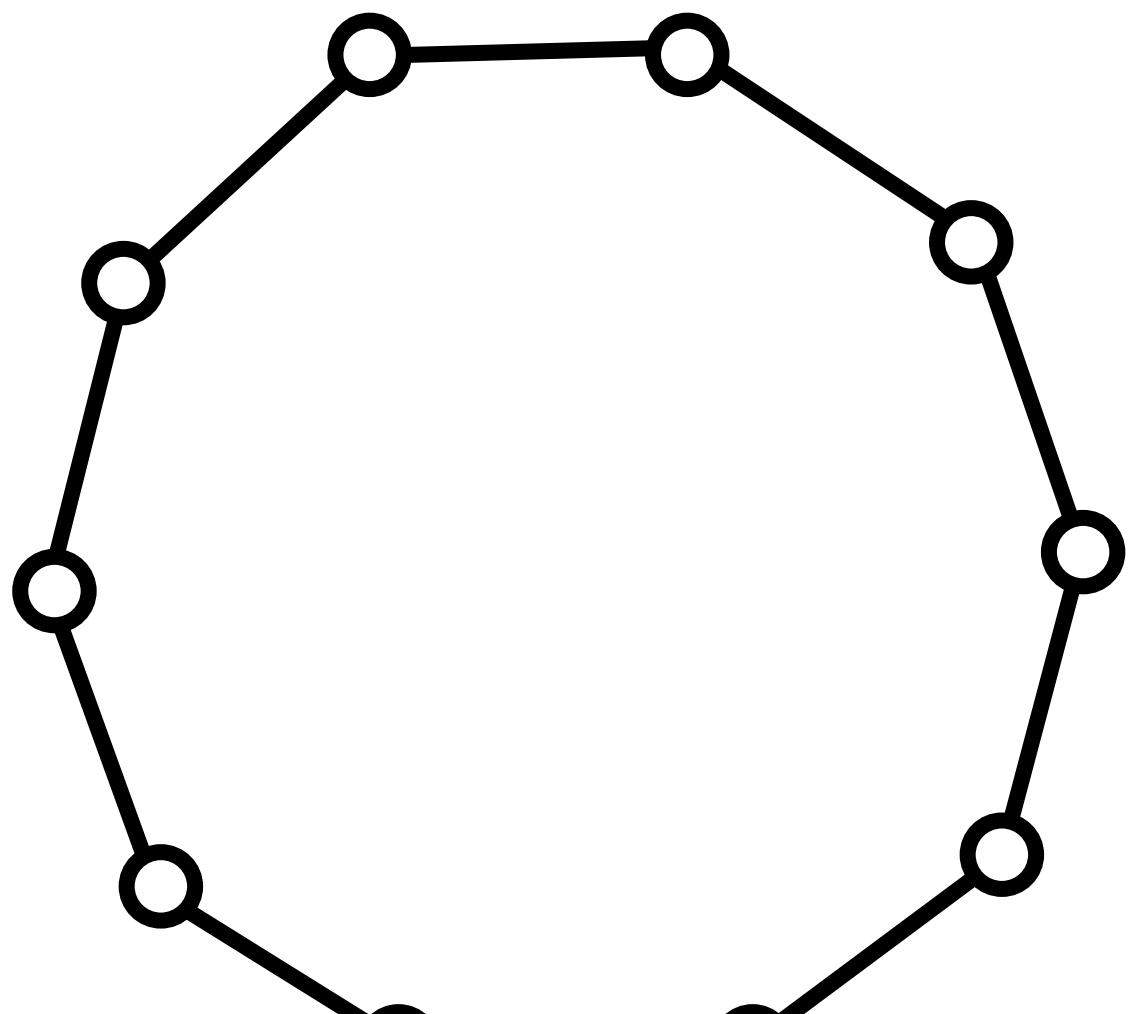


Initial Geometry

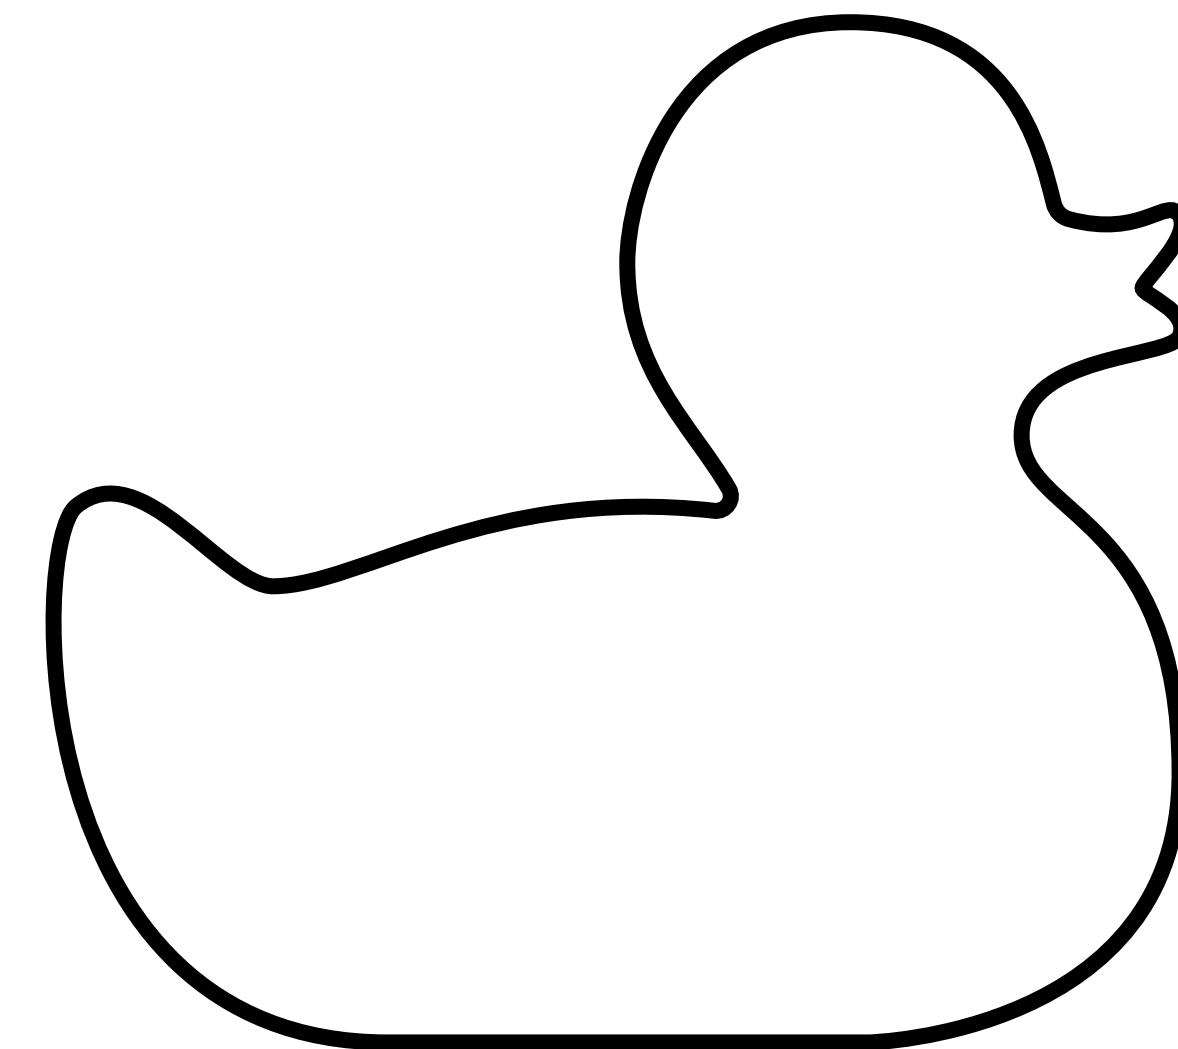


Target Geometry

# Gradient Based Optimization

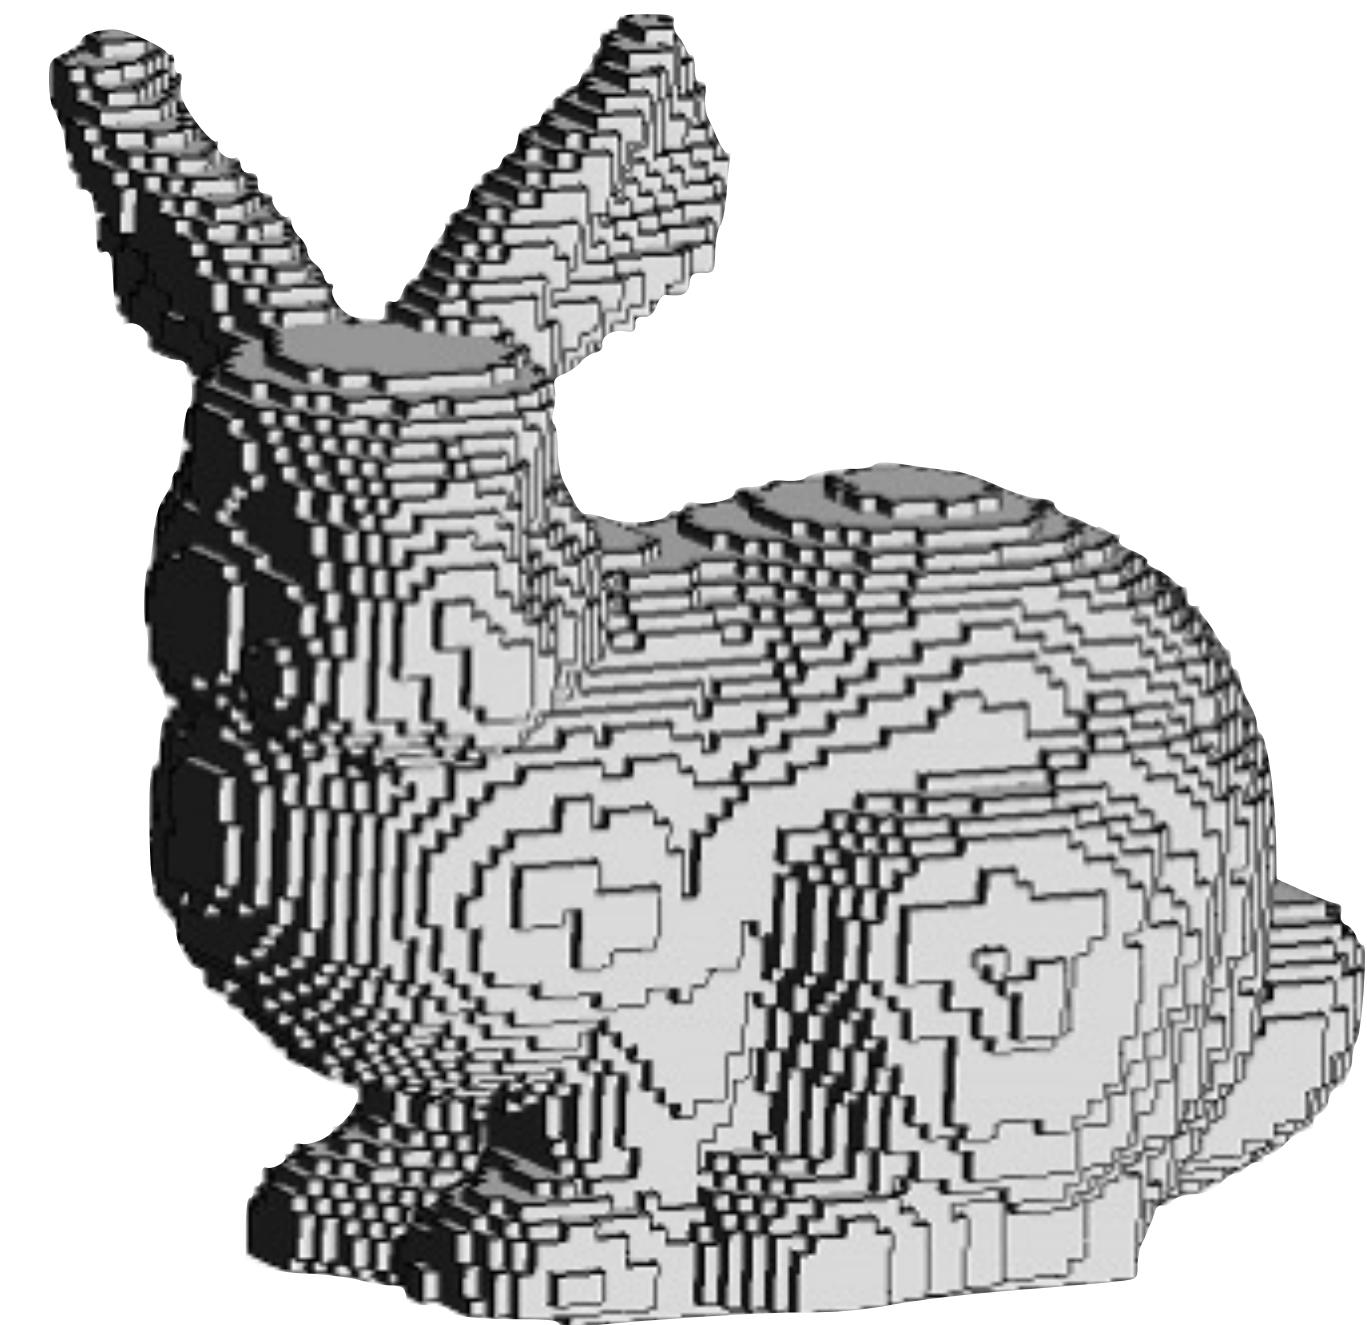


Initial Geometry

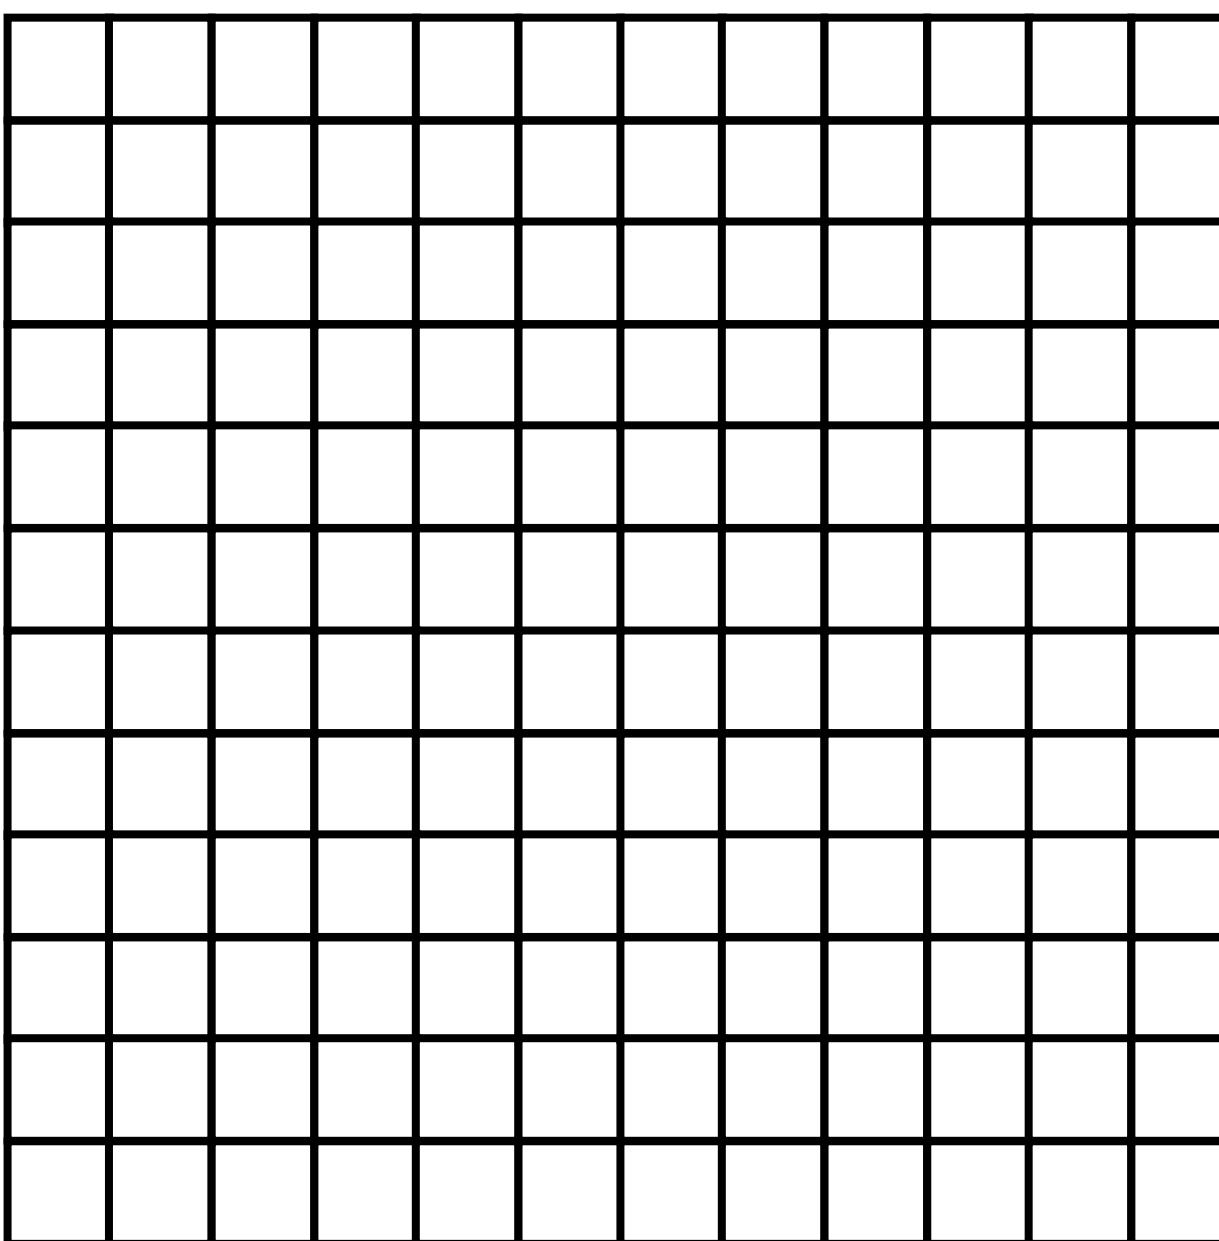


Target Geometry

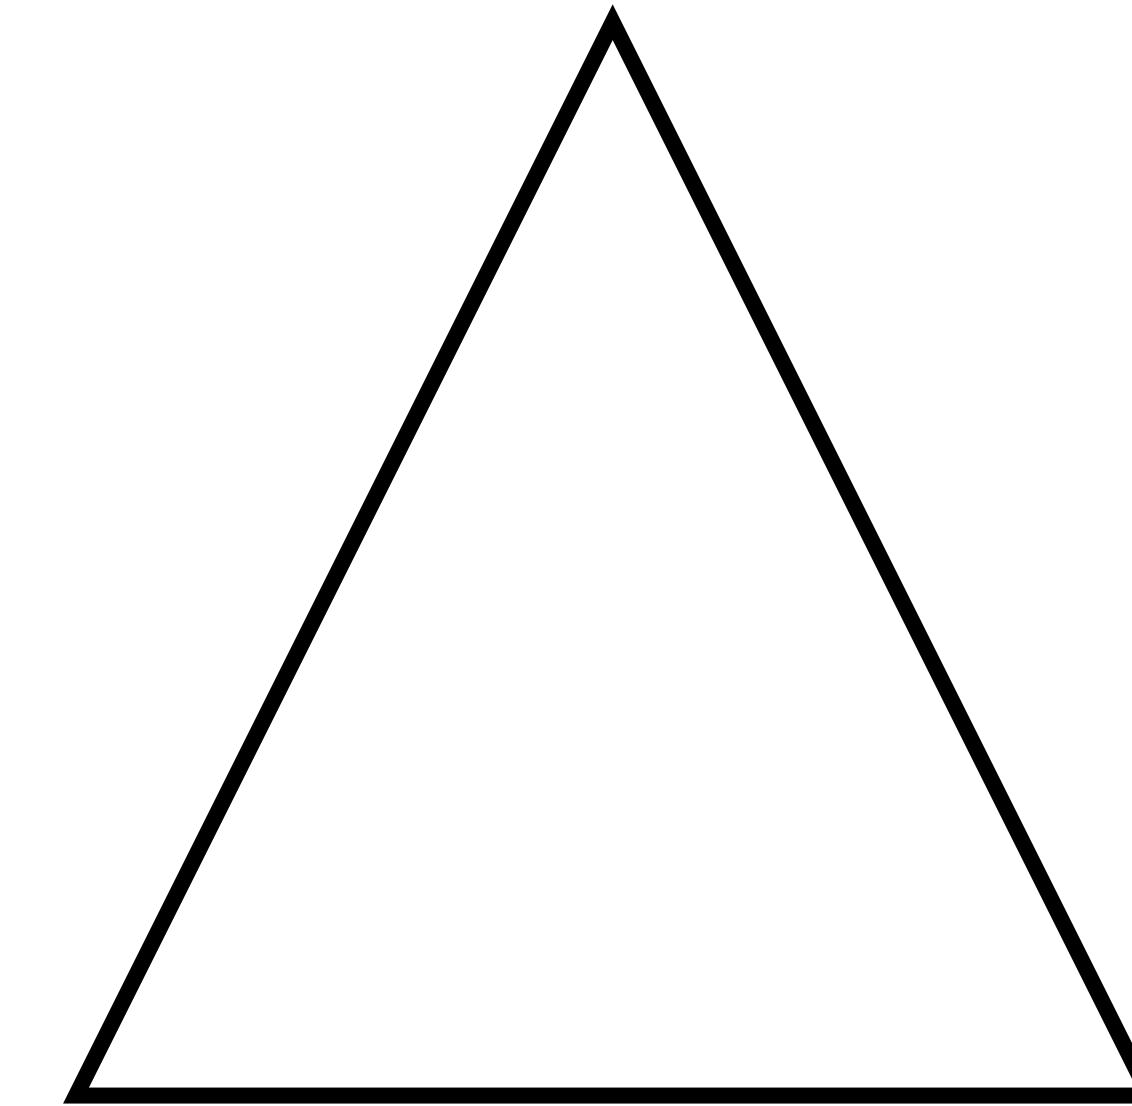
# Voxel Representation



# Gradient Based Optimization

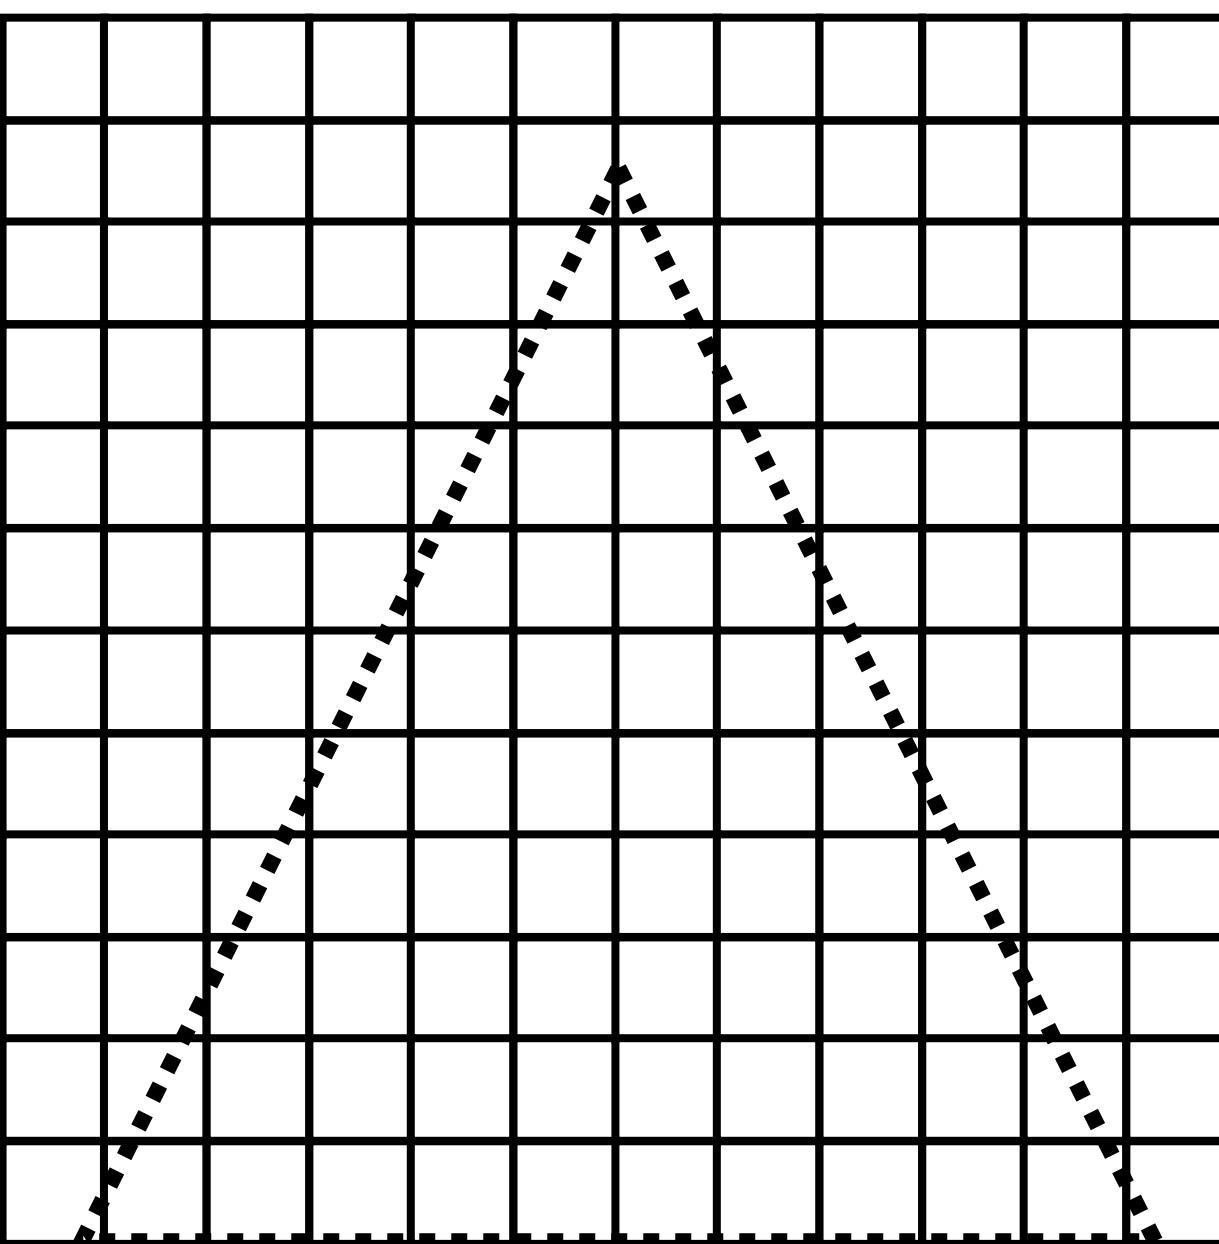


Initialized Grid

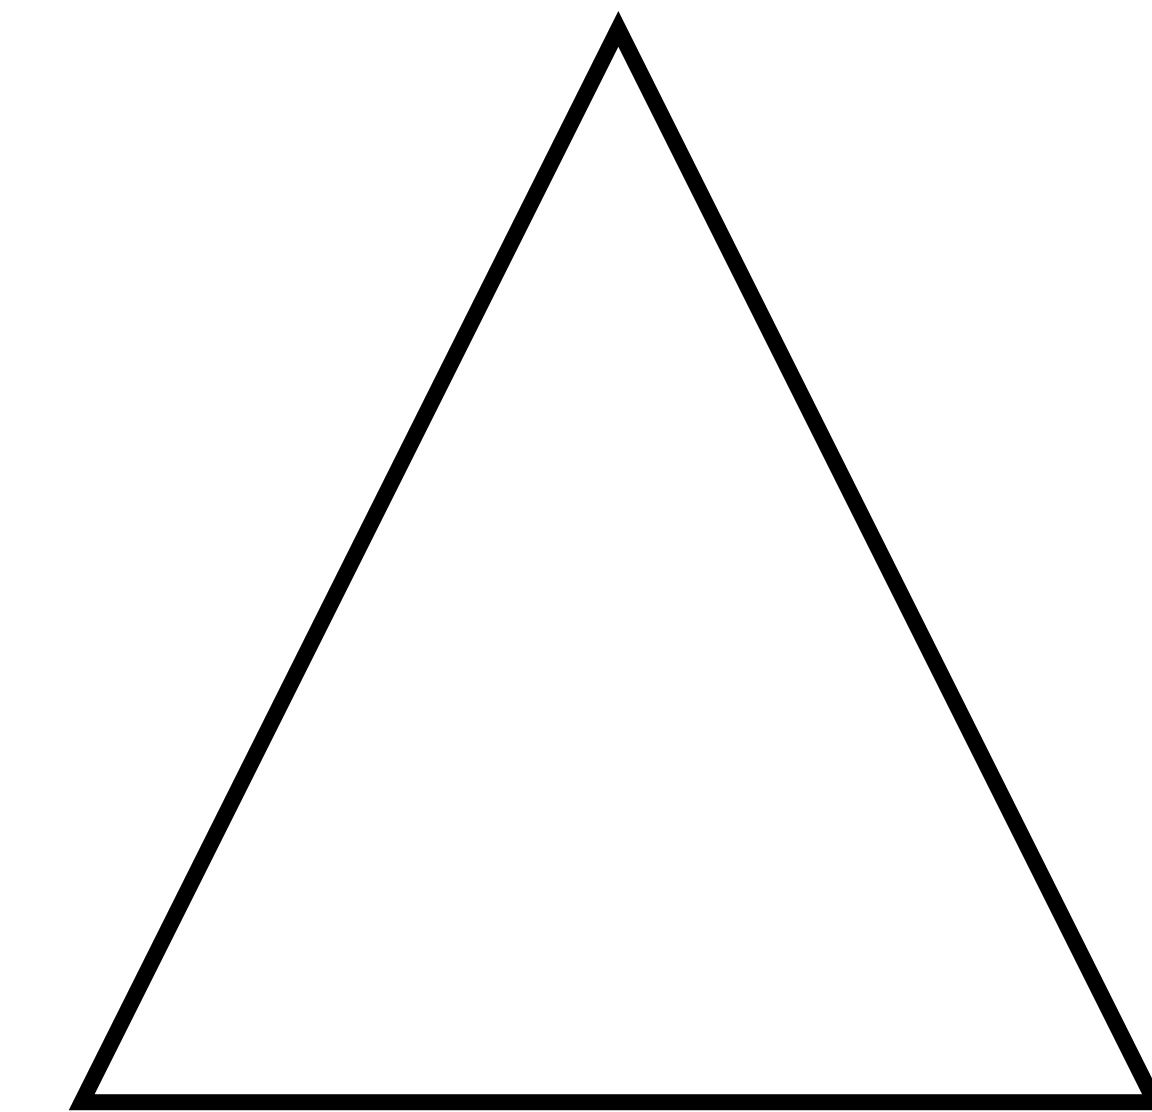


Target Geometry

# Gradient Based Optimization

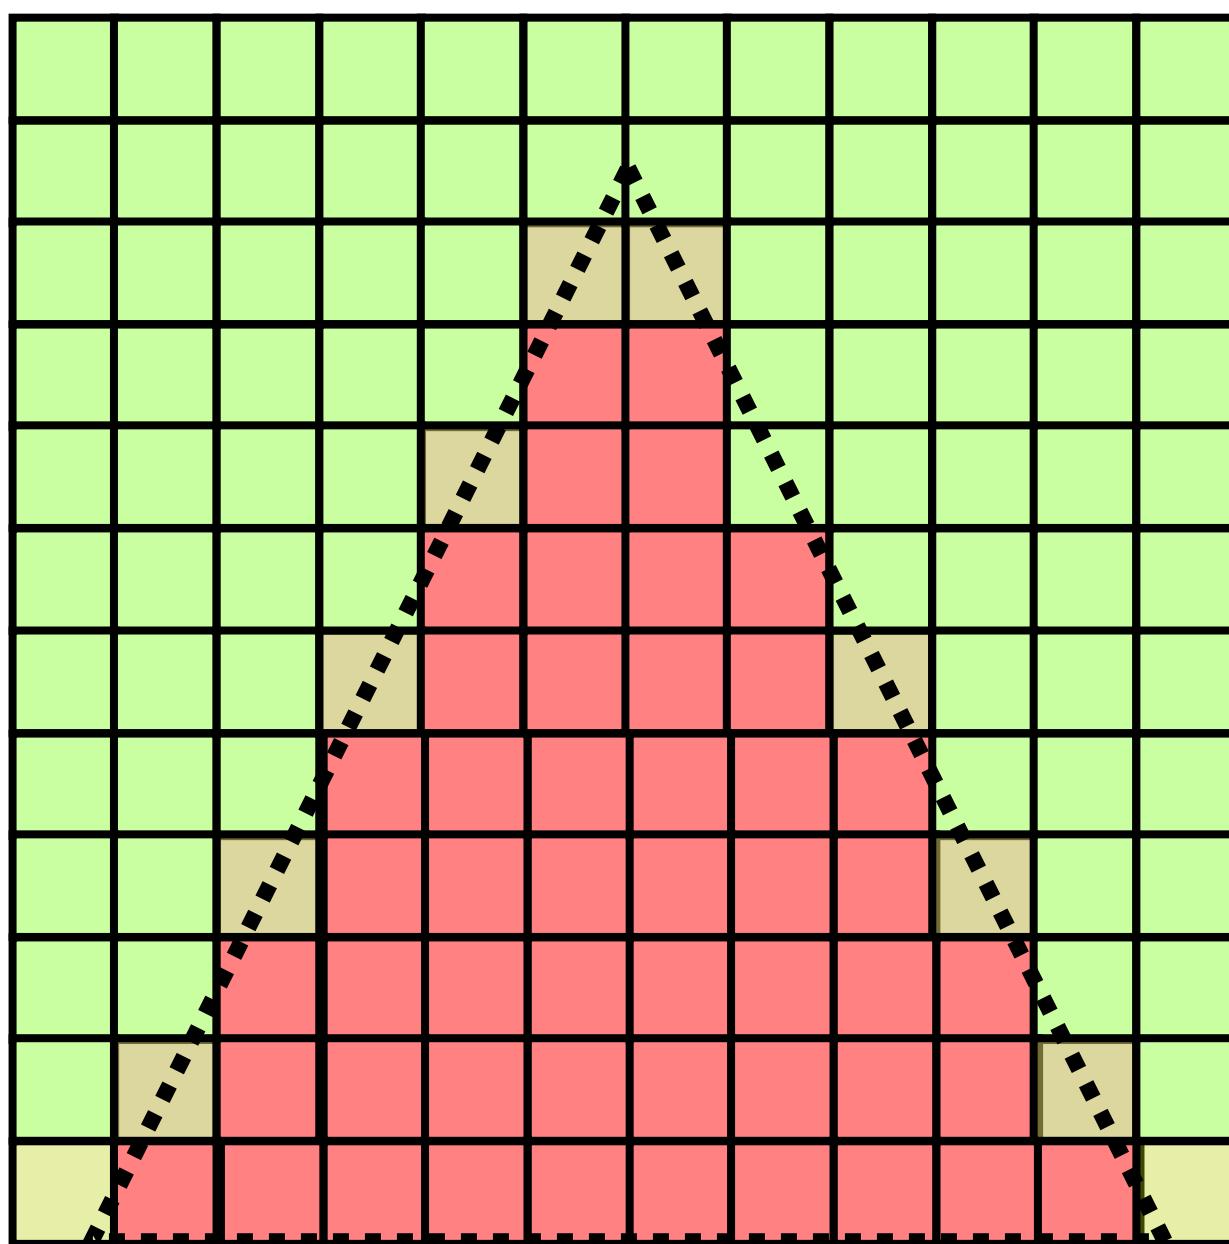


Initialized Grid

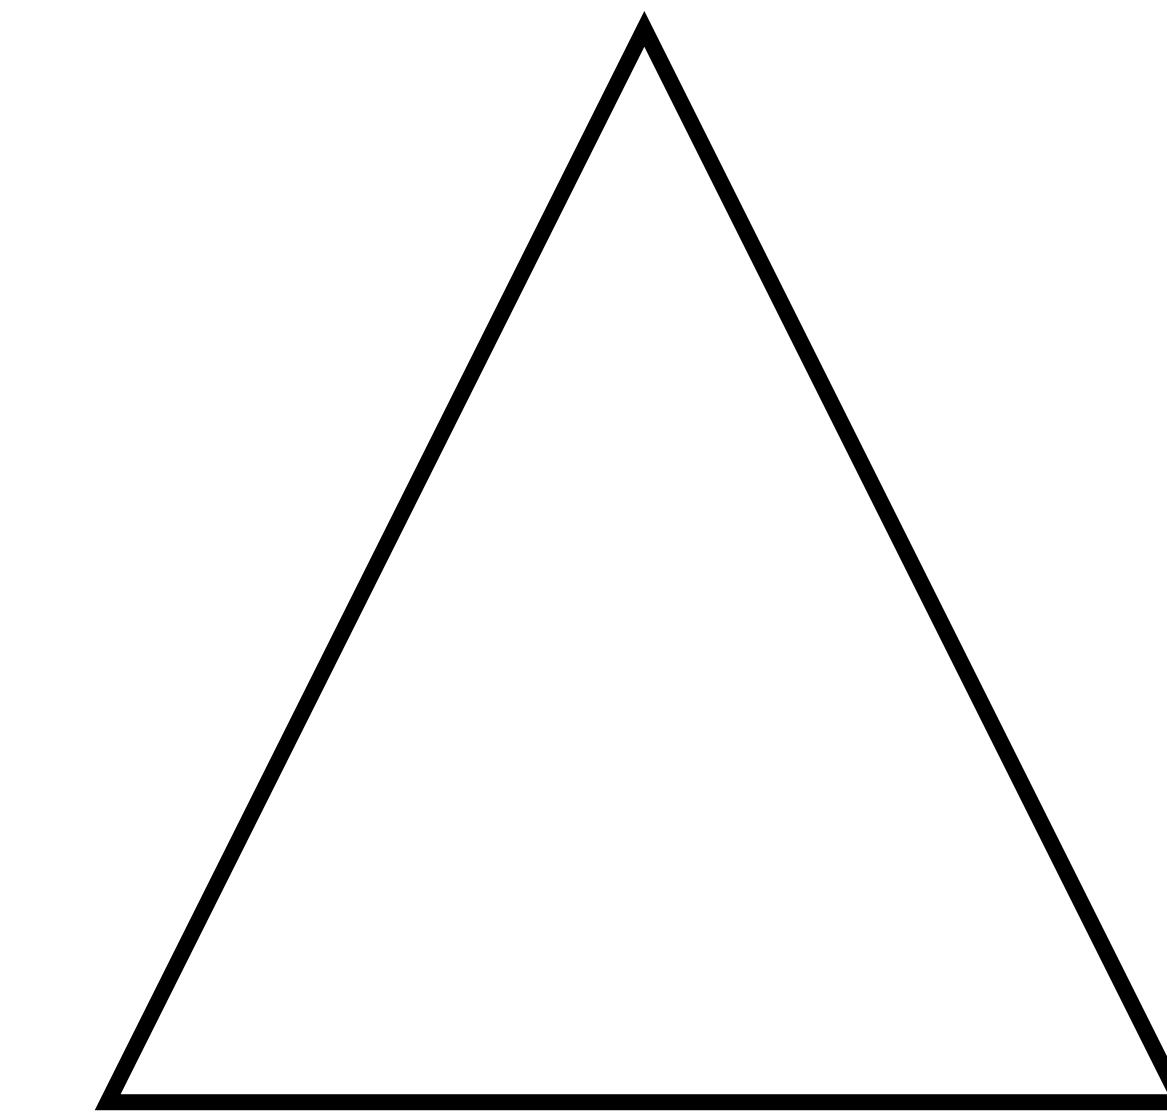


Target Geometry

# Gradient Based Optimization

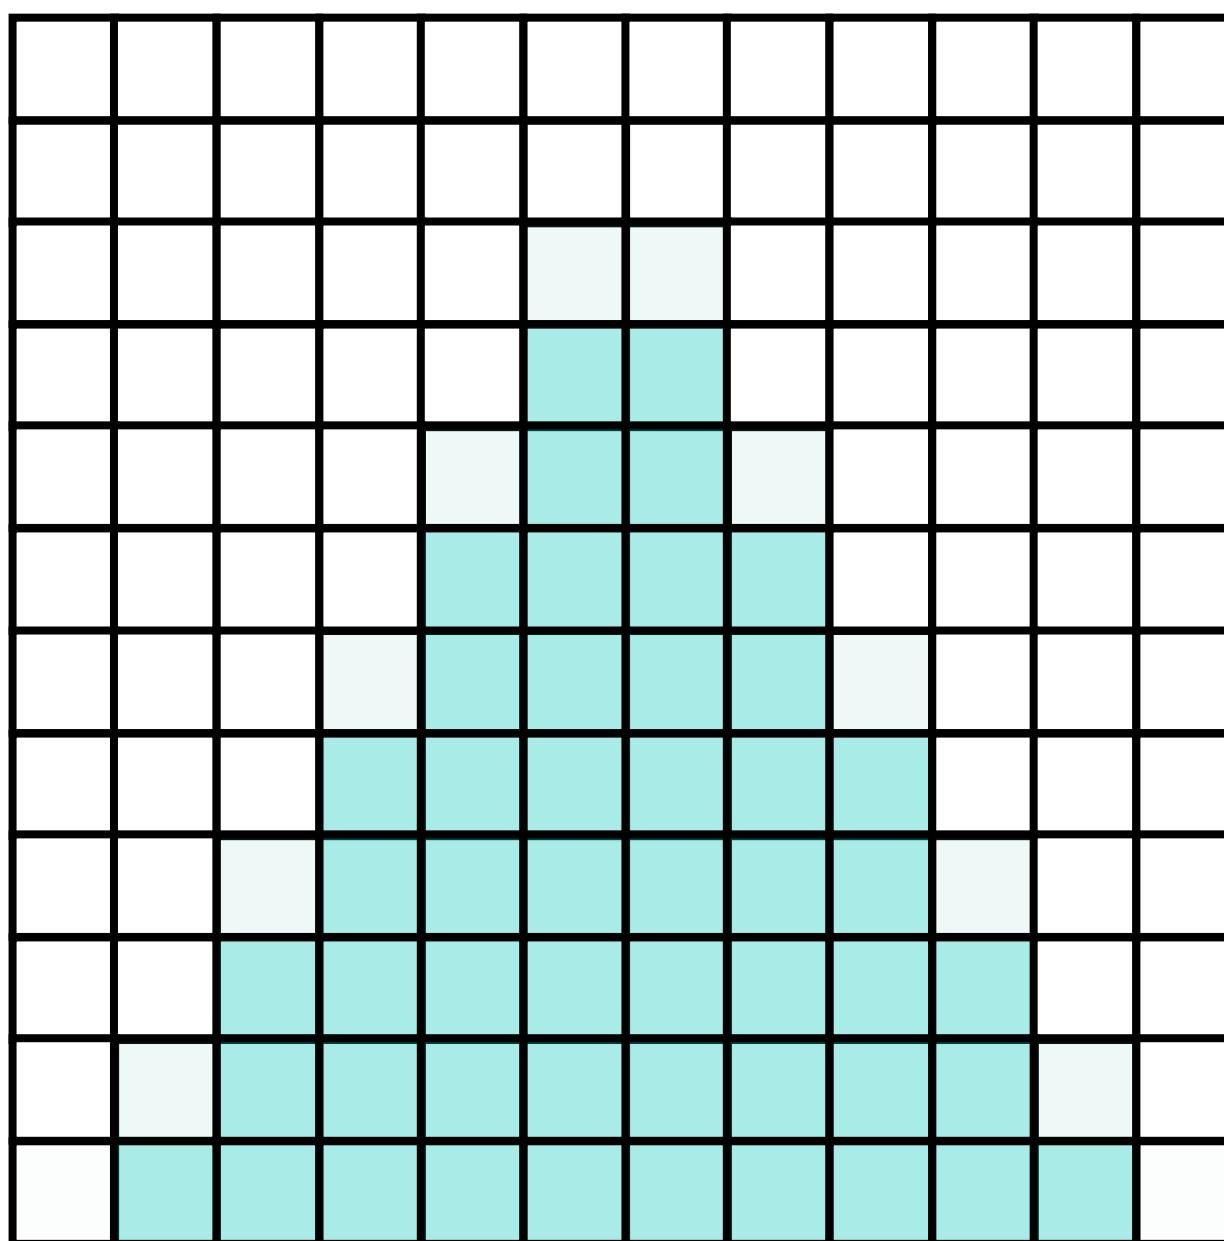


Loss

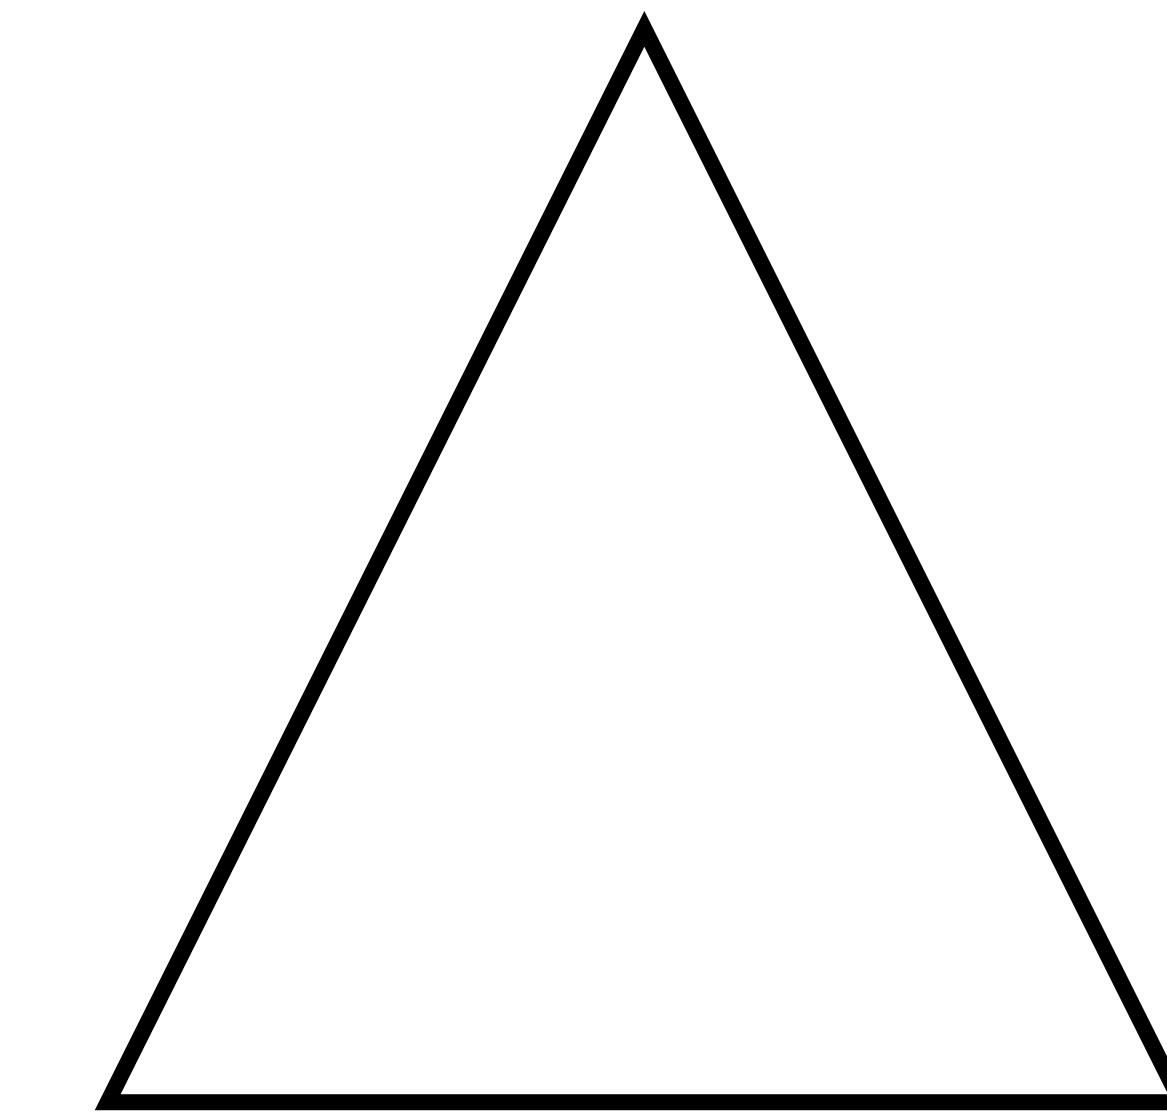


Target Geometry

# Gradient Based Optimization

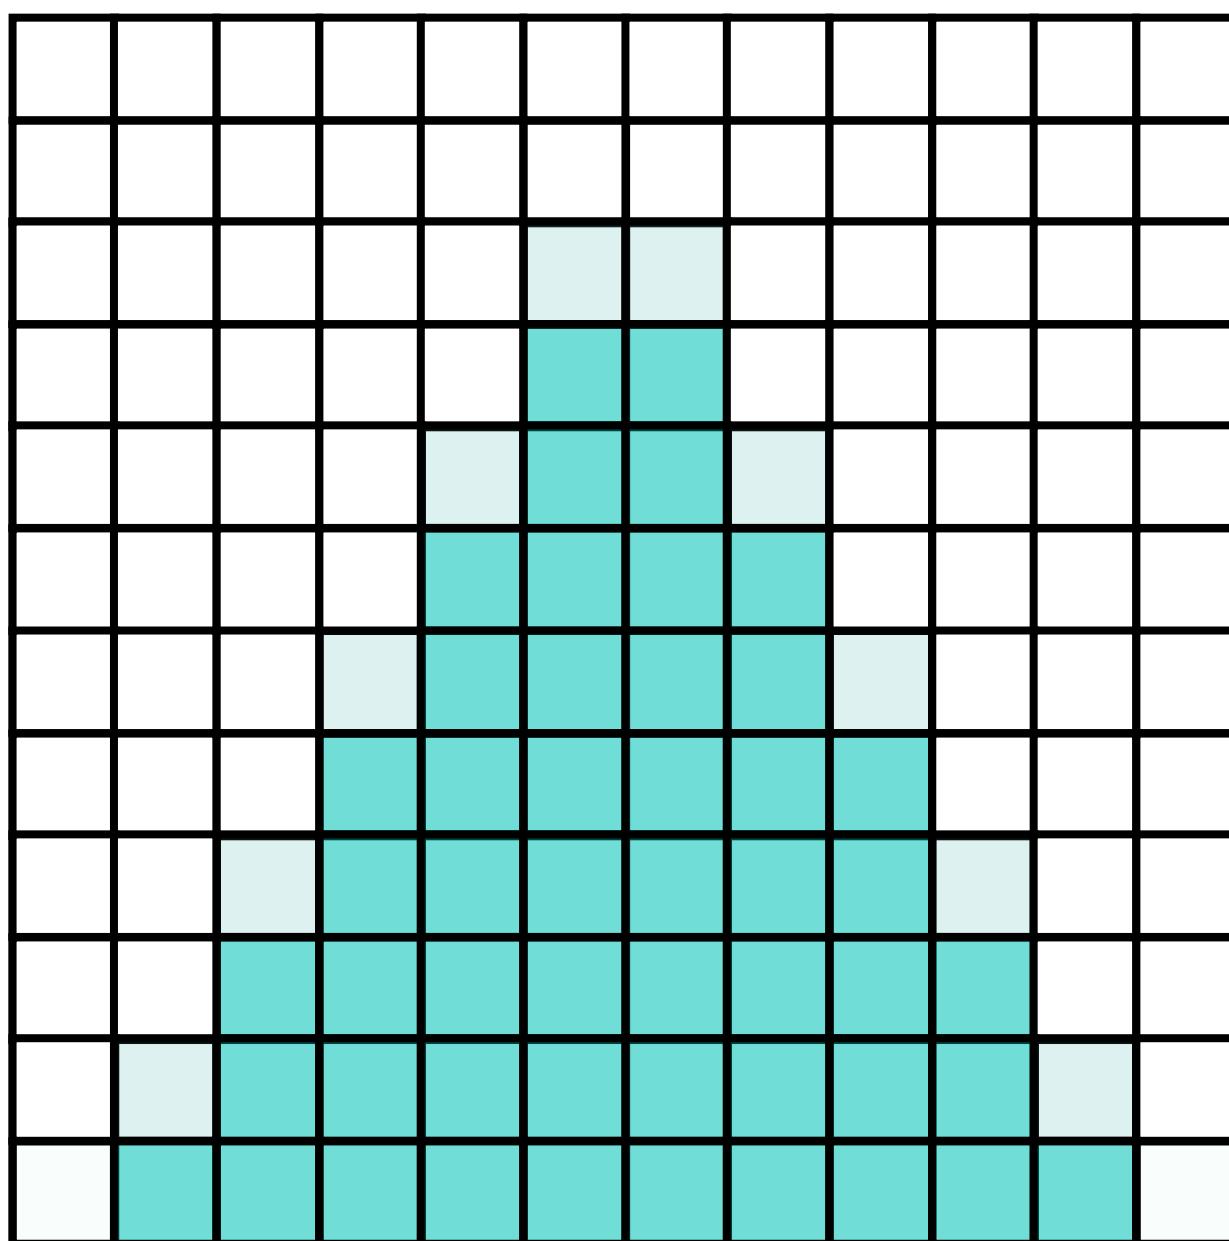


Gradient Step

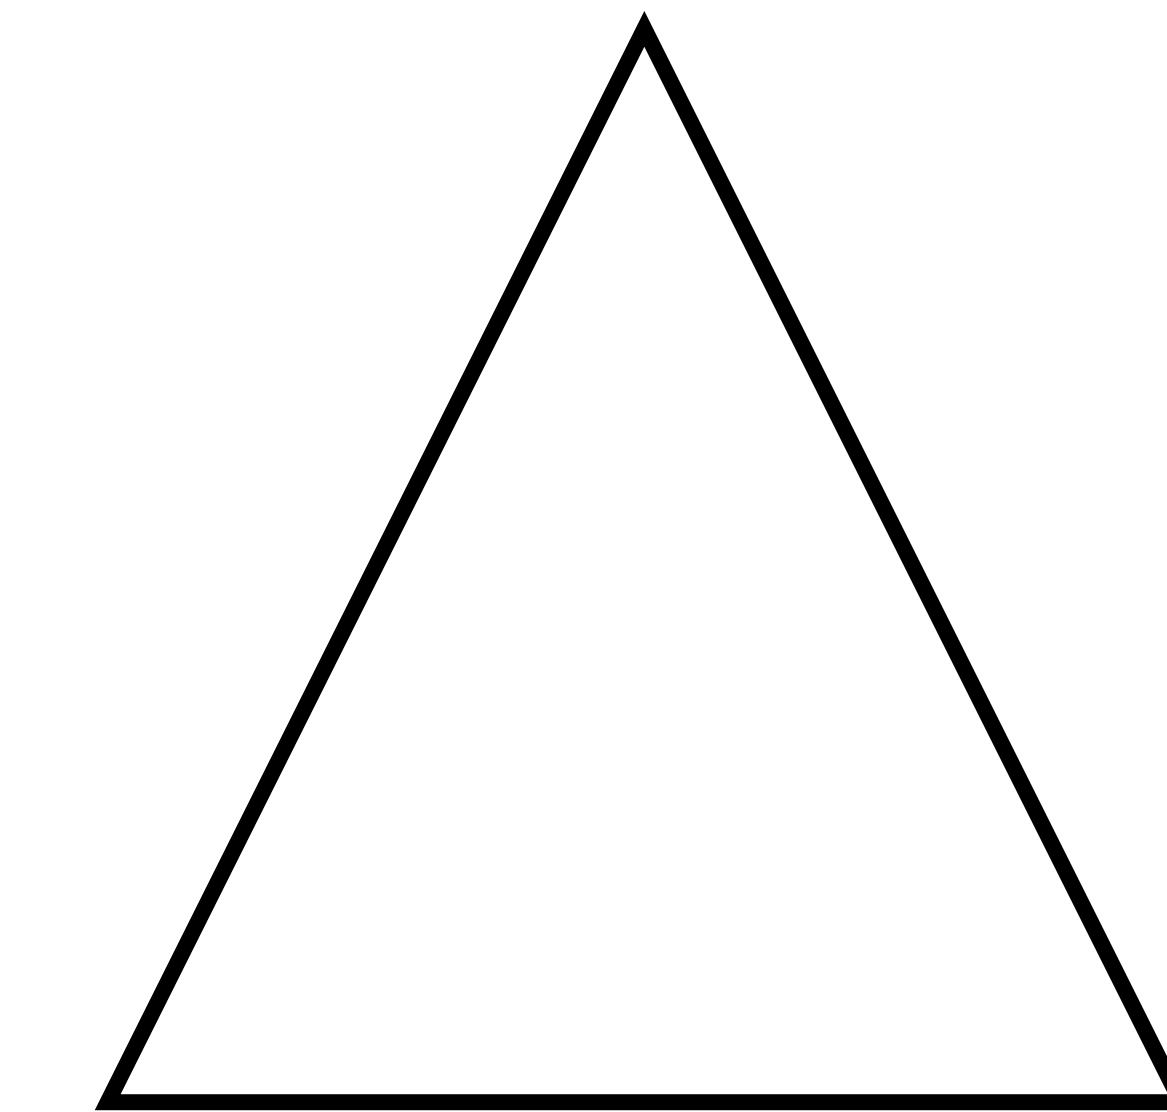


Target Geometry

# Gradient Based Optimization

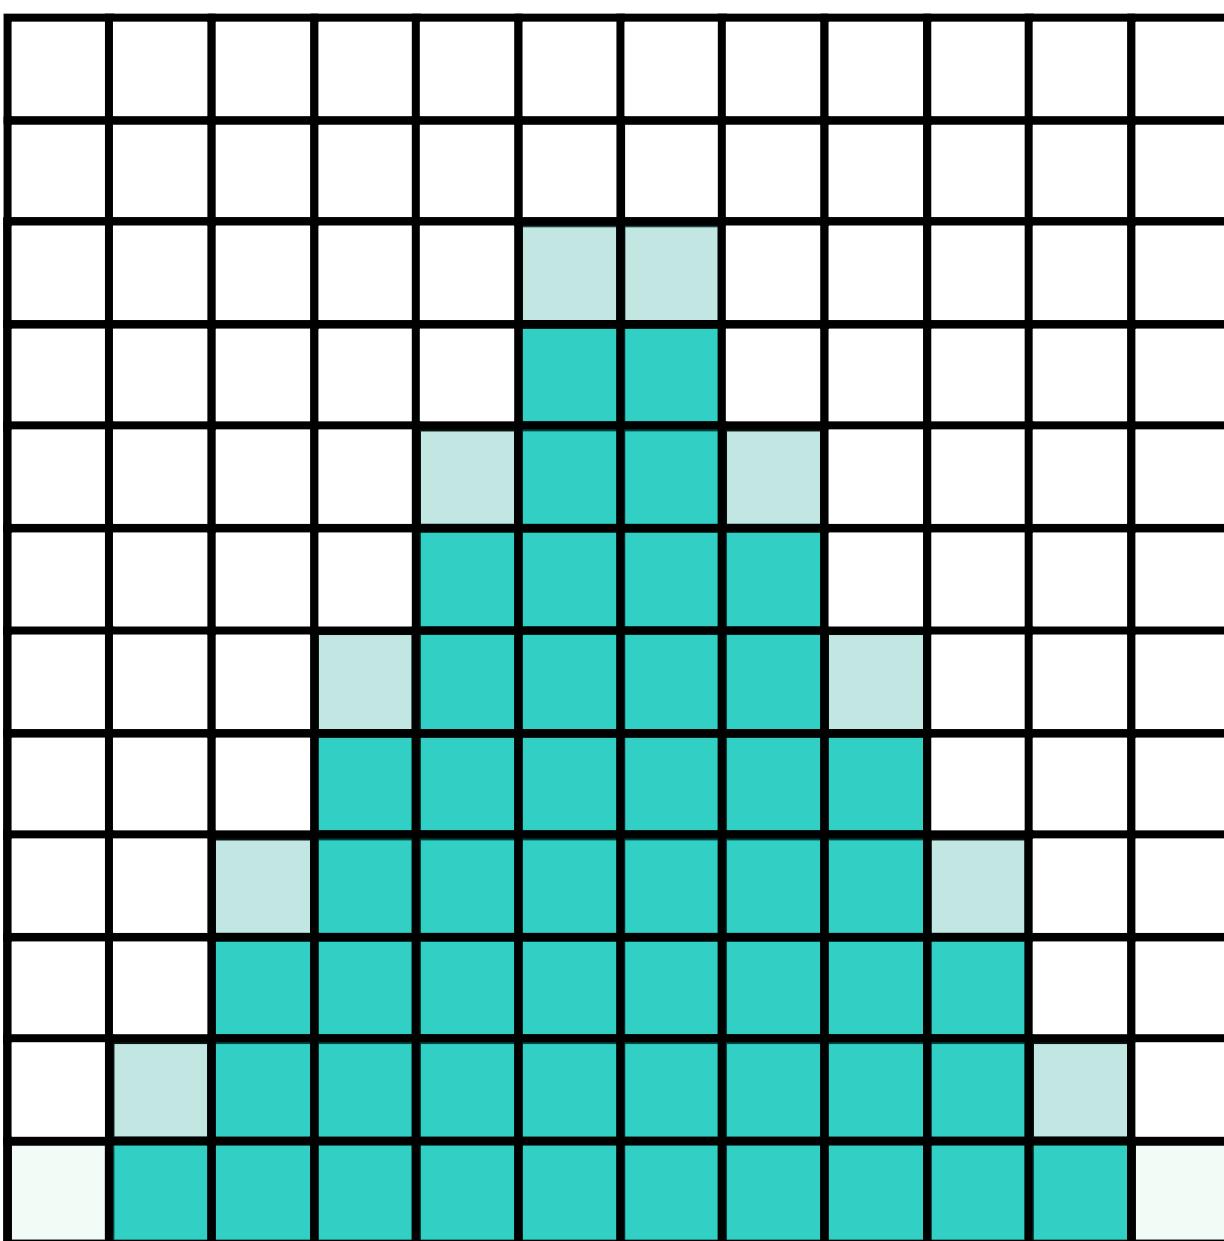


Repeat

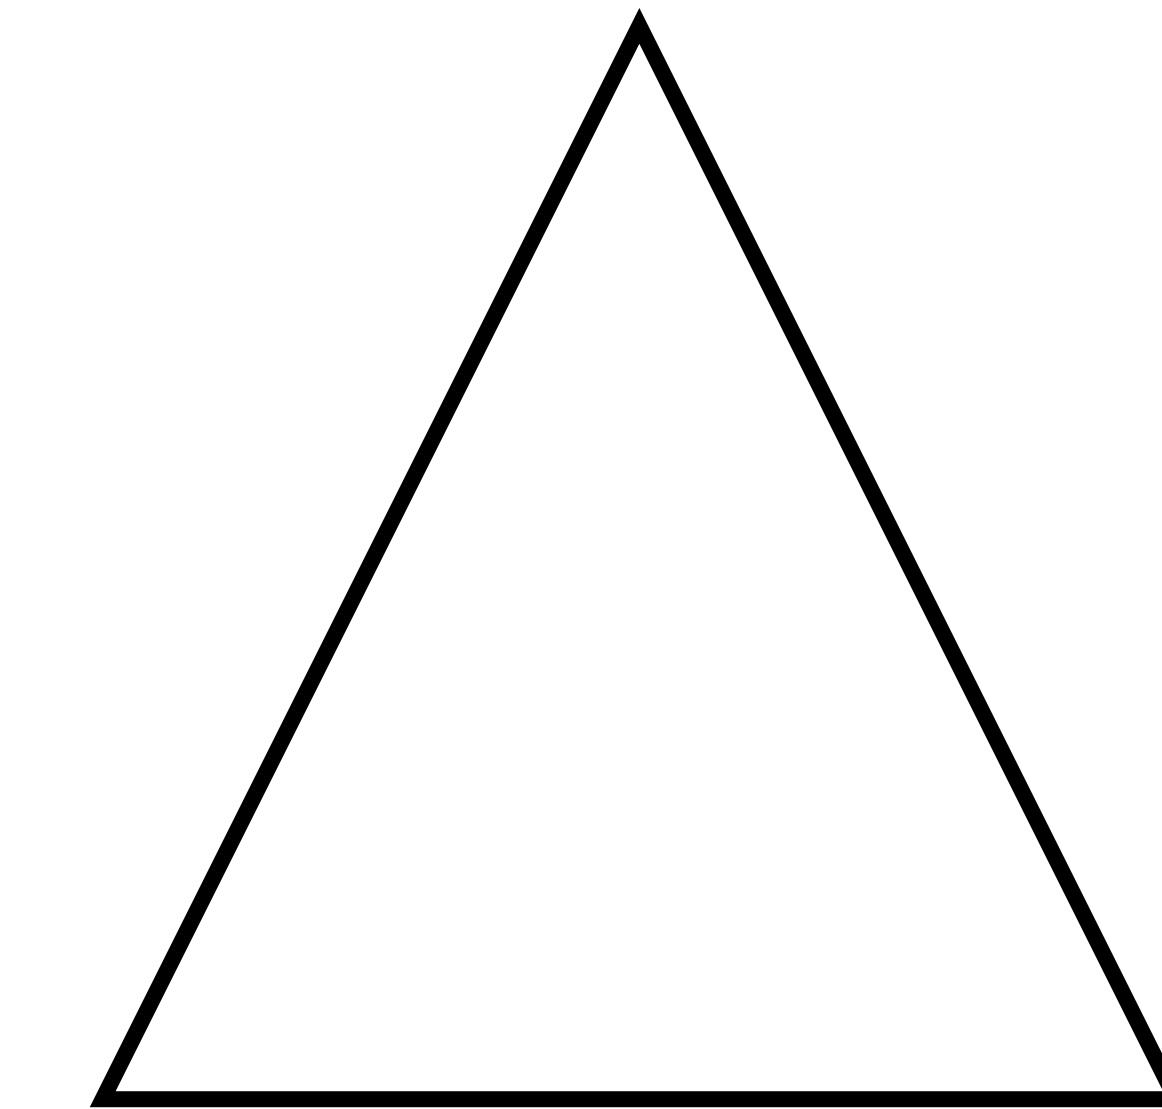


Target Geometry

# Gradient Based Optimization

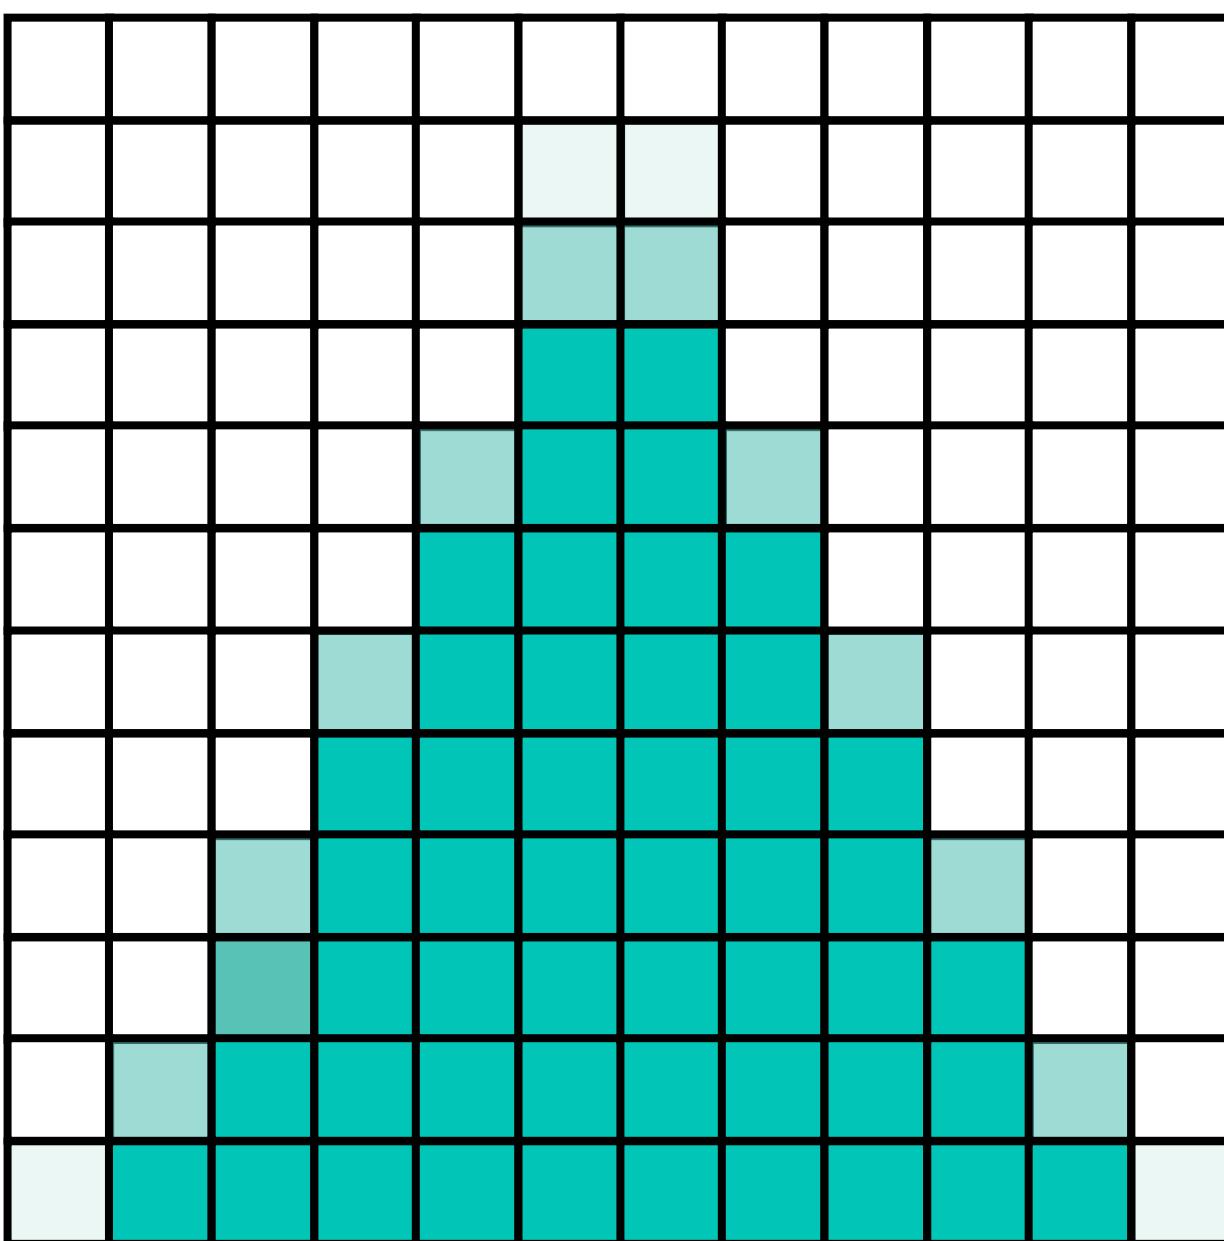


Repeat

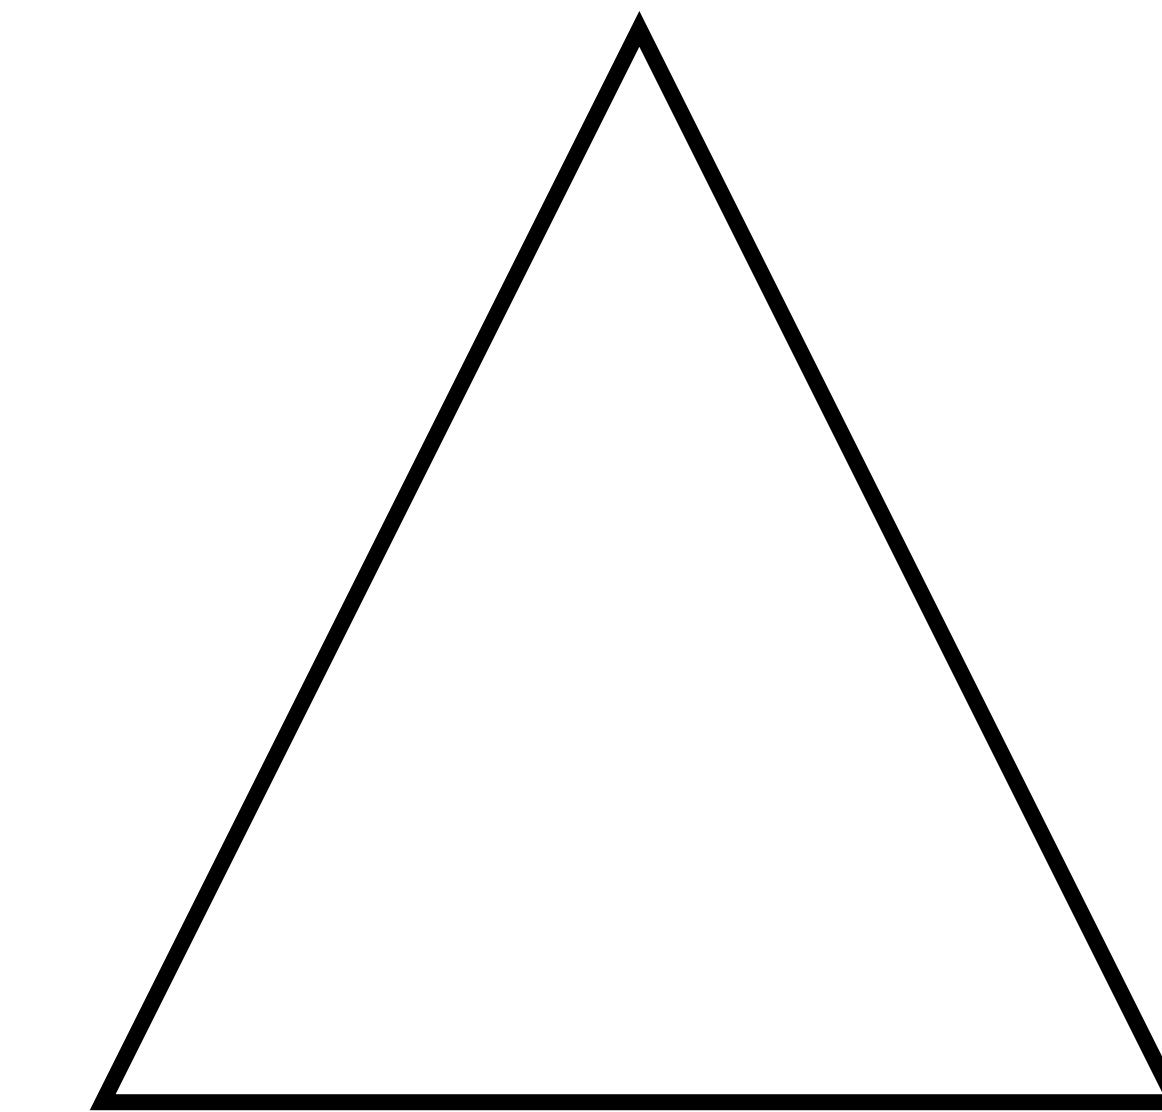


Target Geometry

# Gradient Based Optimization

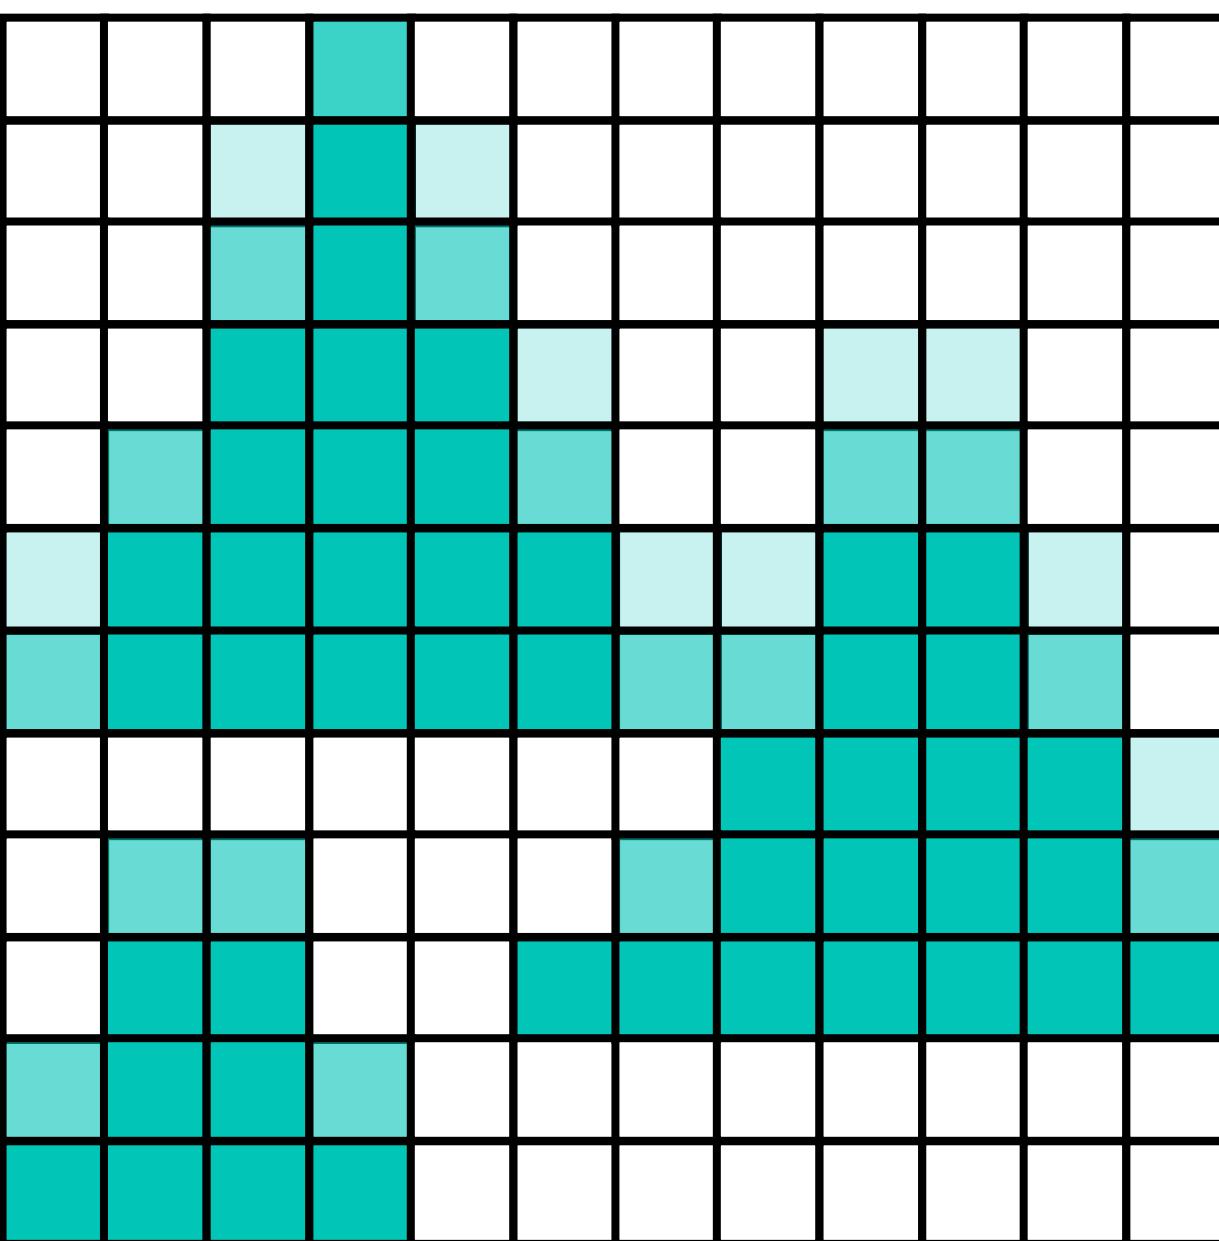


Repeat

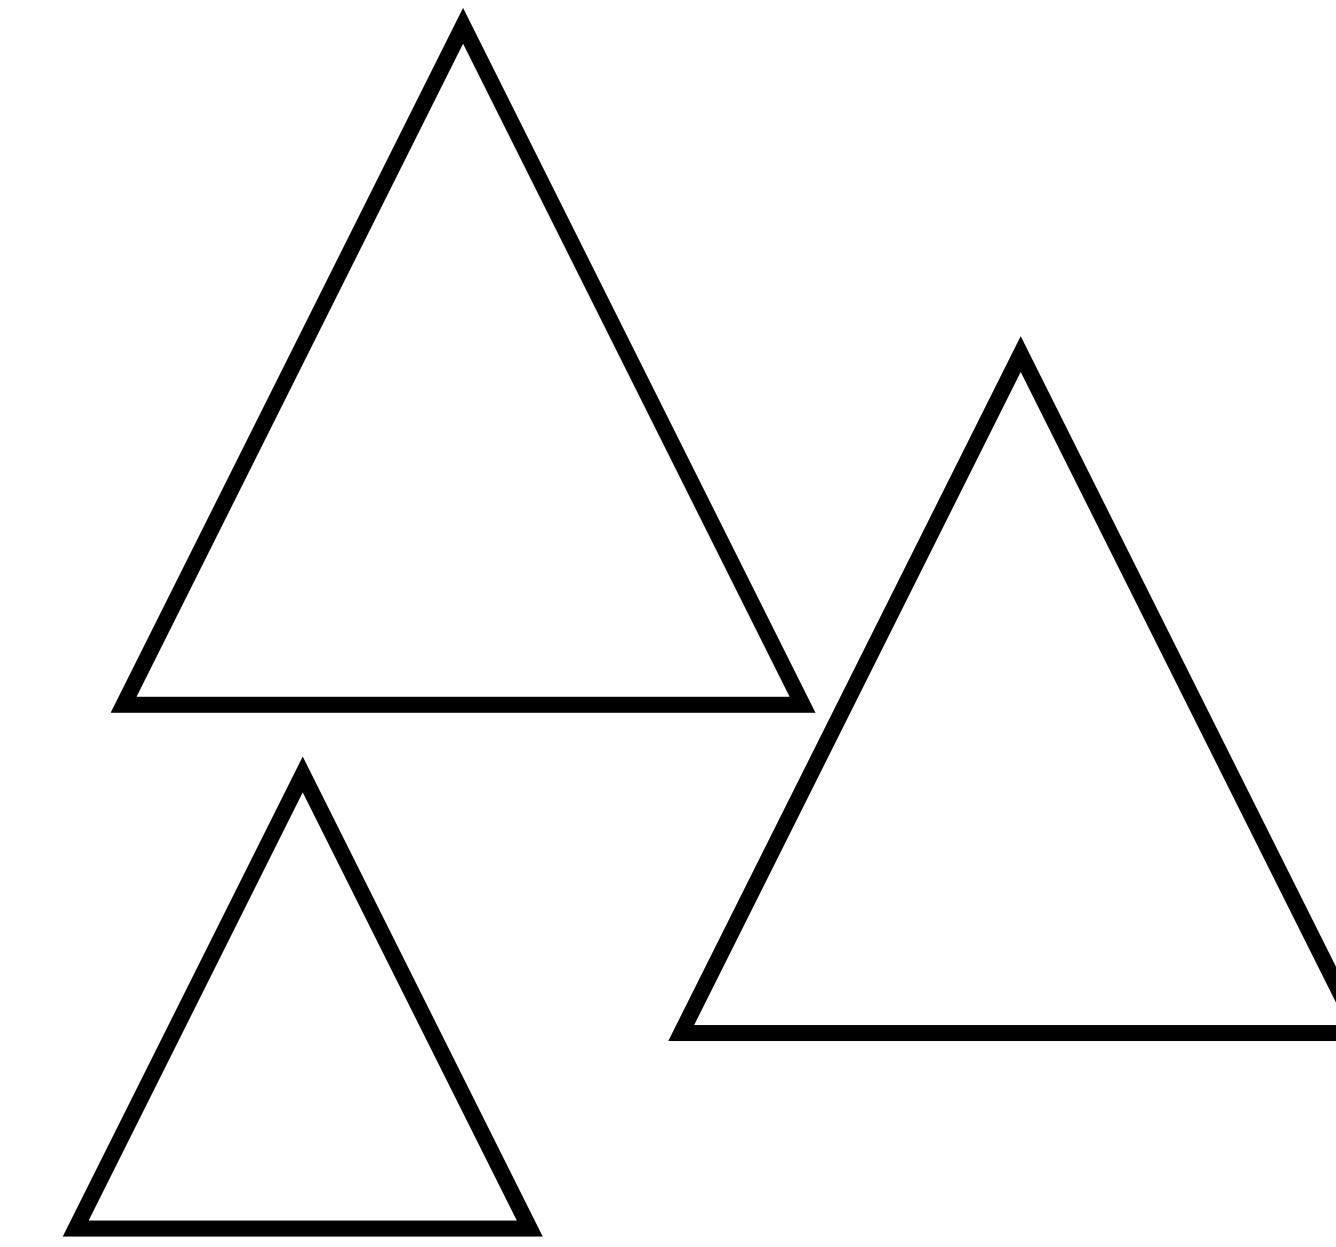


Target Geometry

# Gradient Based Optimization

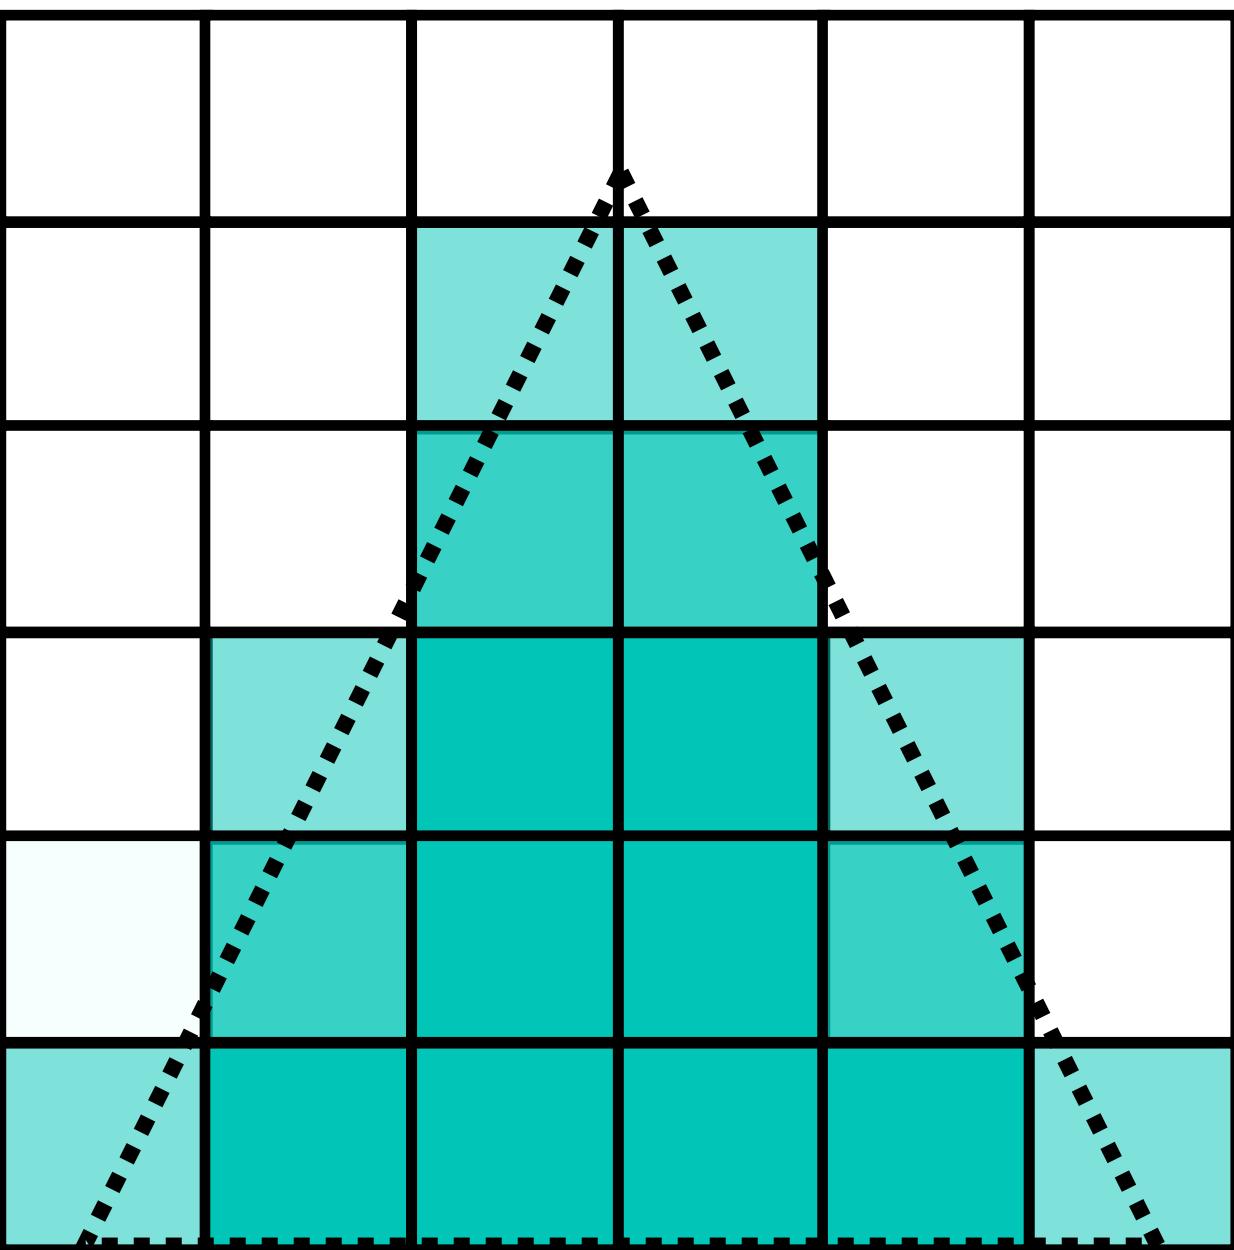


Reconstruction

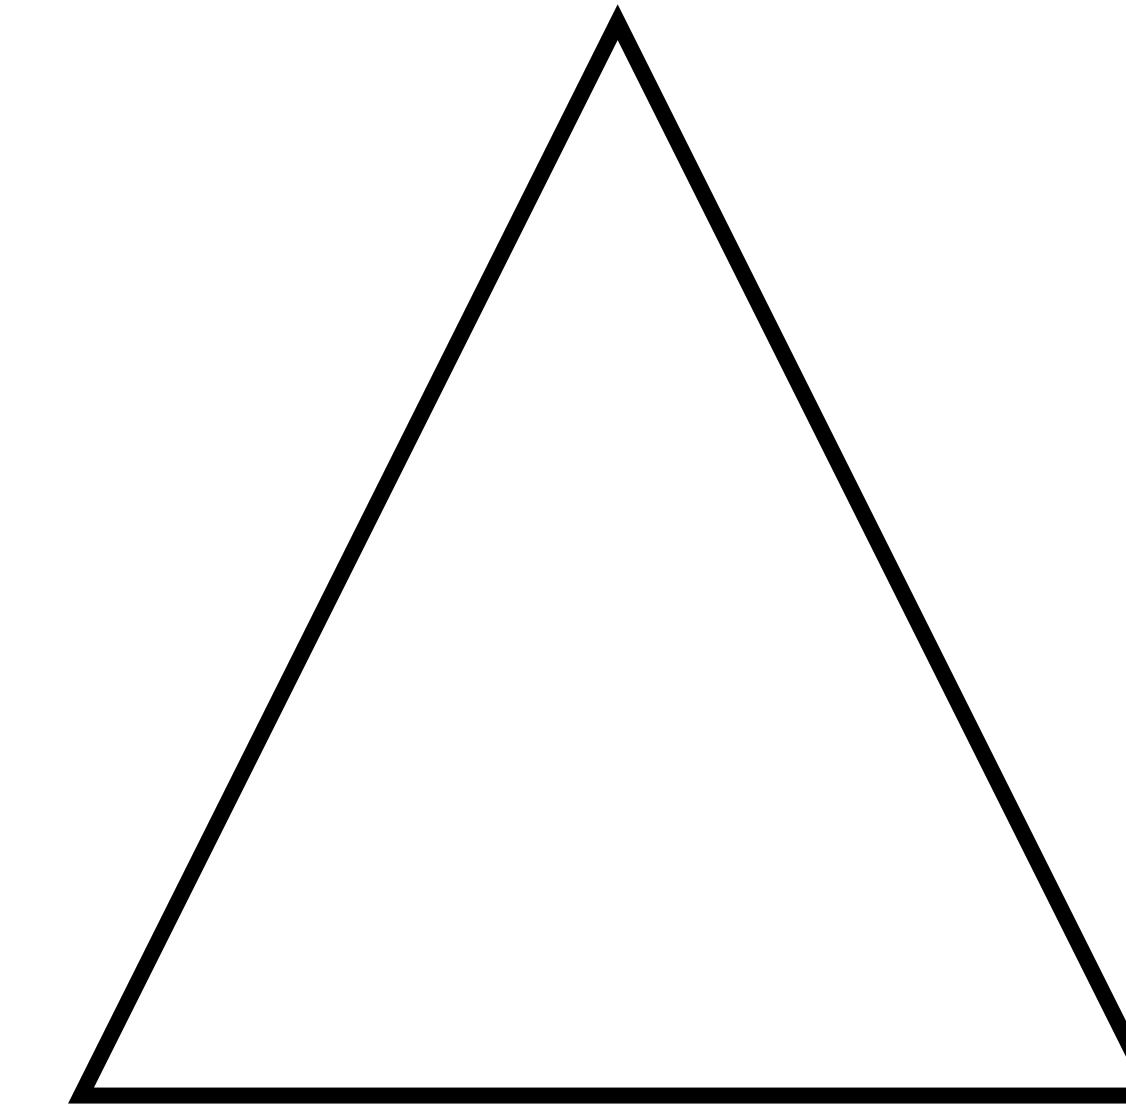


Target Geometry

# Gradient Based Optimization

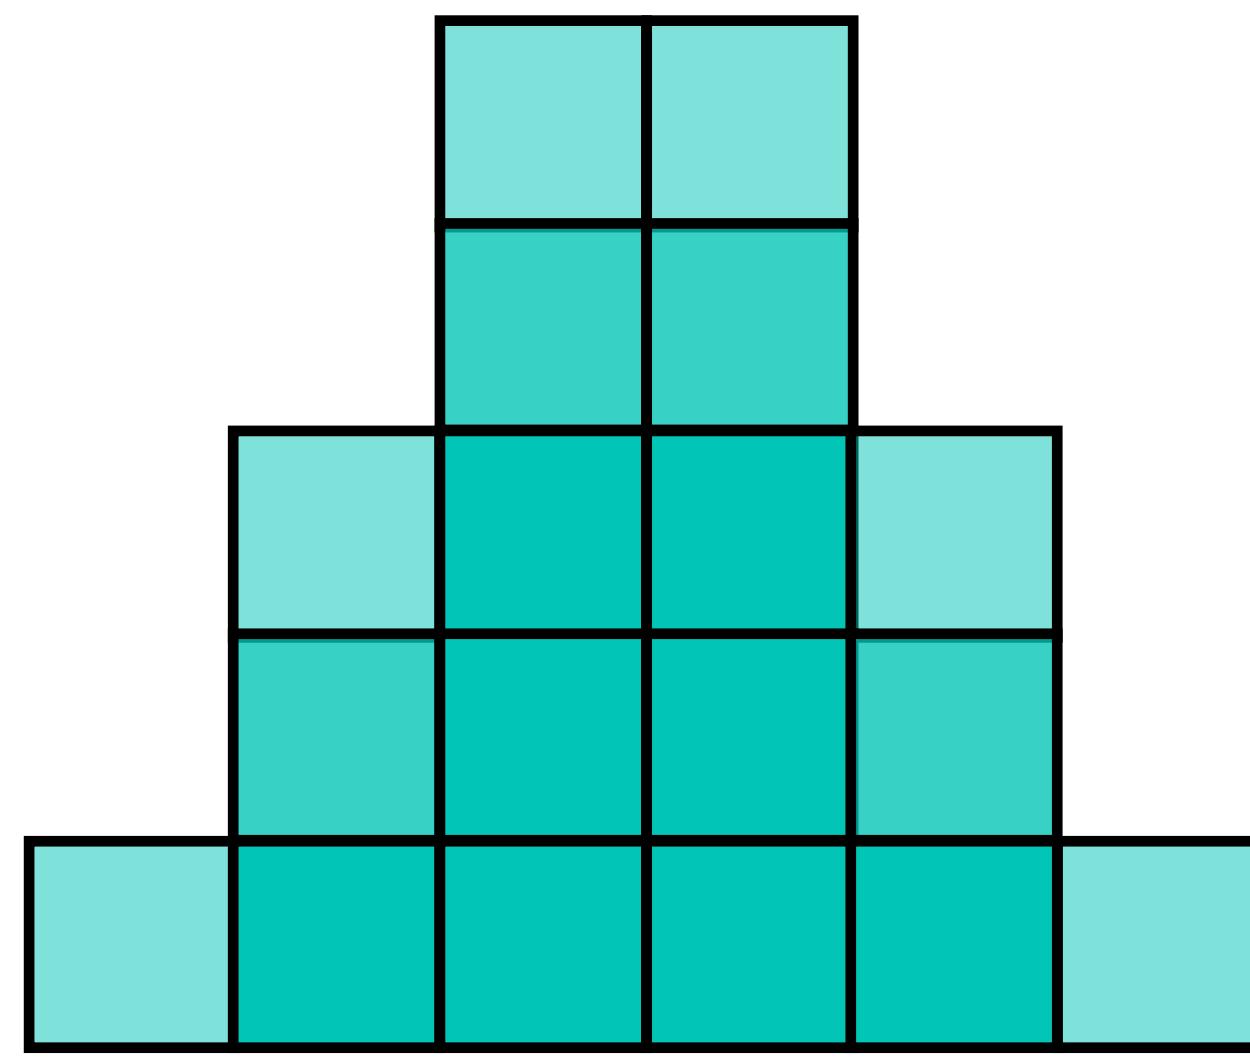


Reconstruction

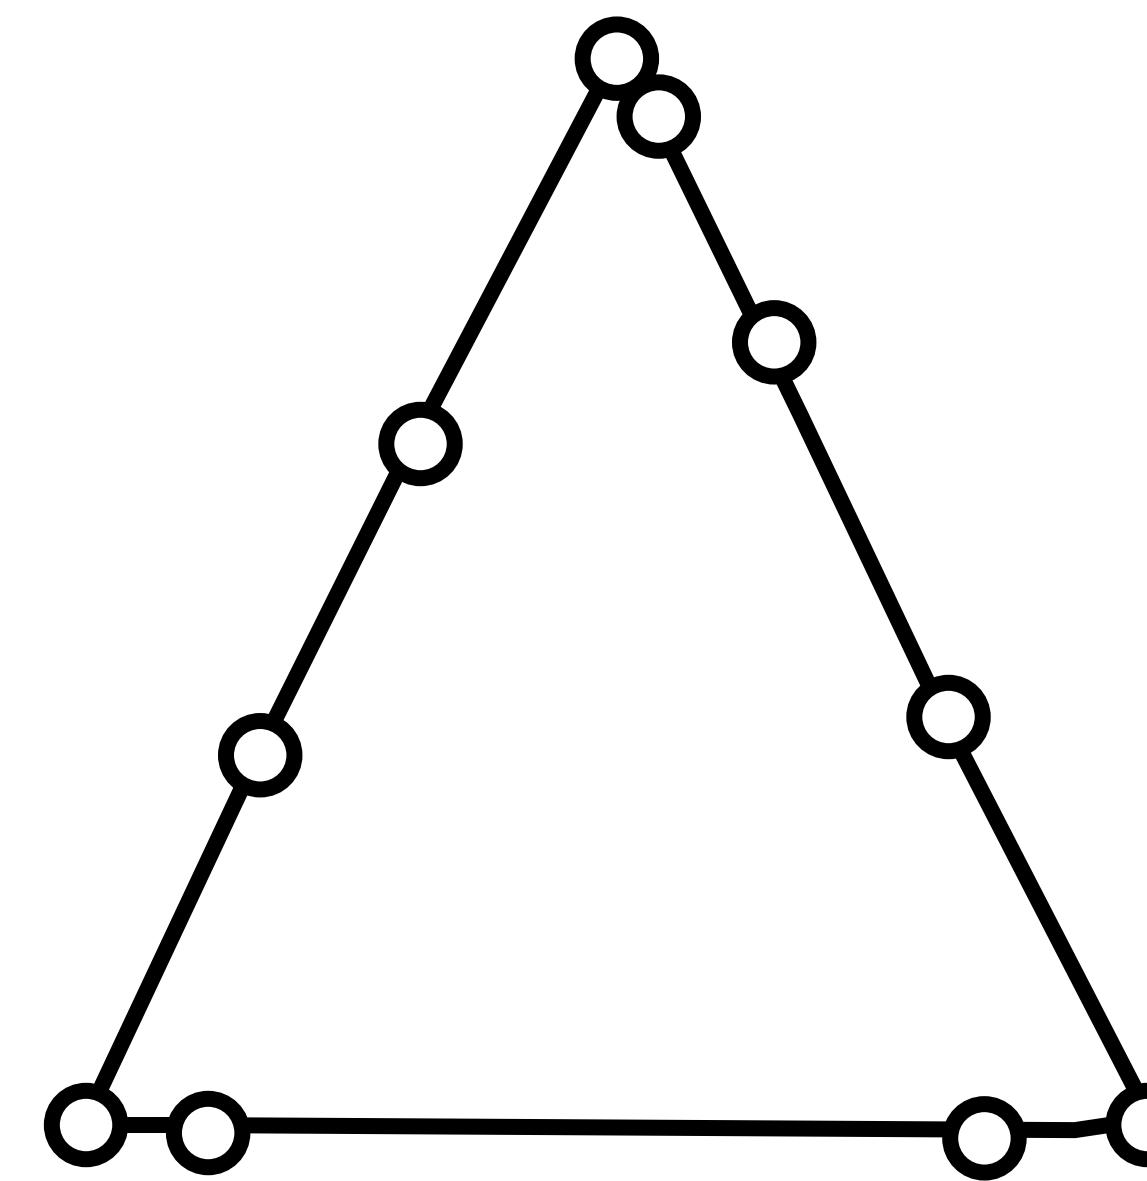


Target Geometry

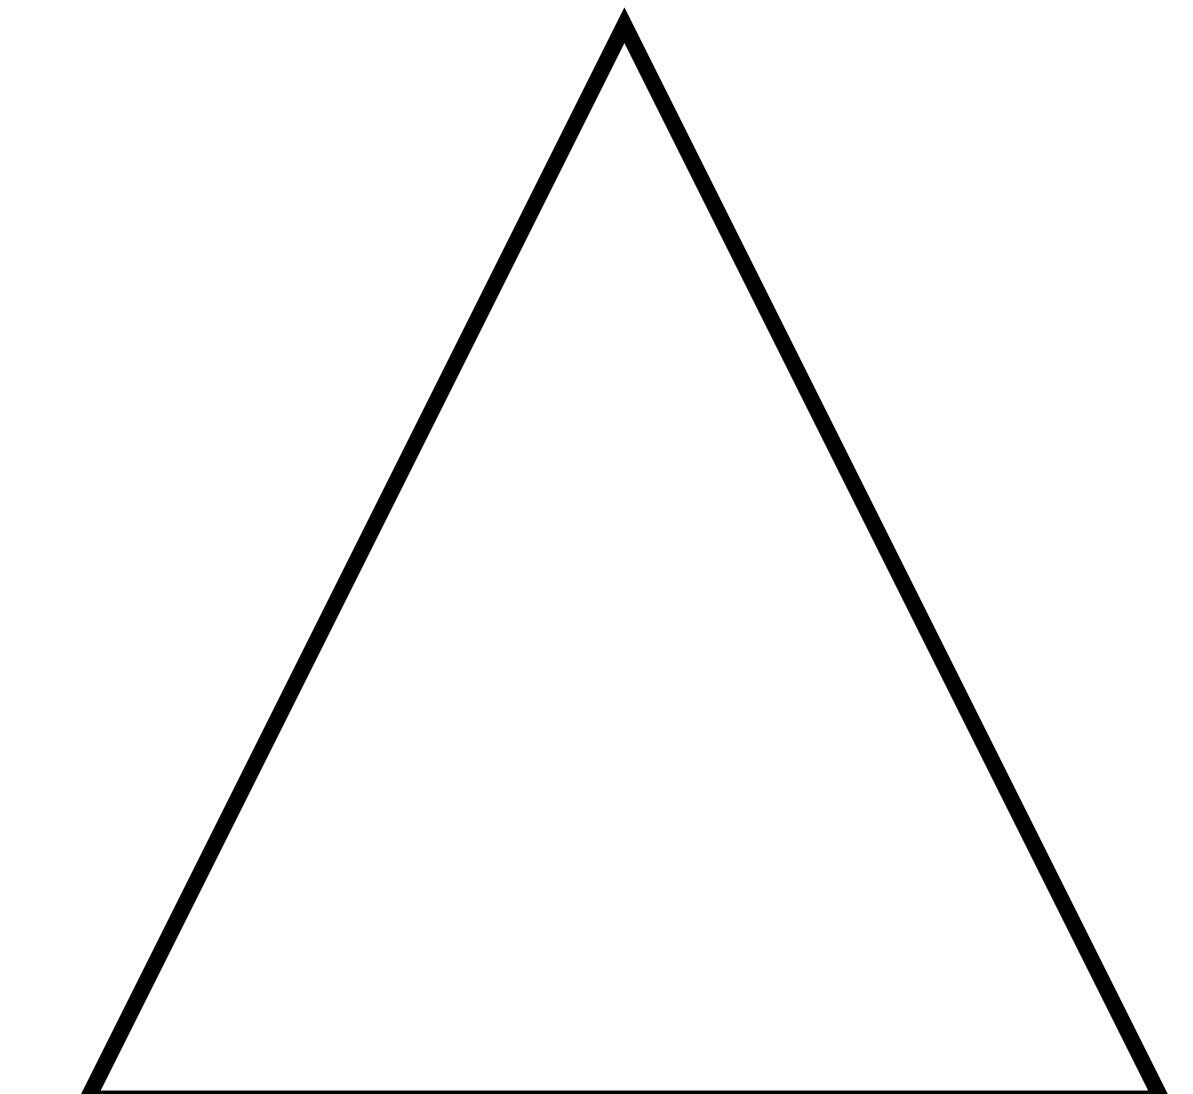
# Gradient Based Optimization



Voxel

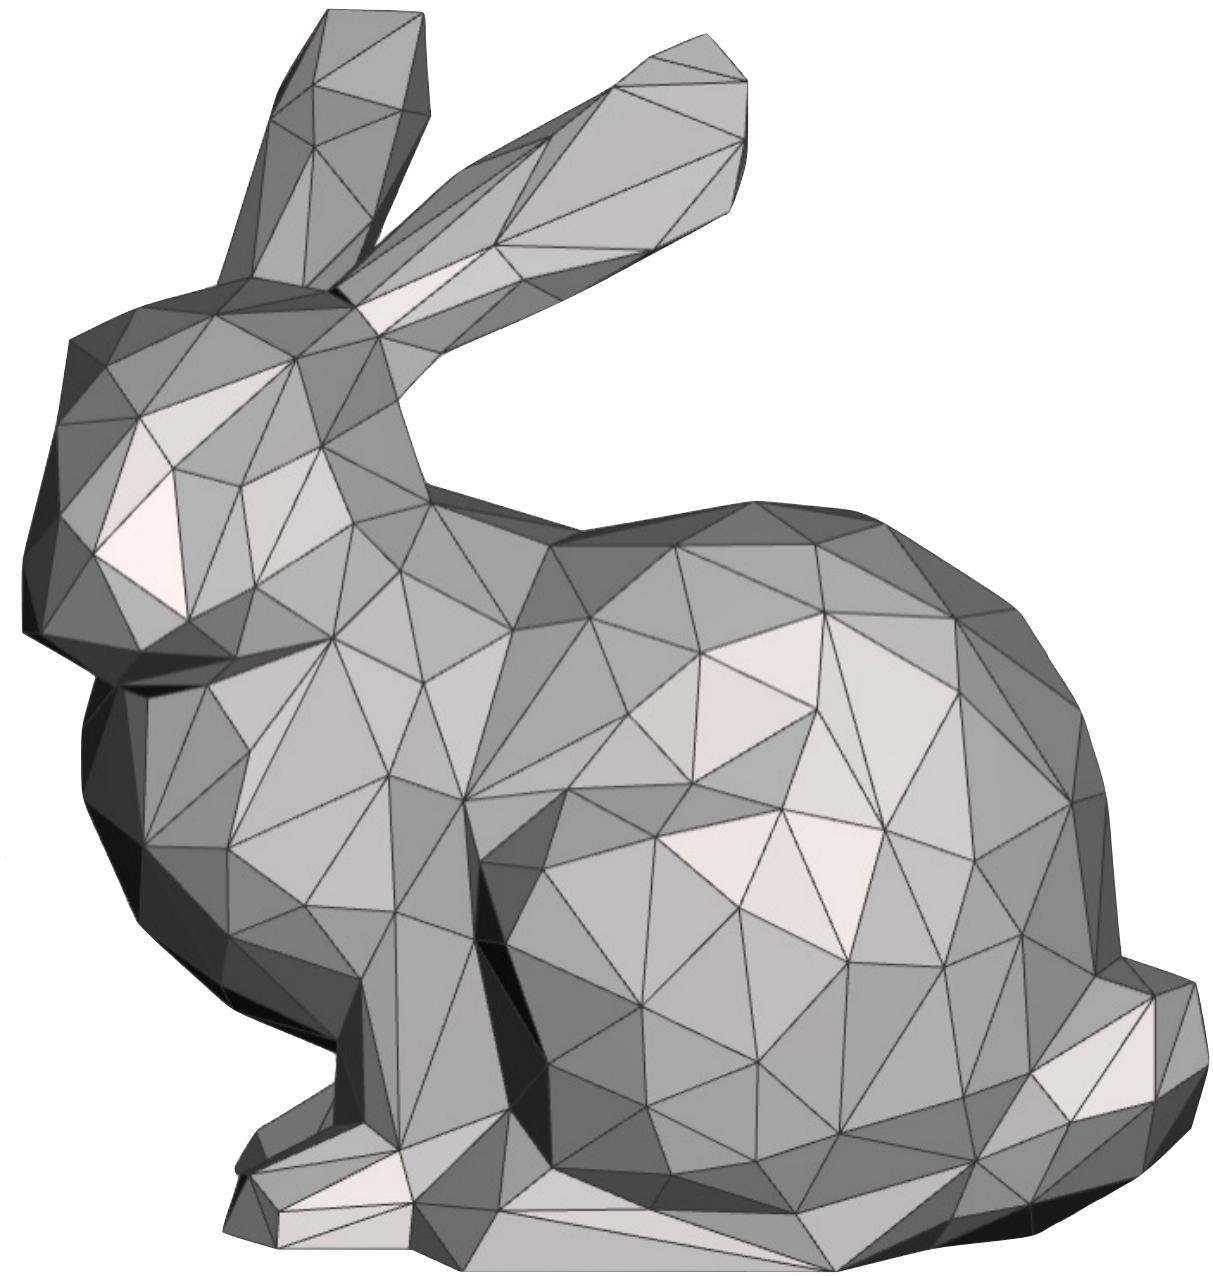


Mesh



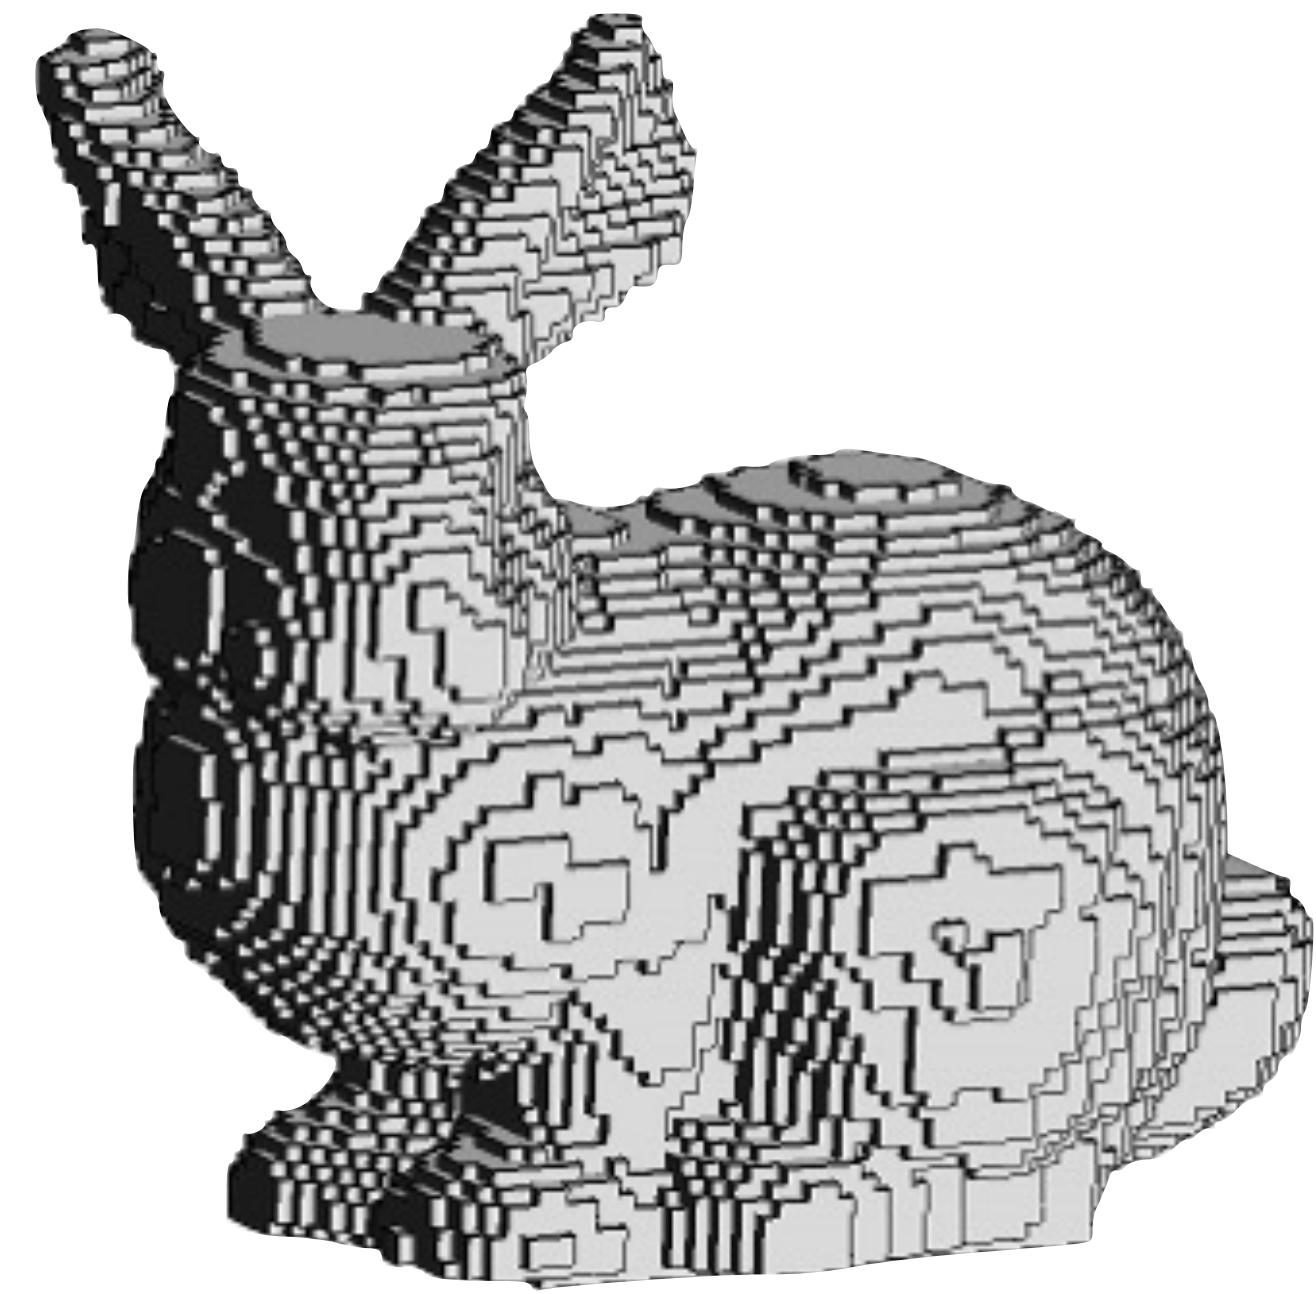
Target Geometry

# Geometry Representations



Mesh Representation

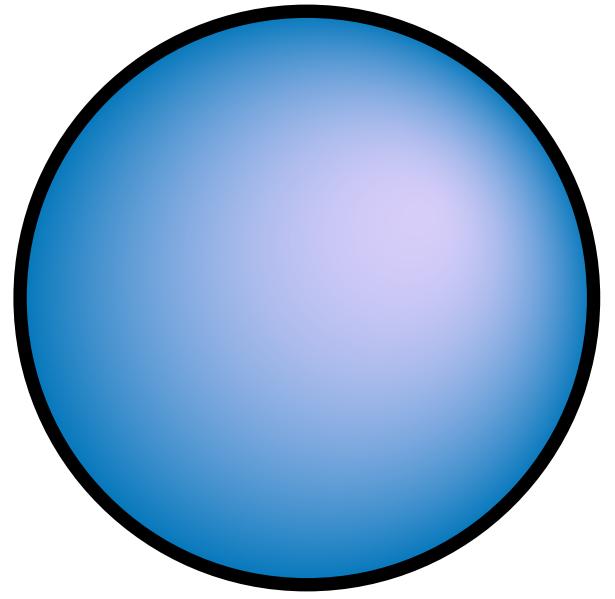
Small memory footprint  
Hard to optimize



Voxel Representation

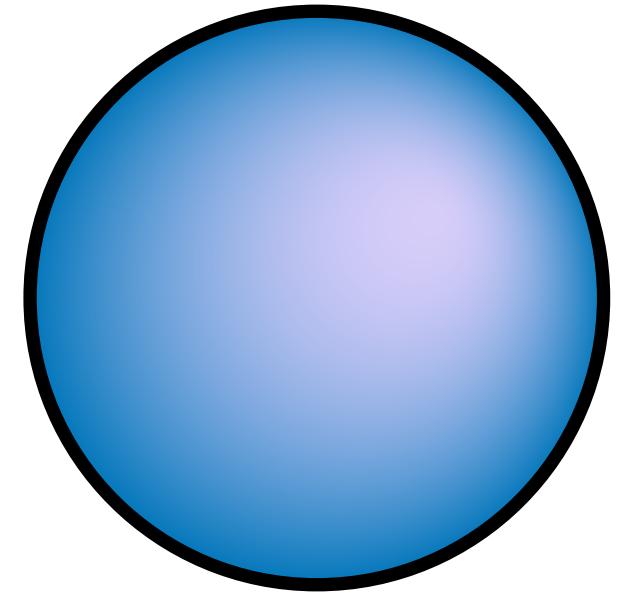
Easy to optimize  
Large memory footprint

# Implicit Functions

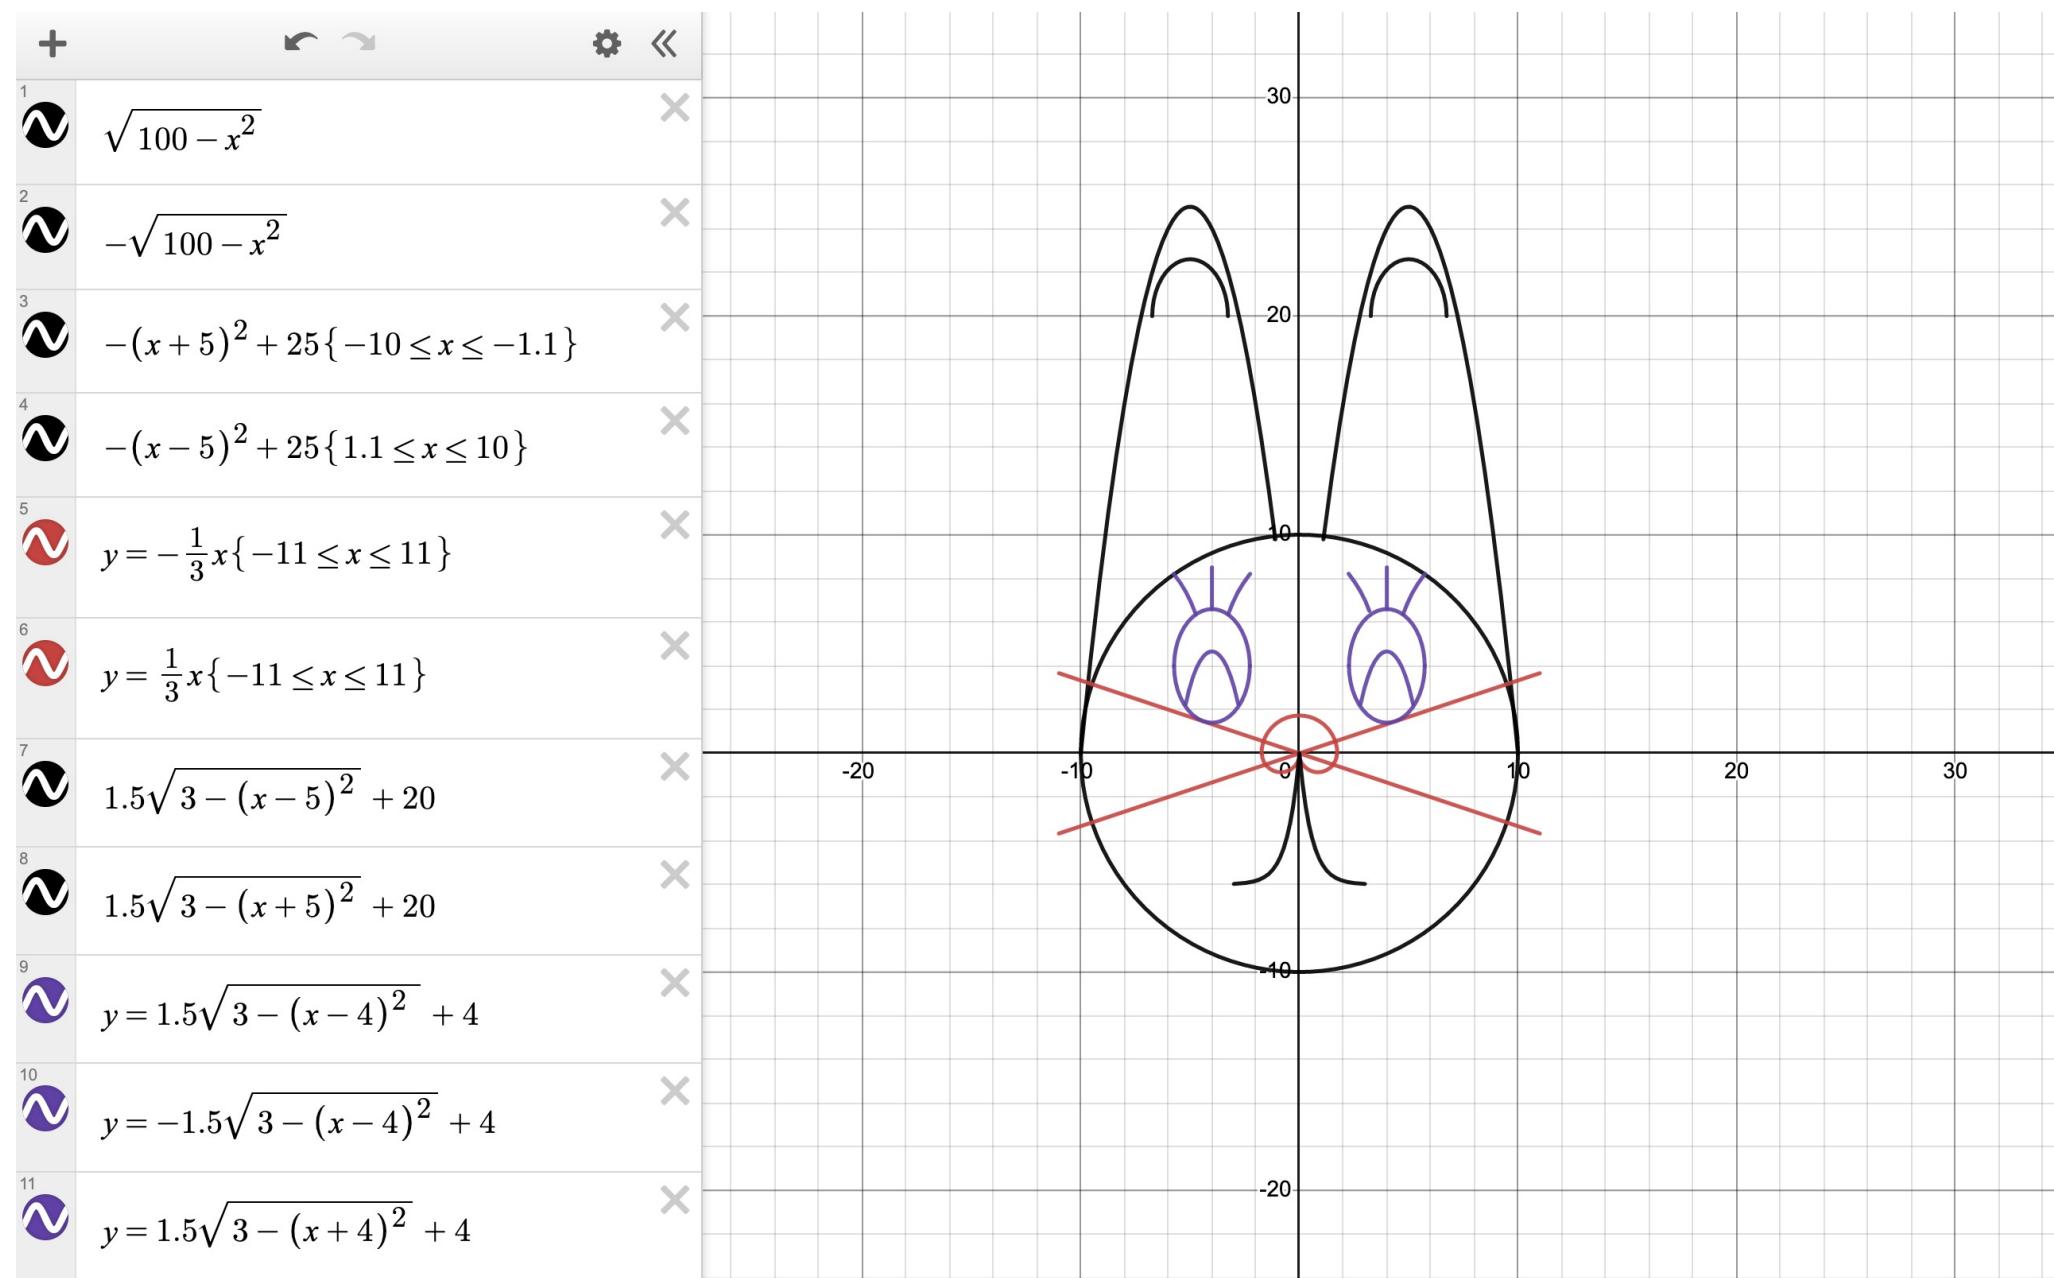


$$x^2 + y^2 + z^2 = 1$$

# Implicit Functions

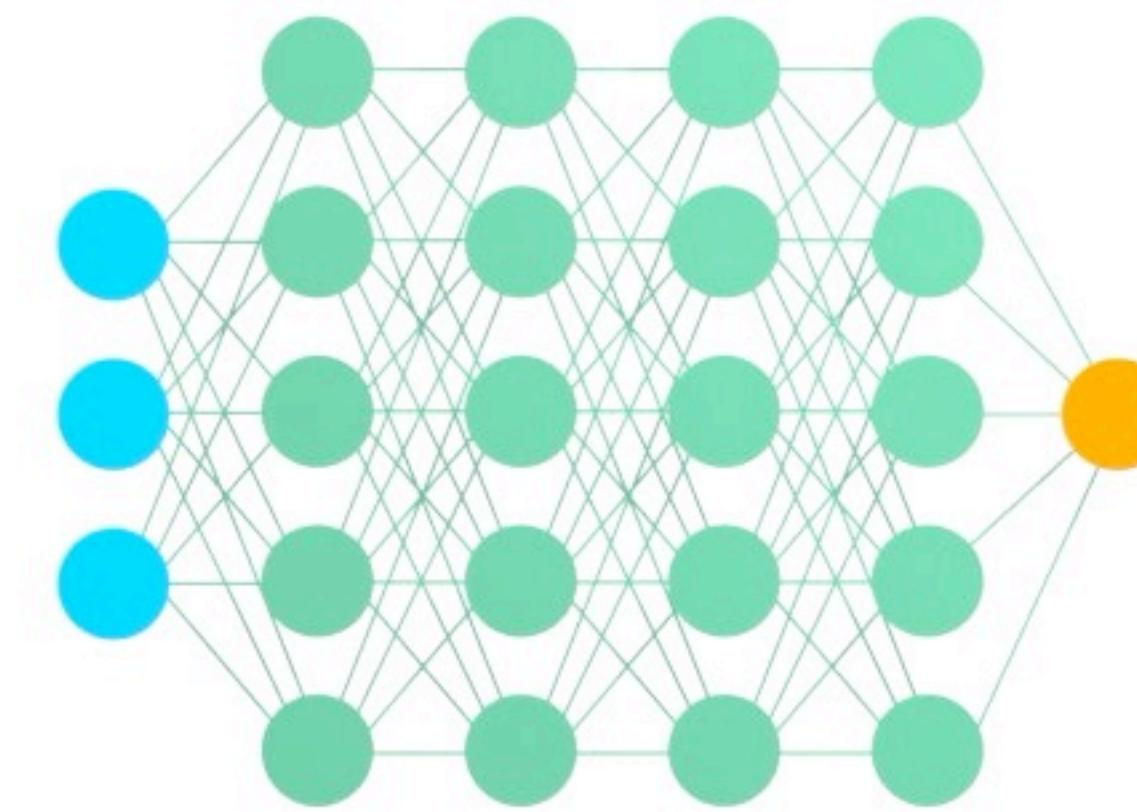


$$x^2 + y^2 + z^2 = 1$$



# Coordinate Based Neural Network

Input  
Coordinate



Value at  
Coordinate

Multi Layer Perceptron  
MLP

# Neural networks as a continuous shape representation

**Occupancy Networks**  
(Mescheder et al. 2019)

$(x, y, z) \rightarrow \text{occupancy}$

**DeepSDF**  
(Park et al. 2019)

$(x, y, z) \rightarrow \text{distance}$

**Scene Representation Networks**  
(Sitzmann et al. 2019)

$(x, y, z) \rightarrow \text{latent vec. (color, dist.)}$

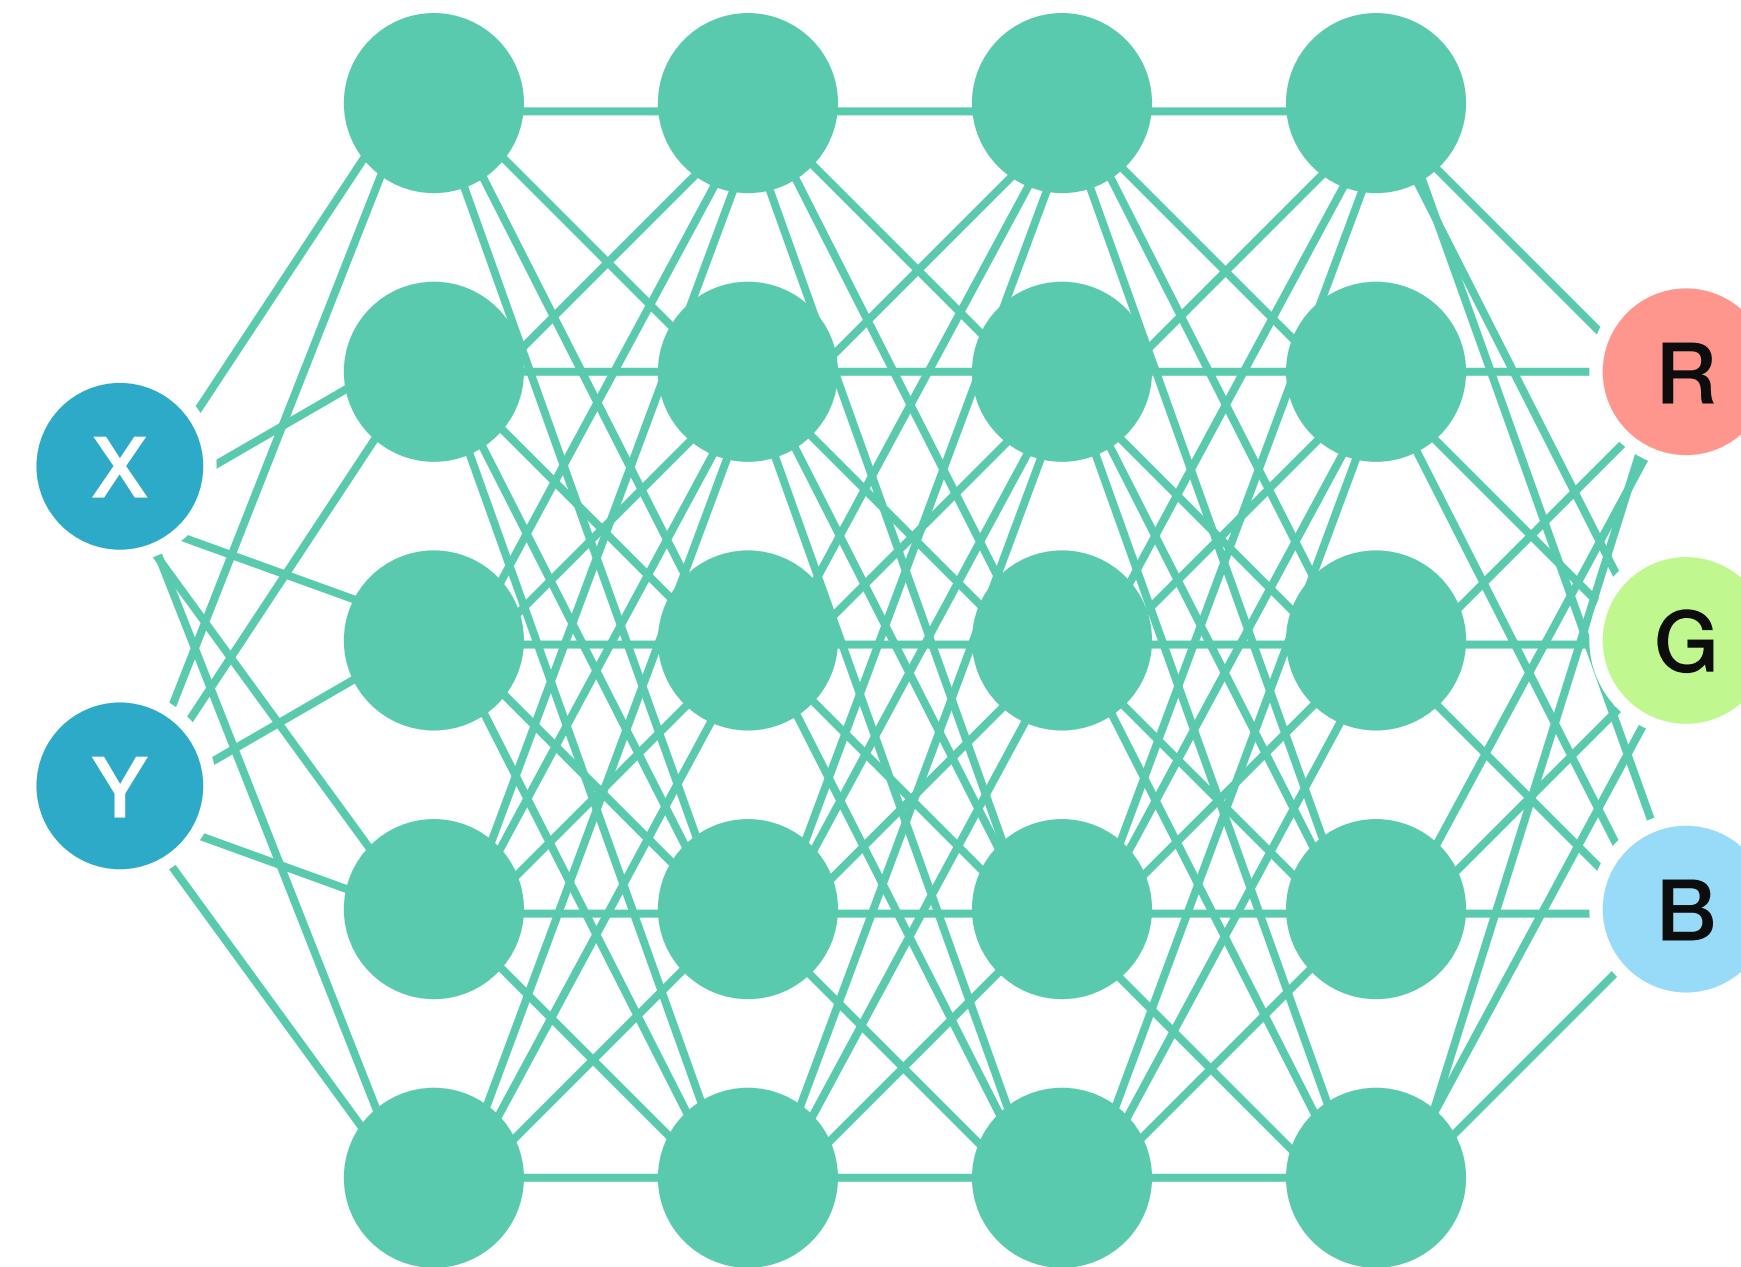
**Differentiable Volumetric Rendering**  
(Niemeyer et al. 2020)

$(x, y, z) \rightarrow \text{color, occ.}$

# Challenge:

- How to get MLPs to represent higher frequency functions?

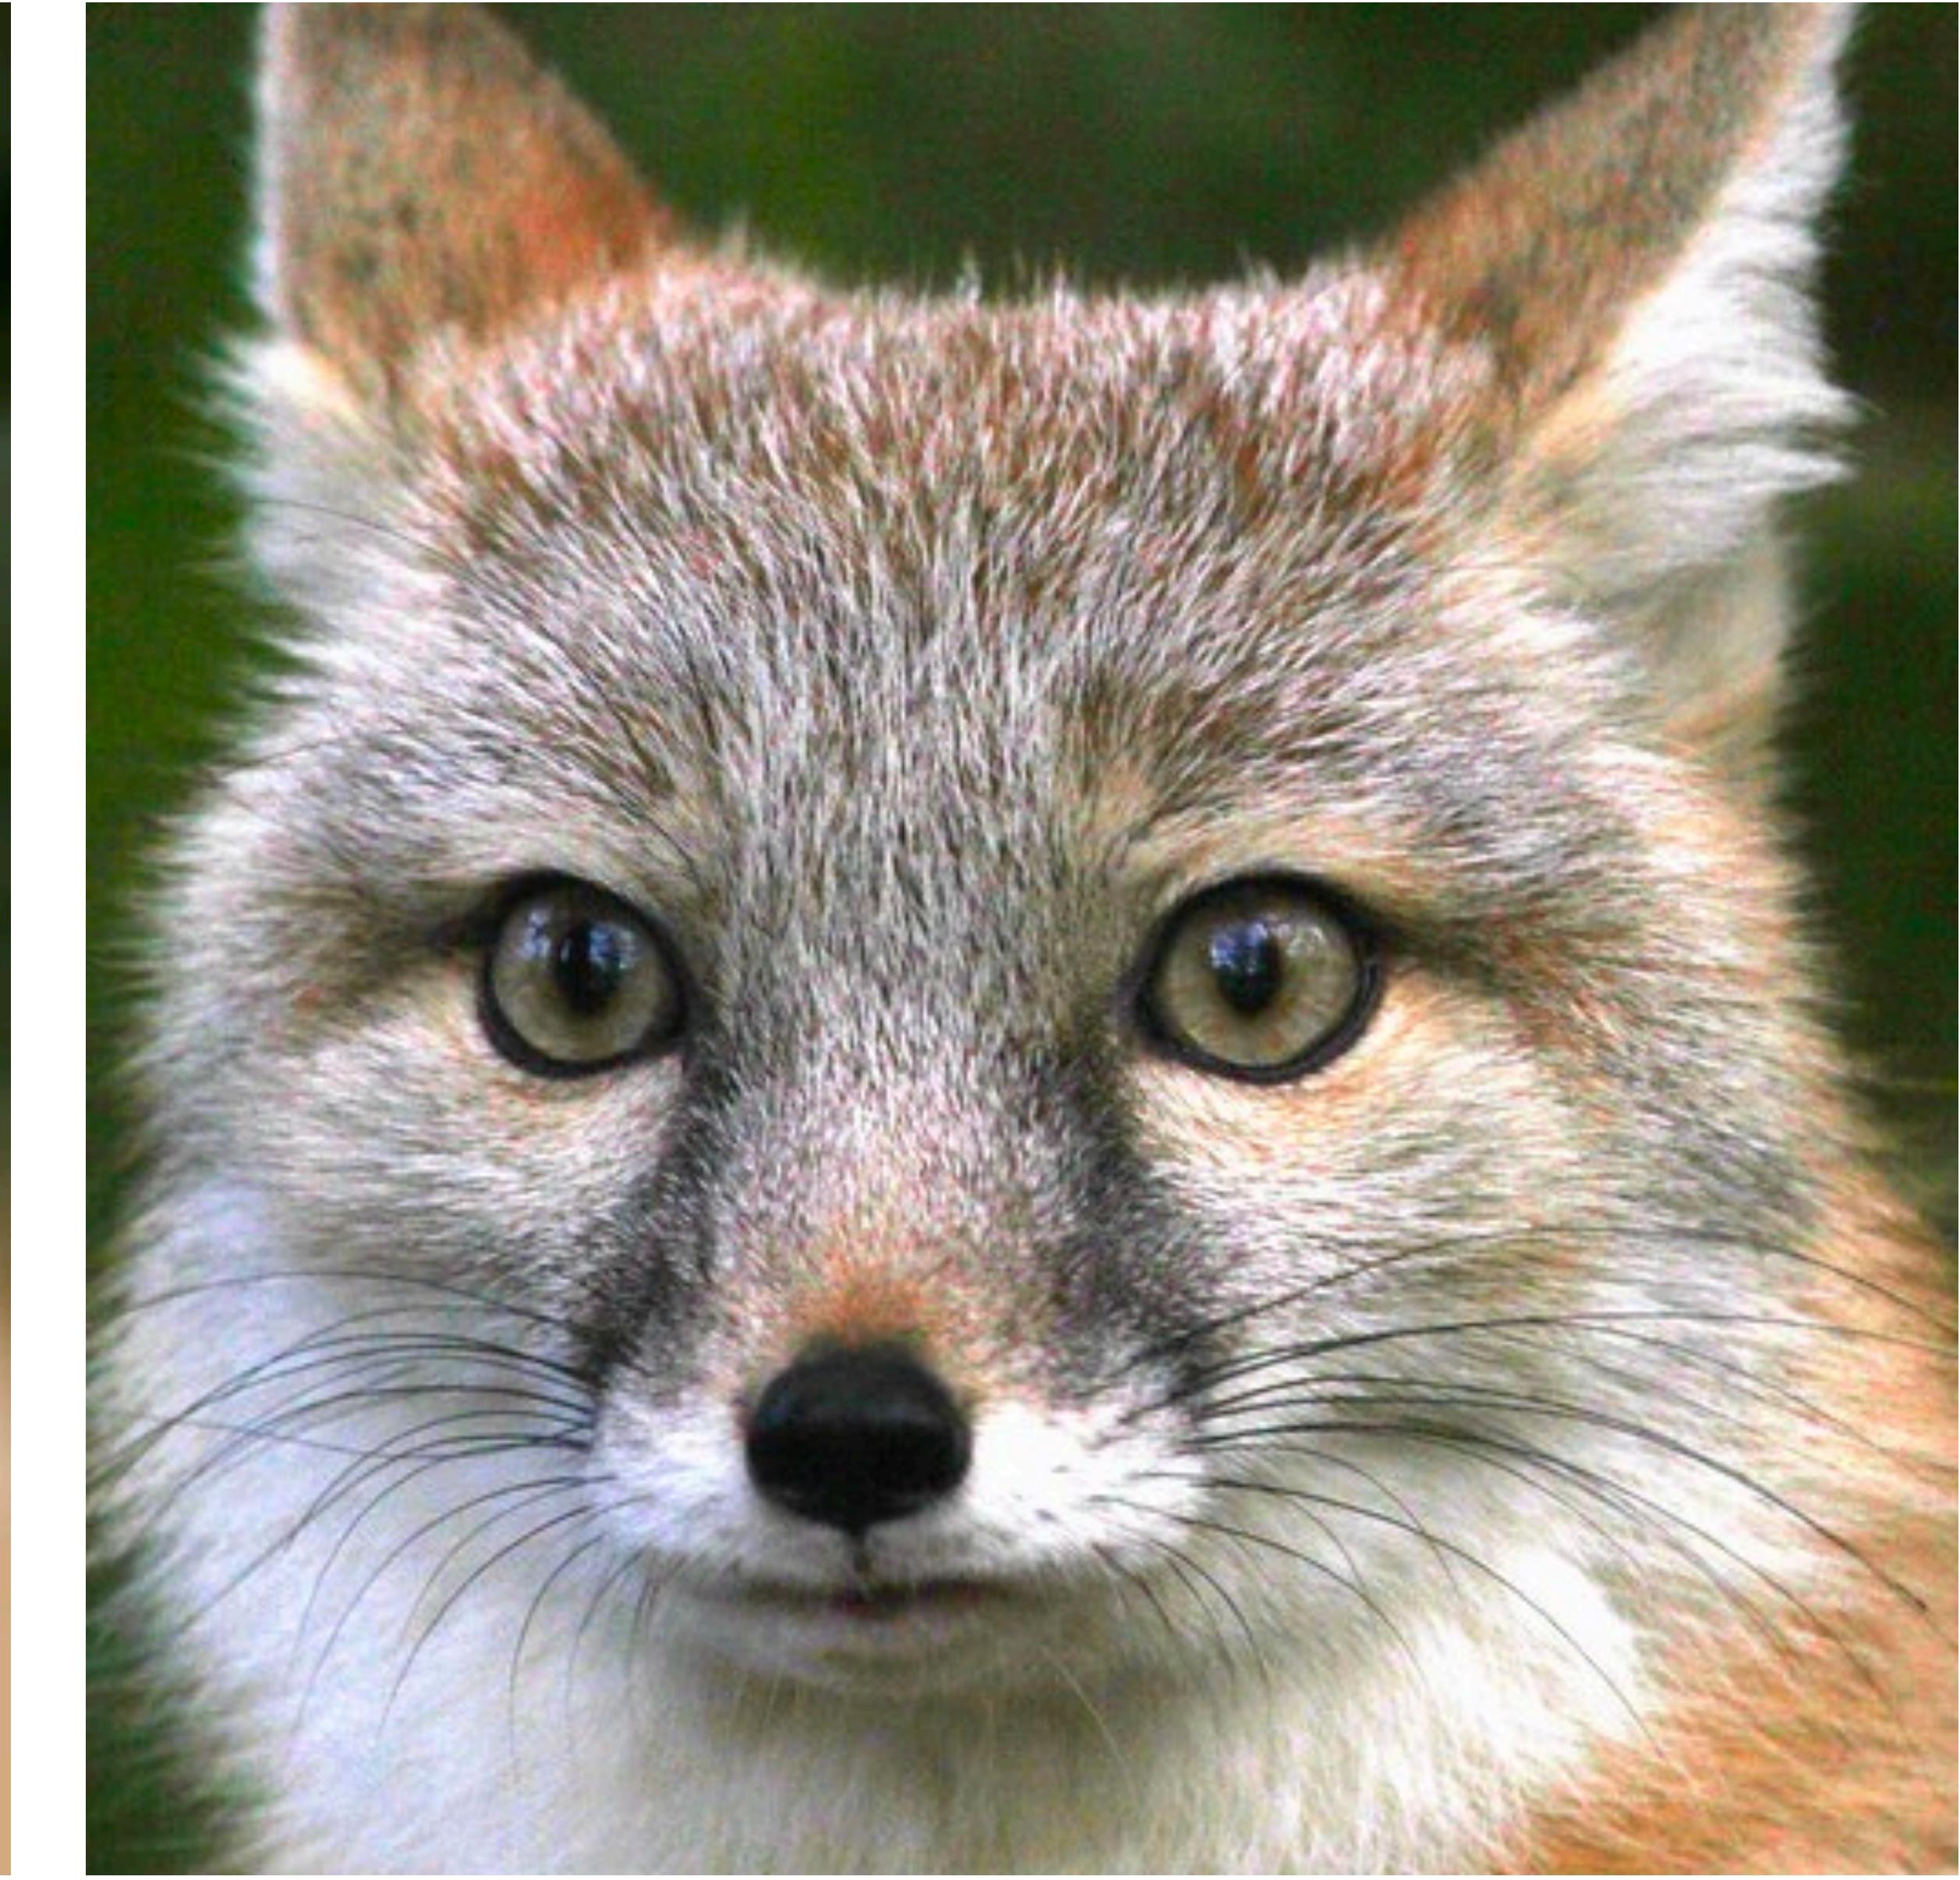
# Image Representation



Iteration 1000



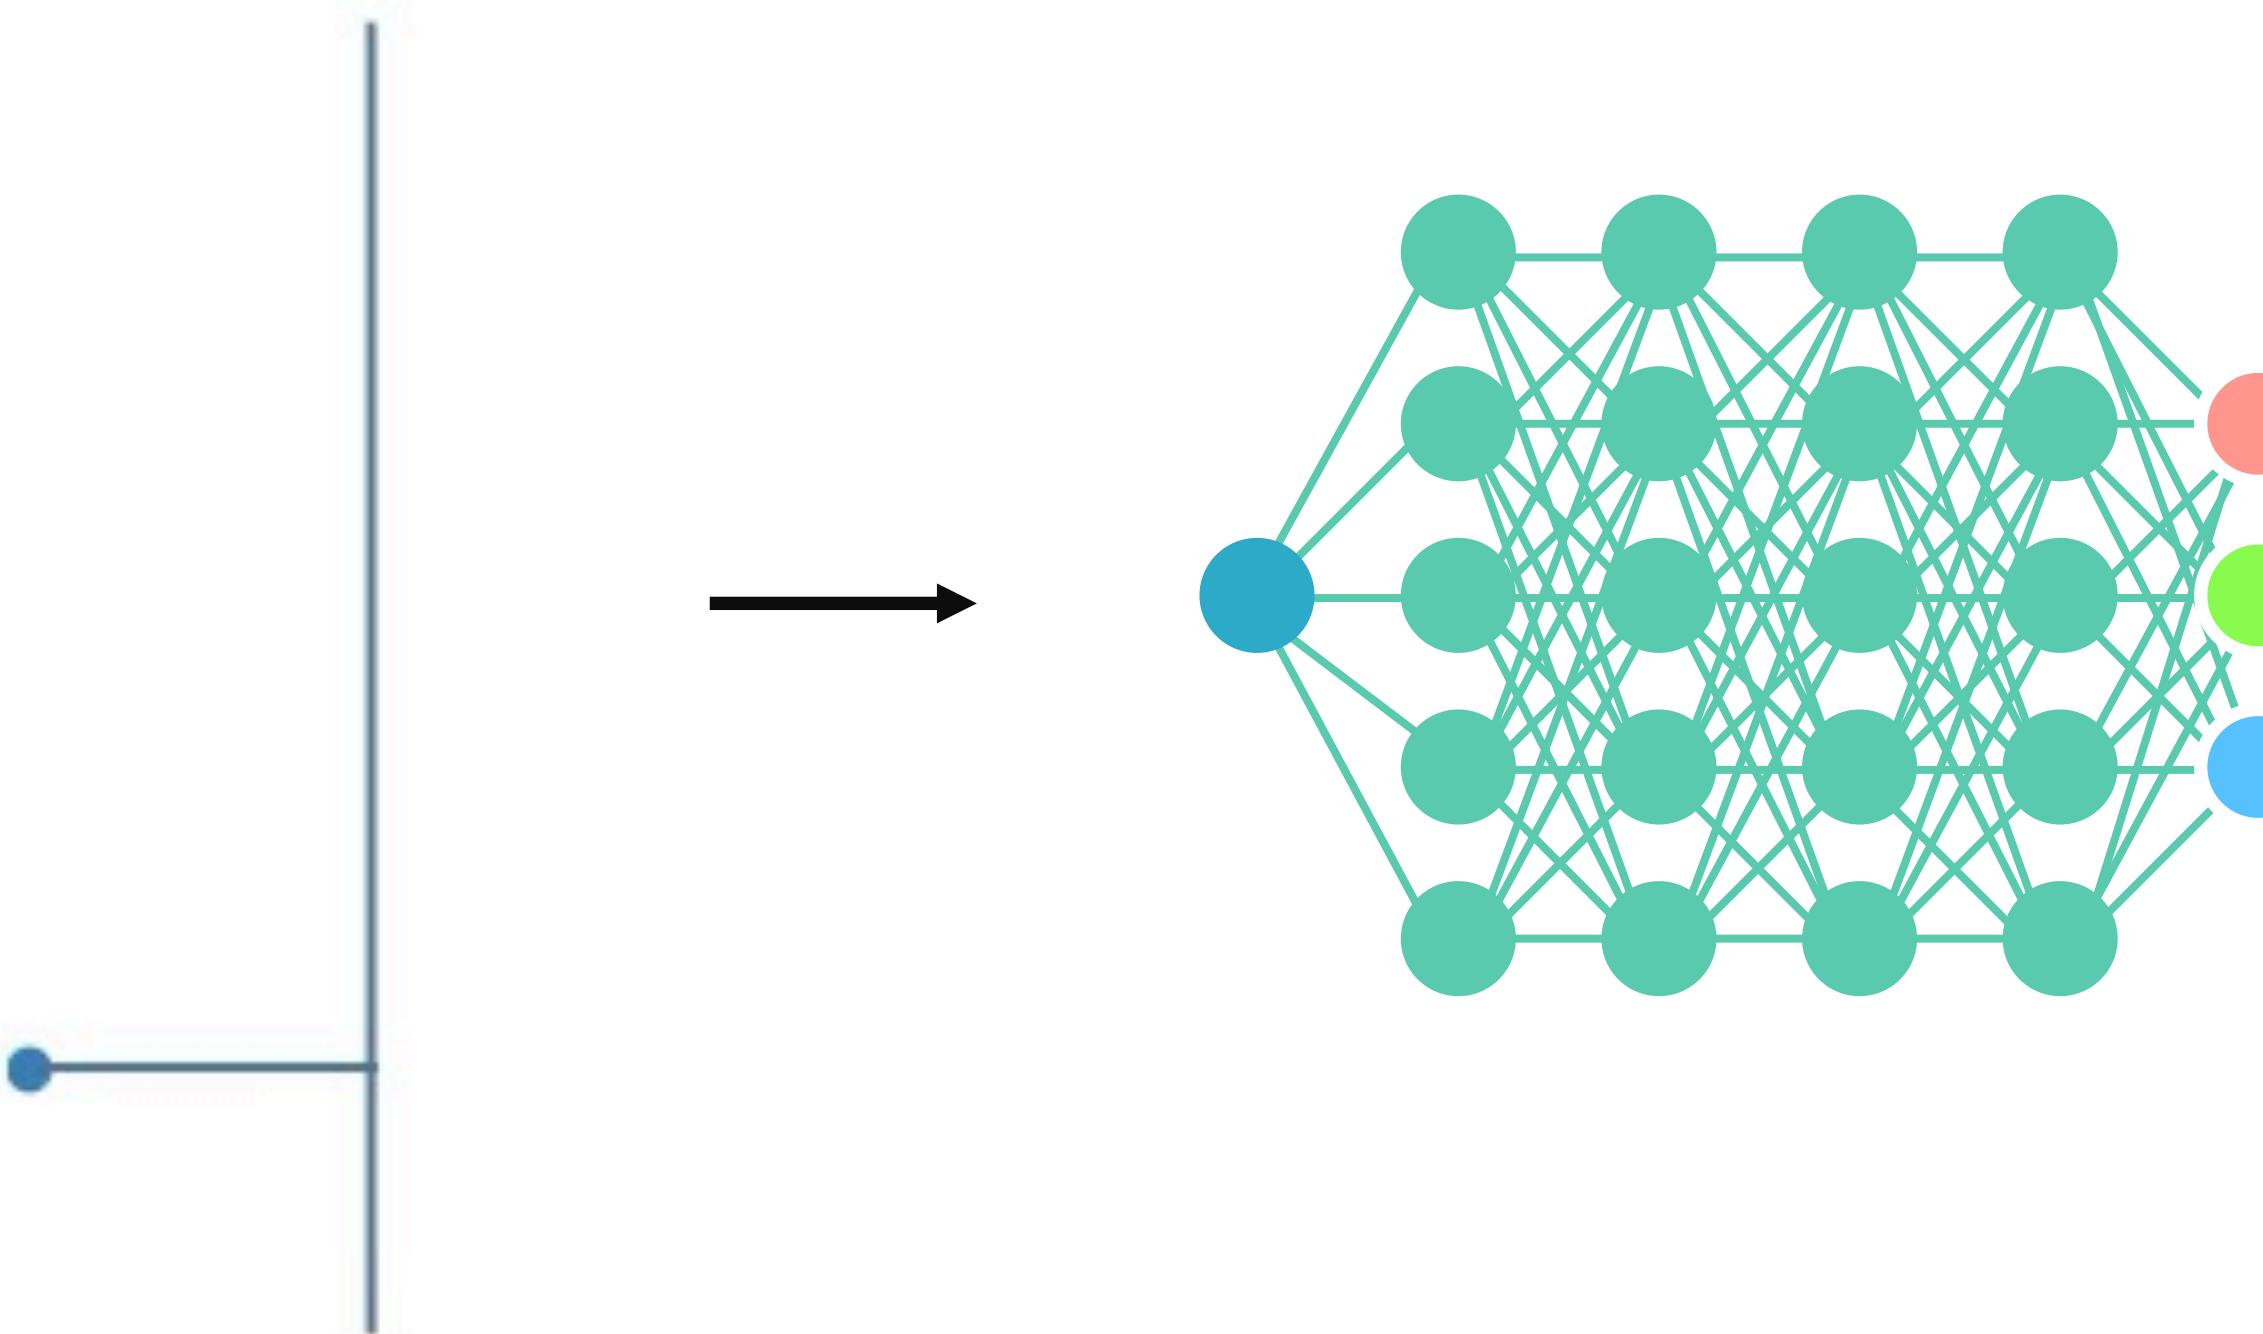
MLP output



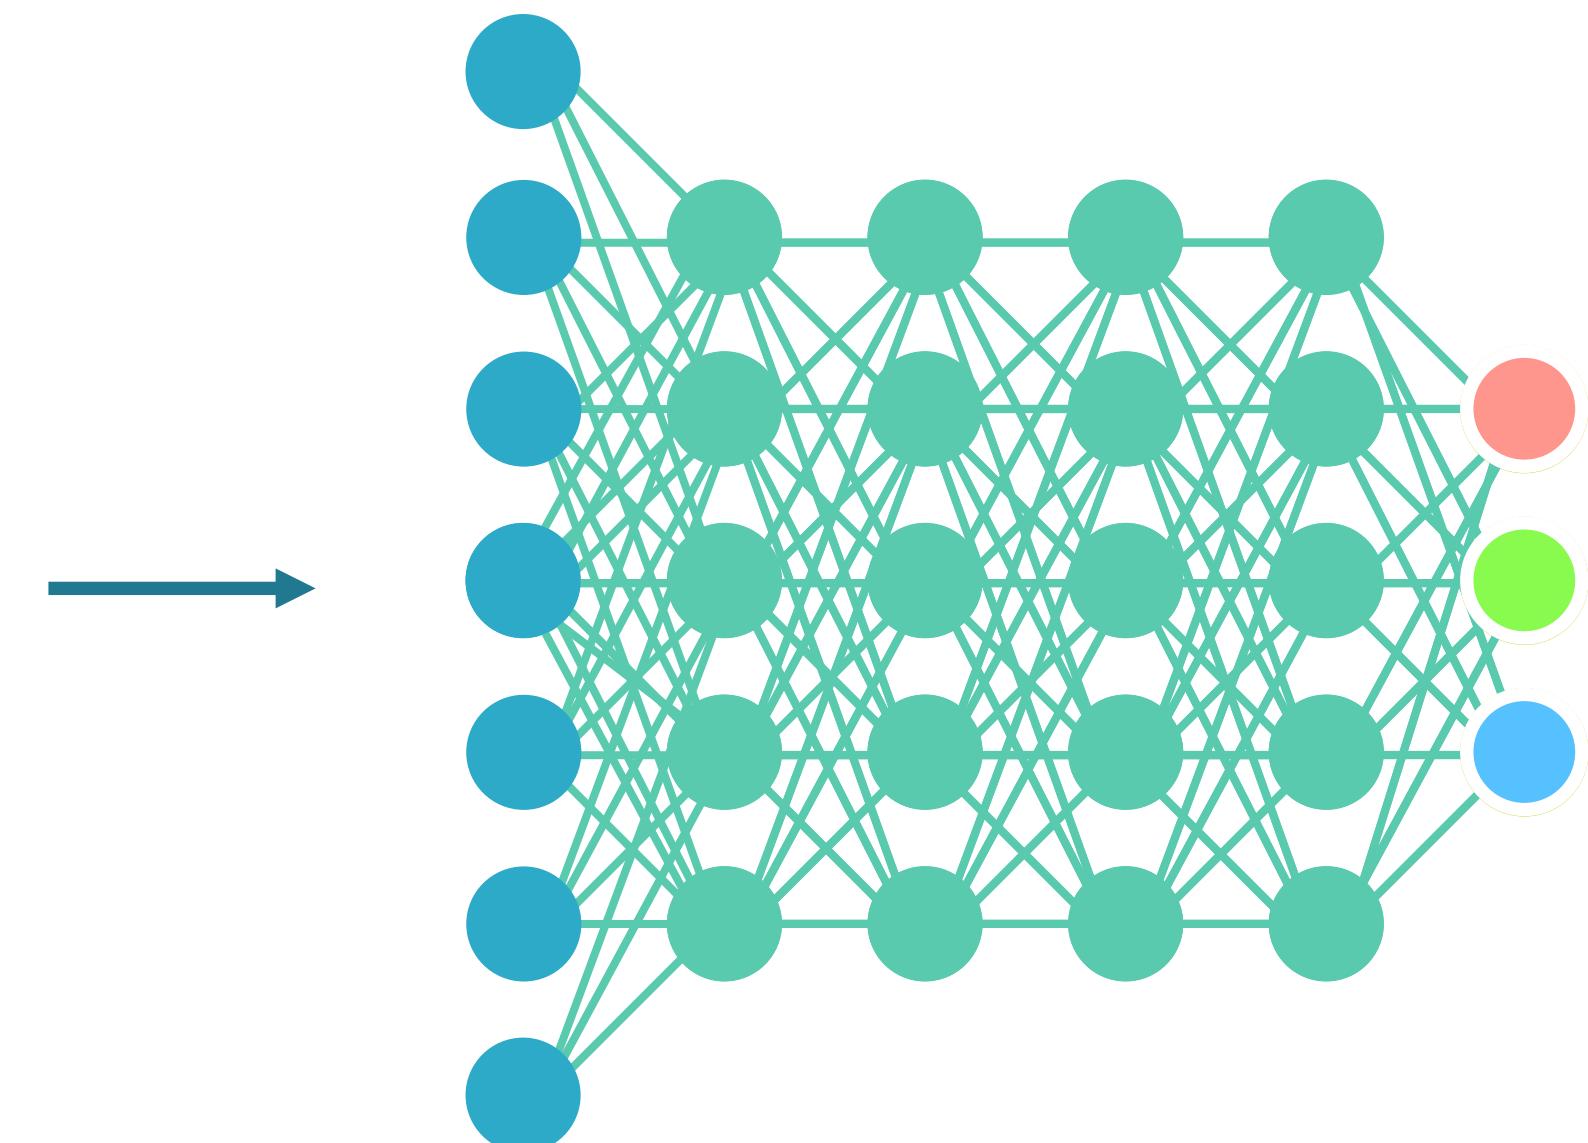
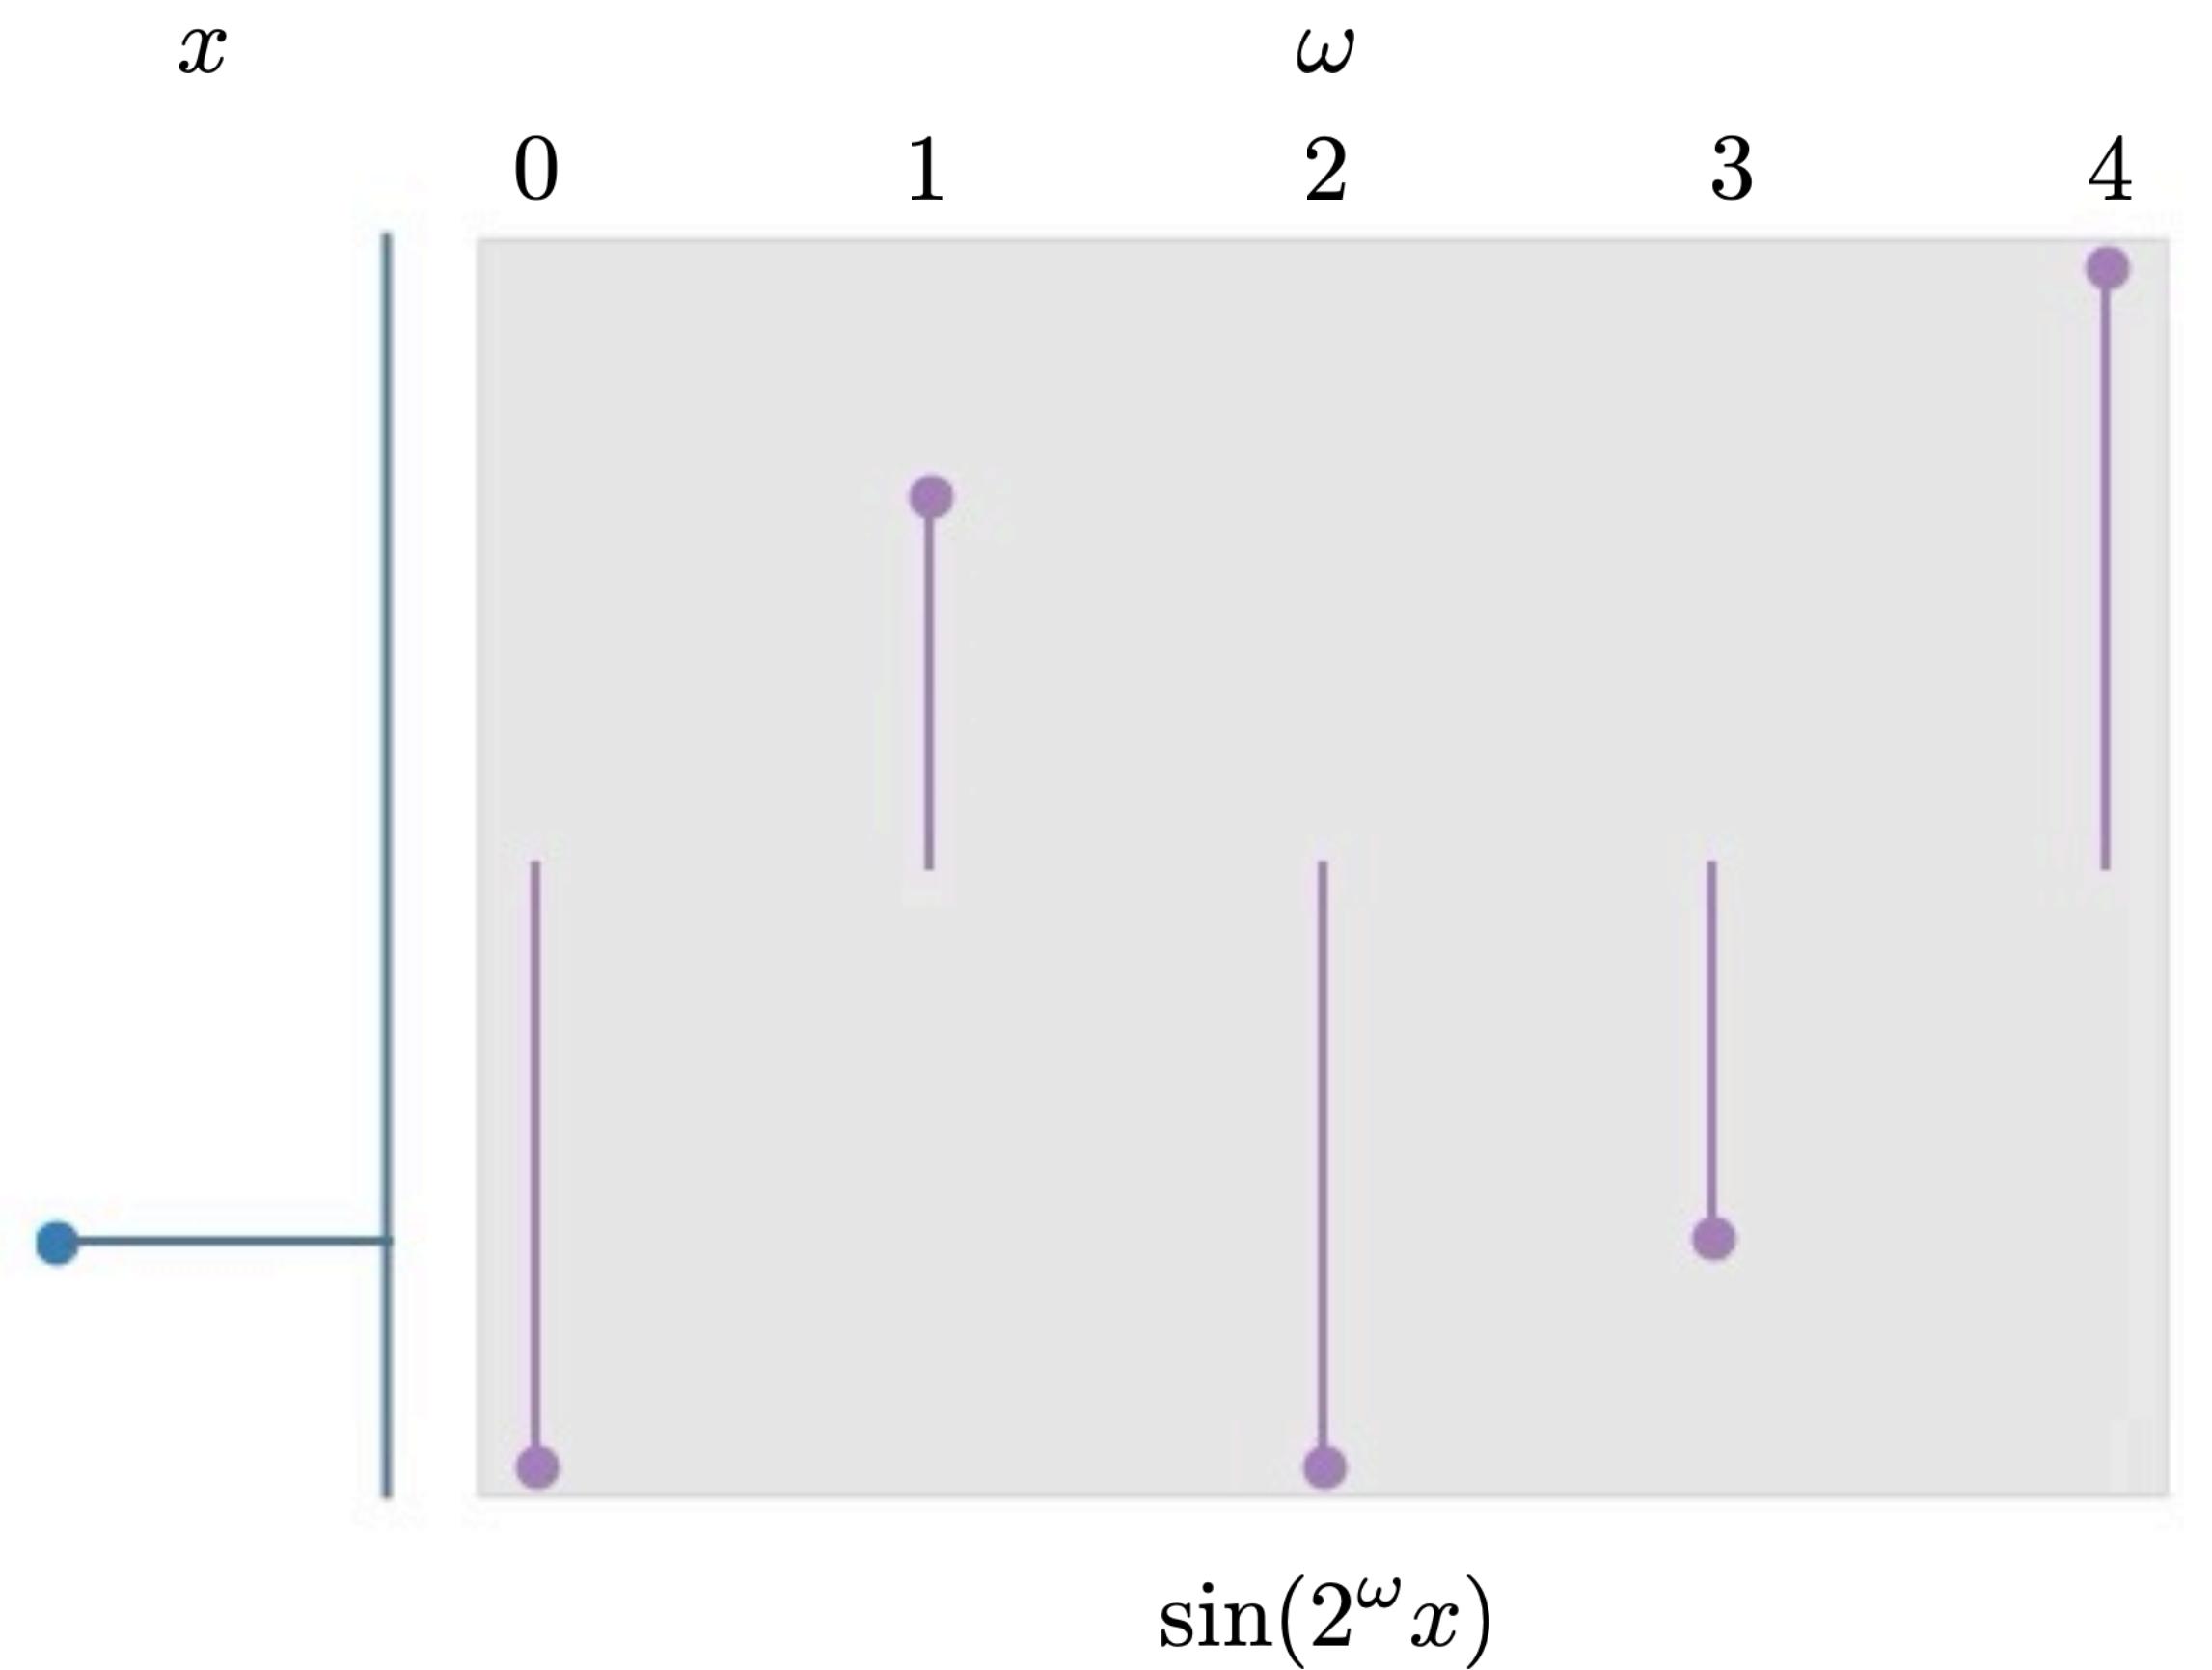
Supervision image

Standard input

$x$

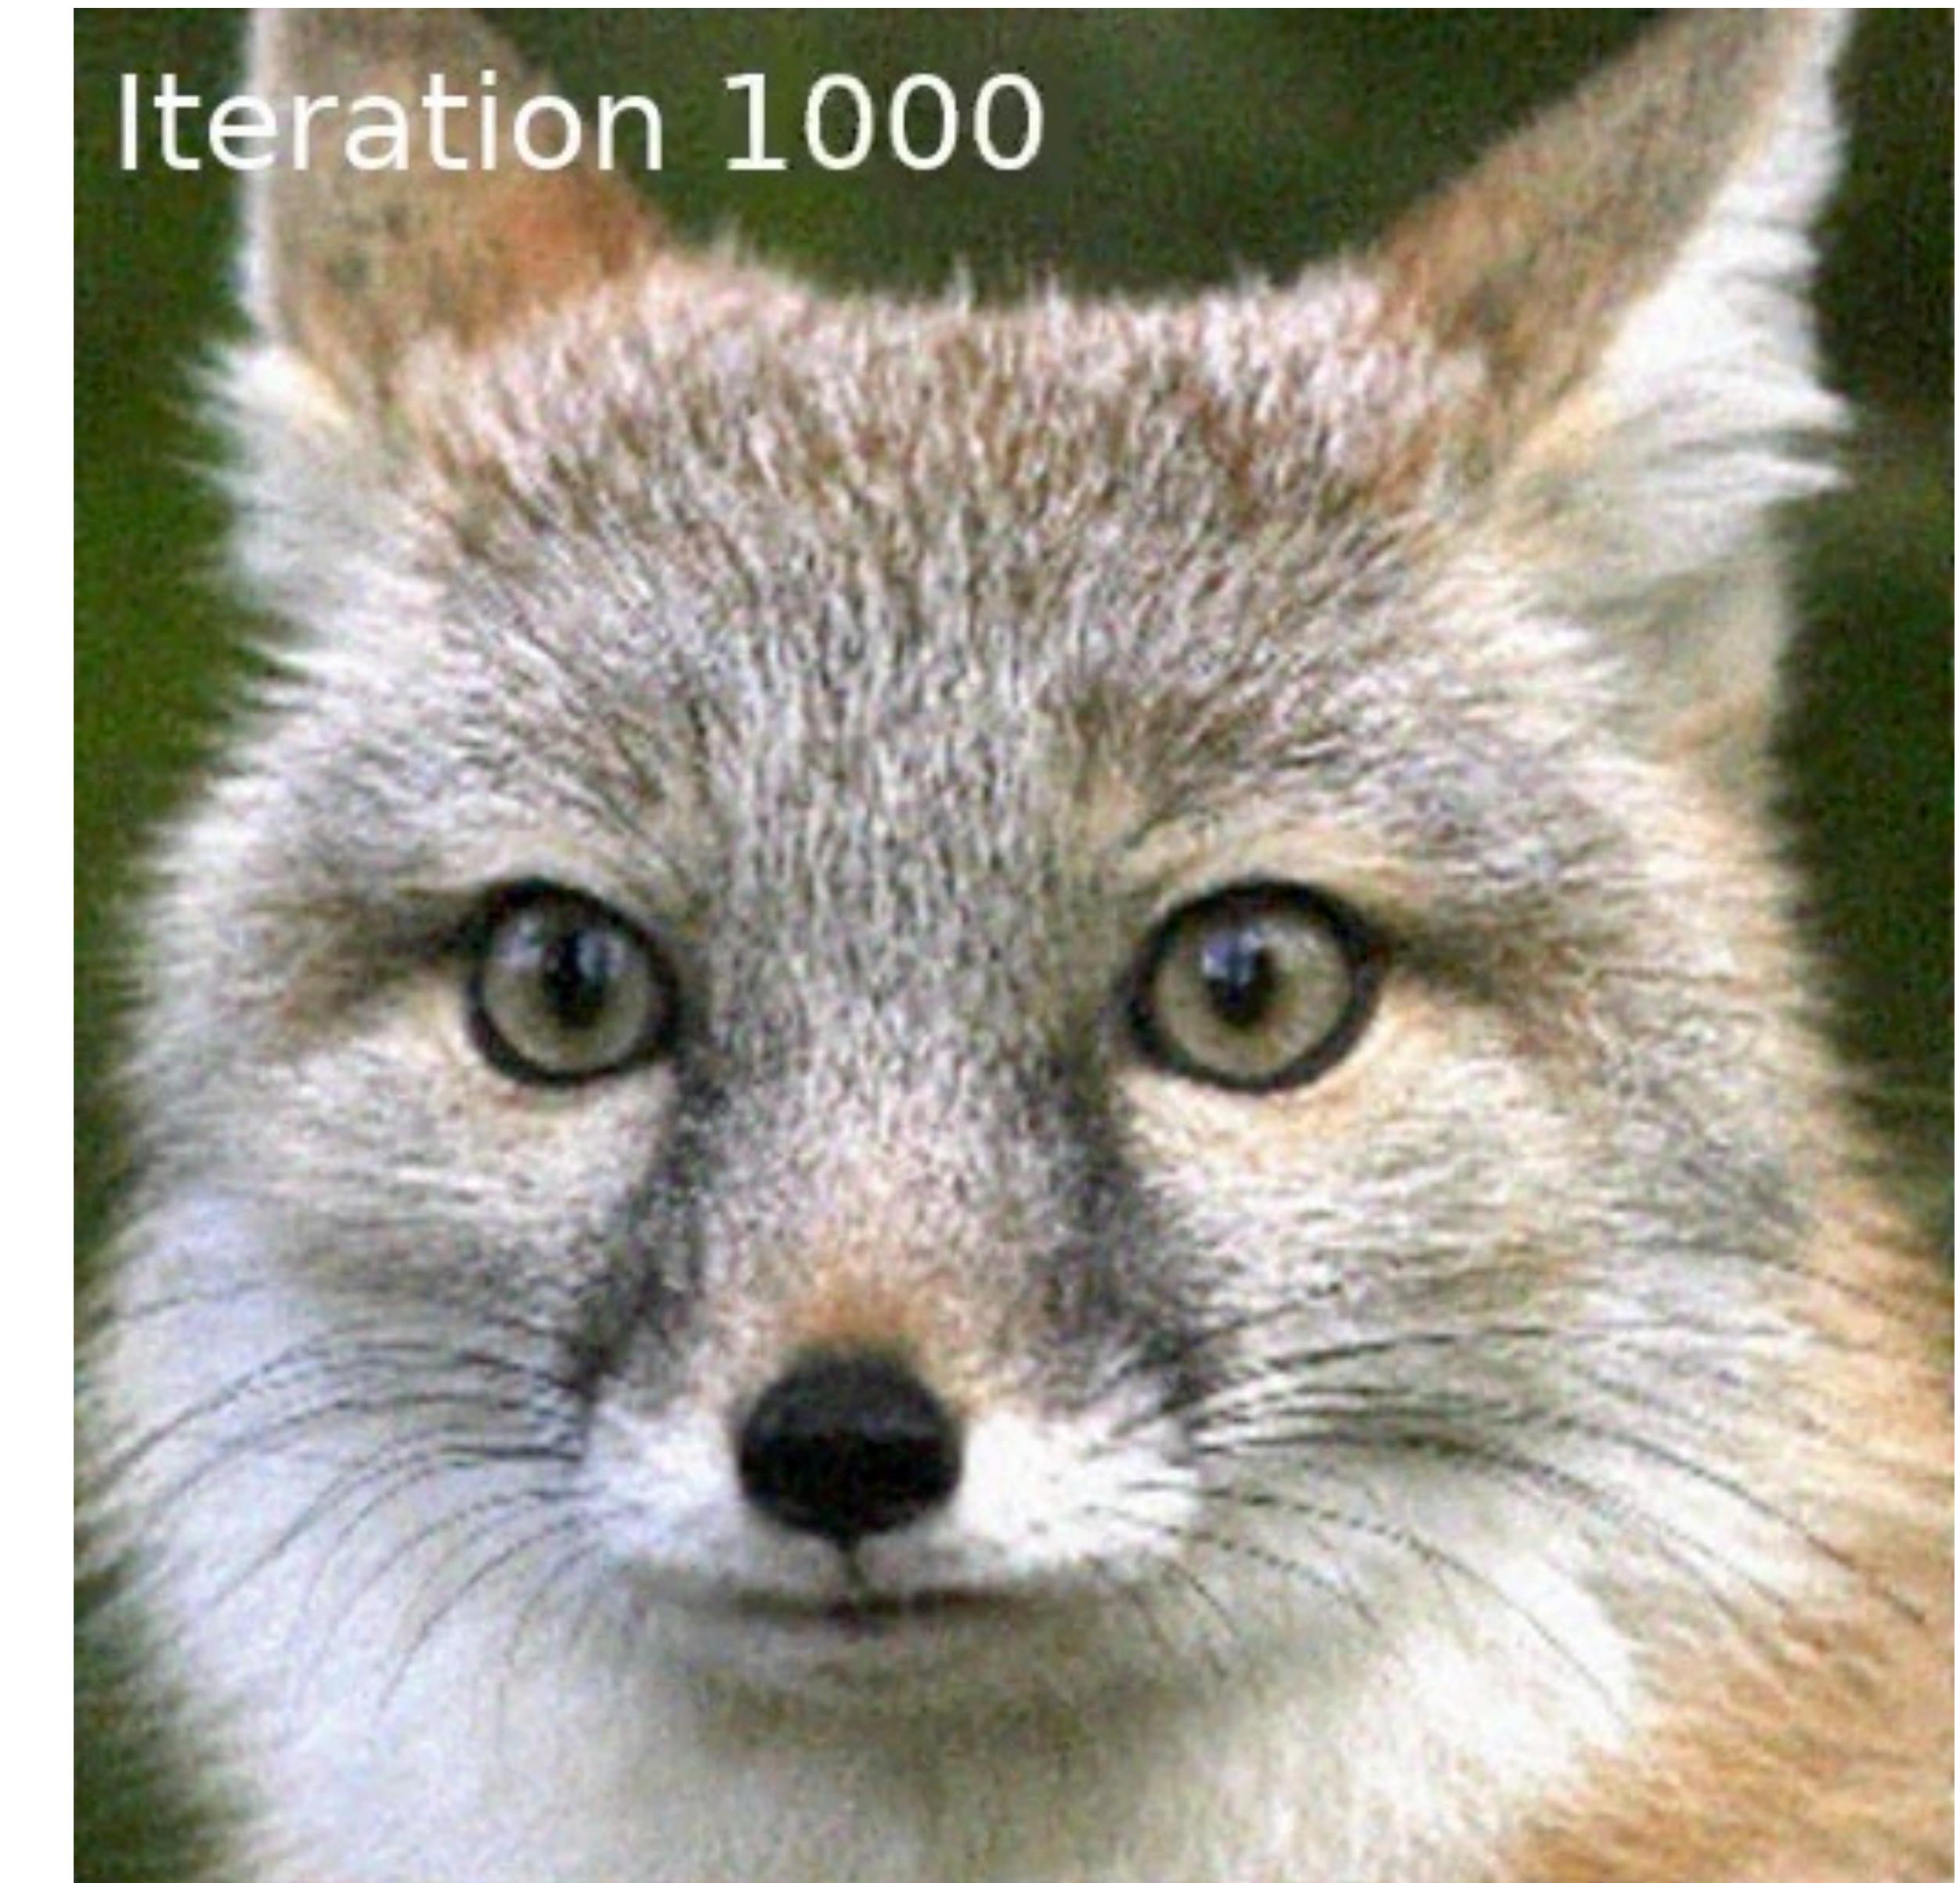


## Standard input      Positionally Encoded input



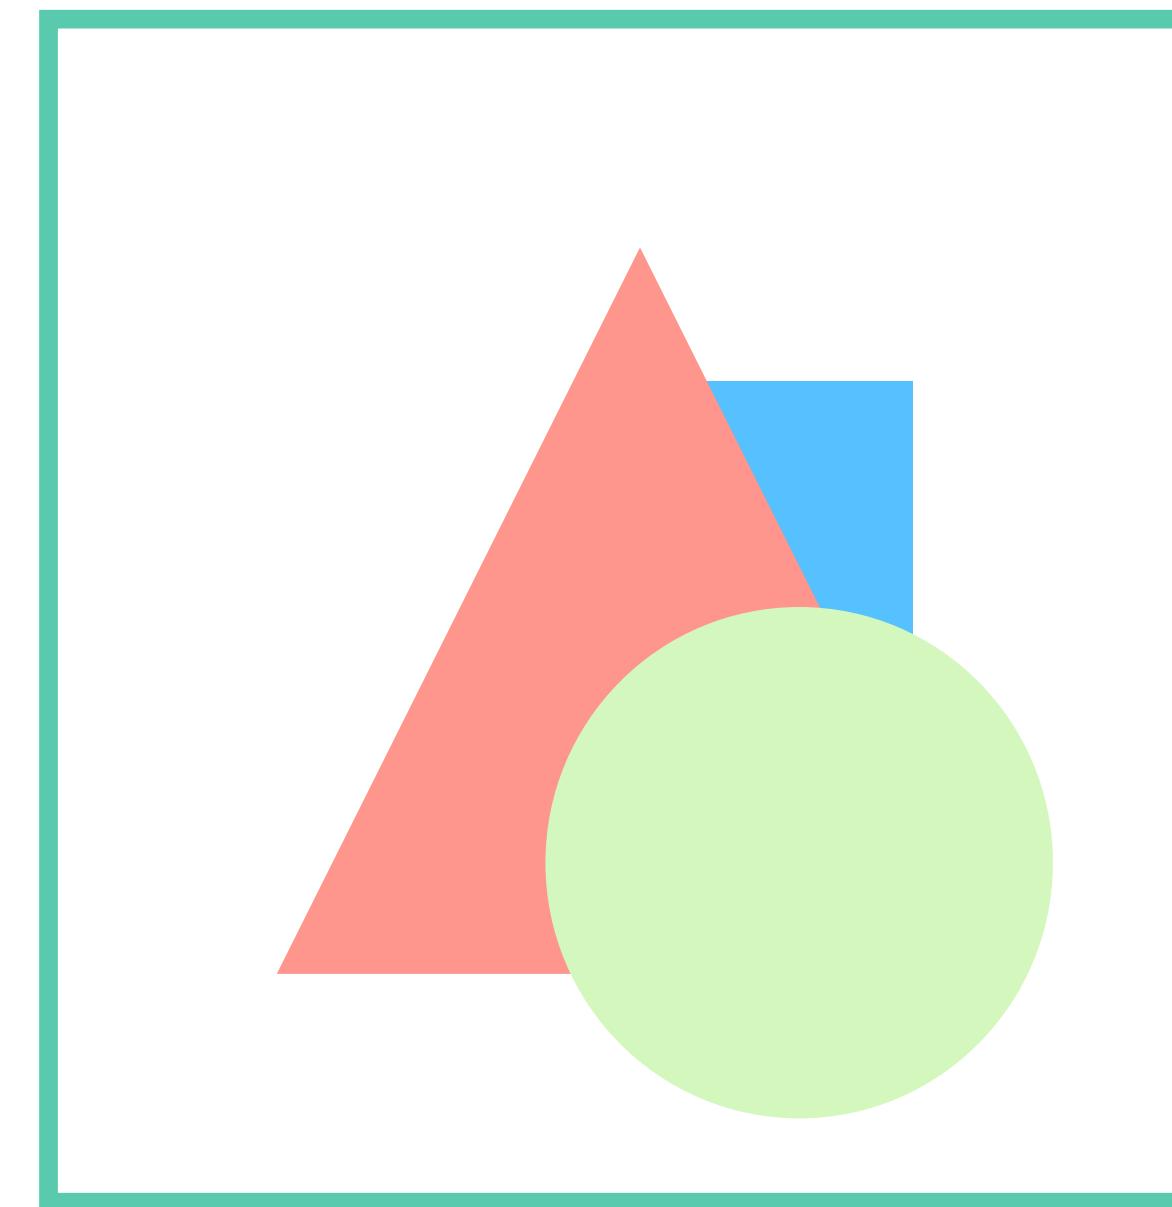


Standard MLP



MLP with Fourier features

# Why does positional encoding help?

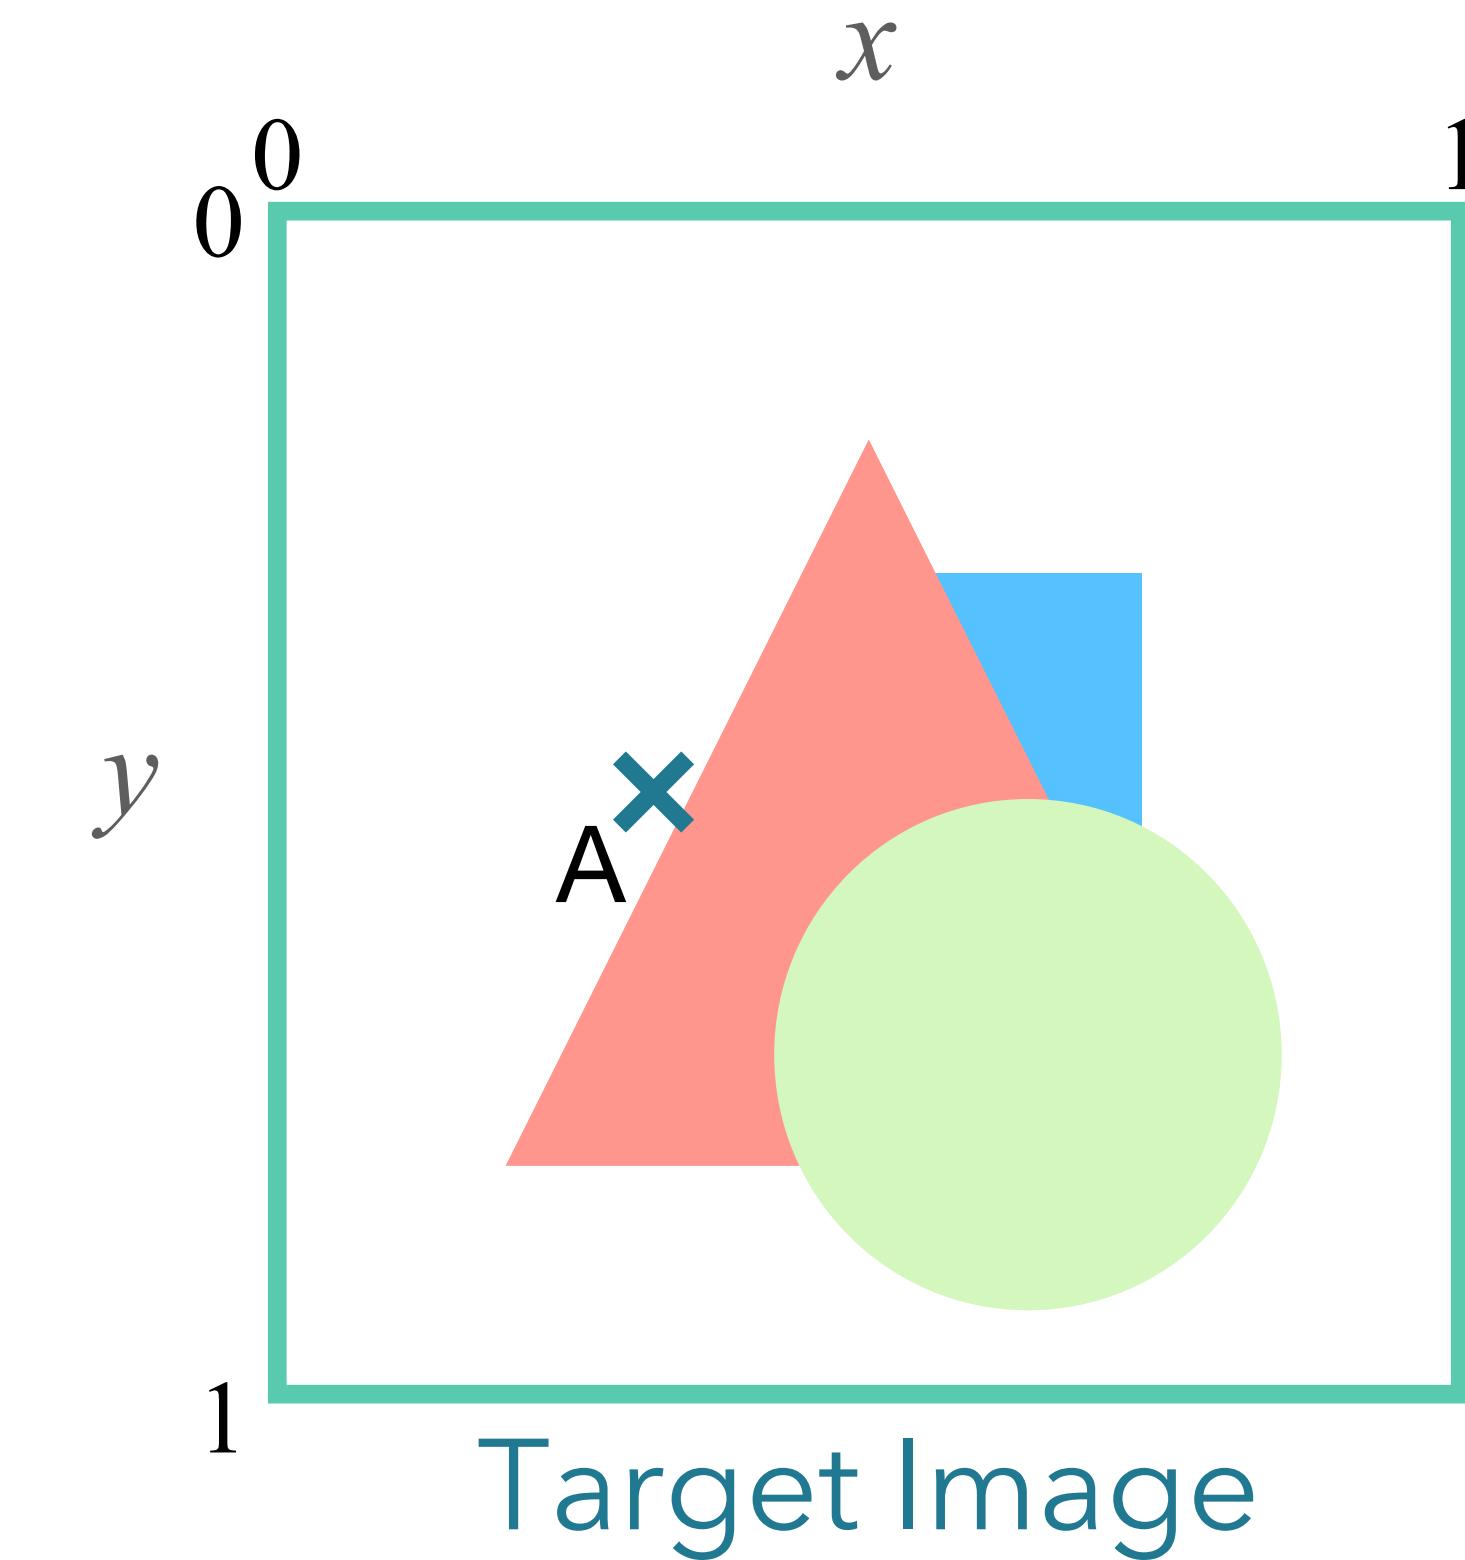


Target Image

# Why does positional encoding help?

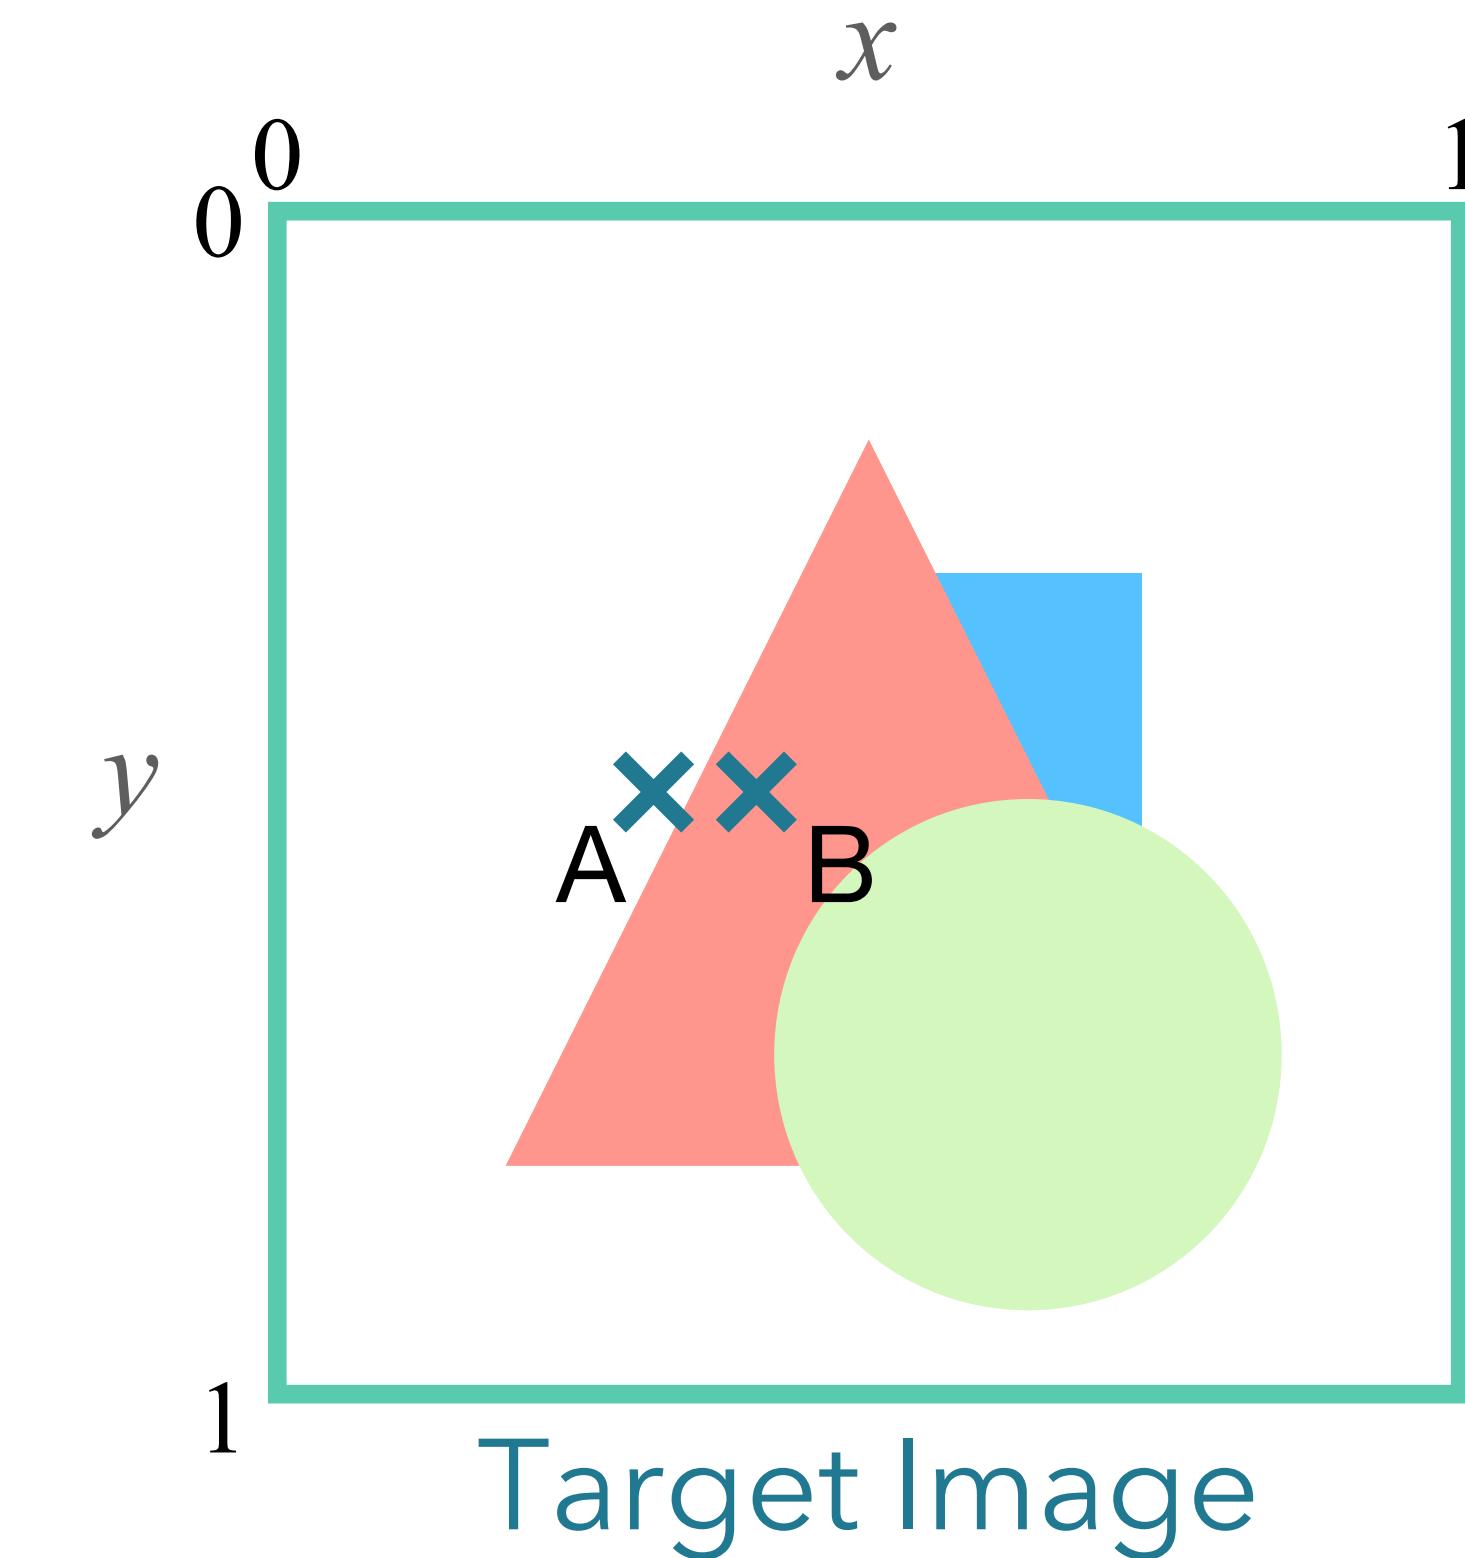
Input                      Target

	$x$	$y$
A	.36	.5

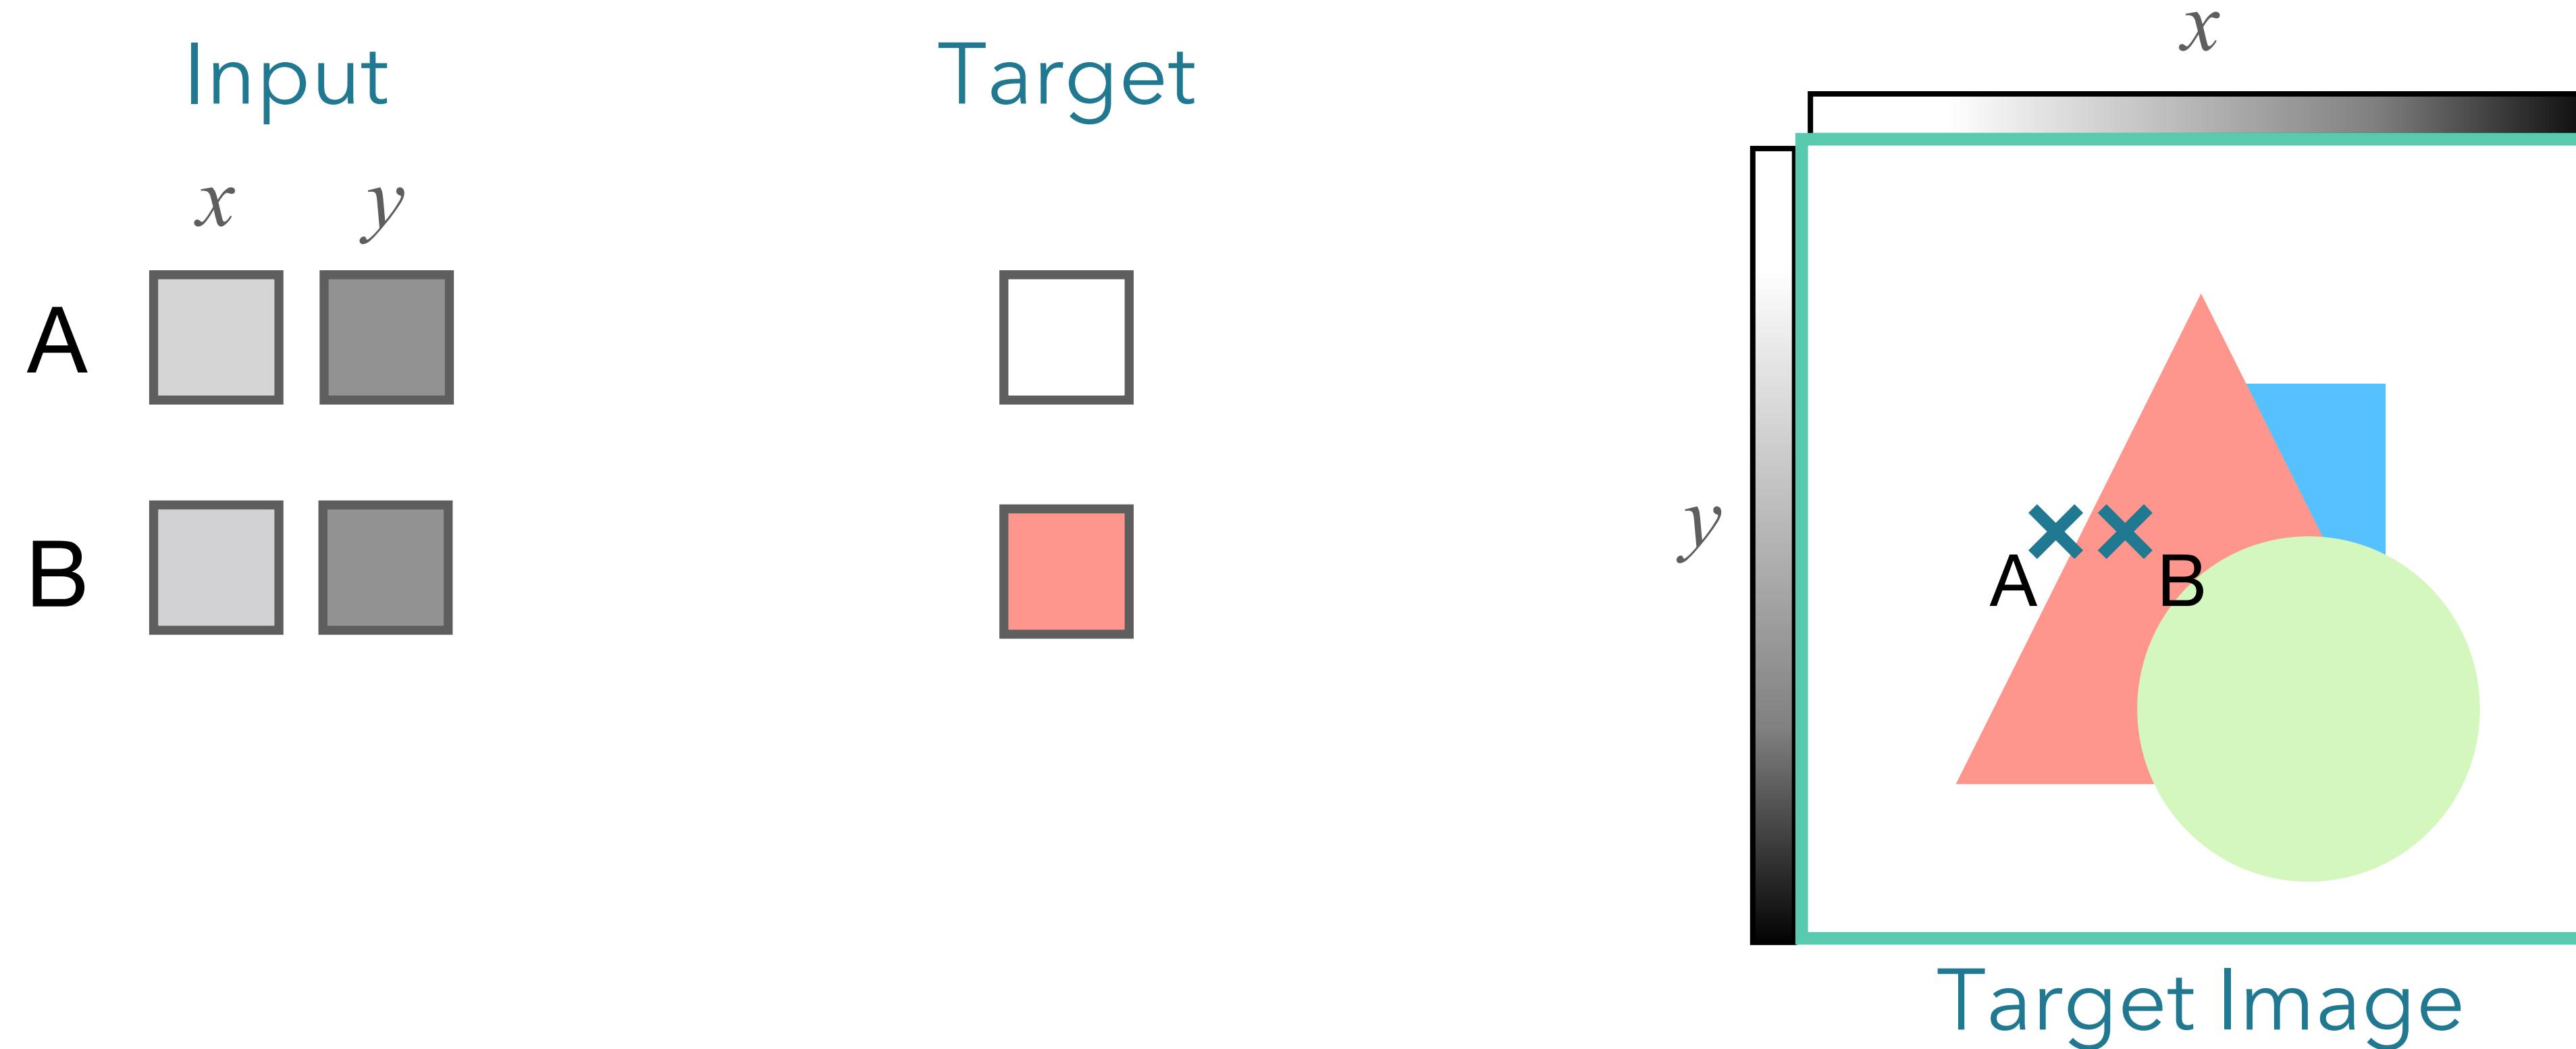


# Why does positional encoding help?

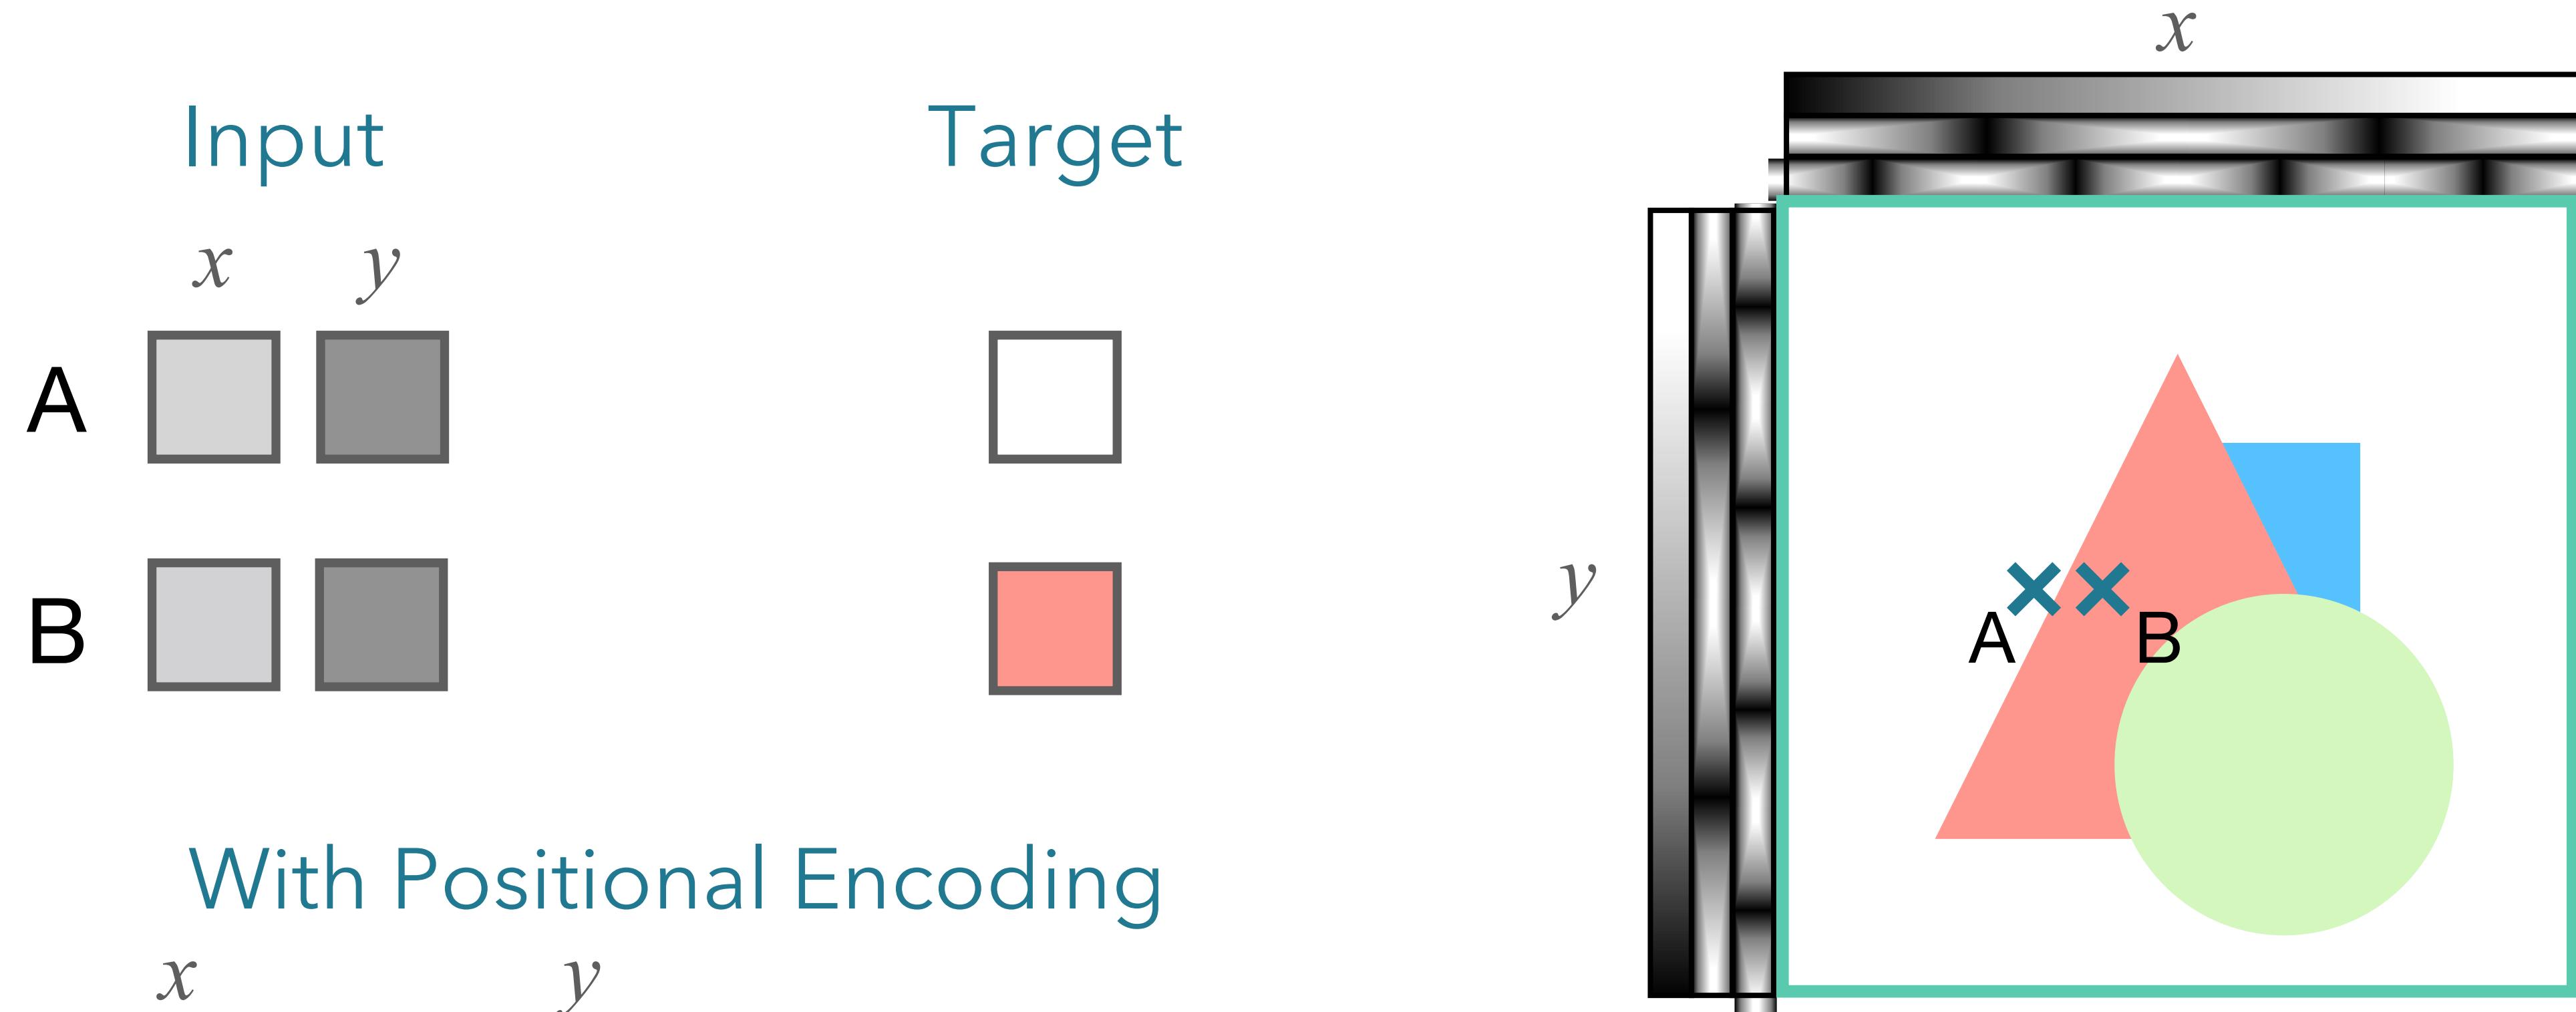
Input		Target
	$x$	$y$
A	.36	.5
B	.38	.5



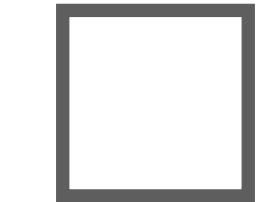
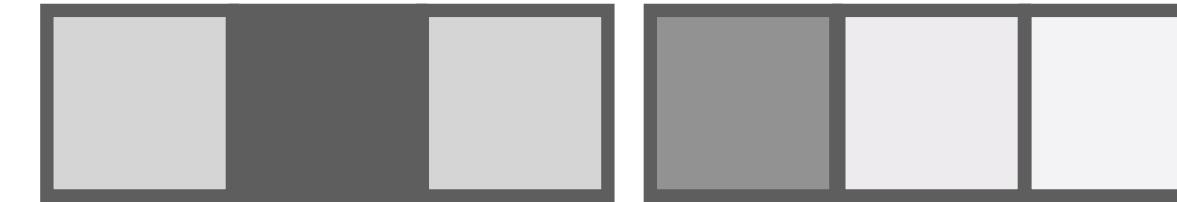
# Why does positional encoding help?



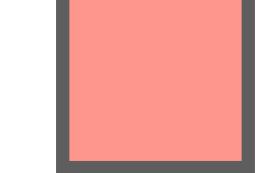
# Why does positional encoding help?



A



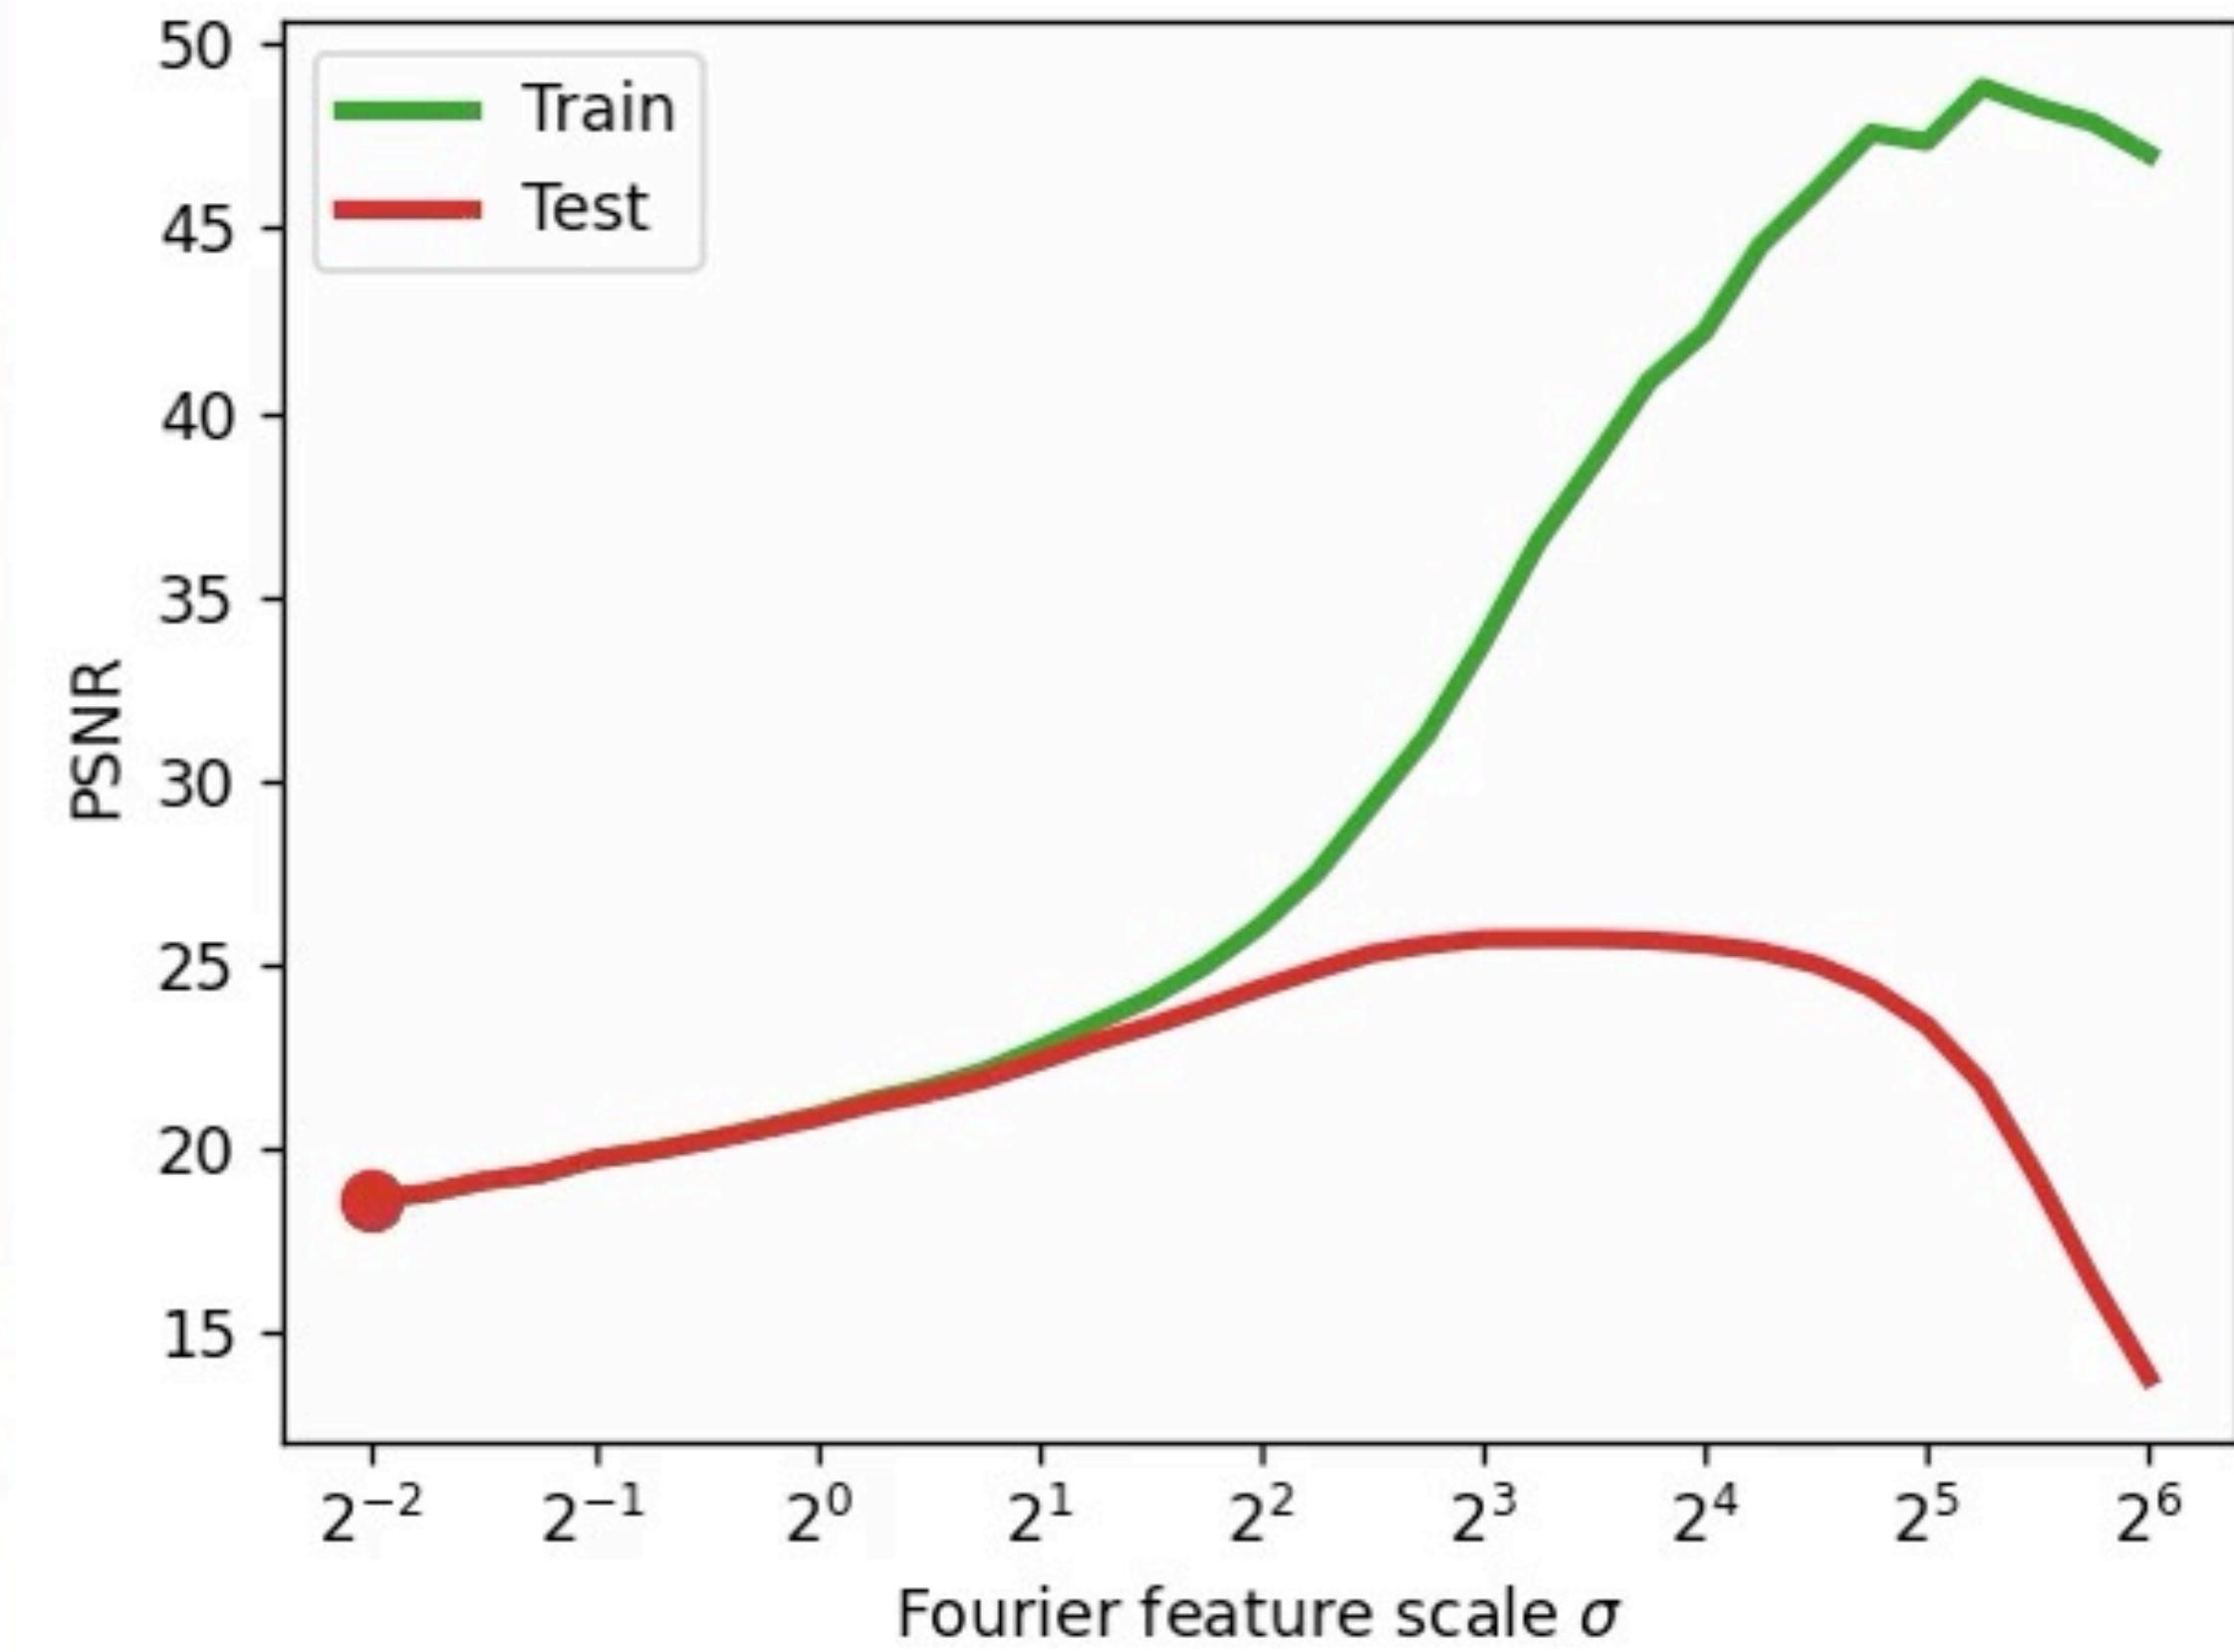
B



# Performance depends on max encoding frequency



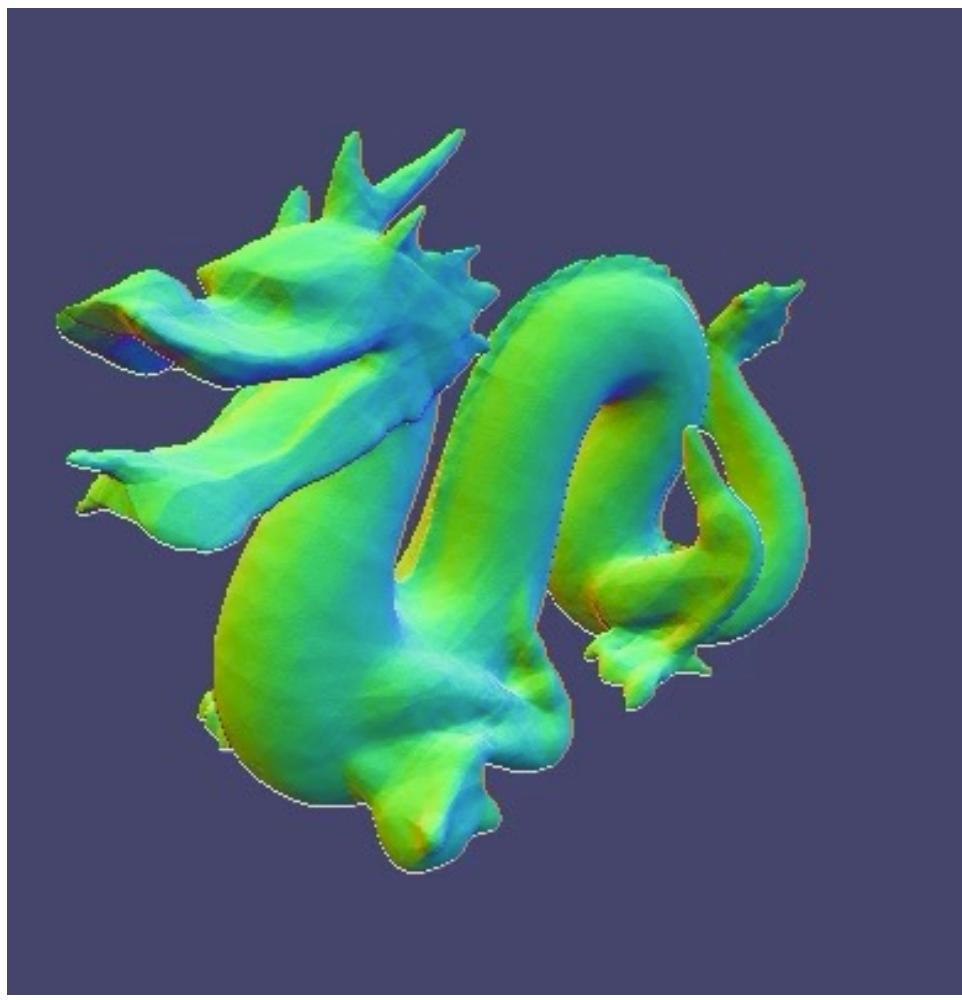
Network output



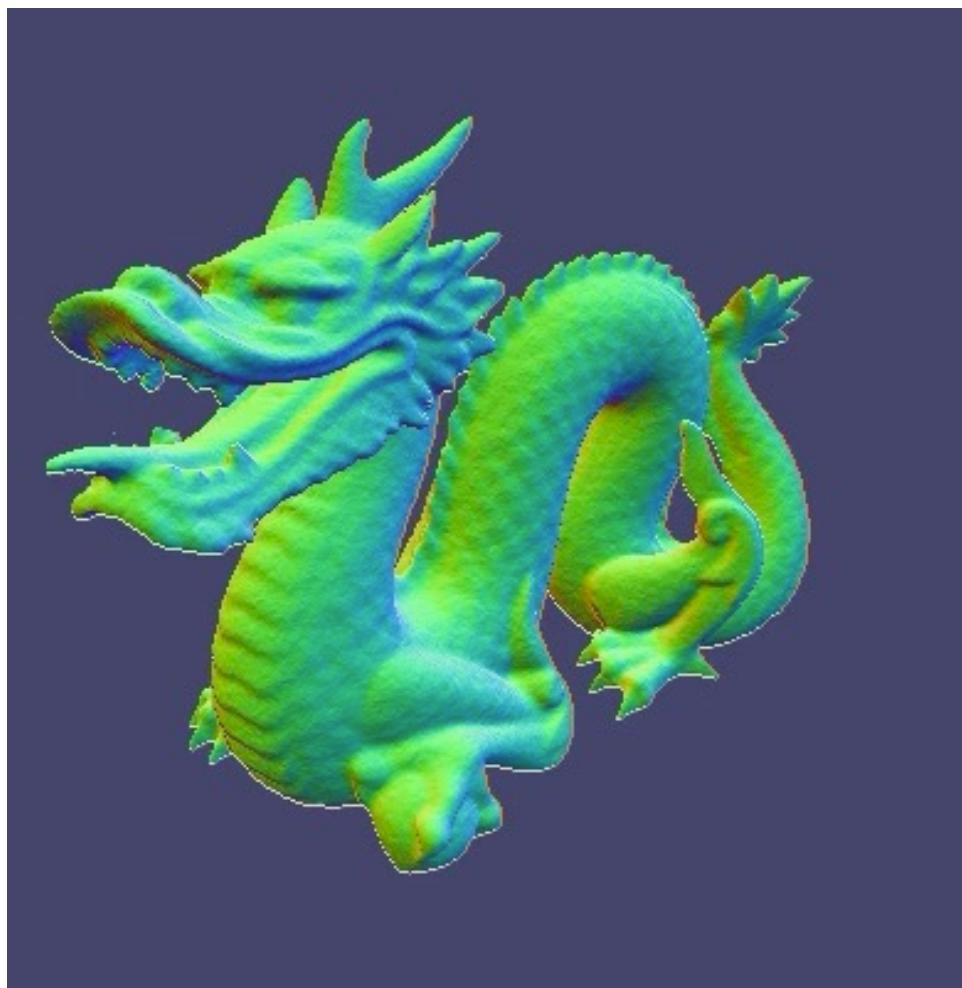
Performance vs. scale value

Coordinate-based MLPs can replace any low-dimensional array

Without Encoding



With Encoding



3D Shape

# NeRF with and without positional encoding

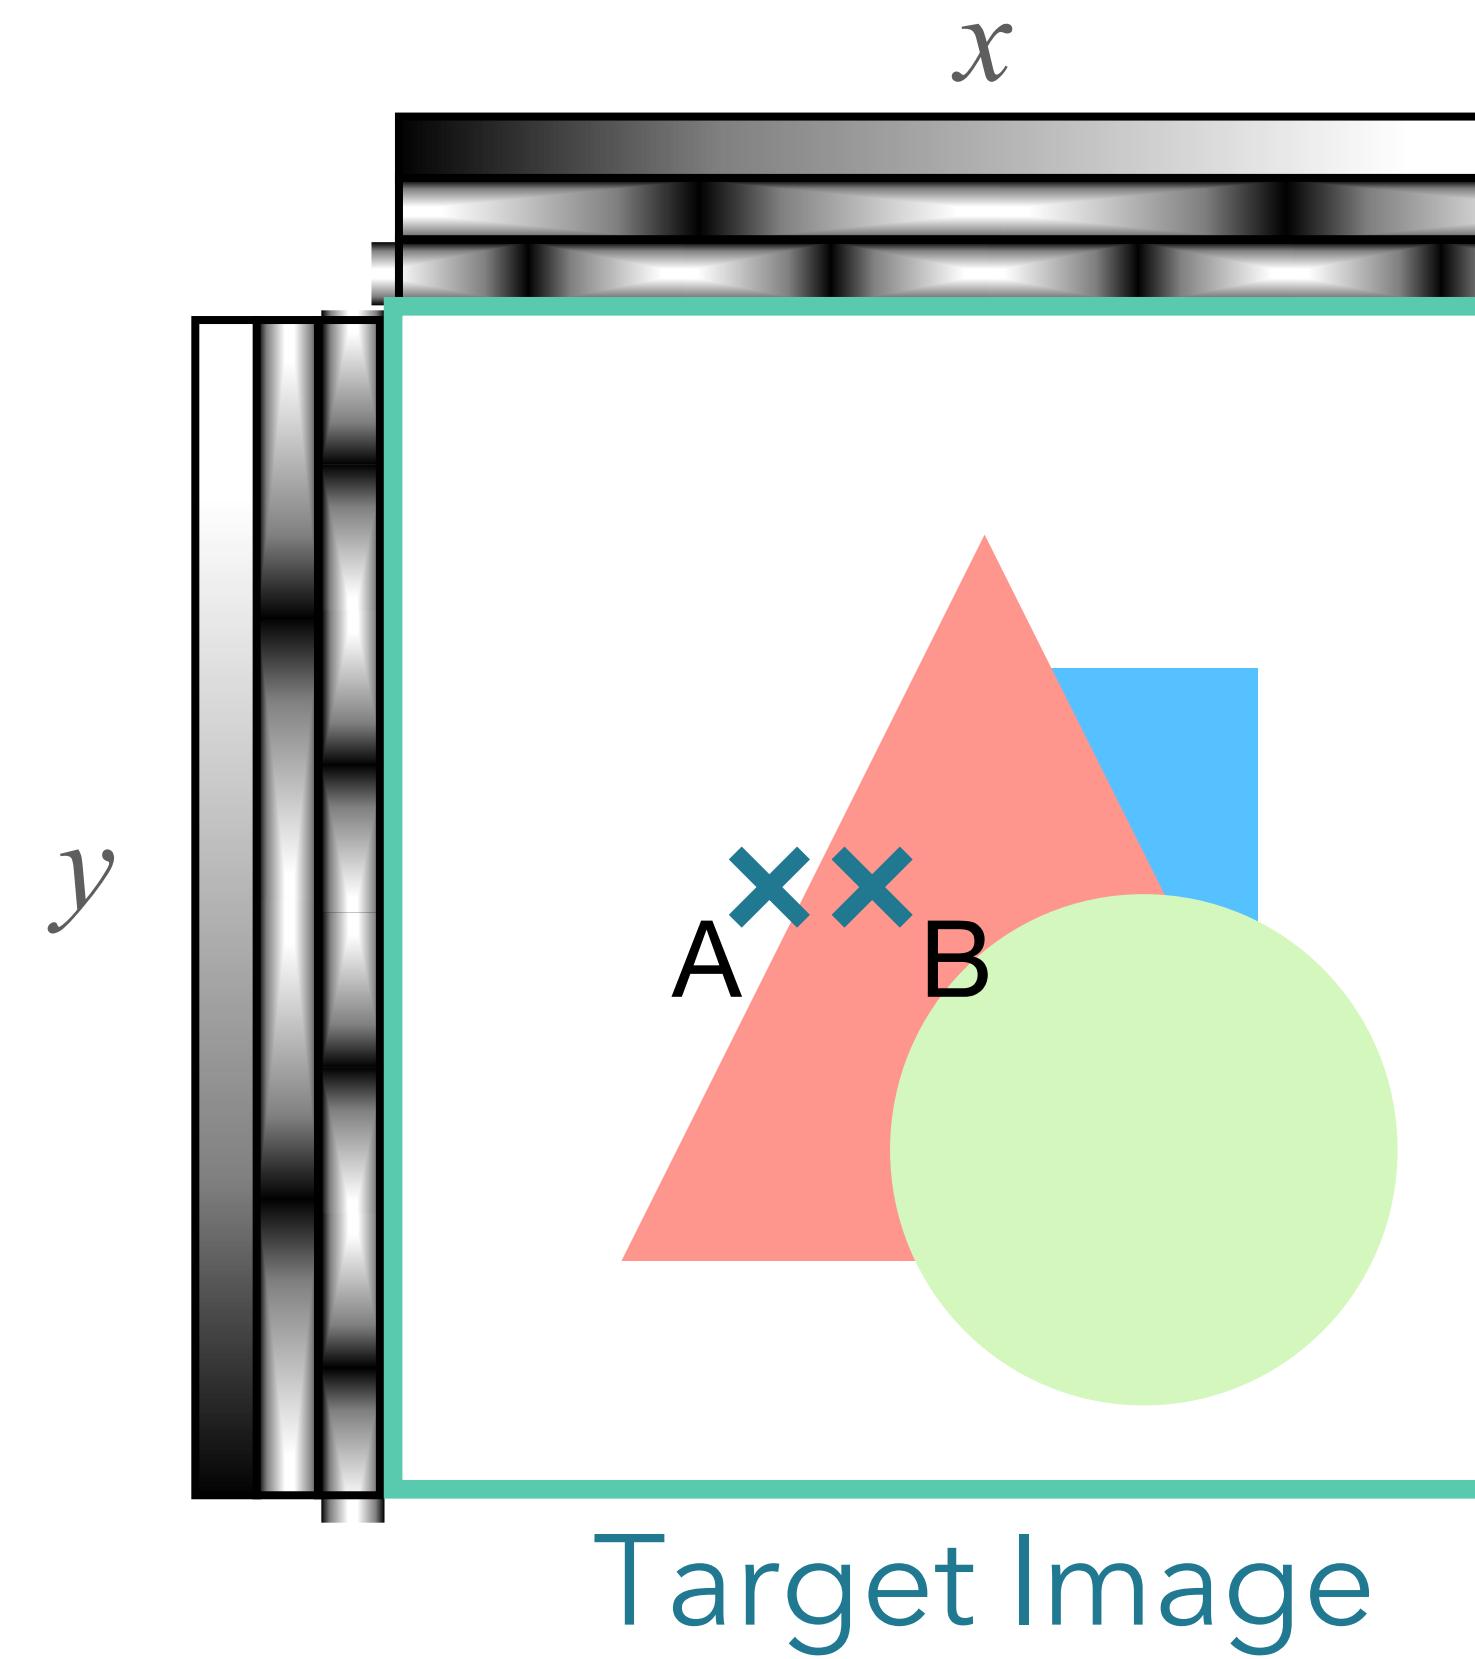


NeRF (Naive)

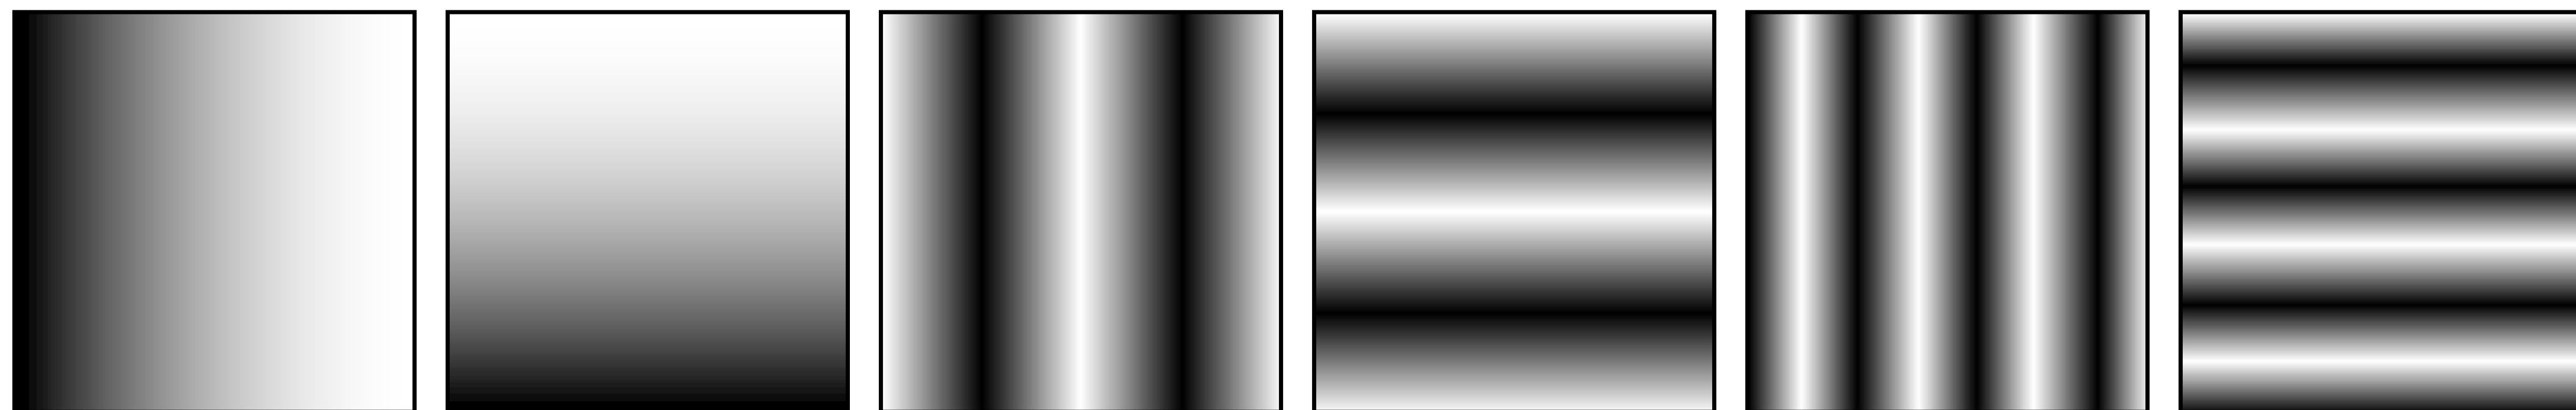
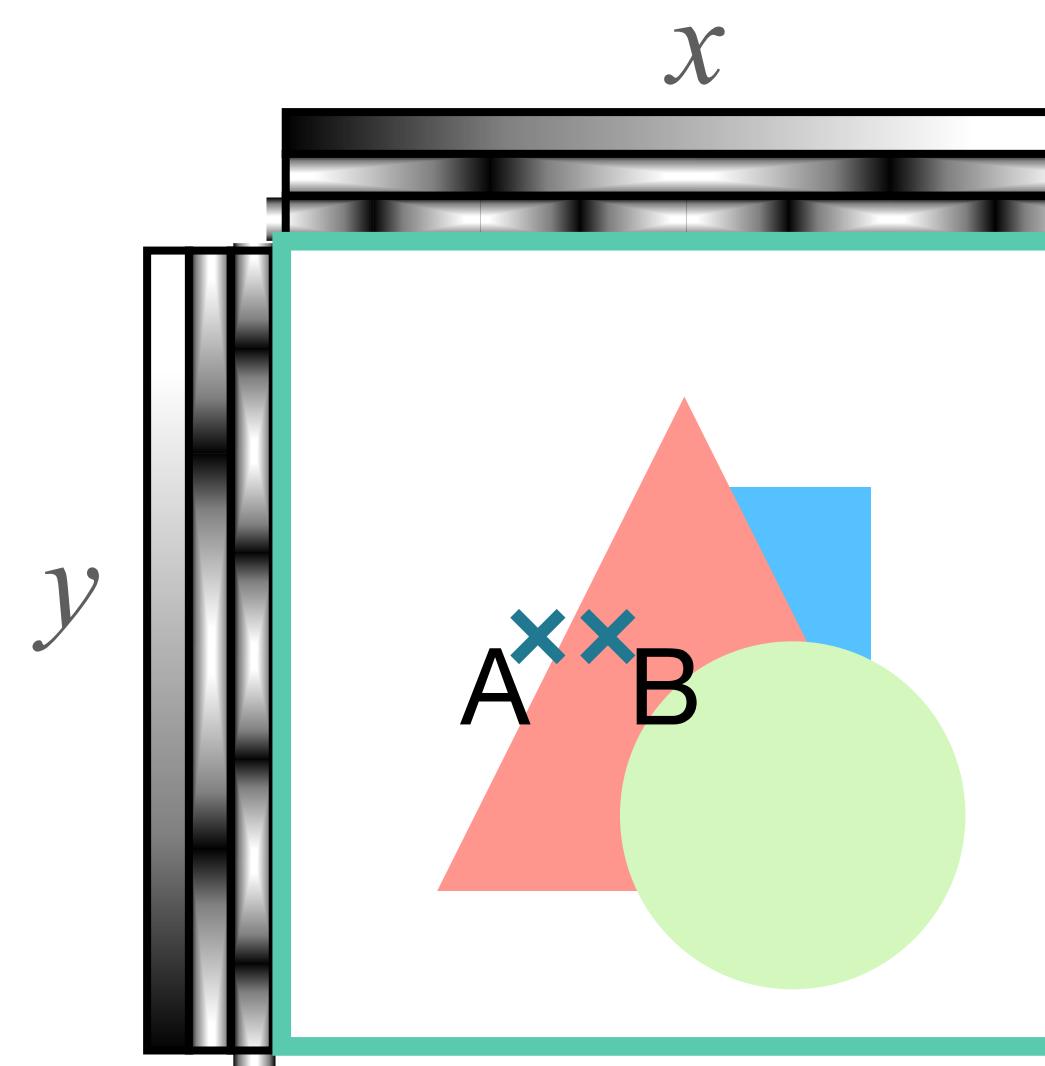


NeRF (with positional encoding)

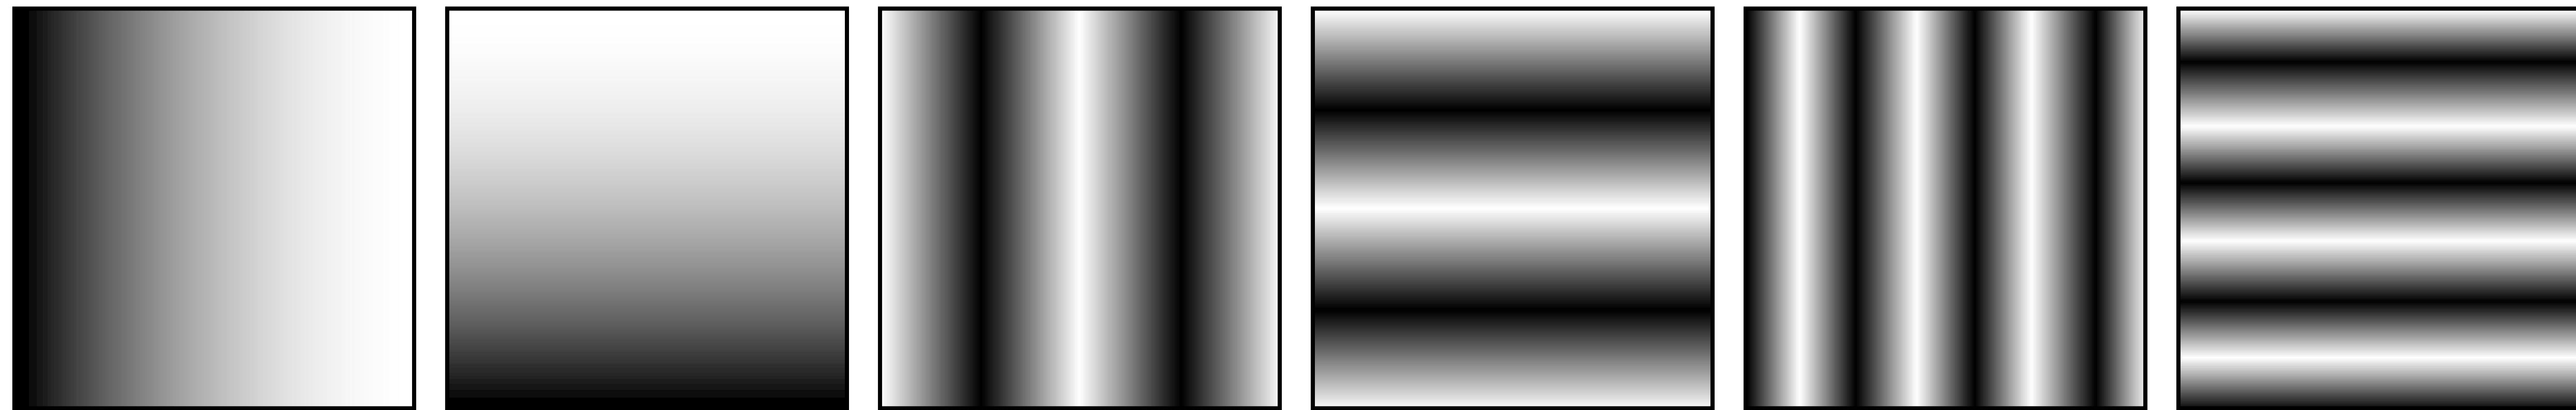
# Other Encoding Considerations



# Other Encoding Considerations

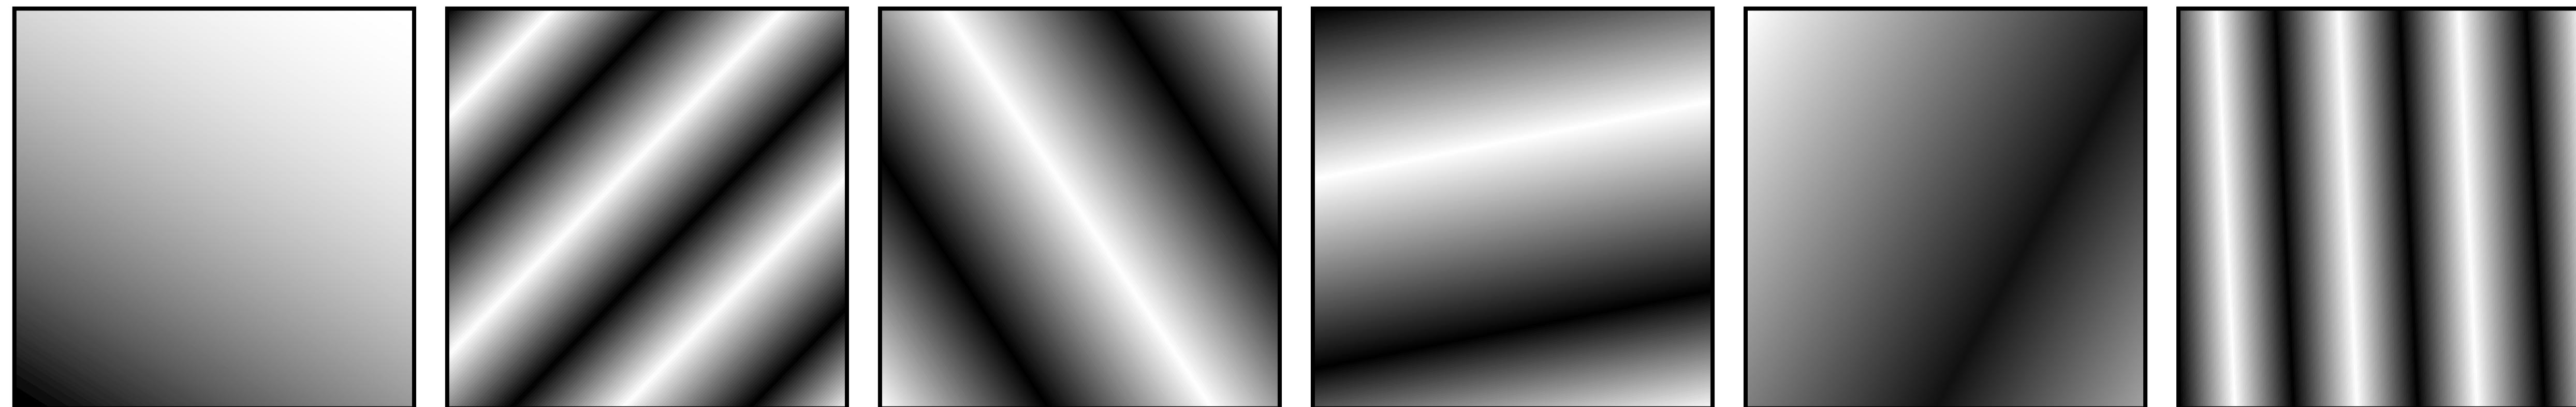


# Other Encoding Considerations



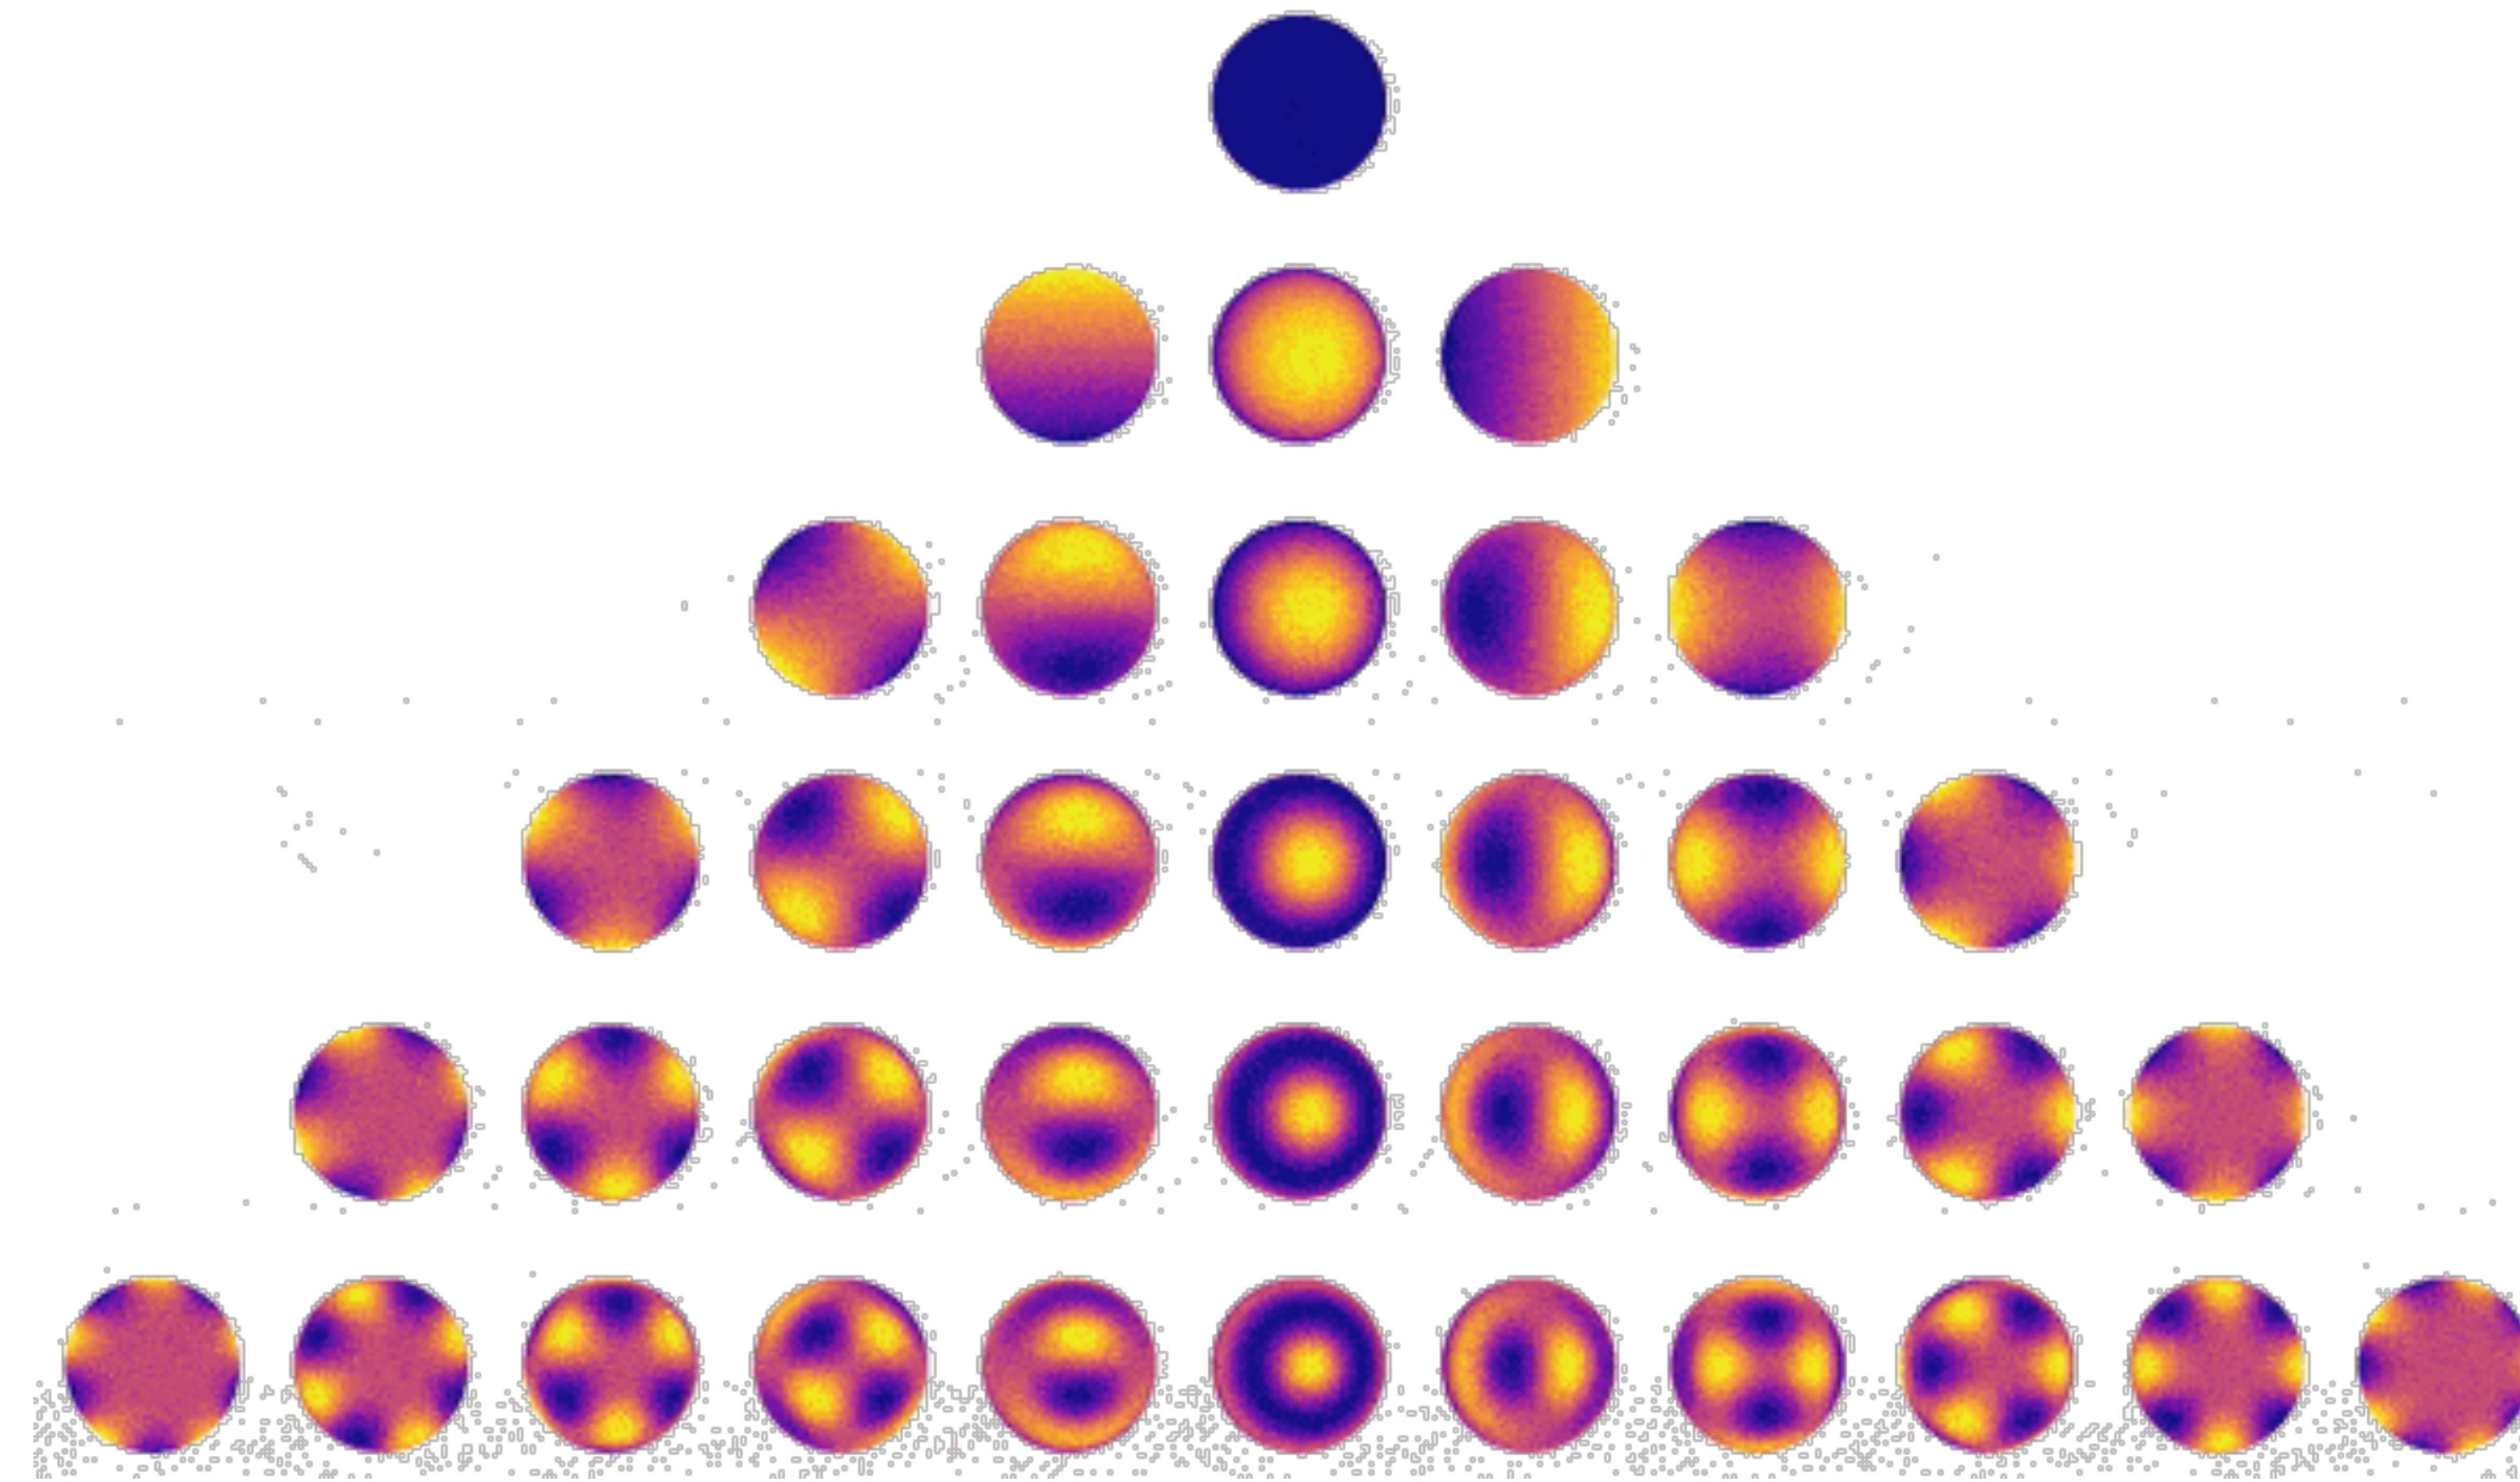
Can lead to "grid" artifacts

# Other Encoding Considerations



Random Fourier Features

# Other Encodings



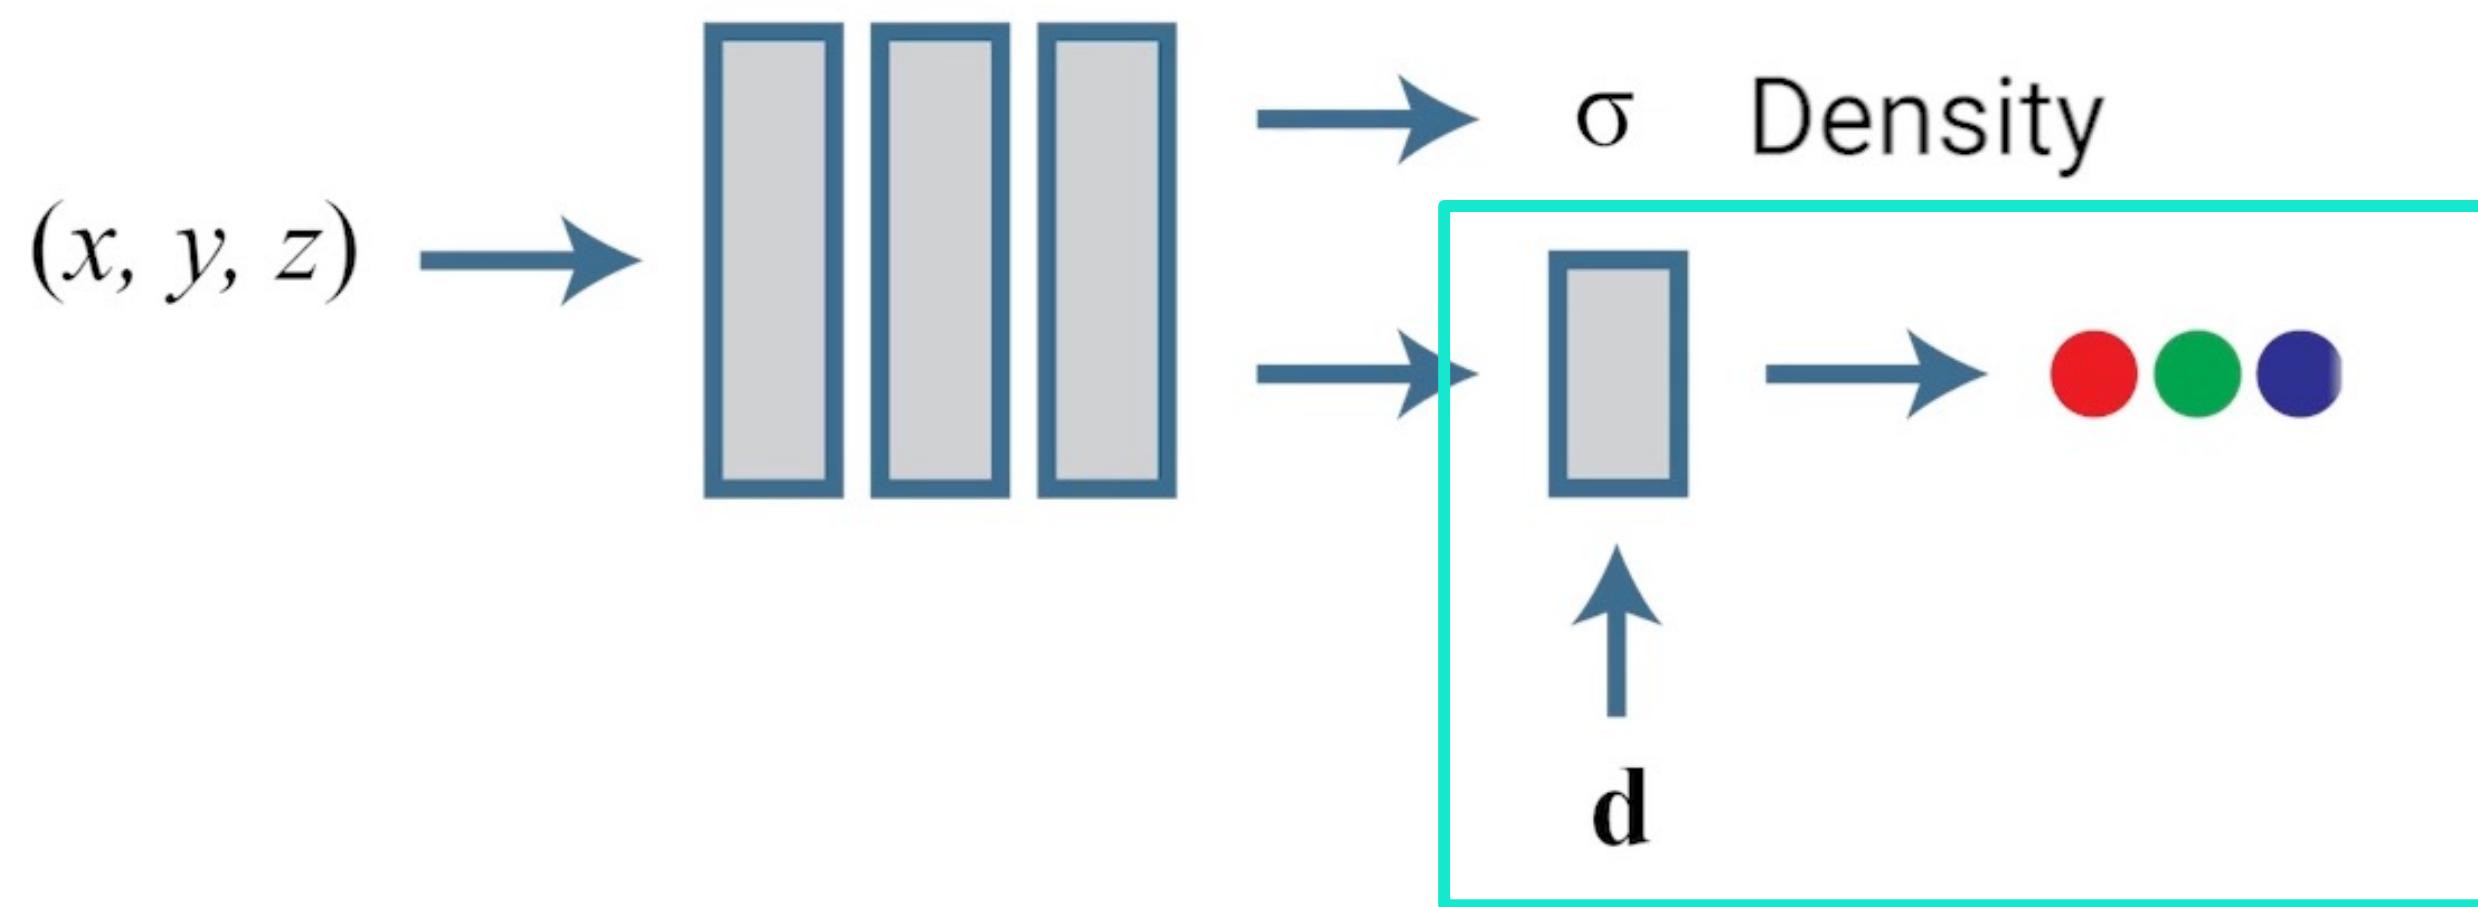
Spherical Harmonics  
fourier basis on the sphere

Approximate a function on the sphere with  
Spherical Harmonics coefficients



# NeRF with Spherical Functions

NeRF



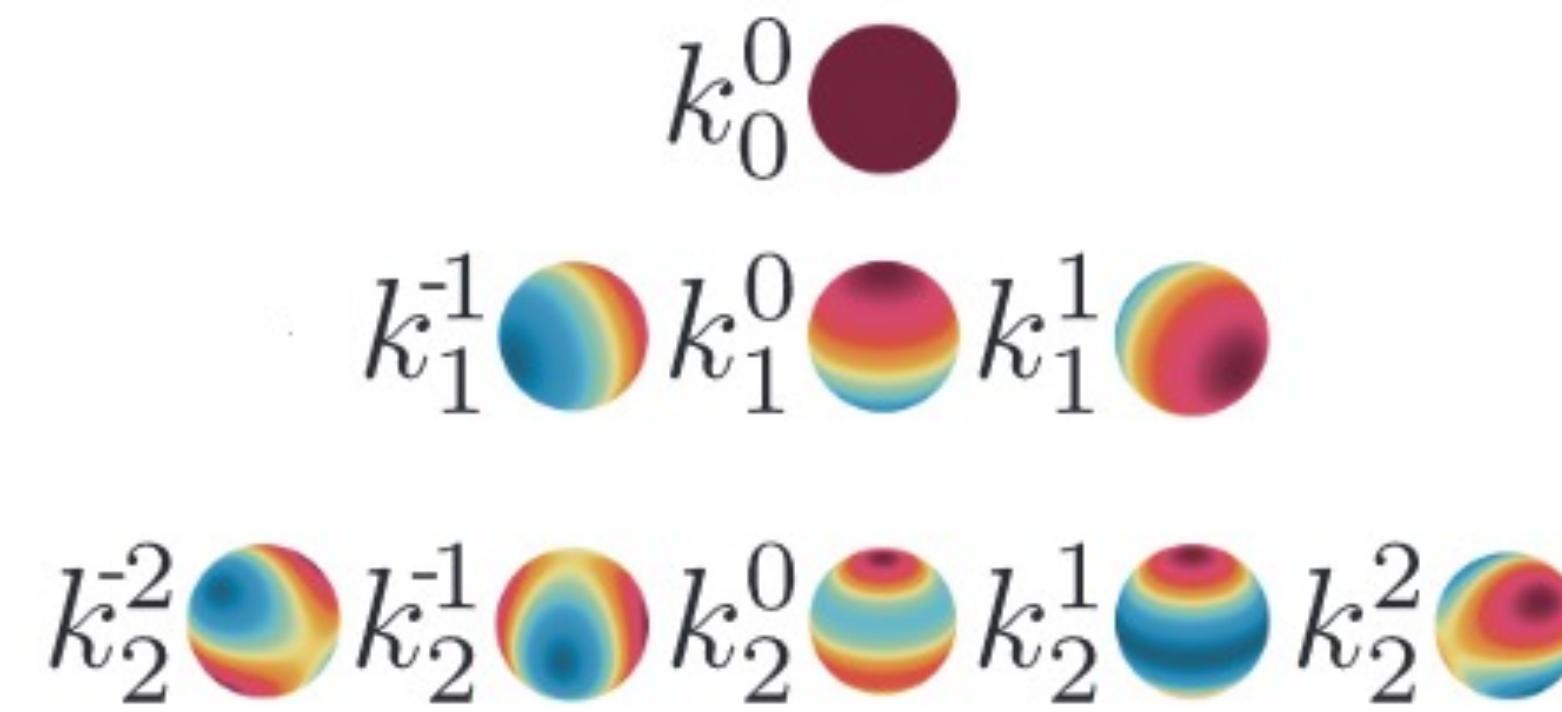
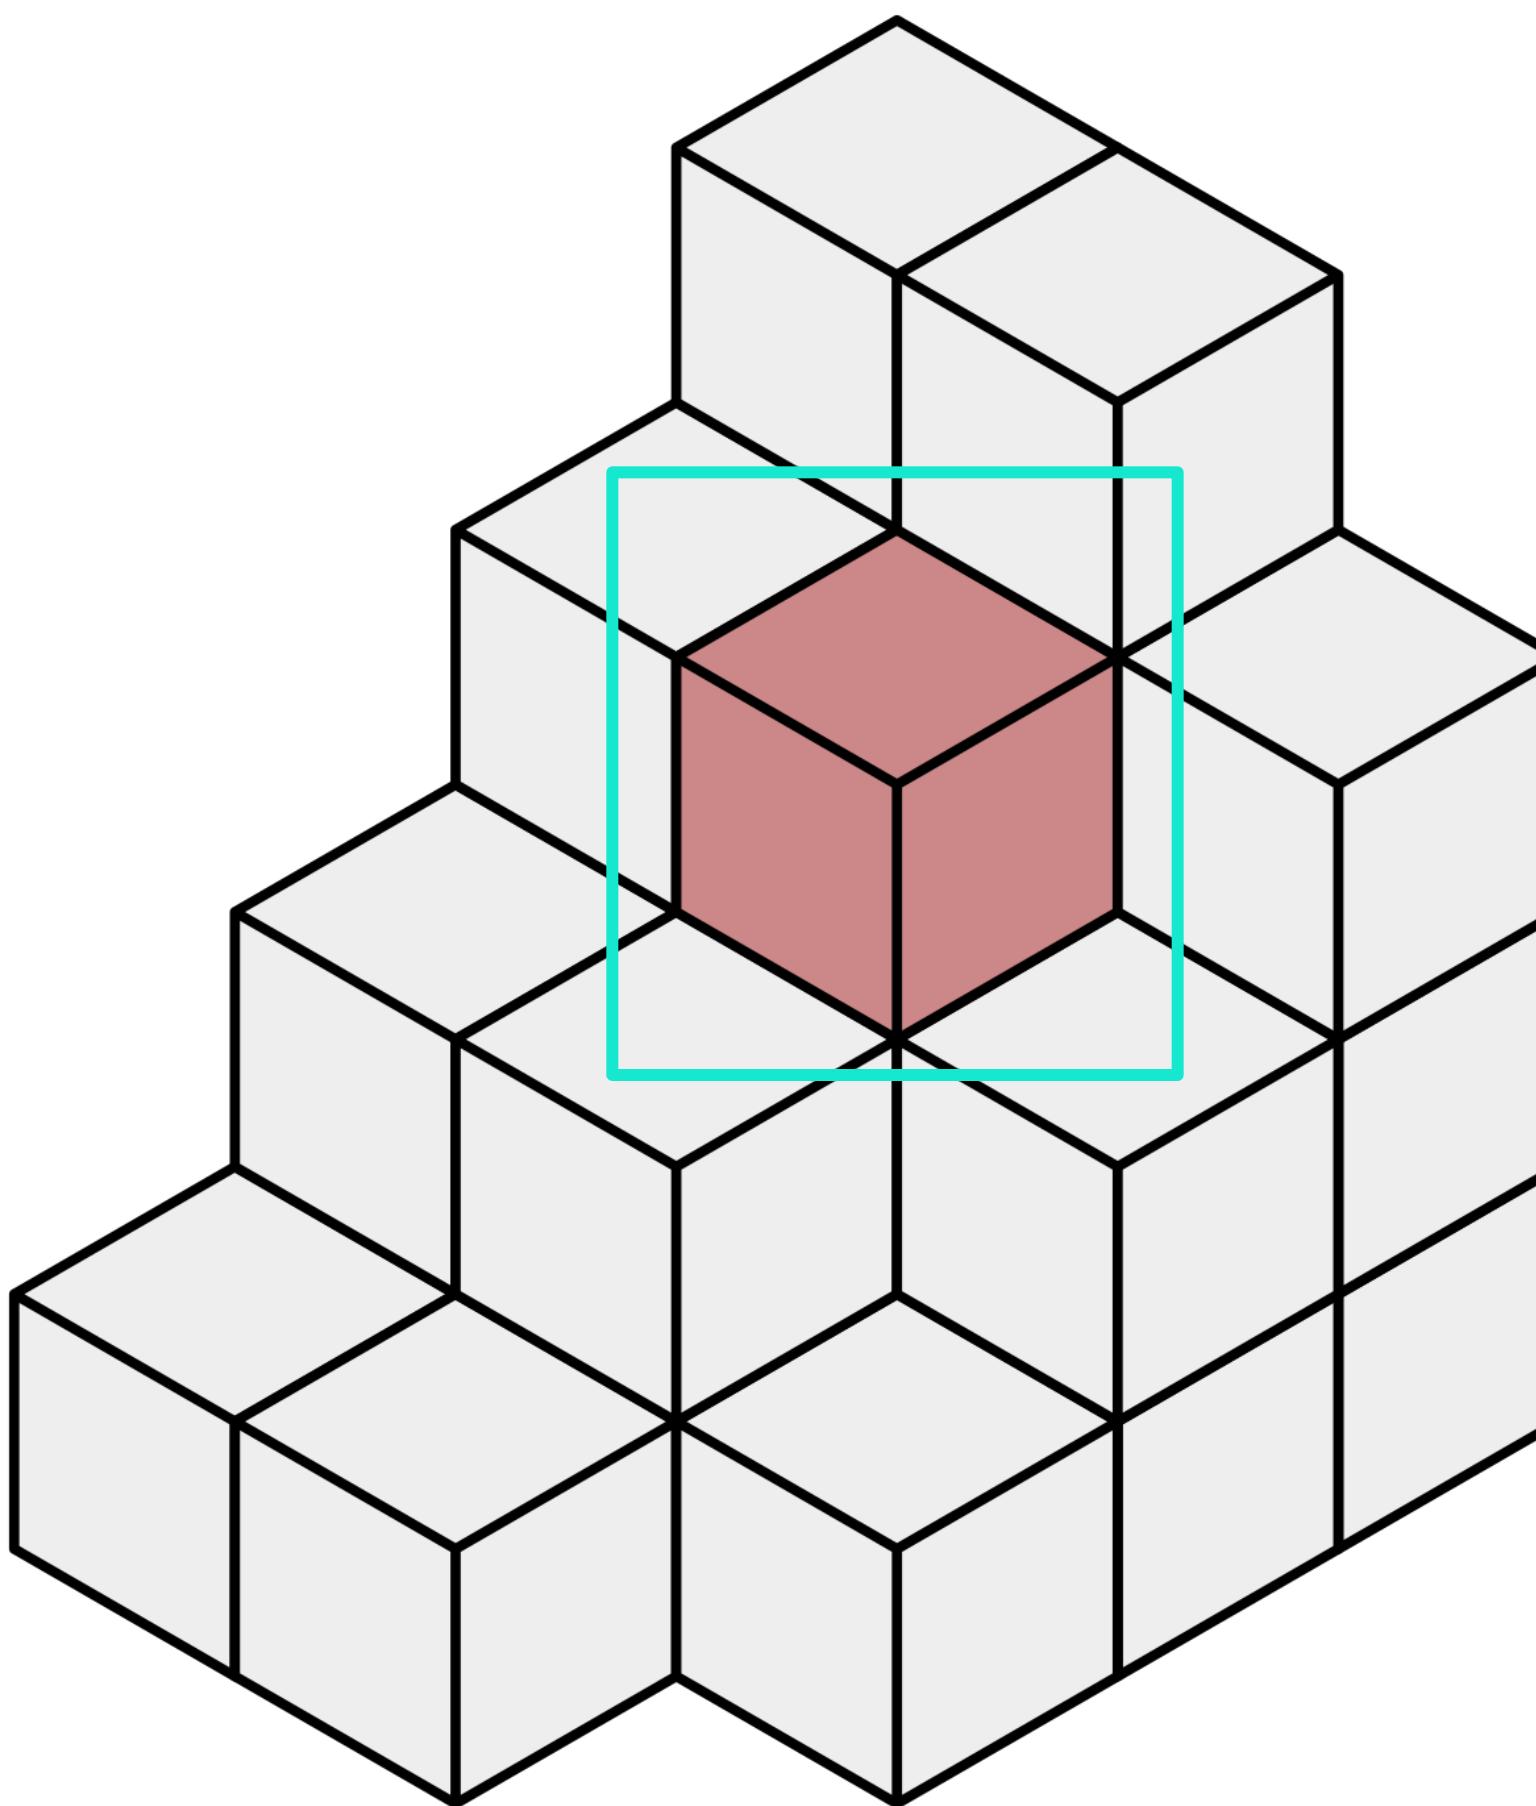
NeRF with Spherical Harmonics (NeRF-SH)



$$\begin{matrix} k_0^0 \\ k_1^1 & k_1^0 & k_1^1 \\ k_2^1 & k_2^0 & k_2^1 & k_2^2 \\ k_3^2 & k_3^1 & k_3^0 & k_3^1 & k_3^2 & k_3^3 \end{matrix}$$

A large bold black letter **K** is positioned between the  $k_1$  and  $k_2$  rows of coefficients.

PlenOctree =  
Sparse Voxels with density + SH coefficients



Skips empty regions  
→ much faster rendering time!

**PlenOctree**  
54.00 FPS



**NeRF**  
0.013 FPS

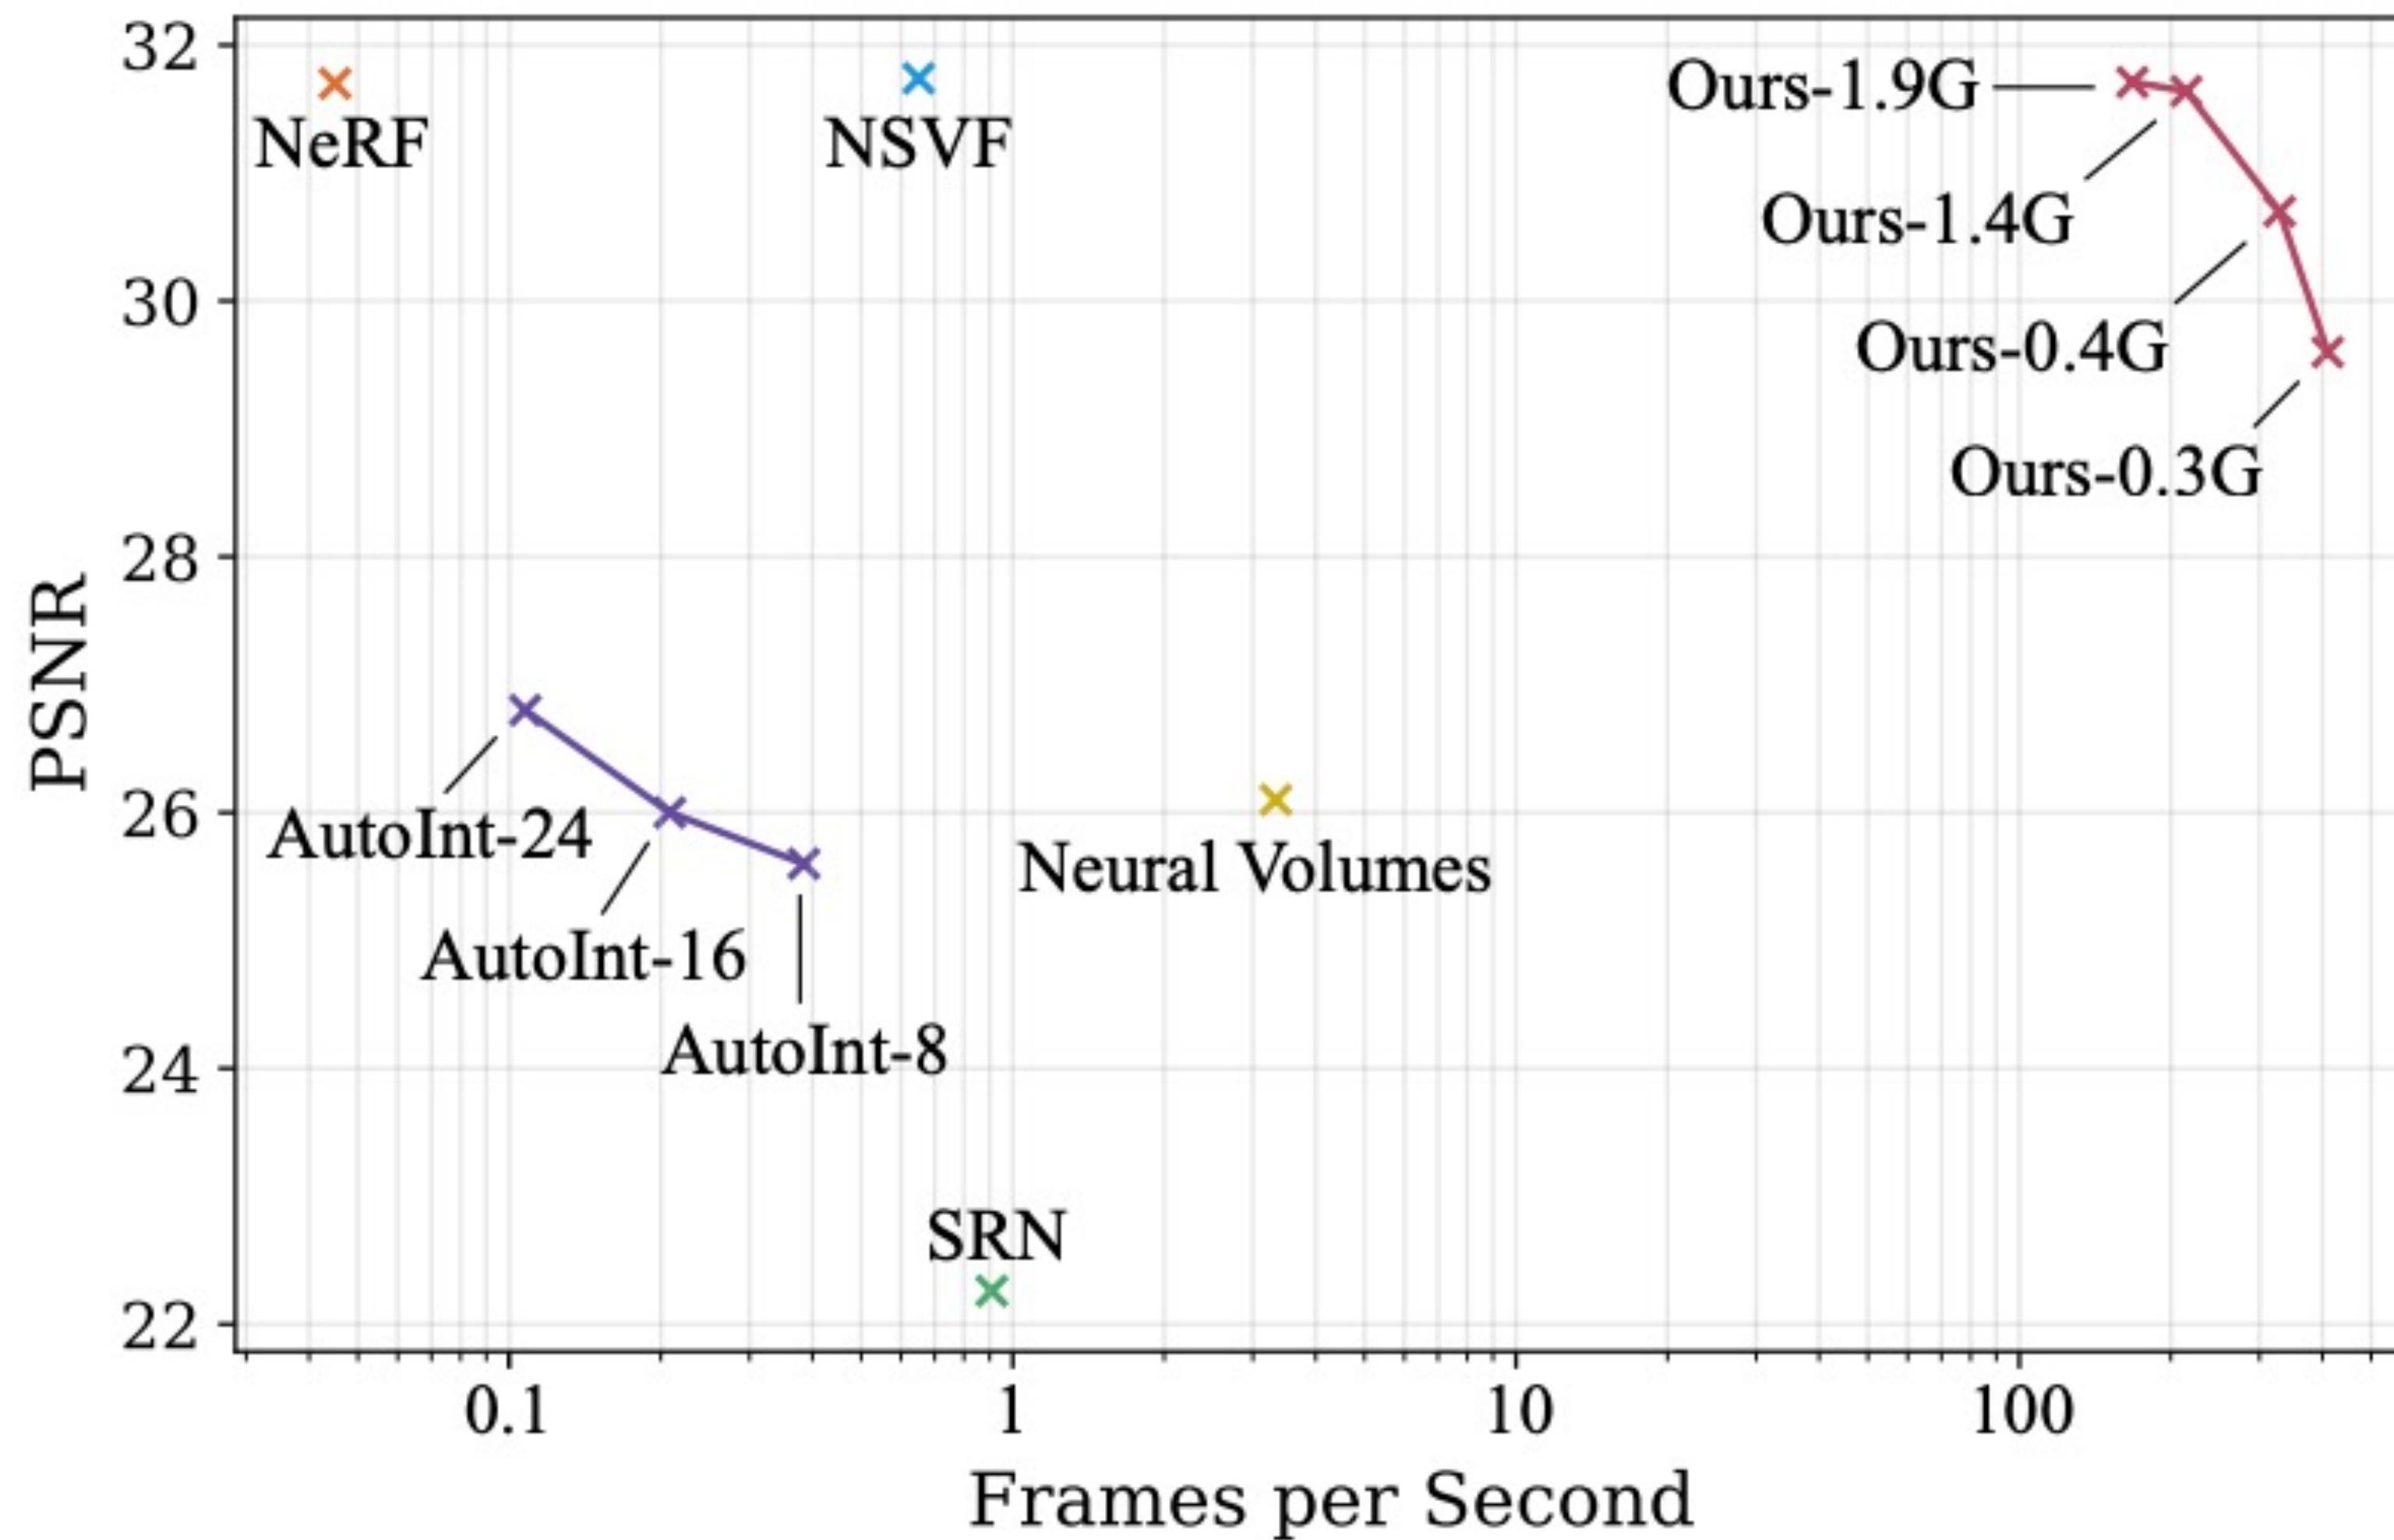


More Results

# Trade-off



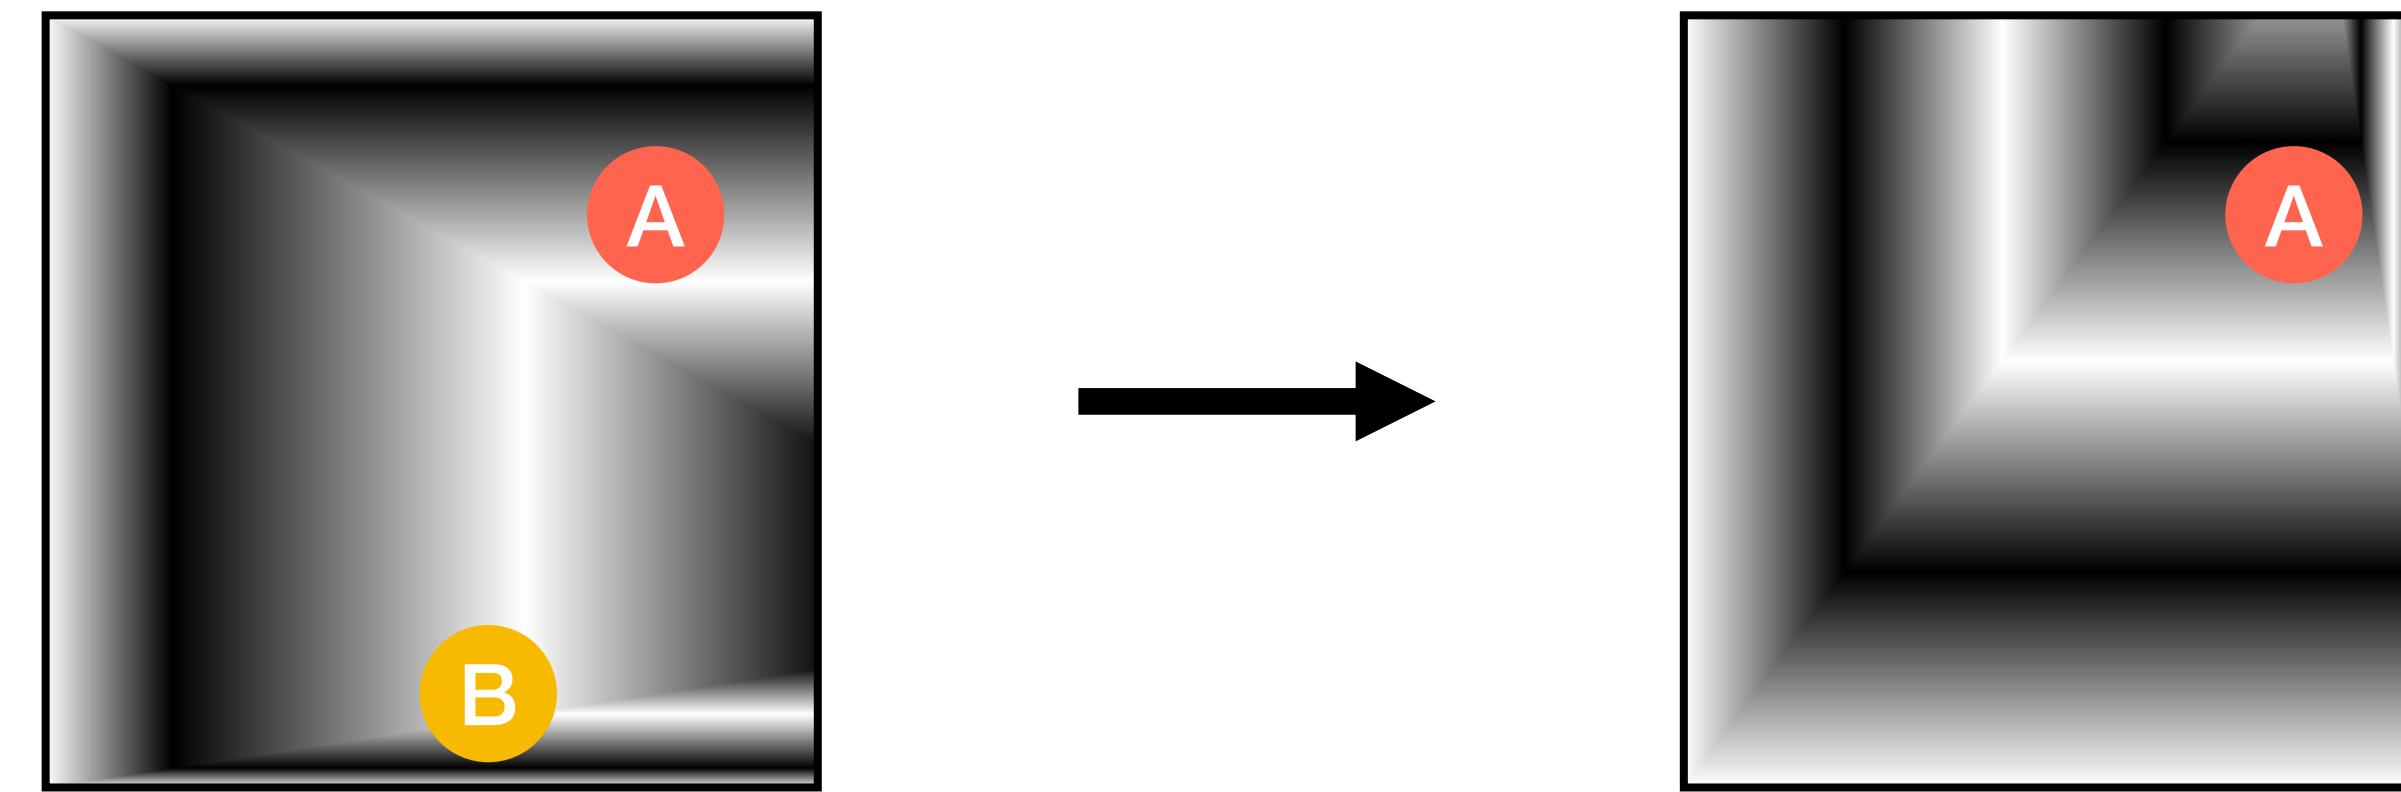
3000x faster than OG NeRF!



**Can we learn the encoding?**

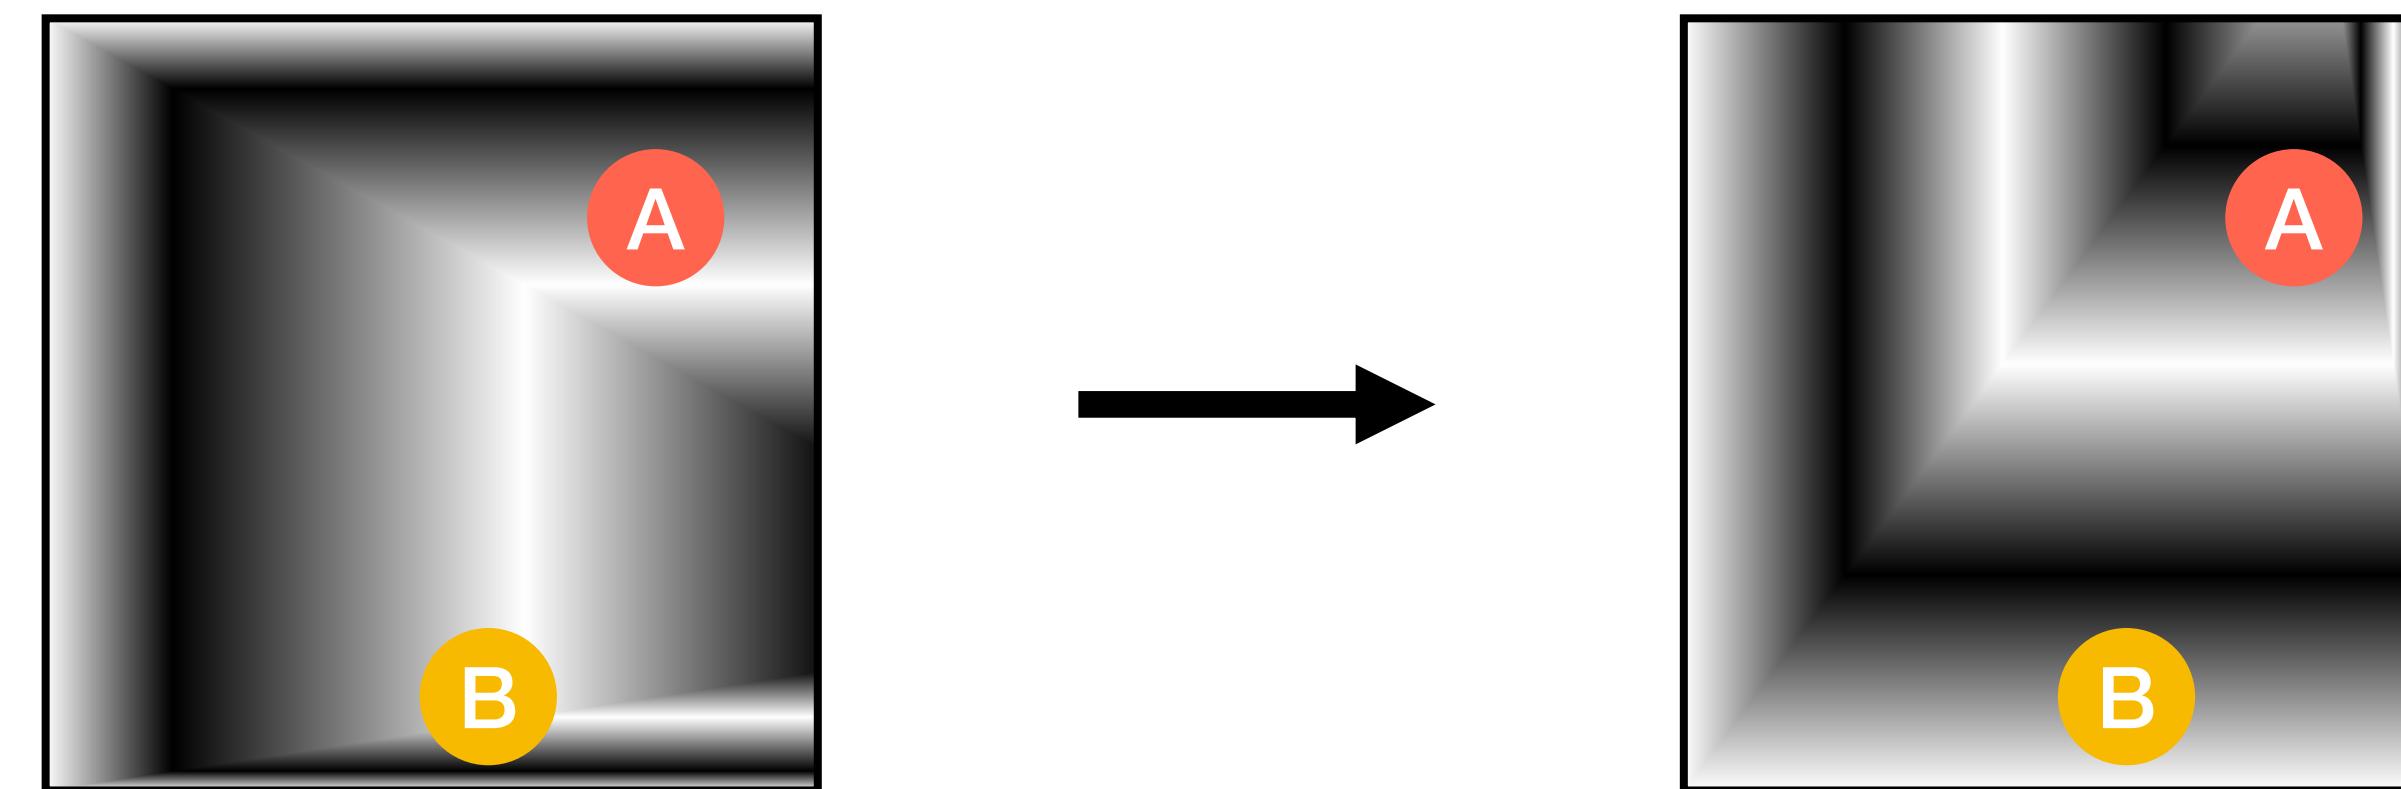
# Learnable Encodings

Lets try optimizing phase and frequency s.t. A goes to ■



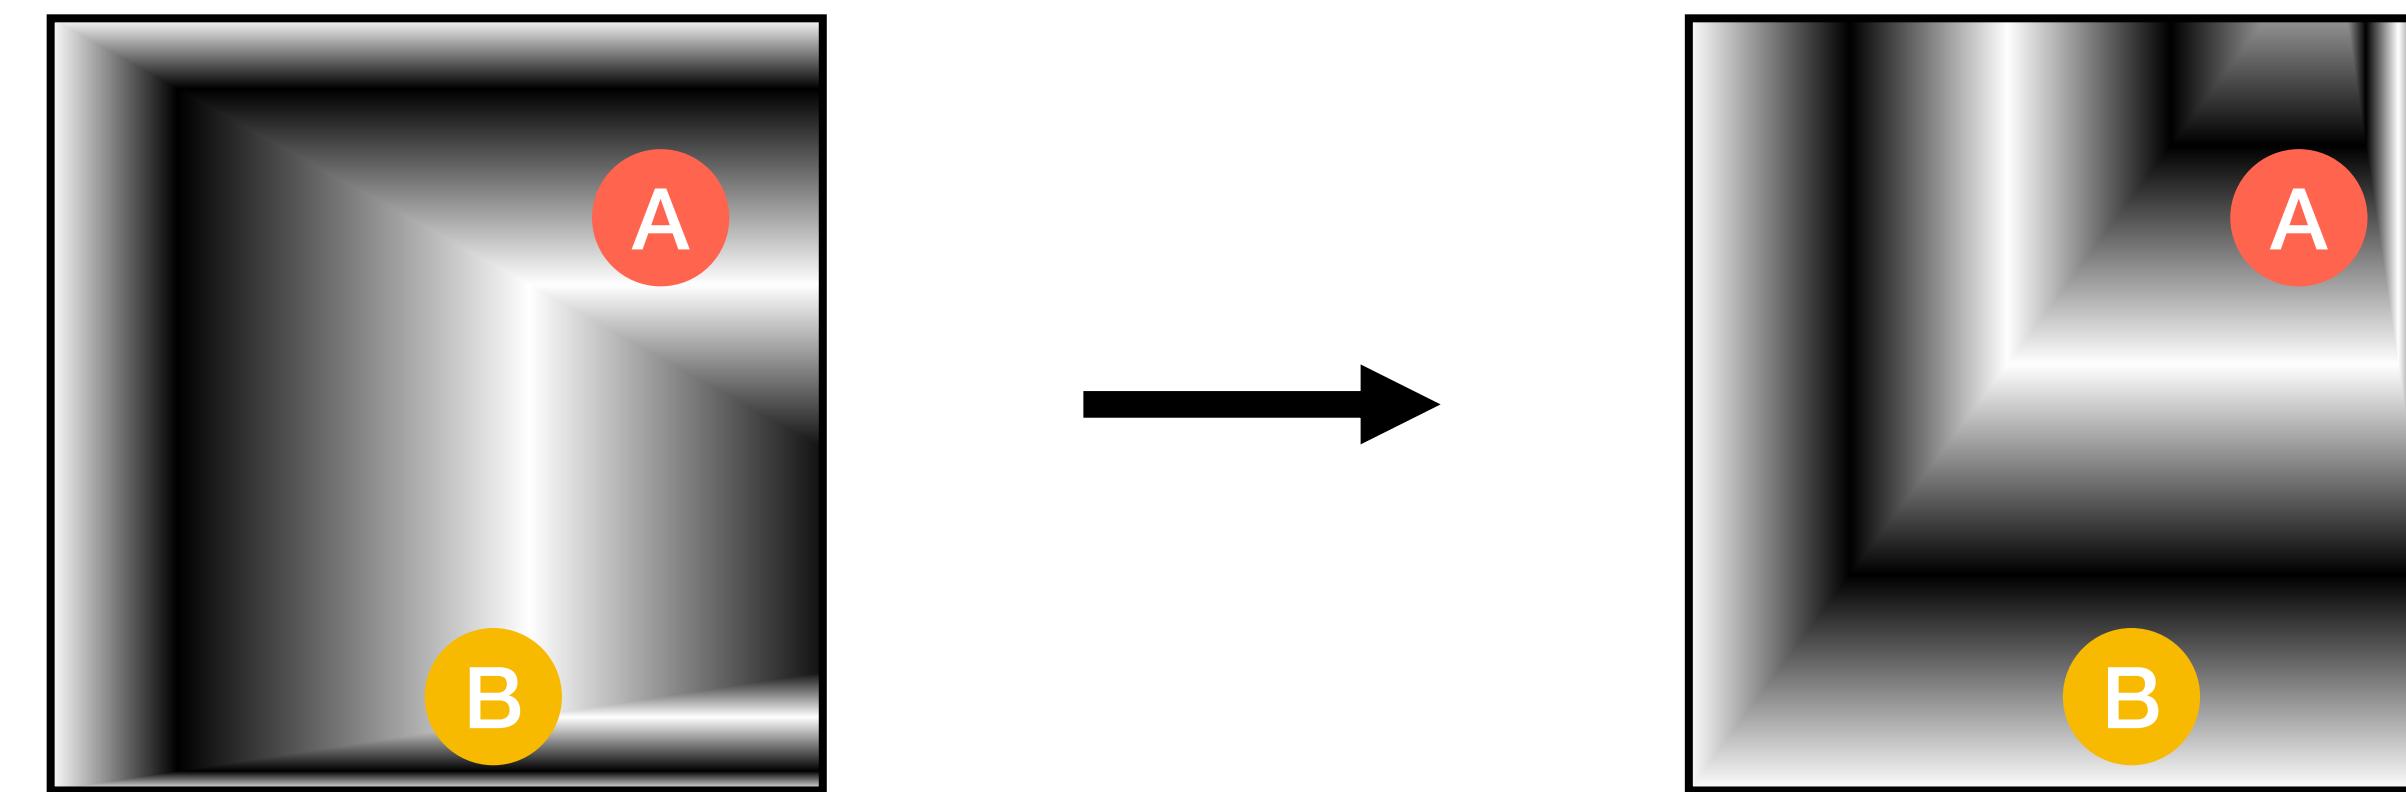
# Learnable Encodings

Lets try optimizing phase and frequency s.t. A goes to ■



# Learnable Encodings

Lets try optimizing phase and frequency s.t. A goes to ■



B also changes which makes optimization difficult

**Desired: Learnable encoding  
with local extent**

**Desired: Learnable encoding  
with local extent**

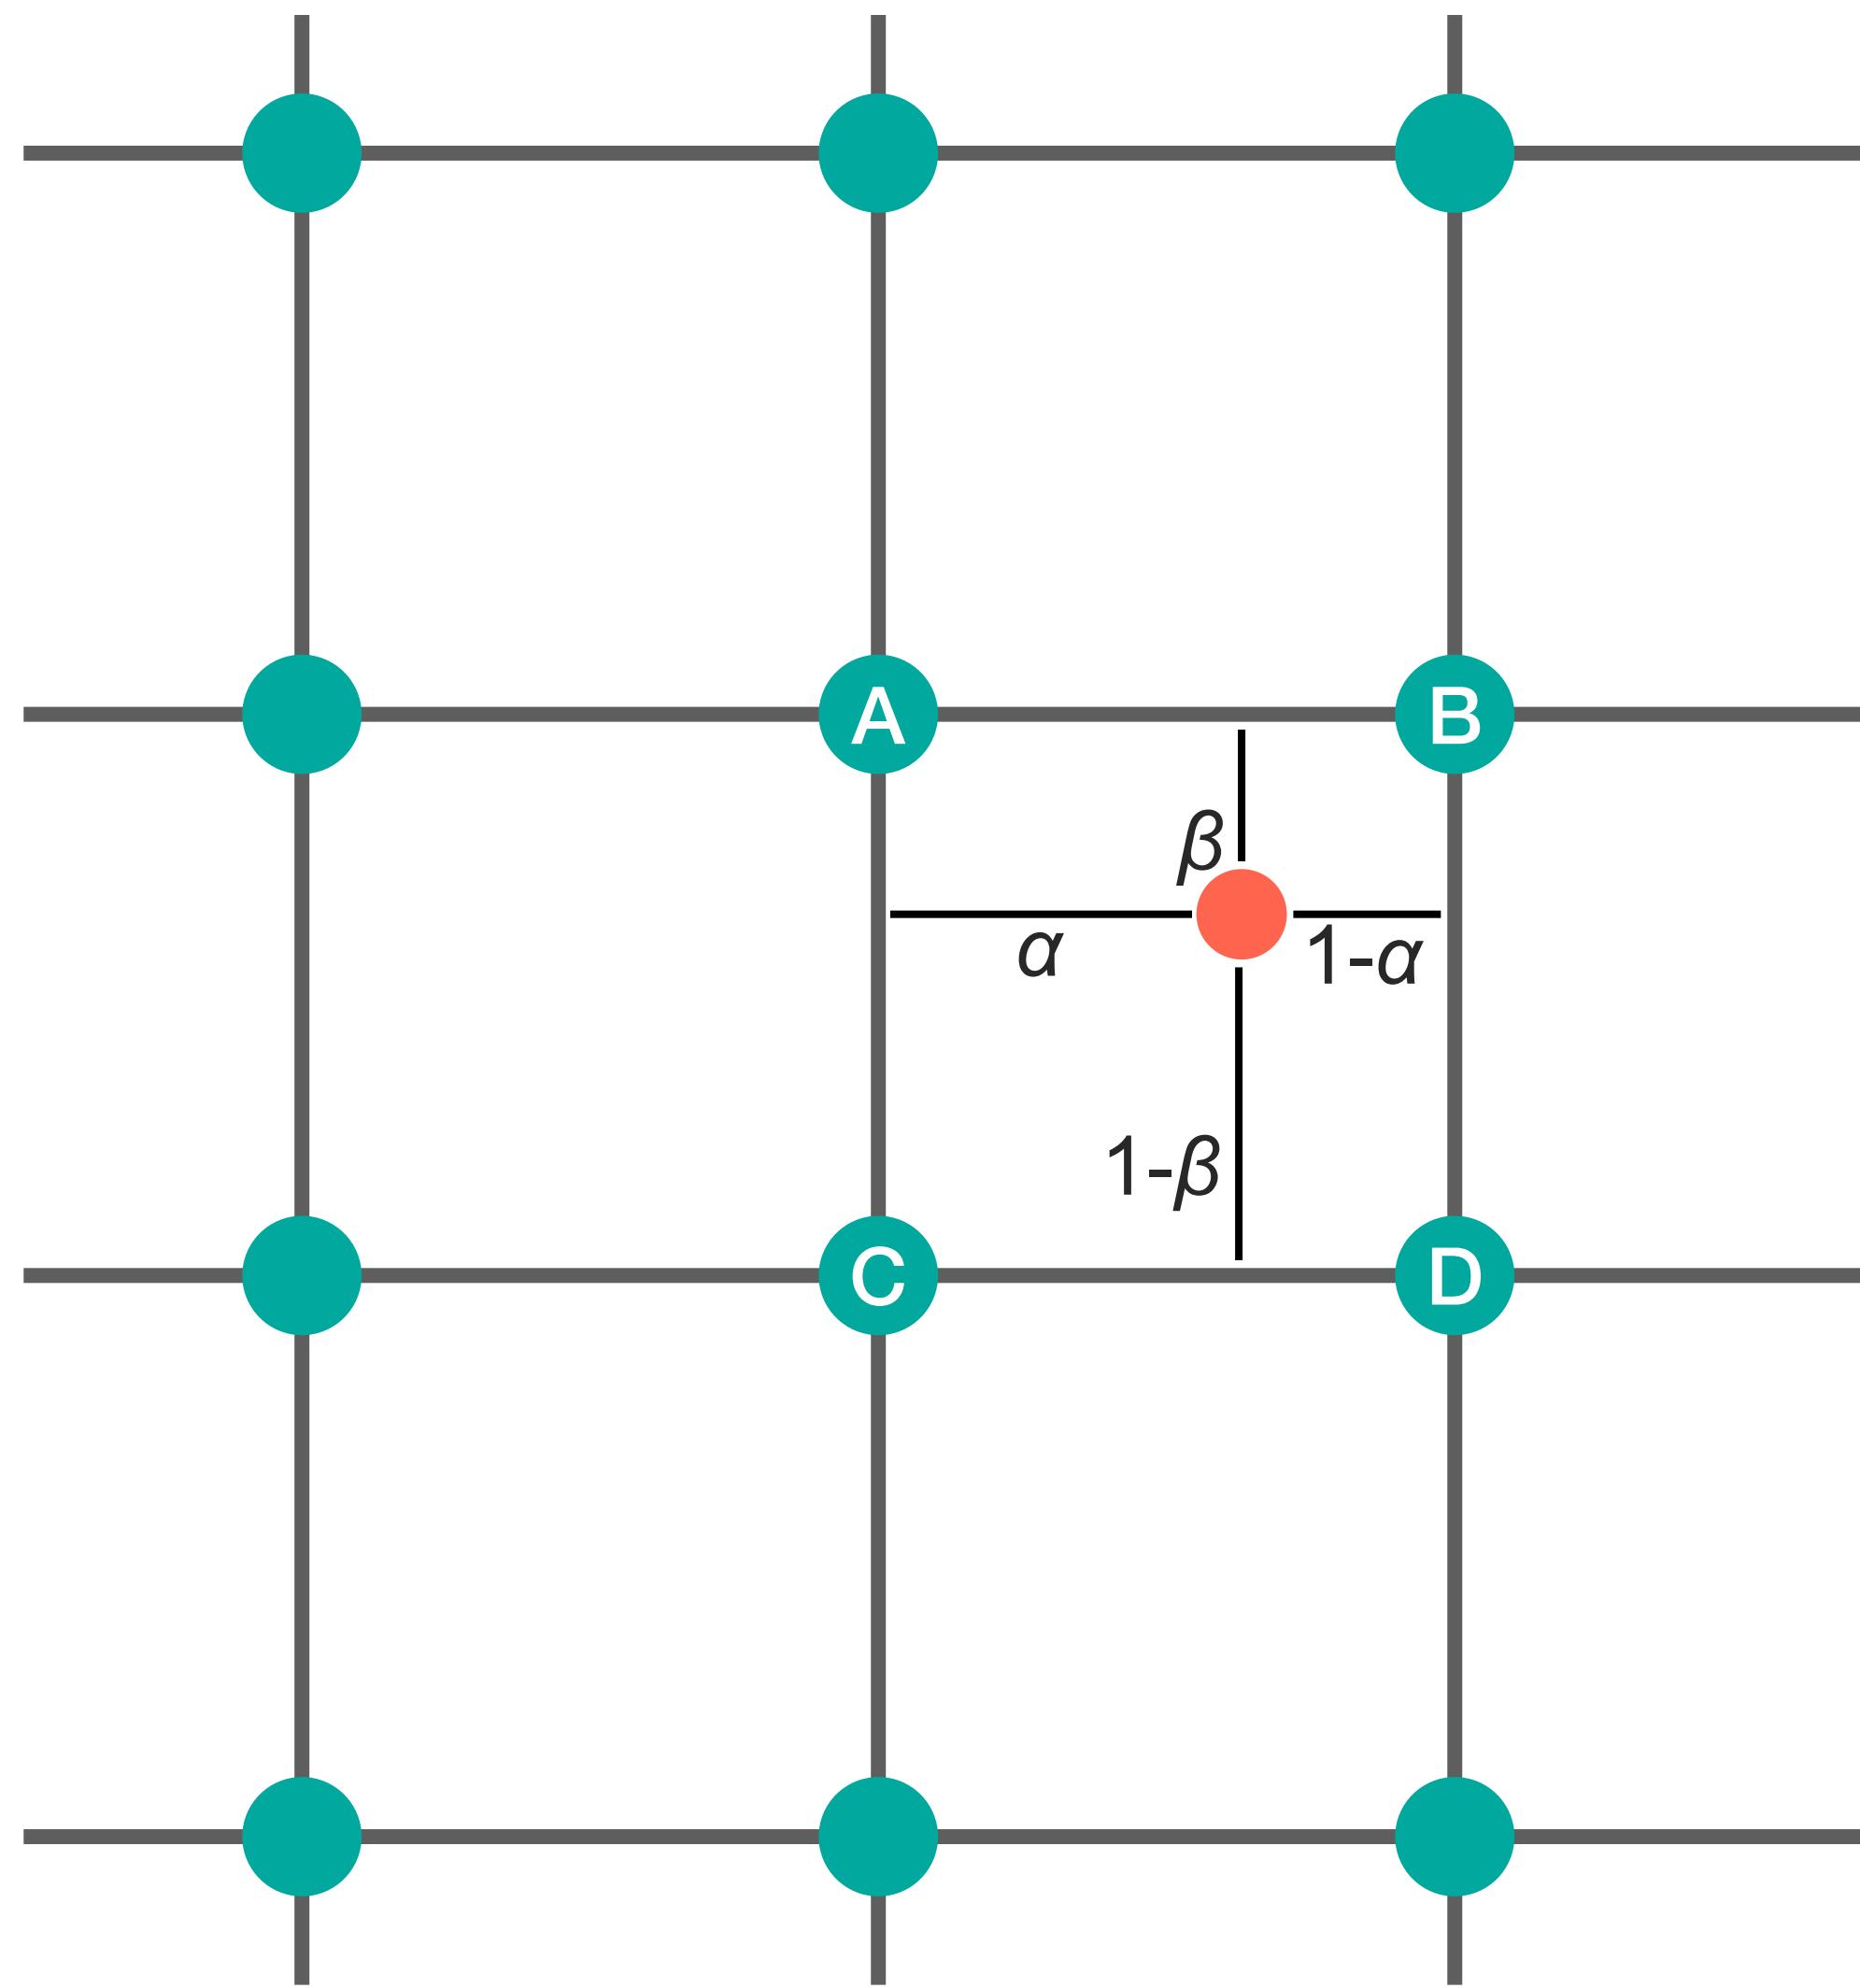
**Solution: Feature Grids (i.e. voxel)**

**Desired: Learnable encoding  
with local extent**

**Solution: Feature Grids (i.e. voxel)**

**Challenge: Resolution**

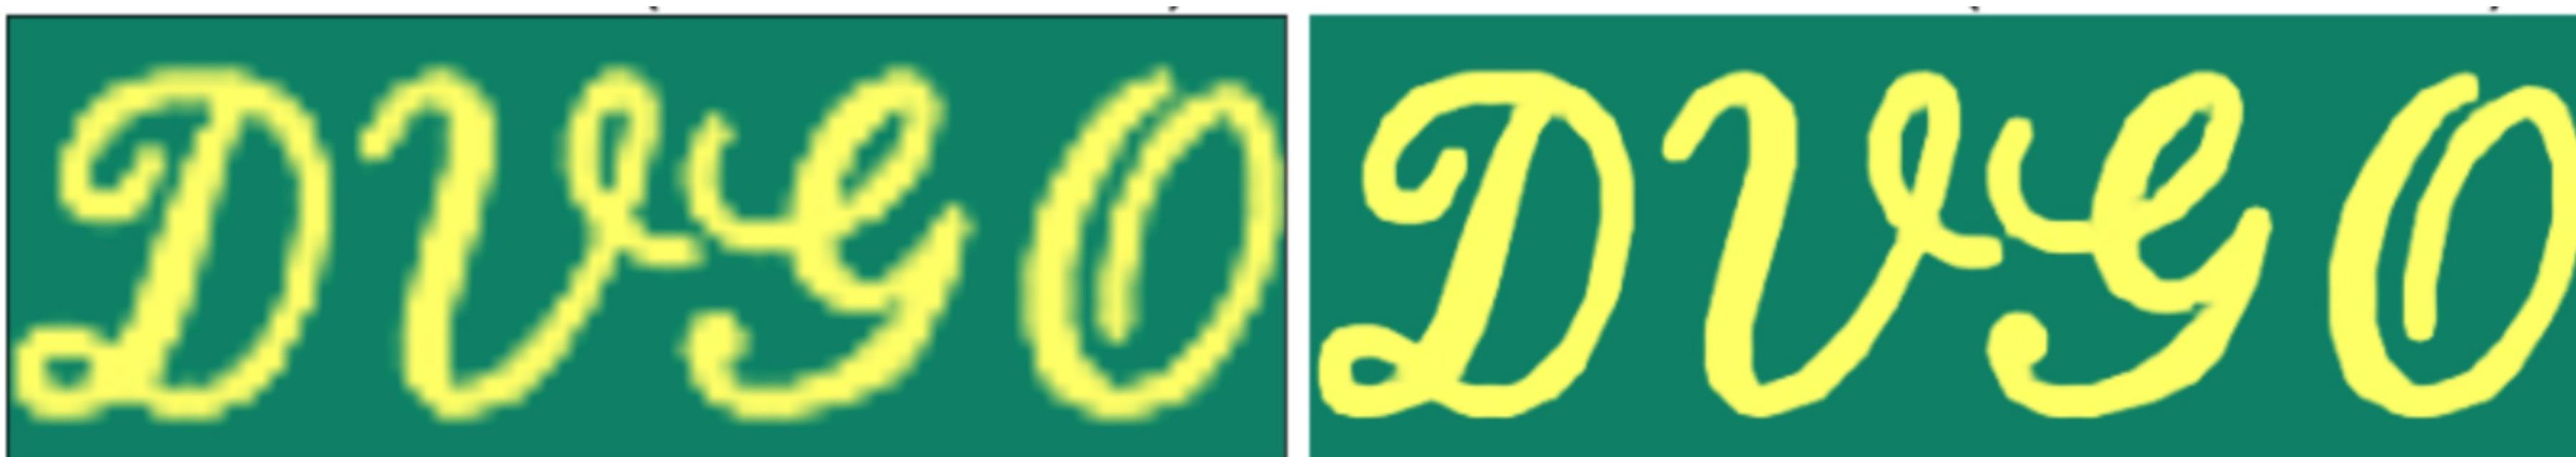
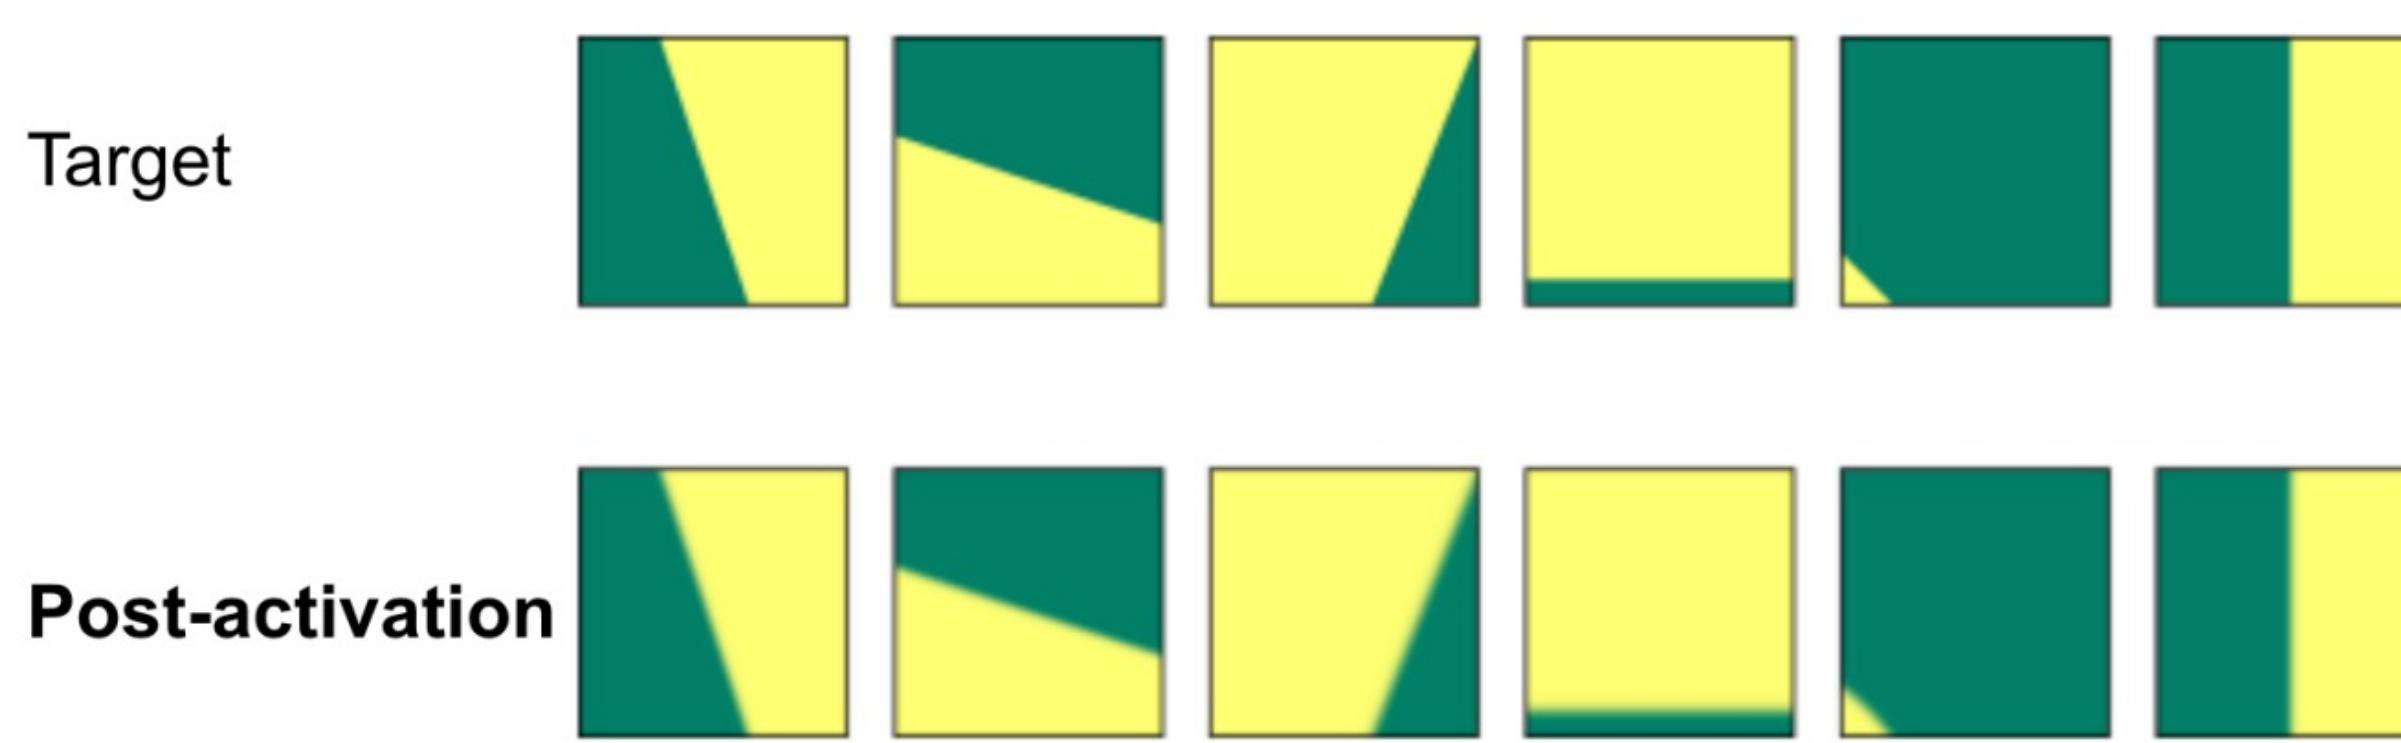
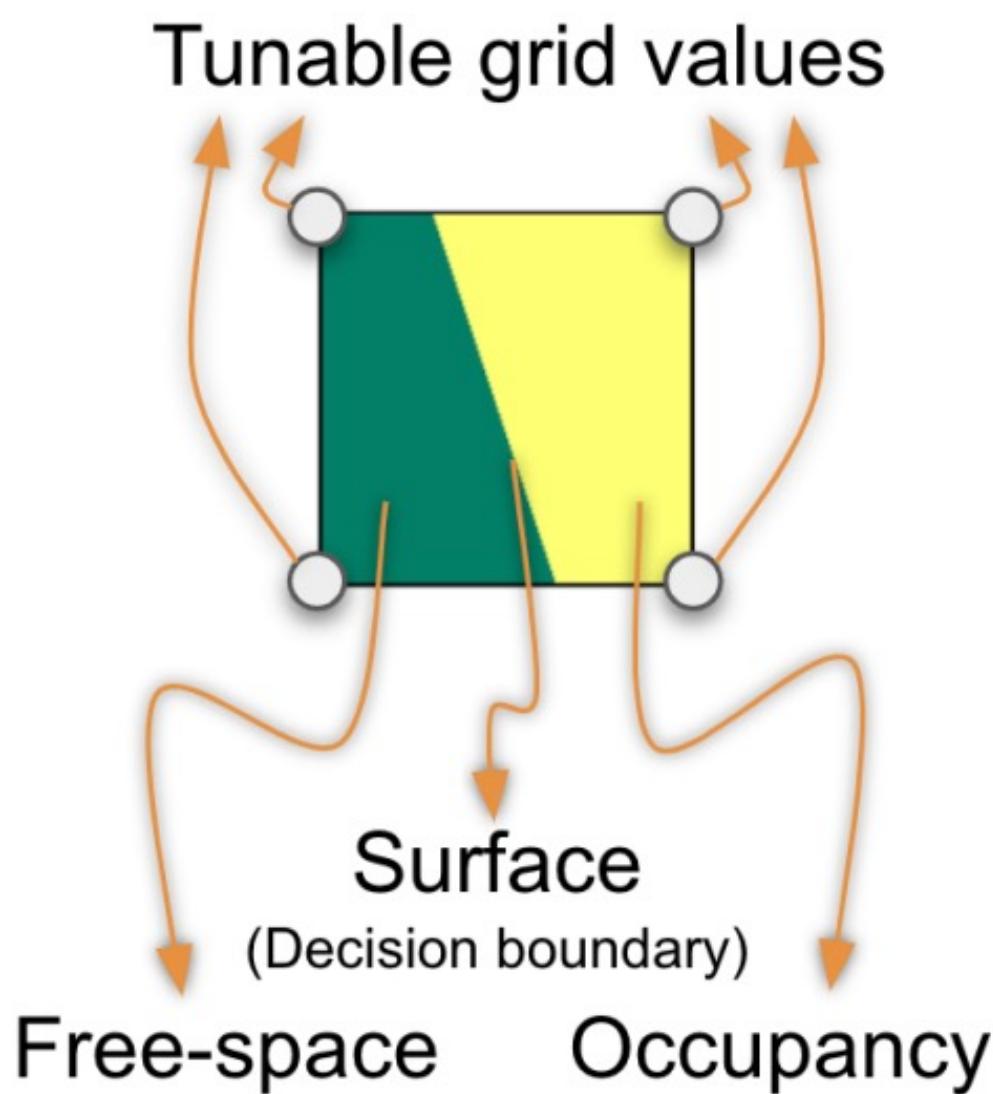
# Make continuous via interpolation



$$\bullet = \beta(\alpha A + (1 - \alpha)B) + (1 - \beta)(\alpha C + (1 - \alpha)D)$$

# Feature grids can be effective

Toy task for a 2D grid cell



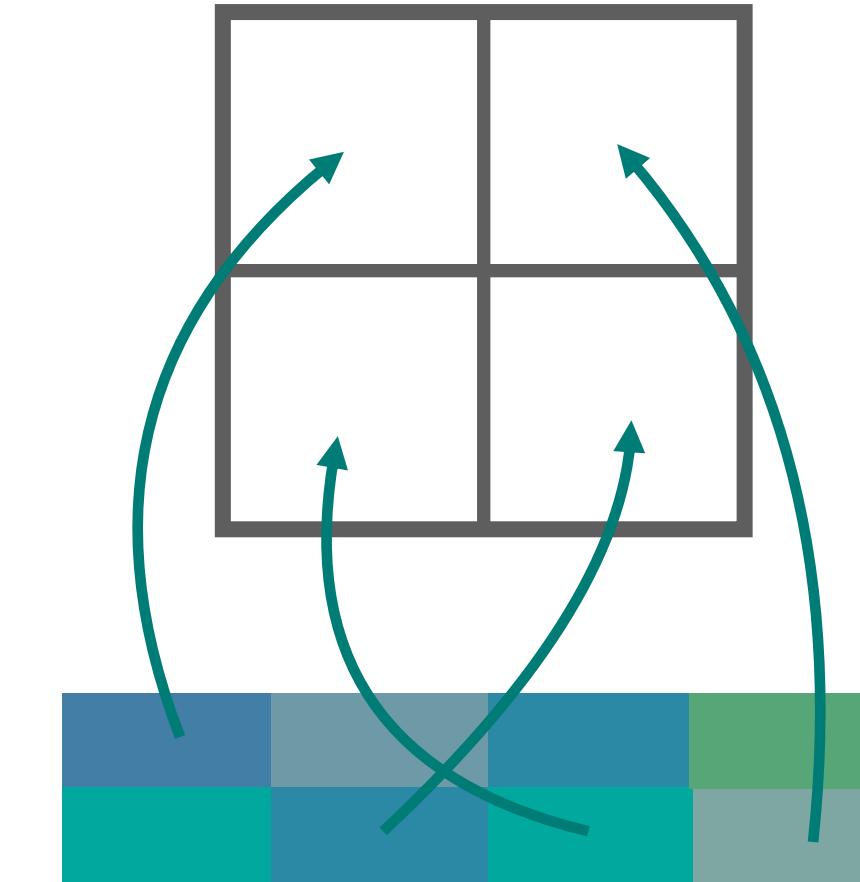
# Compression Techniques



Sparsity

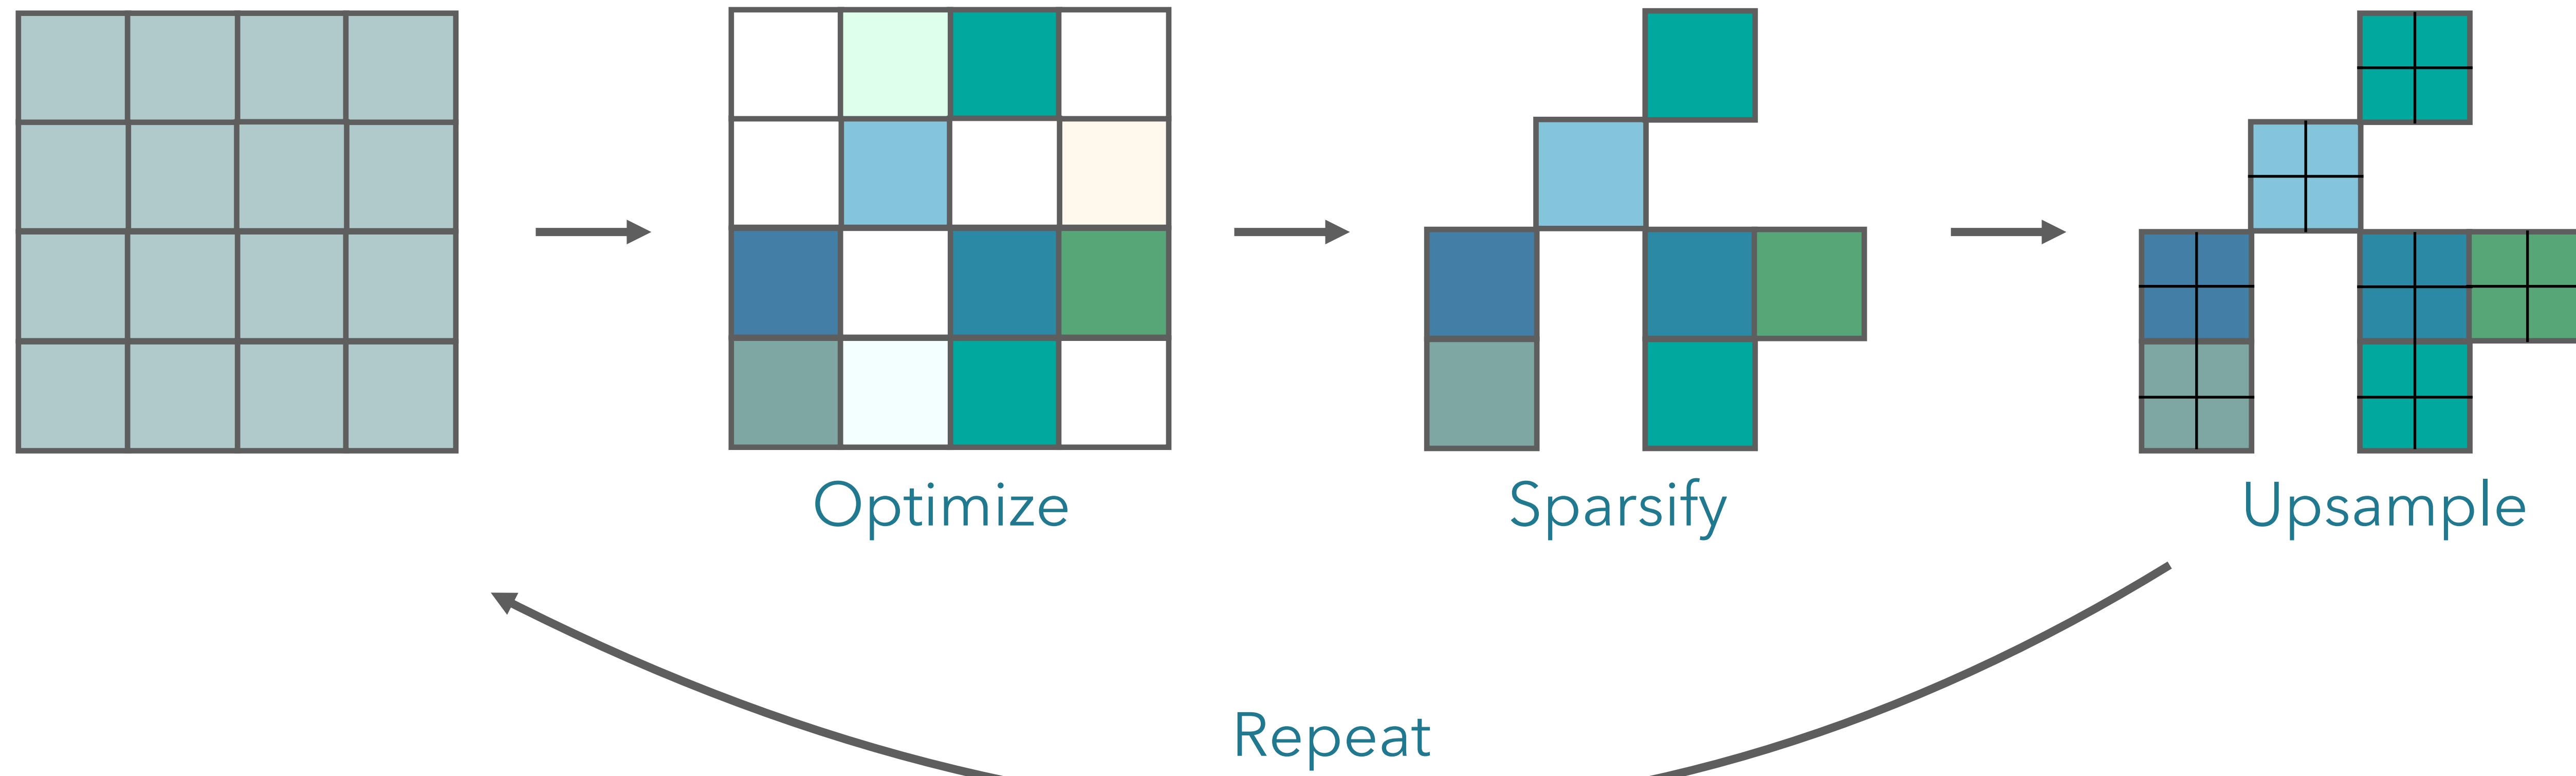


Low Rank



Dictionary

# Sparsity

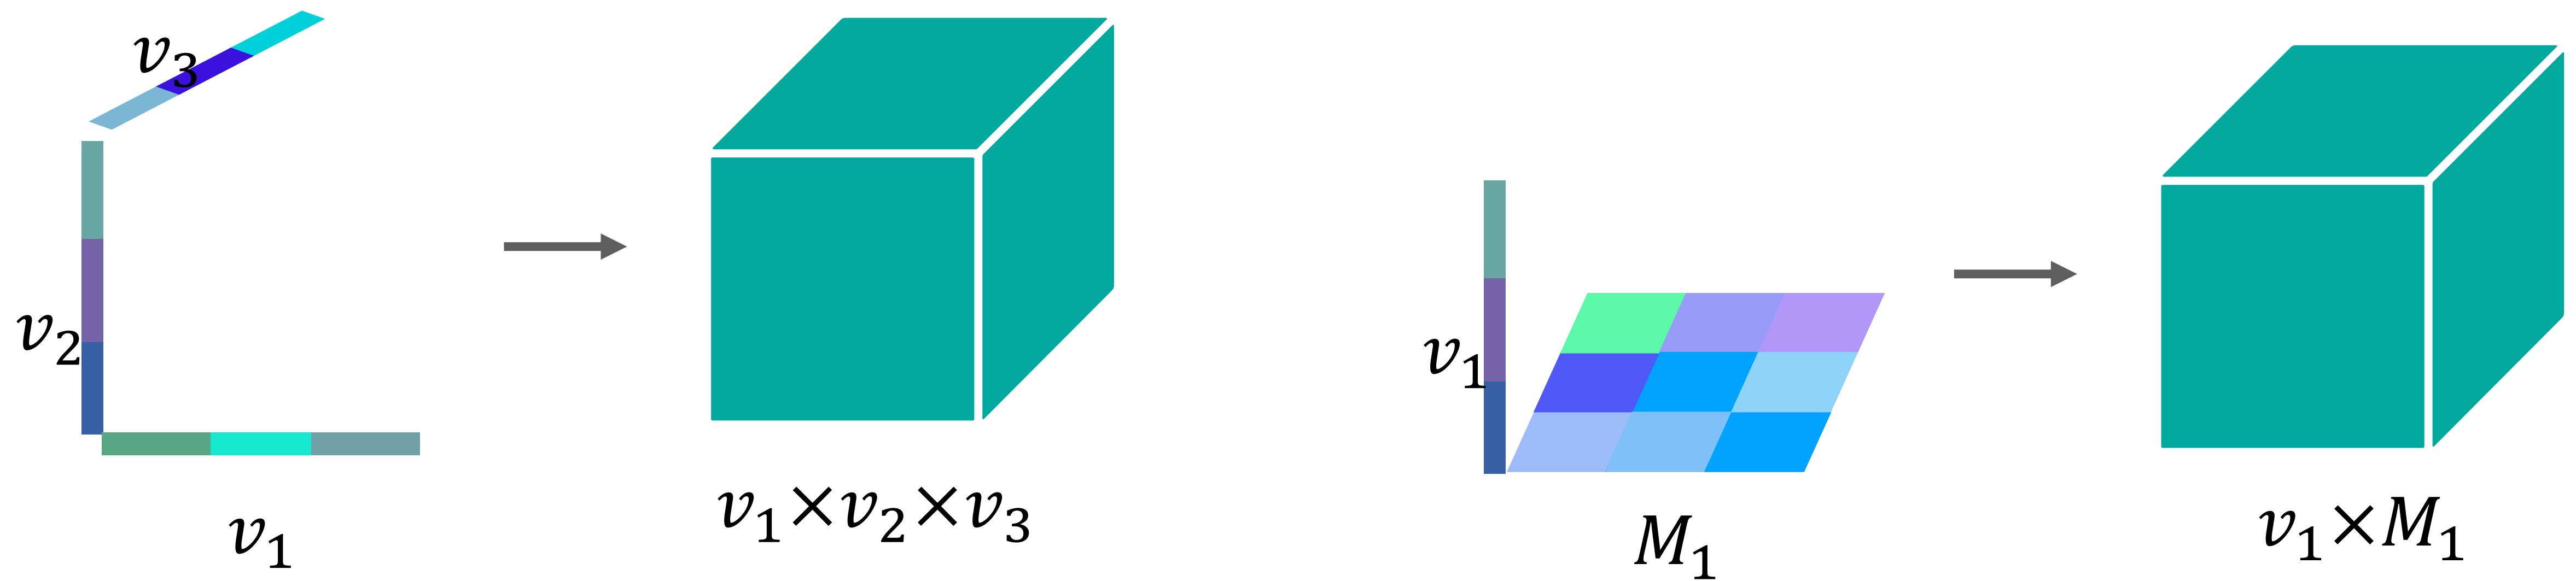


Lui et al. 2020, Neural Sparse Voxel Fields

Yu\*, Friedovich-Keil\* et al. 2021, Plenoxels: Radiance Fields without Neural Networks

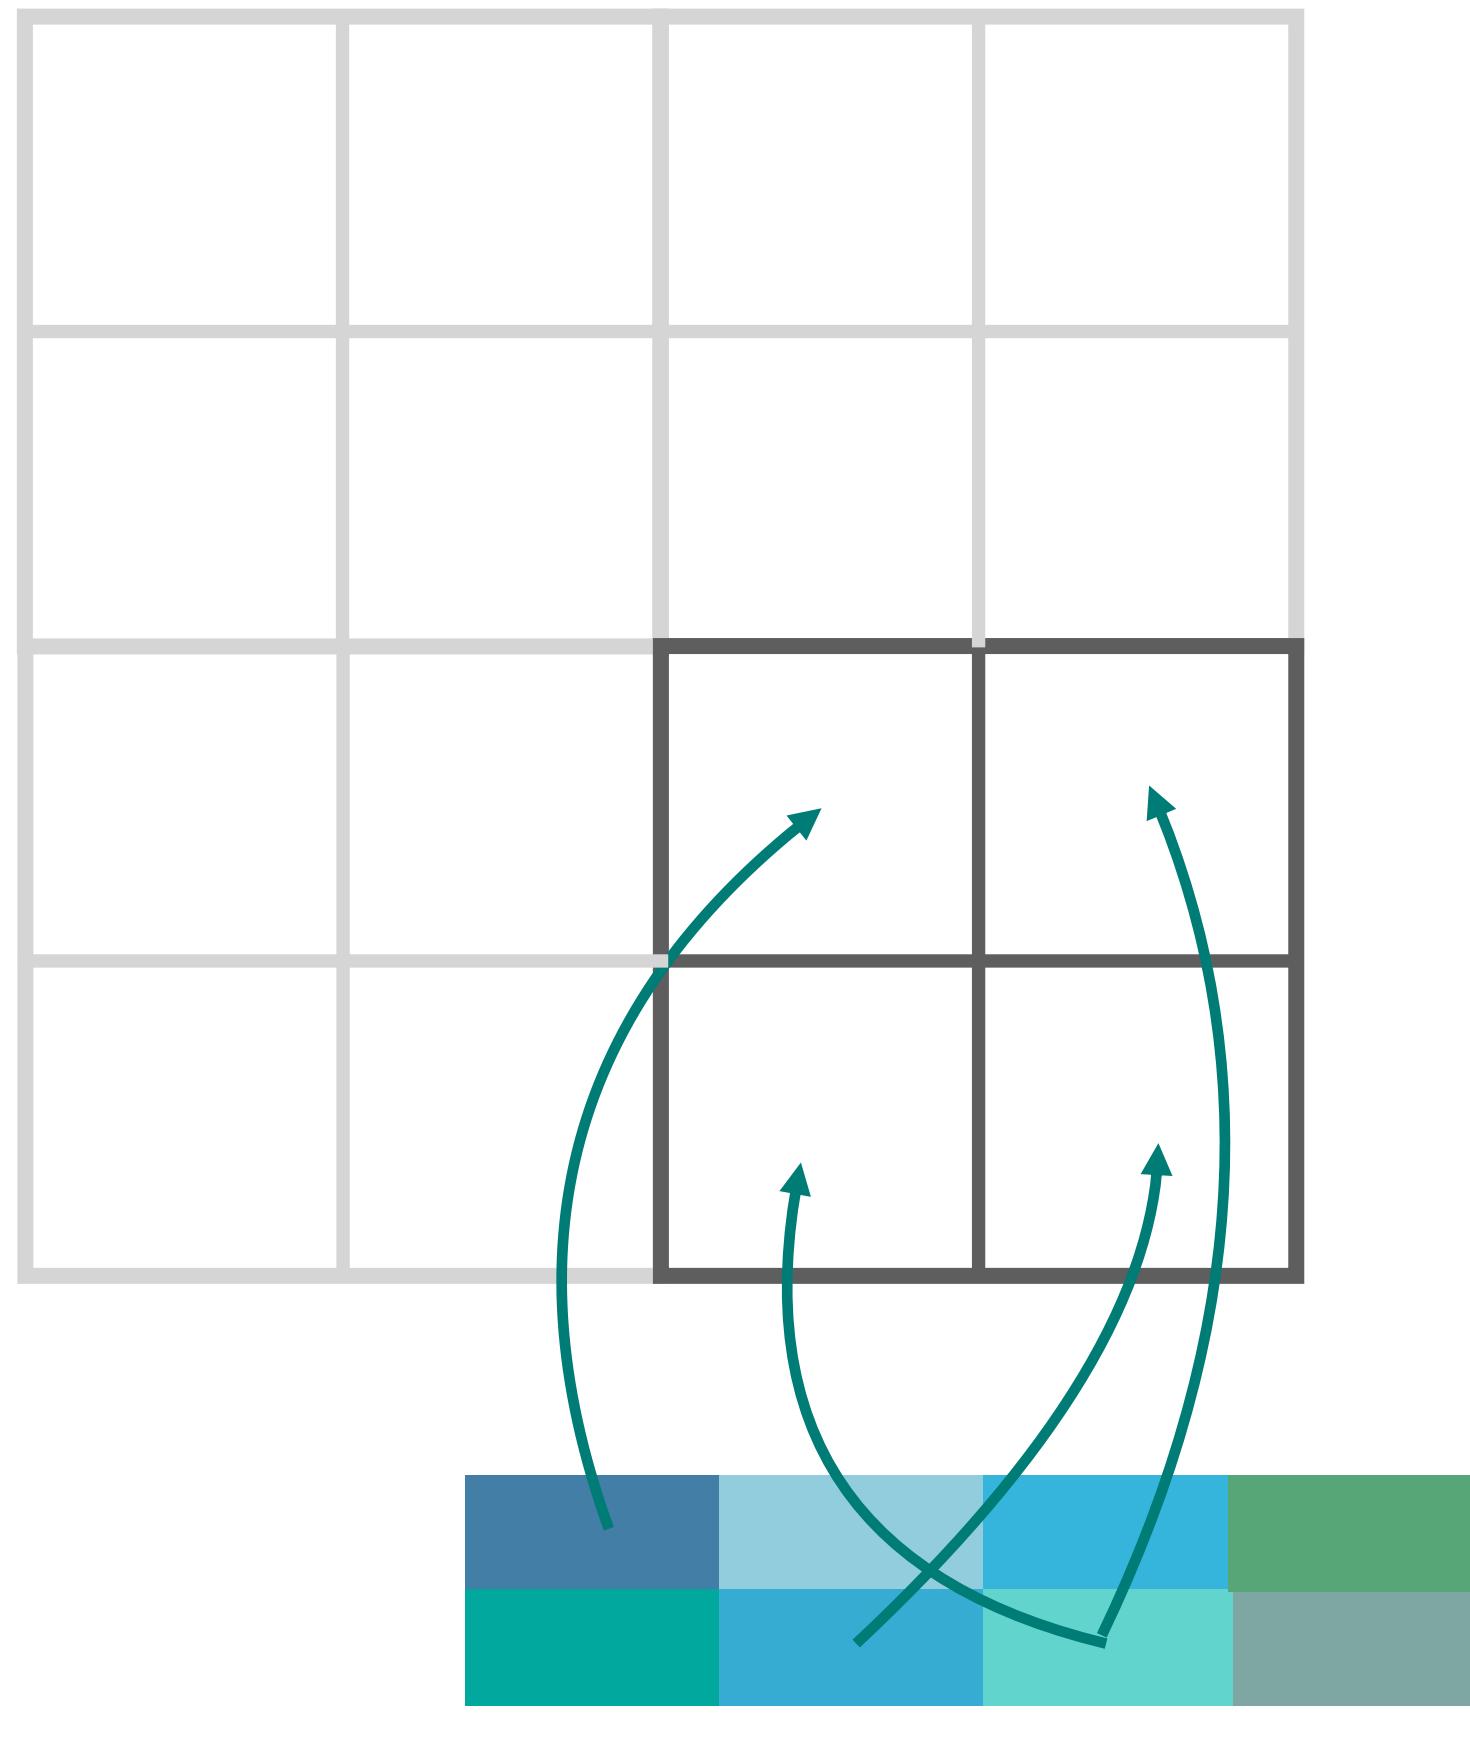
Sun et al. 2021, Direct Voxel Grid Optimization: Super-fast Convergence for Radiance Fields Reconstruction

# Low Rank Approximations



# Dictionary Methods

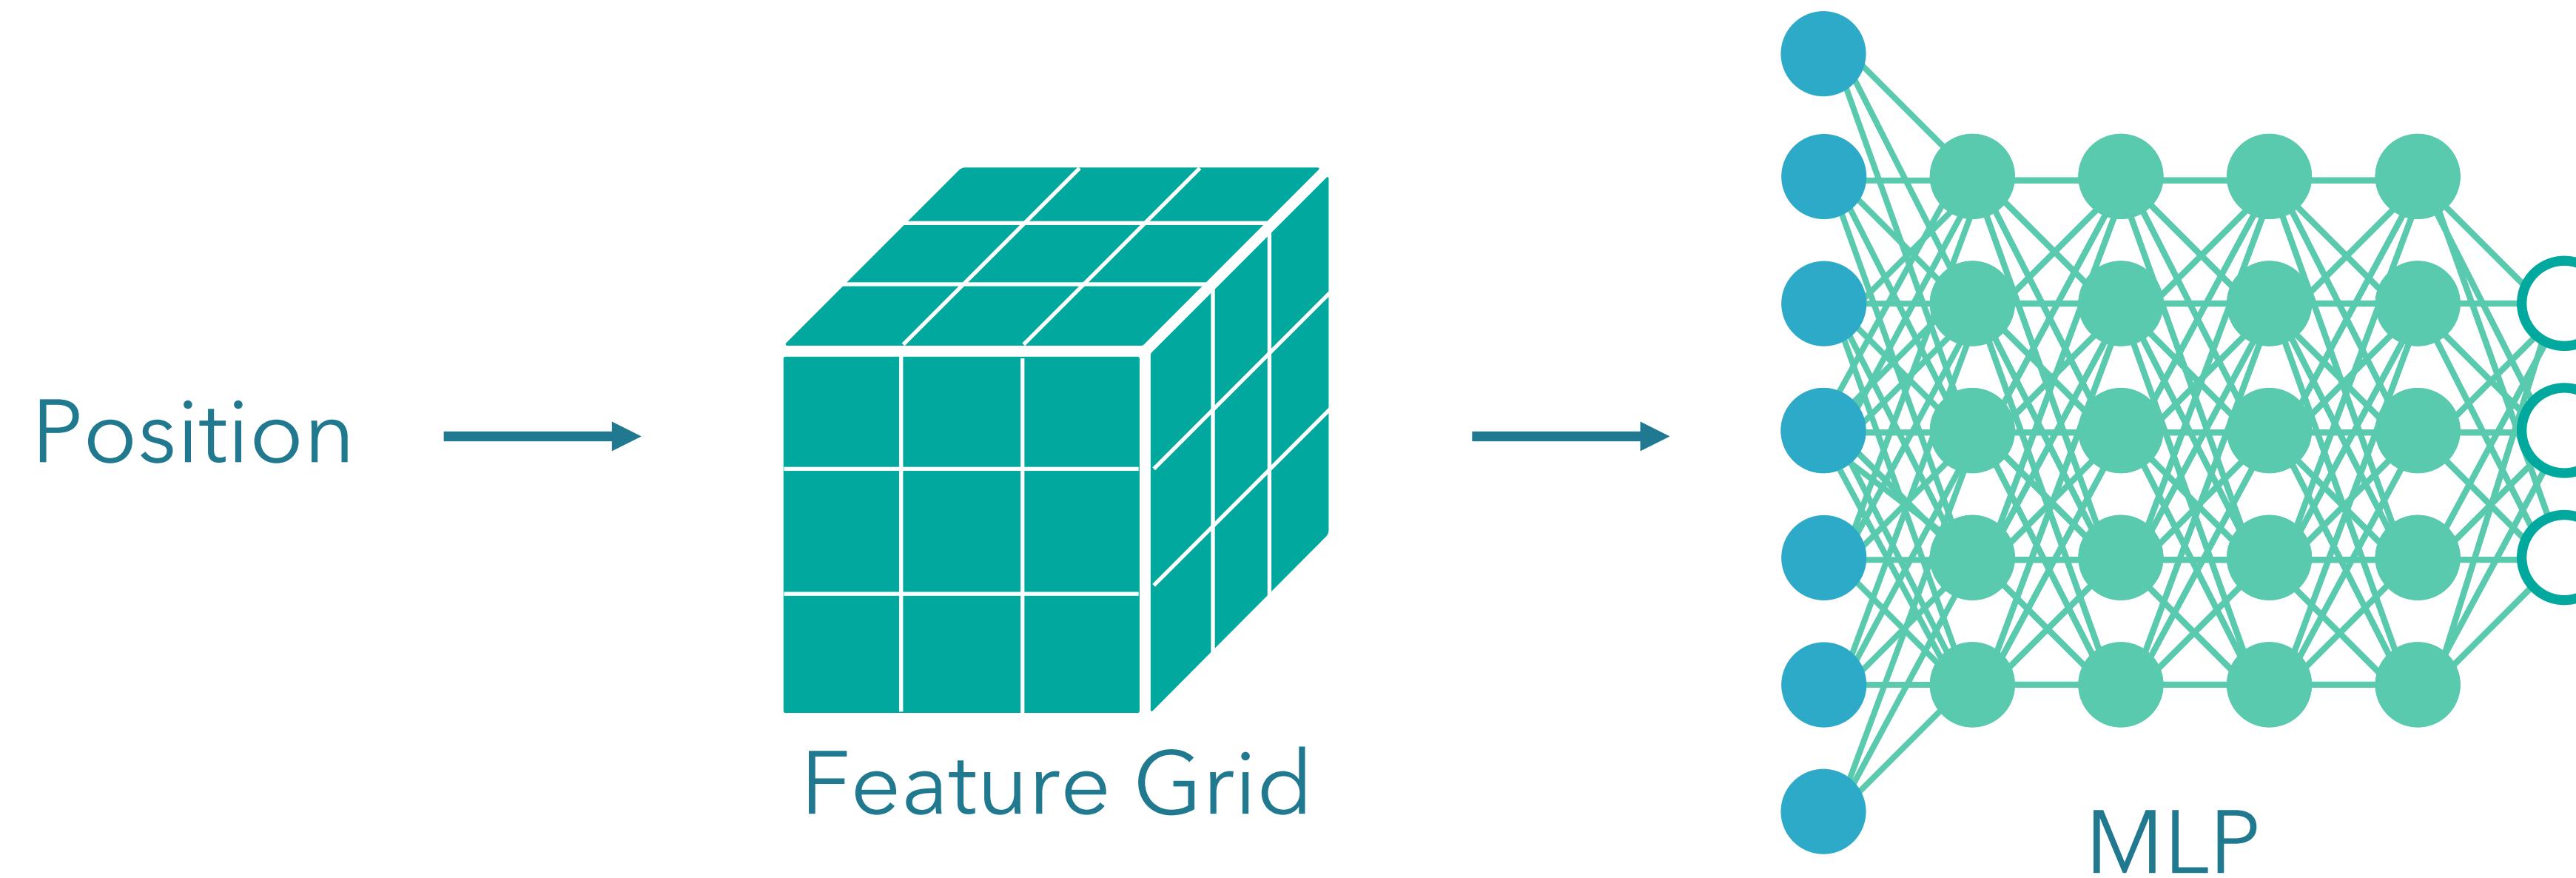
Feature Grid



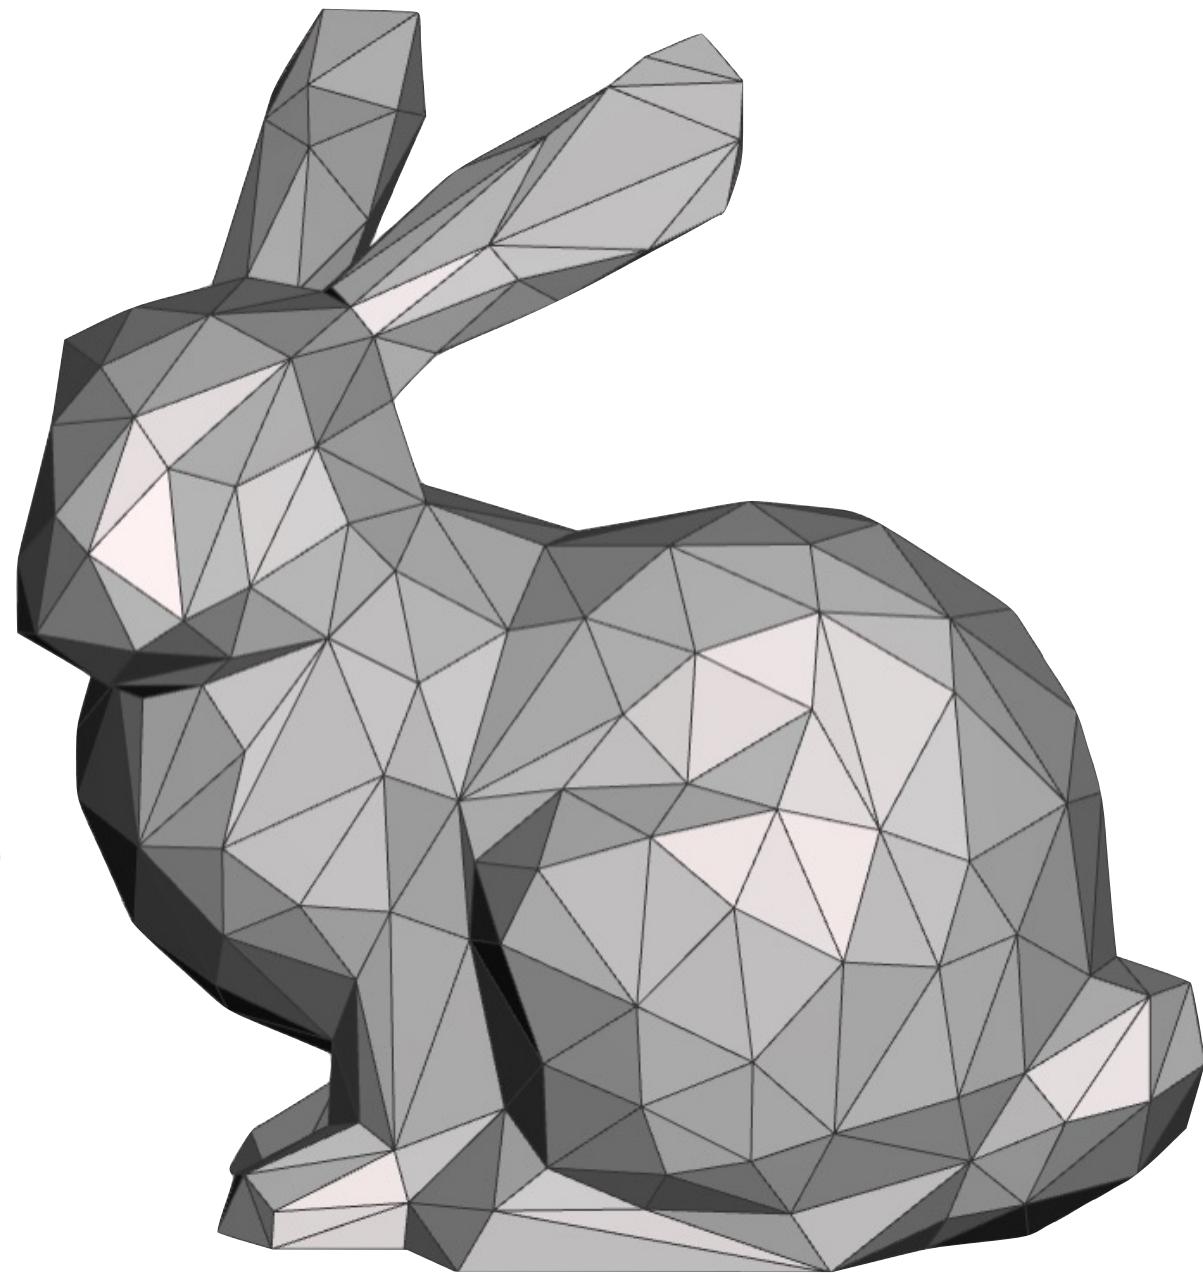
Feature Grid > Dictionary

Mapping with collisions

# Feature Grids as MLP Encodings

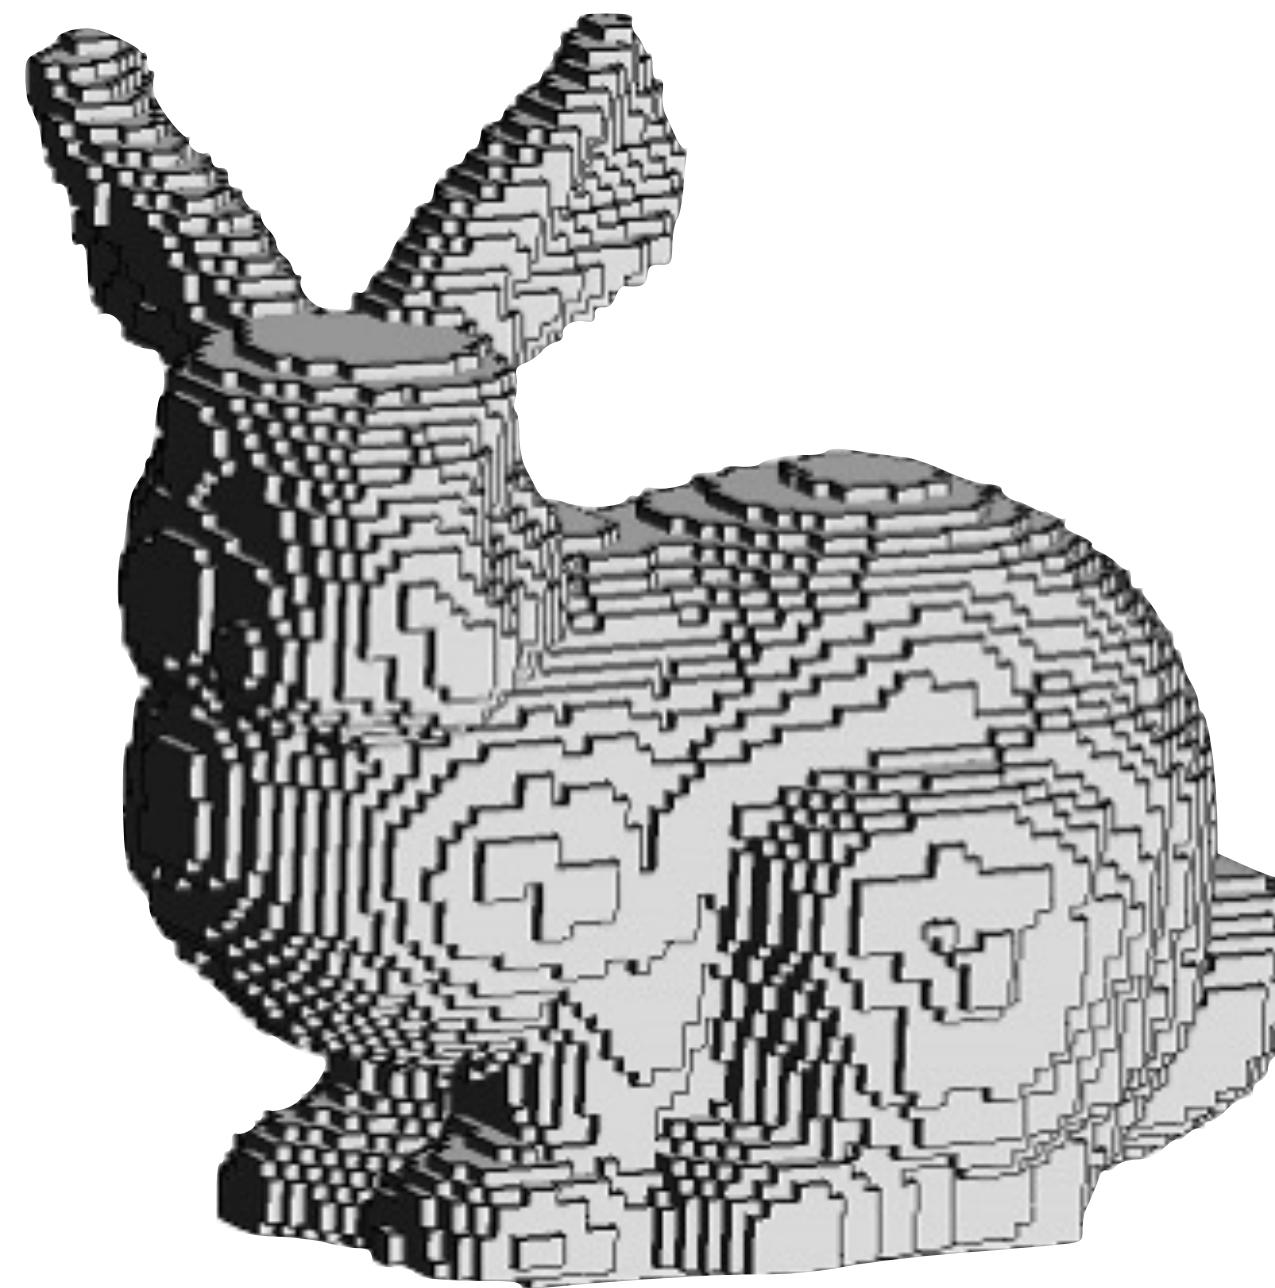


# Revisiting Geometry Representations



Mesh Representation

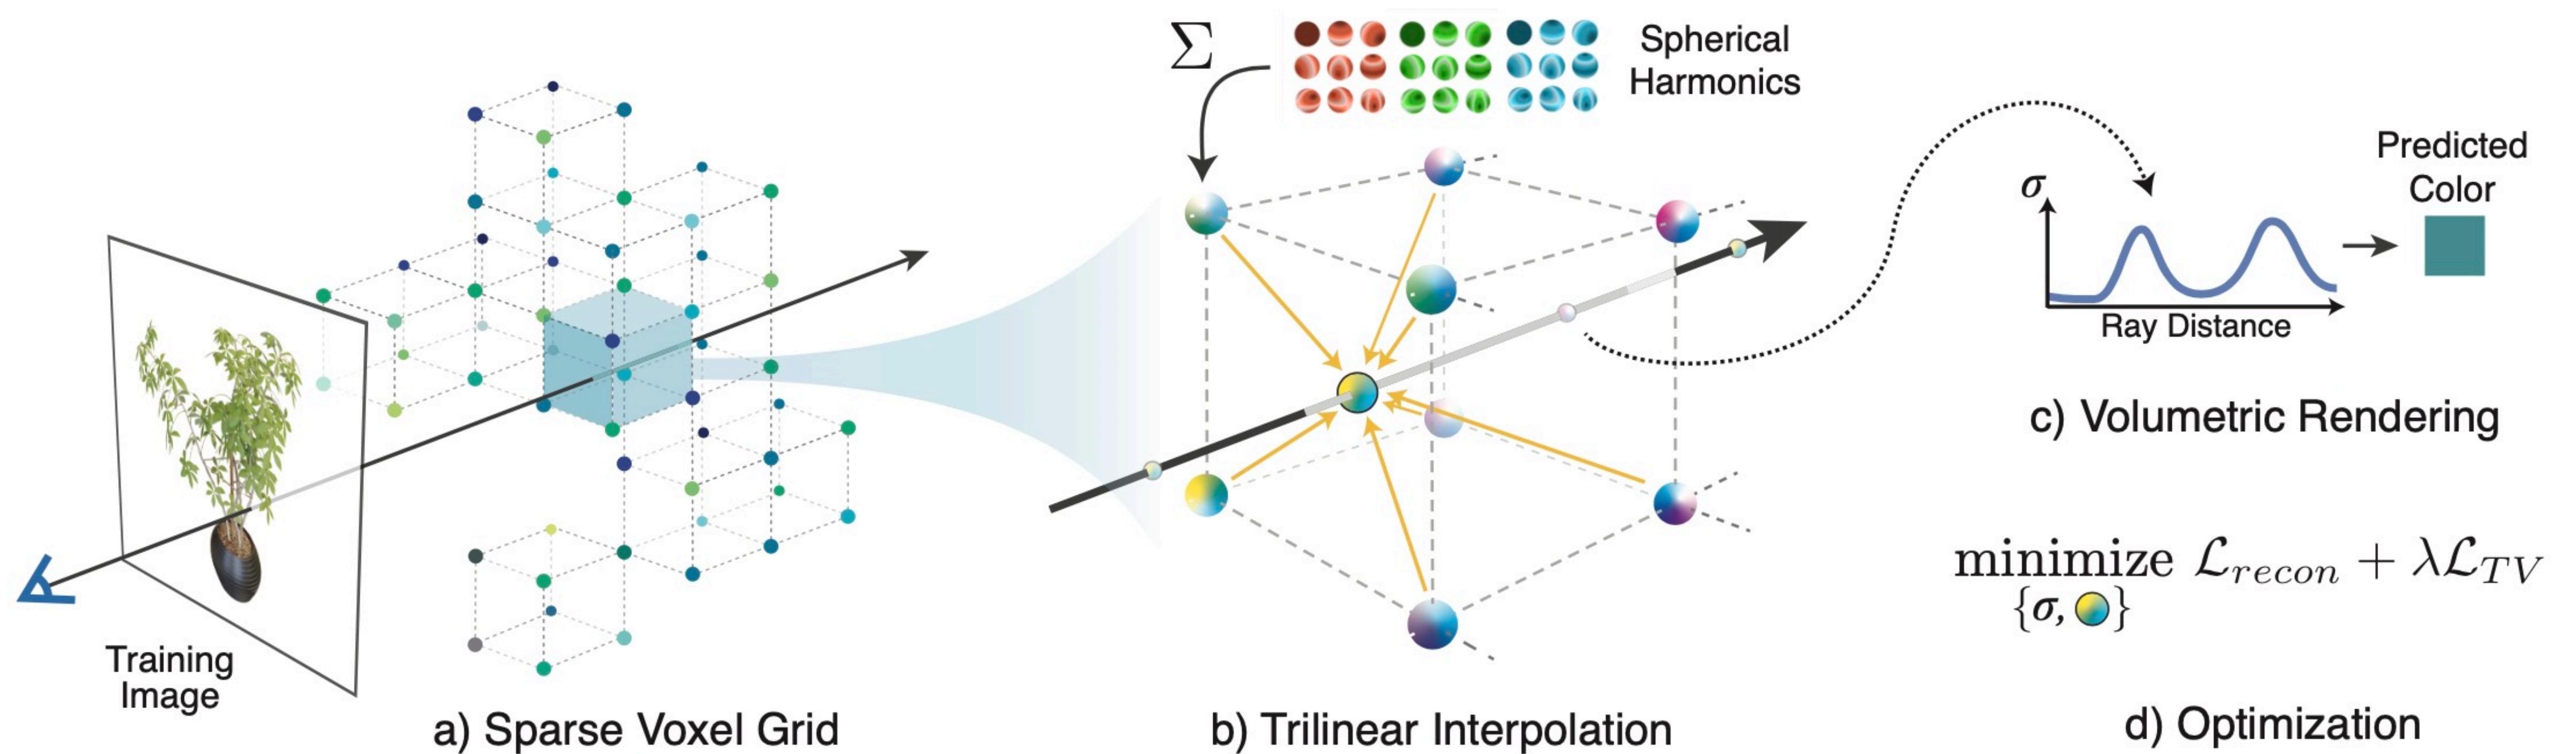
Small memory footprint  
Hard to optimize



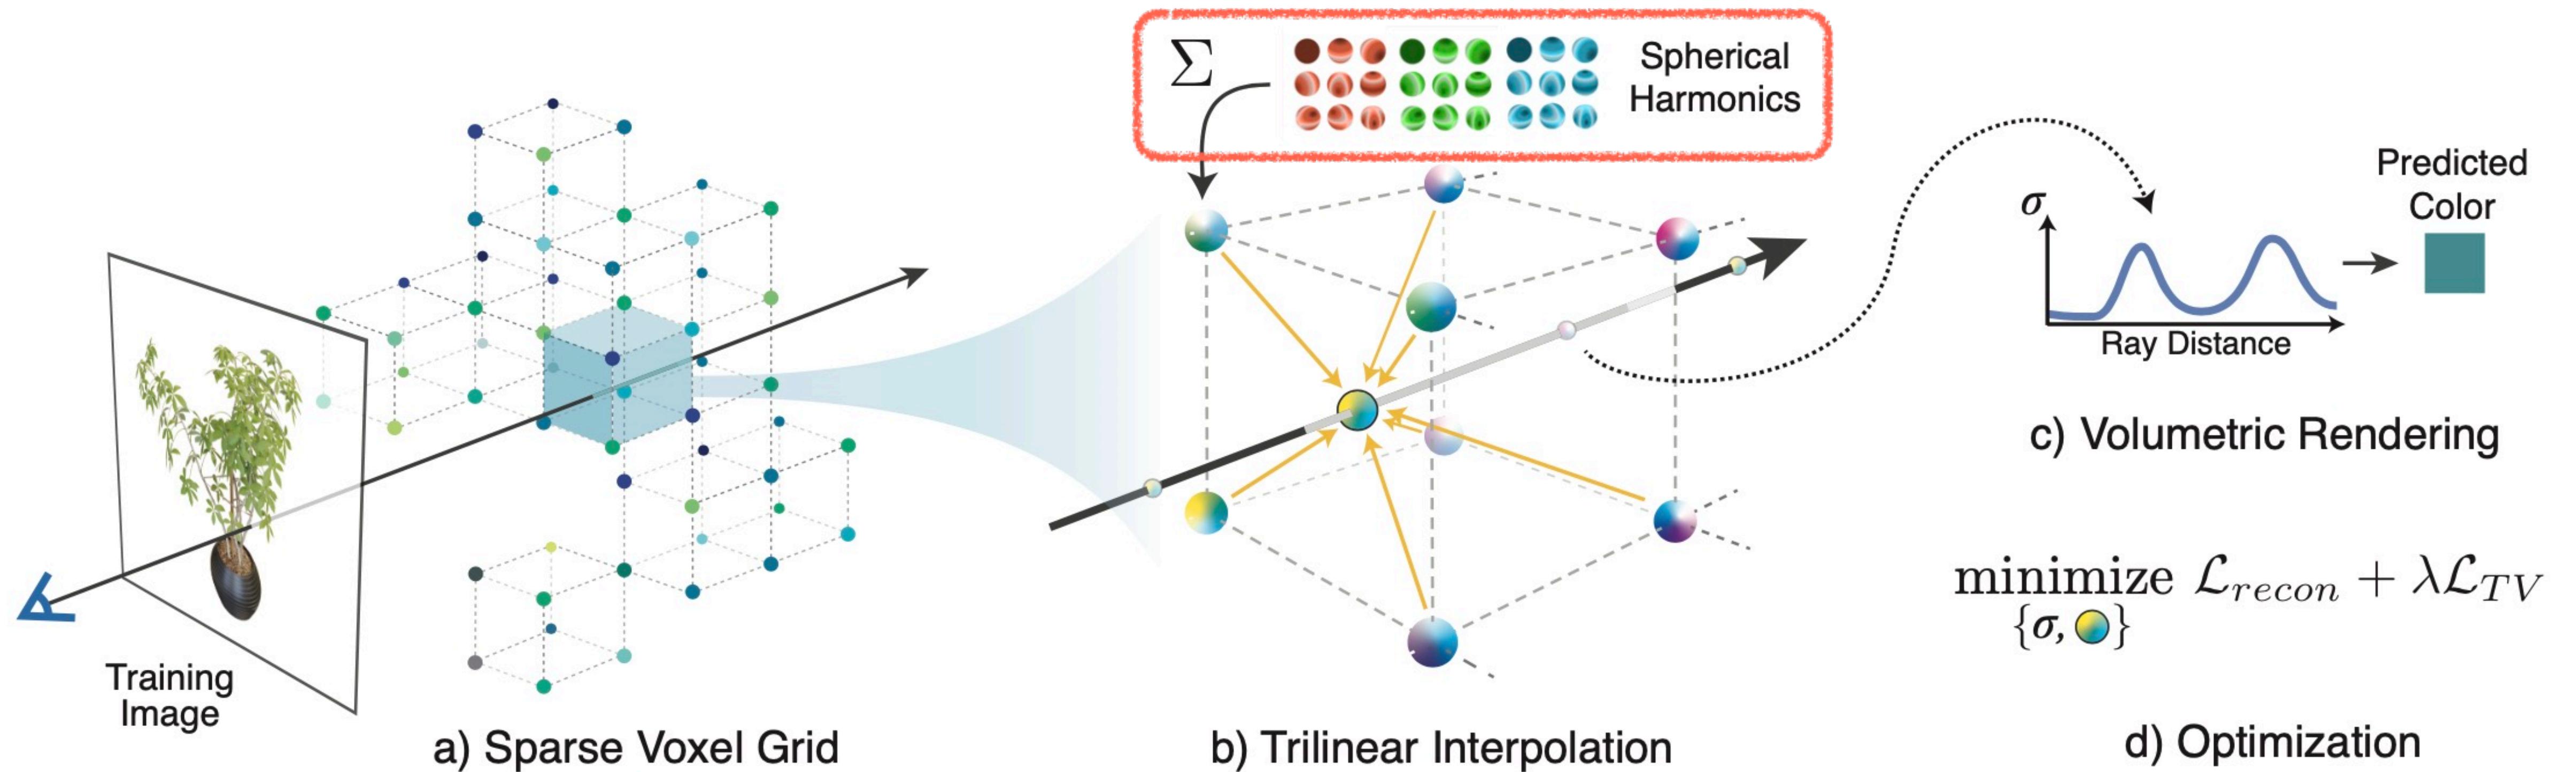
Voxel Representation

? Easy to optimize  
Large memory footprint ?

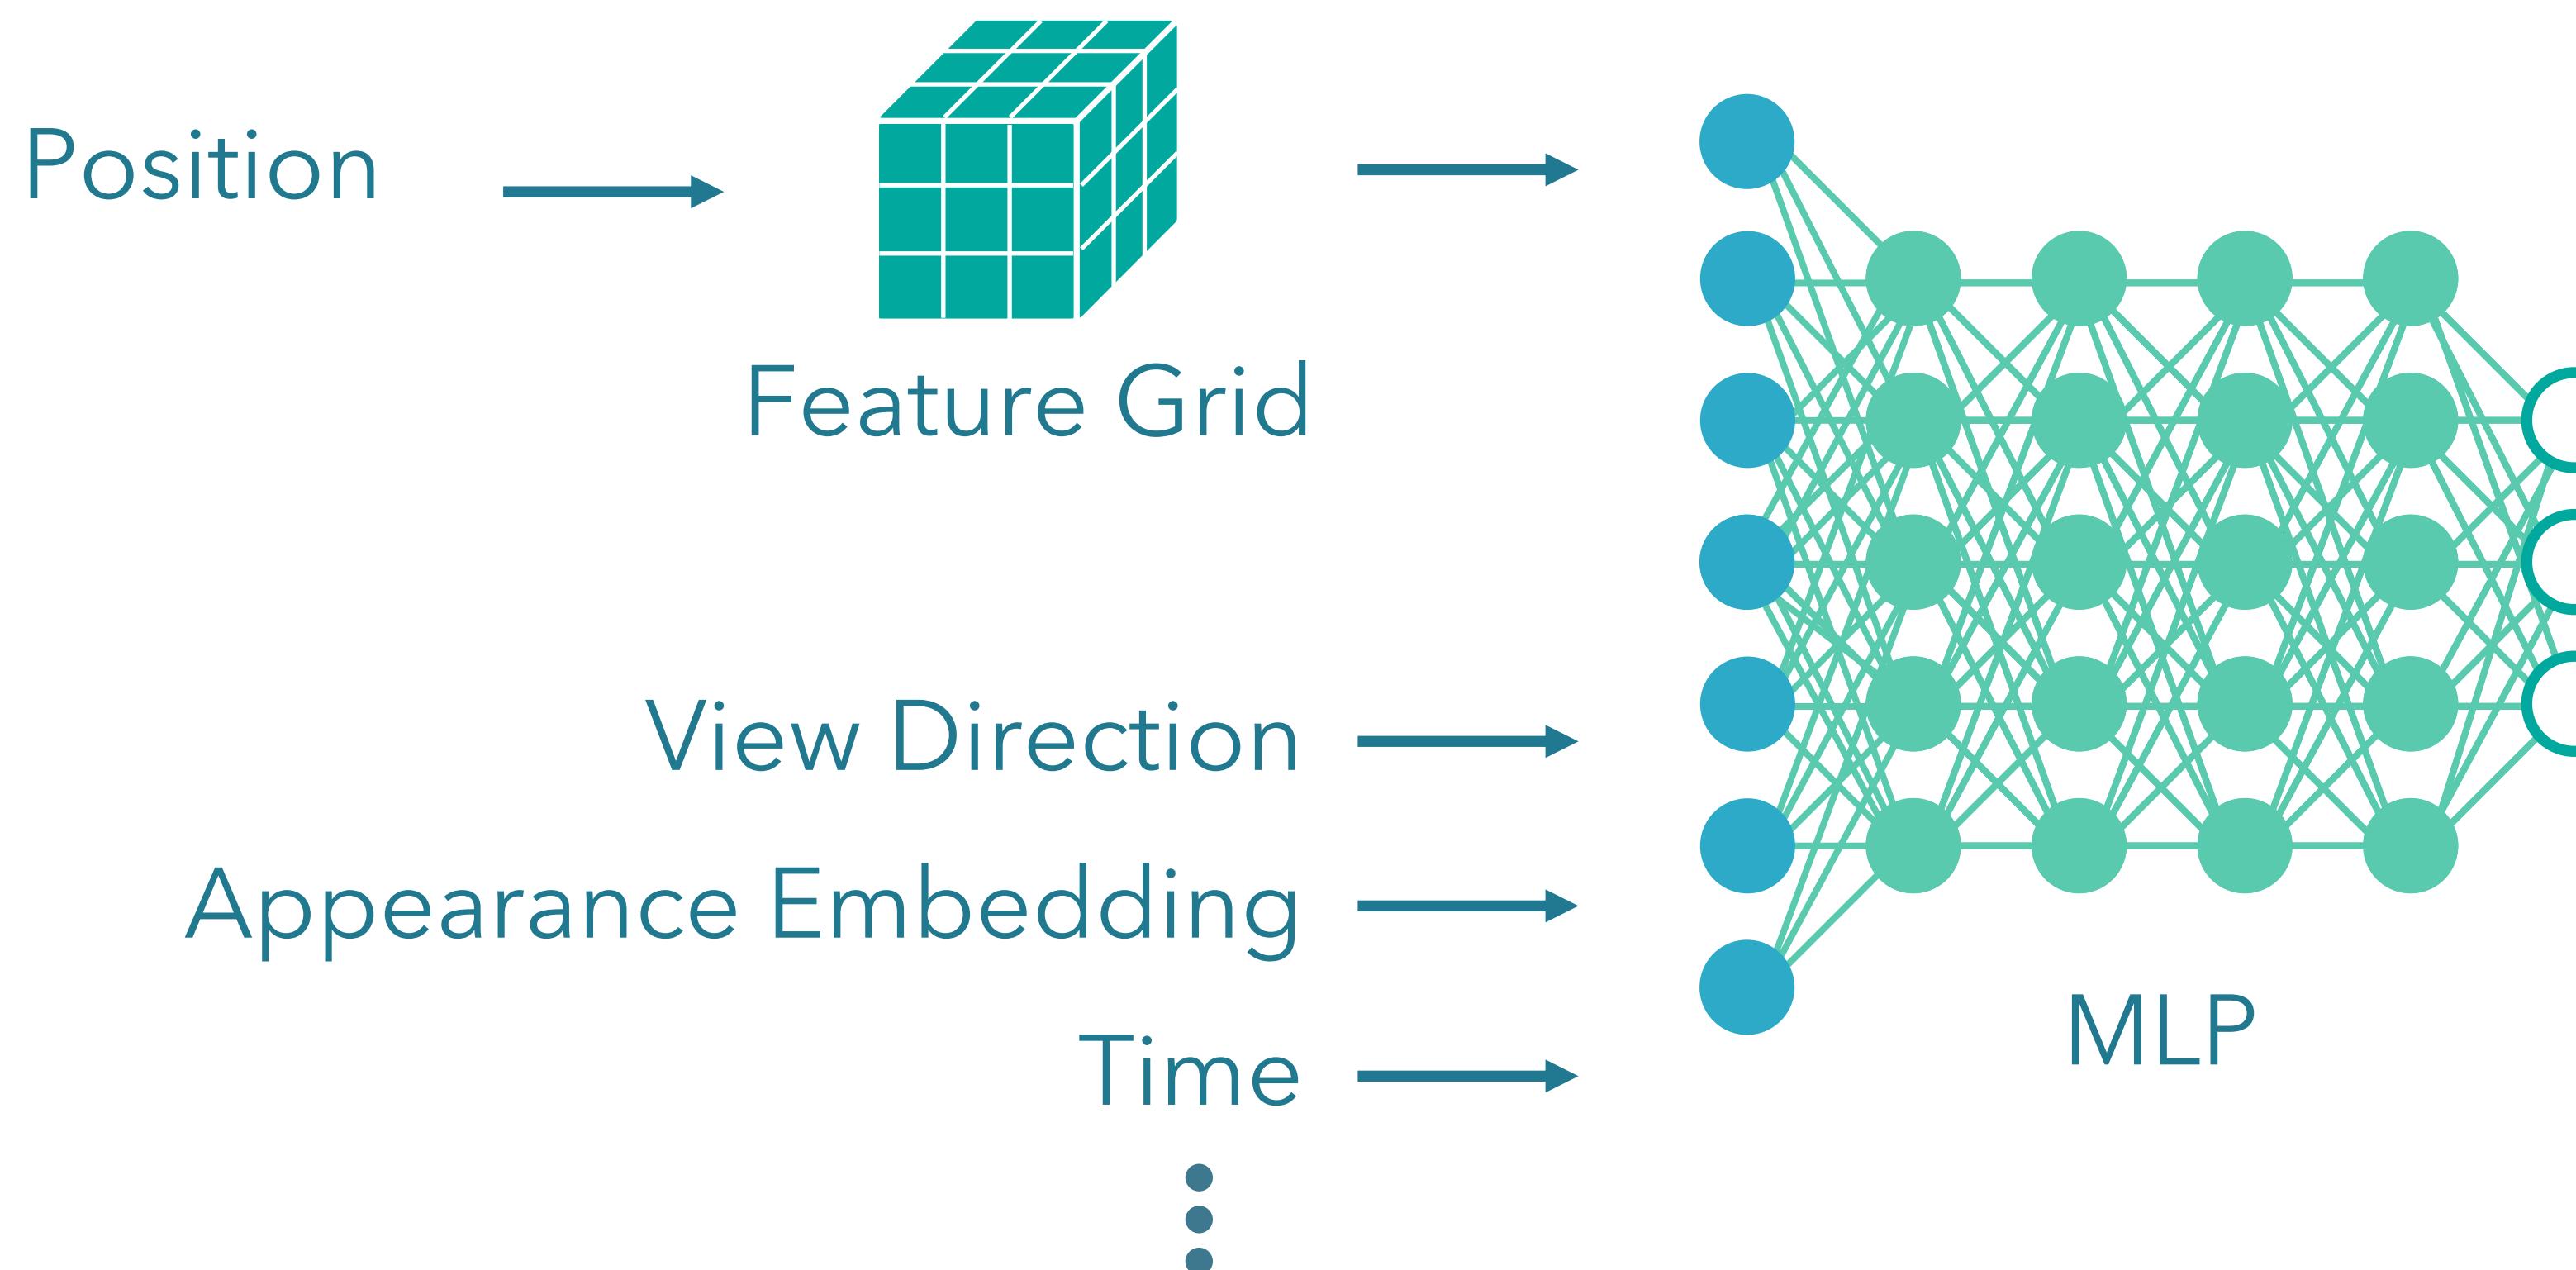
# MLPs are not required...



# MLPs are not required...



# But MLPs are convenient



# Where we are

1. Birds Eye View & Background
2. Volumetric Rendering Function
3. Encoding and Representing 3D Volumes
- 4. Signal Processing Considerations**
5. Challenges & Pointers

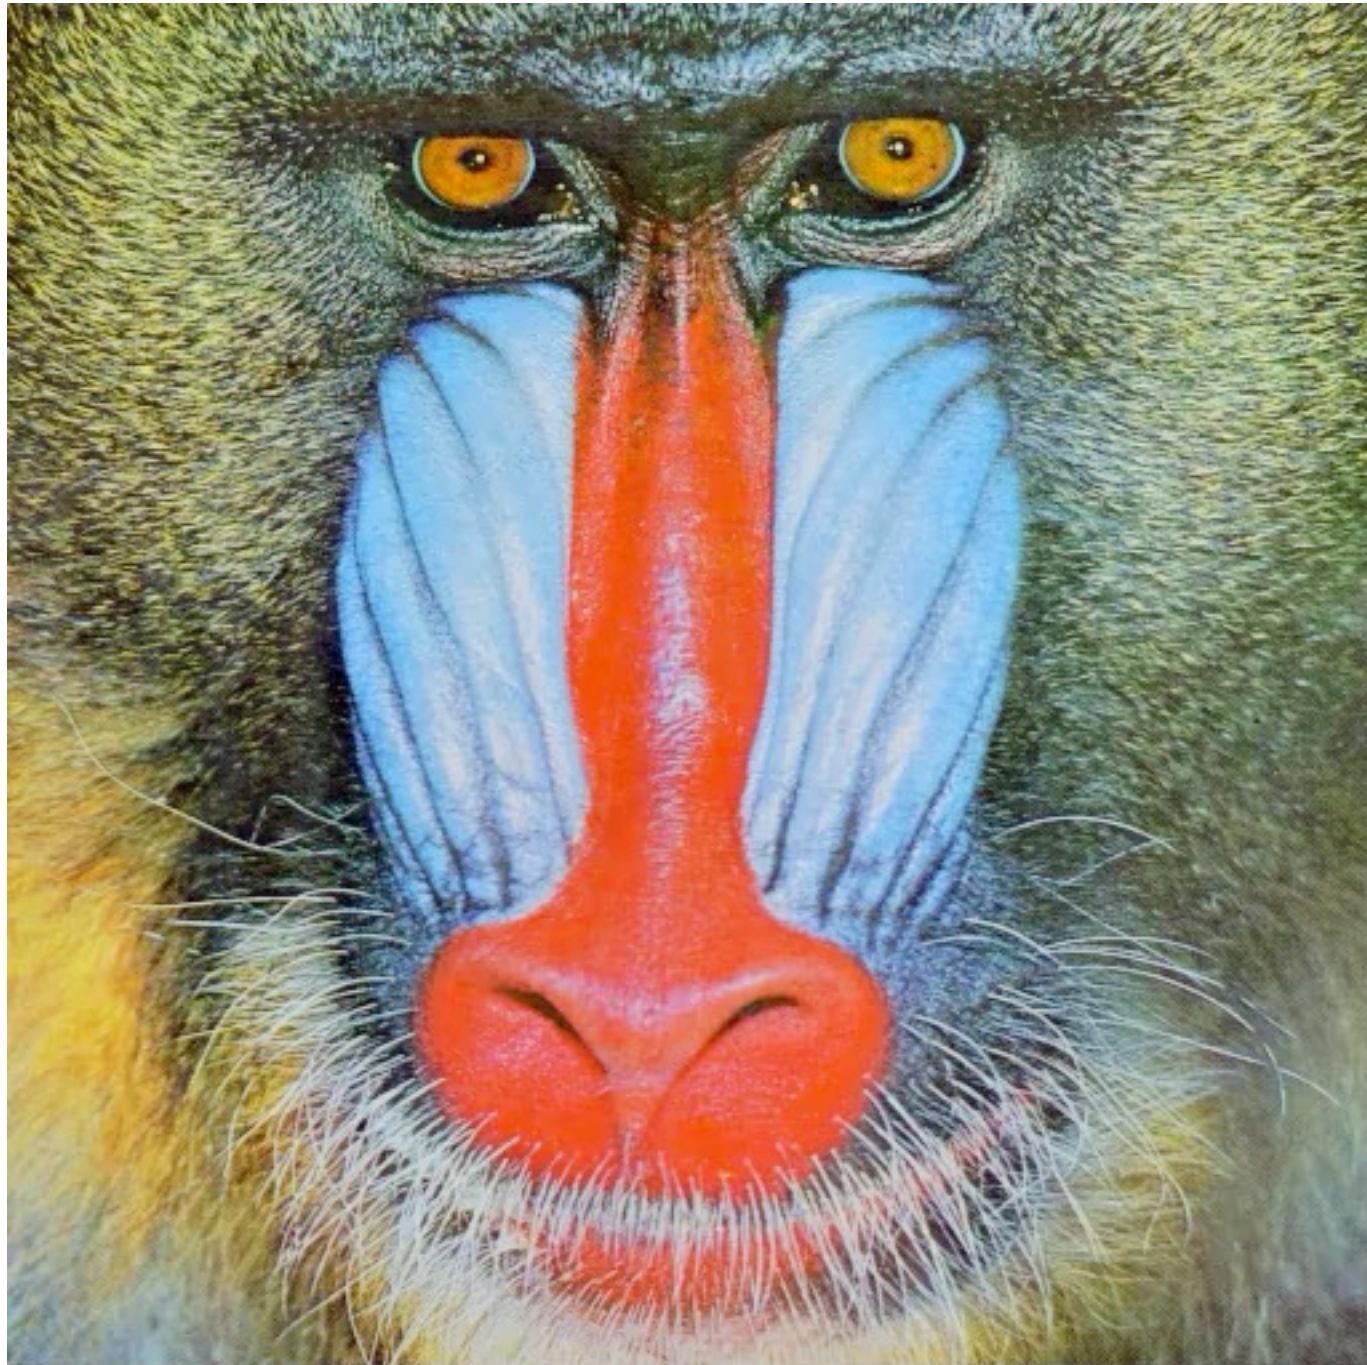
# **Signal Processing Considerations in NeRF**

# What is happening here?

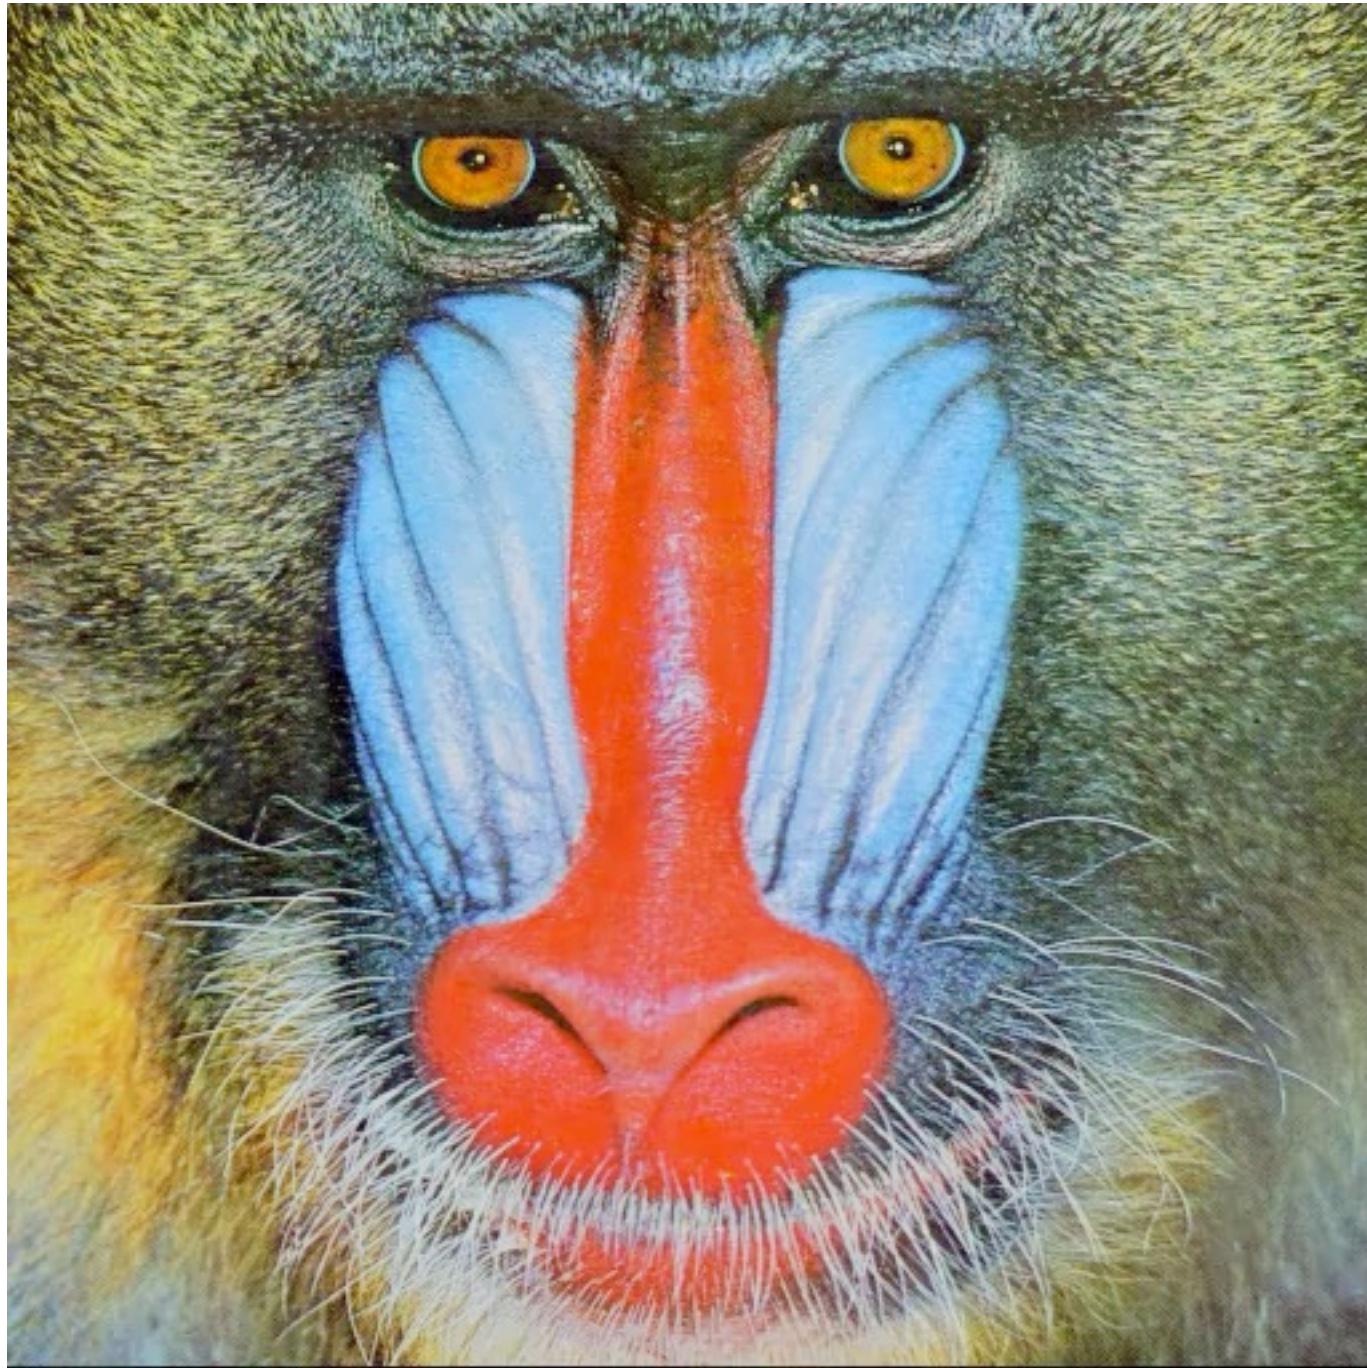


Naïve (original) NeRF

# Review: Aliasing in Image Processing

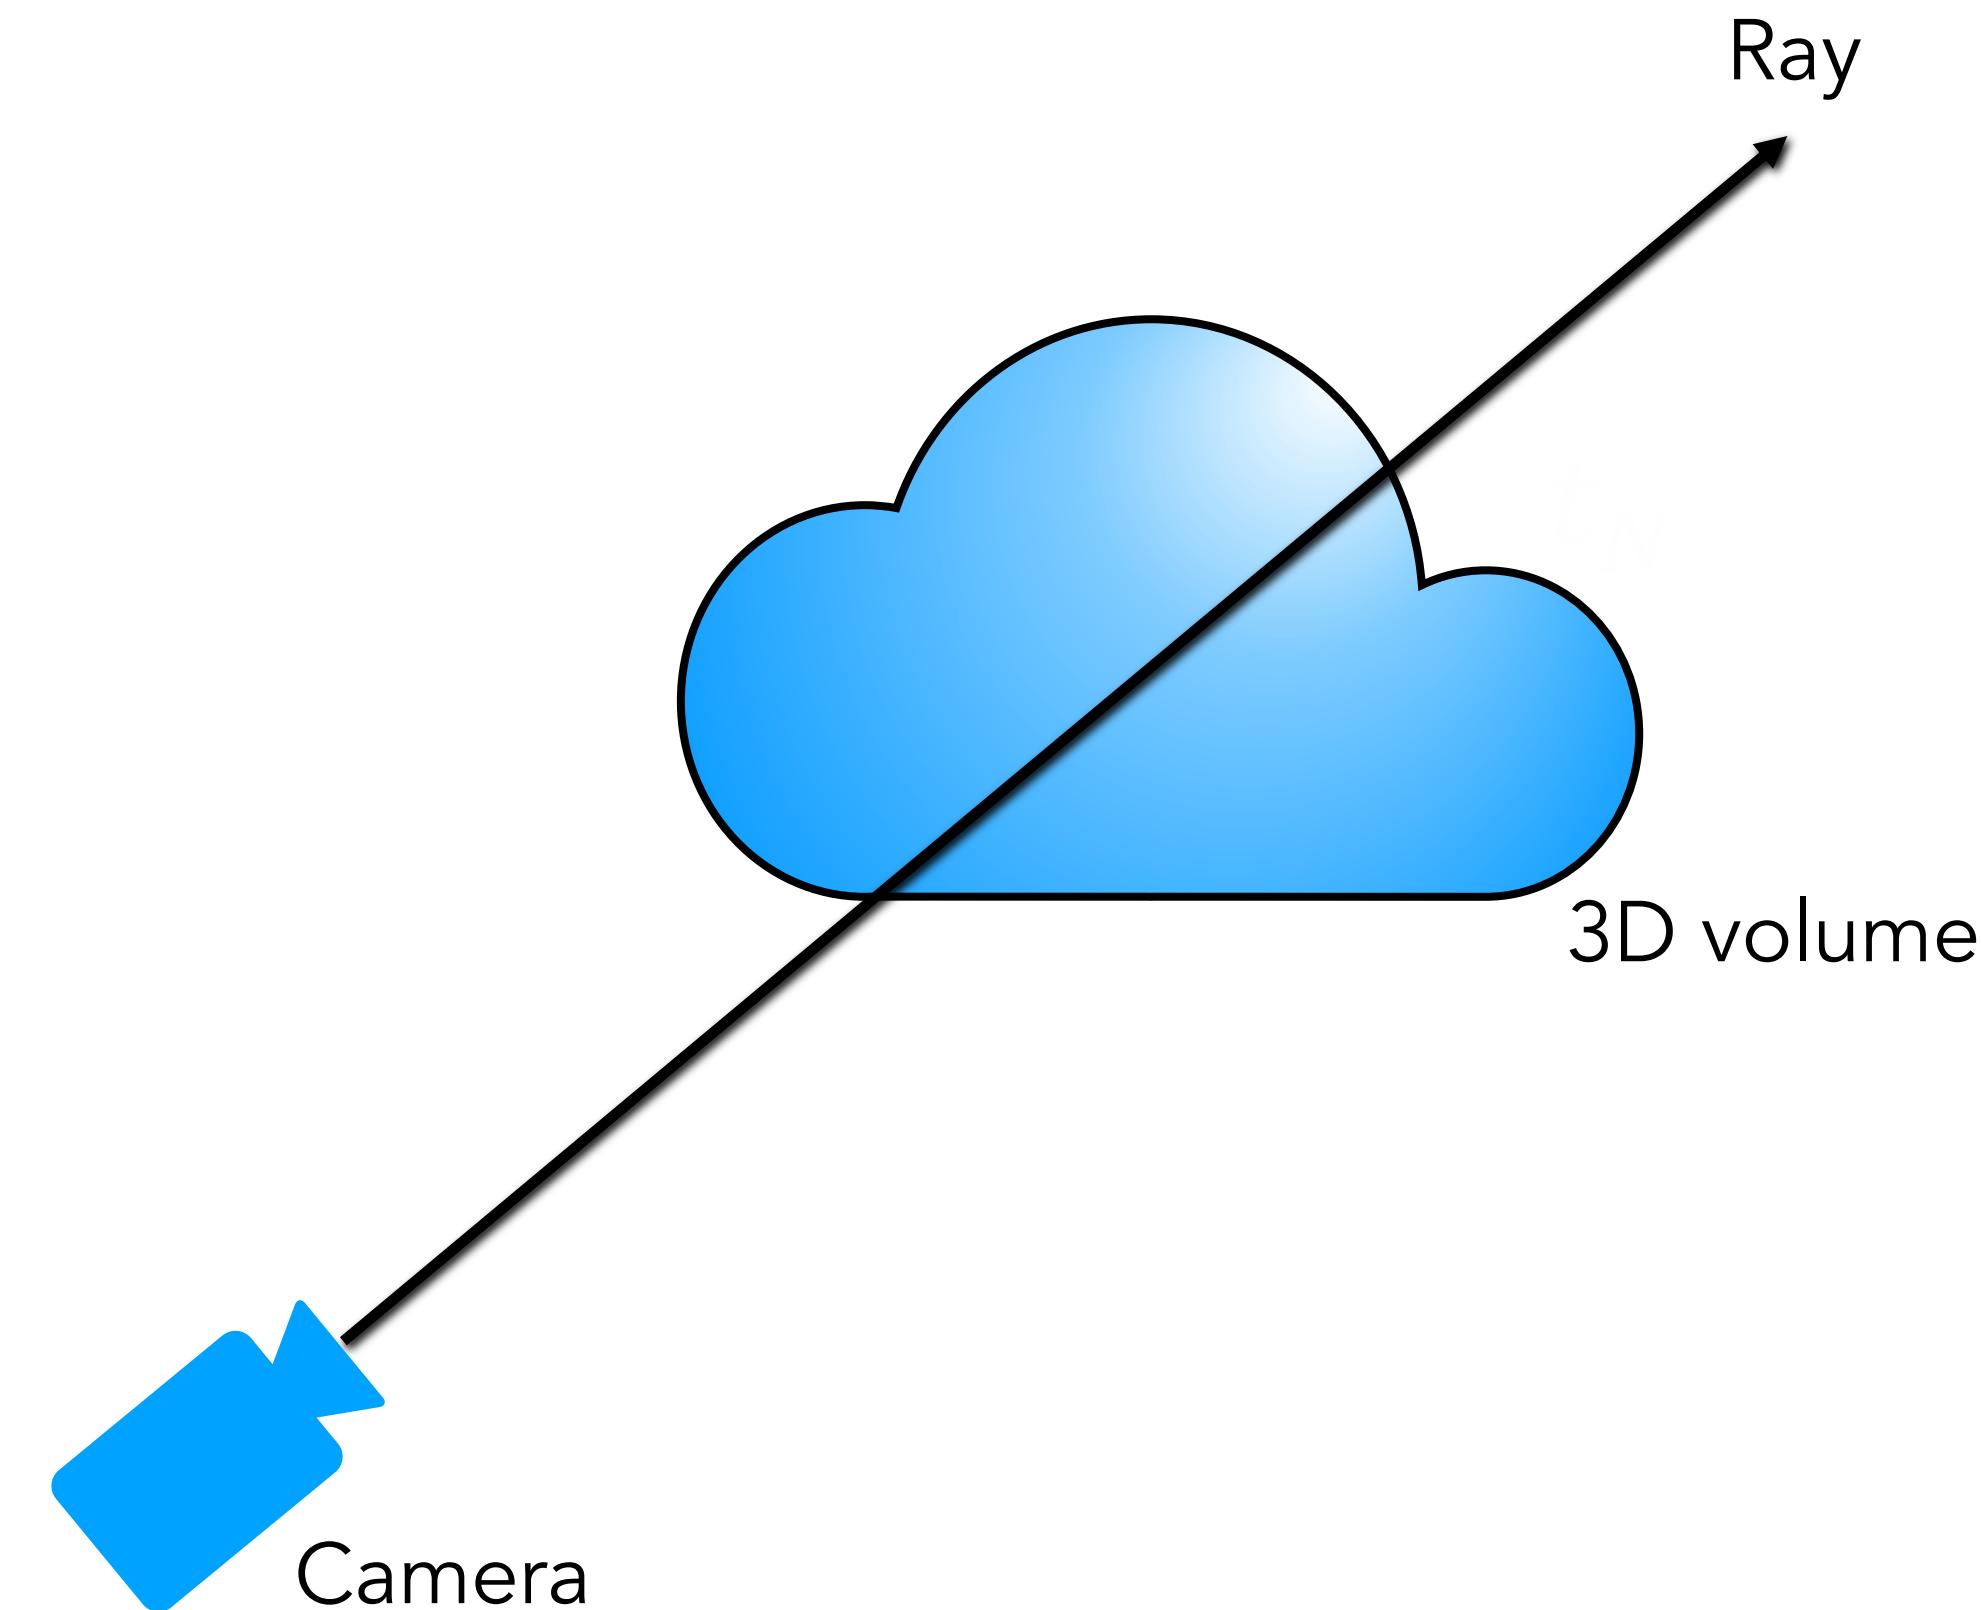


# Review: Aliasing in Image Processing

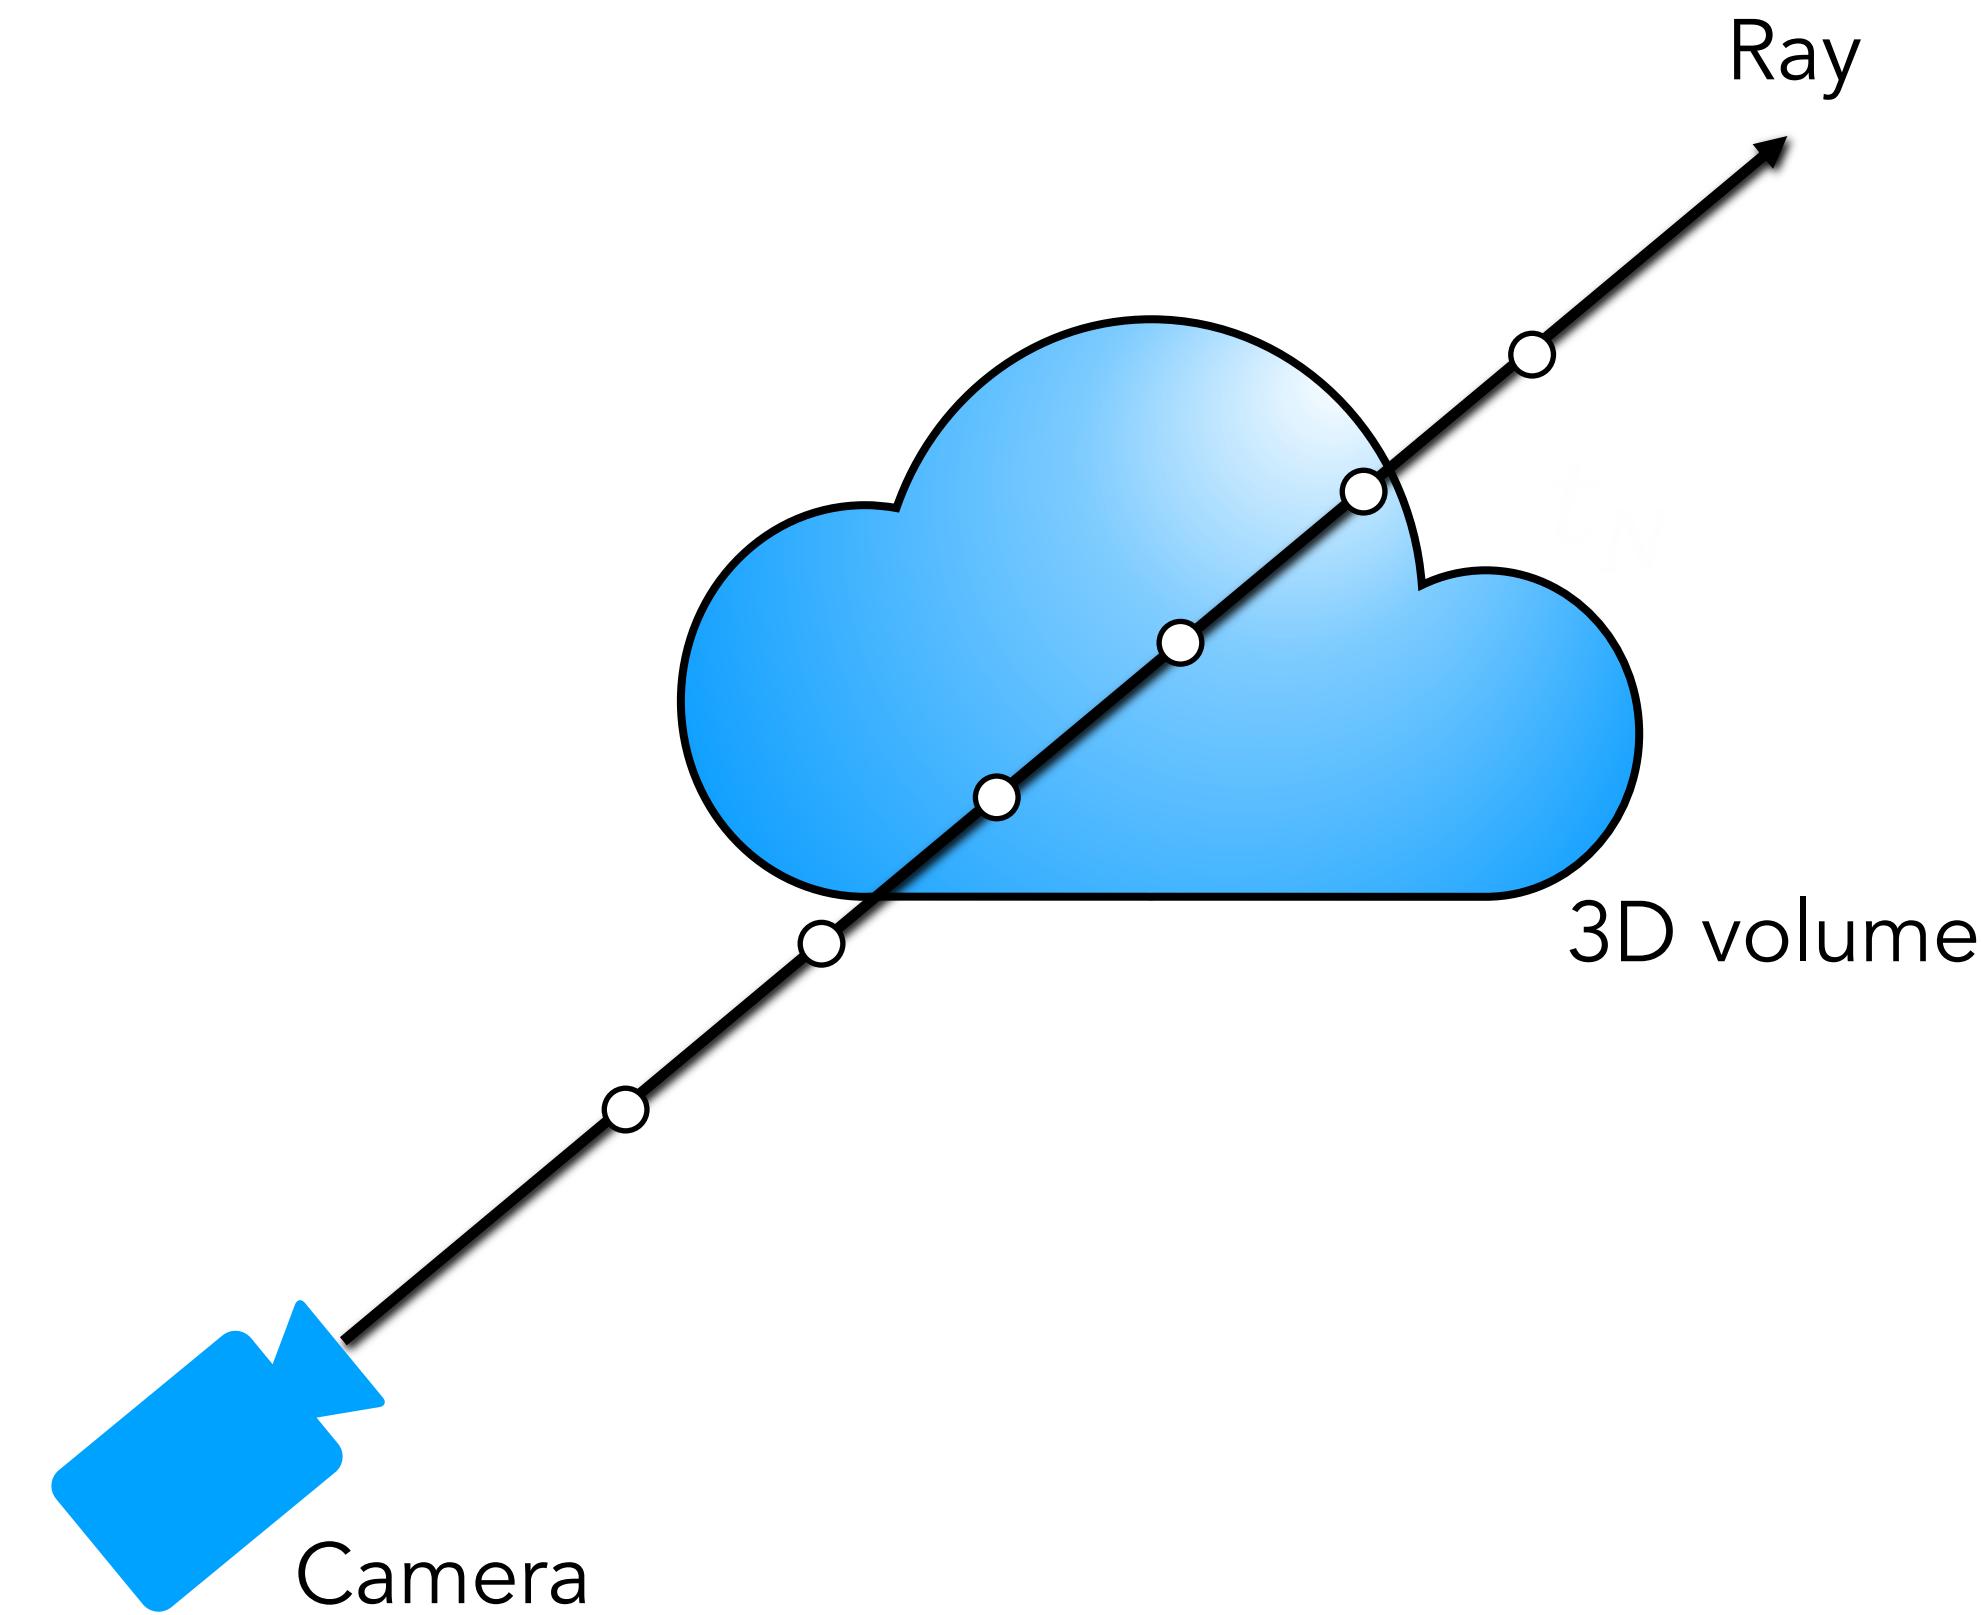


# Sampling Along Rays

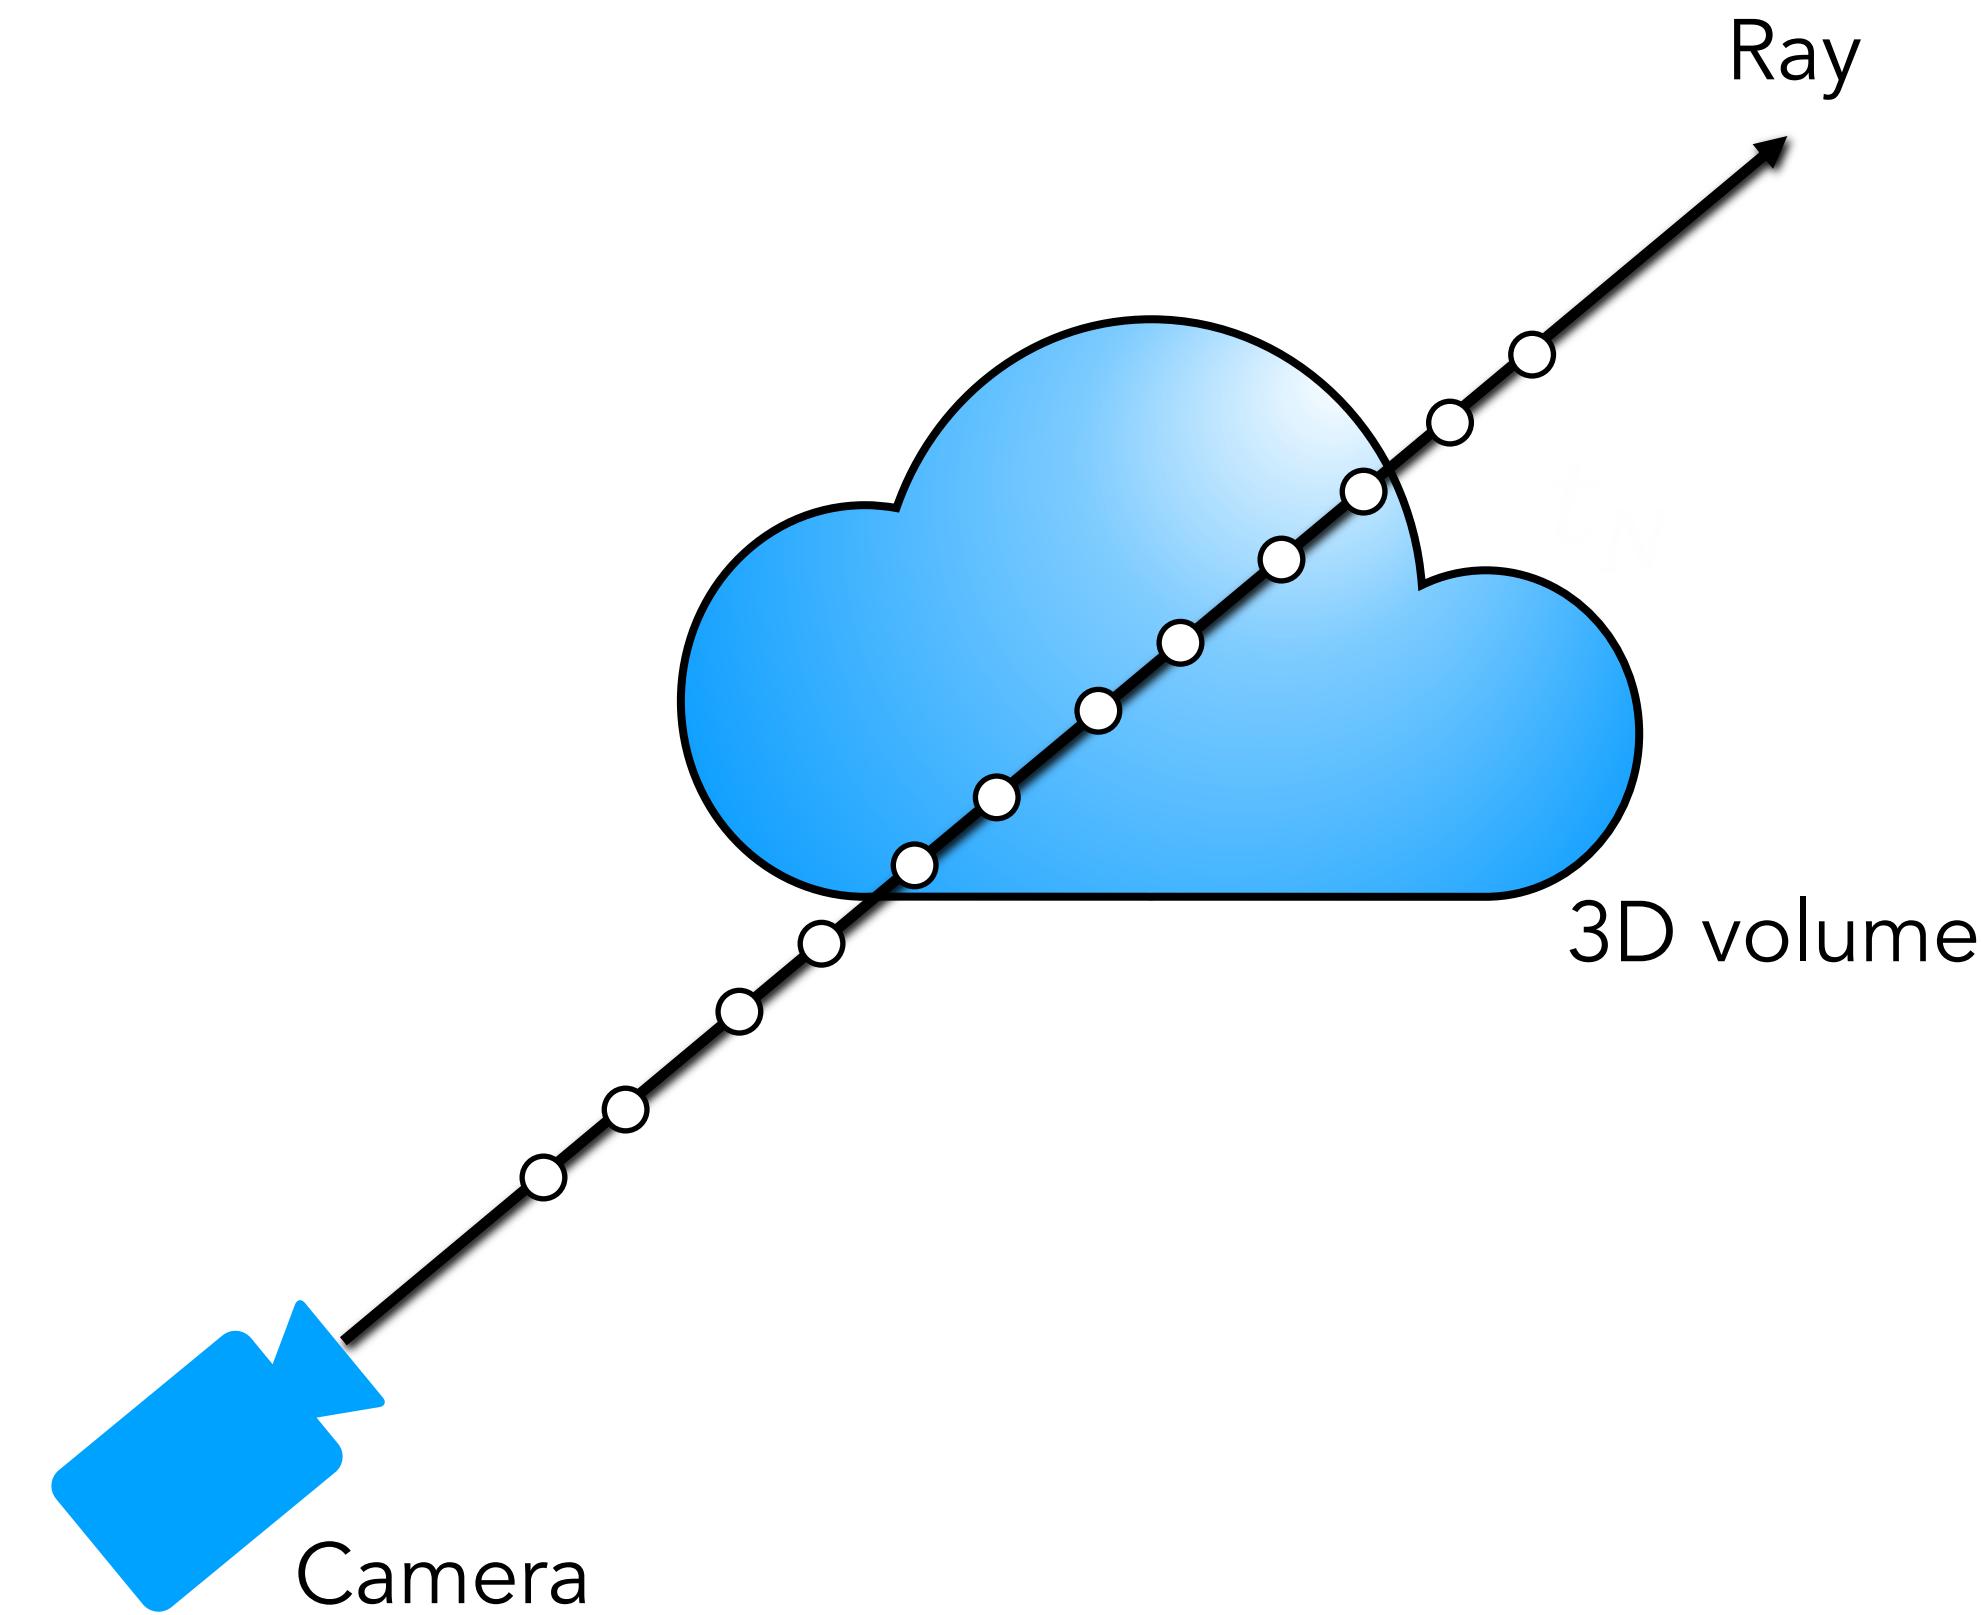
# Where to place samples along rays?



# How to be more efficient than dense sampling?



# How to be more efficient than dense sampling?



# Hierarchical Sampling vs. Acceleration Structures

# Hierarchical Sampling vs. Acceleration Structures

## Hierarchical Sampling

Iteratively use samples from NeRF to more efficiently sample visible scene content

# Hierarchical Sampling vs. Acceleration Structures

## Hierarchical Sampling

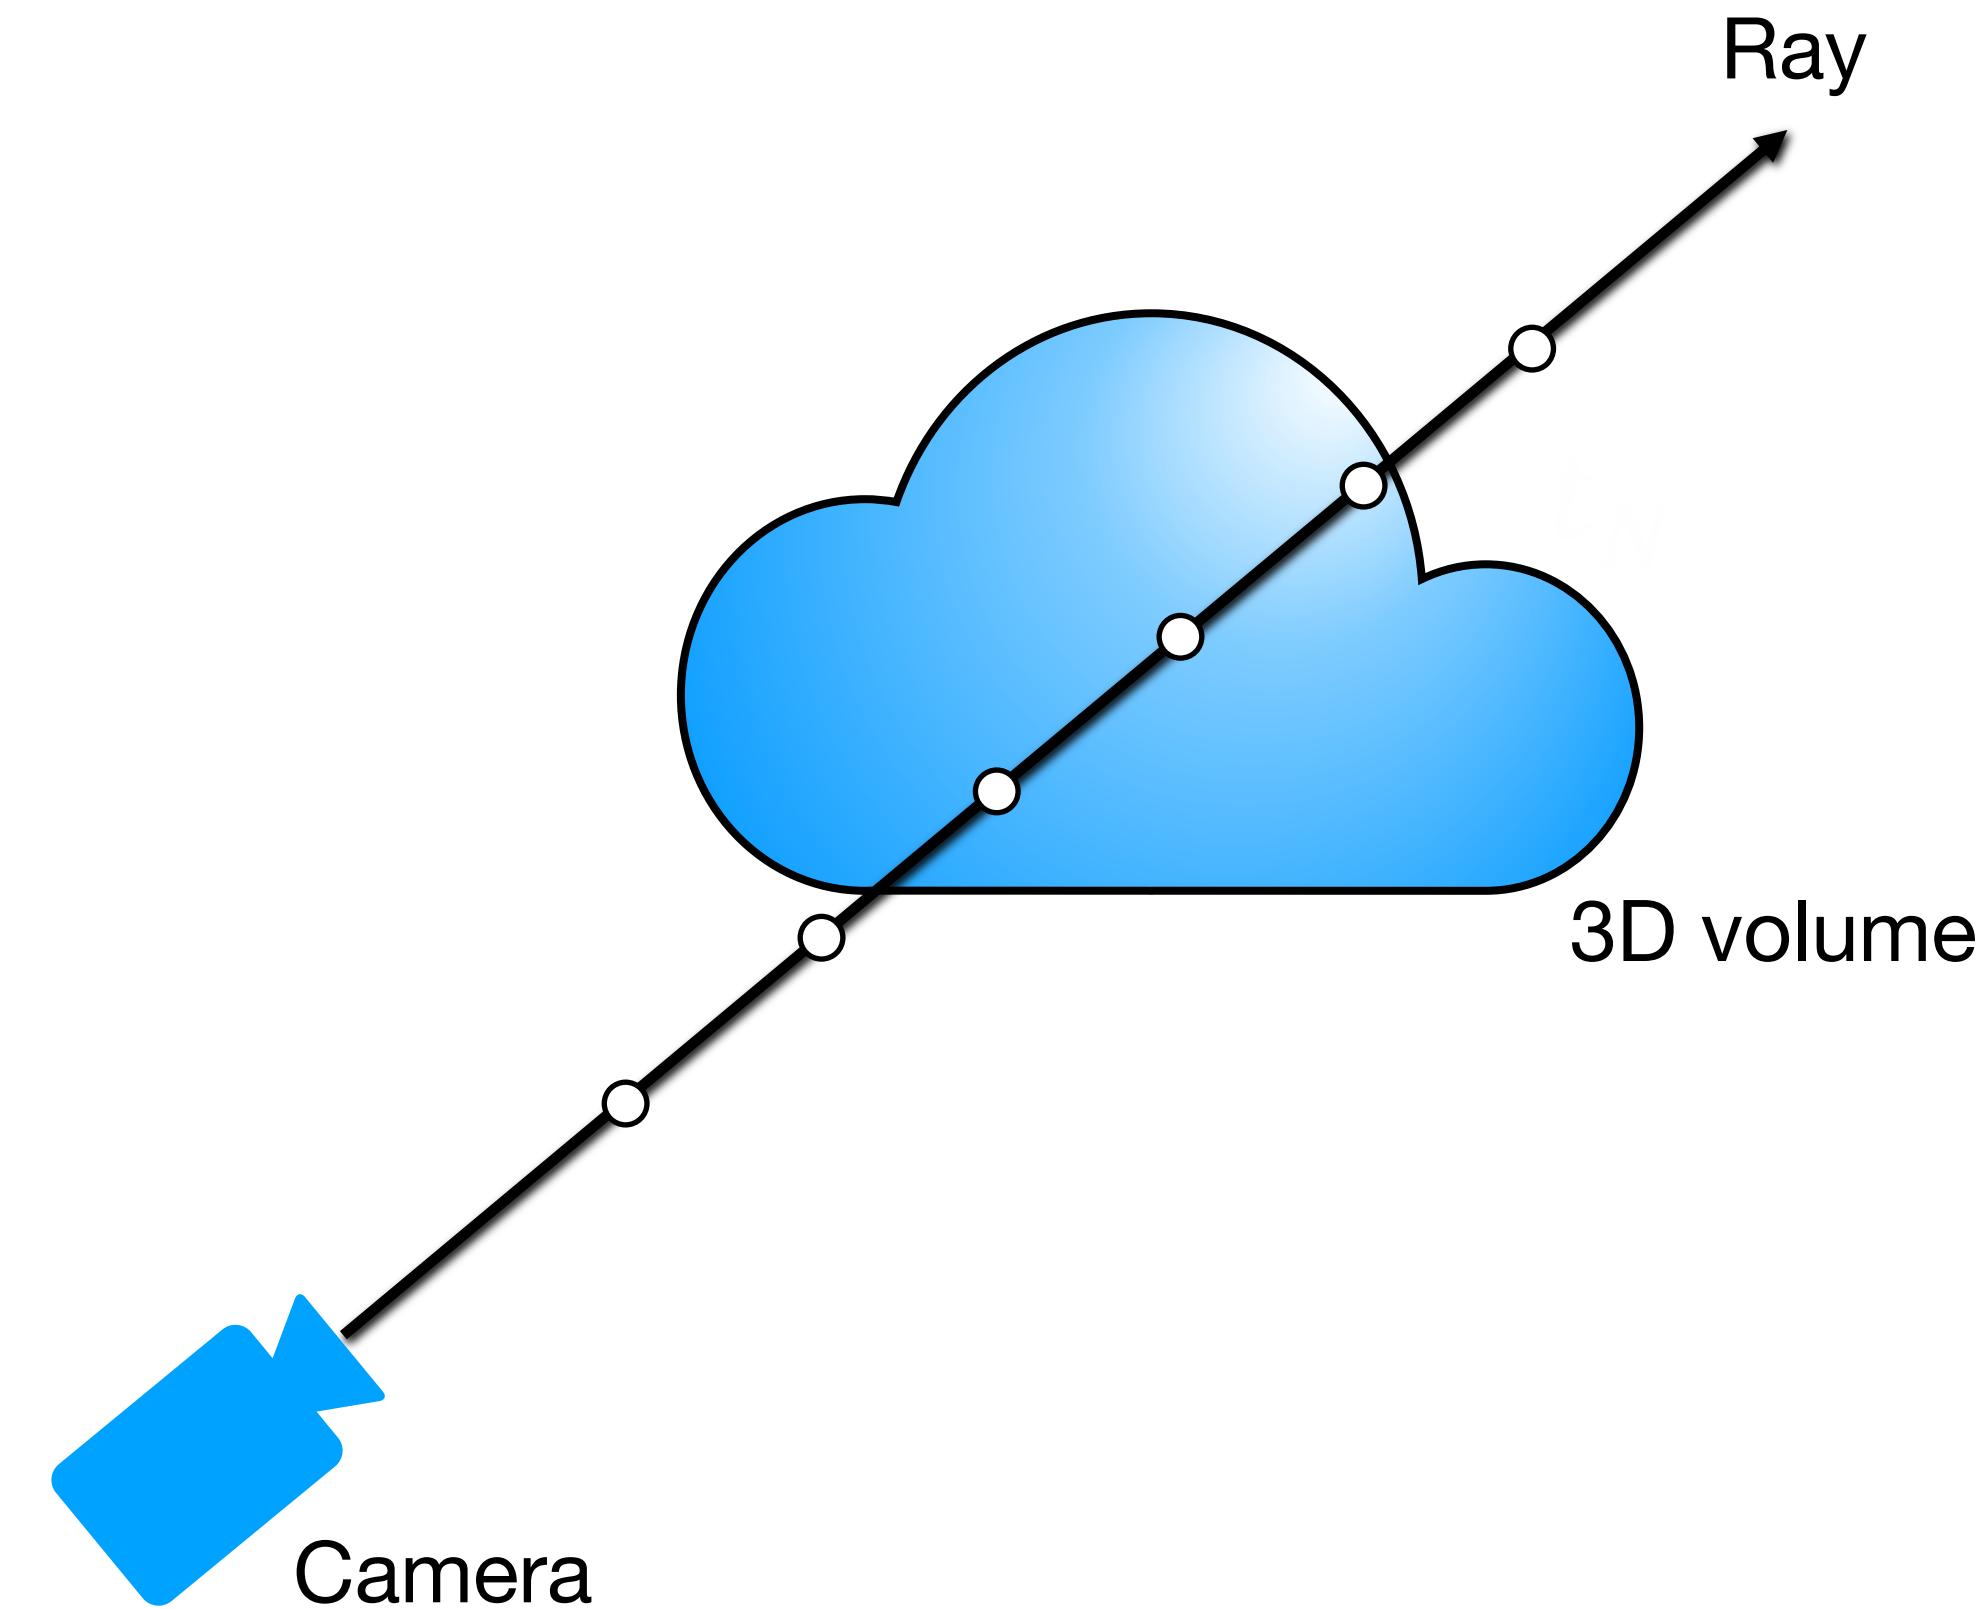
Iteratively use samples from NeRF to more efficiently sample visible scene content

## Acceleration Structures

Distill/cache properties of NeRF into a structure that helps generate samples

# Hierarchical ray sampling

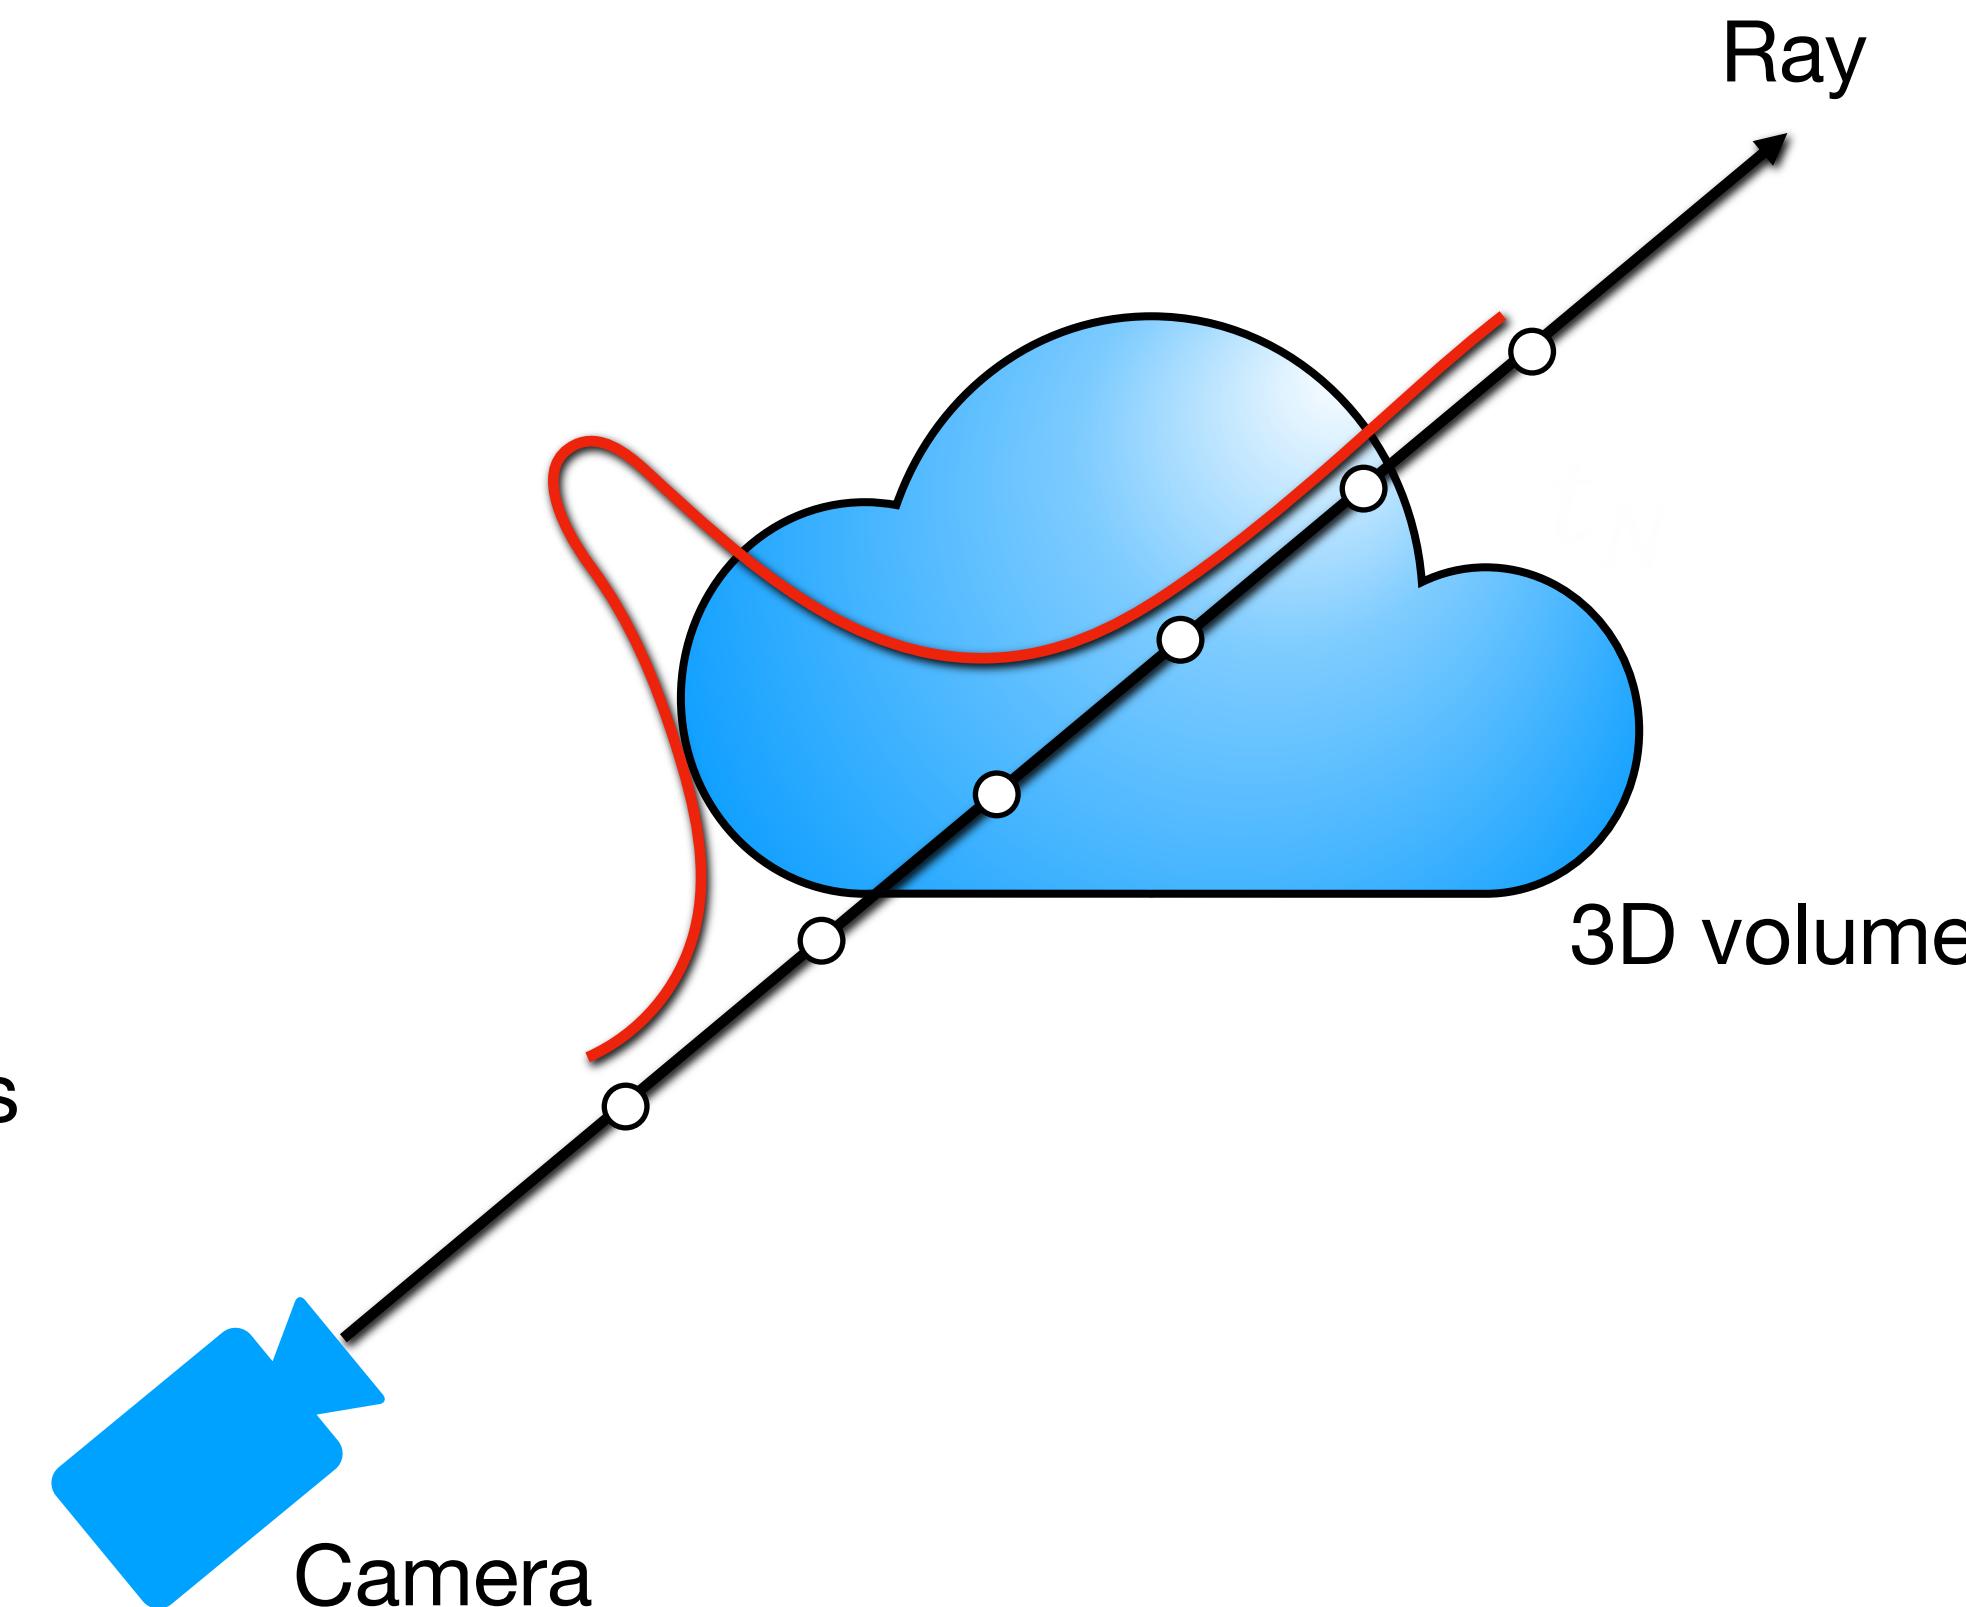
**Key Idea: sample points proportionally to expected effect on final rendering**



# Key Idea: sample points proportionally to expected effect on final rendering

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

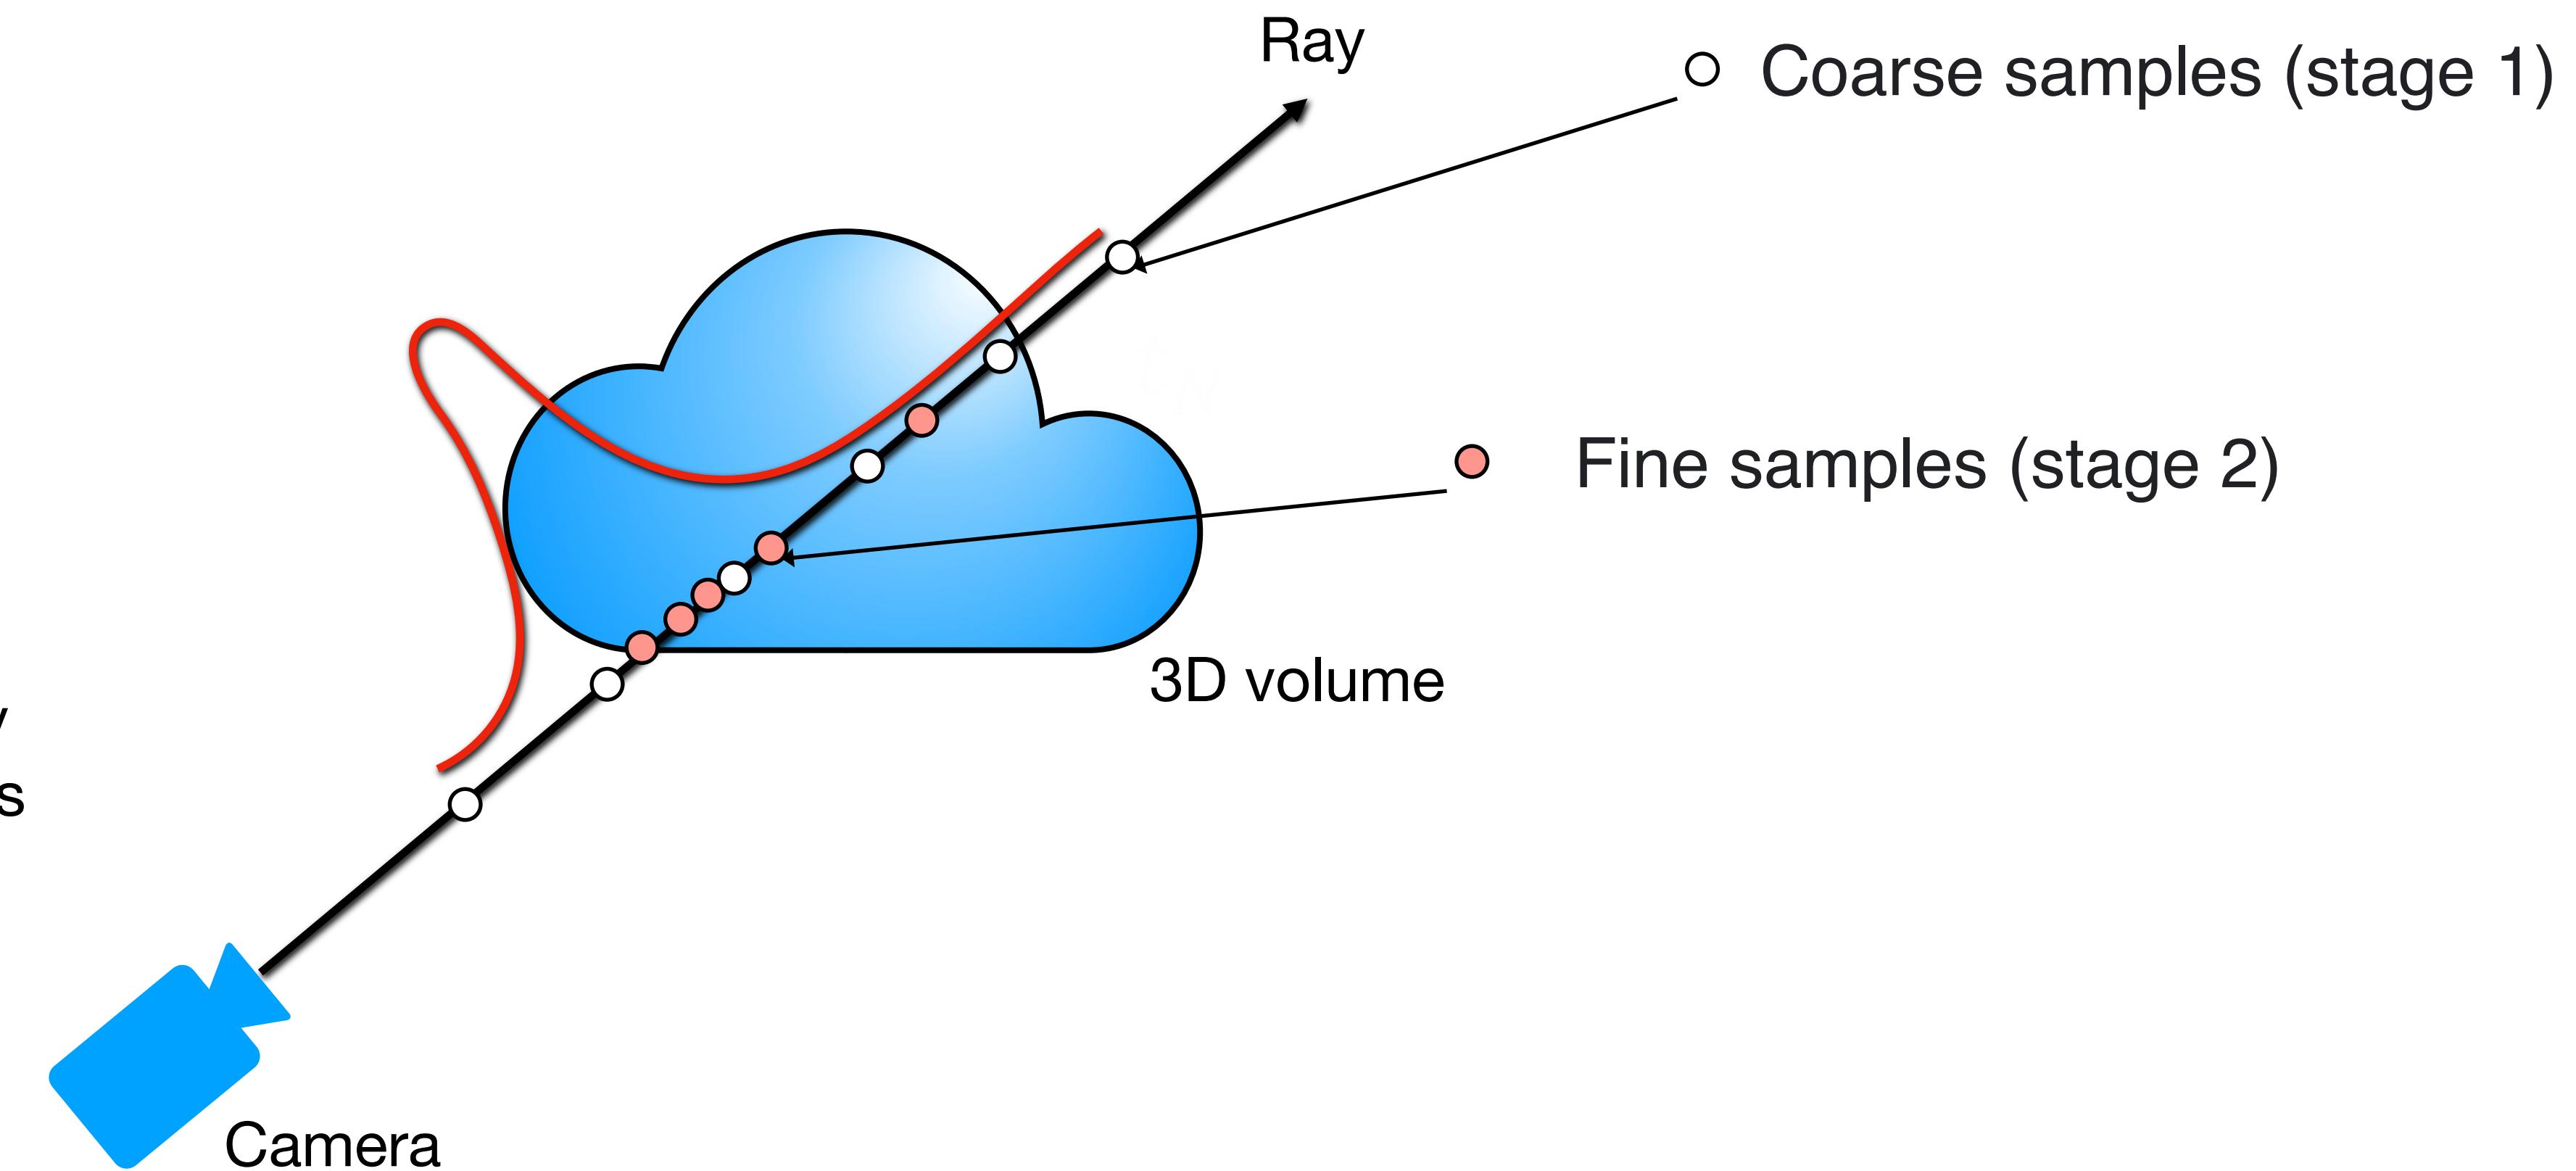
treat weights as probability distribution for new samples



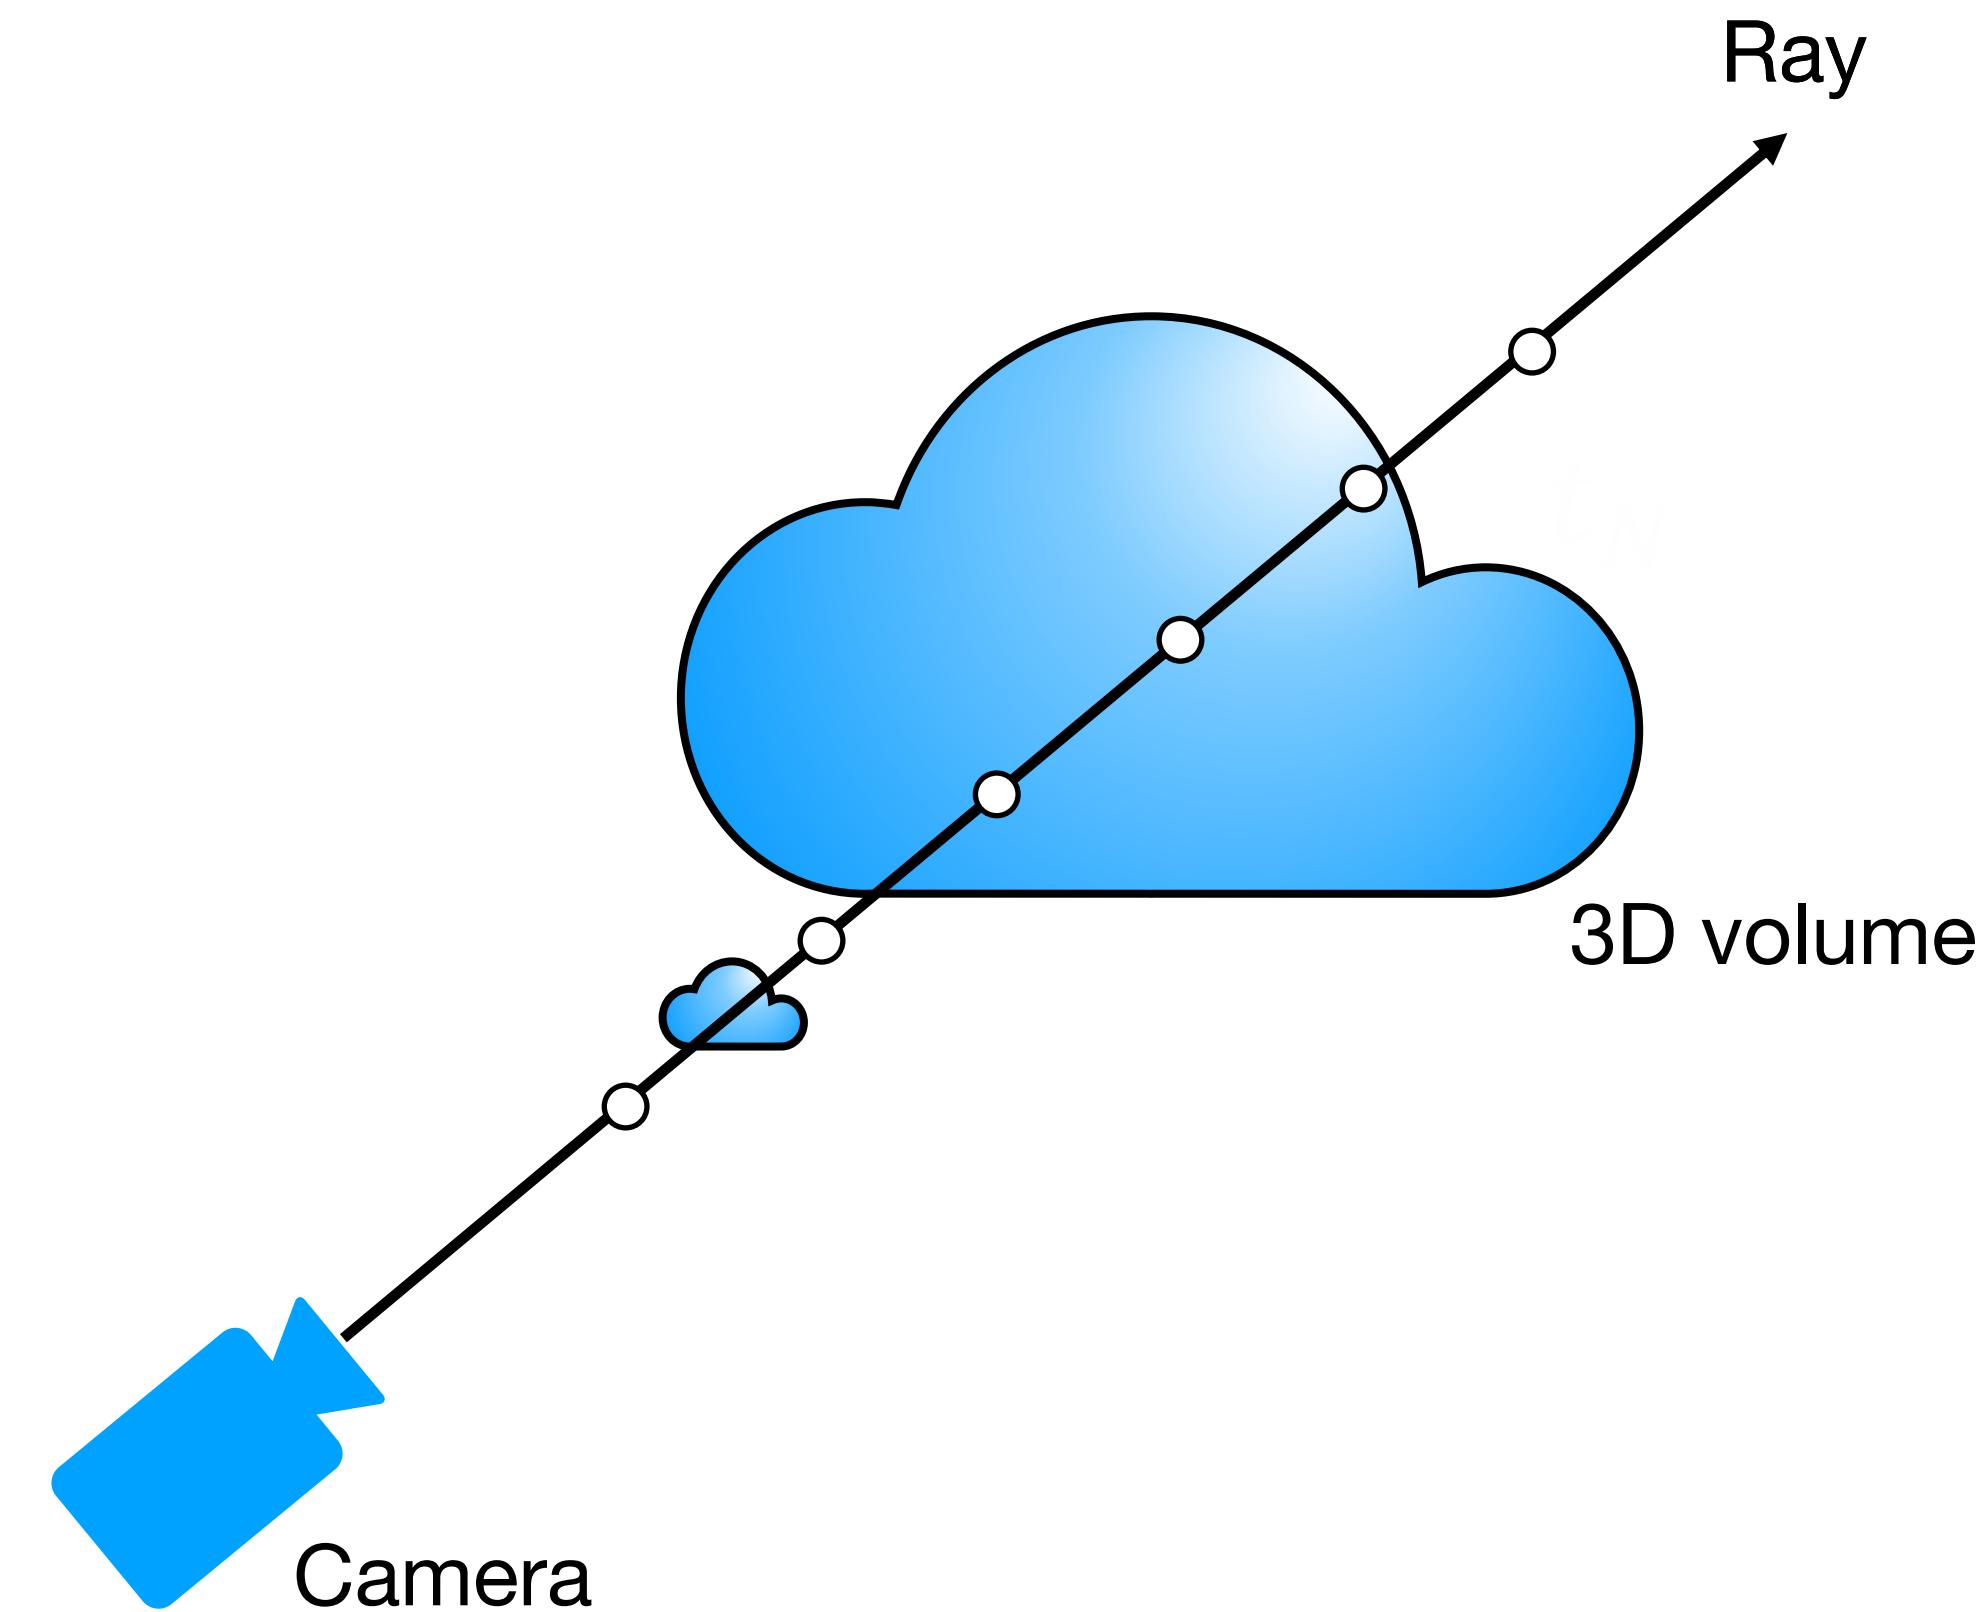
# Key Idea: sample points proportionally to expected effect on final rendering

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

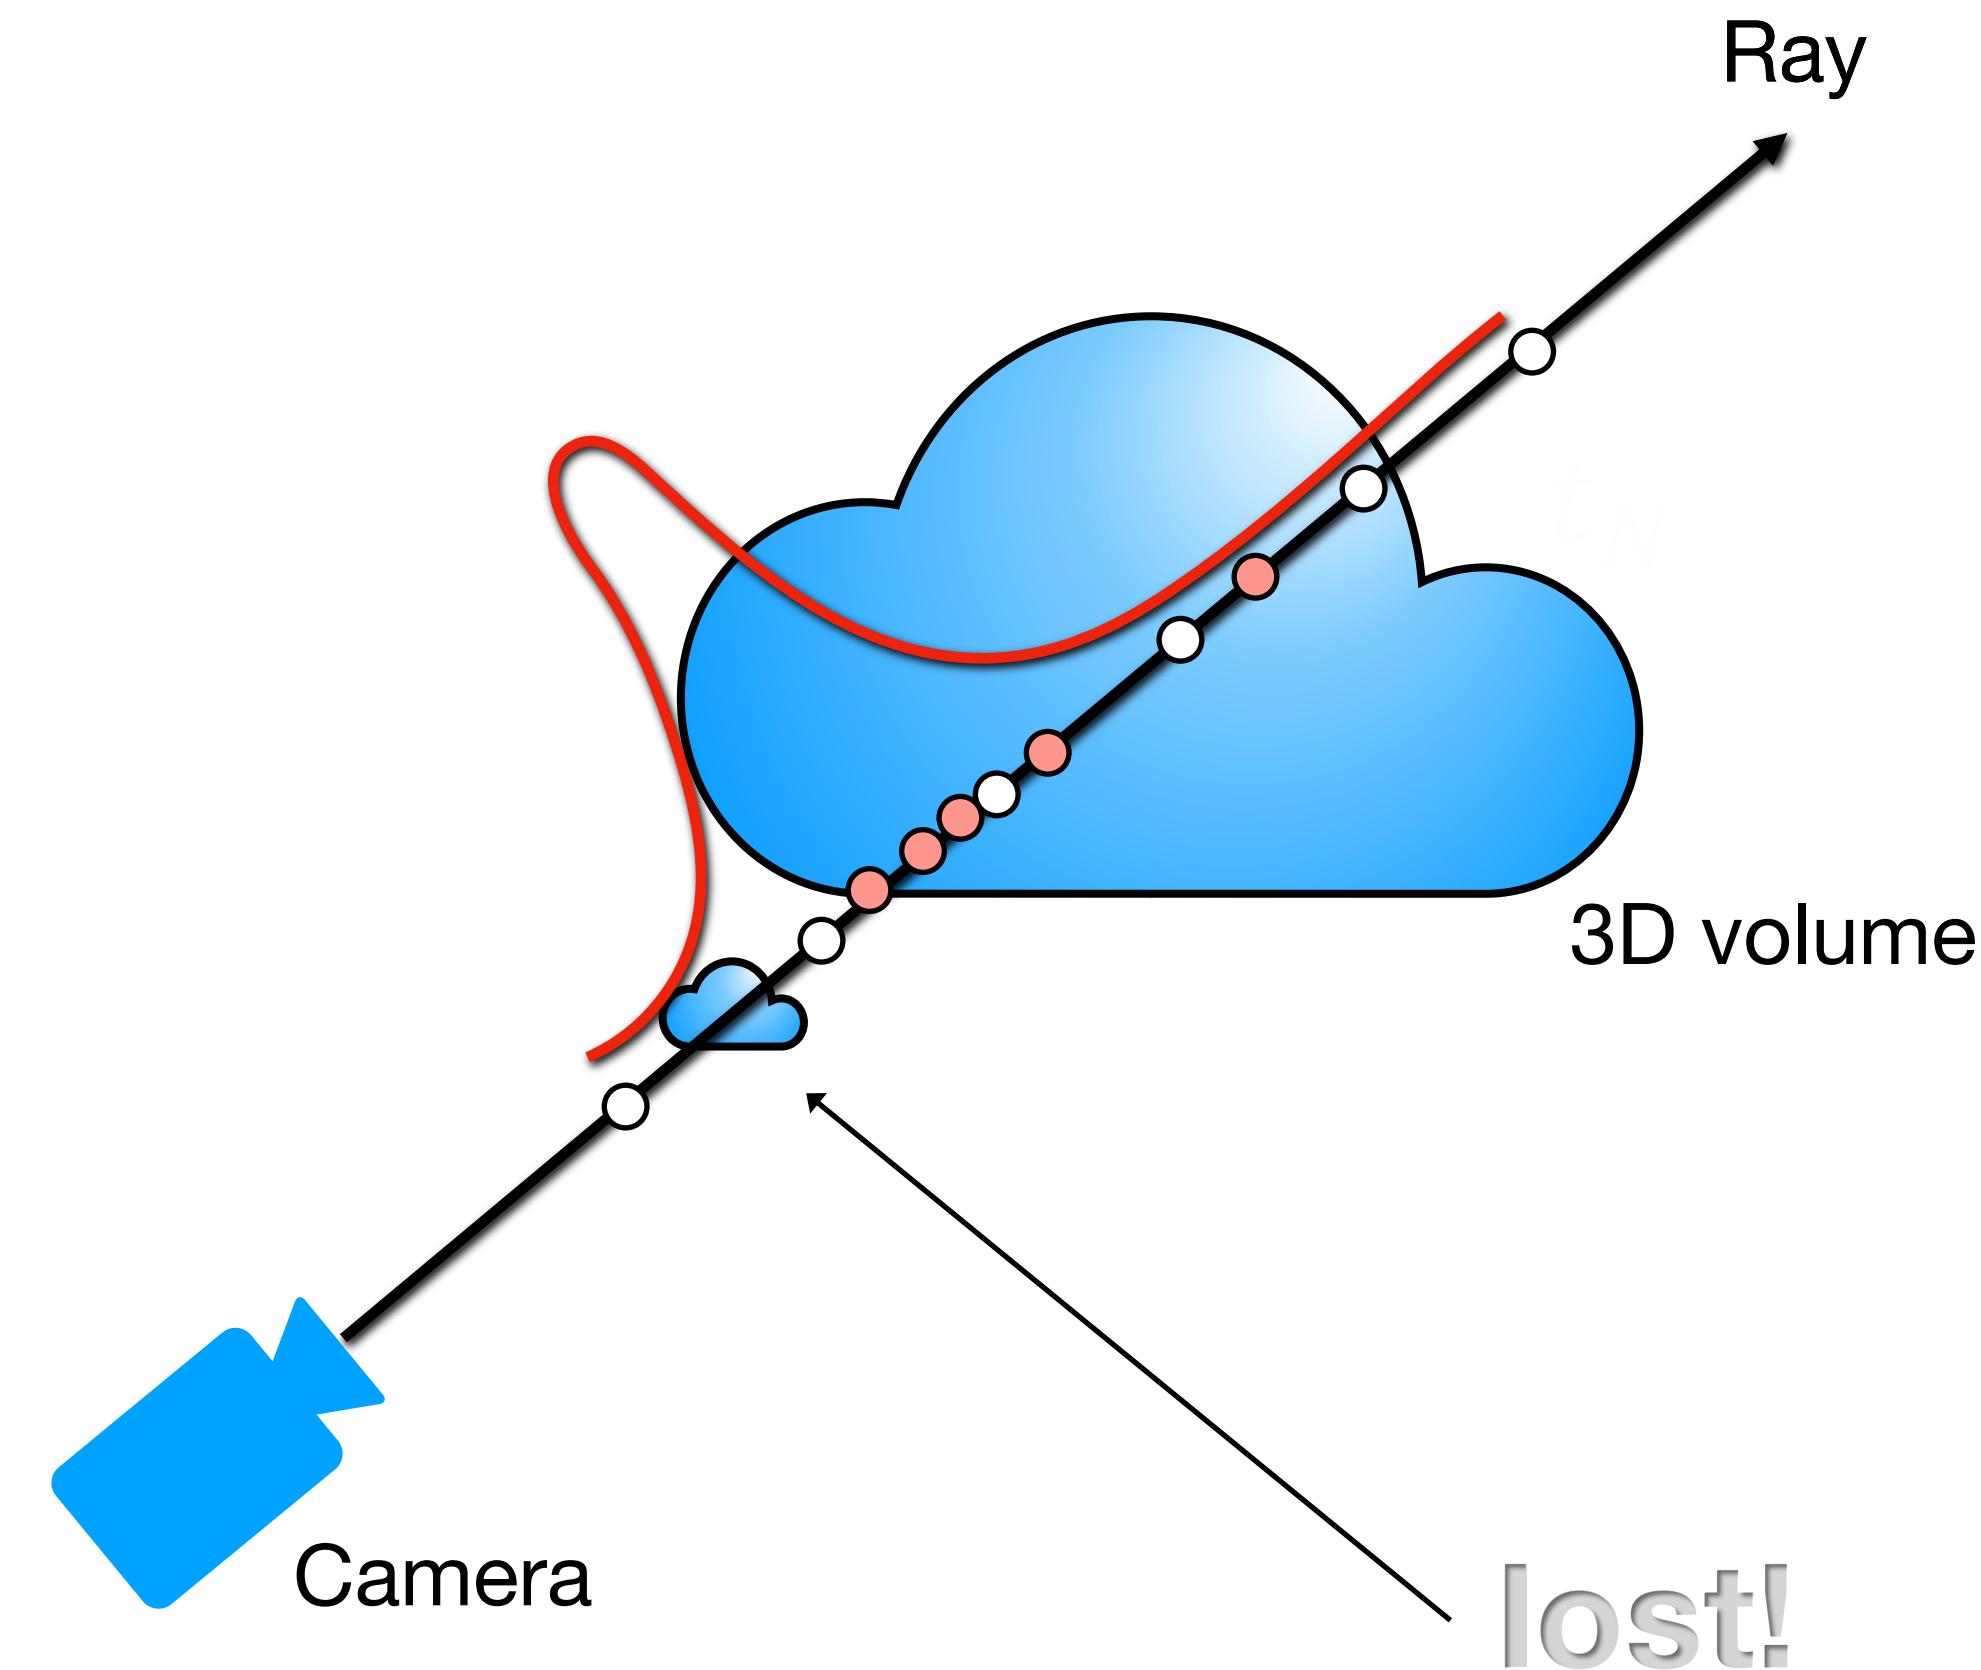
treat weights as probability distribution for new samples



# What about aliasing during coarse sampling?

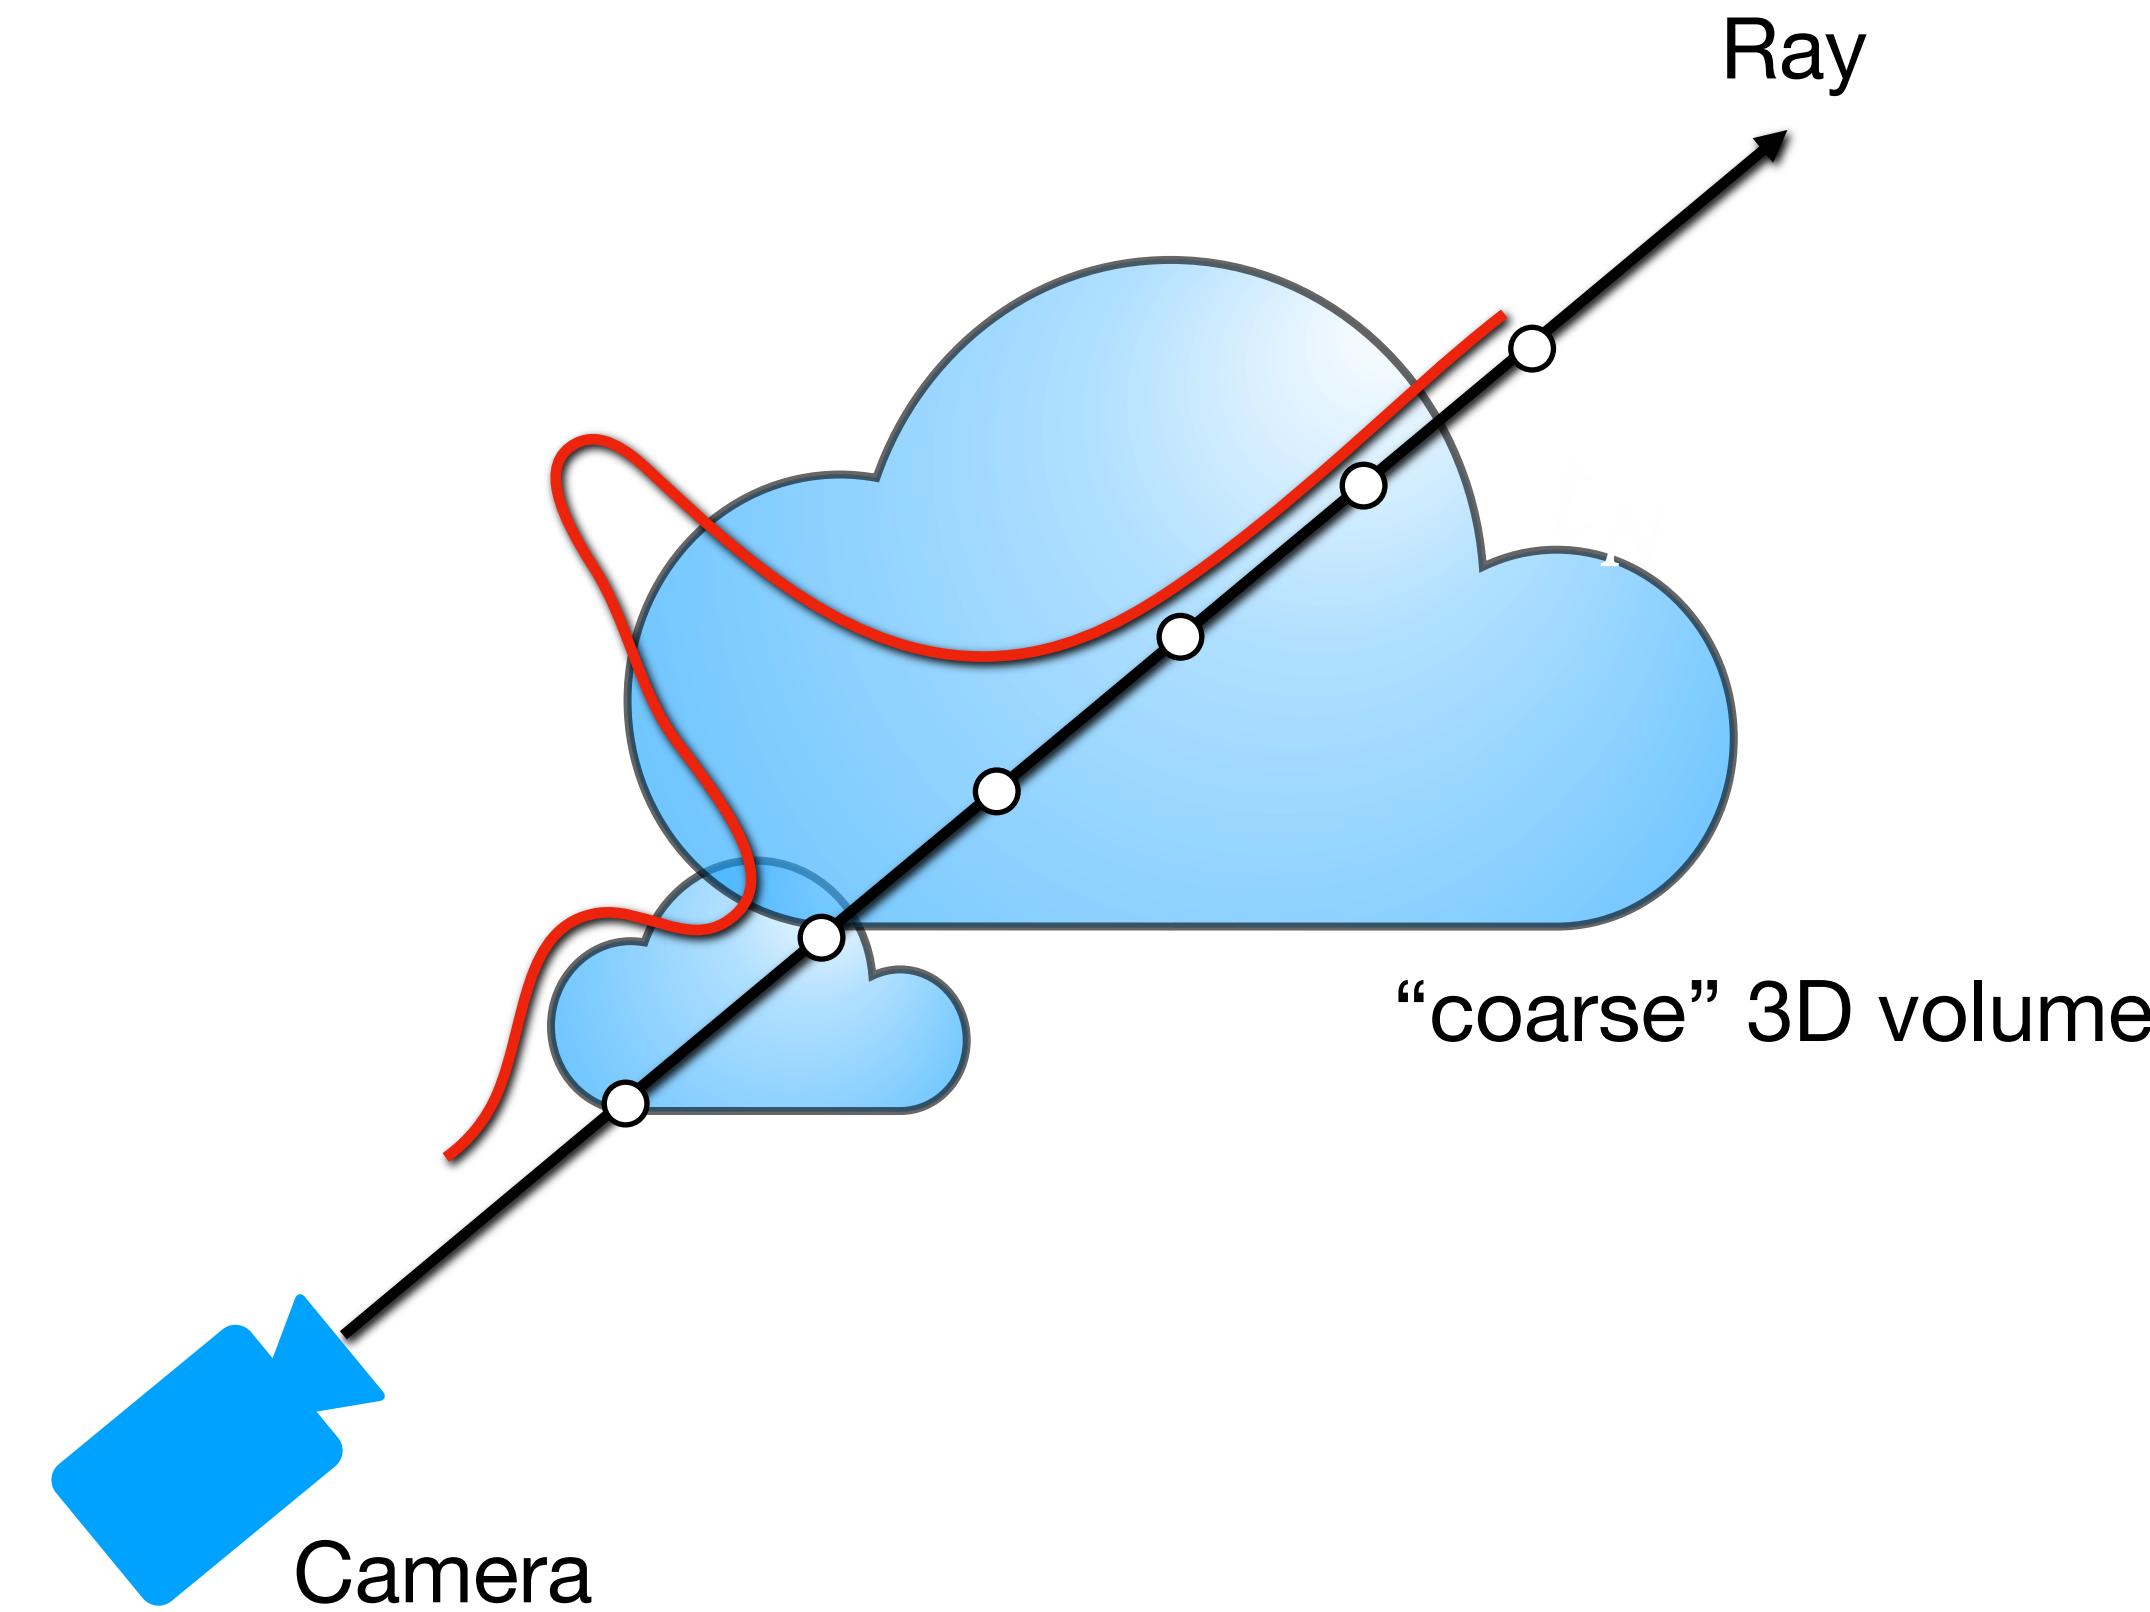


# What about aliasing during coarse sampling?



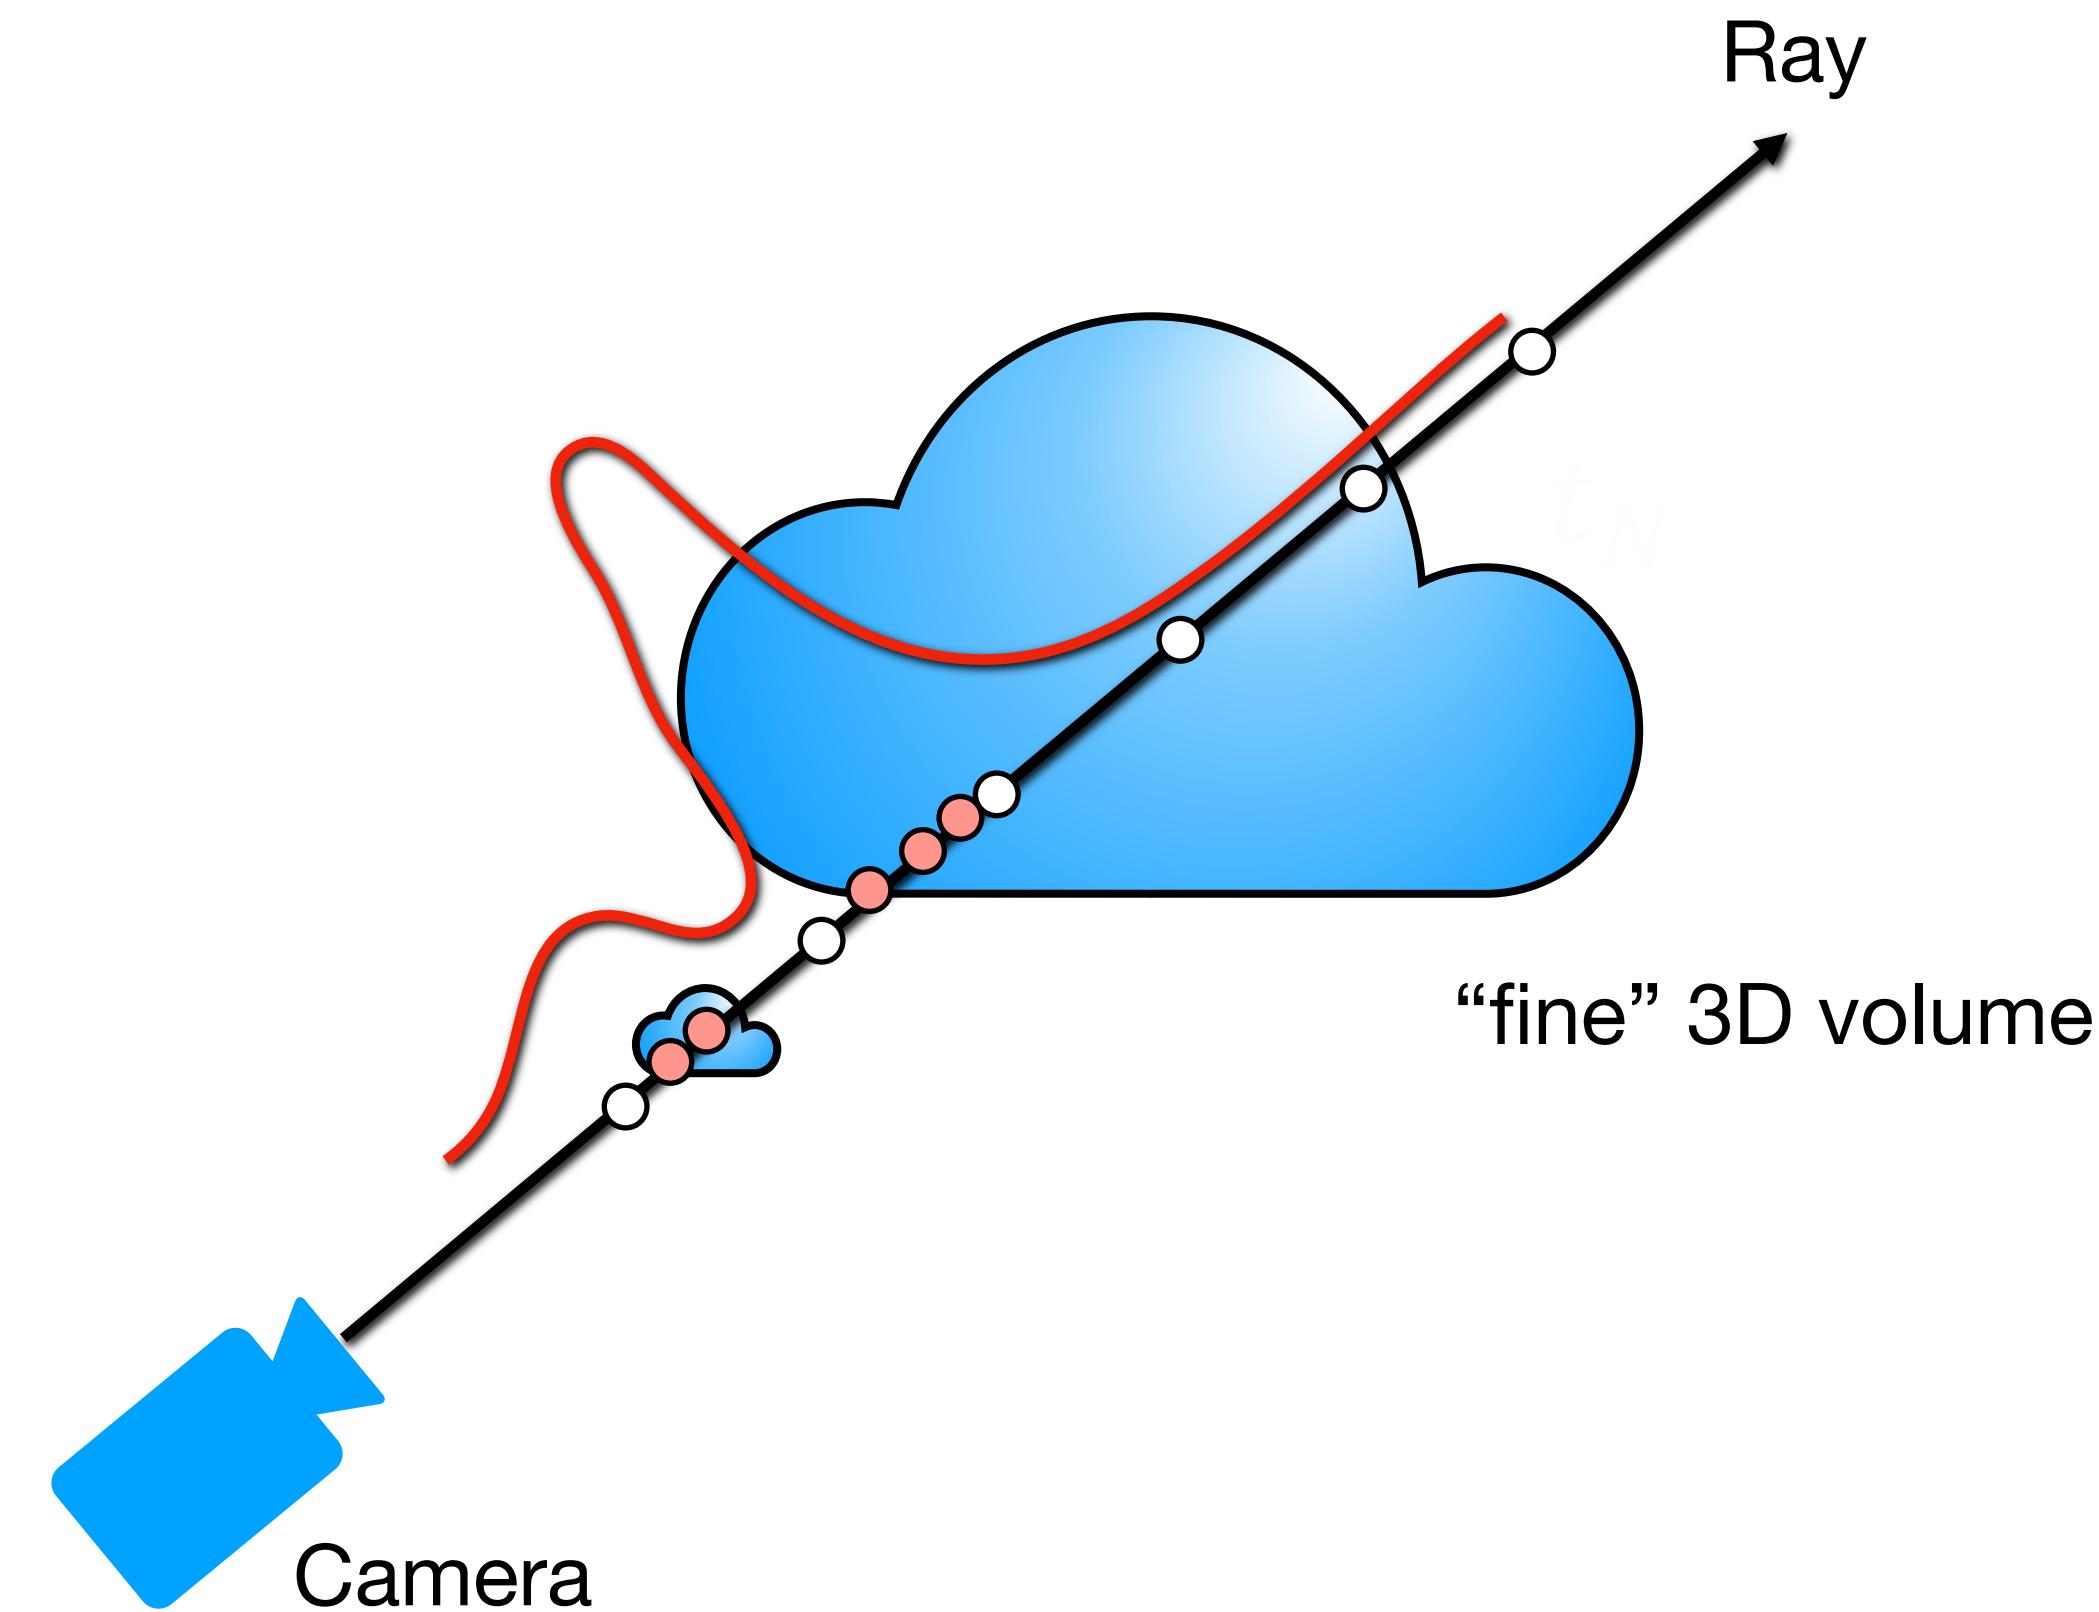
# What about aliasing during coarse sampling?

Solution: train two NeRFs! → lower resolution for first “coarse” level



# What about aliasing during coarse sampling?

Solution: train two NeRFs! → higher resolution for second “fine” level



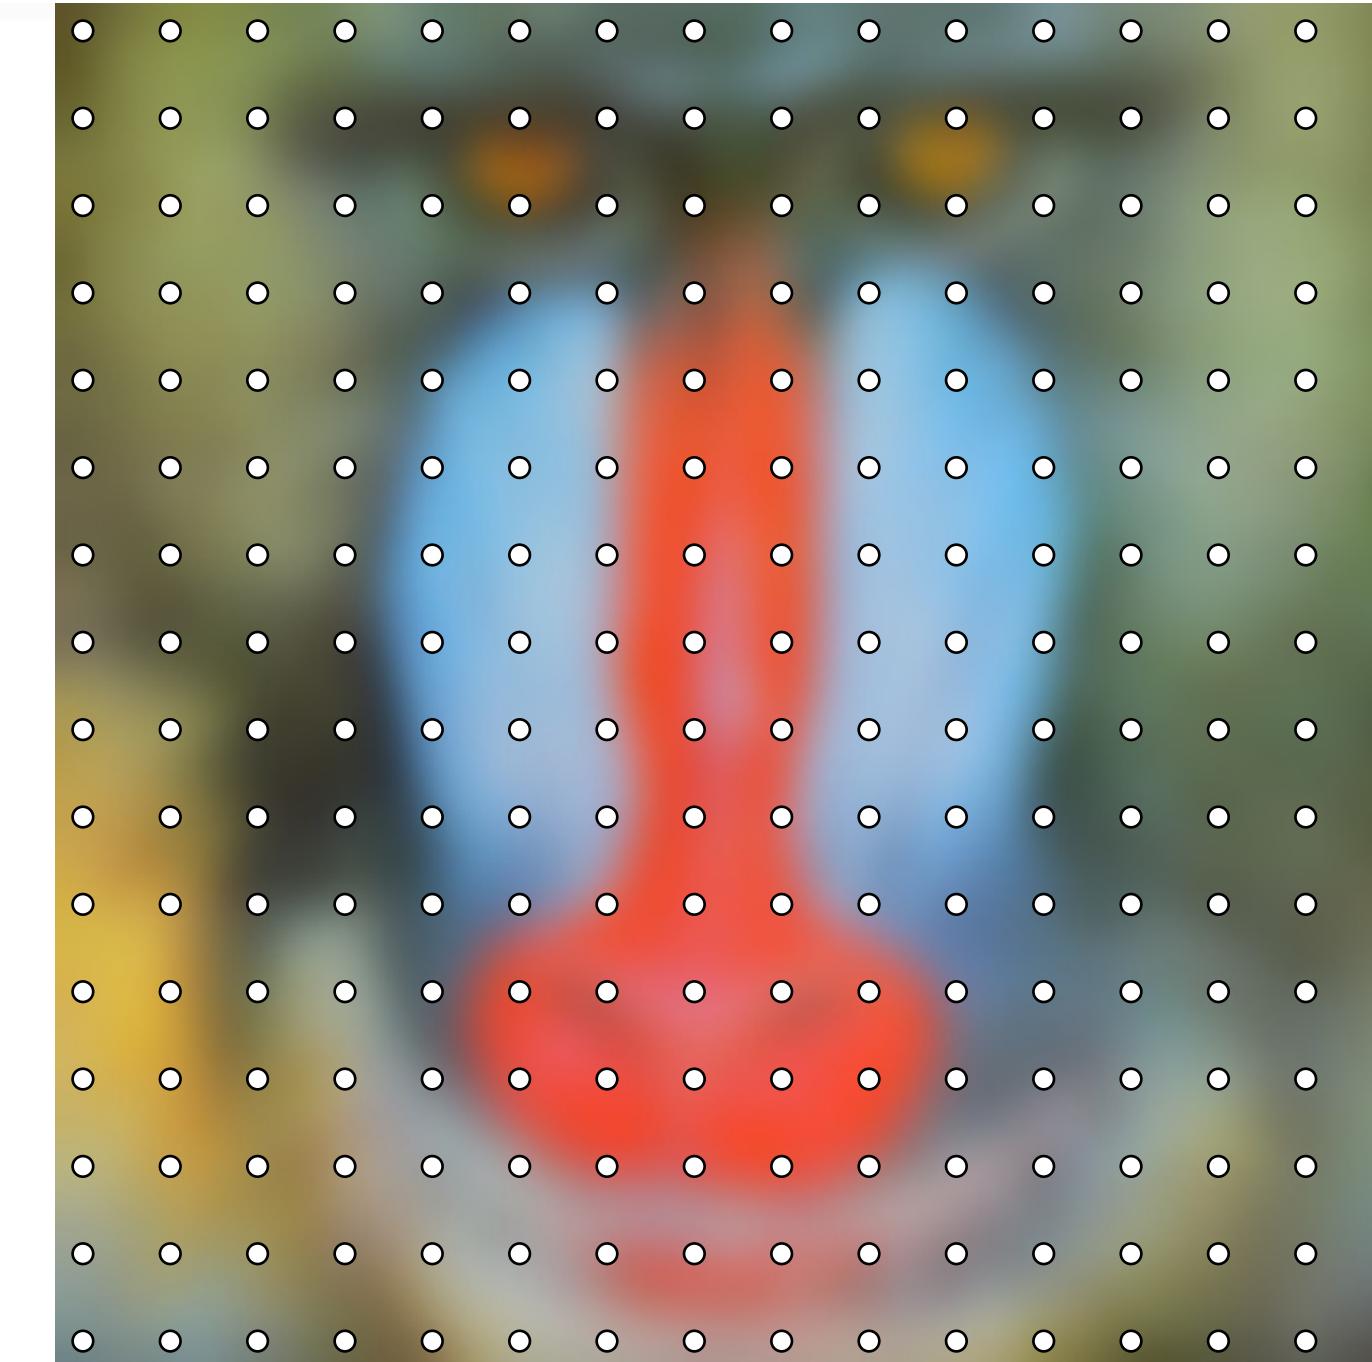
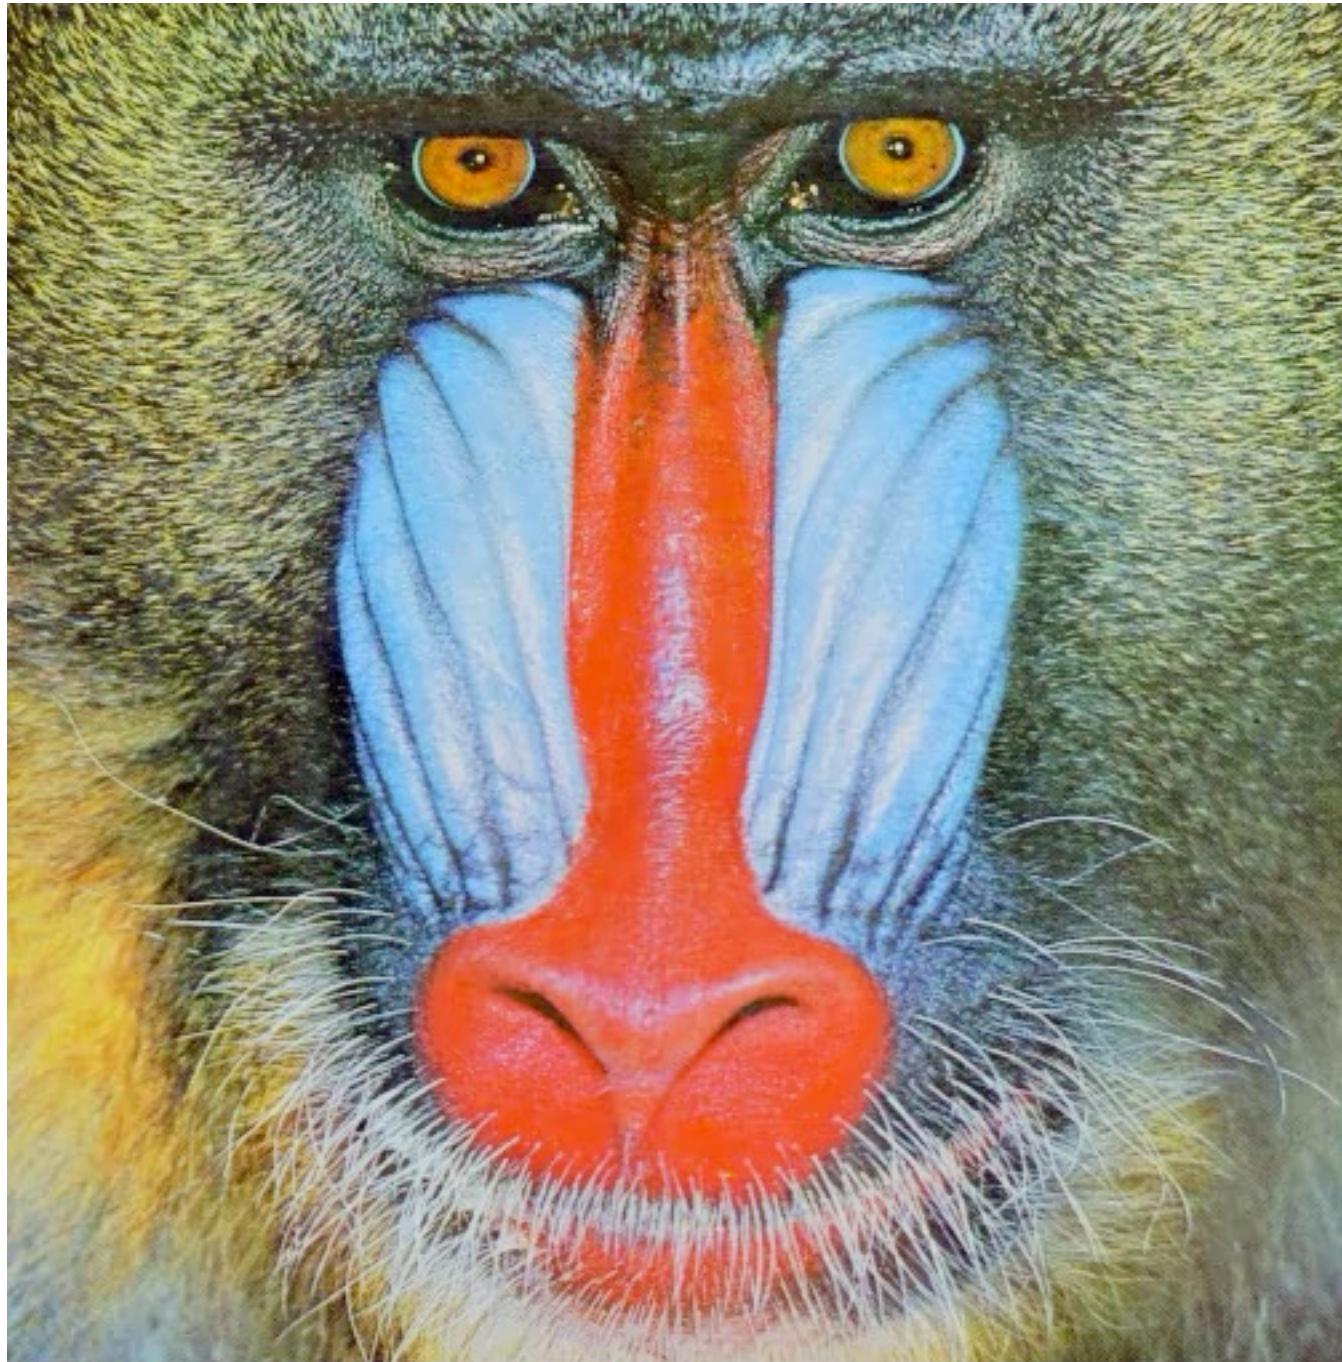
# More anti-aliasing

can we avoid training two networks?

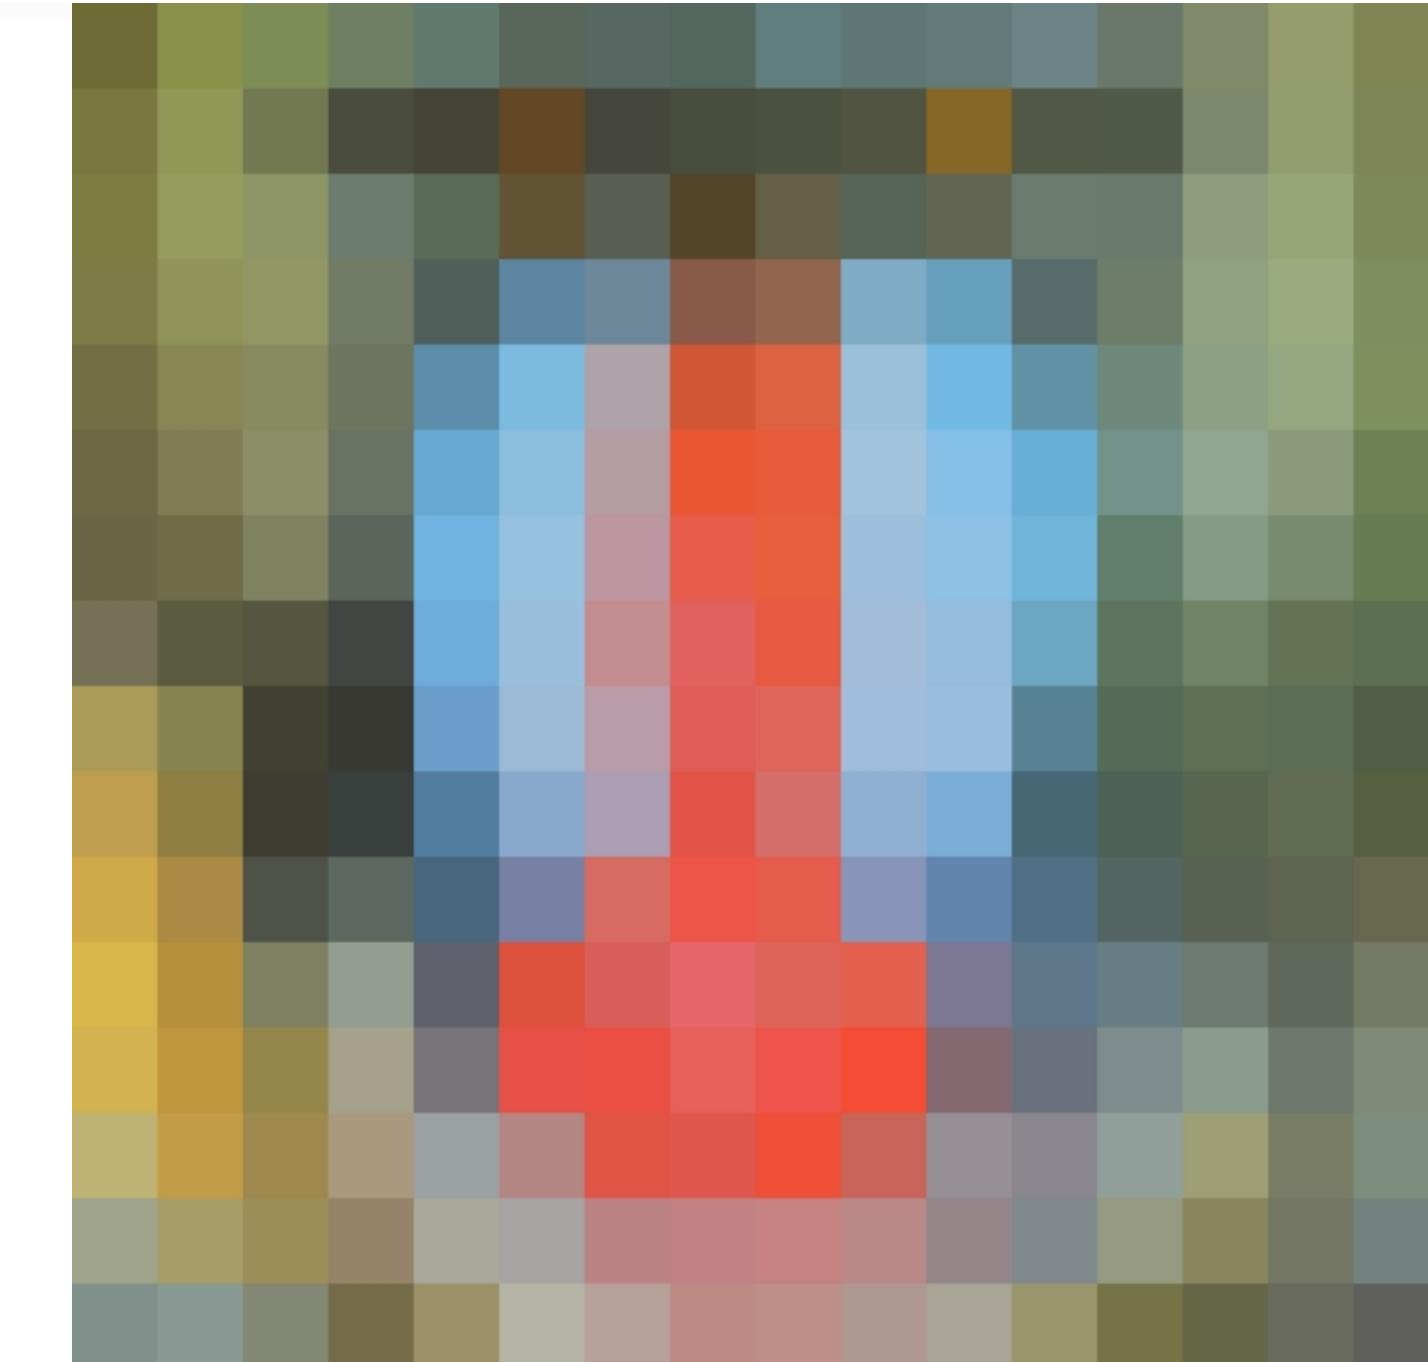
# Aliasing in NeRF renderings



# Recall that averaging reduces aliasing

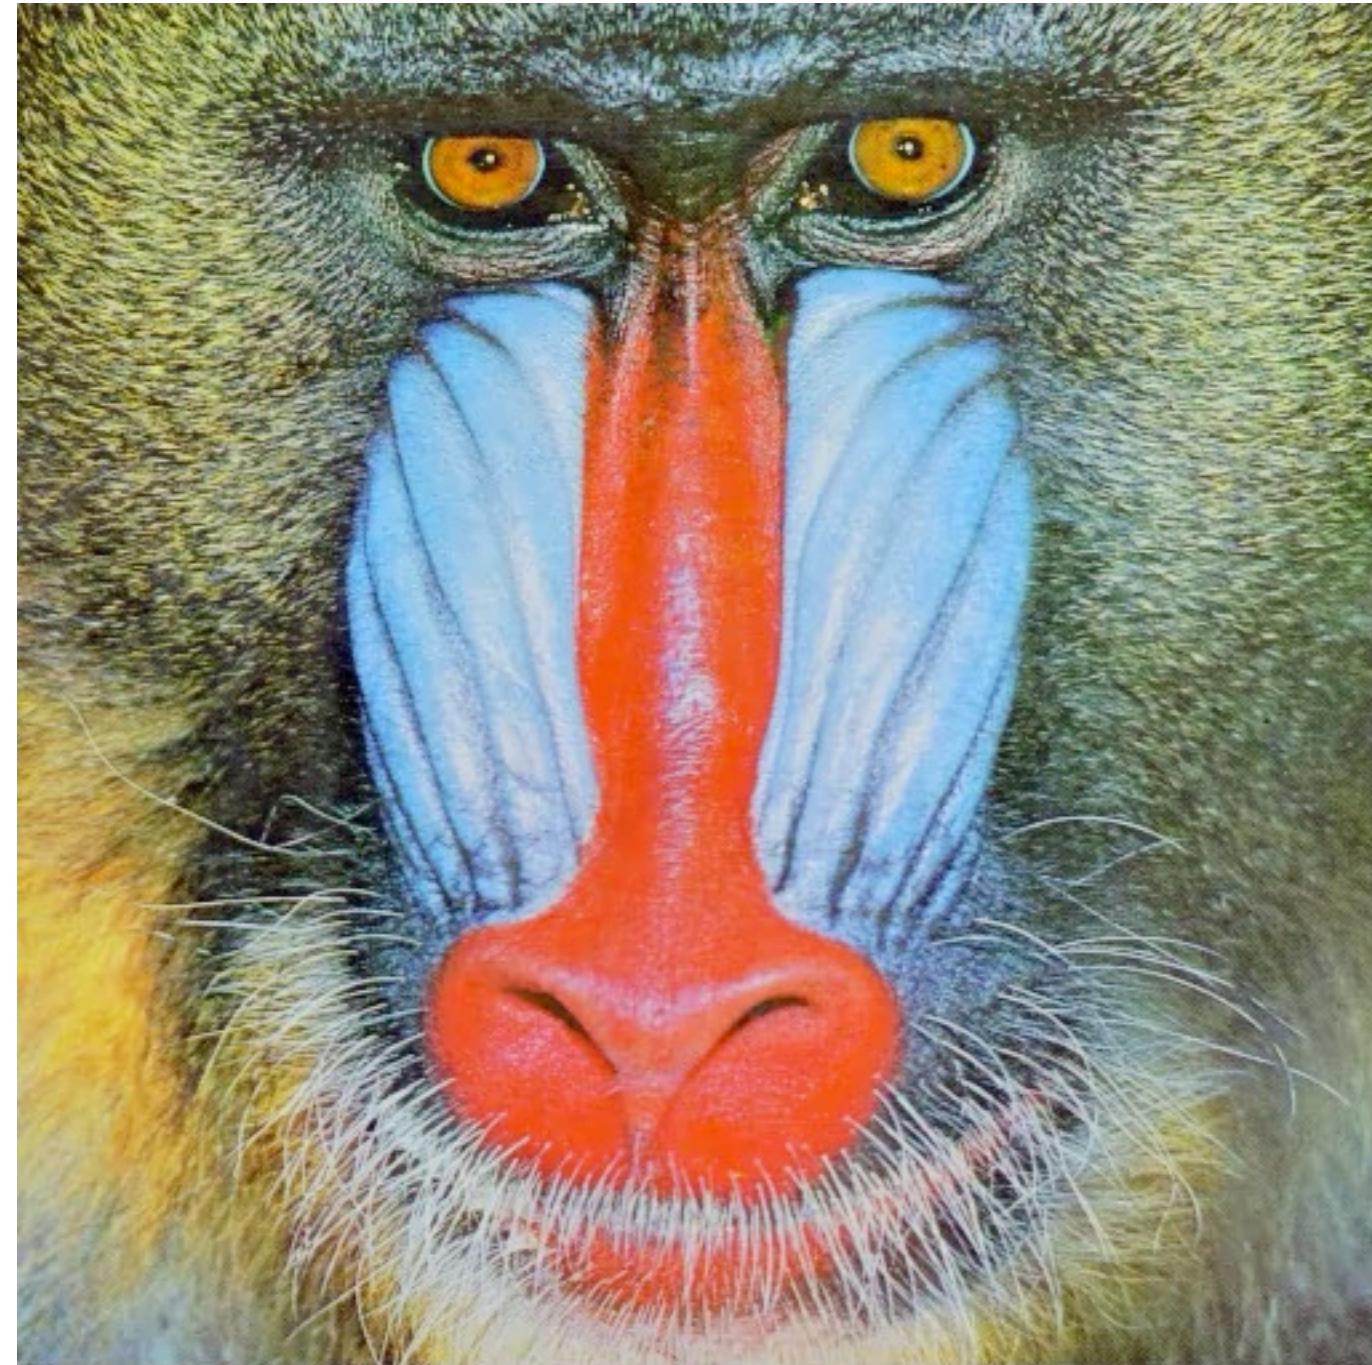


Filter



Sample

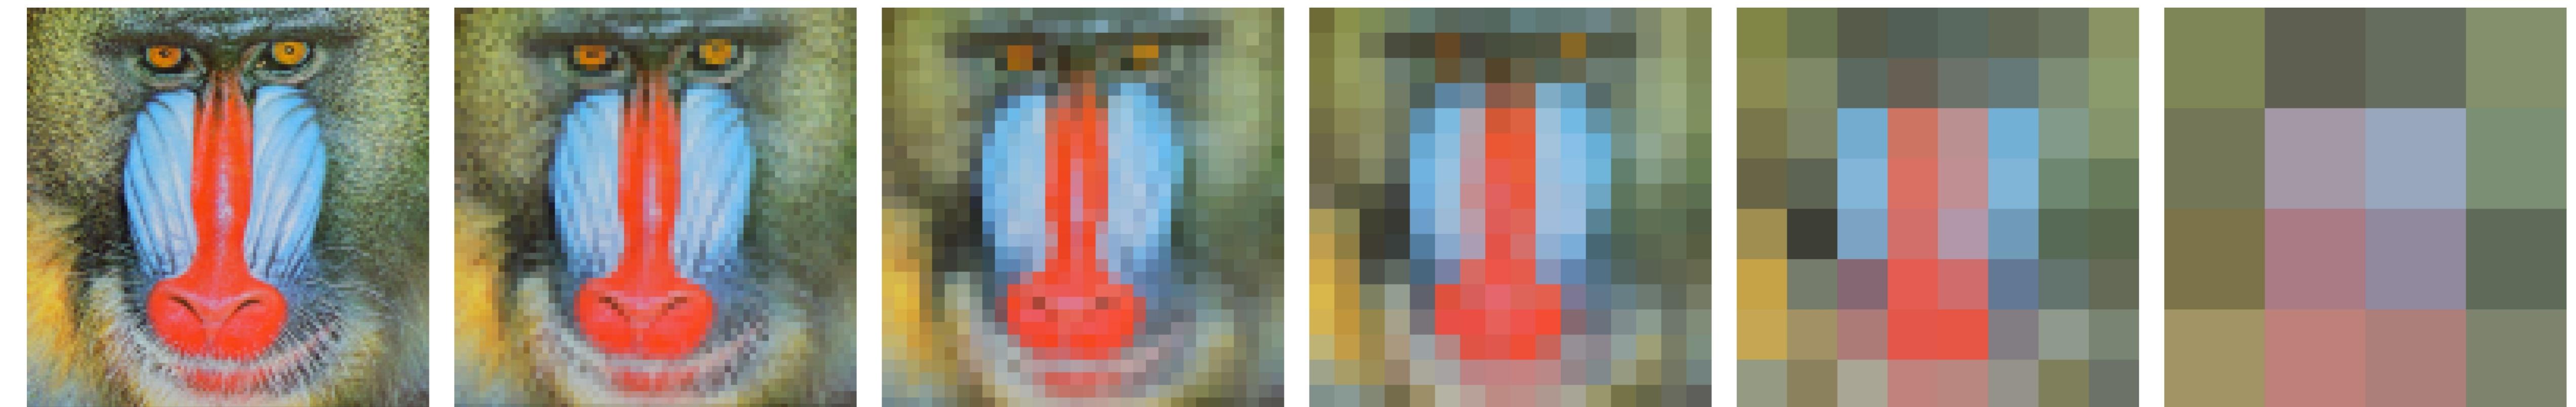
# But repeatedly sampling and averaging is inefficient



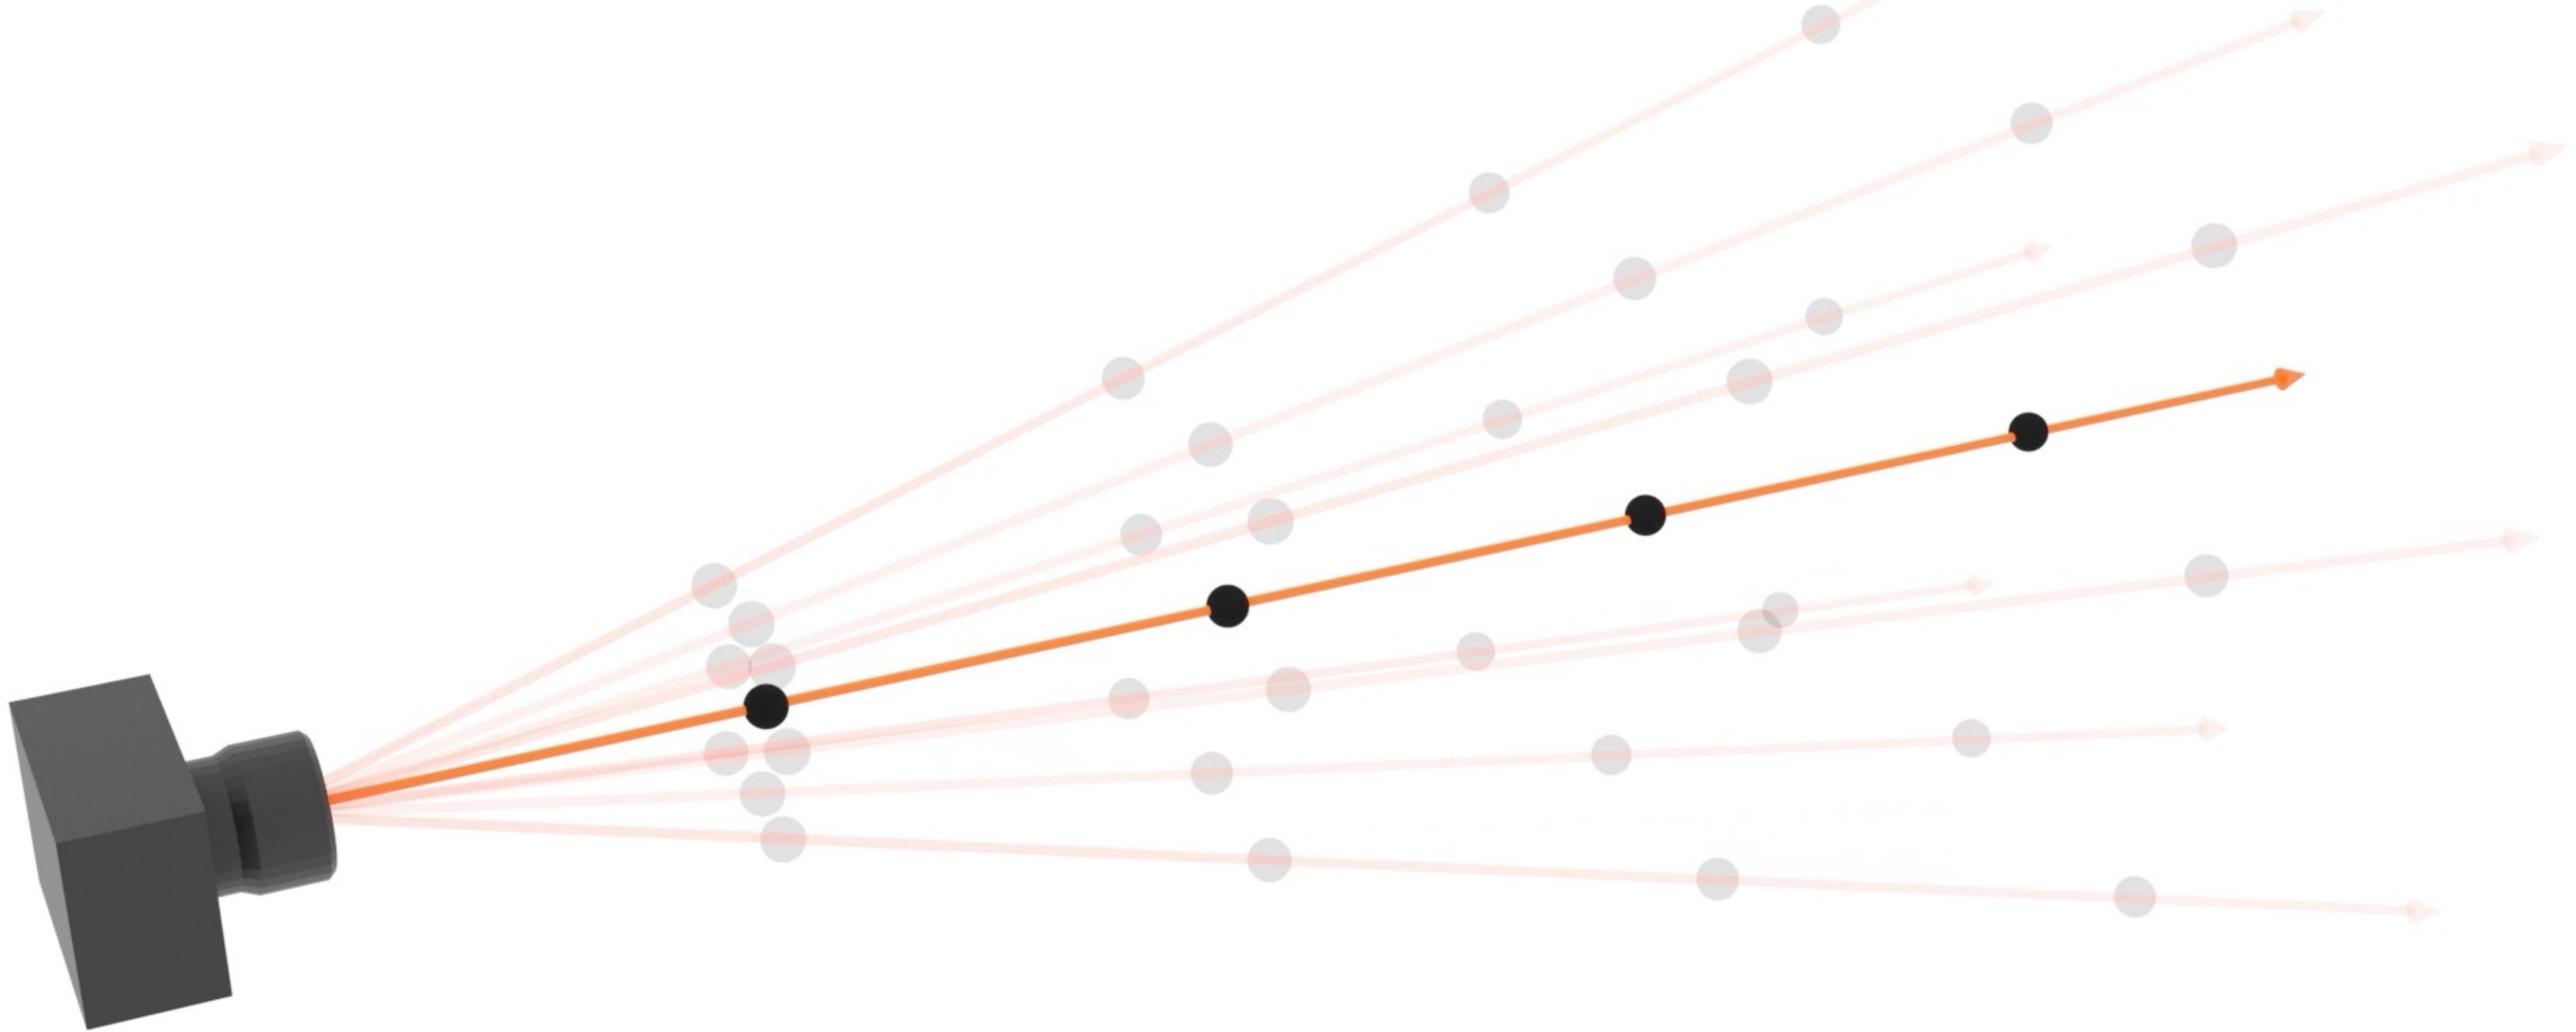
Filter

Sample

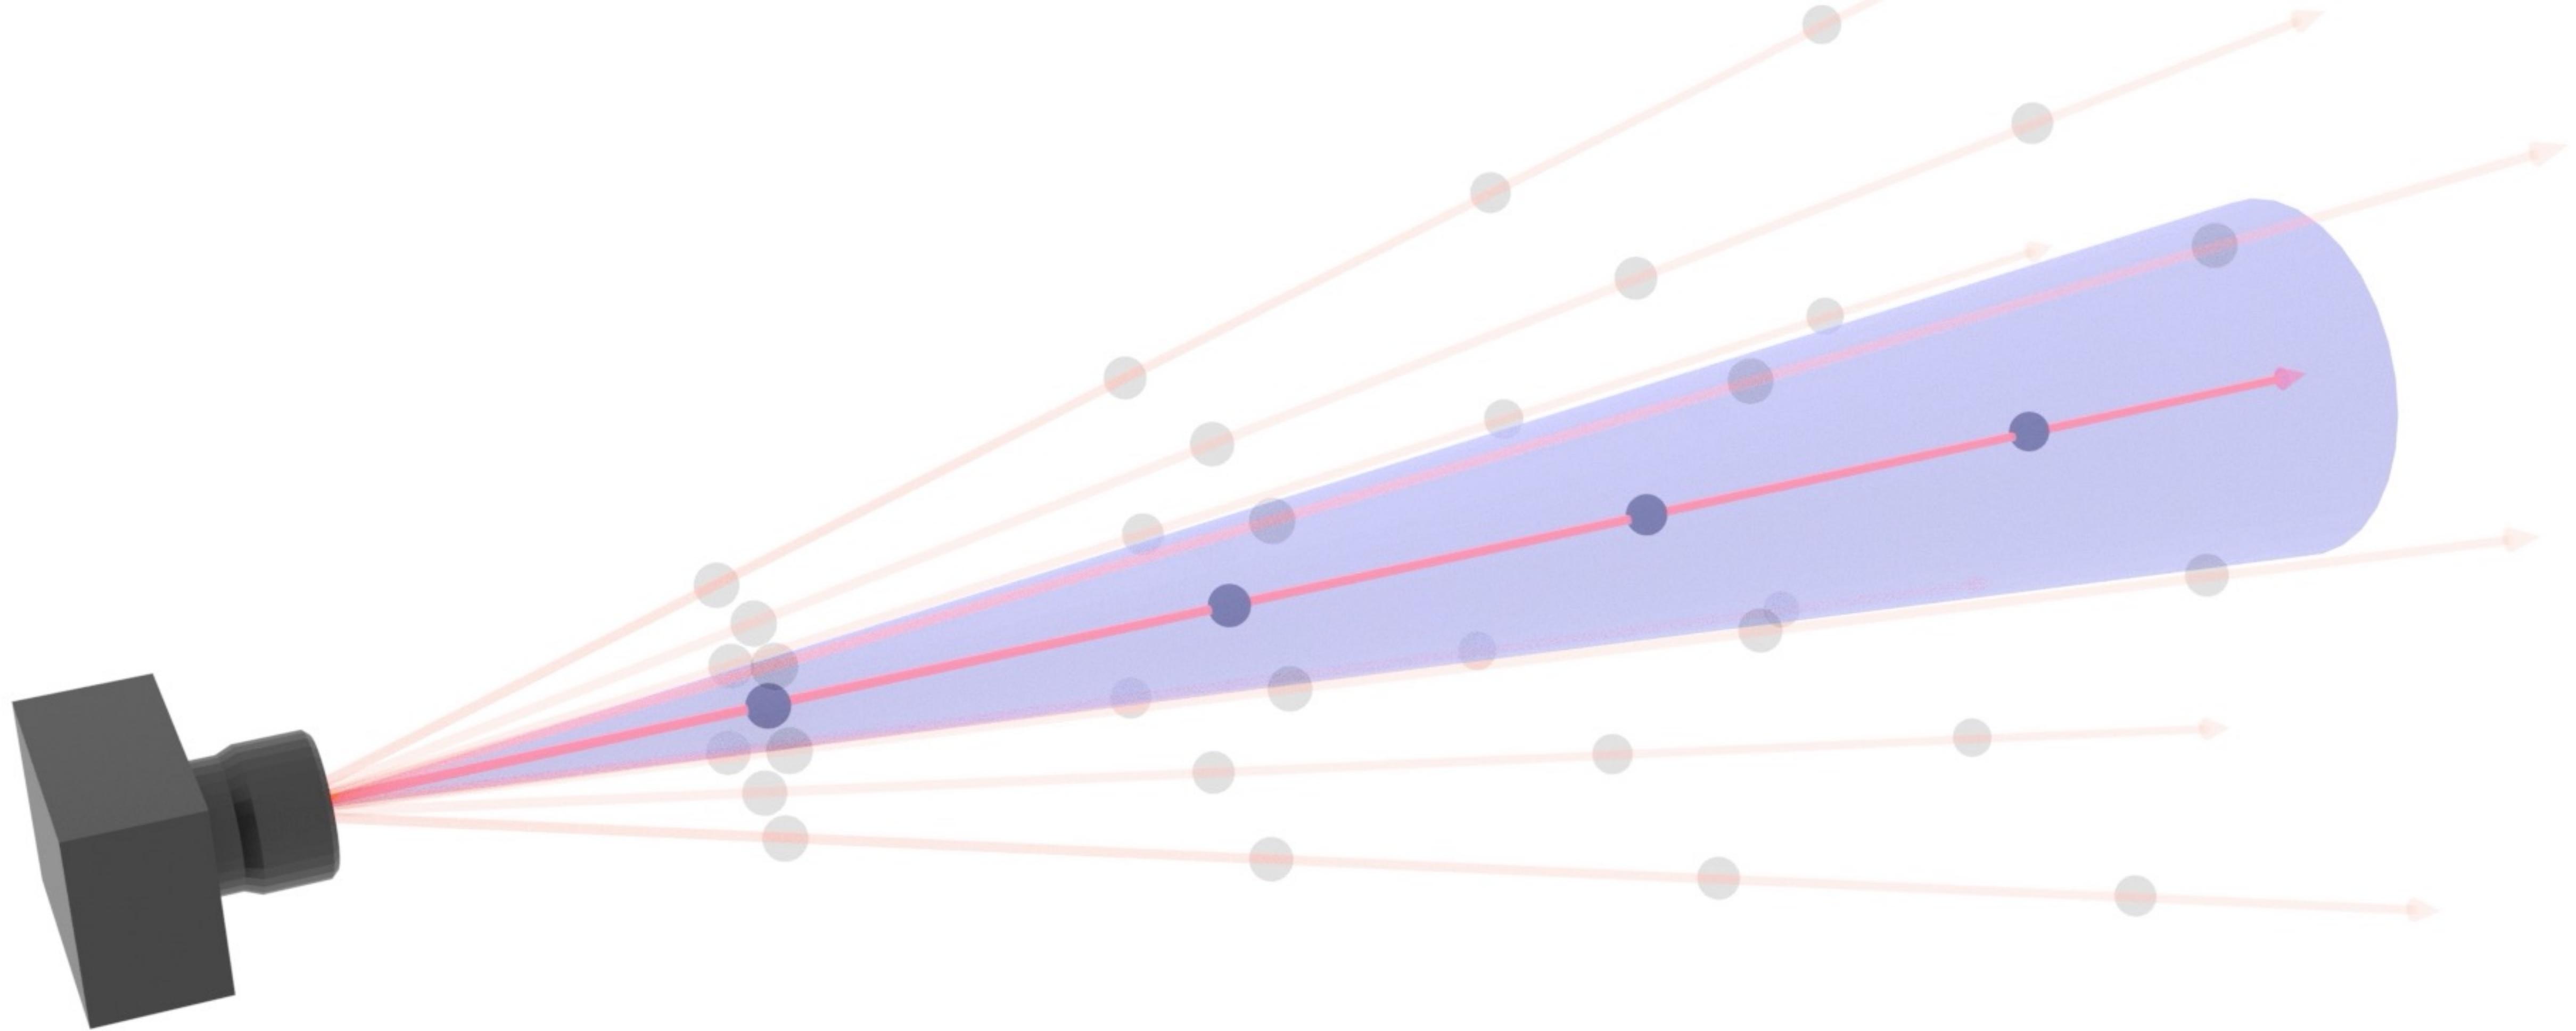
# Standard solution: prefiltering with a mipmap



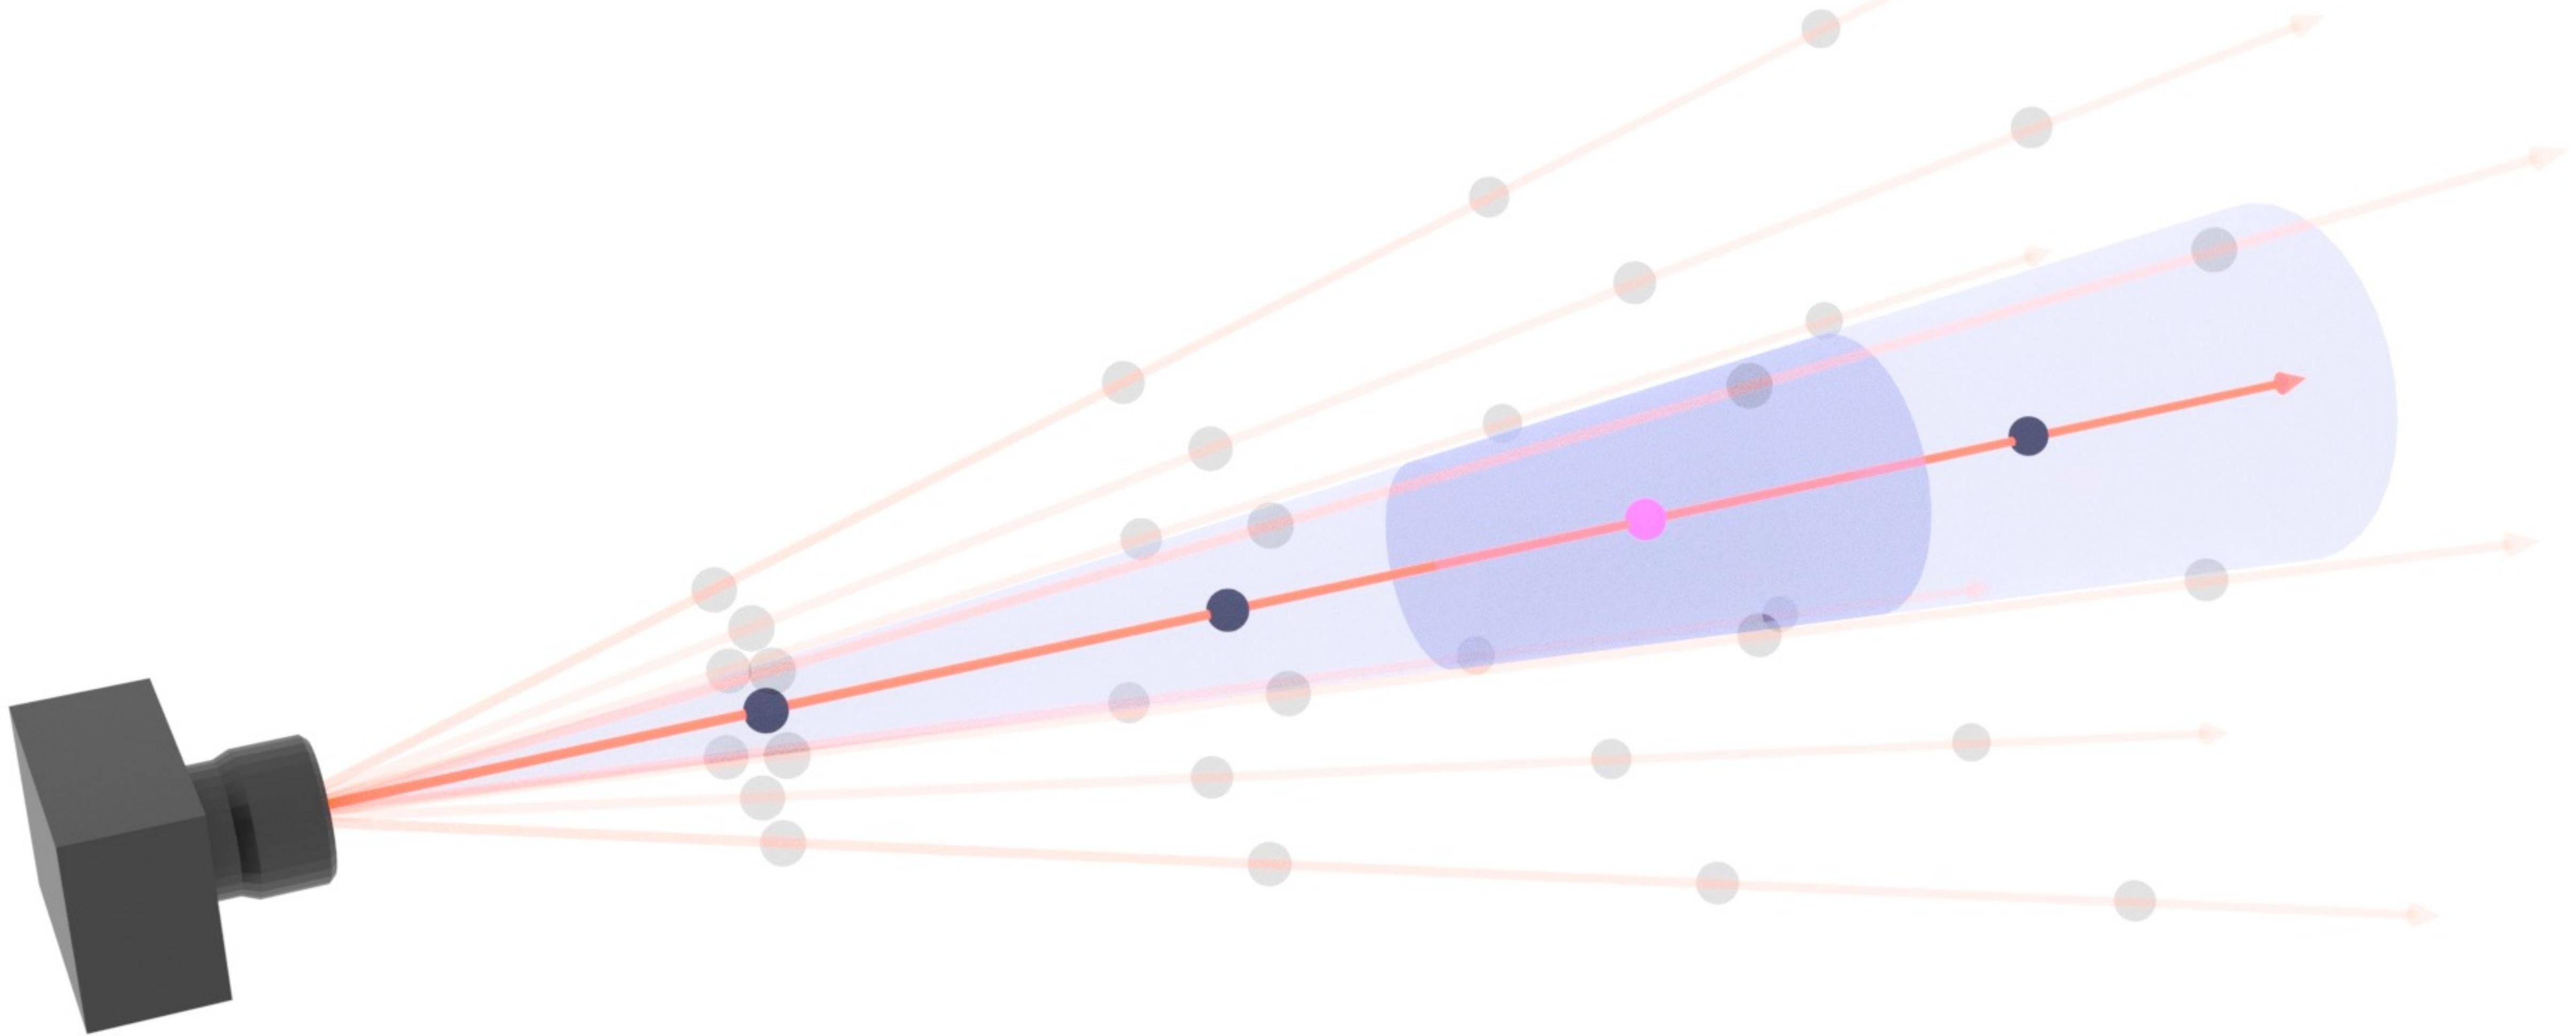
# Antialiasing requires average ray color within pixel



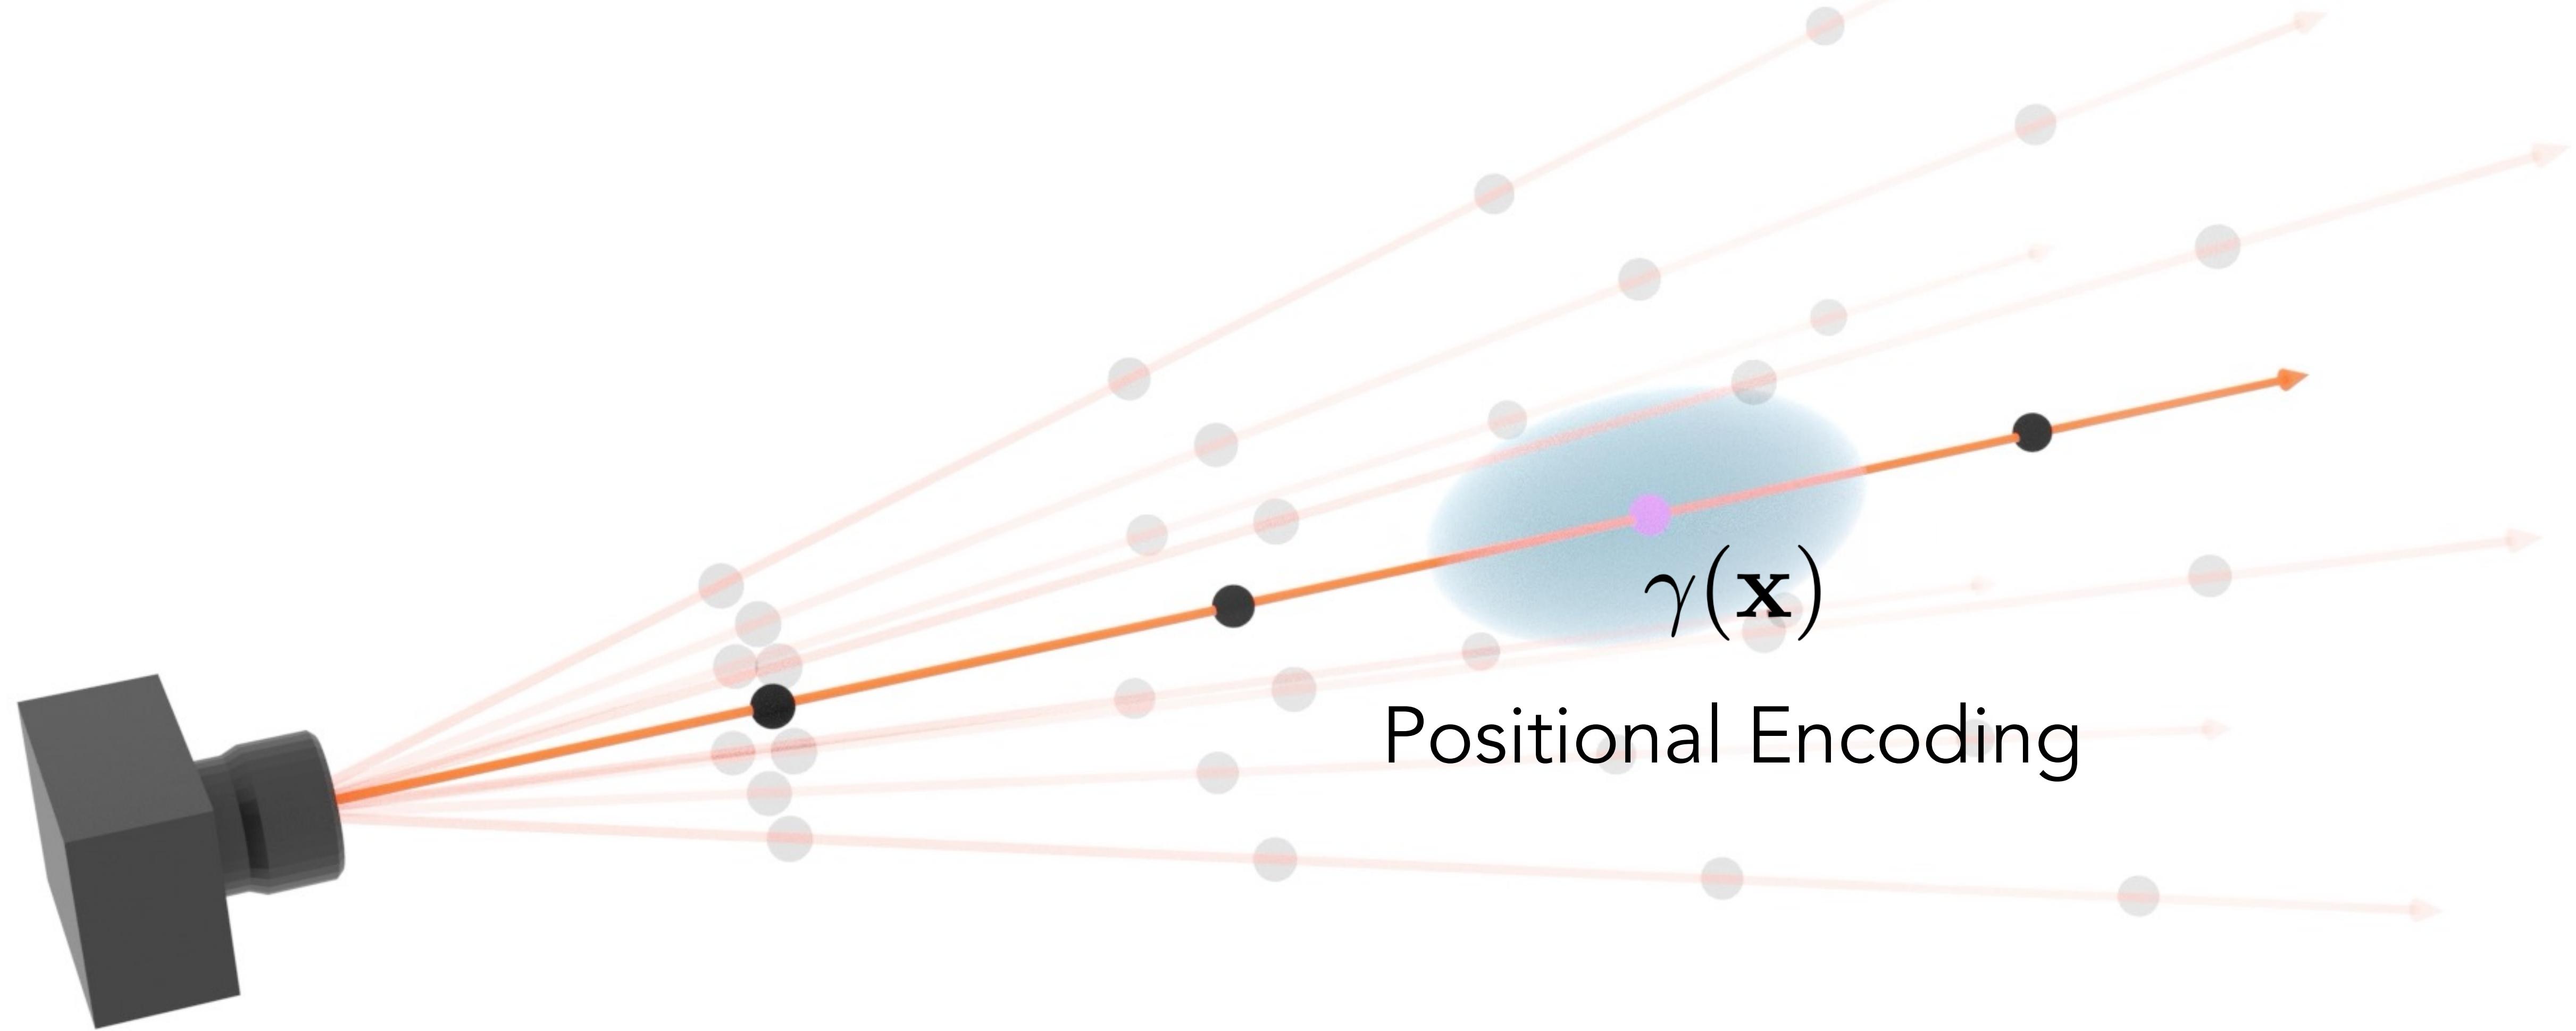
# Supersampling vs. prefiltering



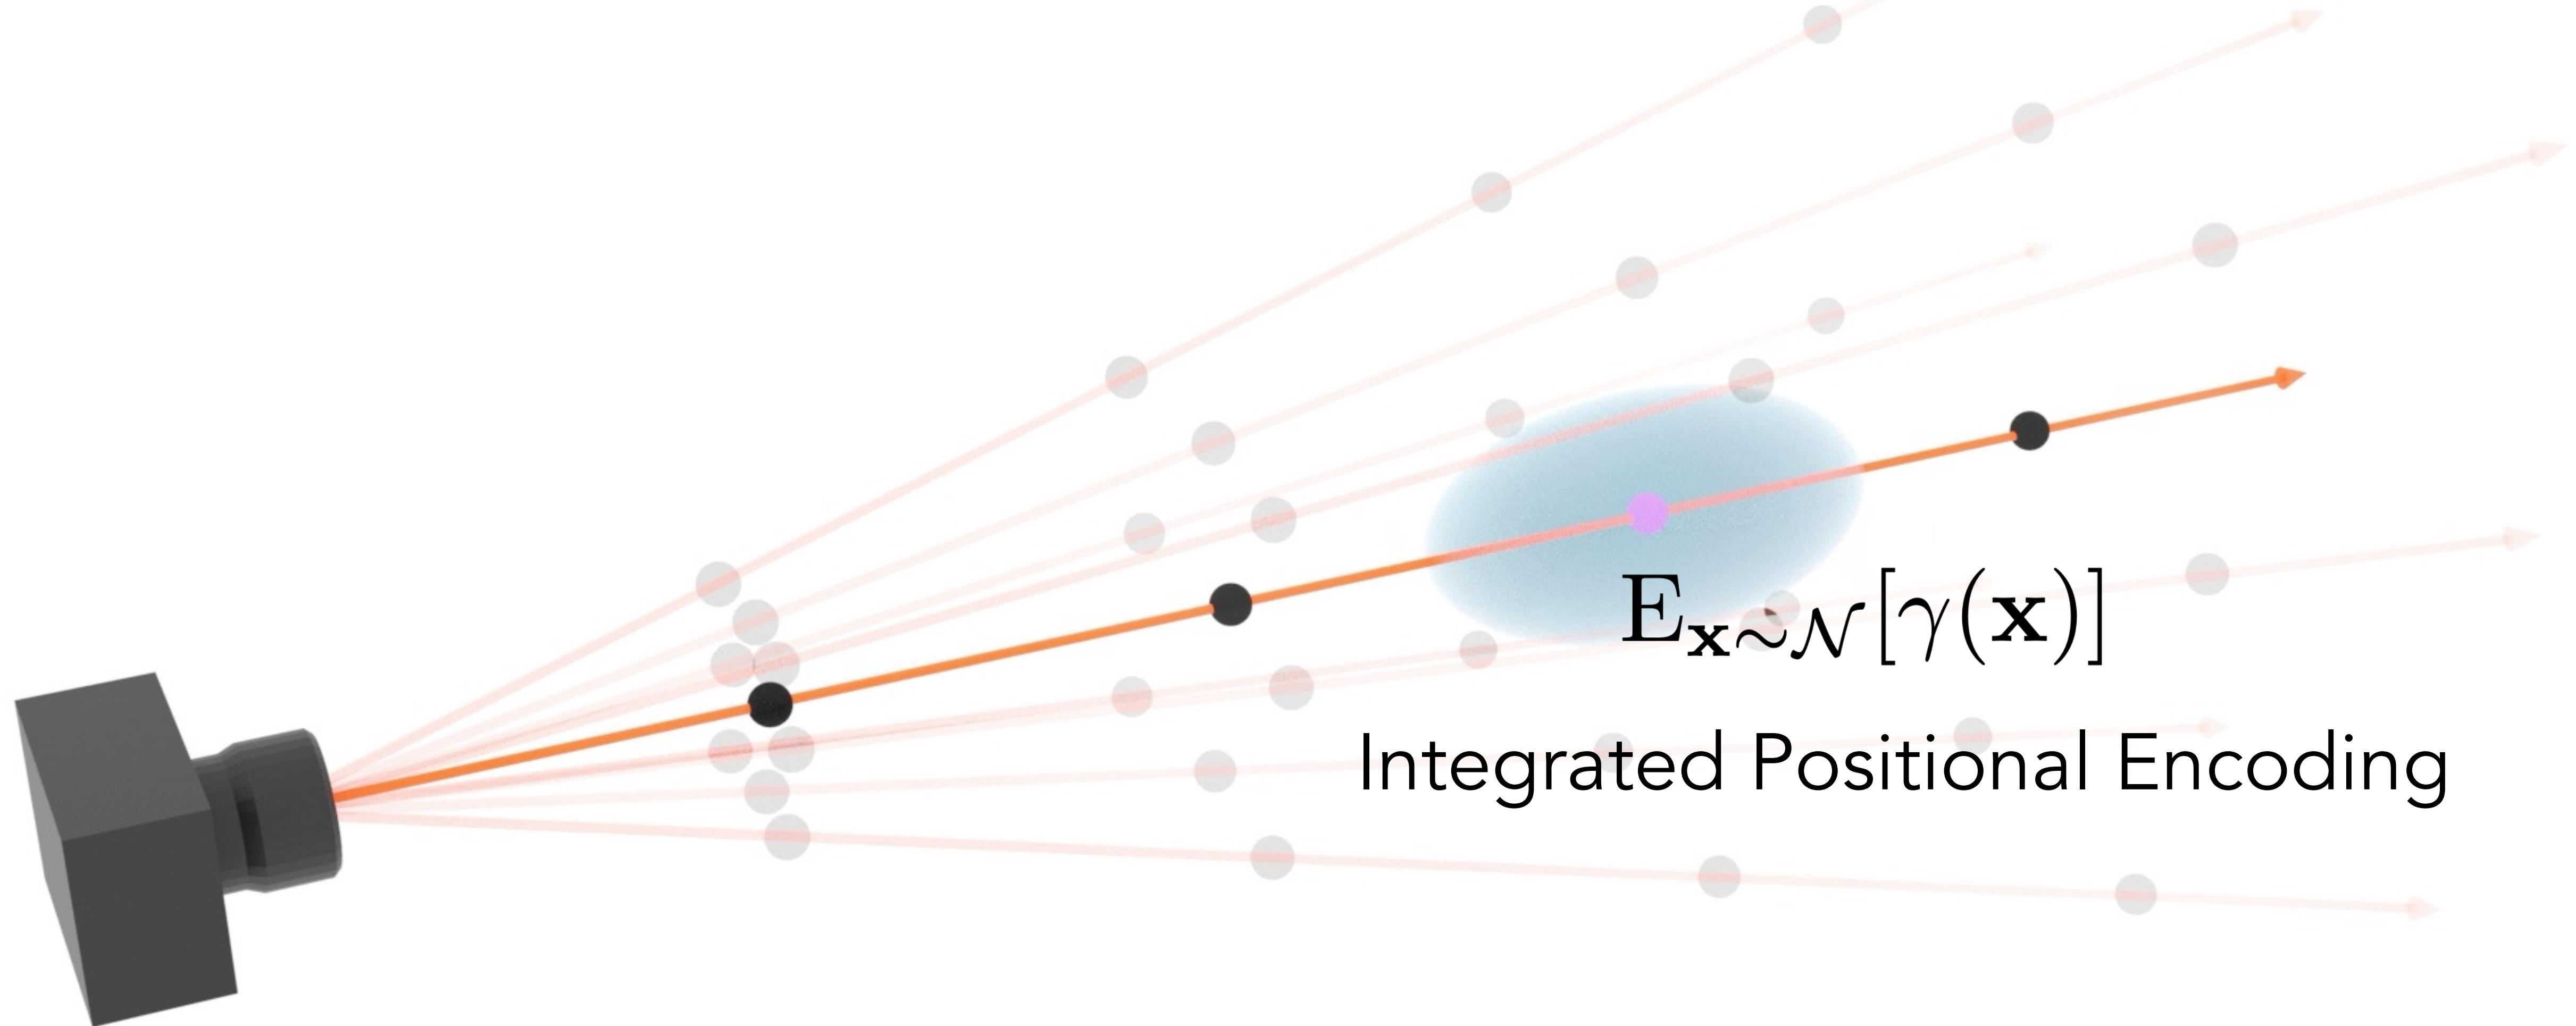
# Want NeRF to represent *integrals* within frustum



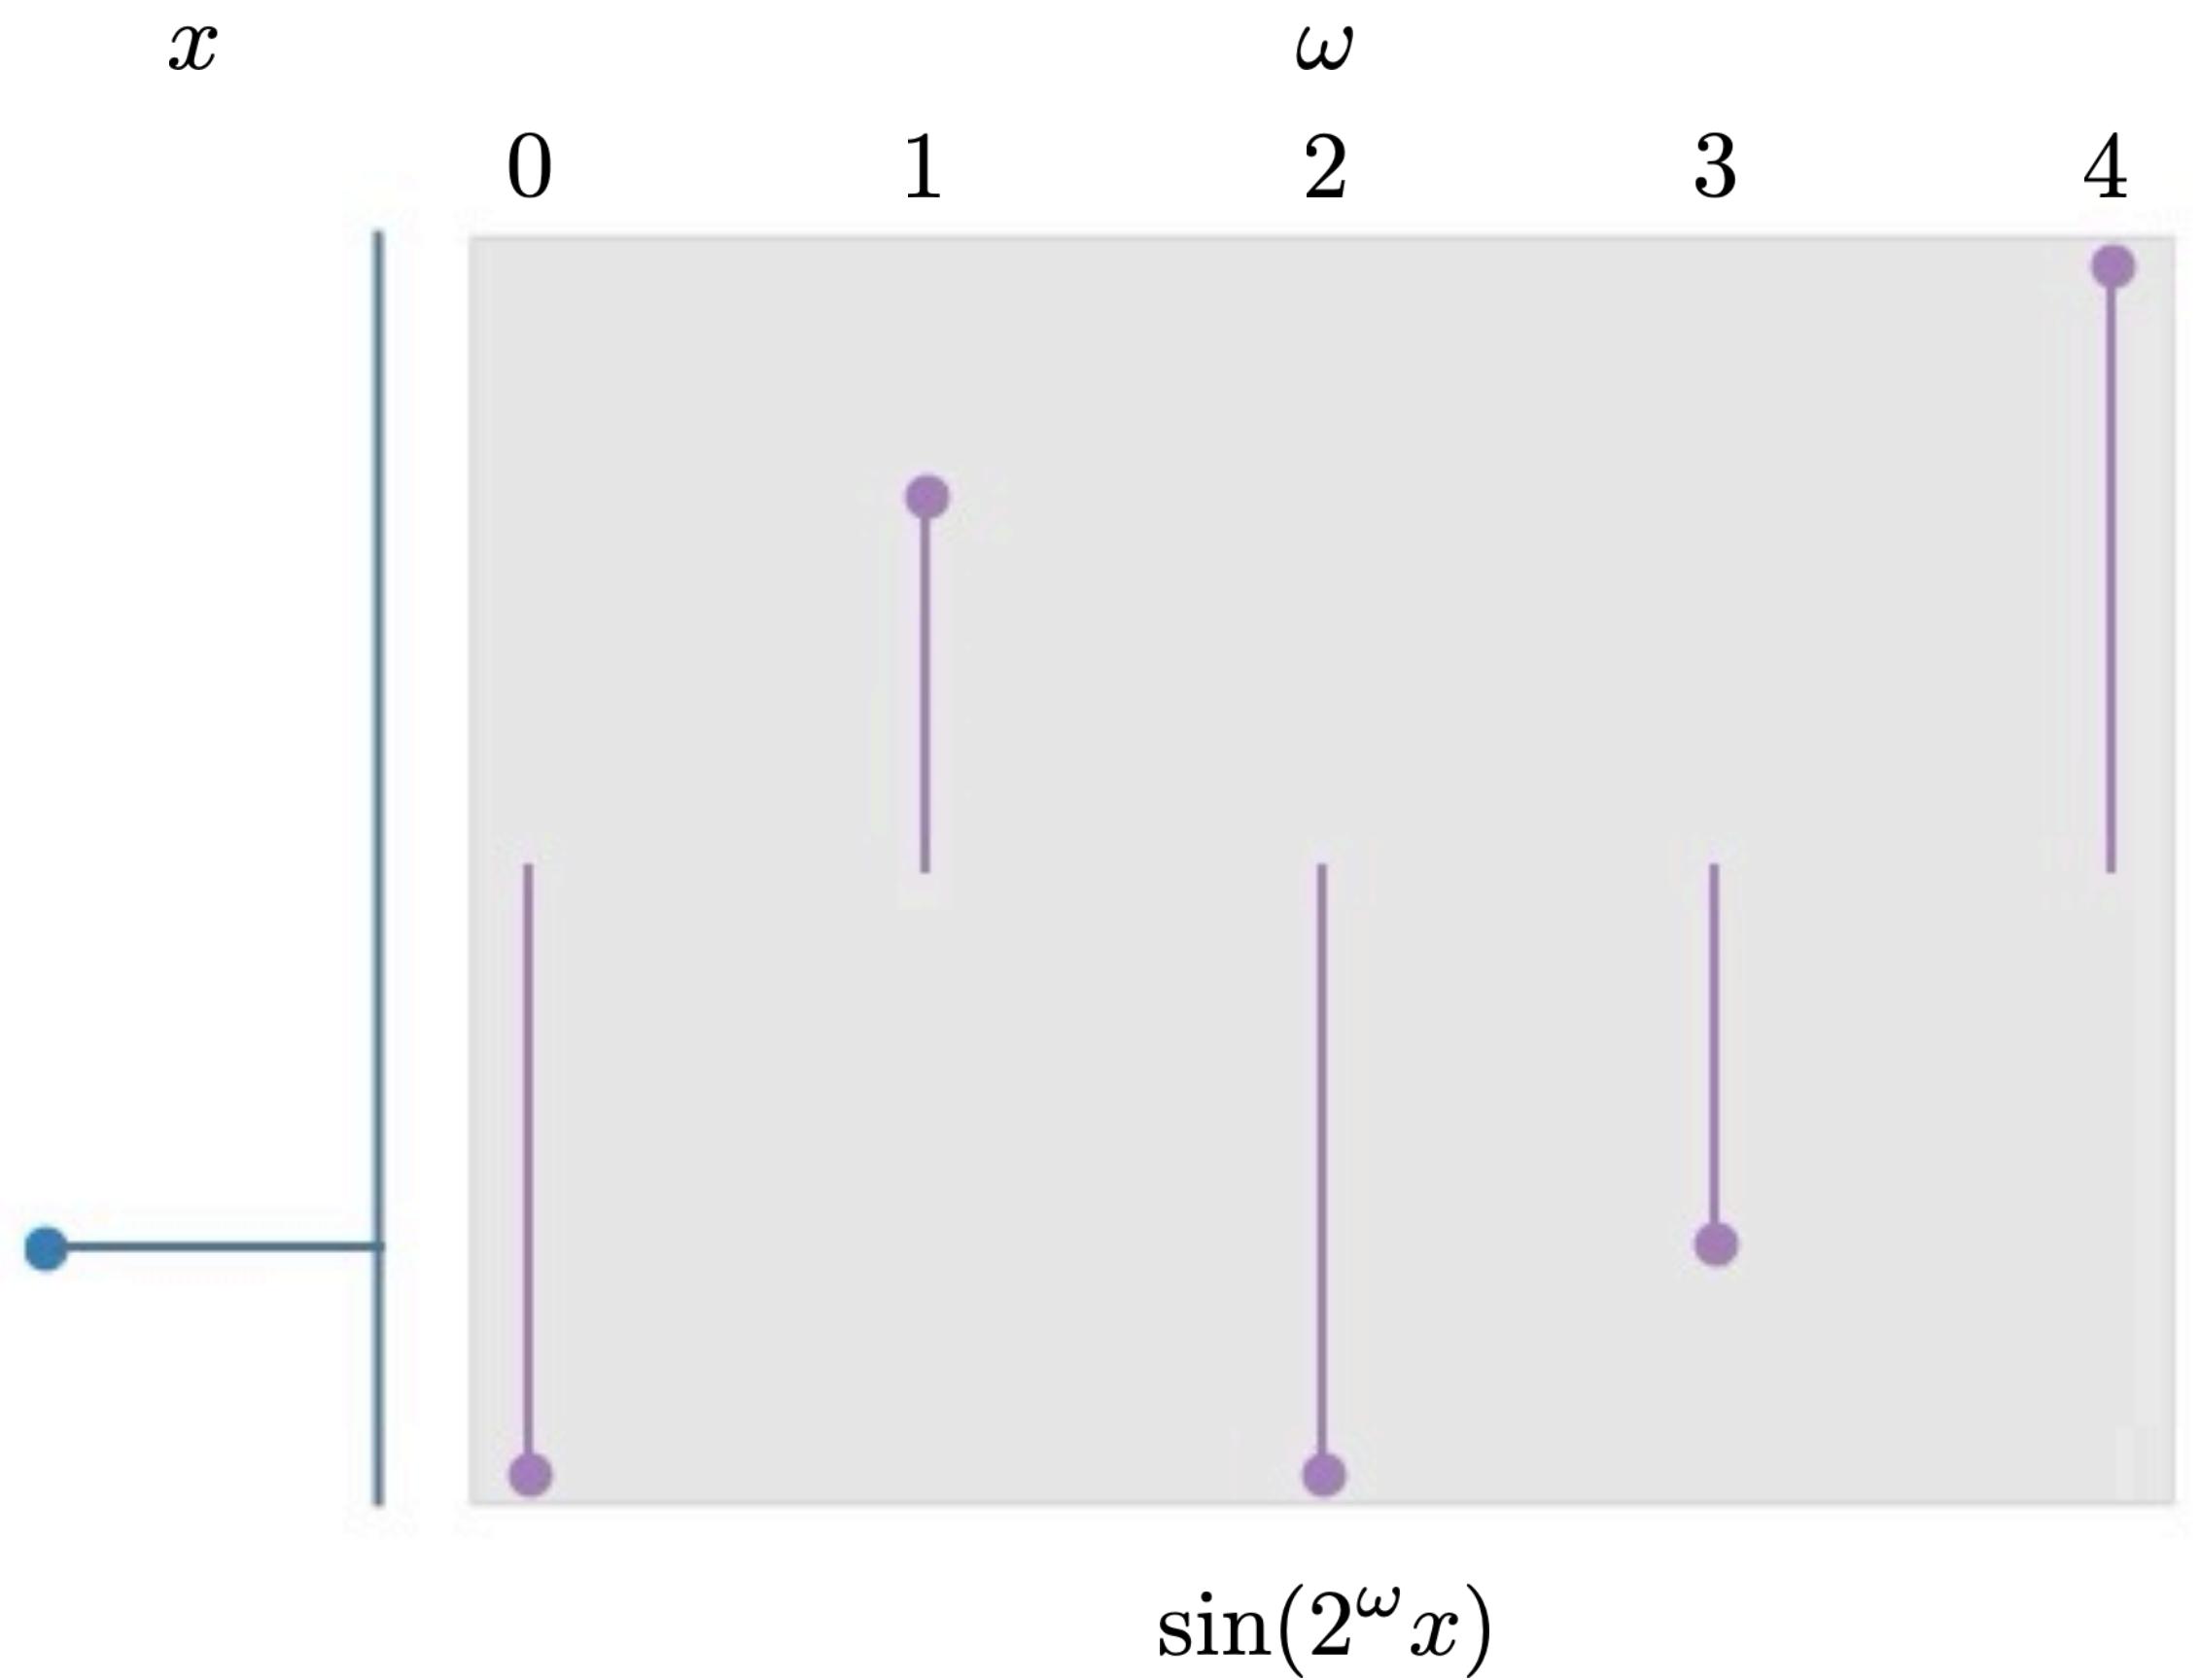
# Instead of using positional encoding of a point...



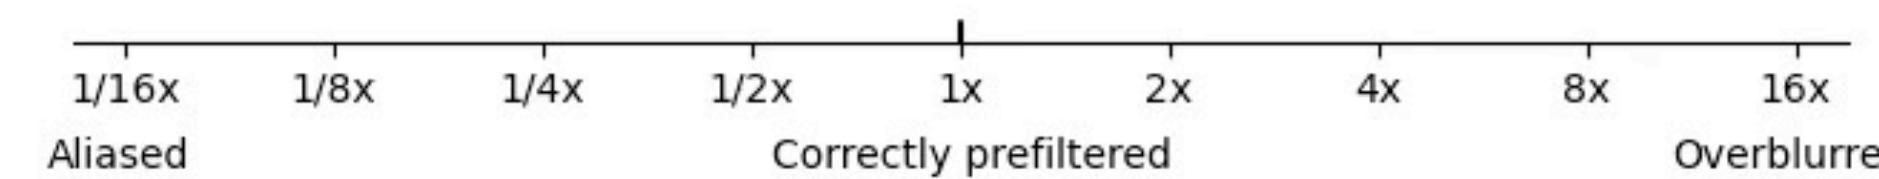
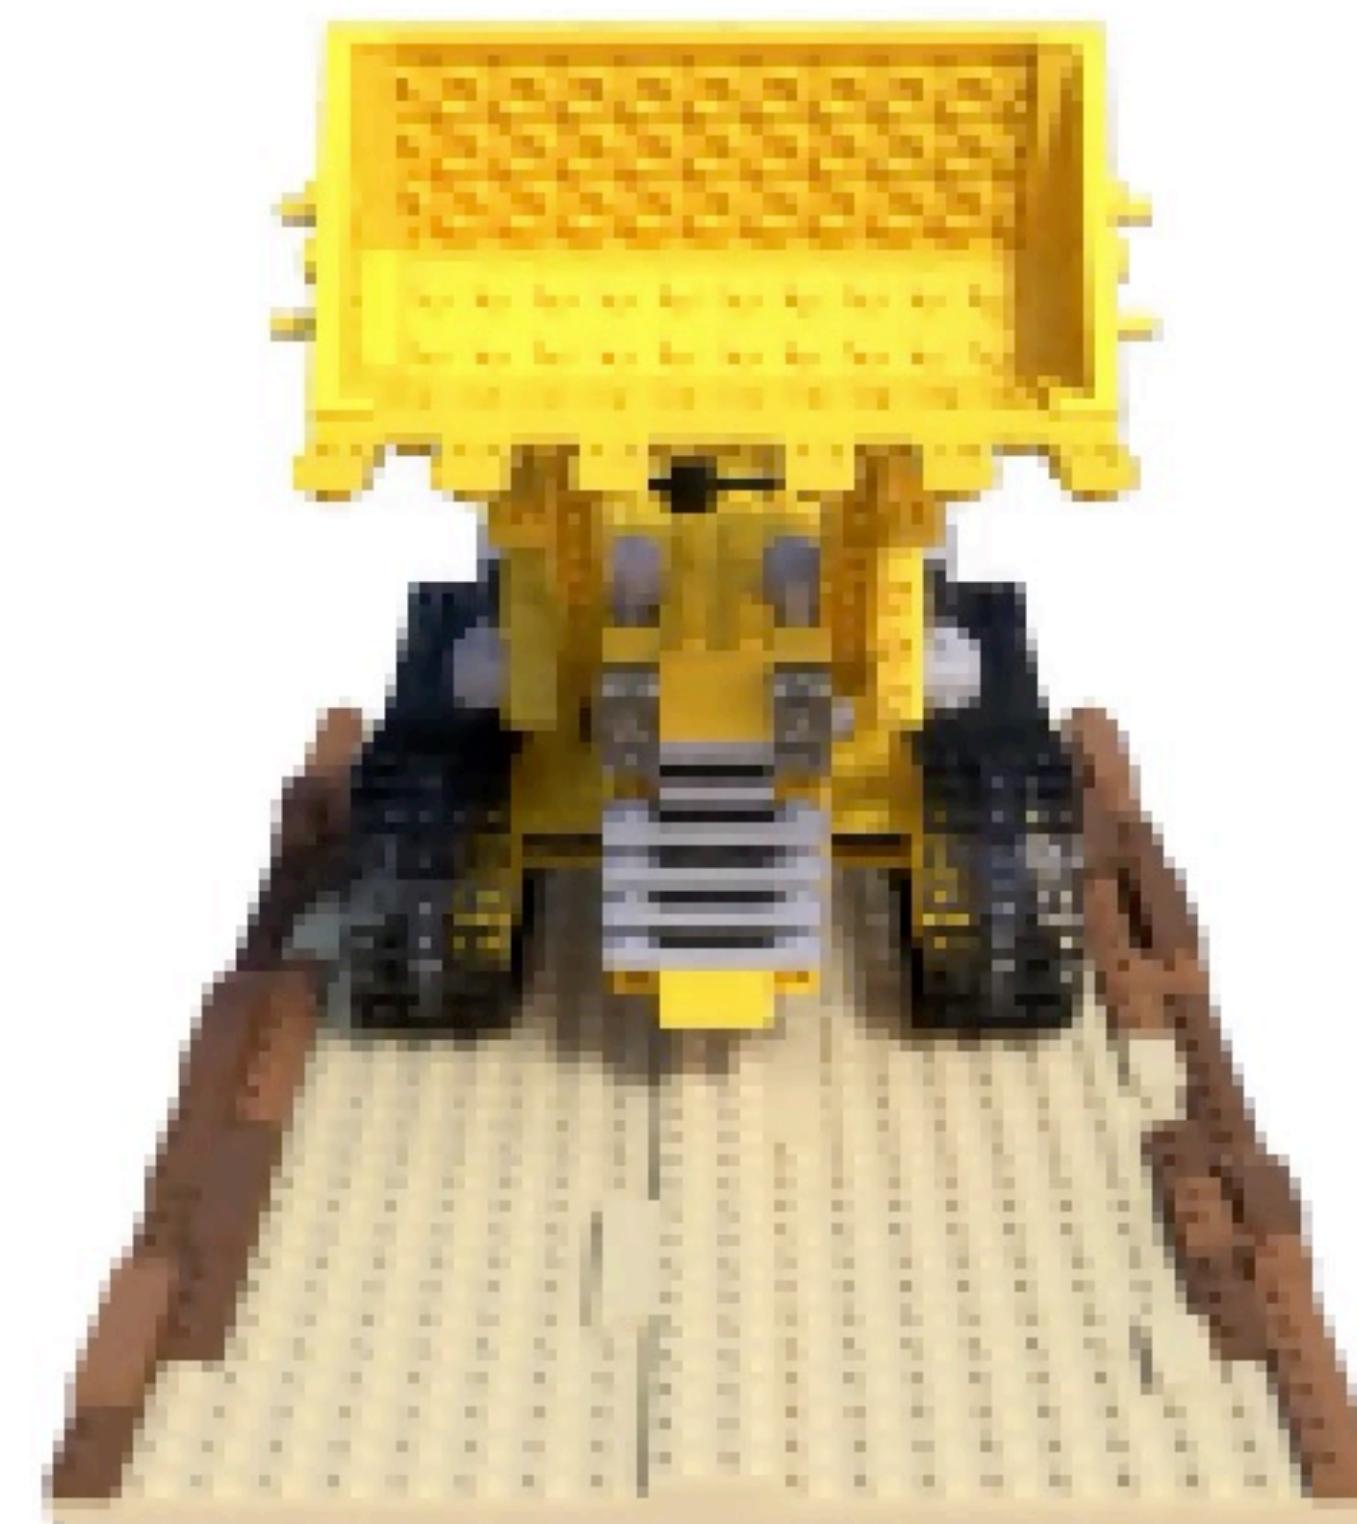
# mip-NeRF uses *integrated* positional encoding



# Positional Encoding

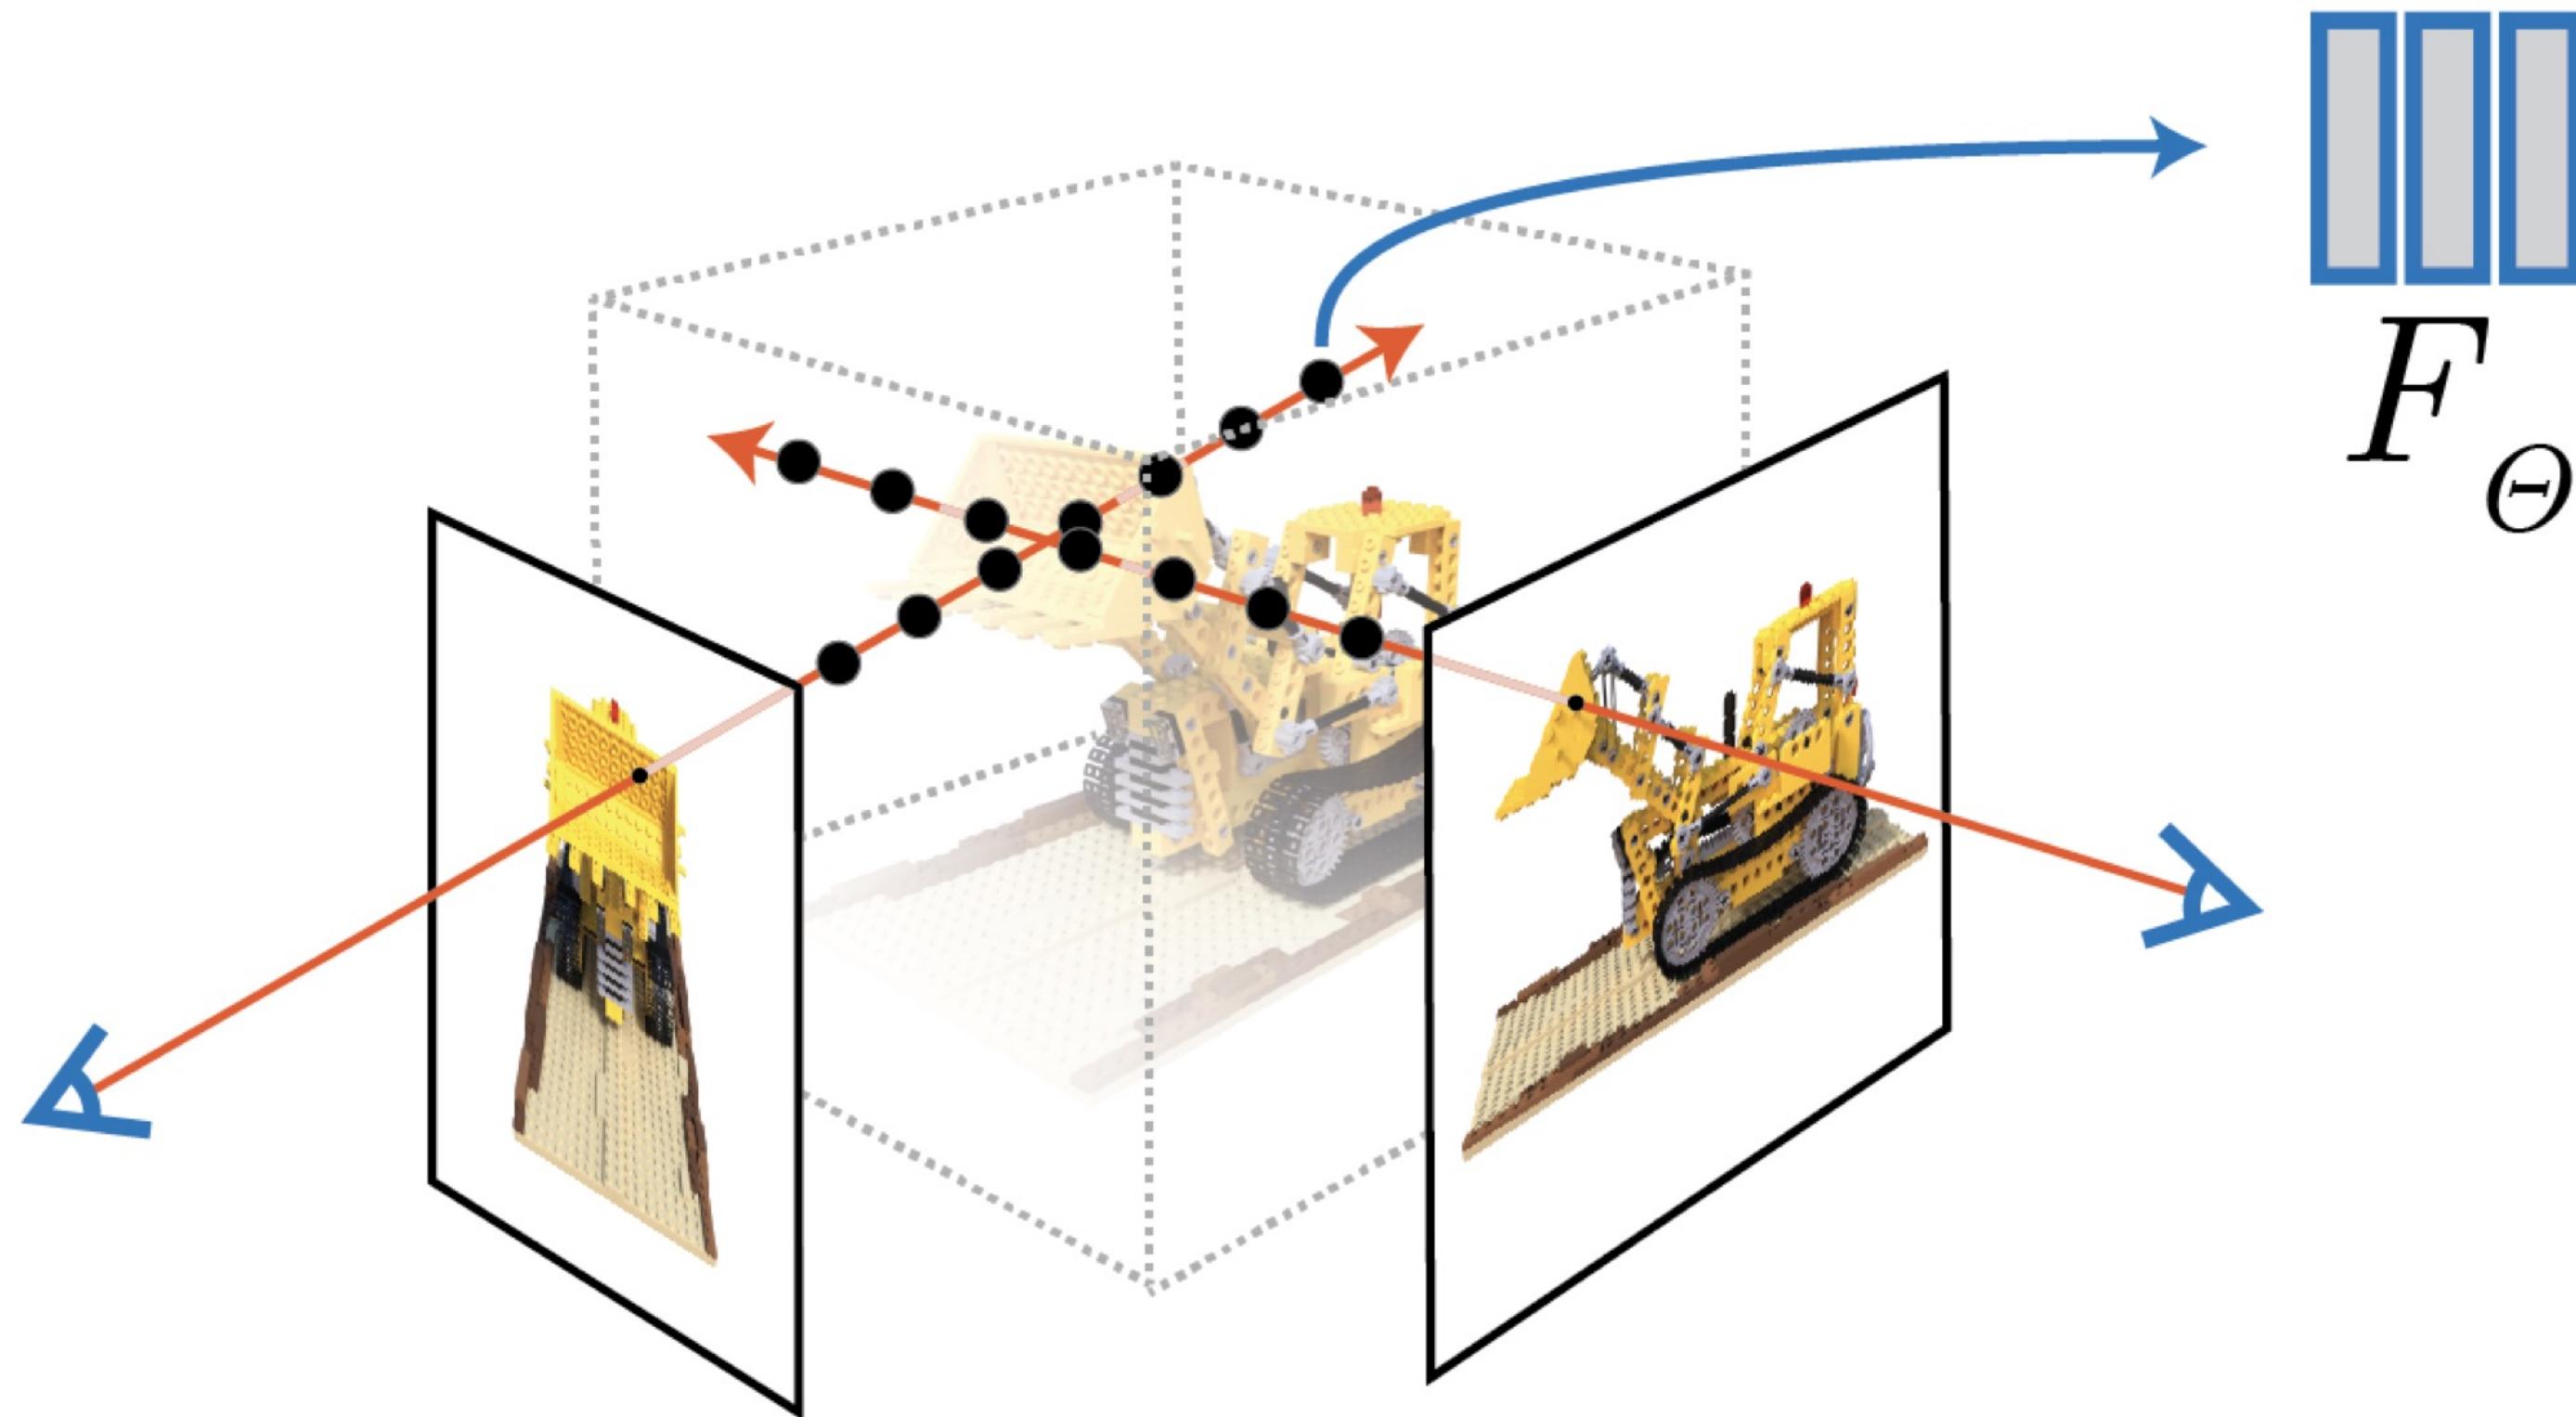


# Integrated positional encoding can reasonably approximate prefiltering

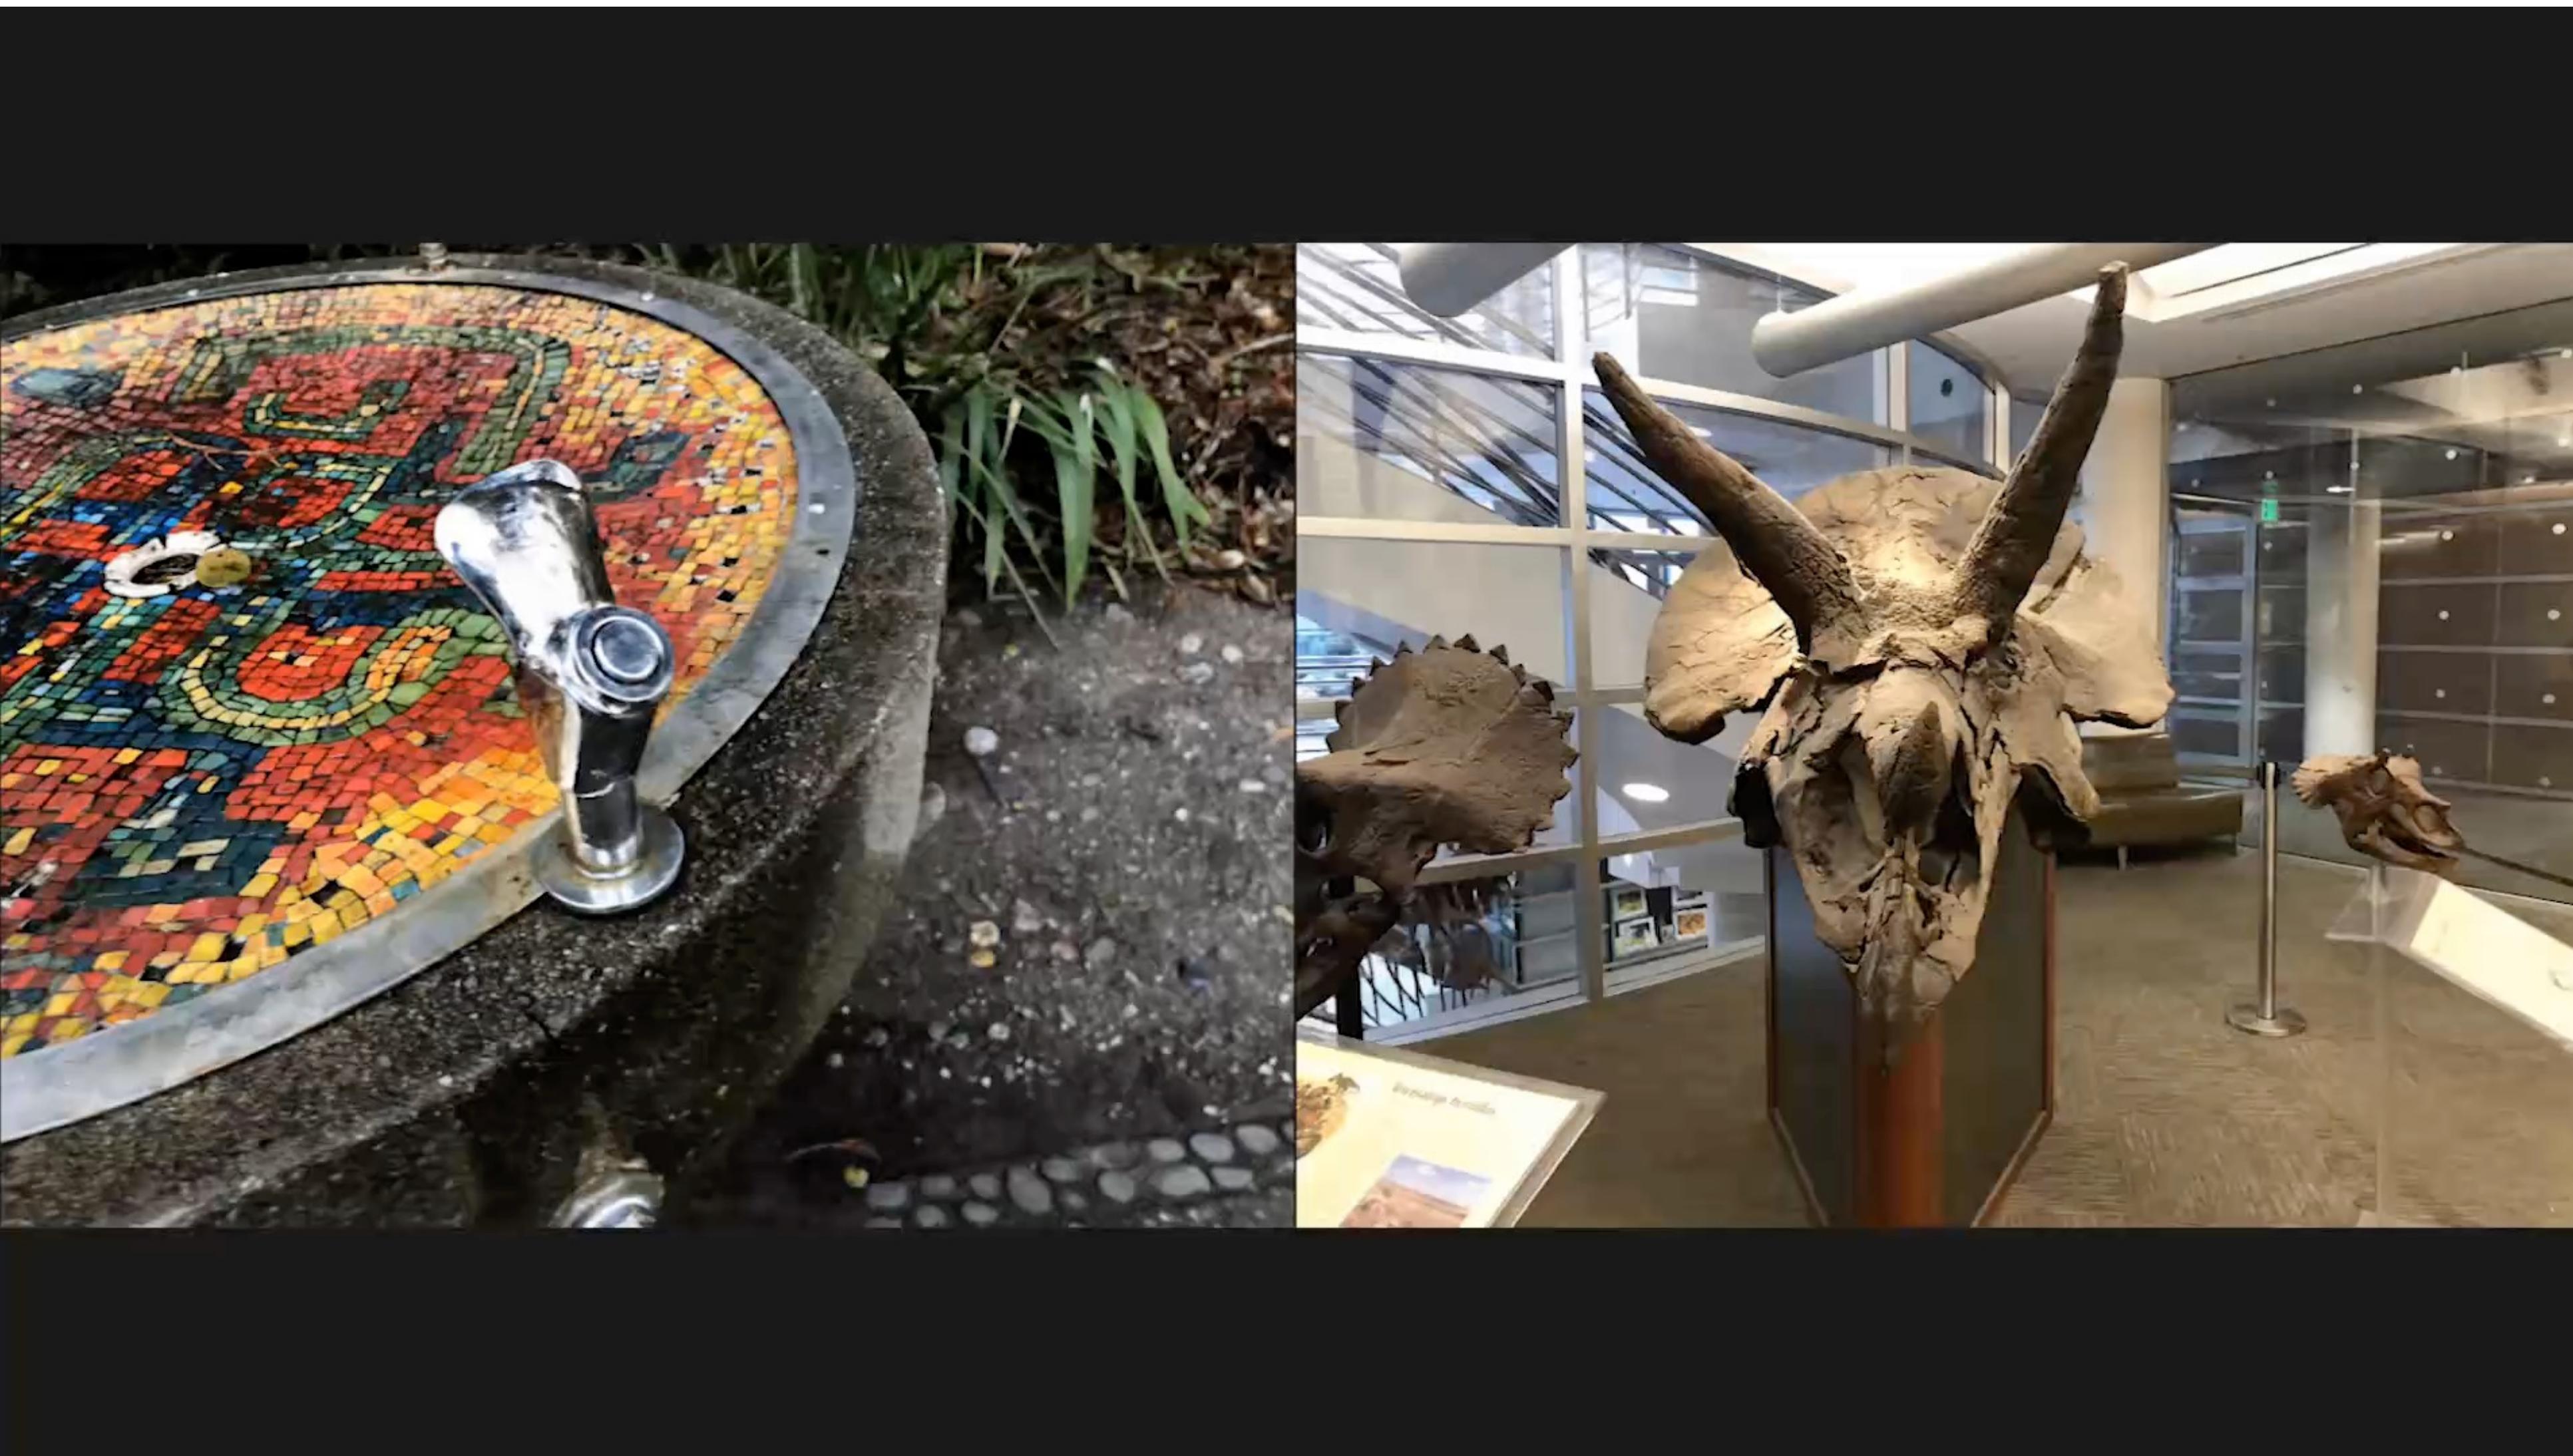


# Parameterizing 3D Space

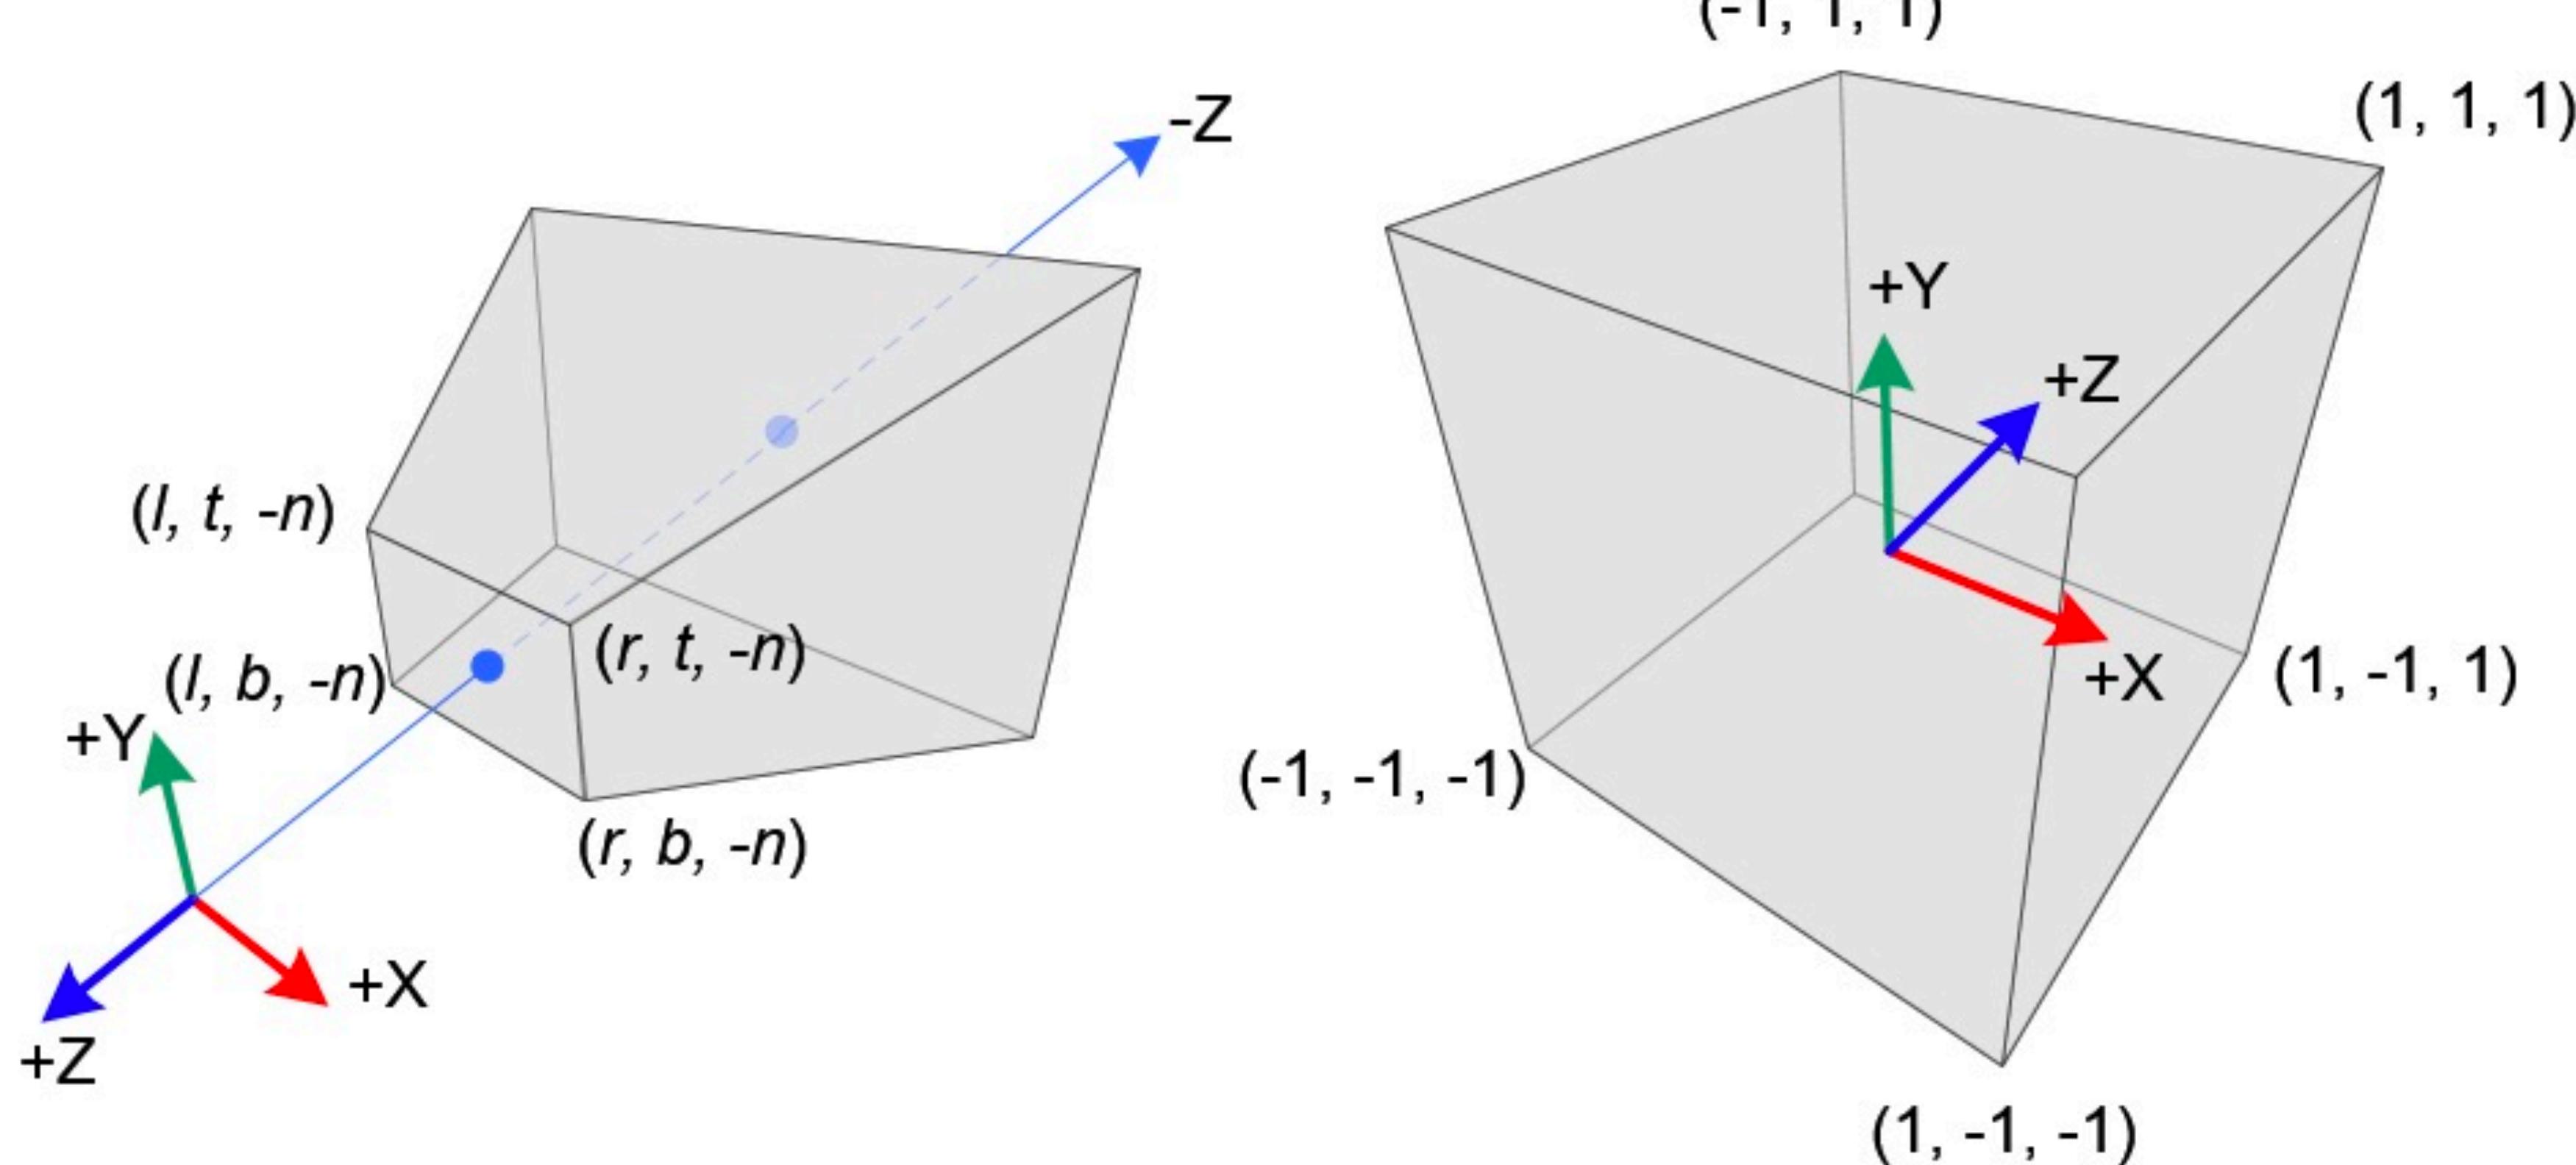
# Standard coordinates for bounded volumes



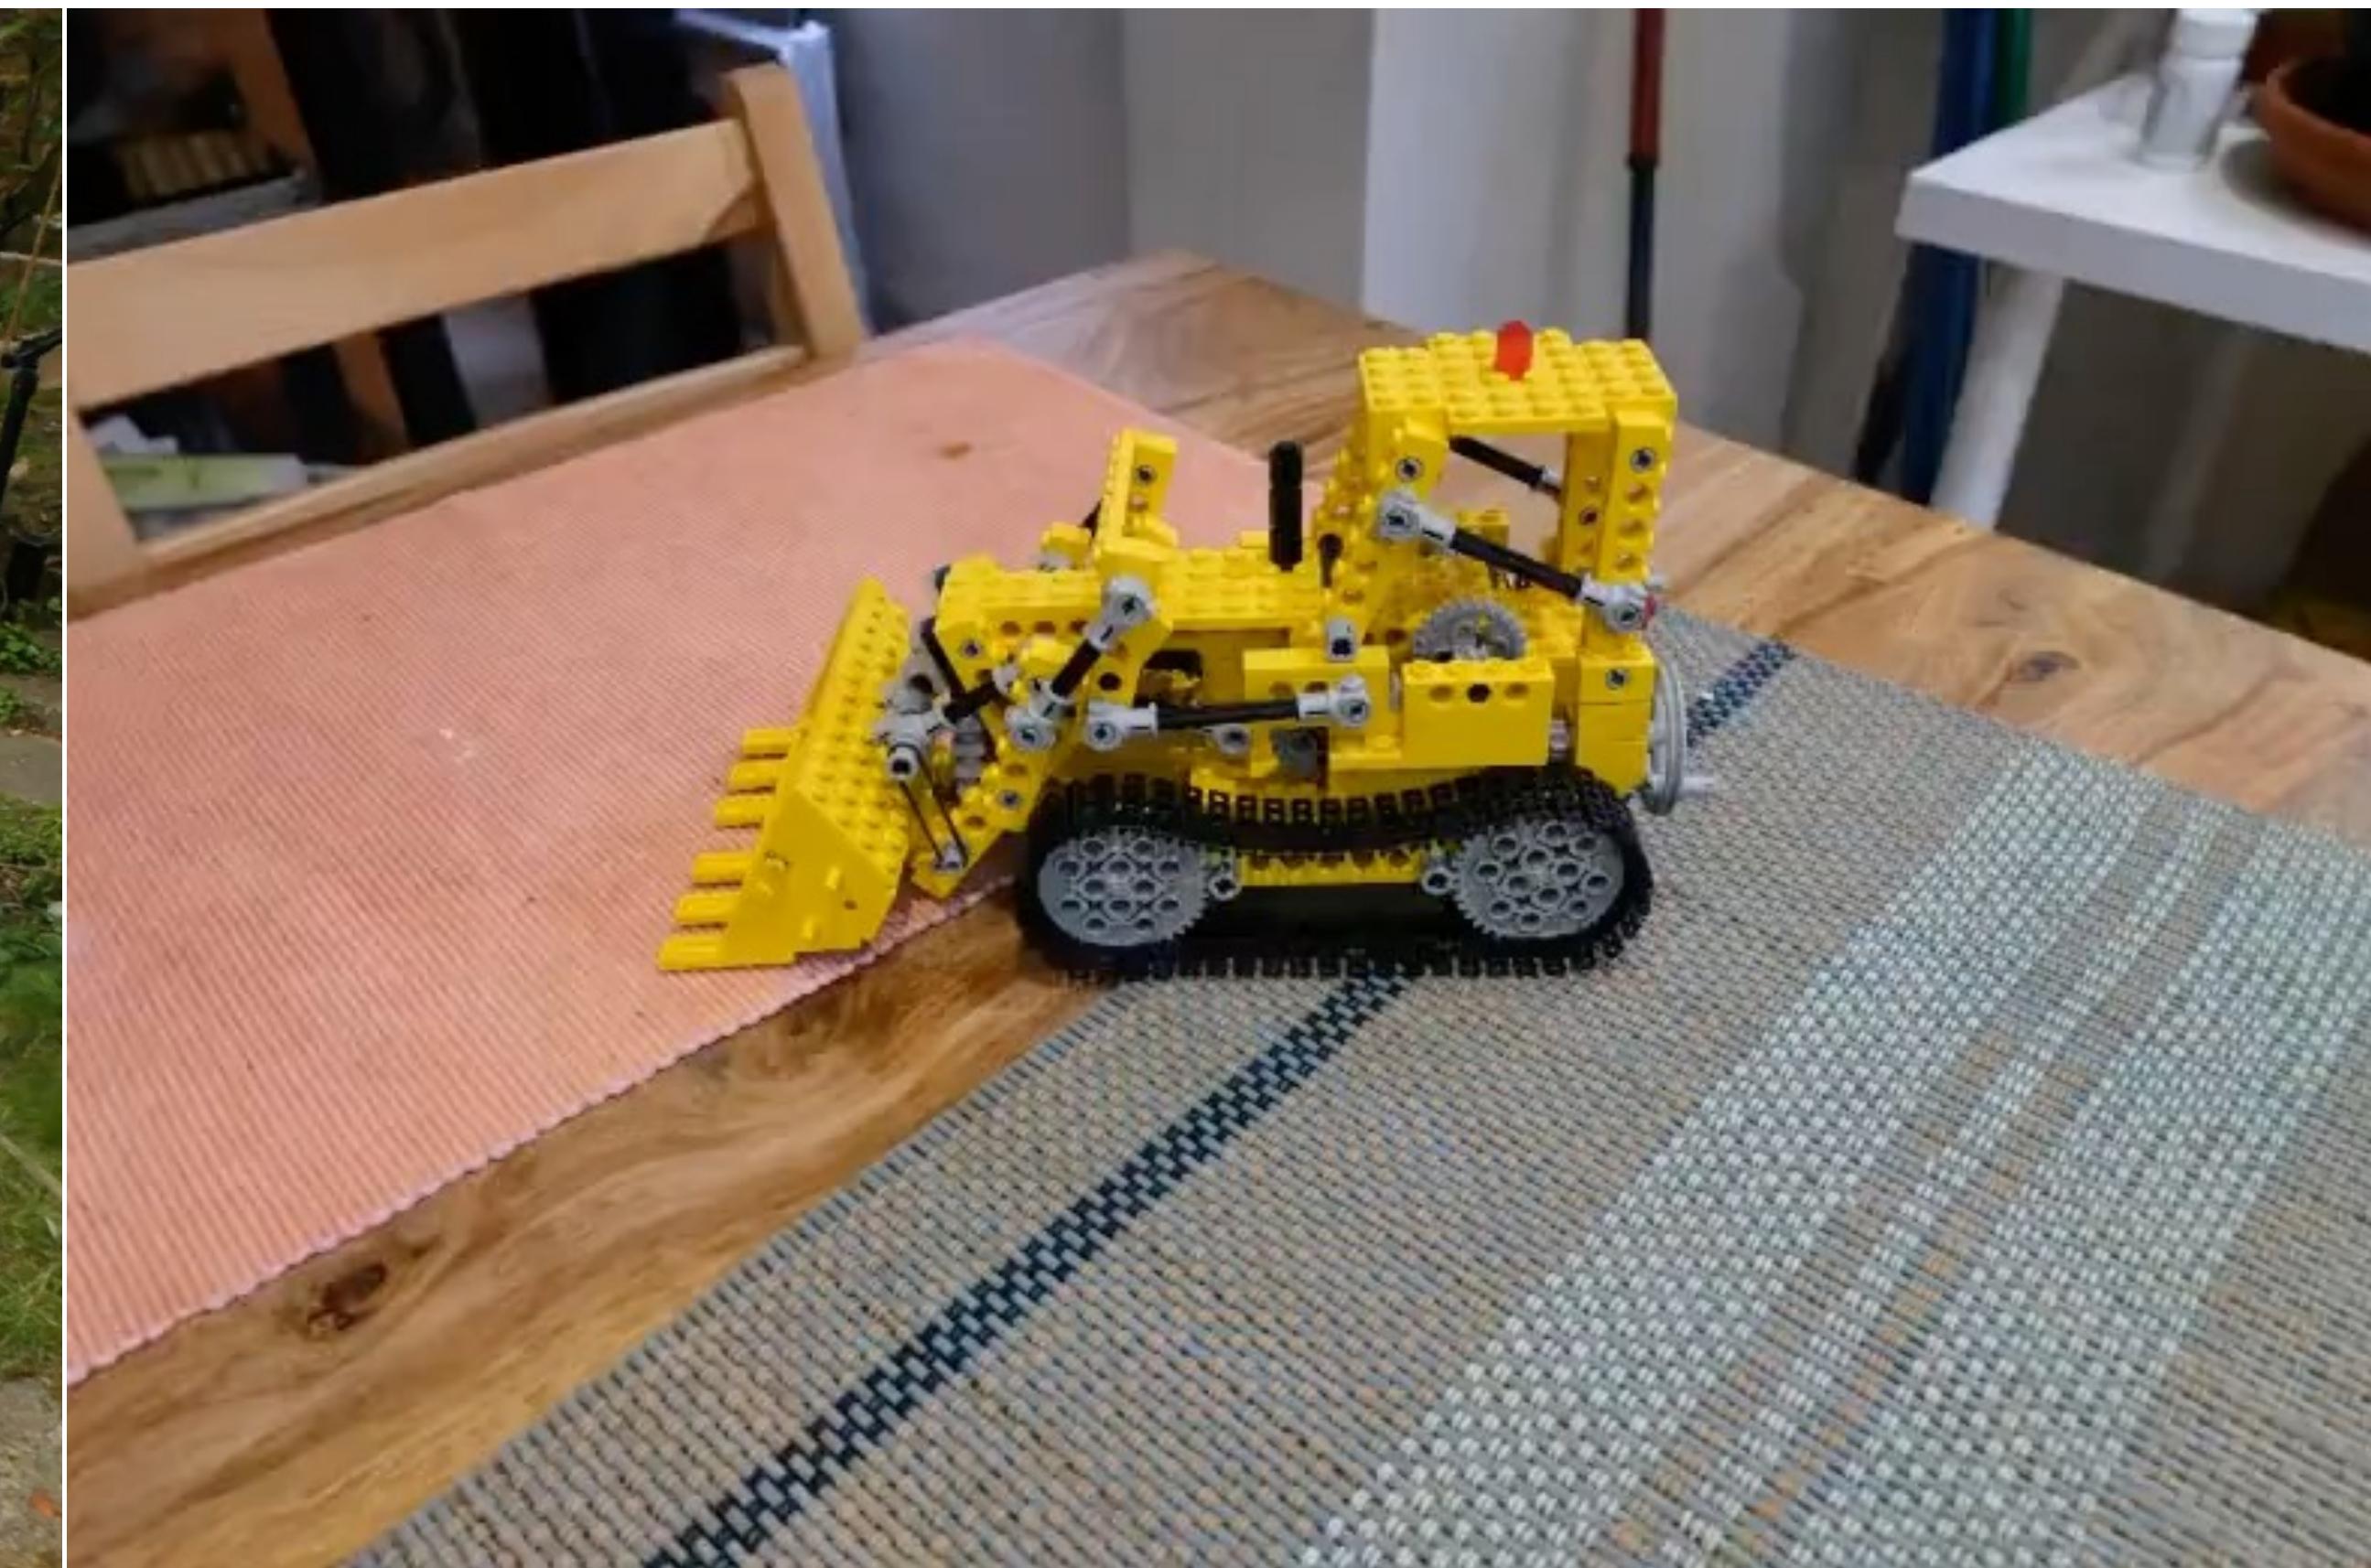
# Normalized device coordinates for unbounded “forwards-facing” volumes



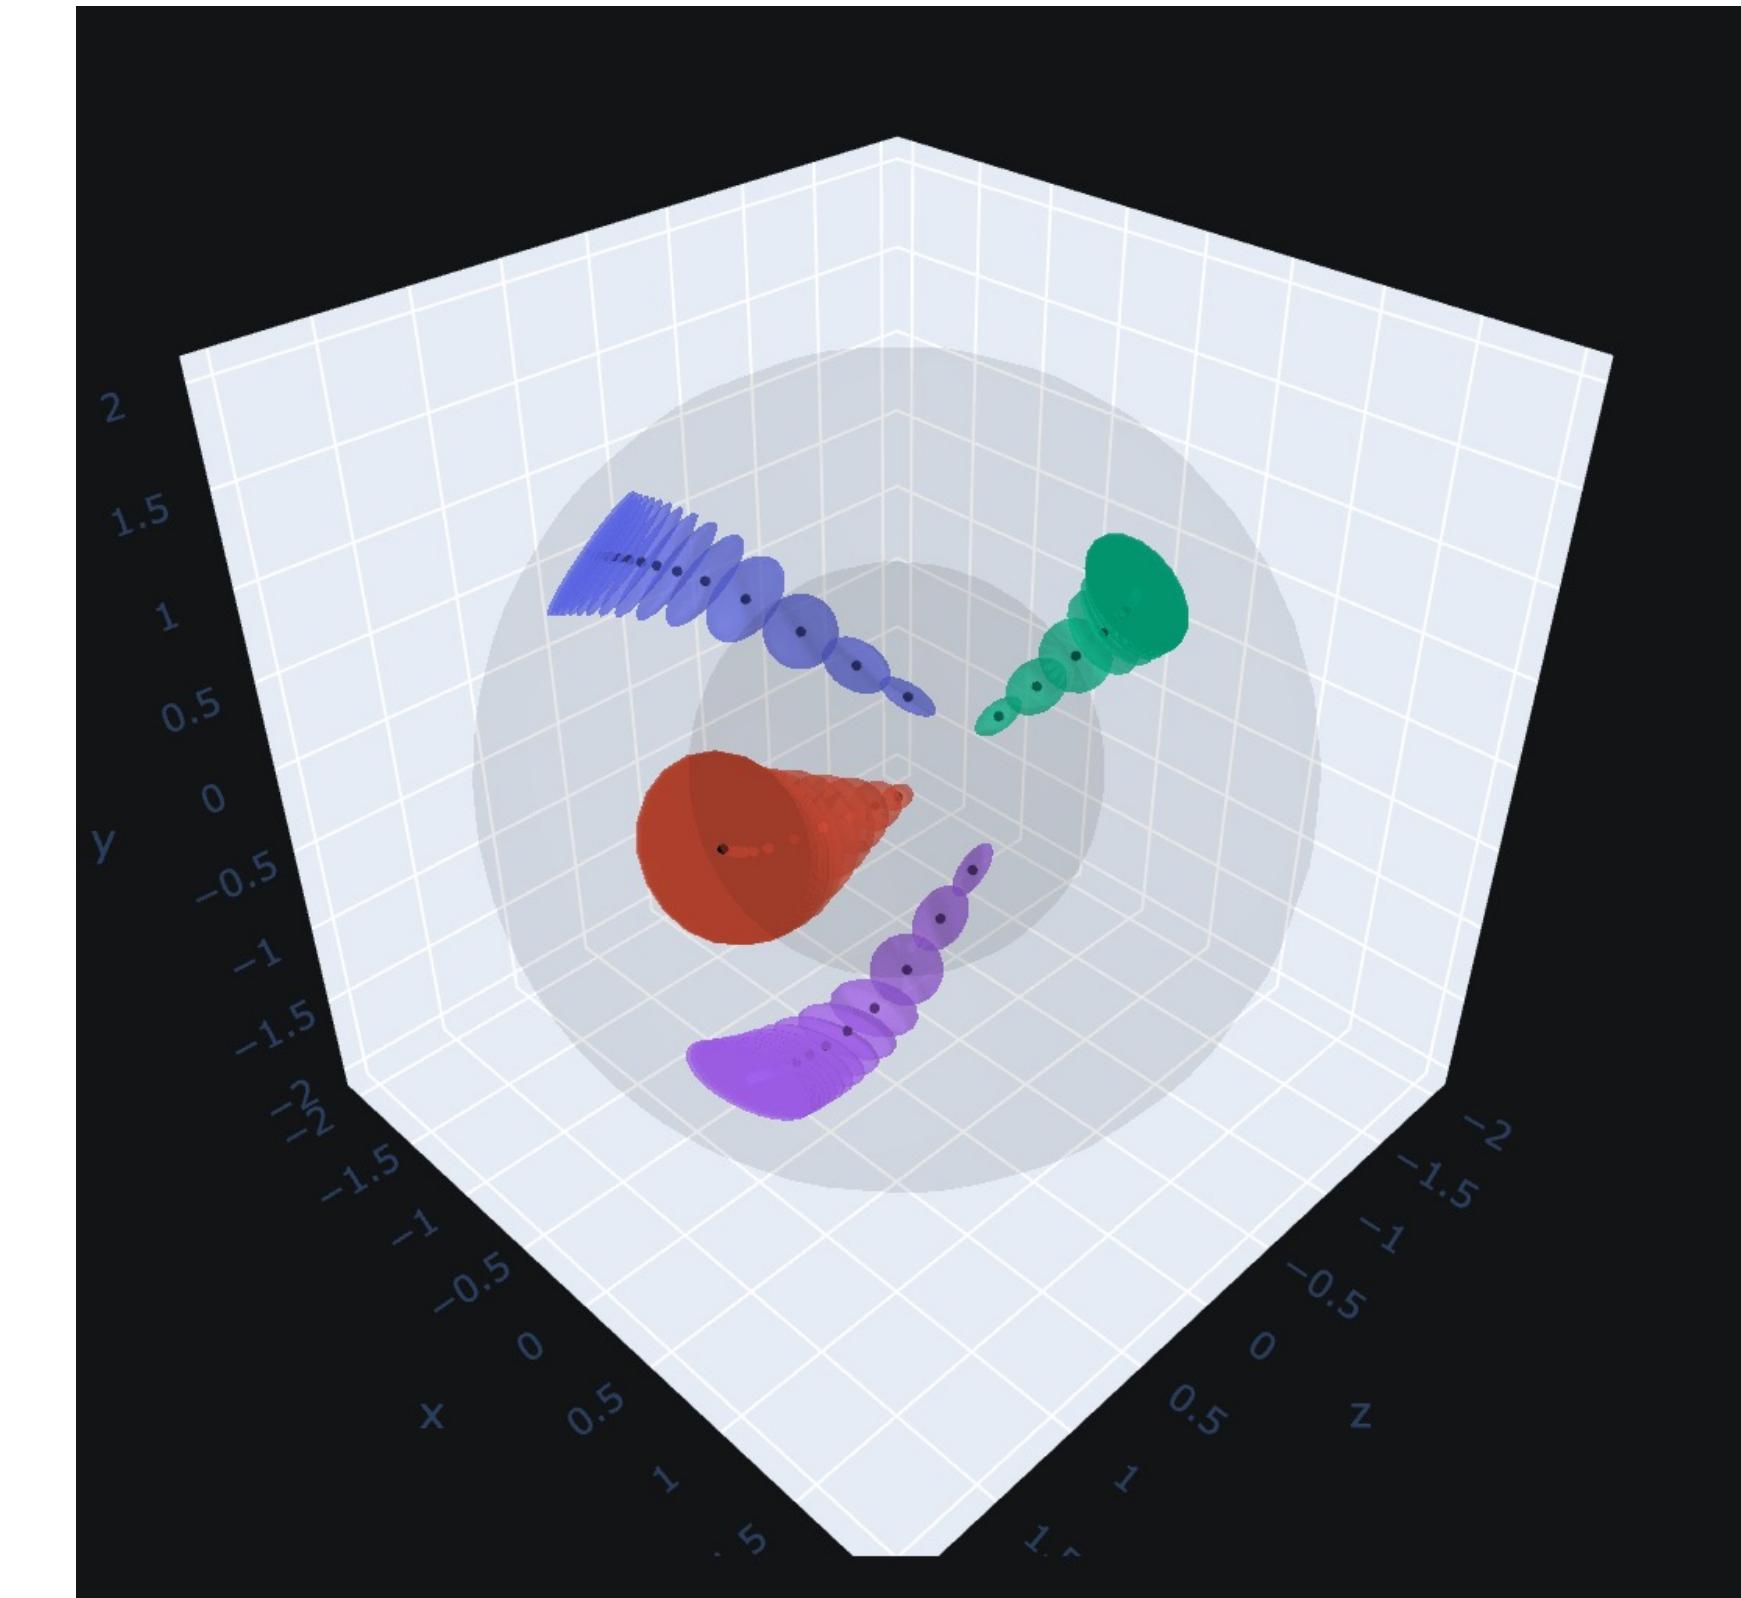
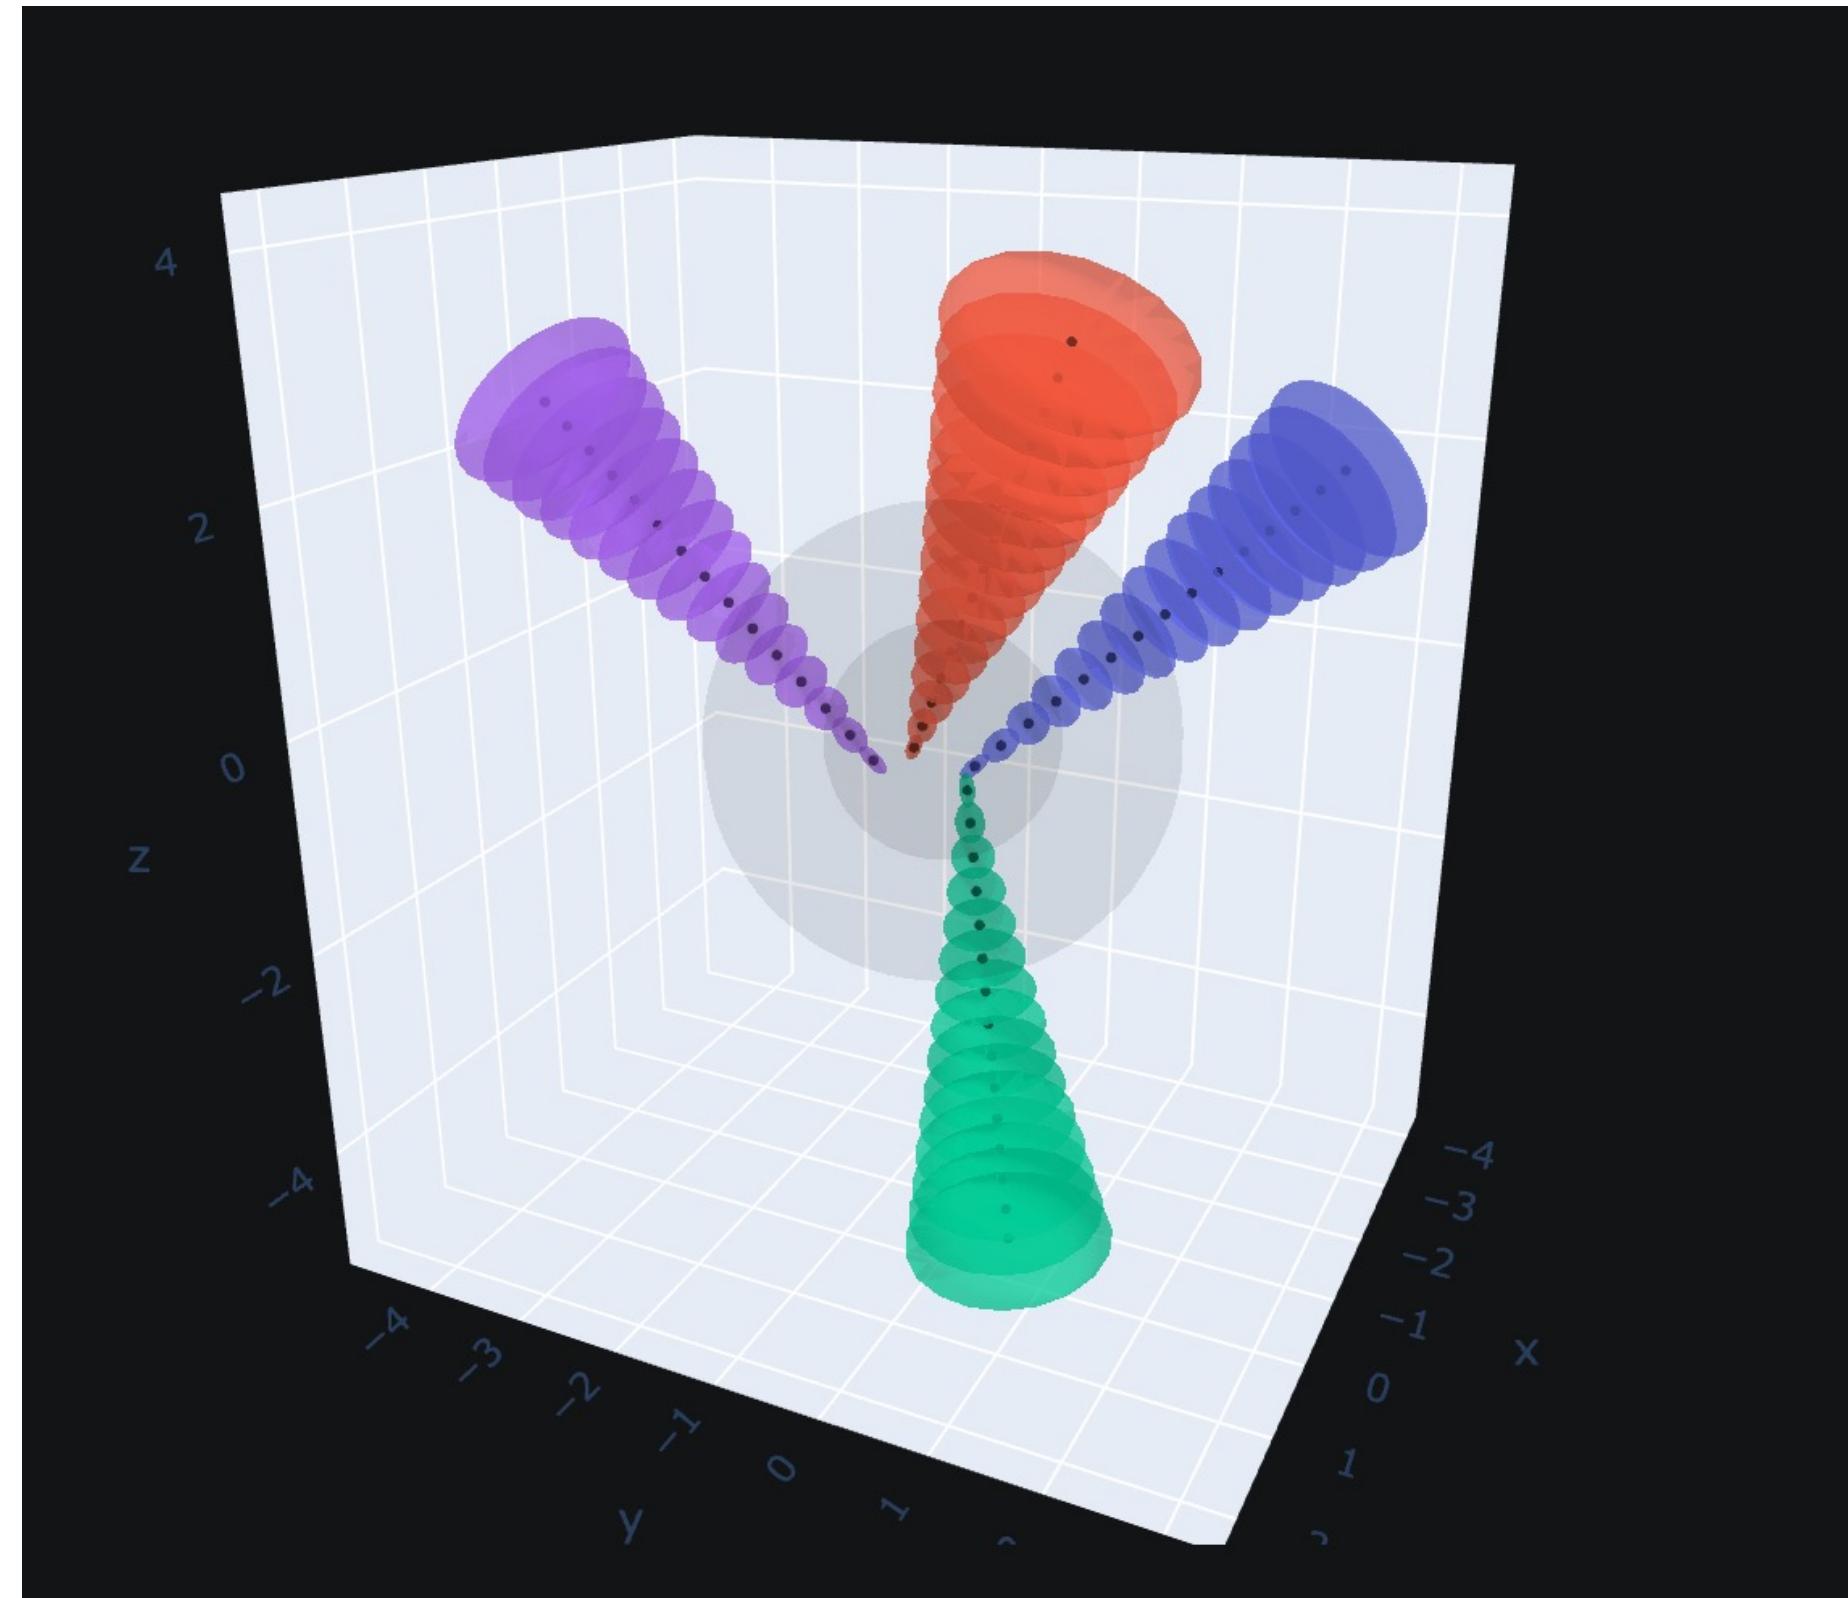
# Normalized device coordinates for unbounded “forwards-facing” volumes



# How to parameterize fully unbounded volumes?



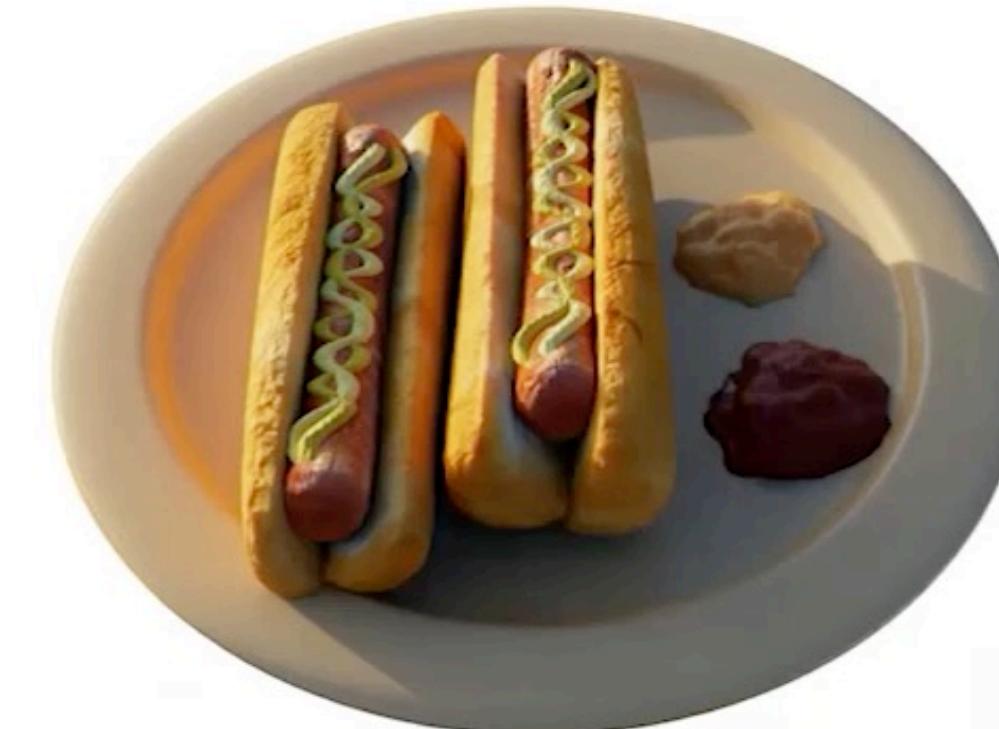
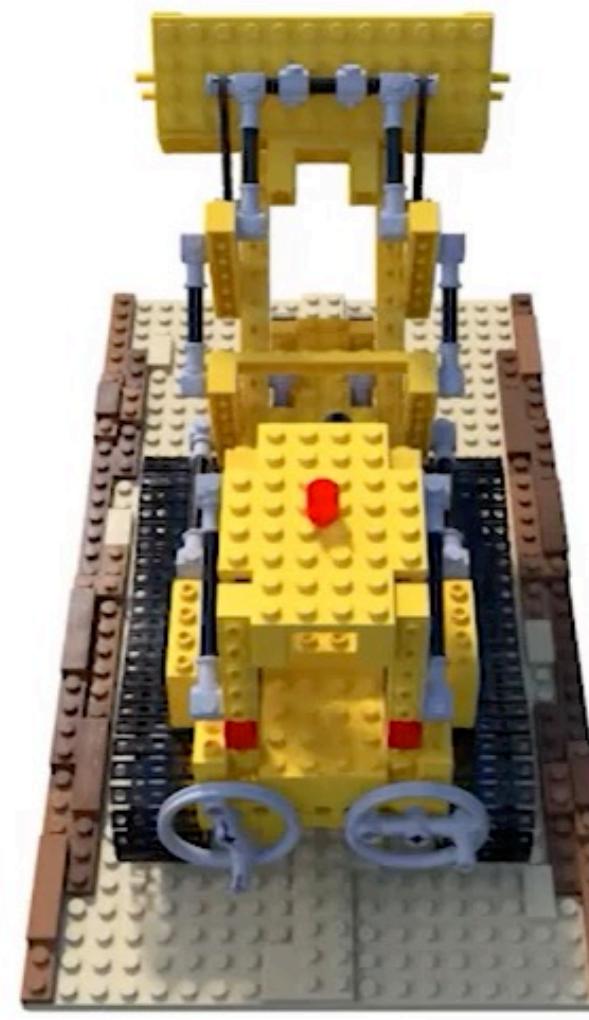
# Continuous warping of space



# Where we are

1. Birds Eye View & Background
2. Volumetric Rendering Function
3. Encoding and Representing 3D Volumes
4. Signal Processing Considerations
5. Challenges & Pointers

# Caught up with the core NeRF pieces!



**What are the remaining challenges?**

**Don't worry there is a lot!**

# The neverending list of NeRF limitations (back in 2020)

- Expensive / slow to train
- Expensive / slow to render
- Sensitive to sampling strategy
- Sensitive to pose accuracy
- Assumes static lighting
- Not a mesh
- Assumes static scene
- Does not generalize between scenes

# The neverending list of NeRF limitations (back in 2020)

- ~~Expensive / slow to train~~
- ~~Expensive / slow to render~~
- Sensitive to sampling strategy
- Sensitive to pose accuracy
- Assumes static lighting
- Not a mesh
- Assumes static scene
- Does not generalize between scenes

# The neverending list of NeRF limitations (back in 2020)

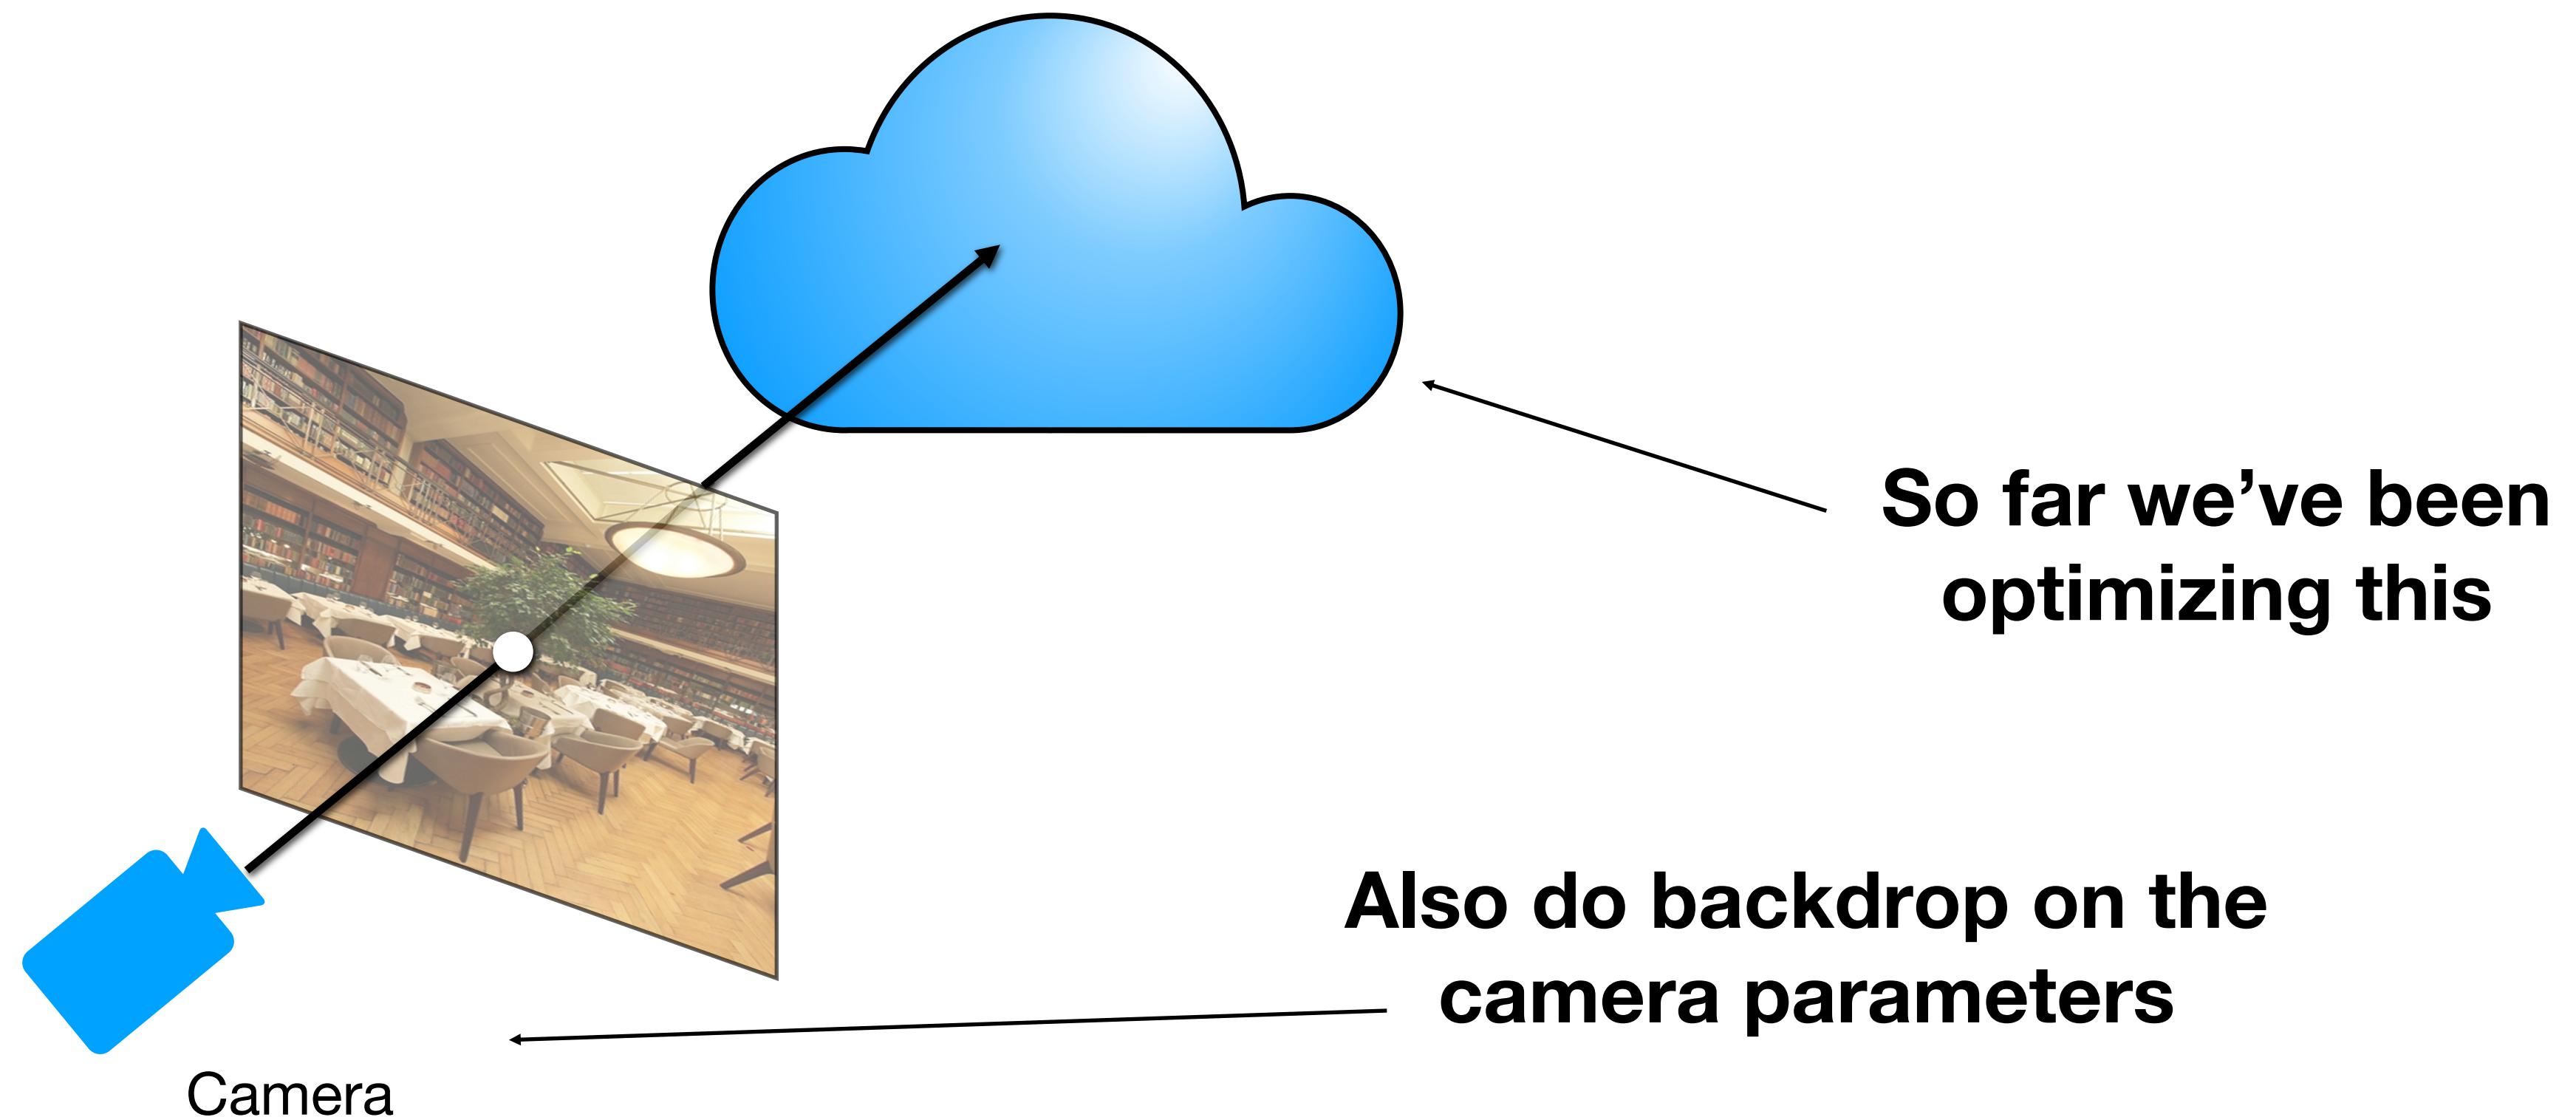
- ~~Expensive / slow to train~~
- ~~Expensive / slow to render~~
- ~~Sensitive to sampling strategy~~
- Sensitive to pose accuracy
- Assumes static lighting
- Not a mesh
- Assumes static scene
- Does not generalize between scenes

# The neverending list of NeRF limitations (back in 2020)

- ~~Expensive / slow to train~~
- ~~Expensive / slow to render~~
- ~~Sensitive to sampling strategy~~
- Sensitive to pose accuracy
- Assumes static lighting
- Not a mesh
- Assumes static scene
- Does not generalize between scenes

# Camera Quality

Small noise in the camera can be made robust by also optimizing the camera



# No Pose Optimization



# Block-NeRF

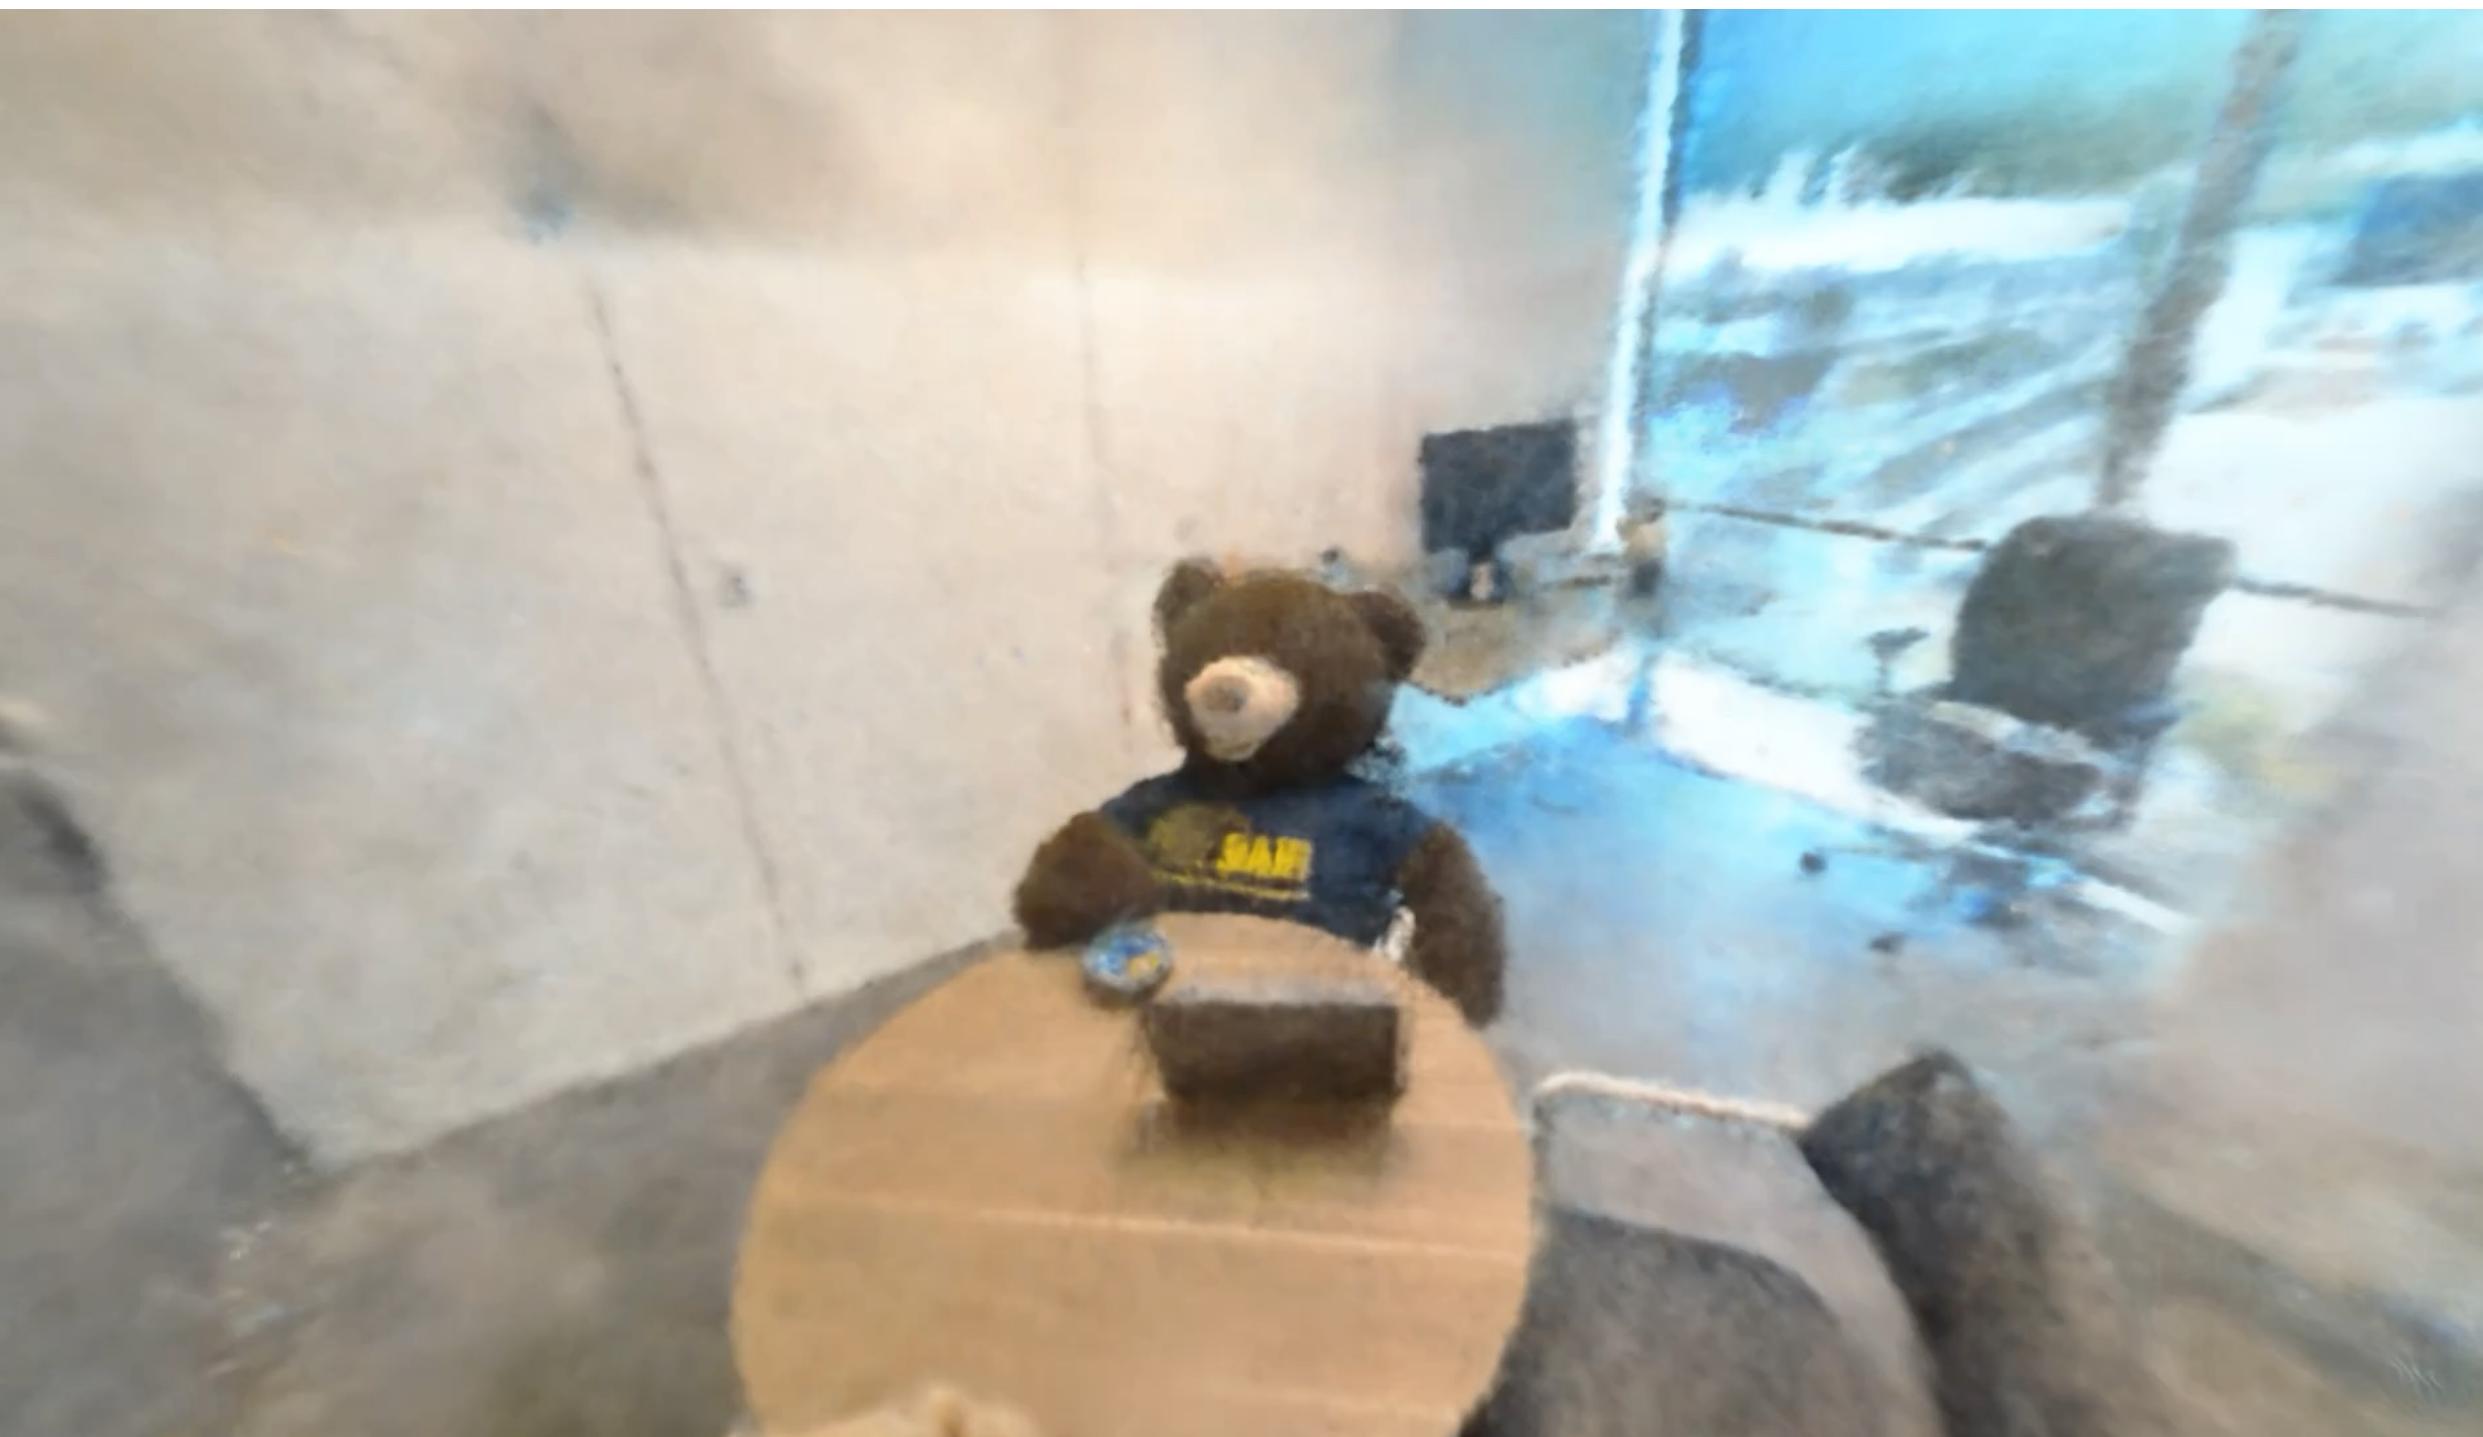


# Camera Optimization

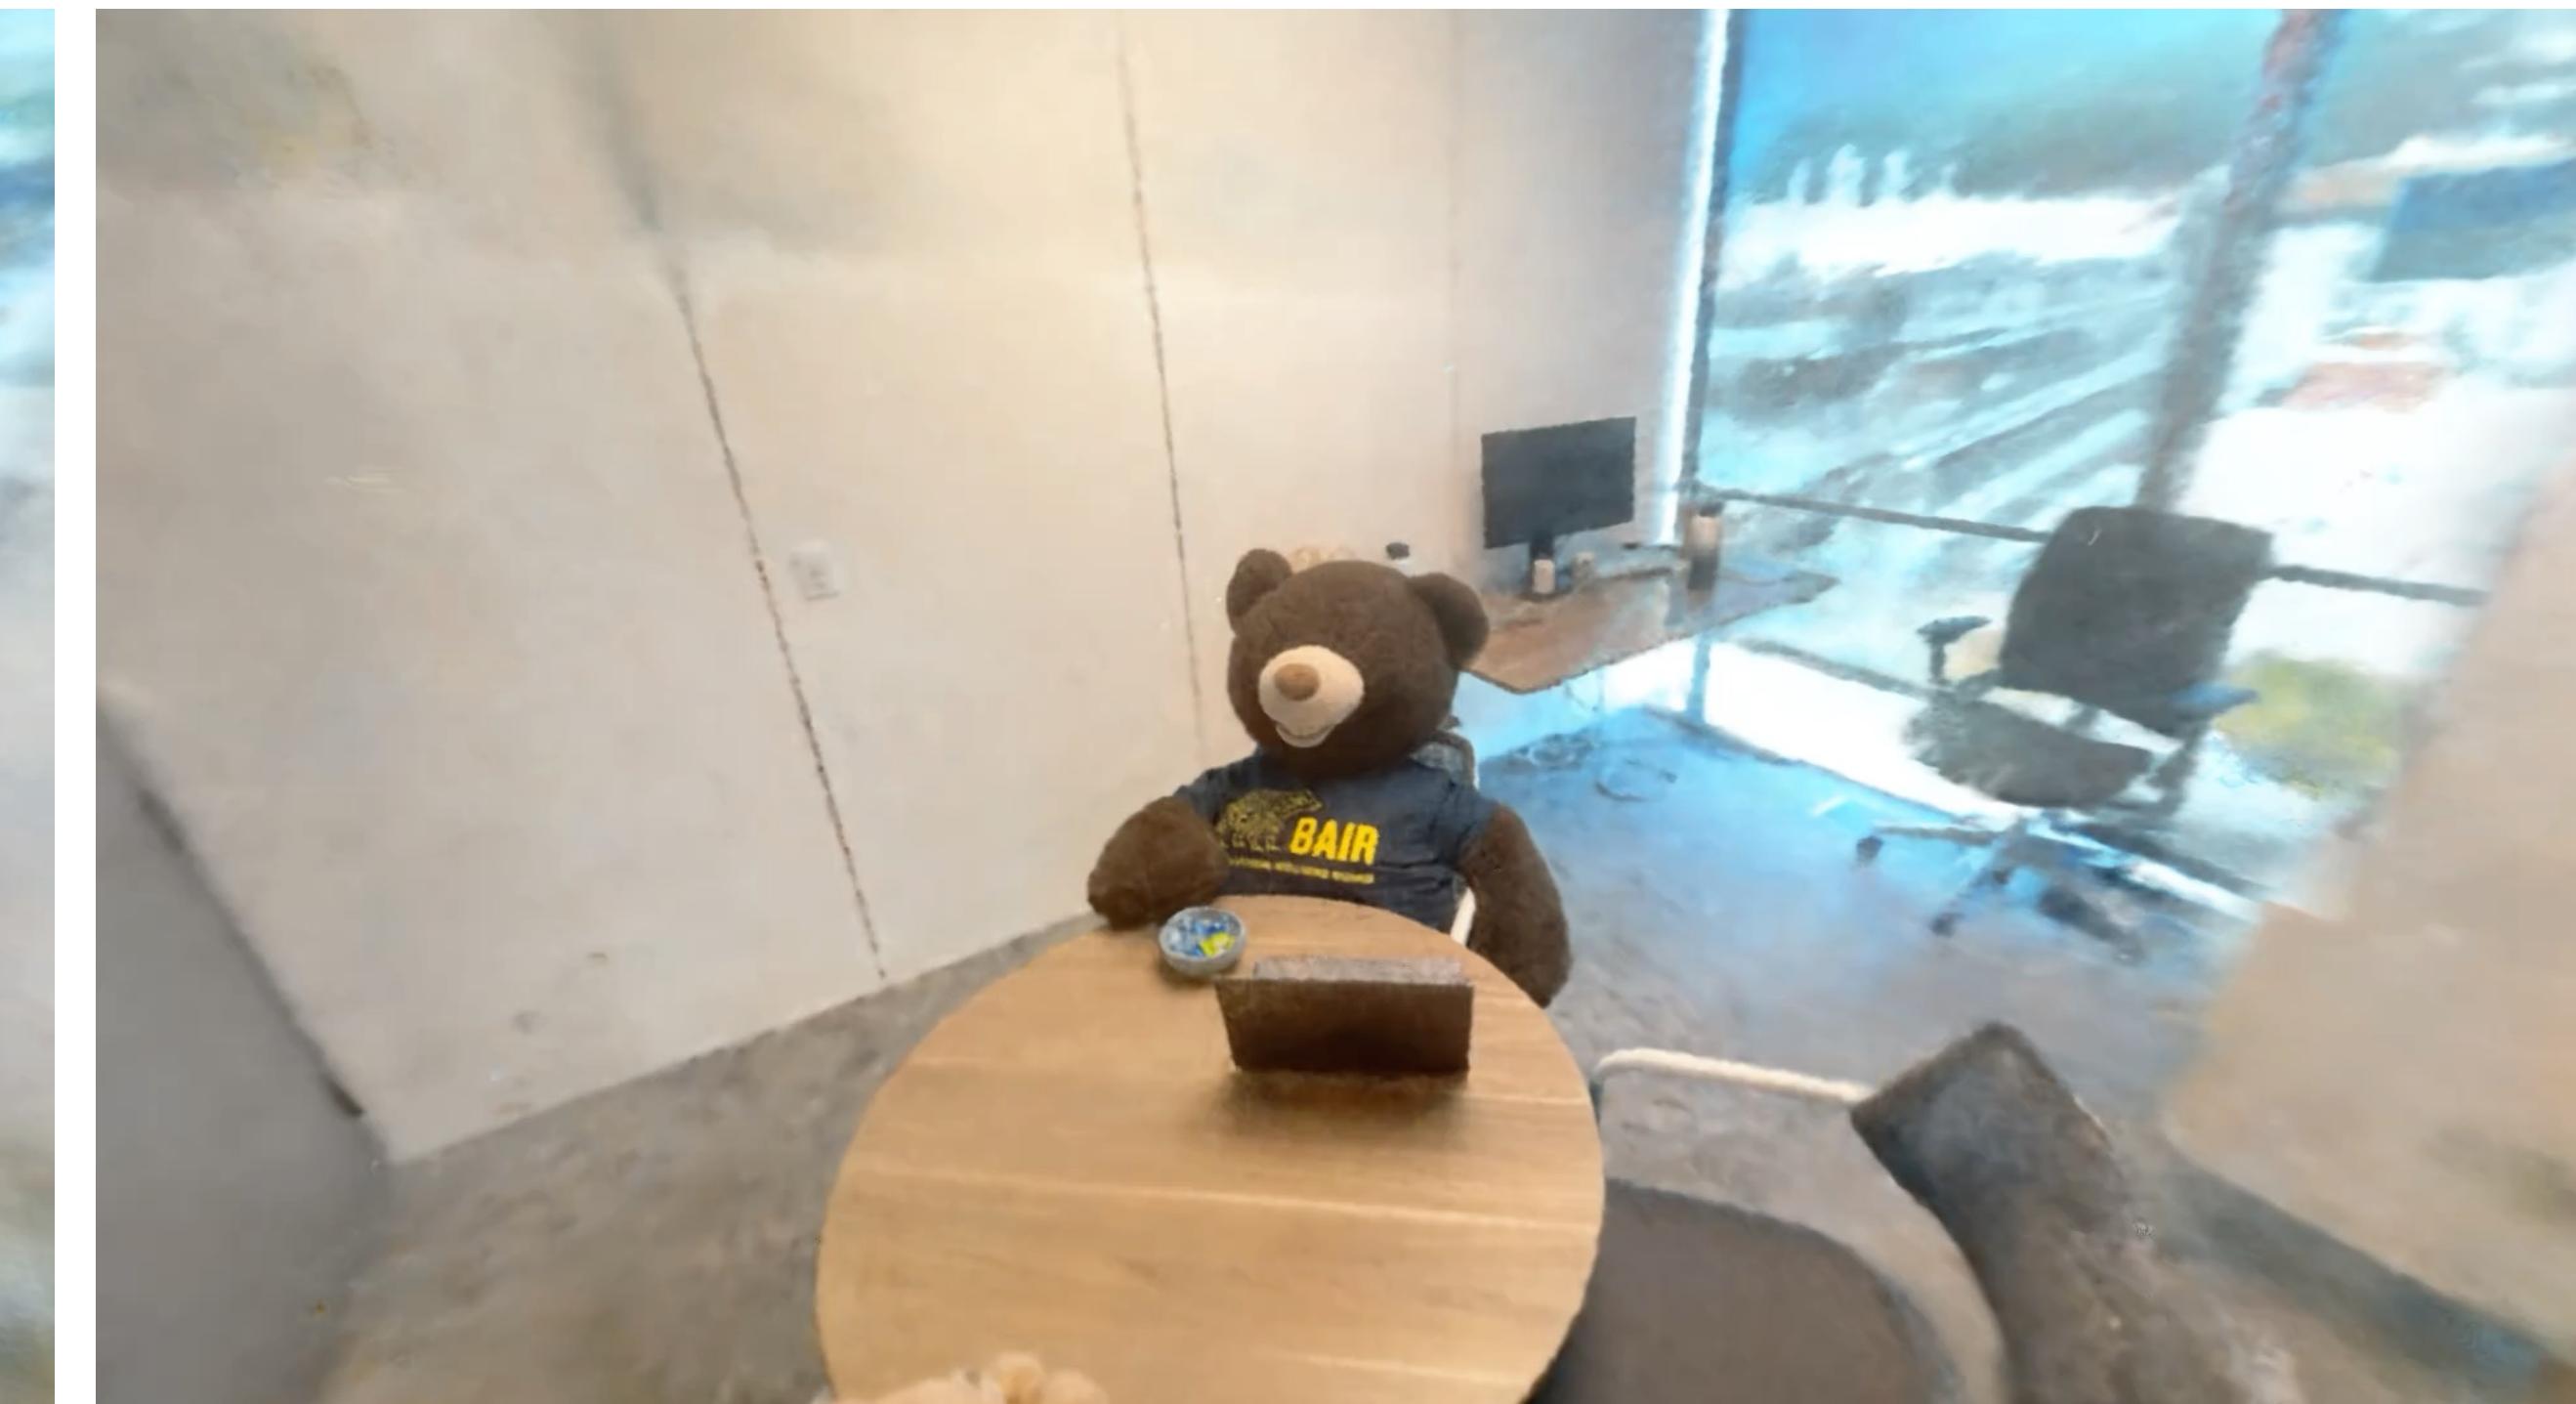
Small noise in the results can be improved

Starting from scratch is still an active area of research [Barf Lin et al. 2021, NeRF— ... ]

Noisy Camera from IMU/Lidar



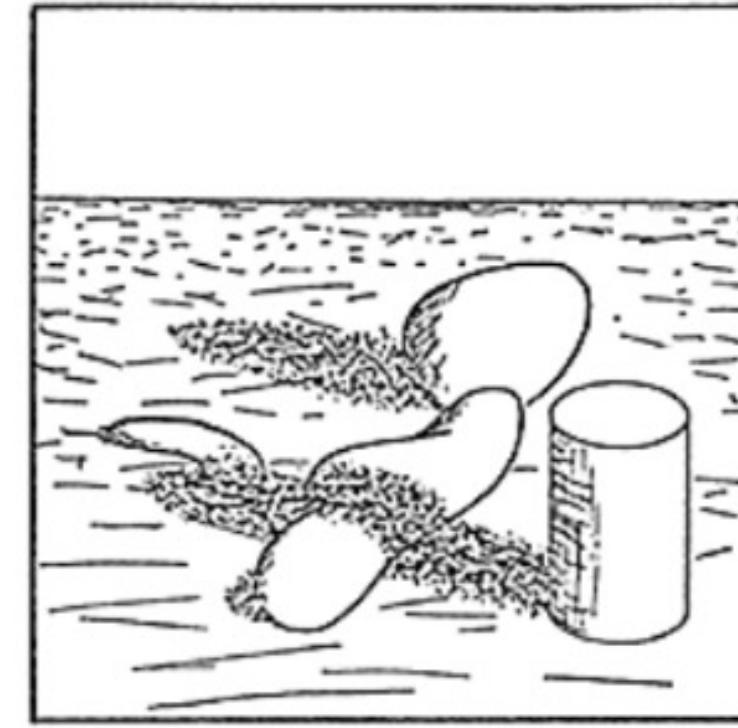
Result with Camera Optimization



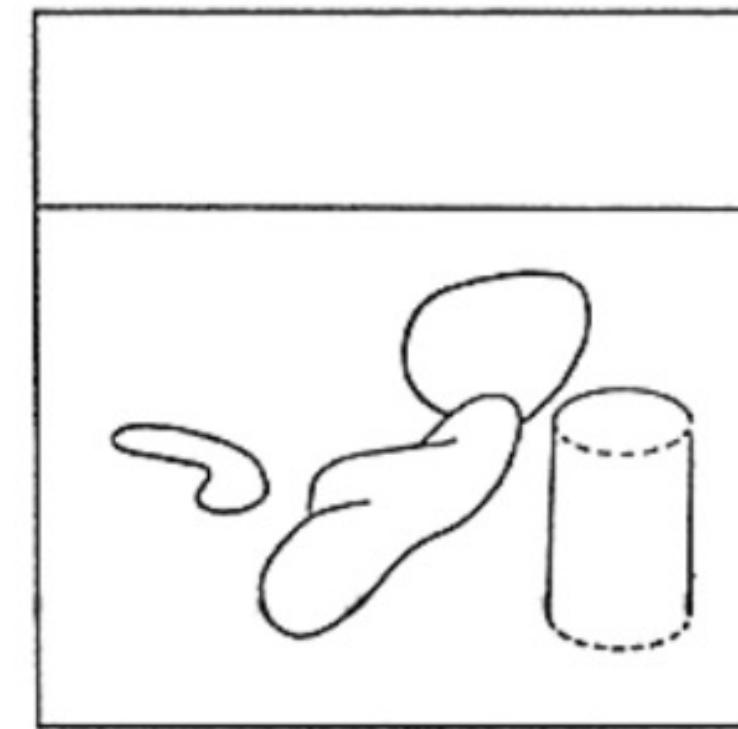
# The neverending list of NeRF limitations (back in 2020)

- ~~Expensive / slow to train~~
- ~~Expensive / slow to render~~
- ~~Sensitive to sampling strategy~~
- Sensitive to pose accuracy
- Assumes static lighting
- Not a mesh
- Assumes static scene
- Does not generalize between scenes

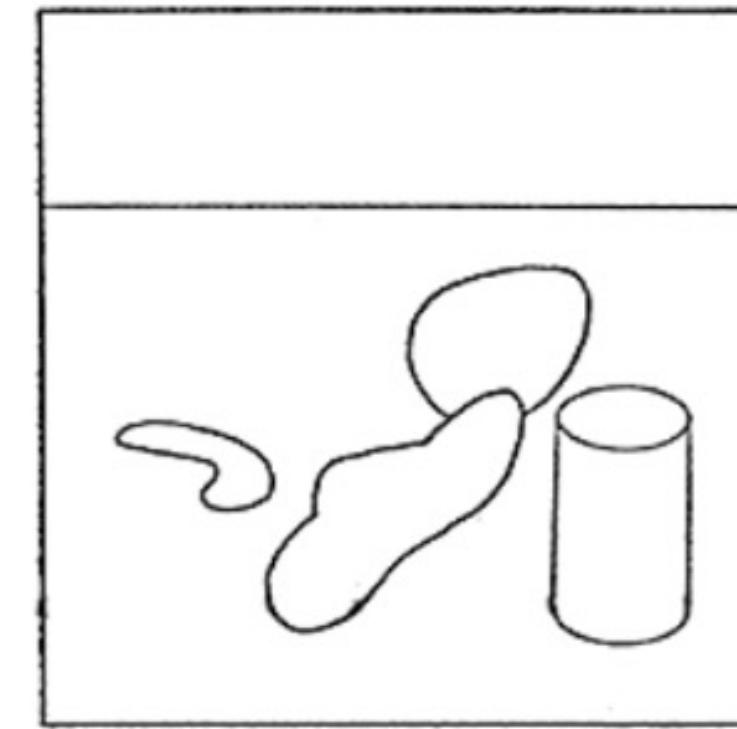
# Inverse Gi



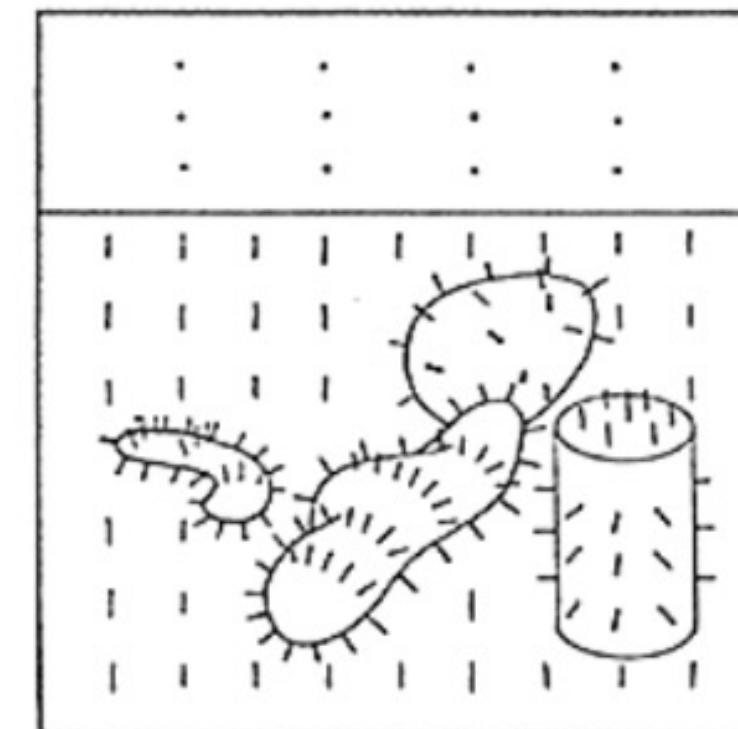
(a) ORIGINAL SCENE



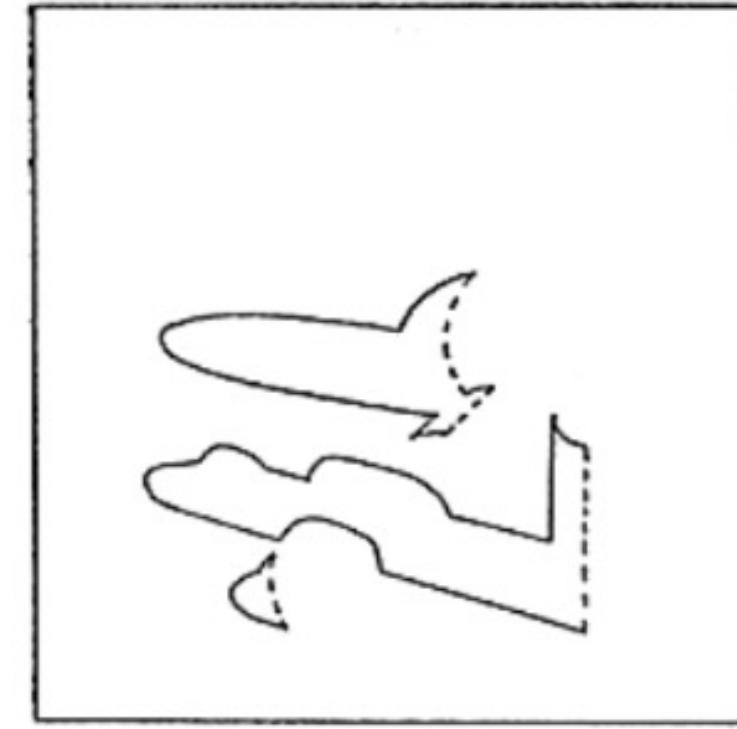
(b) DISTANCE



(c) REFLECTANCE



(d) ORIENTATION (VECTOR)



(e) ILLUMINATION

Figure 3 A set of intrinsic images derived from a single monochrome intensity image. The images are depicted as line drawings, but, in fact, would contain values at every point. The solid lines in the intrinsic images represent discontinuities in the scene characteristic; the dashed lines represent discontinuities in its derivative.

# Inverse Graphics

Some physical world  
created this image.

What was it?

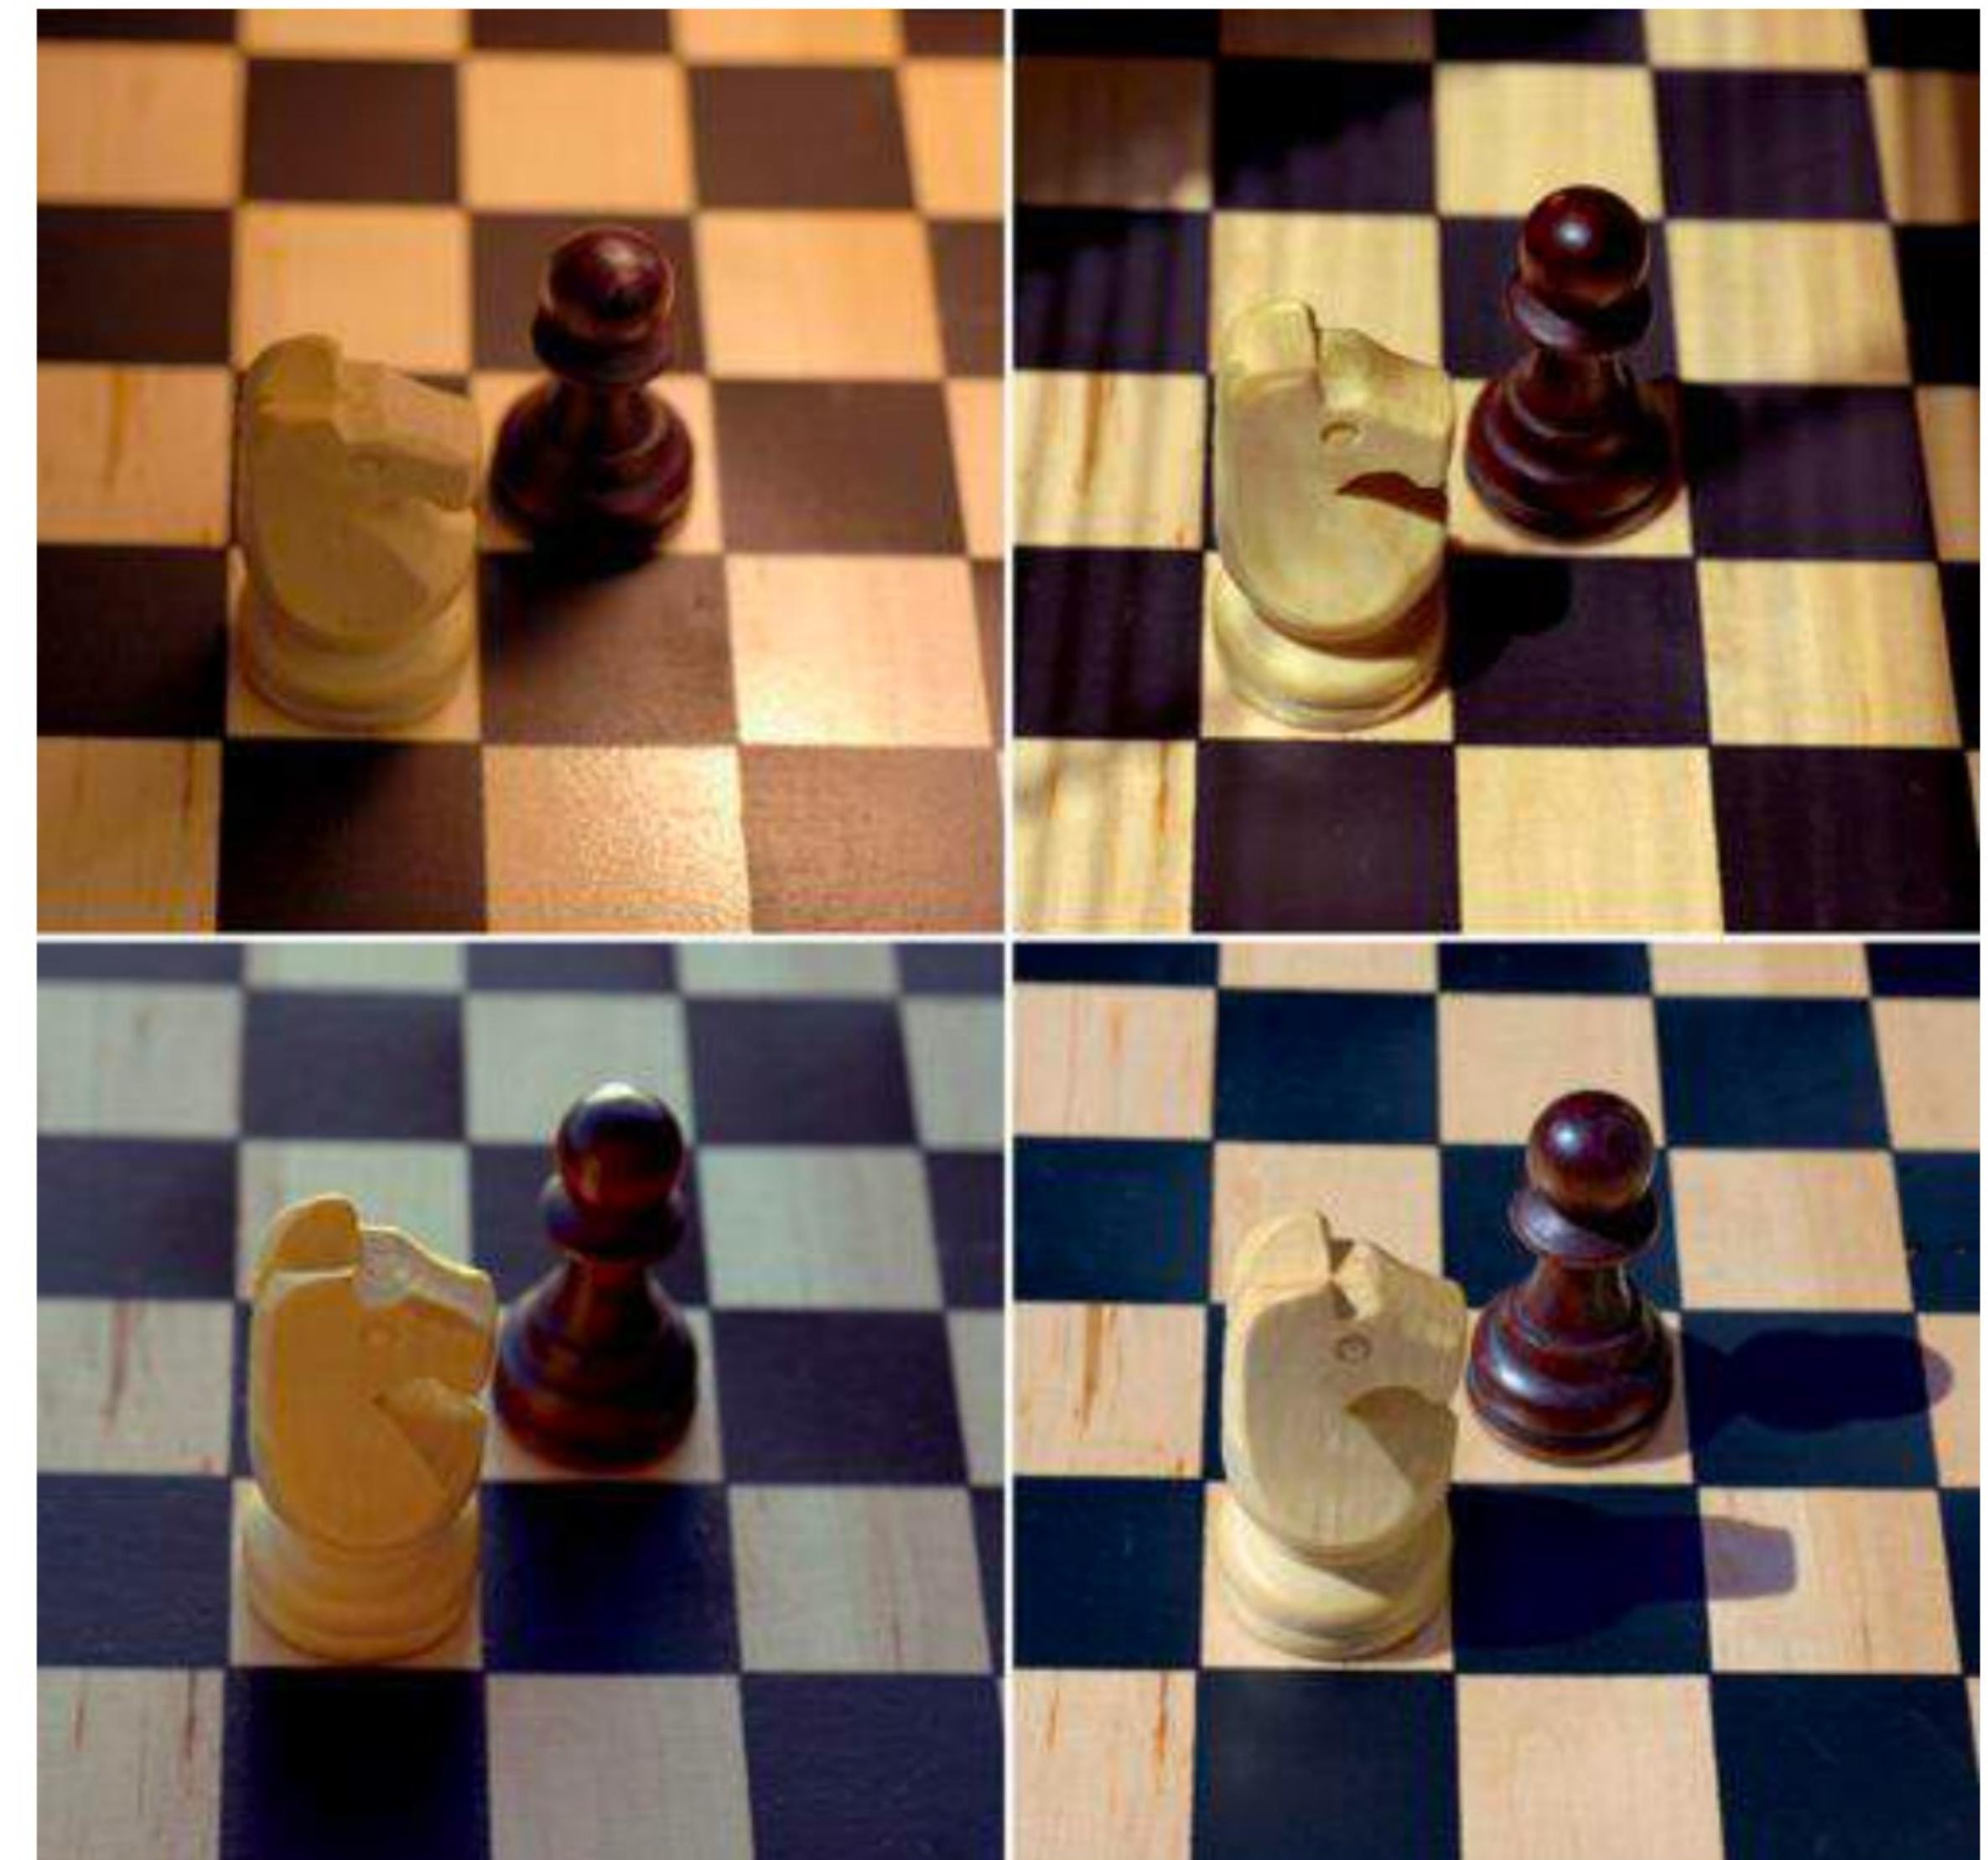


# Problem with Baked Lighting

- As you now see, NeRF bakes in the lighting effects in the scene
- That's what allows it to model the non-Lambertian effects, but it's not always ideal

# Why you want light separated

- Necessary for Relighting & Editing
  - Changing light
  - Inserting objects into another scene (with different lighting)
  - Changing material properties
  - Edit the appearance without changing light
  - ...



# Recall: we simplified by ignoring scattering

Absorption



<http://commons.wikimedia.org>

Scattering



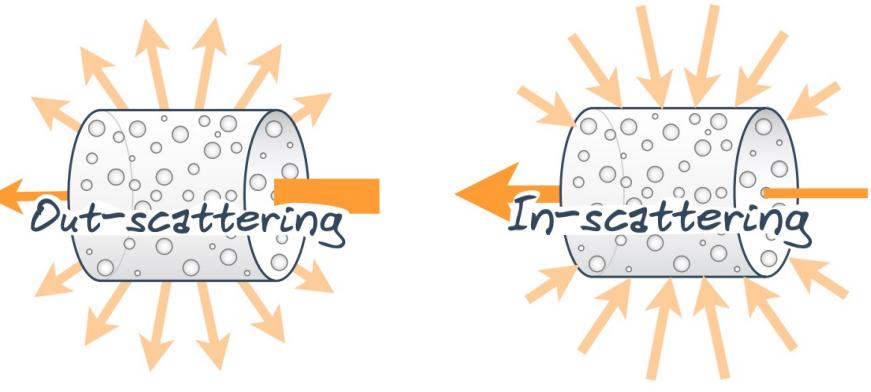
Emission



<http://wikipedia.org>



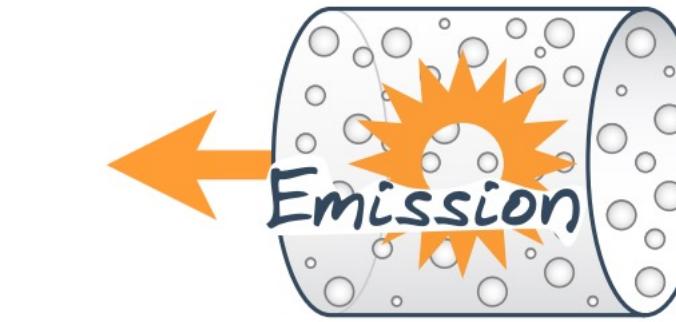
## Absorption



## Scattering



<http://commons.wikimedia.org>



## Emission



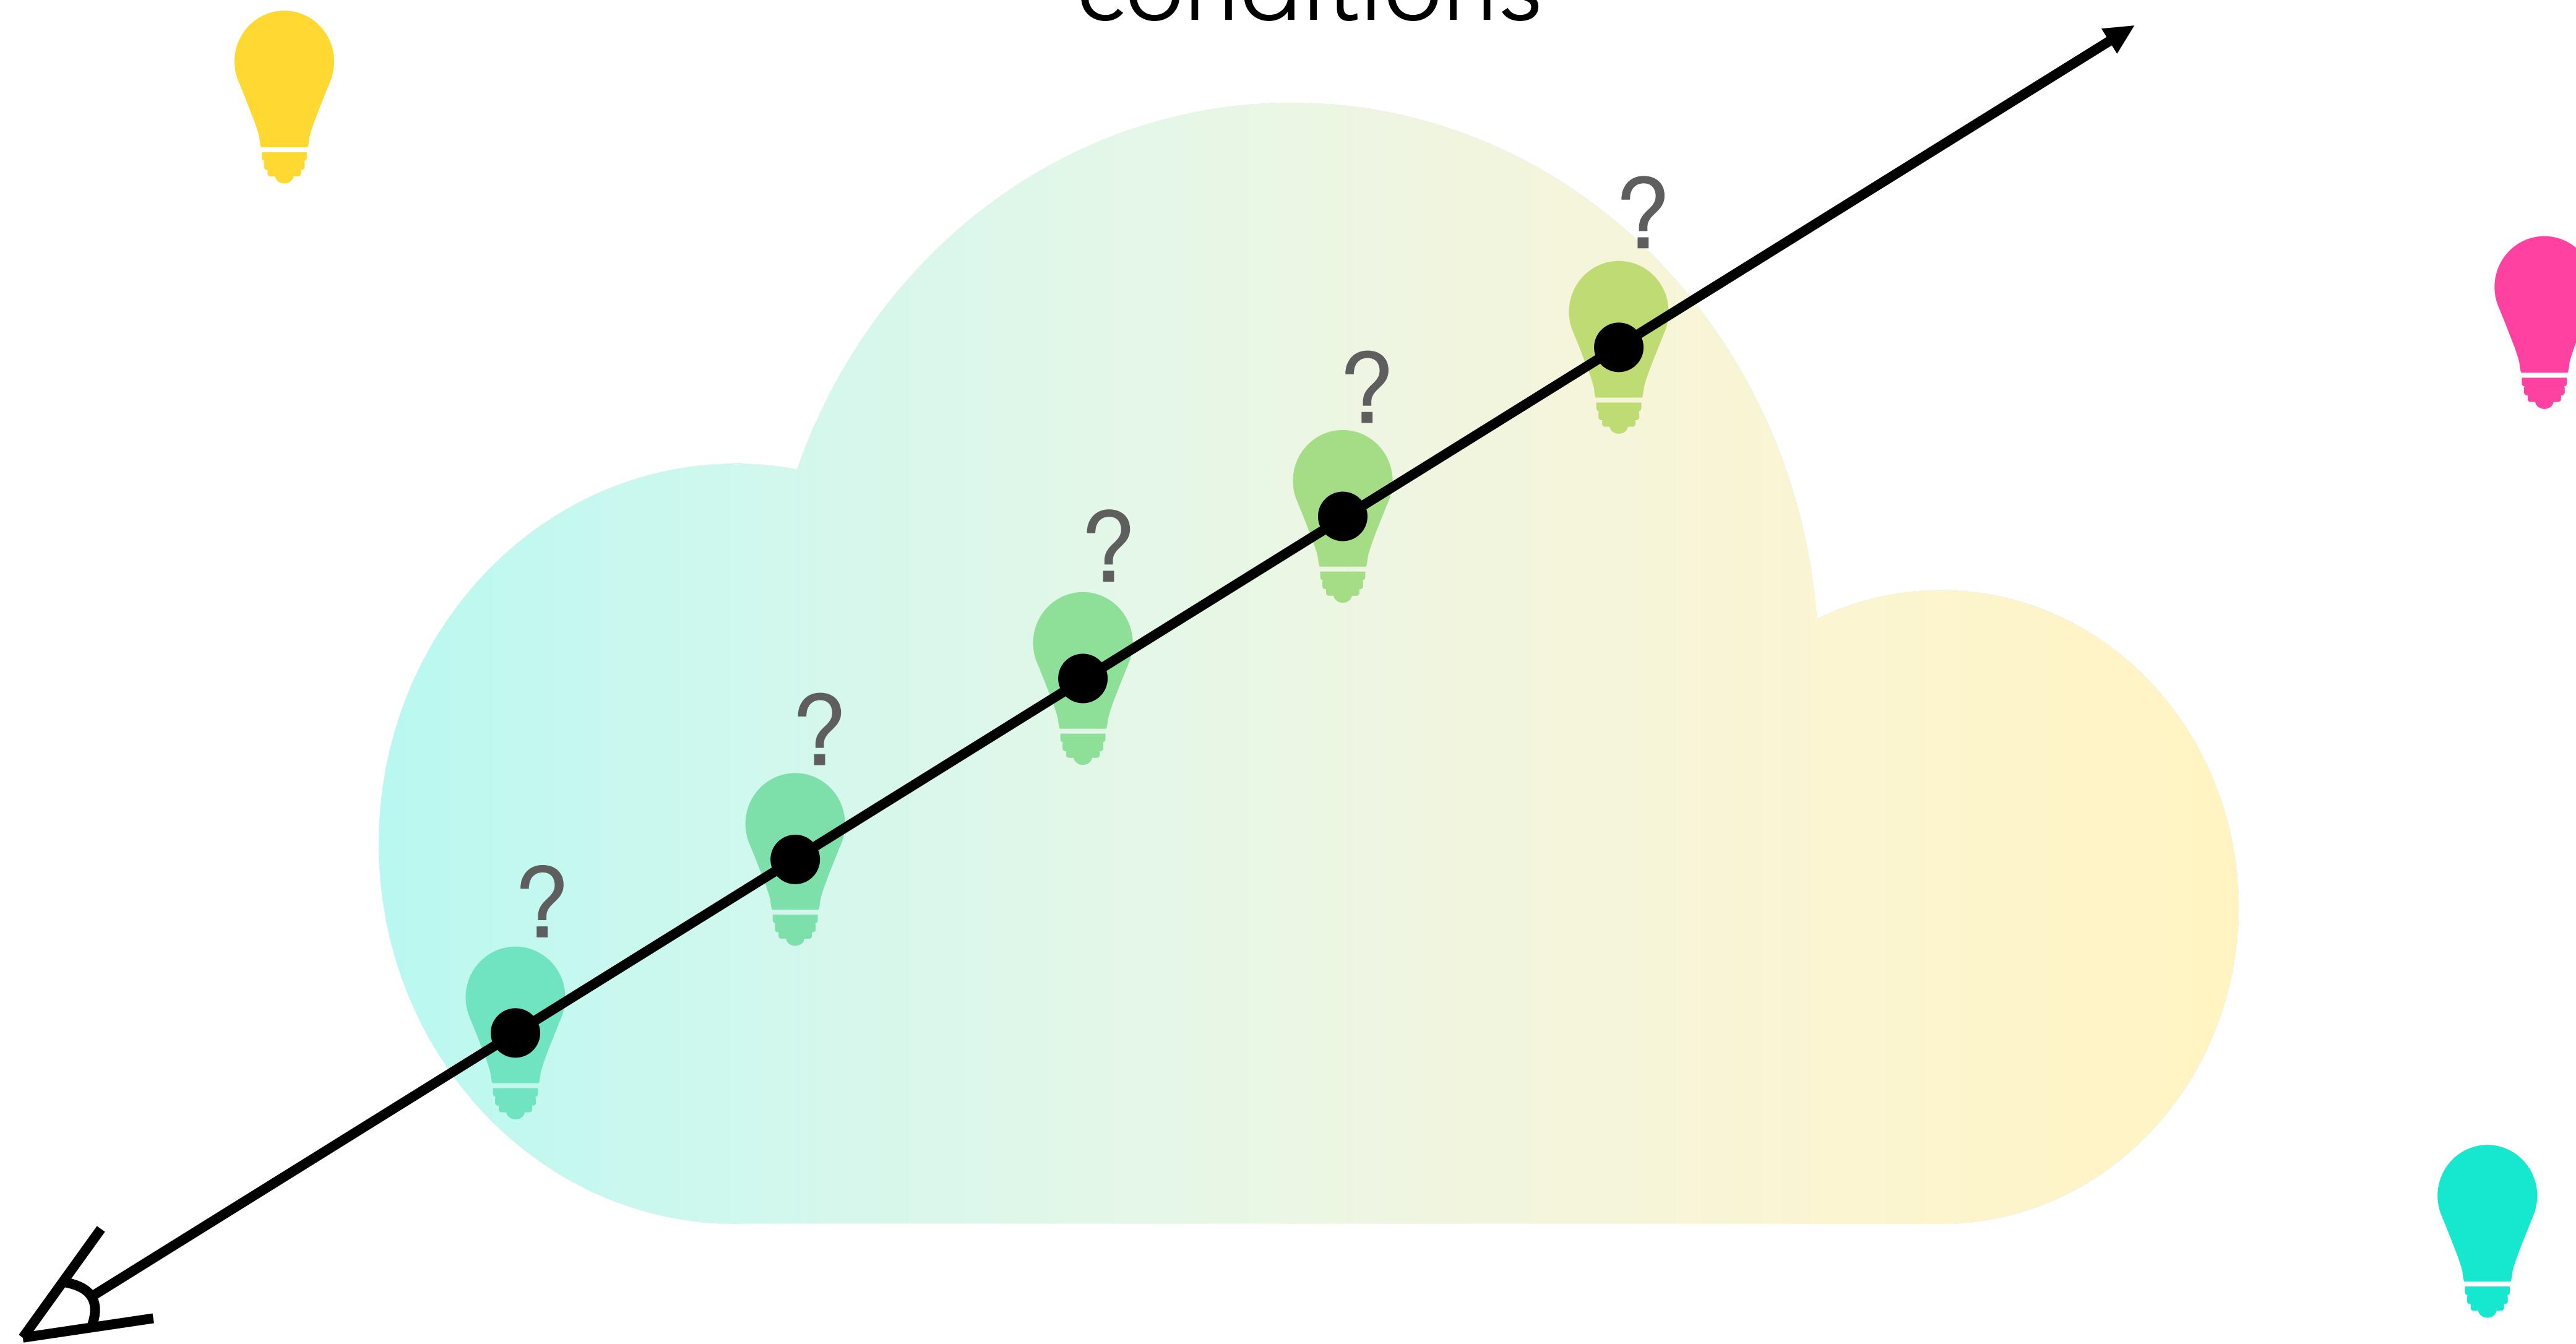
<http://wikipedia.org>

# NeRF represents a volume of particles that emit light

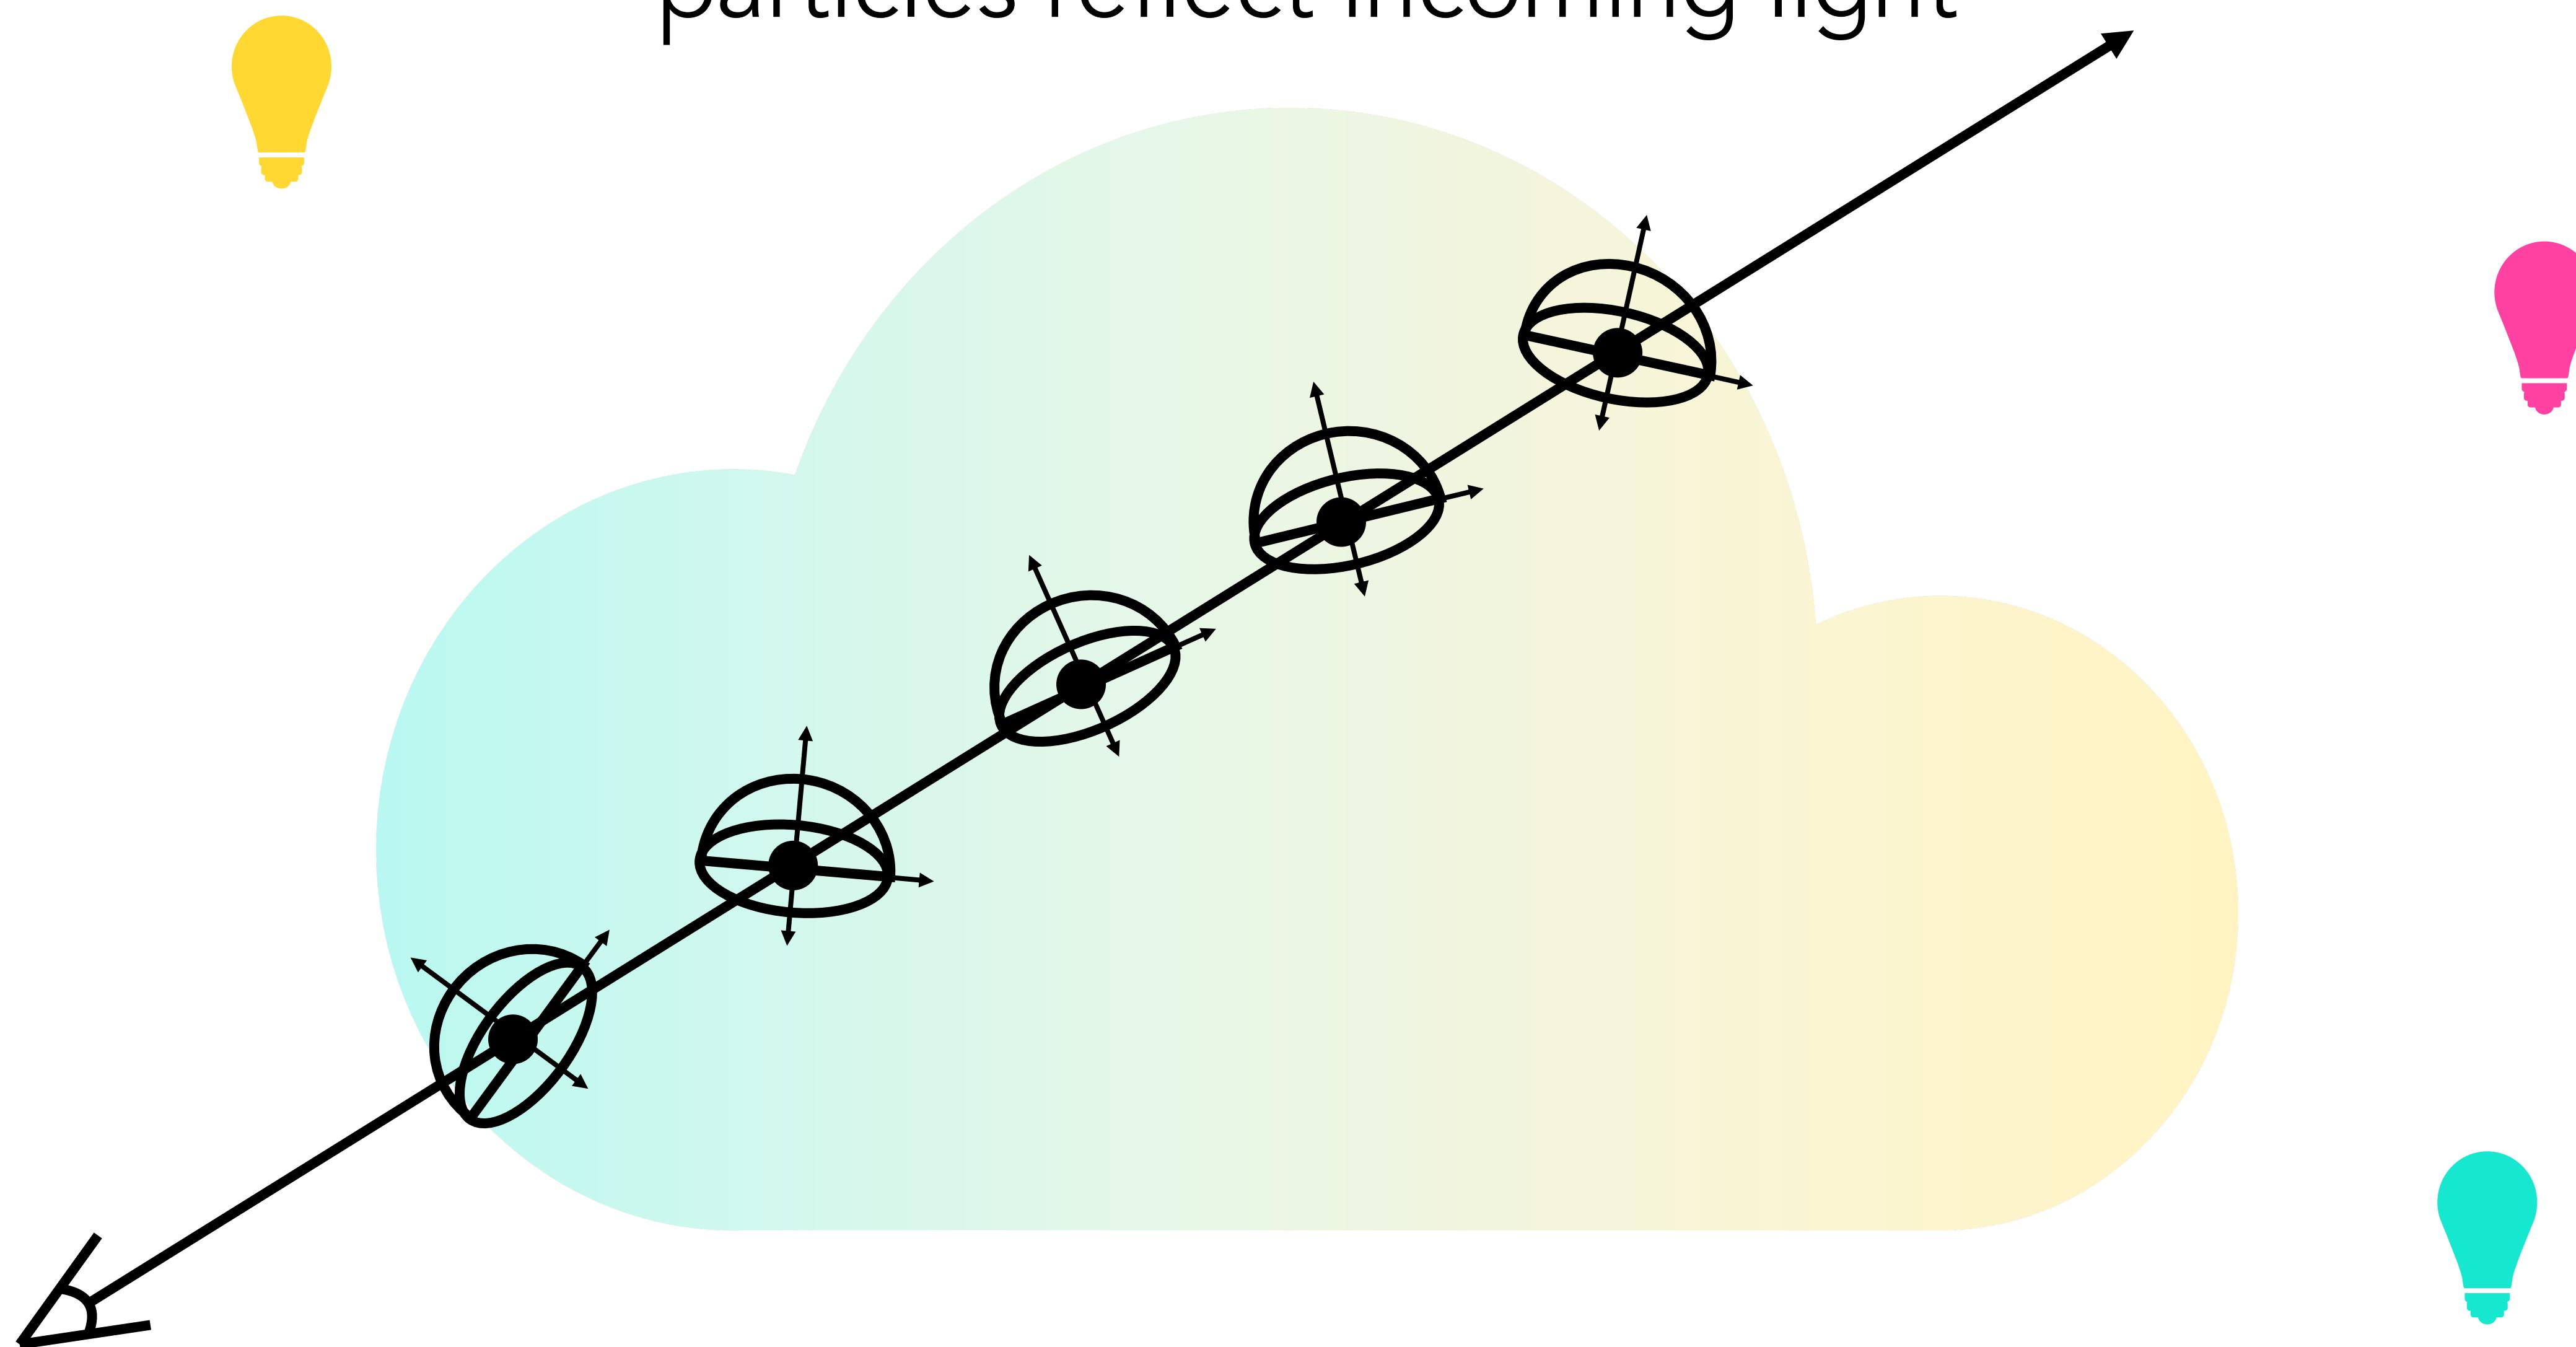


NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis.  
Ben Mildenhall\*, Pratul Srinivasan\*, Matt Tancik\*, Jon Barron, Ravi Ramamoorthi, Ren Ng. ECCV 2020.

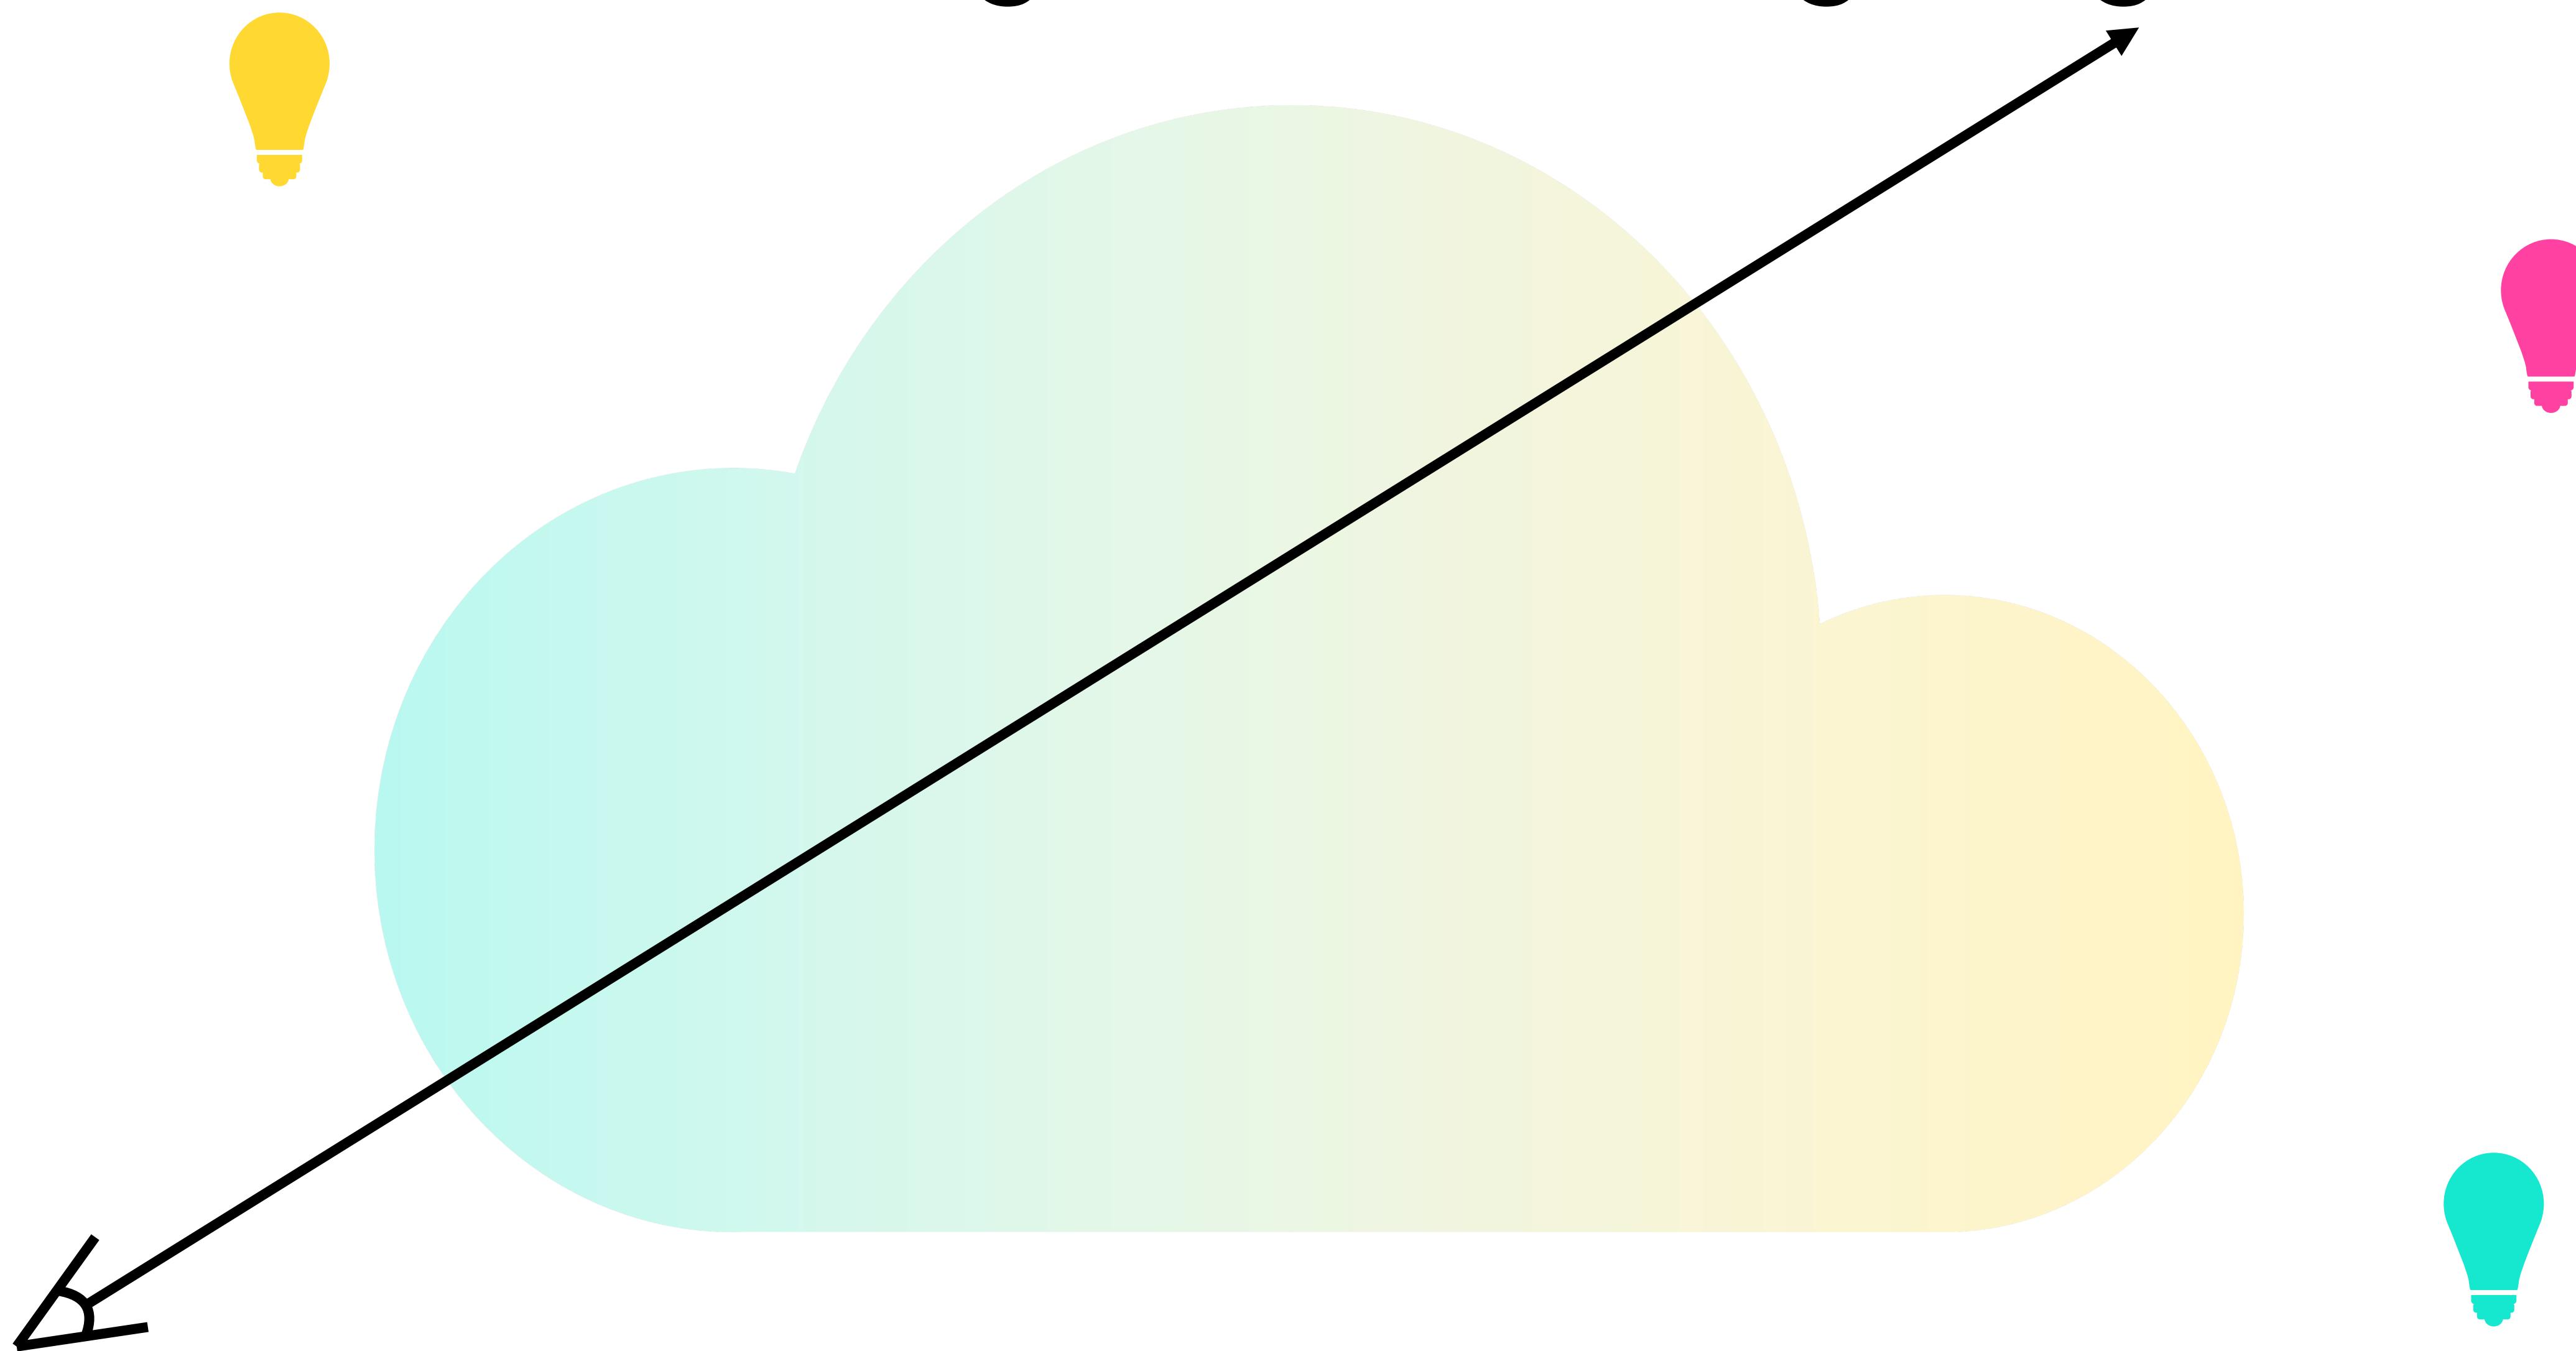
But doesn't let us simulate how light changes with new lighting conditions



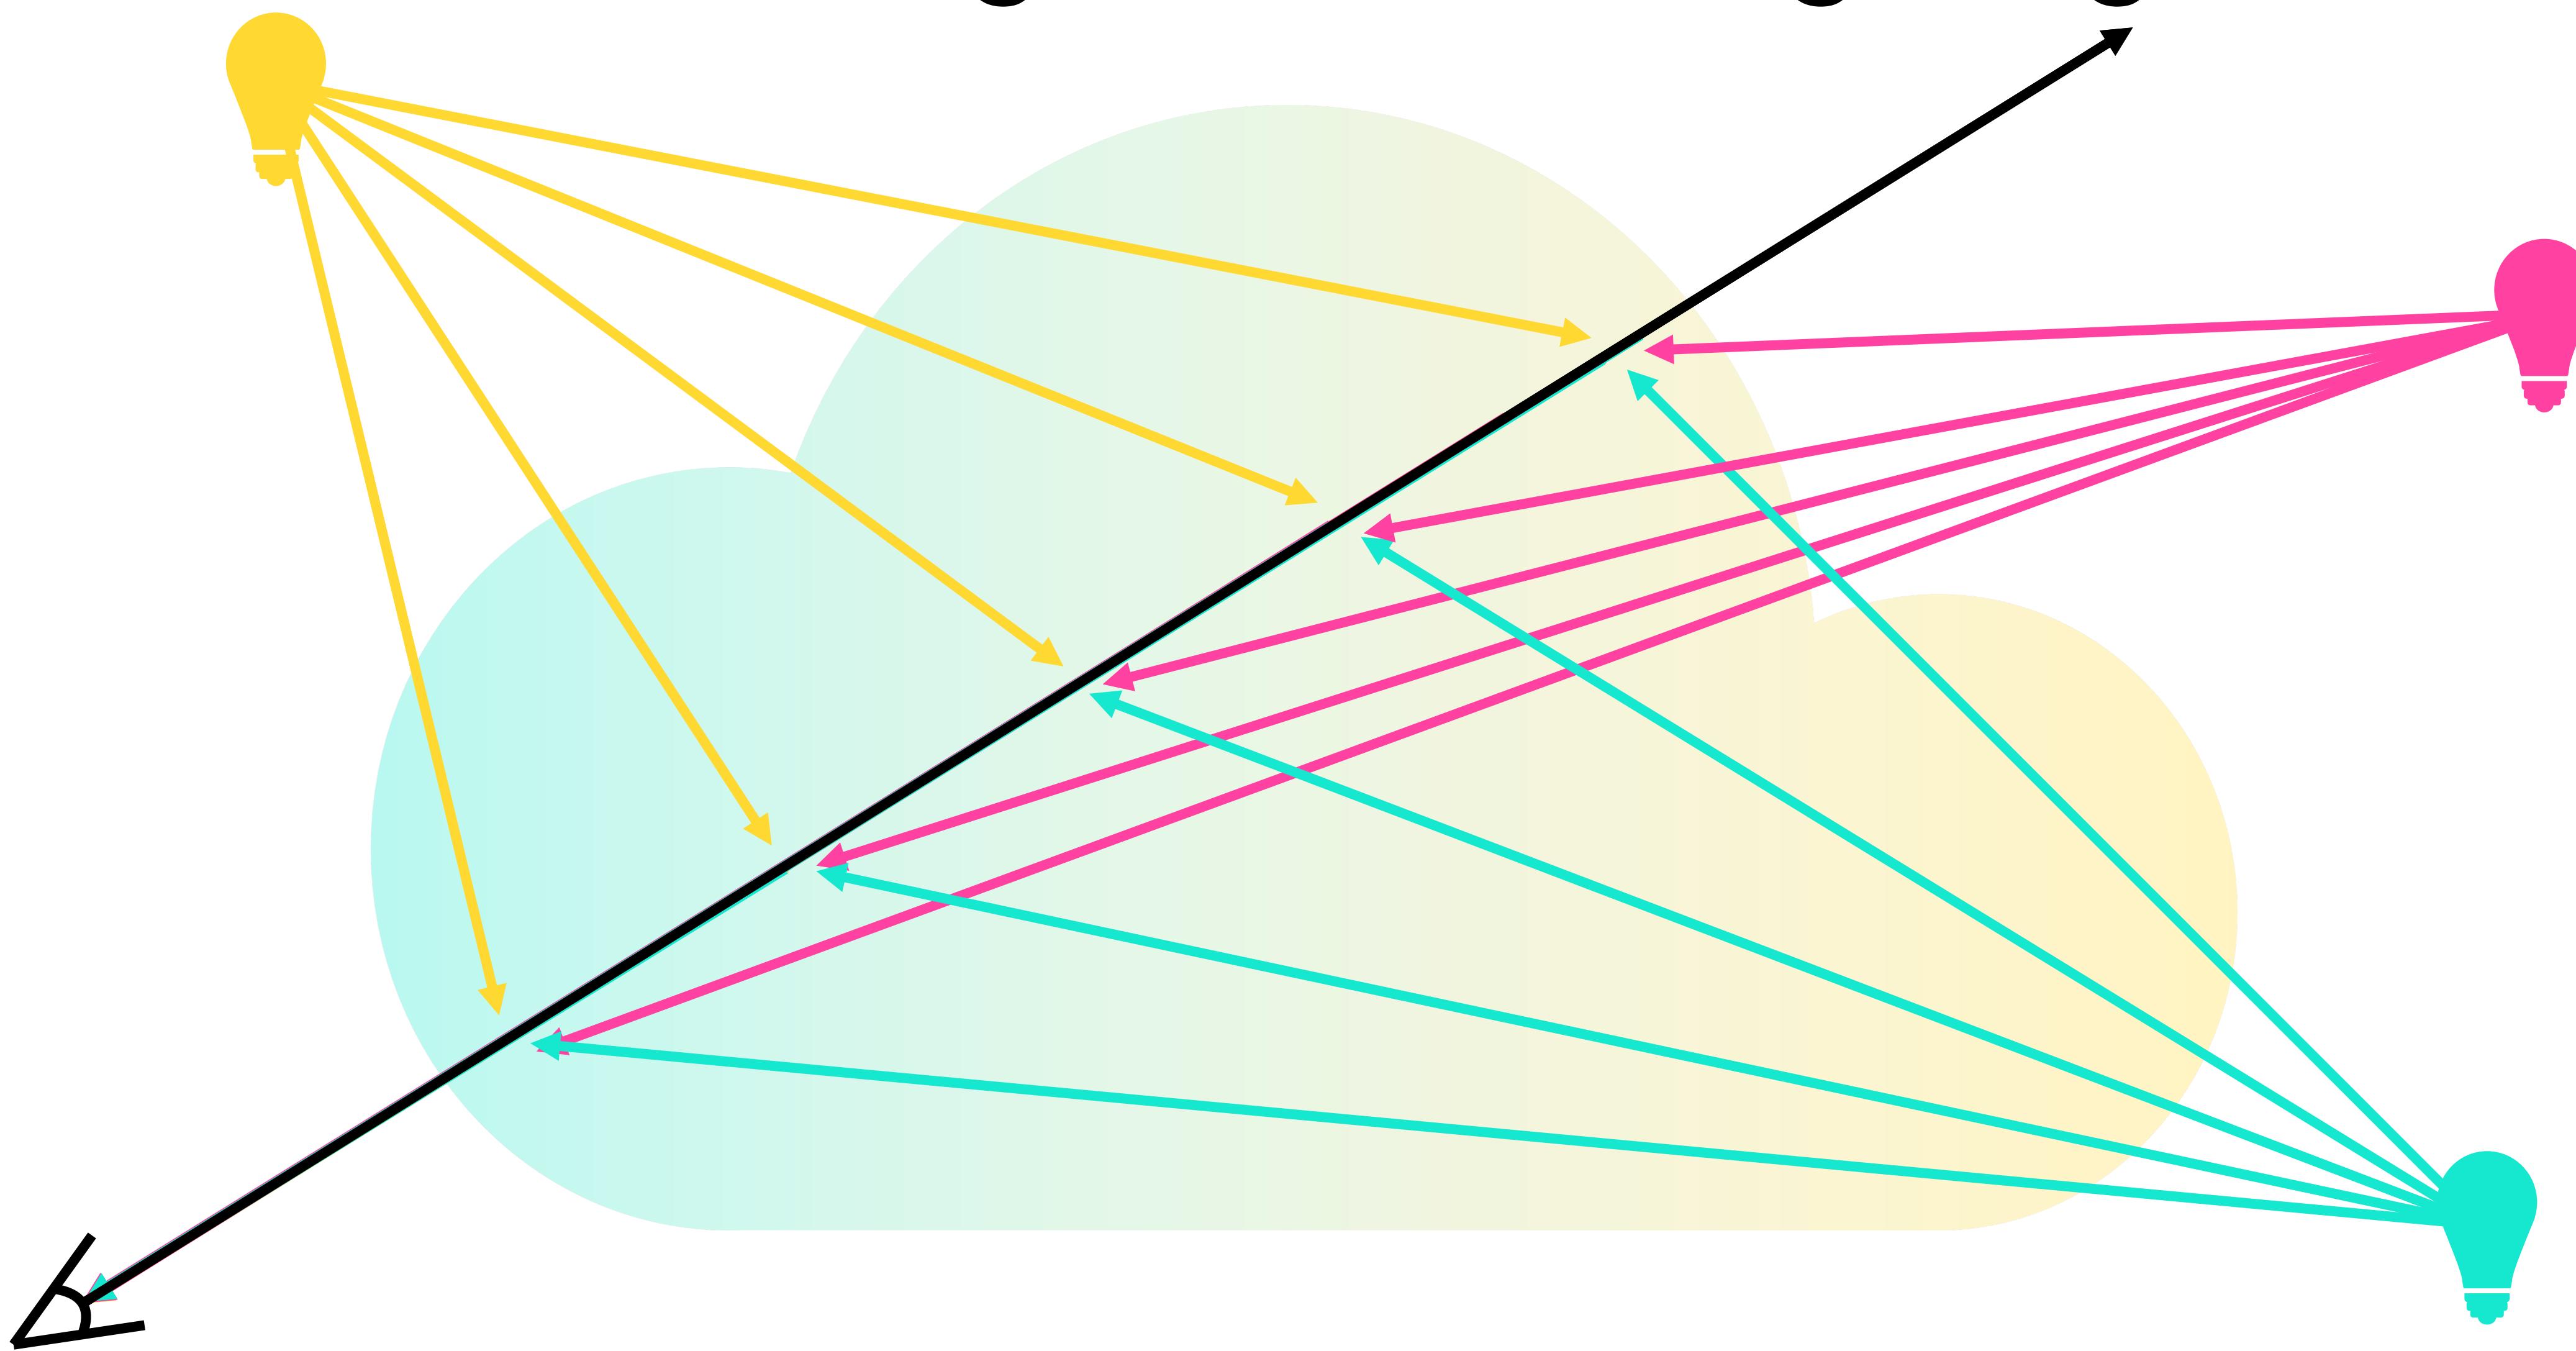
First step: replace emitted light with BRDFs that describe how particles reflect incoming light



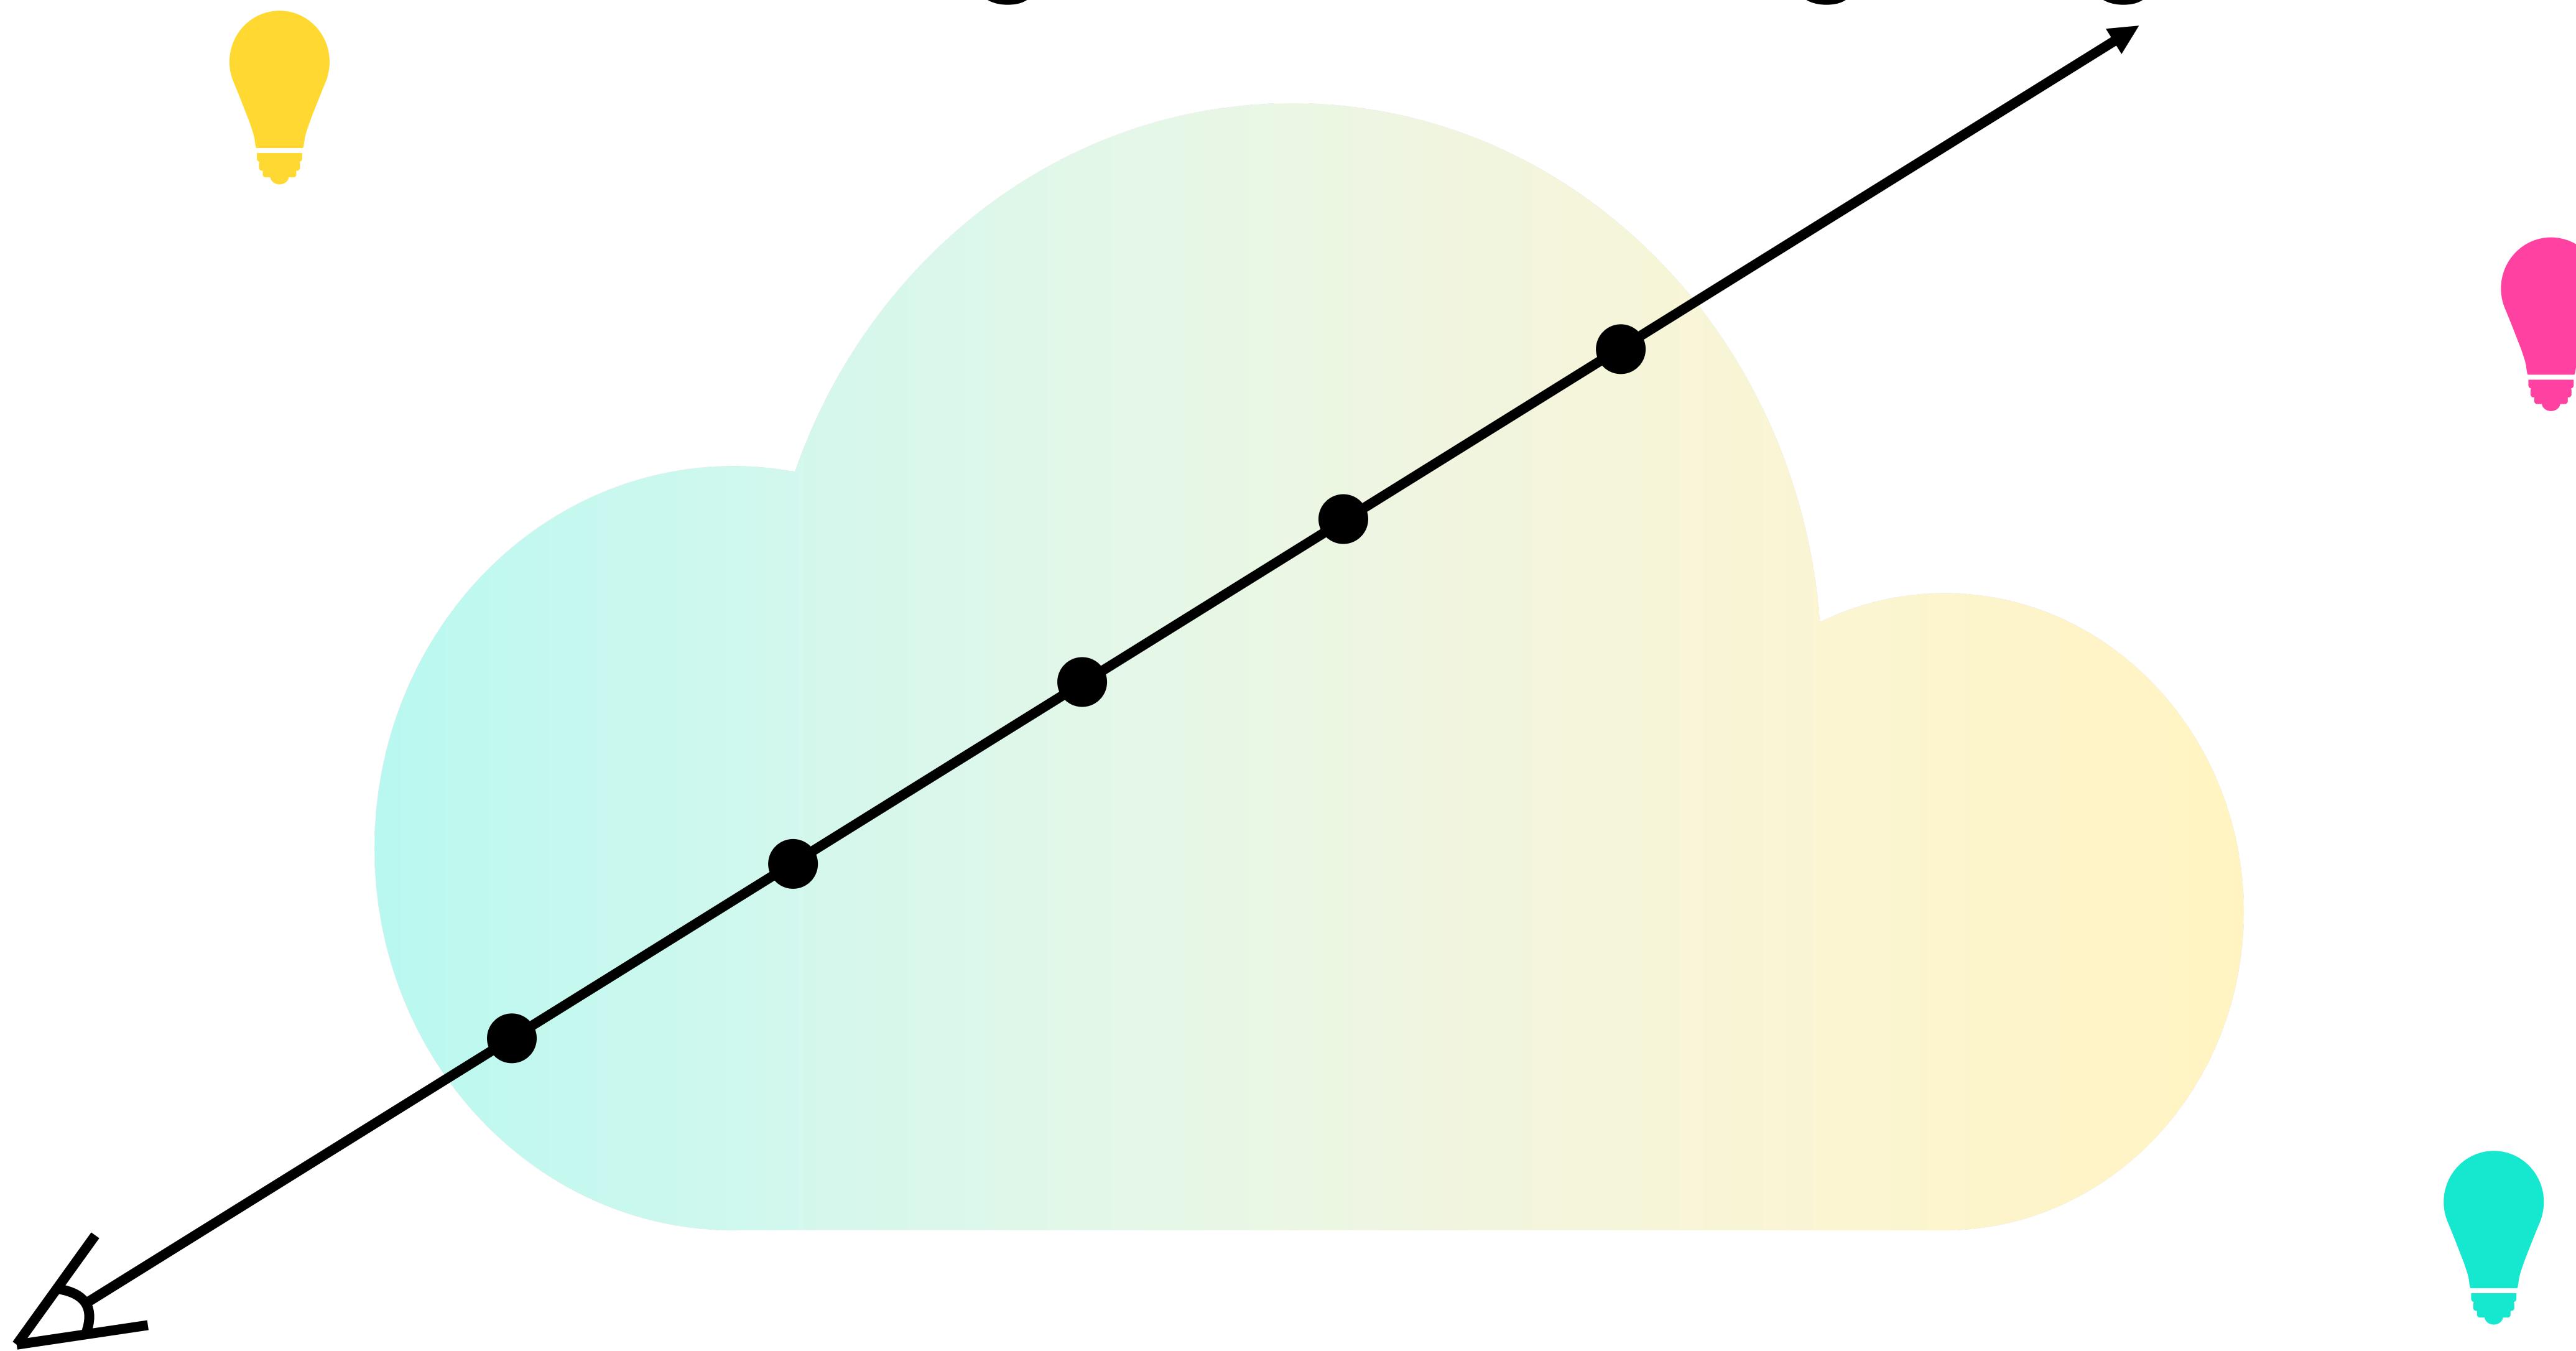
# Rendering with direct lighting



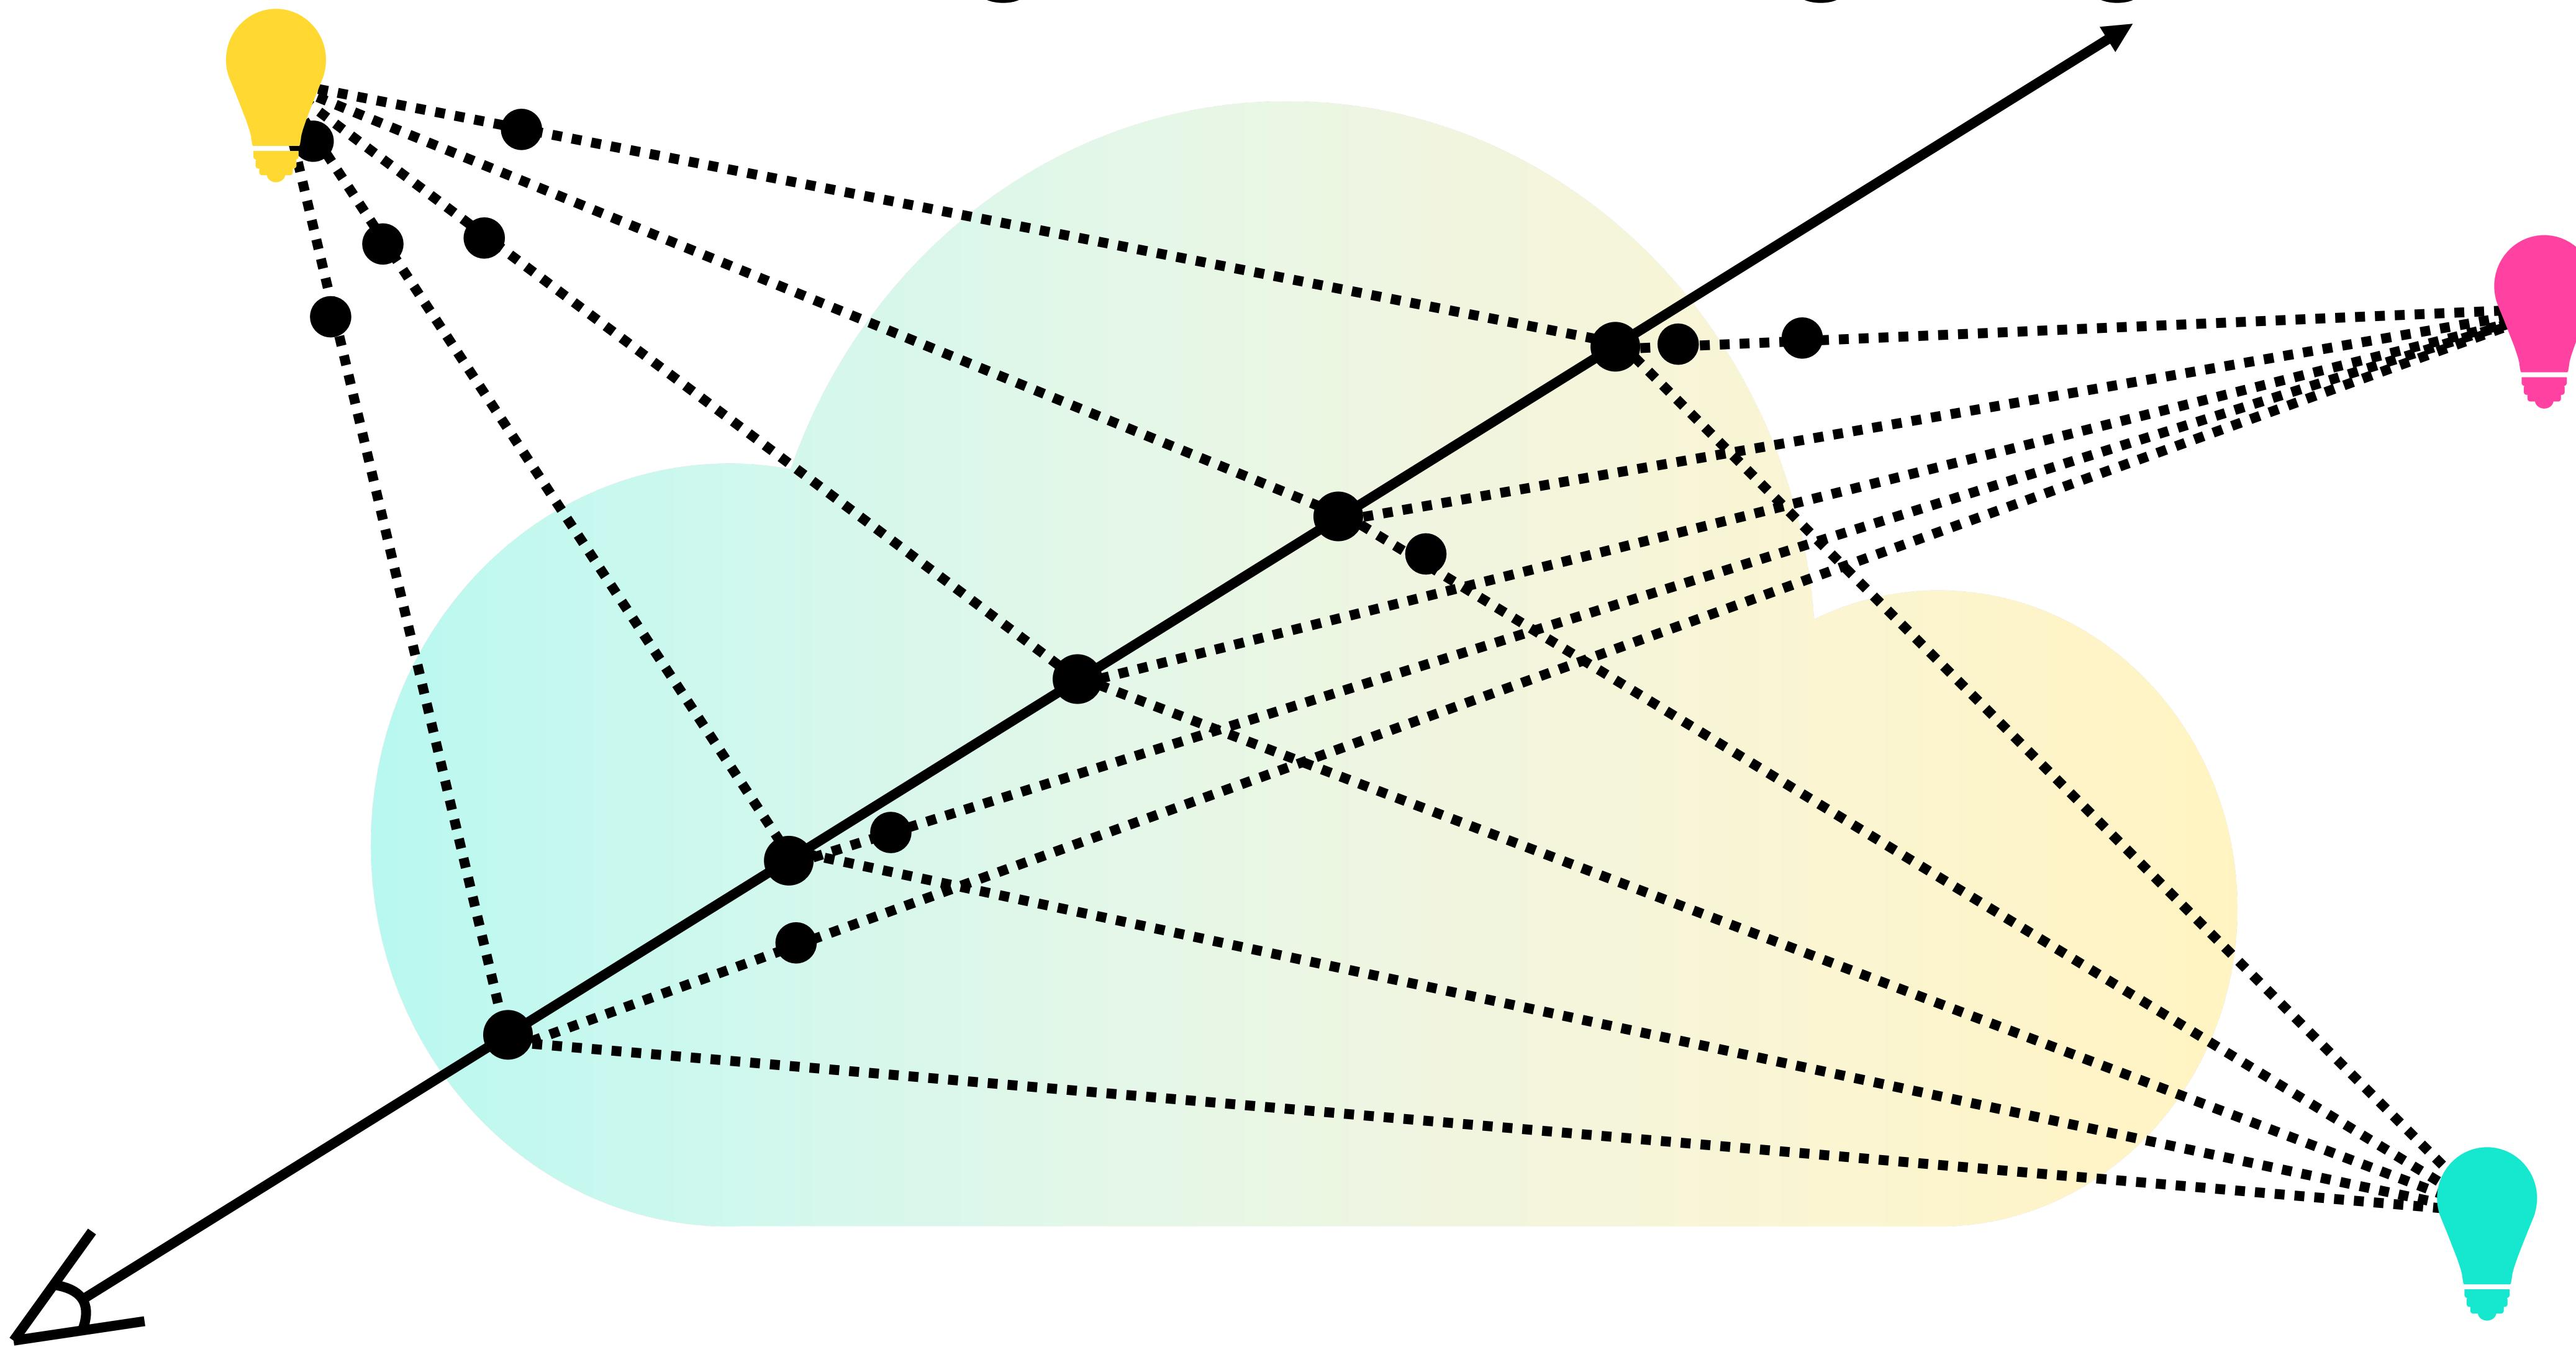
# Rendering with direct lighting



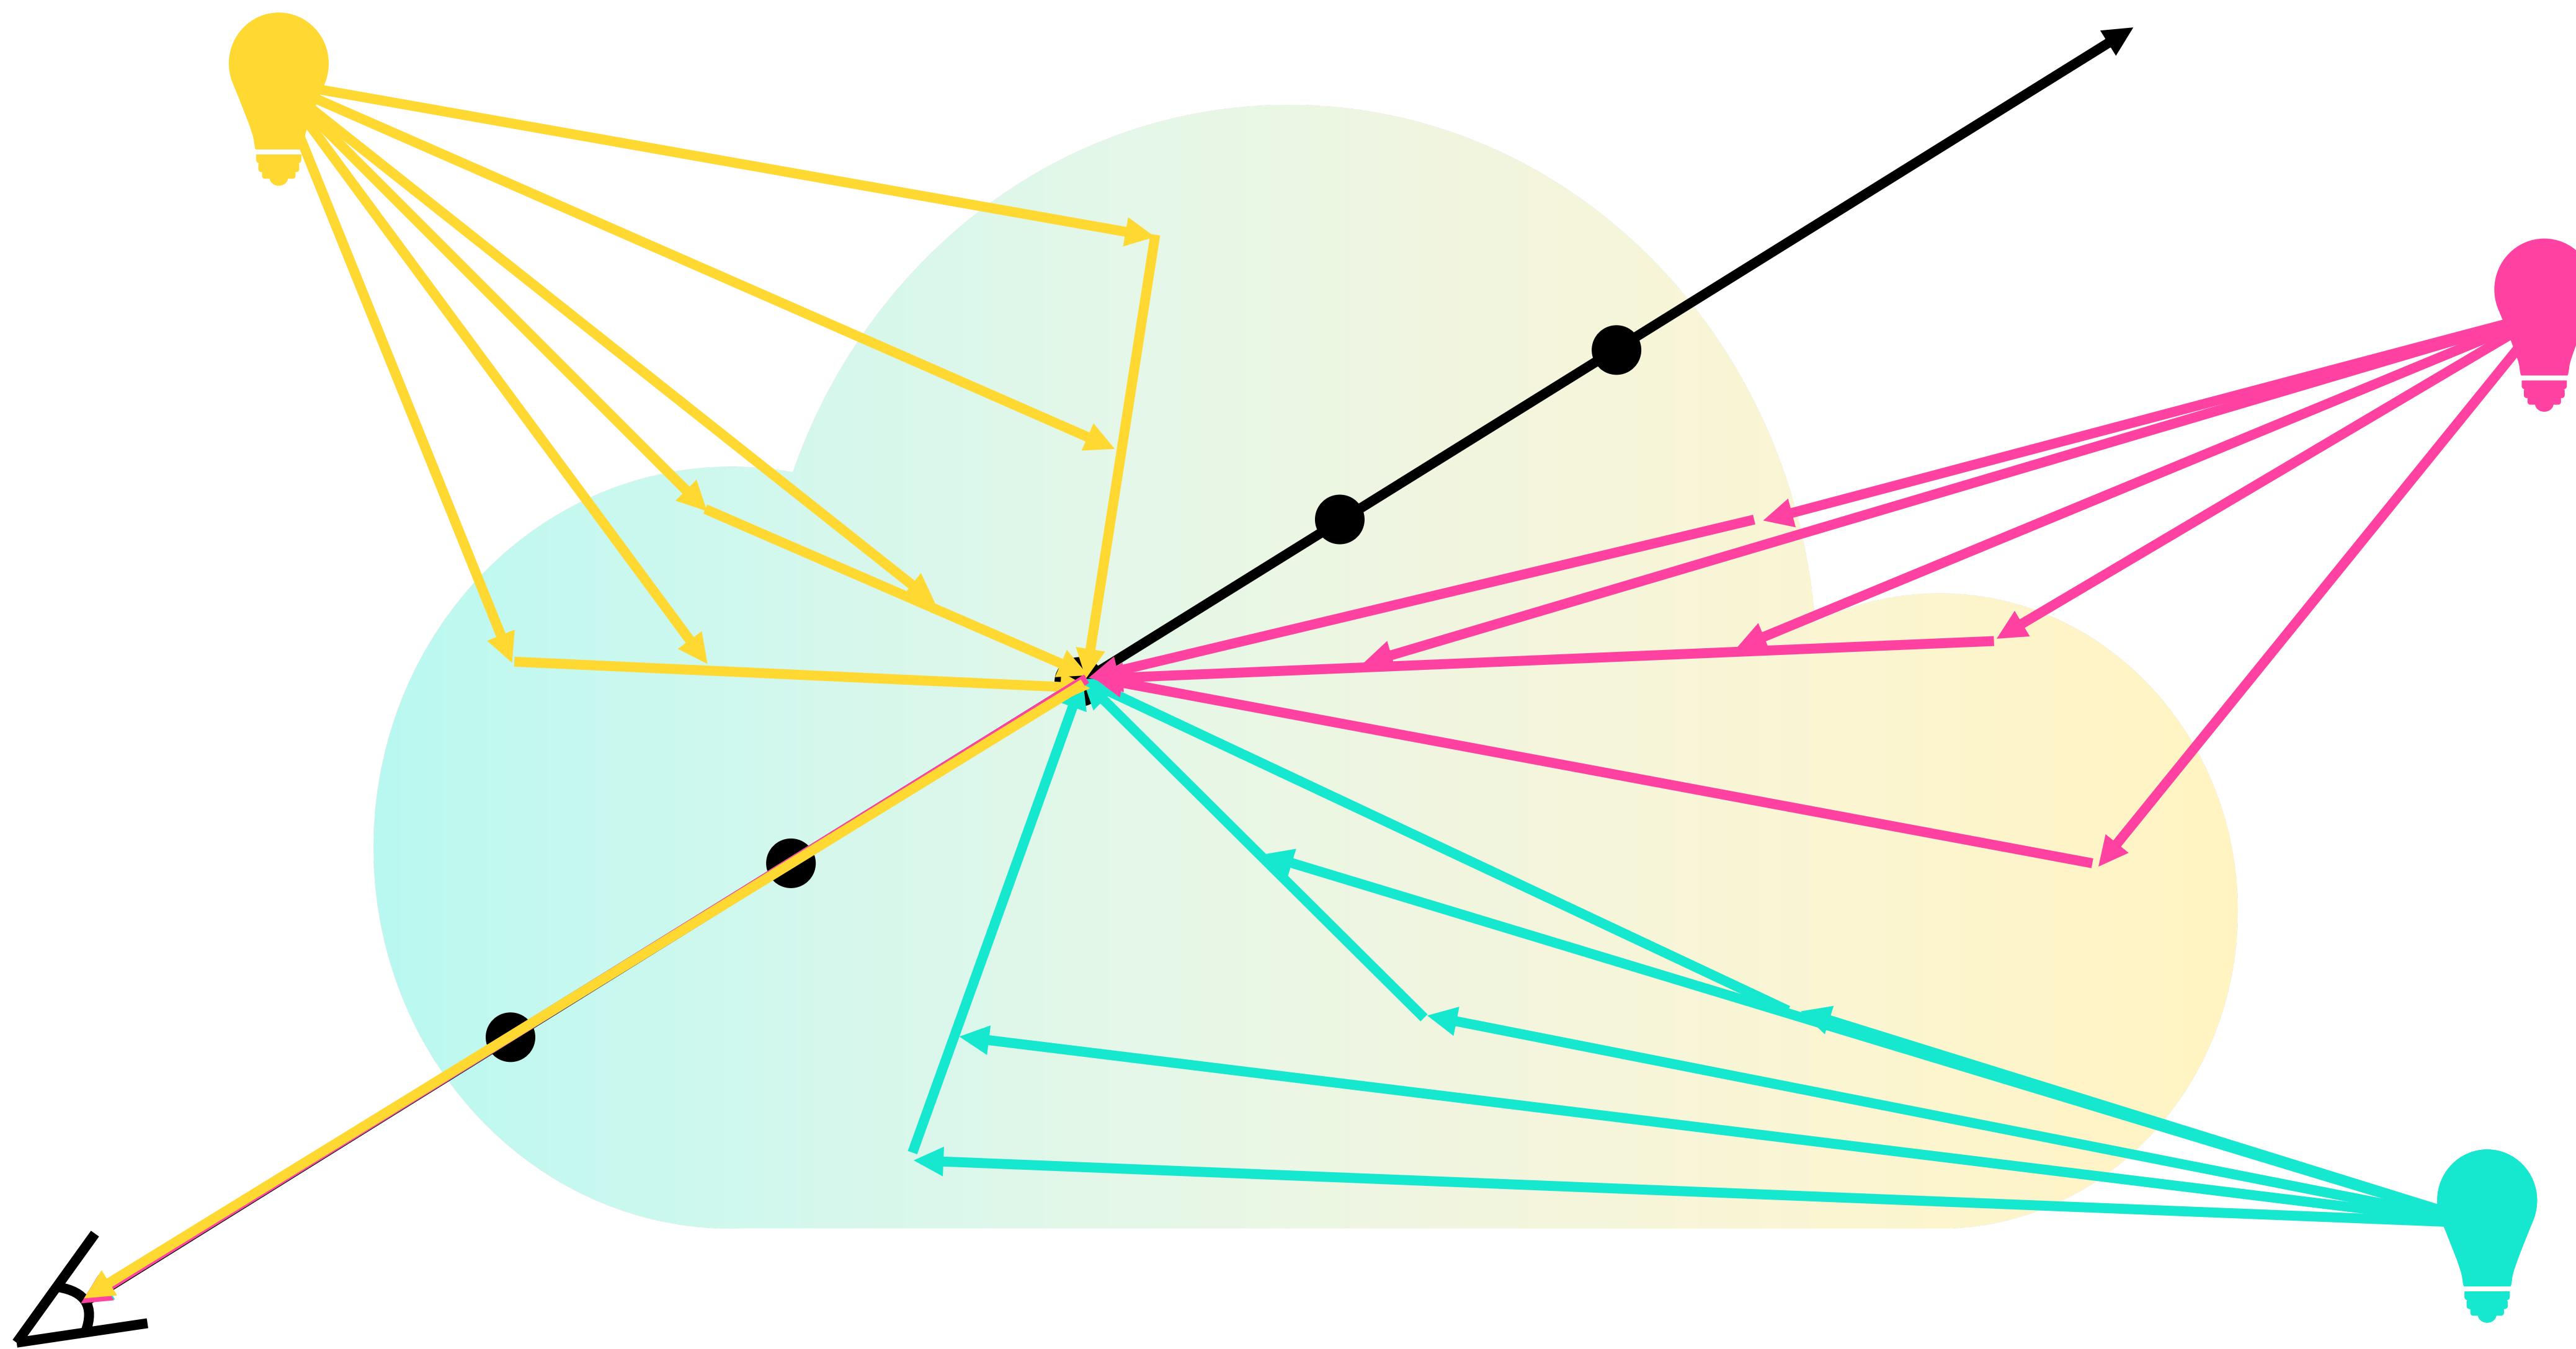
# Rendering with direct lighting



# Rendering with direct lighting

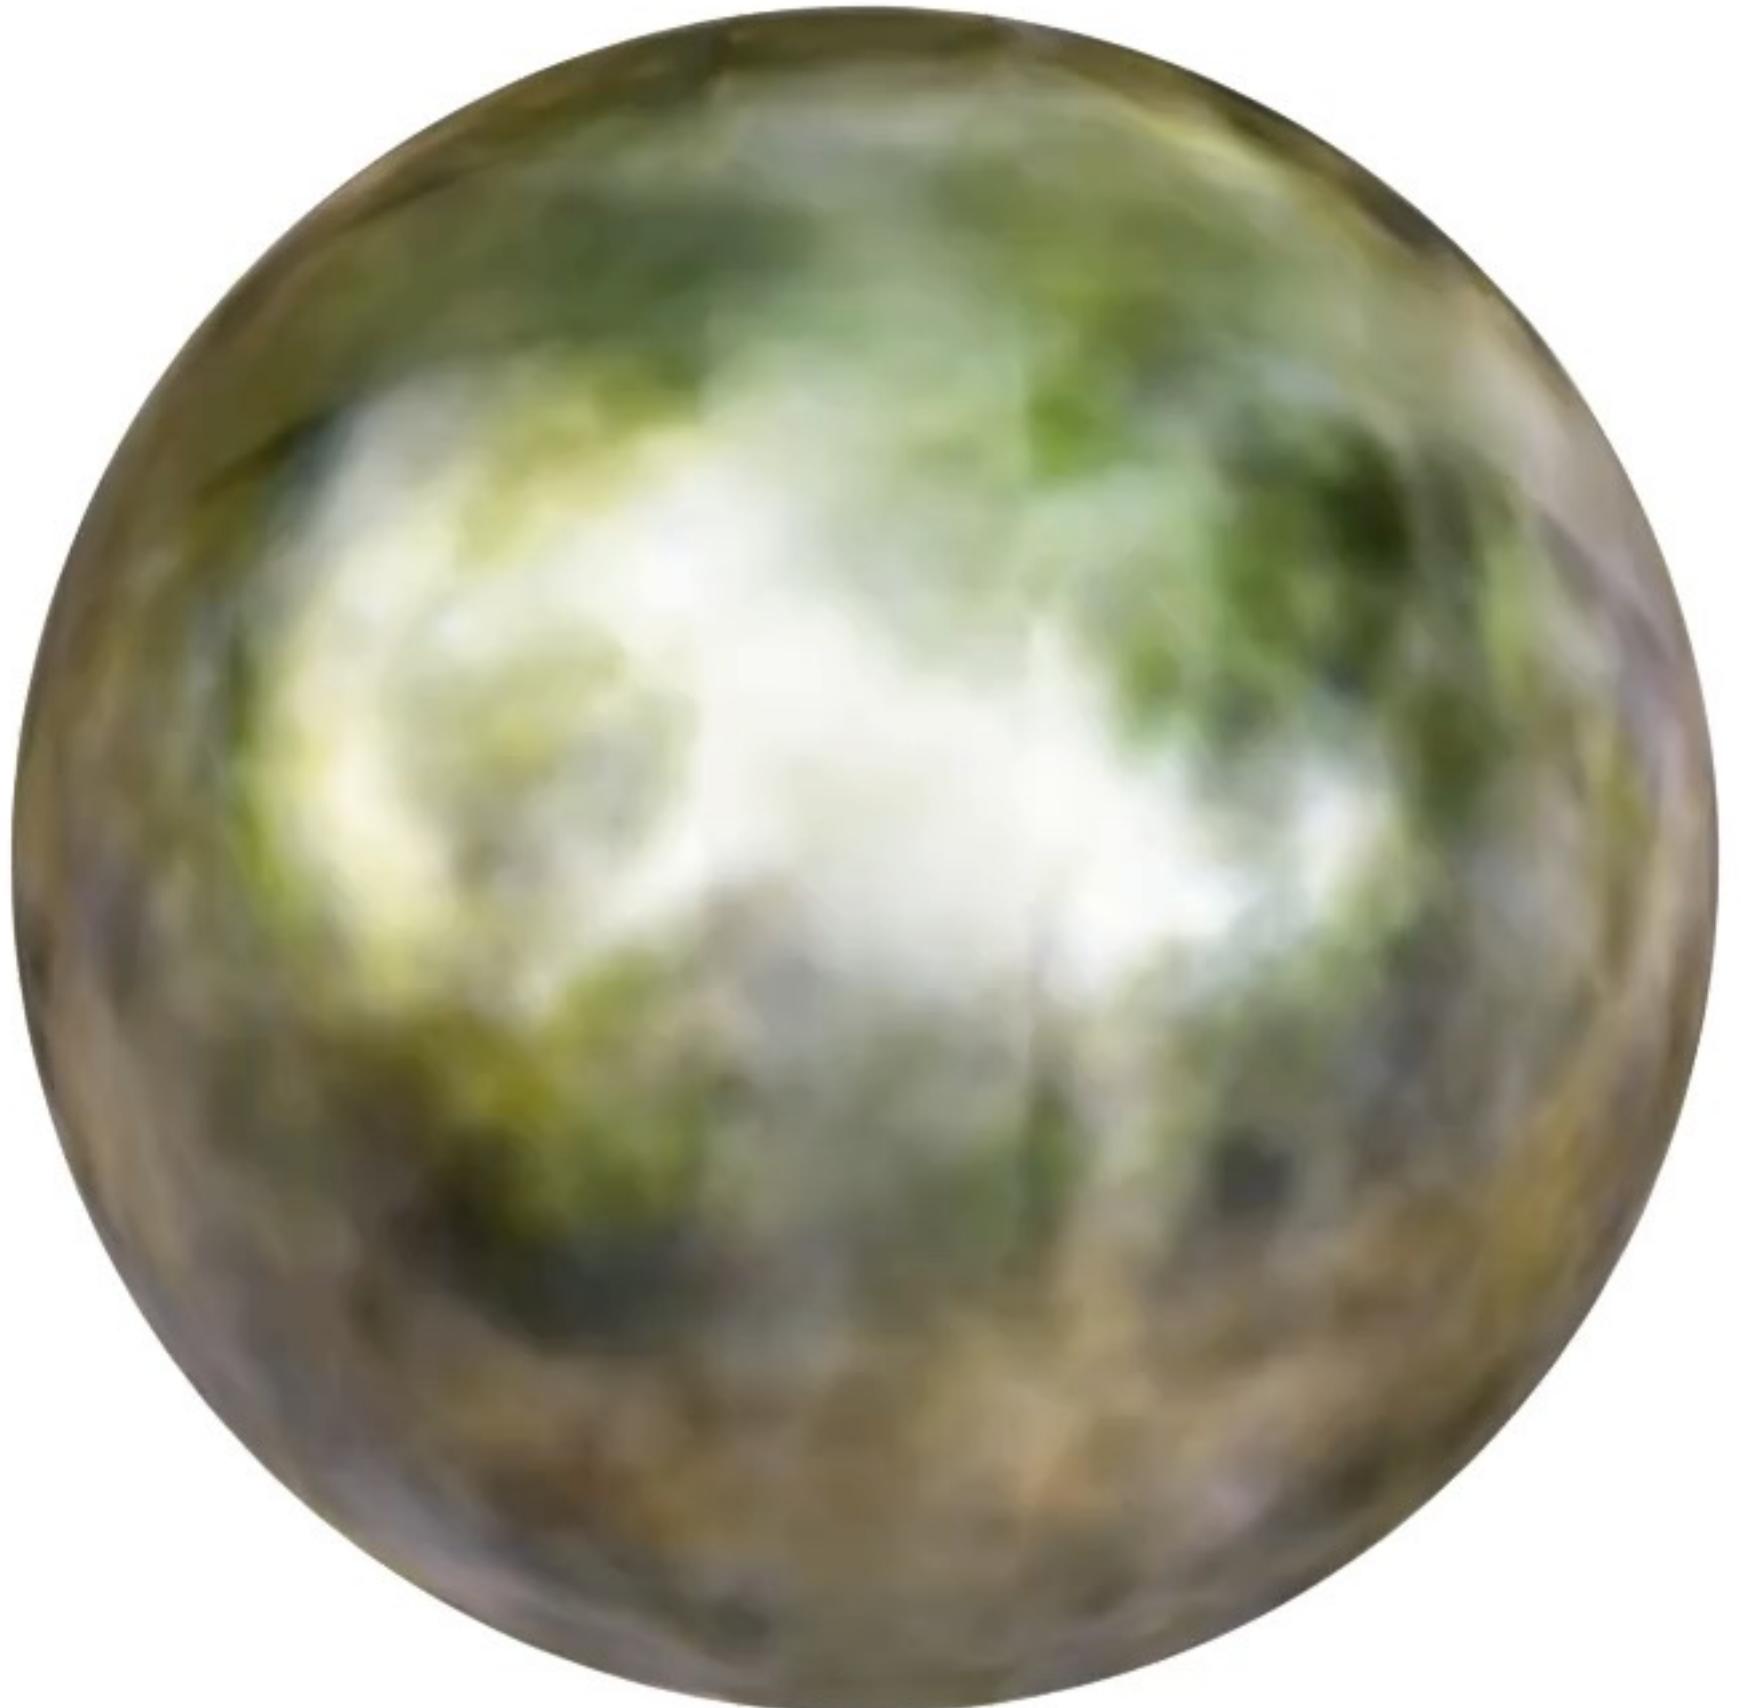


Indirect illumination is even more computationally-expensive



# Modeling light can recover better surfaces

NeRF

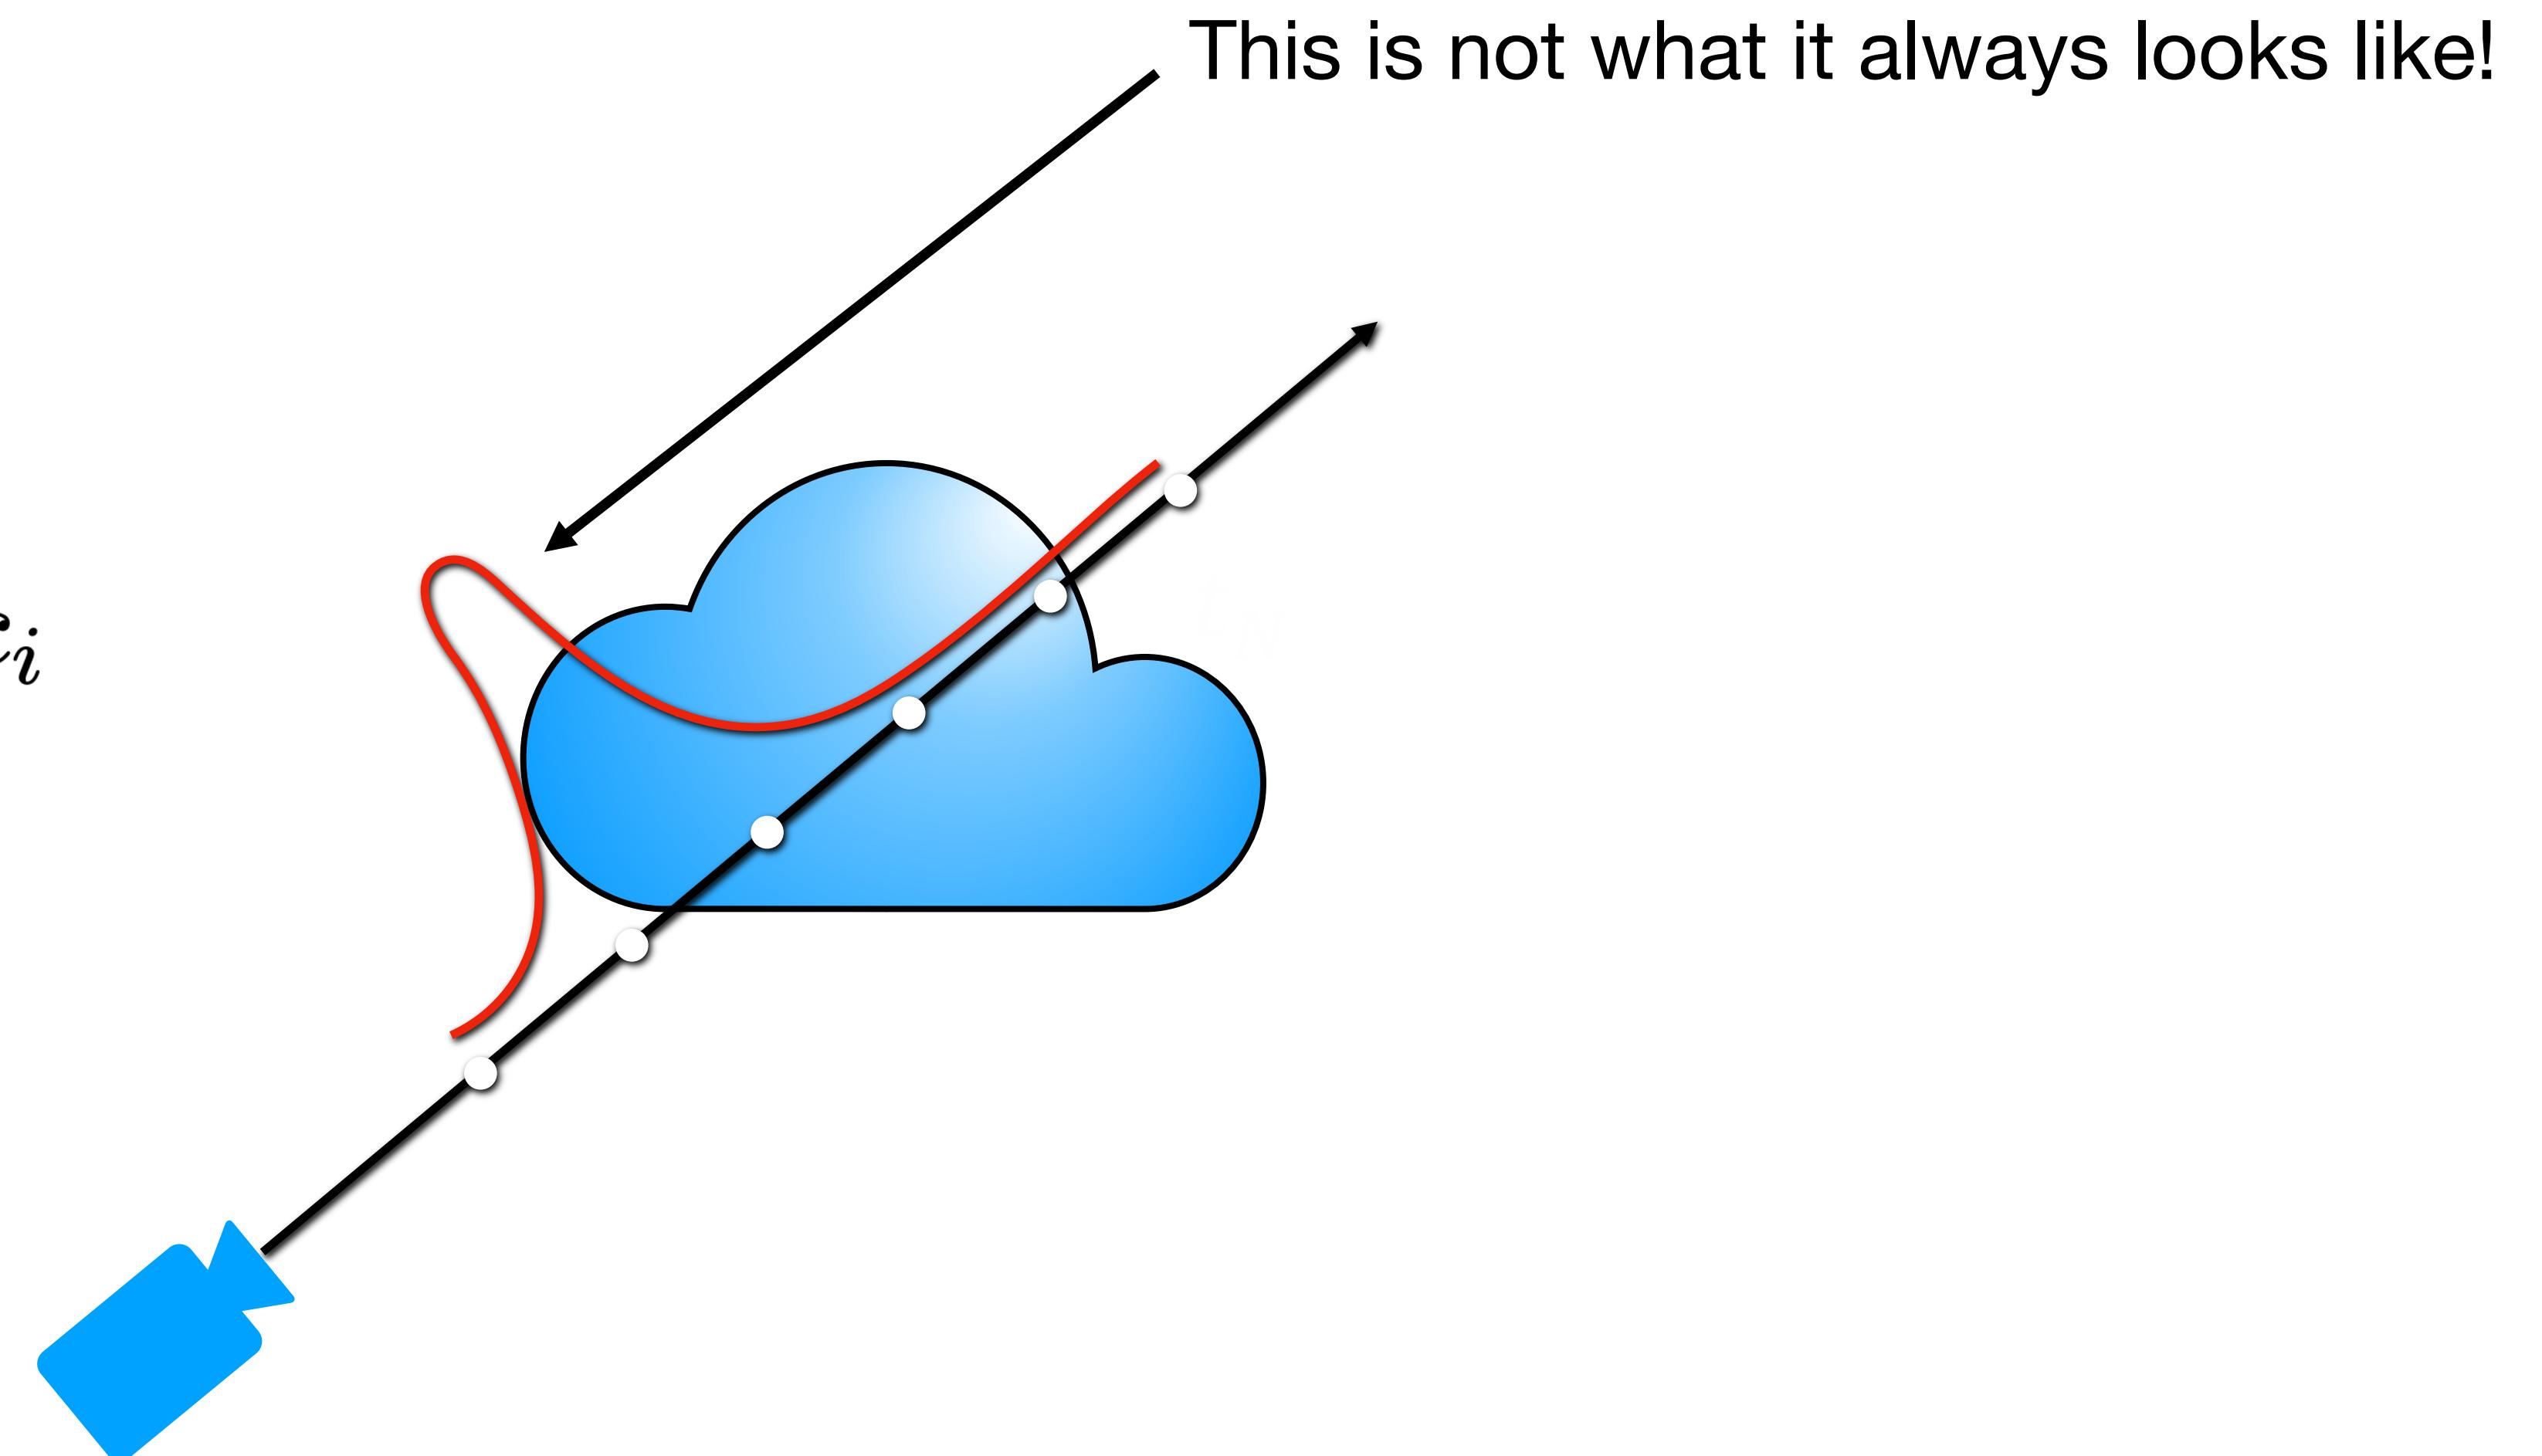


Ref-NeRF



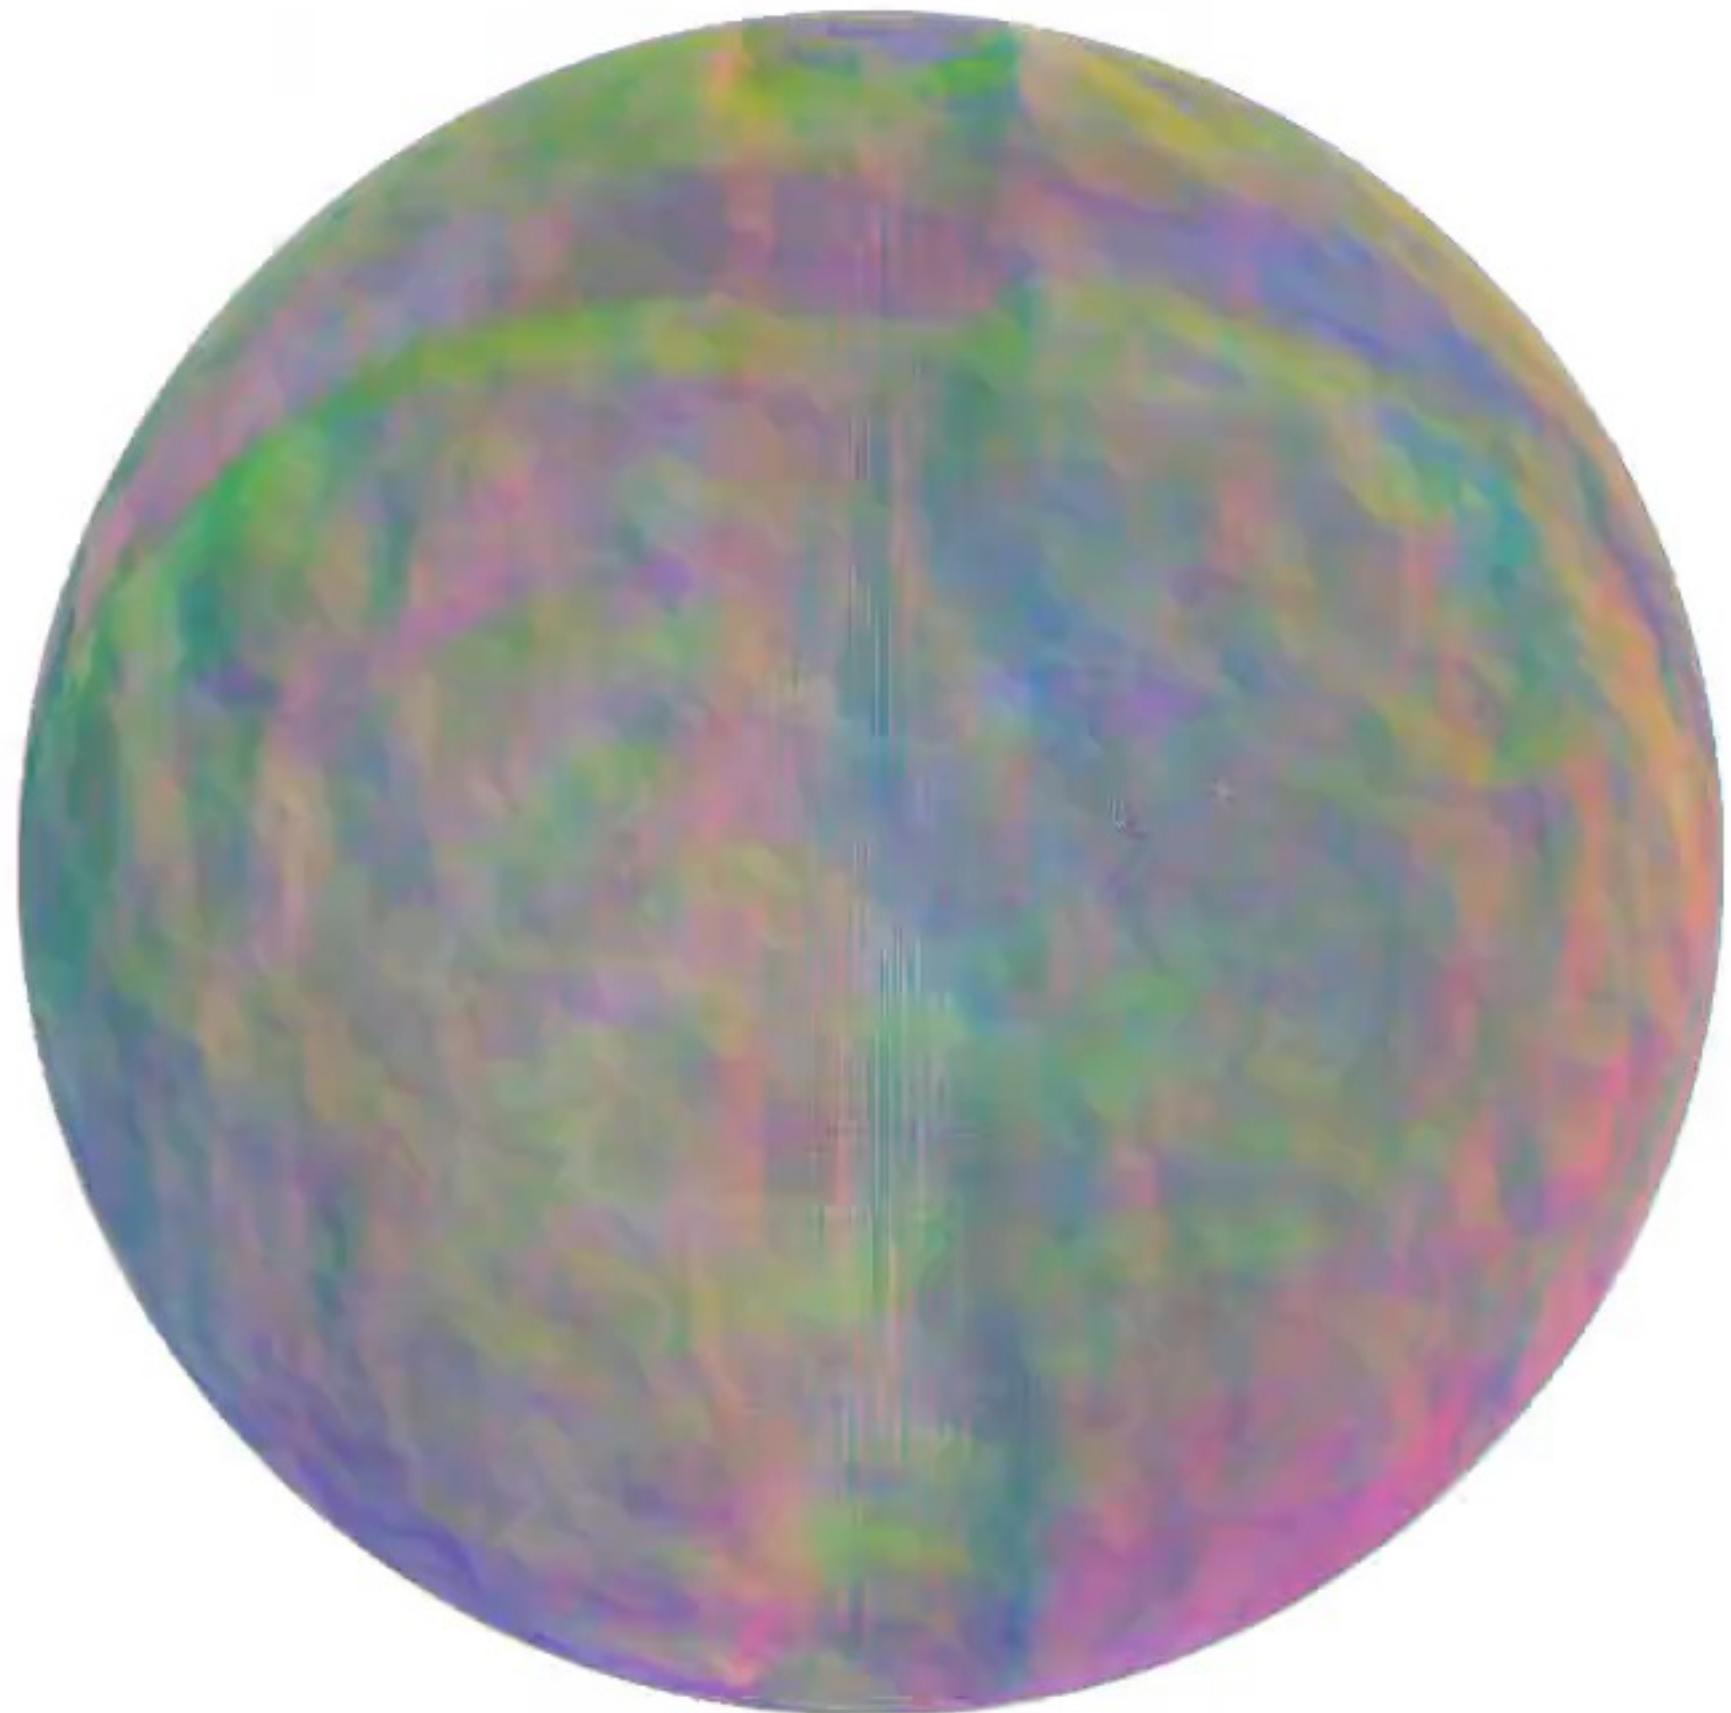
# Decomposing light helps recover sharper surfaces

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

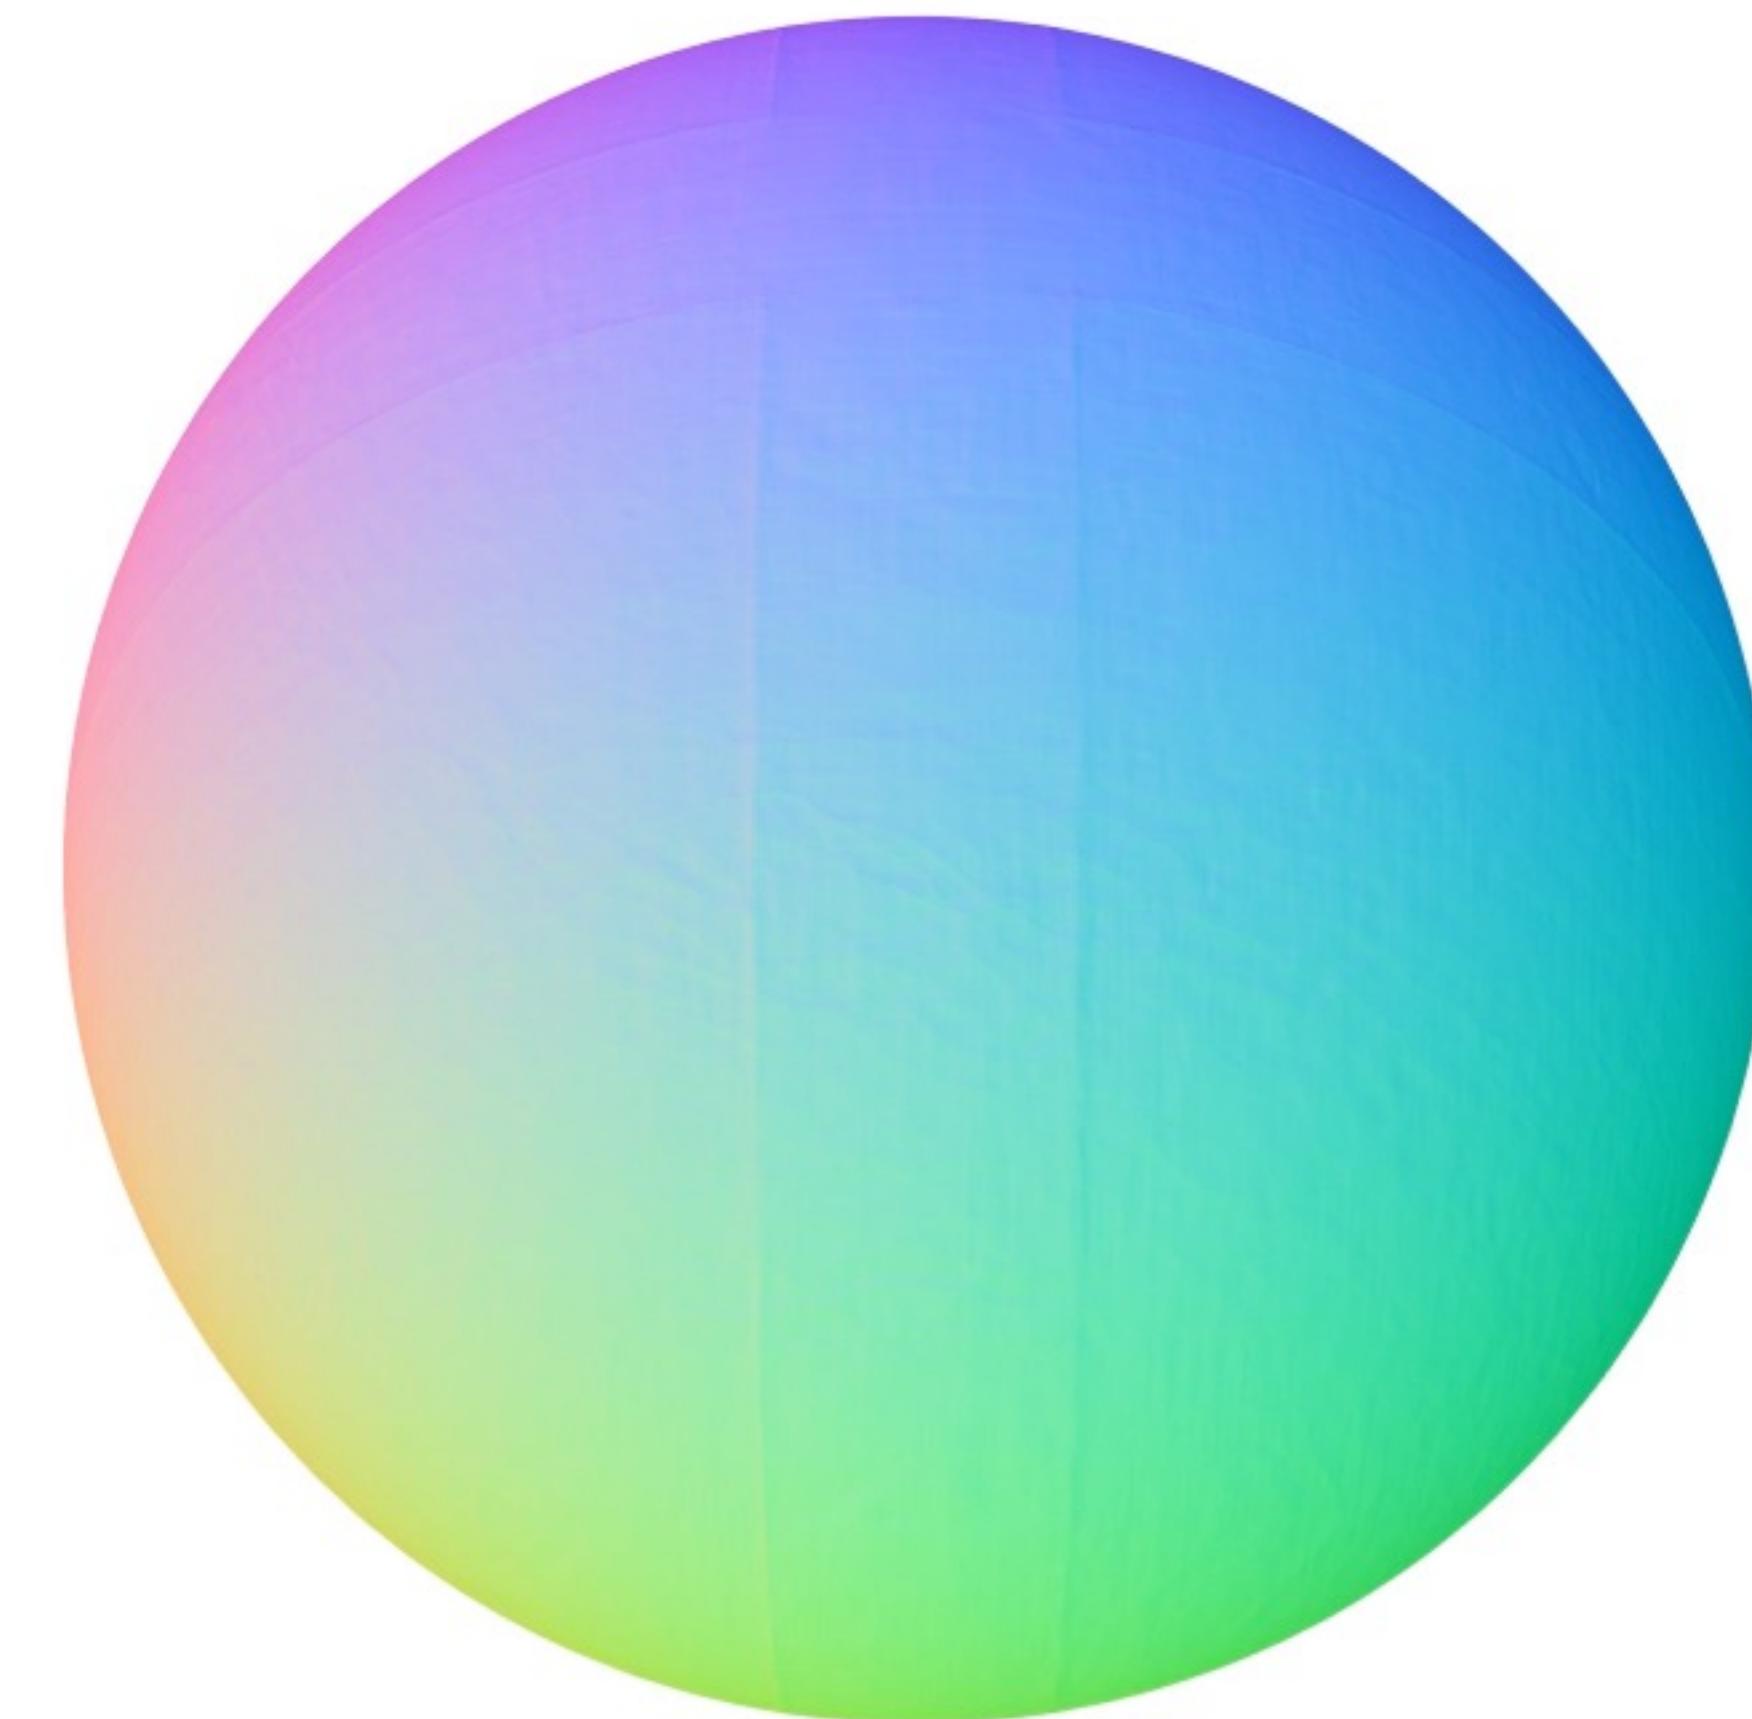


# Better modeling of light helps recover sharper surfaces

NeRF



Ref-NeRF



# Modeling light = better specularities

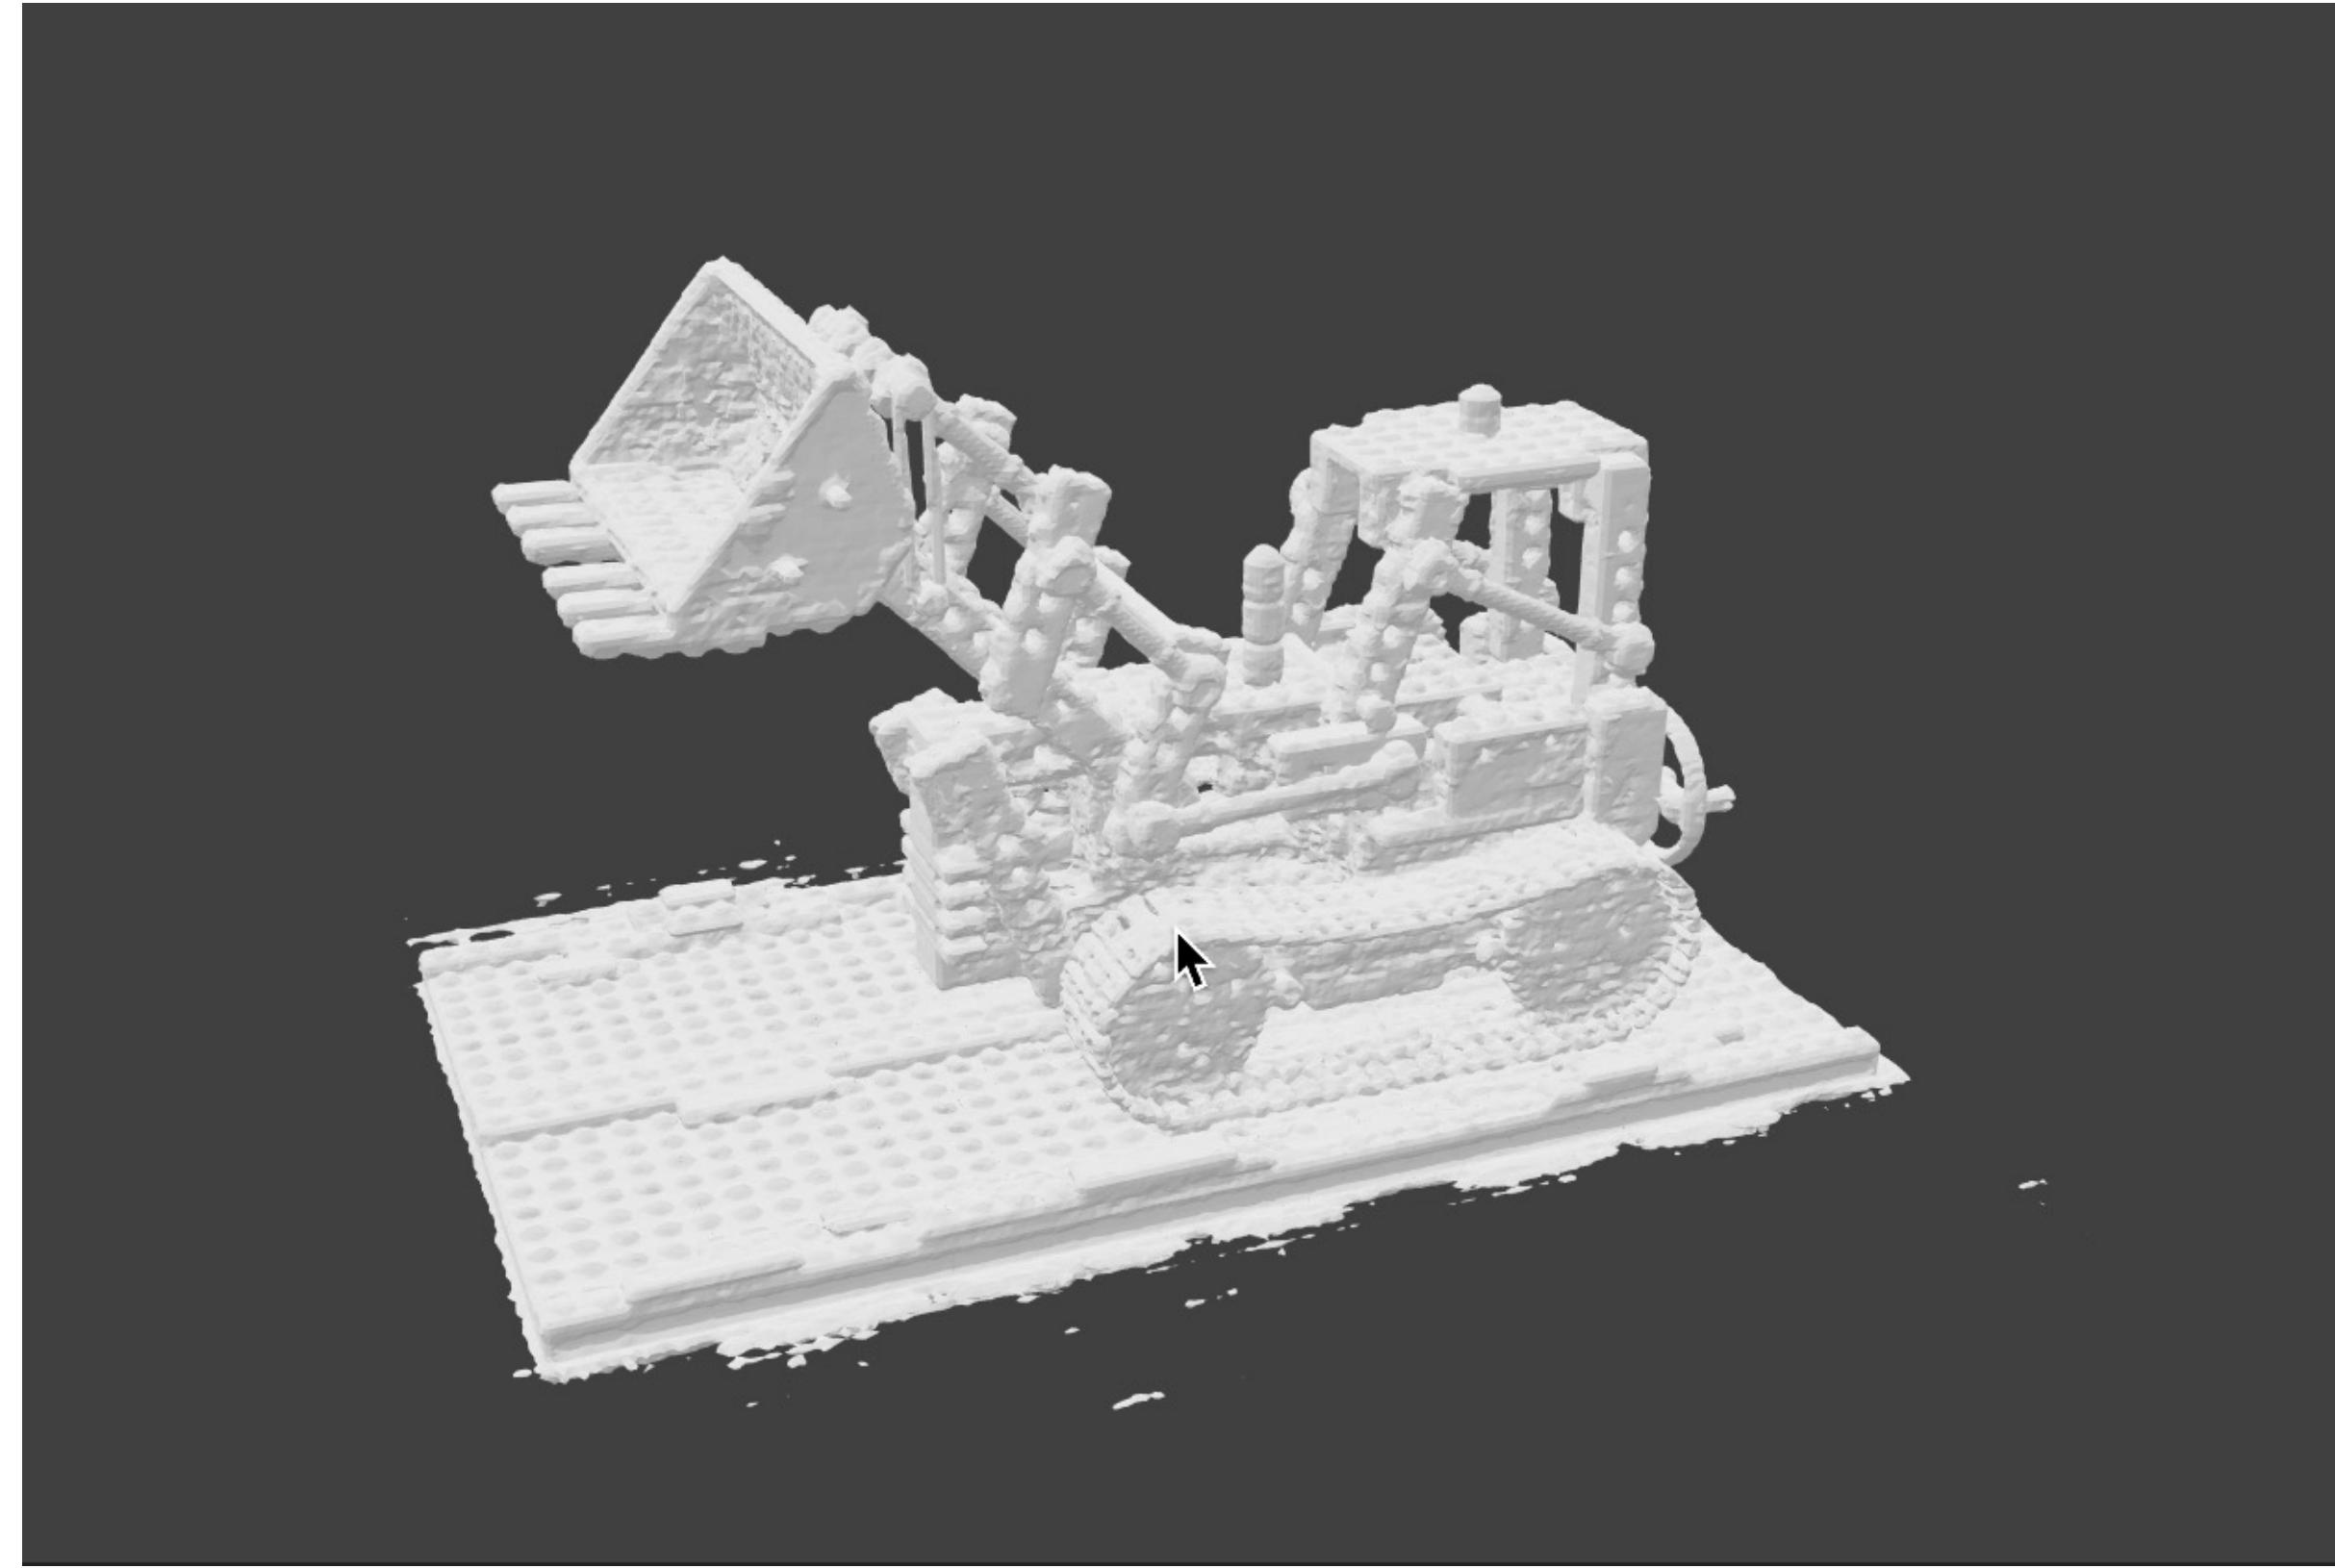
Mip-NeRF



# Editing specular and diffuse colors



# Related Challenge: Extracting Surfaces



- Needed for adoption into existing gaming engines/VFX lifecycle
- Challenges: What if the density recovered isn't peaky (surface) and is not clean? What to do? What about complex scenes, how to group objects?

# The Dynamic World

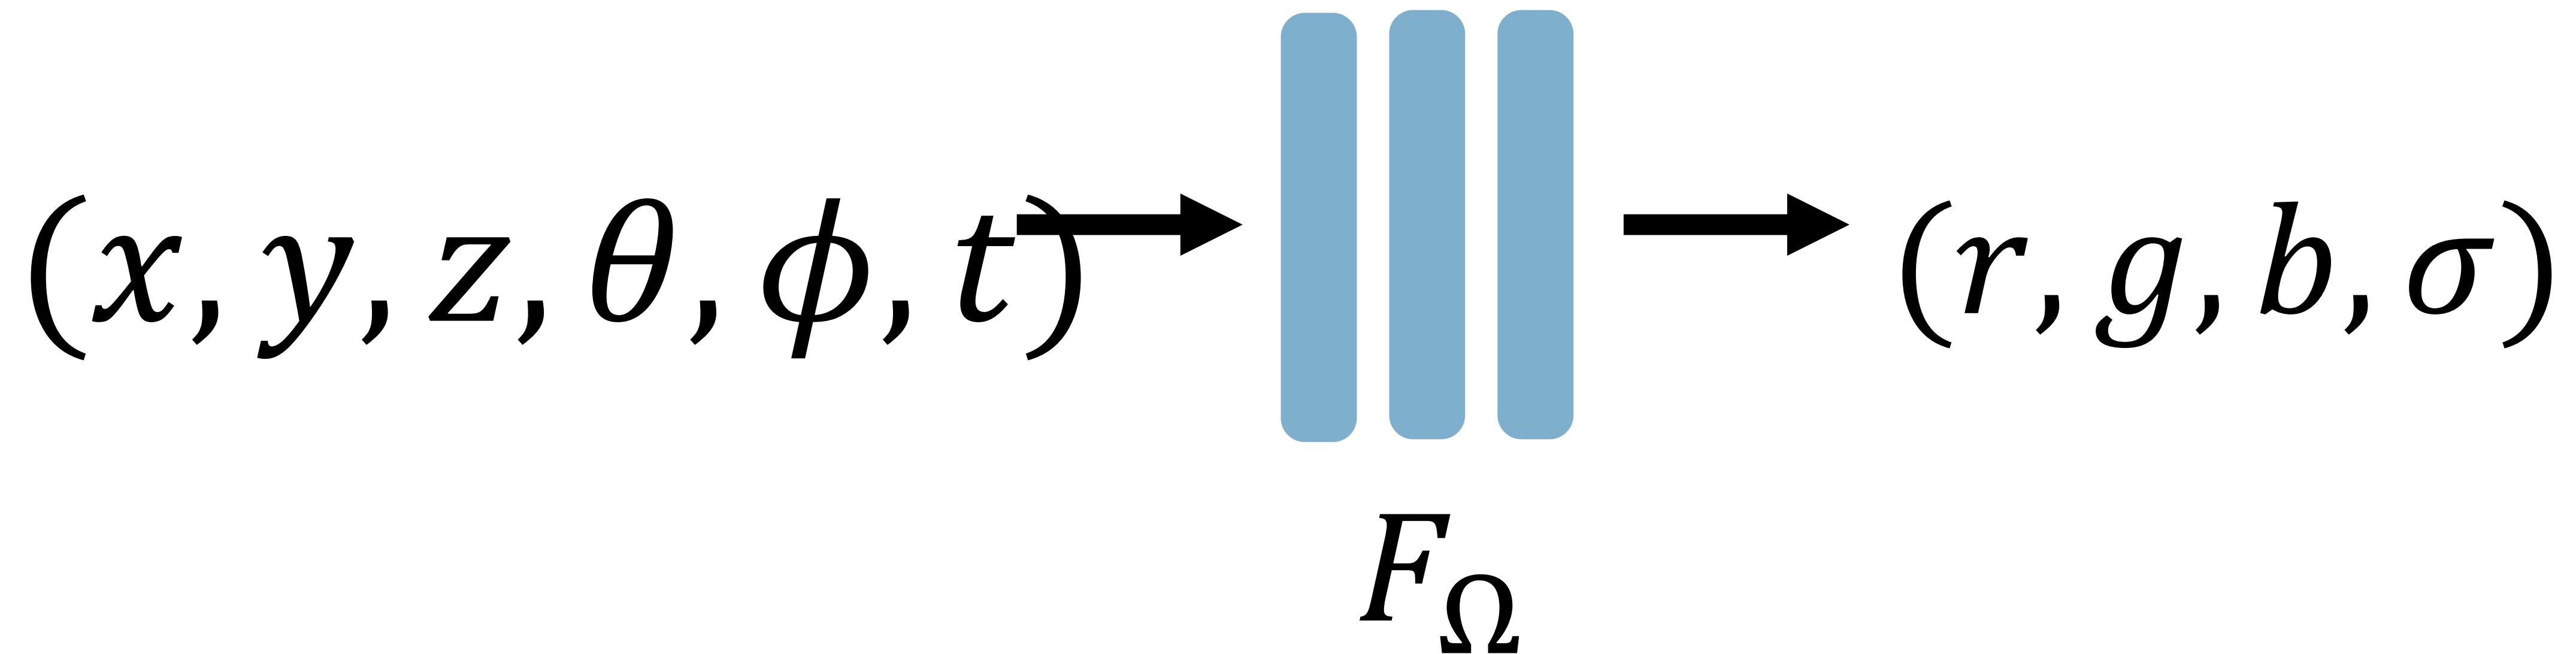


Memories of Australia –Andrew S. Hamilton

# Holy grail

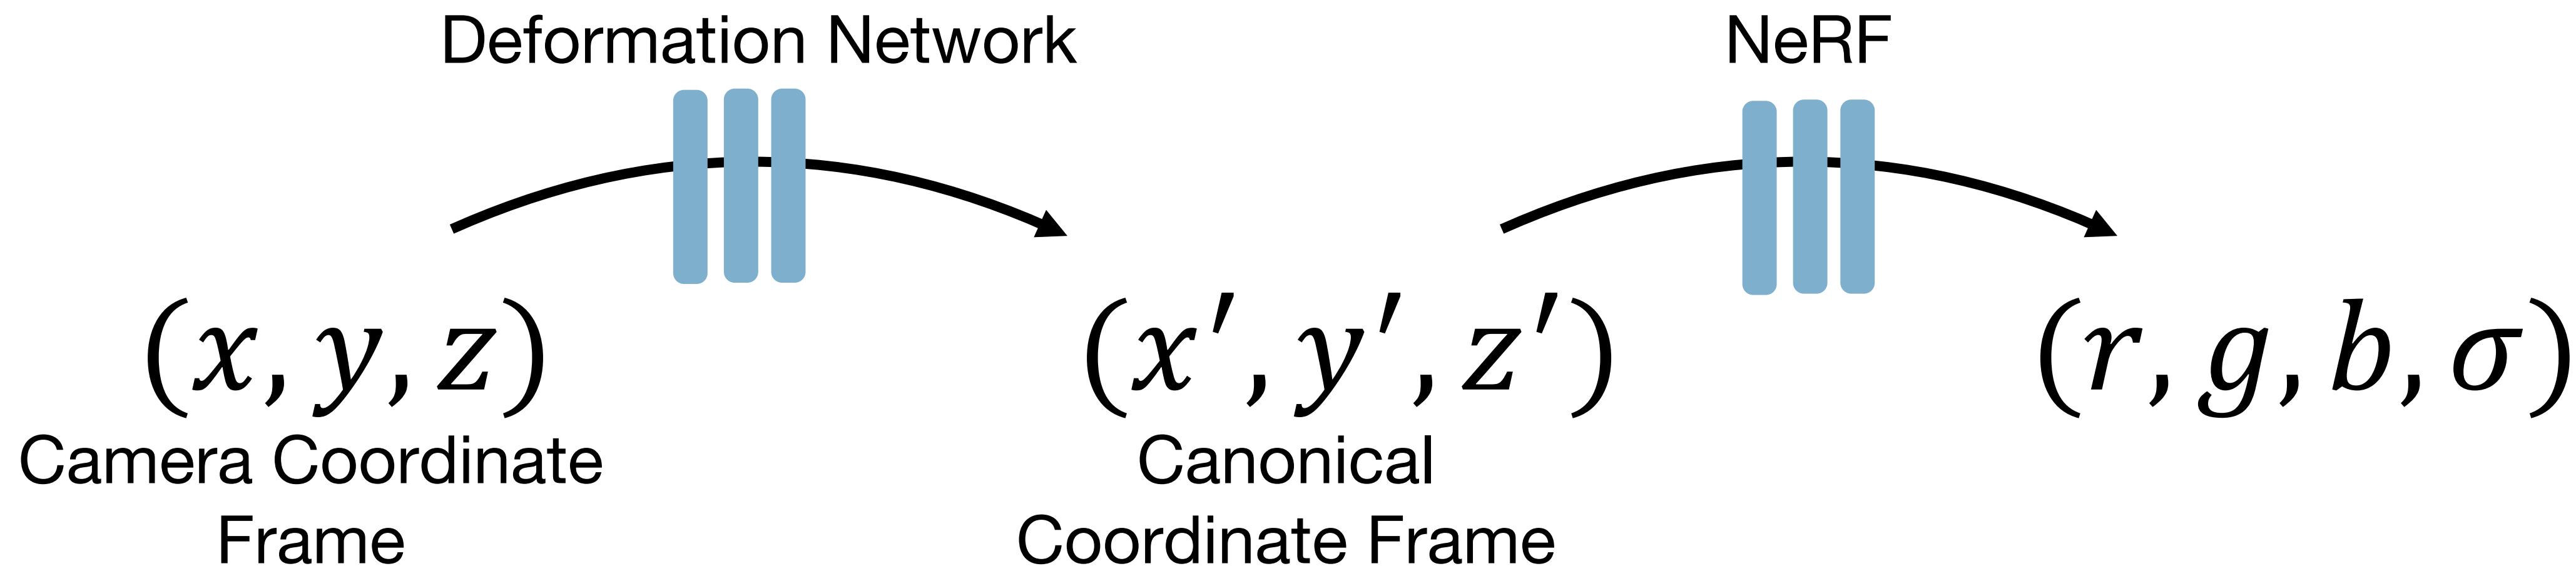
- Dynamic Novel View Synthesis from Monocular Camera
- Very difficult! Extremely under constrained problem

# Simple baseline for adding time



Hard without simultaneous multiple view!

# Through a deformation network



Still very under constrained

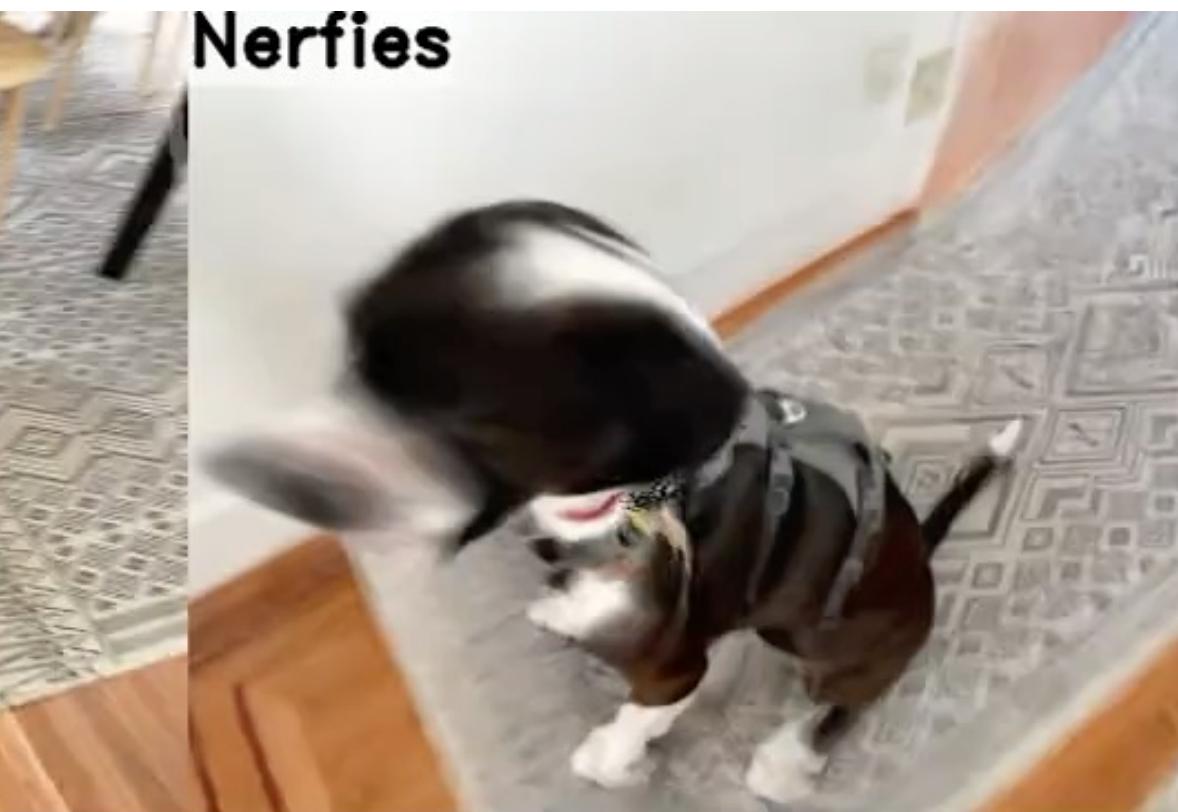
# Dynamic View Synthesis: Monocular is hard



HyperNeRF (Ours)

D-NeRF [Pumarola et al. CVPR 2021], NSFF [Li et al., CVPR 2021], HyperNeRF [Park et al. SIGASia 2021].....

- But performance on in-the-wild monocular capture still far [Gao et al. NeurIPS 2022]



train view

Nerfies



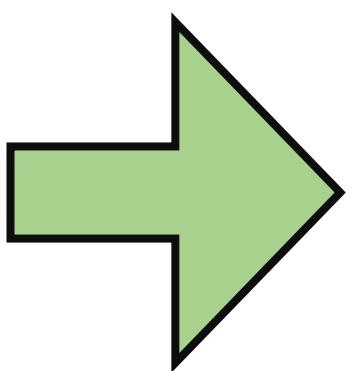
train view

Nerfies

# What if we knew how they deform?



HMMR, Kanazawa et al.  
CVPR 2019



# **Other kinds of dynamic changes**

# Appearance Changes

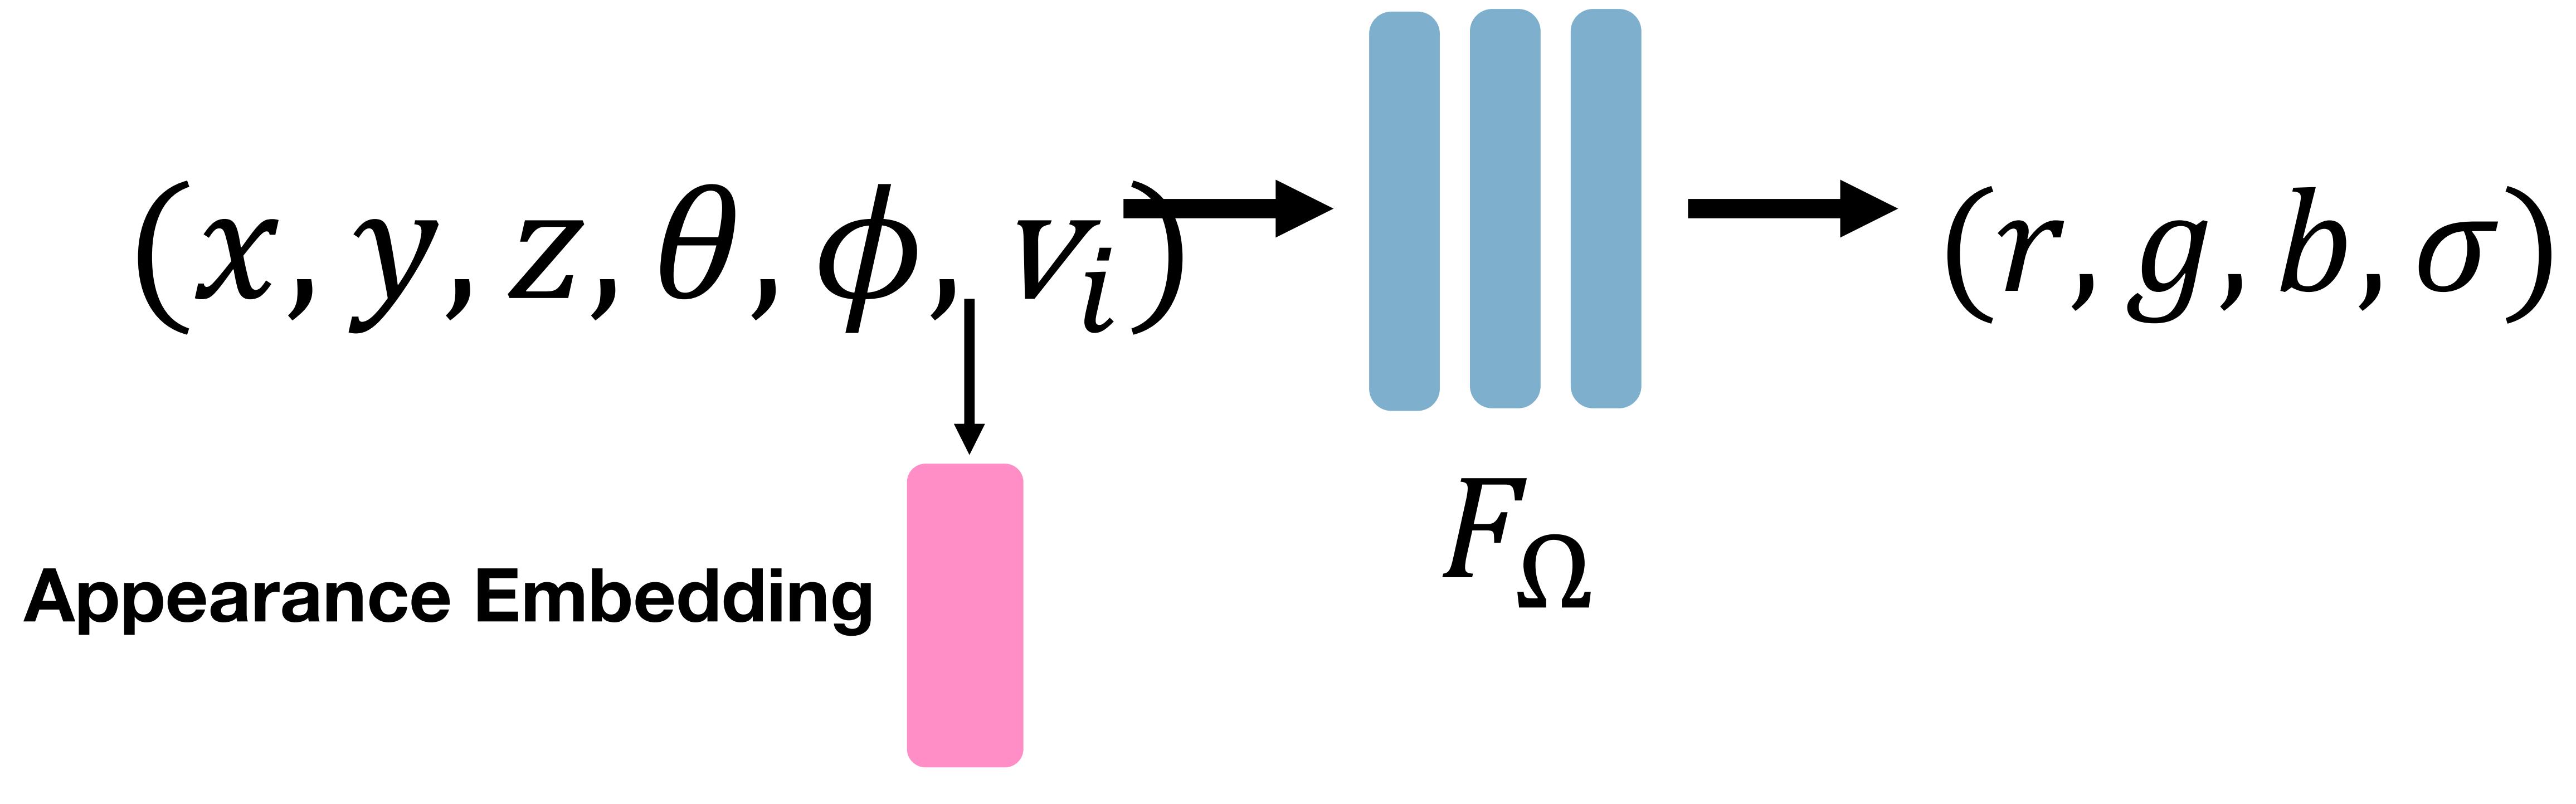
Exposure differences

Lighting changes (day, night)..

Clouds passing by..



# Appearance Embedding: Pretty Robust Solution

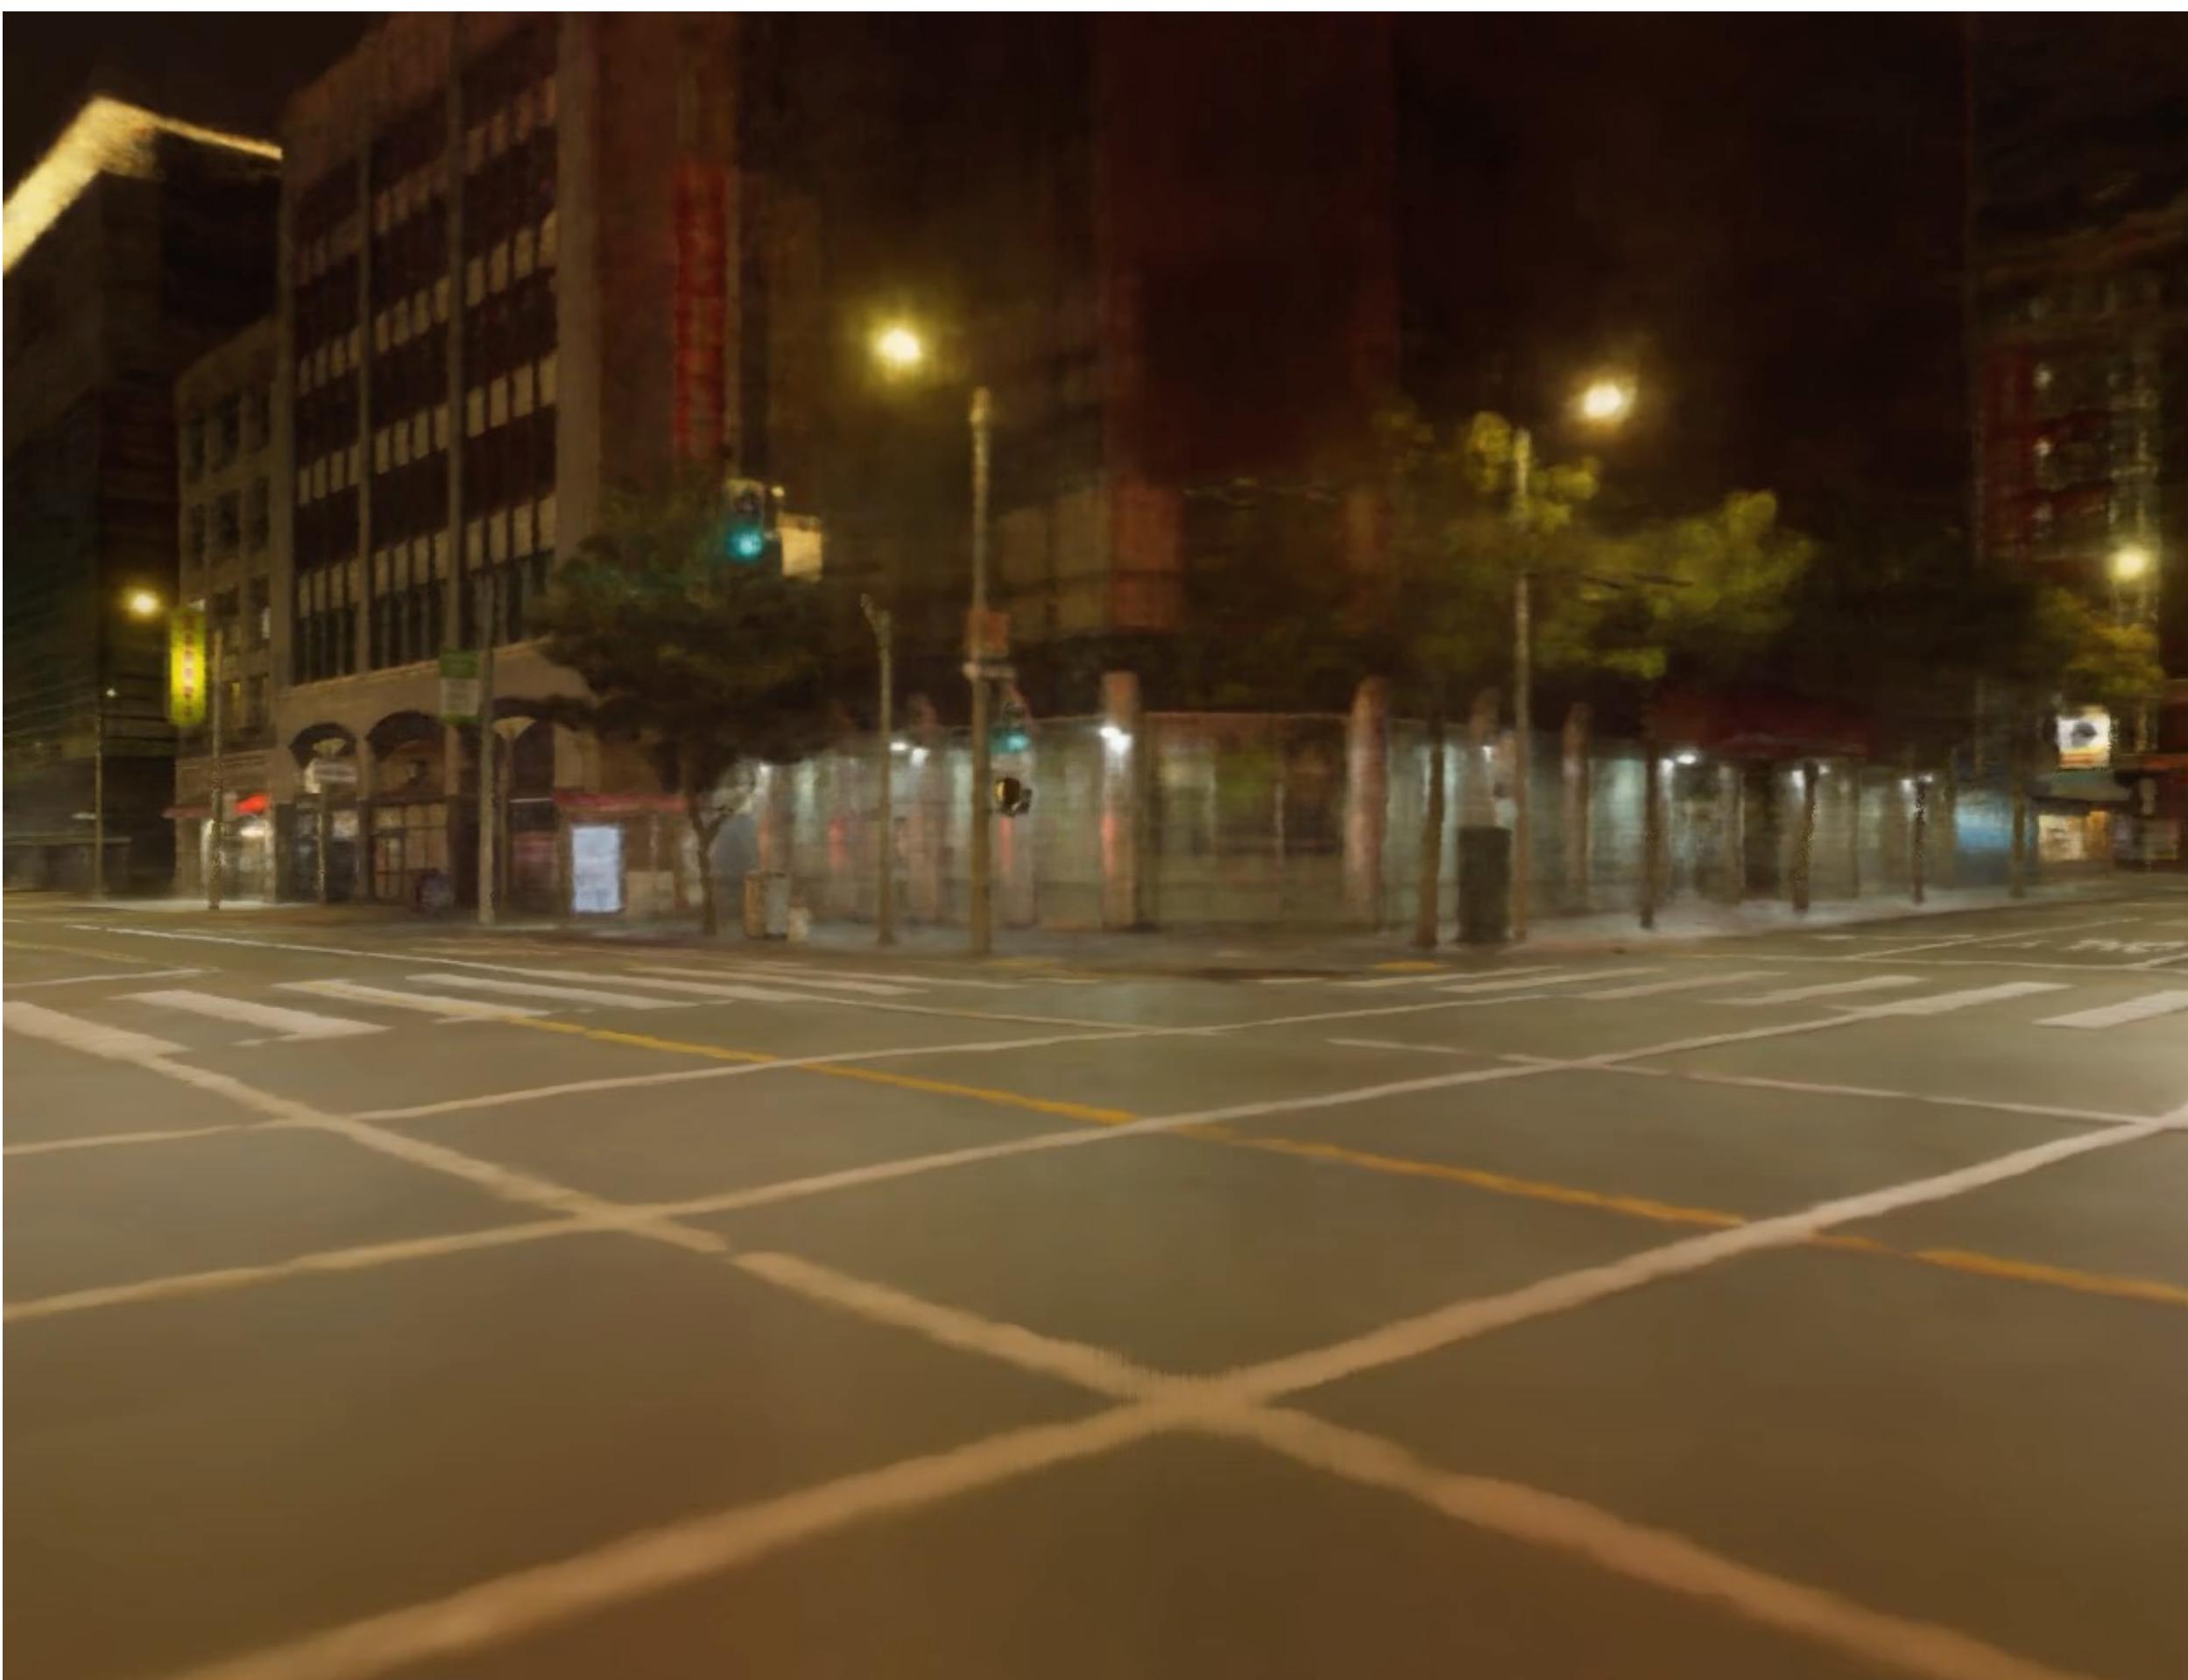


Optimized *per* image: “Auto-Decoding”

ie GLO: Generative Latent Optimization [[Bojanowski et al.](#) ICML 2018]

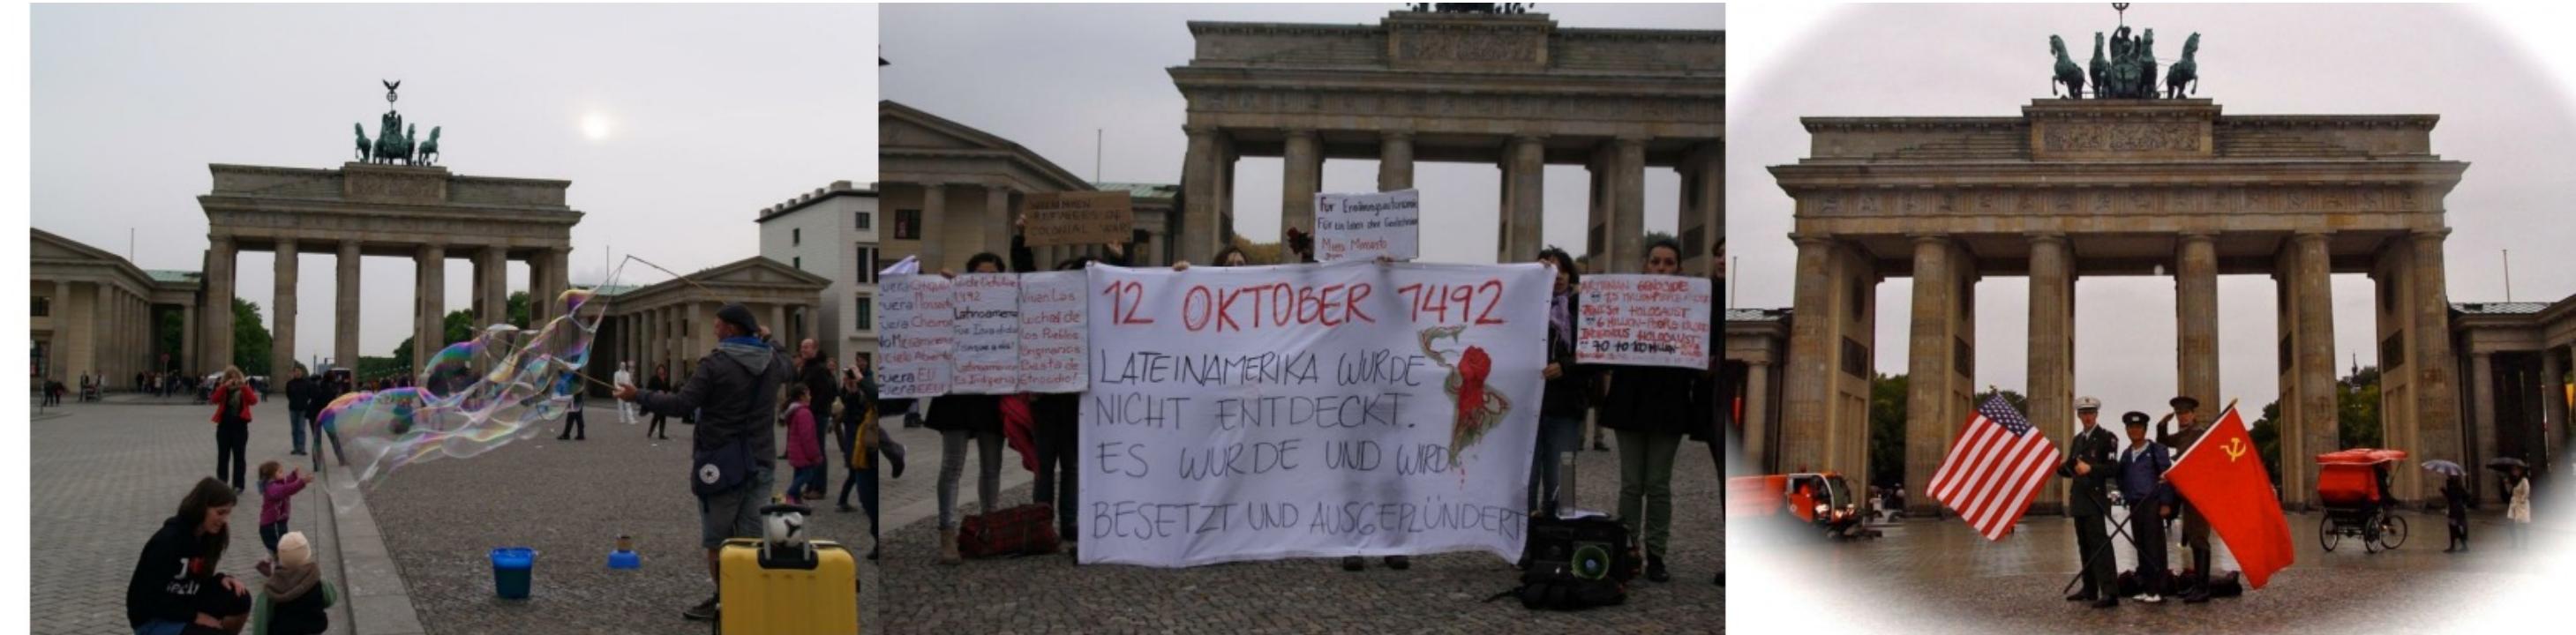
# Appearance Changes

Appearance Encoding is Effective

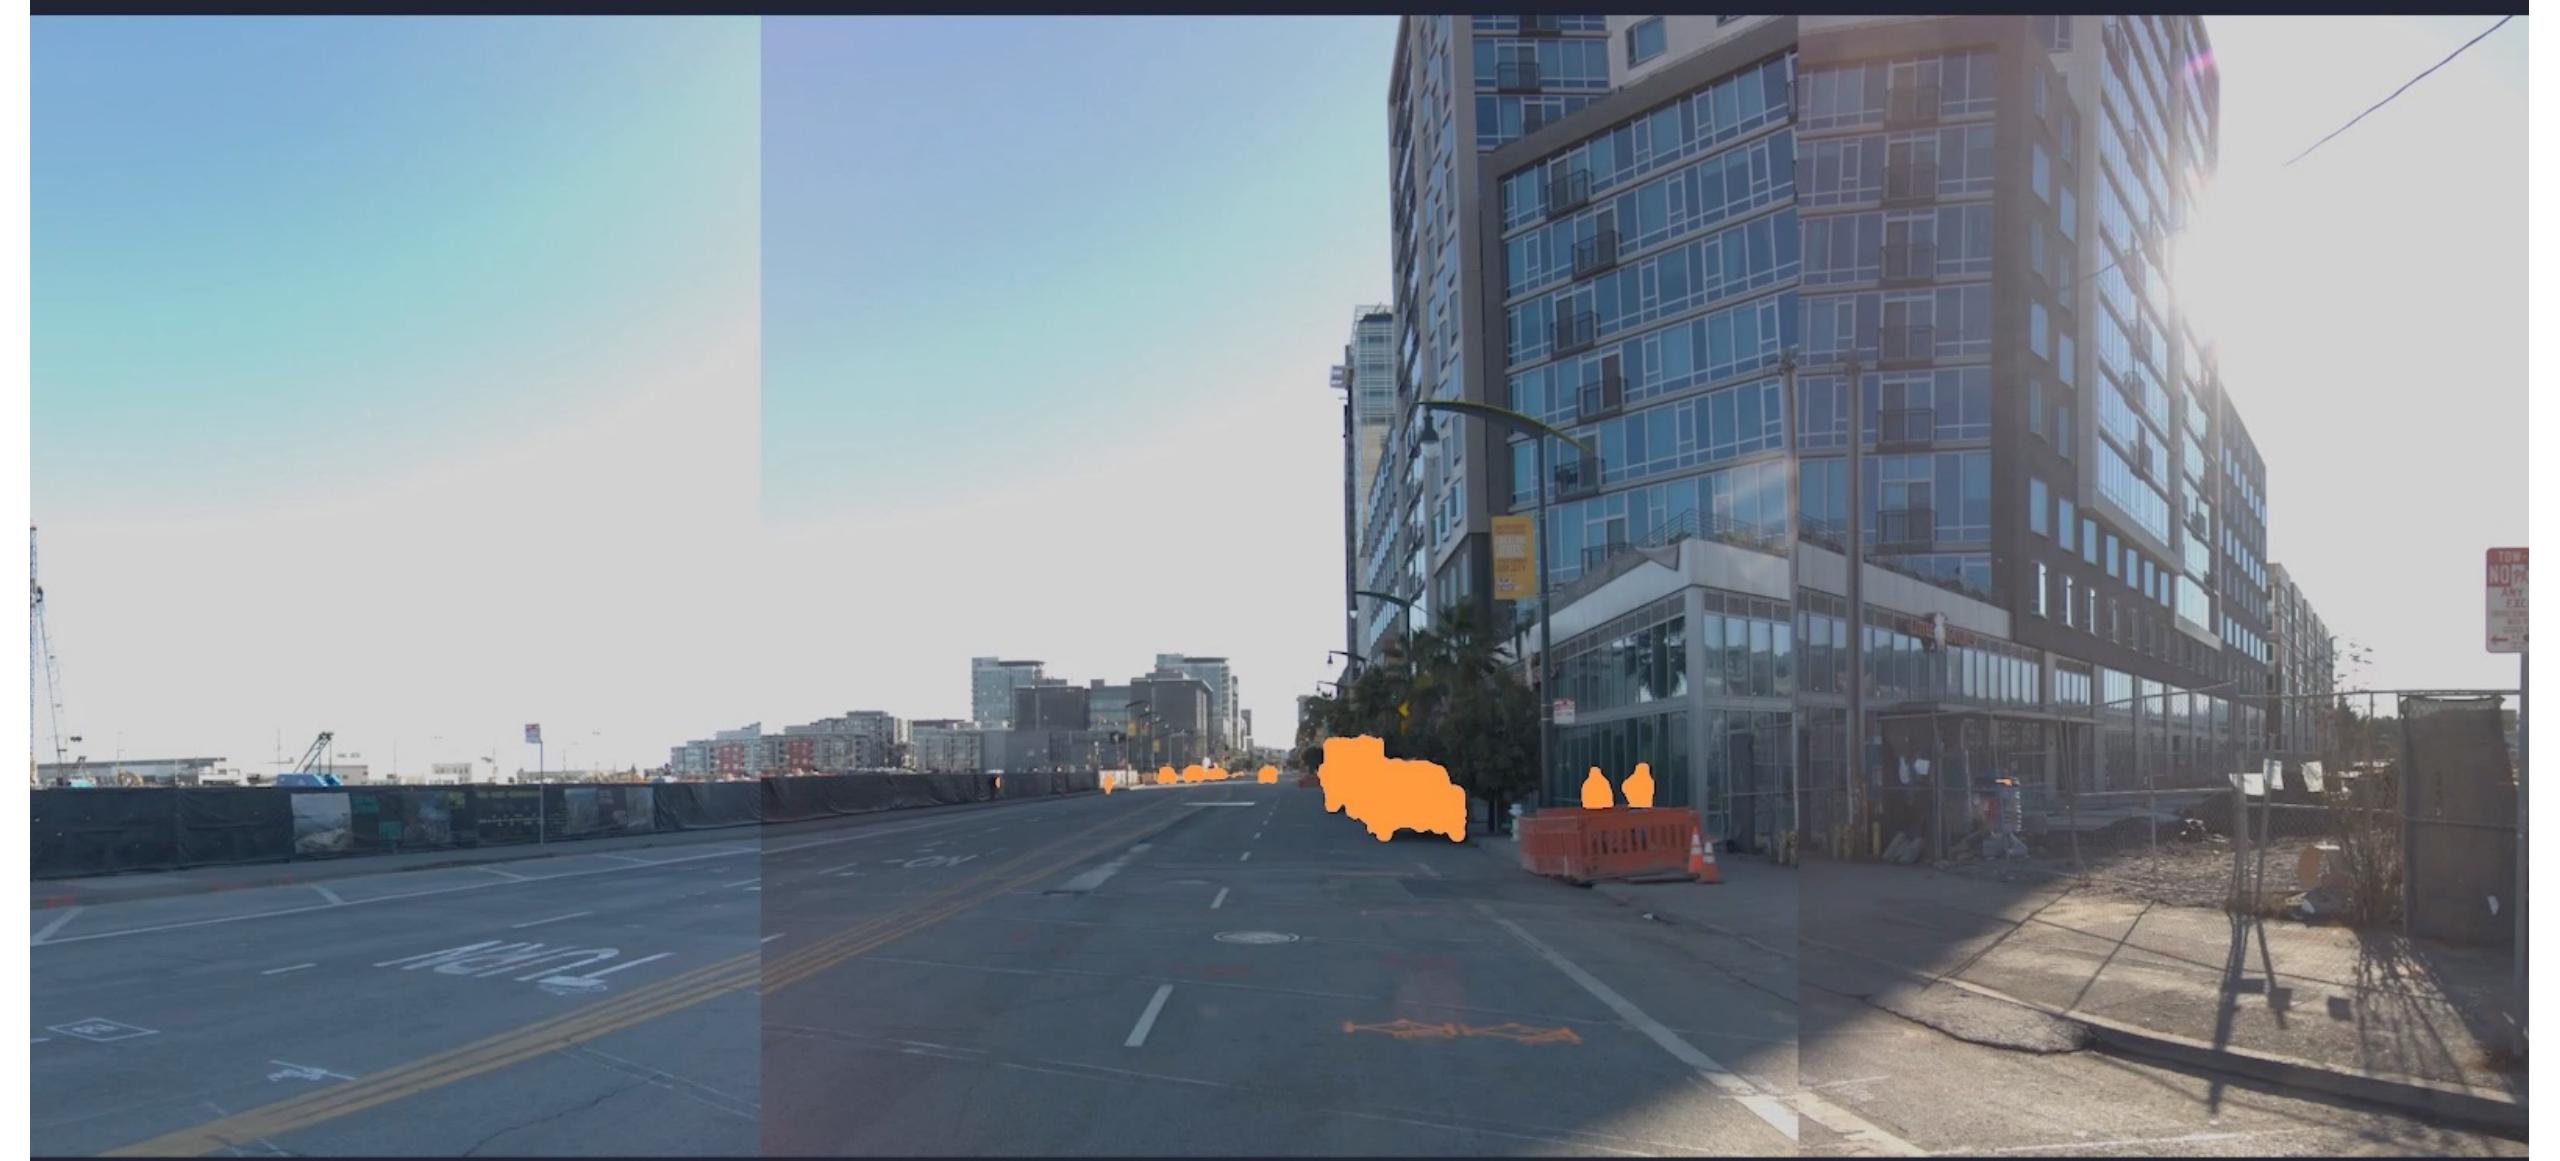


# Transient objects

- Happens all the time! People moving around, interacting with the world
- Difficult! Problem of Grouping
  - how do you know which part is connected or
  - Can use two NeRFs, one global, one per-image, but this often leads to degenerate solutions
- Current solution: Ignore (mask out)



Filter Dynamic Objects Using Segmentation Masks



# Why is dynamic scenes hard?

- Unless you have a light dome
- Essentially you only have a single-view

# **Building & Reusing Prior Knowledge**

Machine Learning

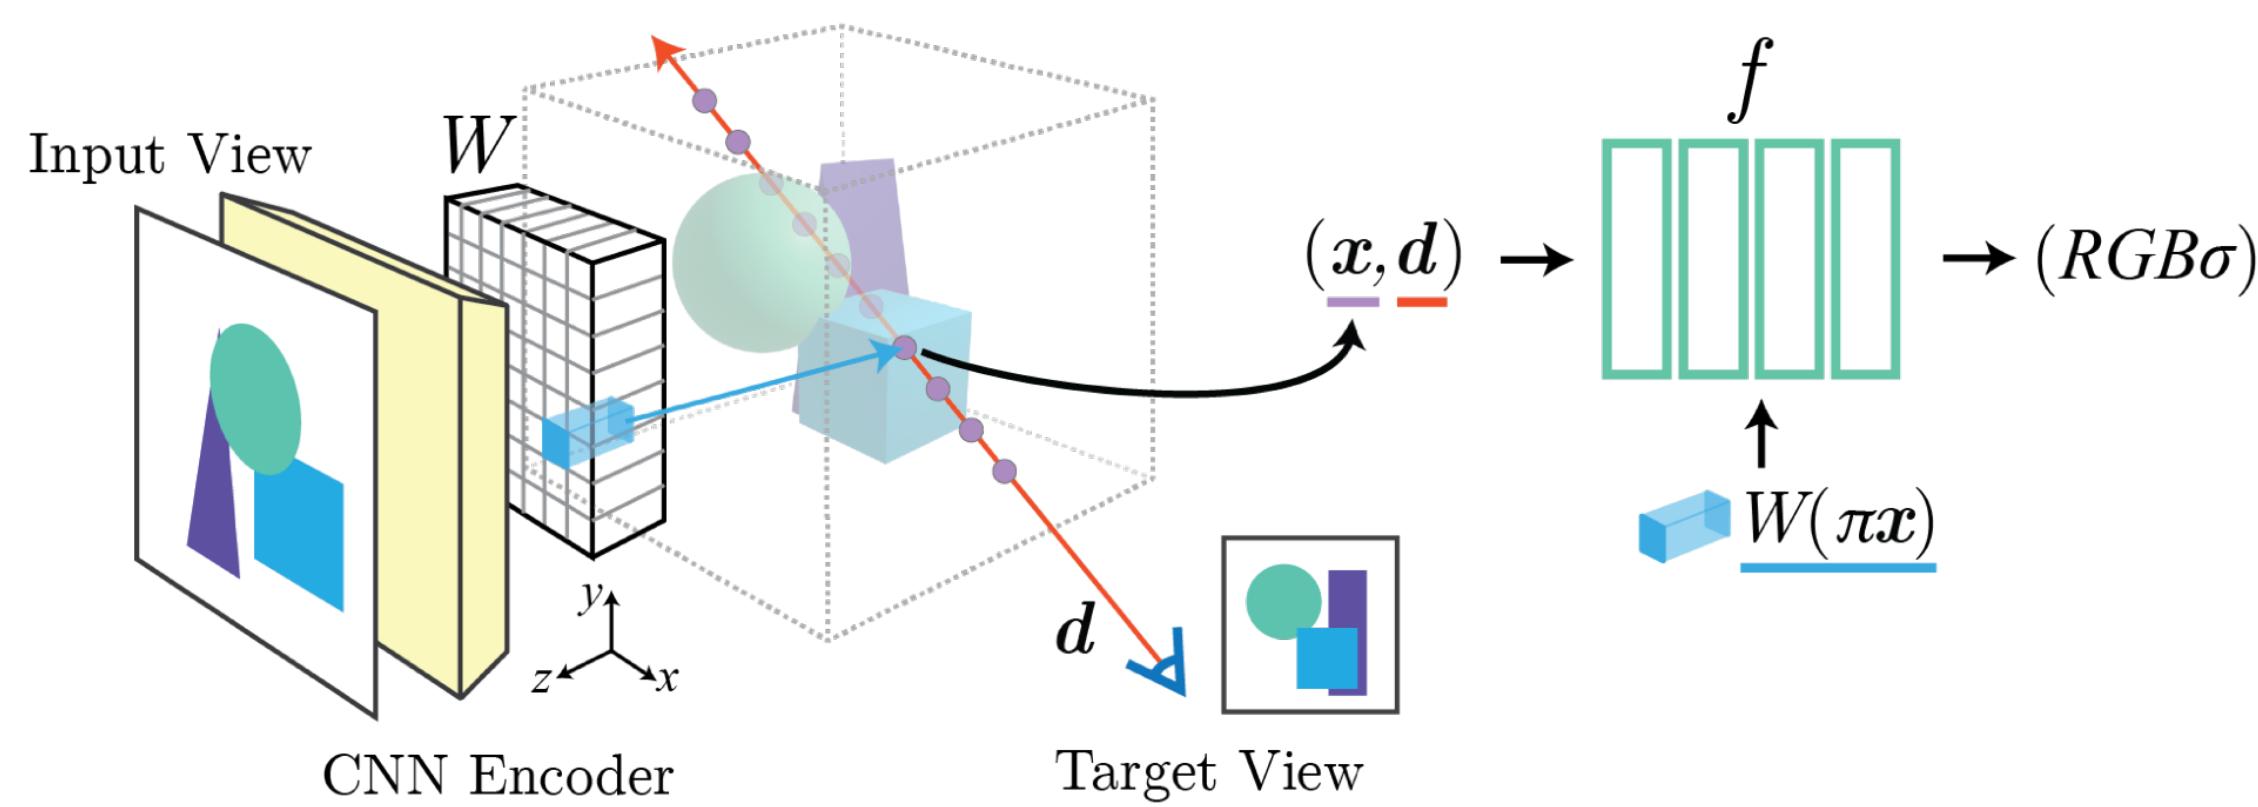
# NeRF is per-scene optimization

- We need lots of images to get good view synthesis!!
- Also there's no knowledge reused from prior scene reconstructions
- How to bring learning in the picture?



# Few-shot NeRF

- One-shot (single-view): pixelNeRF [Yu et al. CVPR'19]



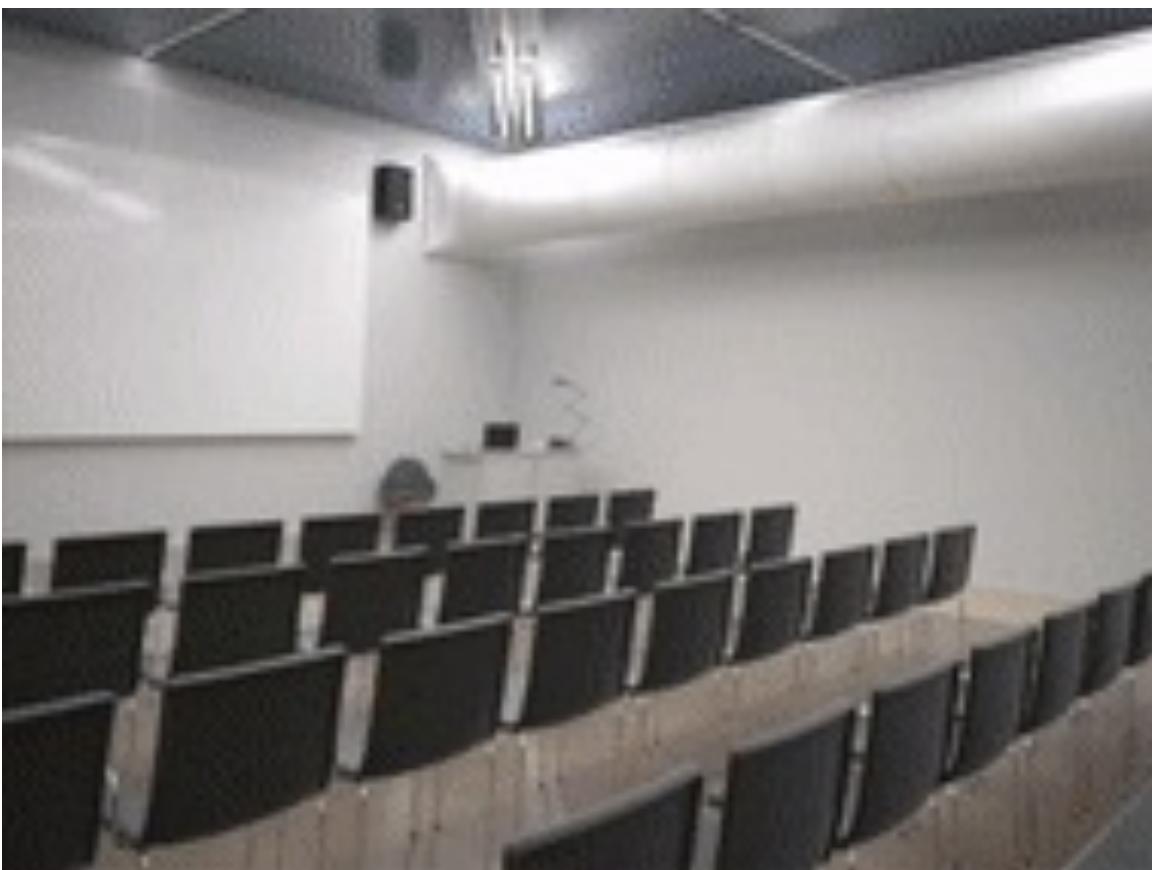
- Few-shot (3~10 views): pixelNeRF, IBRNet [Wang et al. CVPR'21], MVSNet [Chen et al. ICCV'21], etc...
- Challenging for predicting completely unseen real scenes



IBRNet

# Data is the bottleneck

- Large-scale Real-World Multi-view Data is hard to collect:  
CO3D [Reizenstein ICCV 2021]
- A lot to learn from other single-view 3D prediction models:



# Generating NeRFs from 2D Generative Models



DreamFusion  
[Poole et al.  
arXiv 2022]

# Enabling specific edits

**Input Image**



**Edited Image**



“A bird spreading wings”

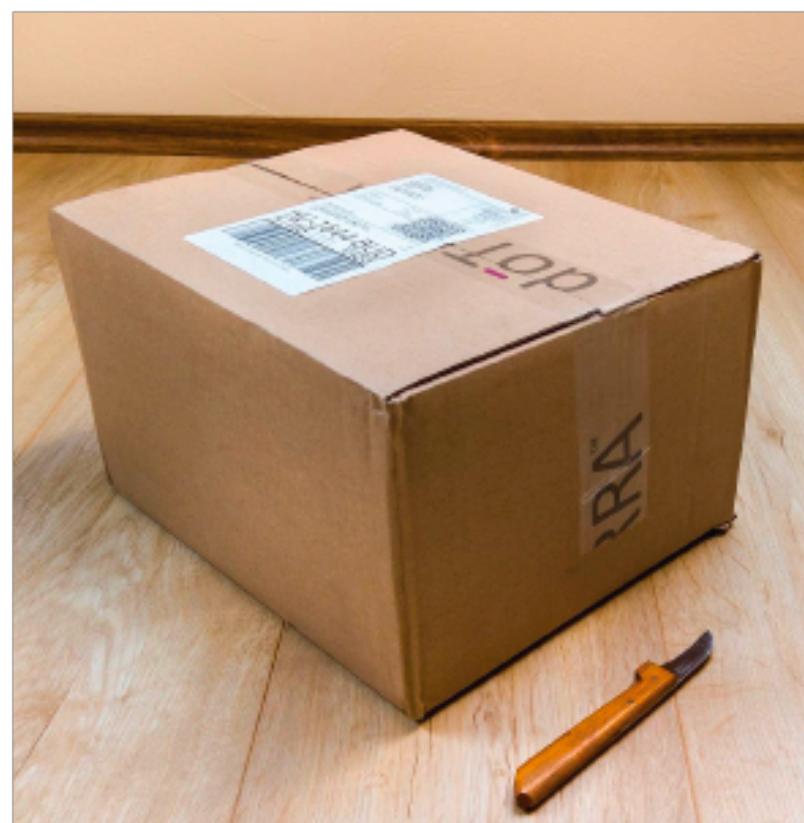
**Input Image**



**Edited Image**



“Two kissing parrots”



“A photo of an open box”



“A photo of a sitting dog”

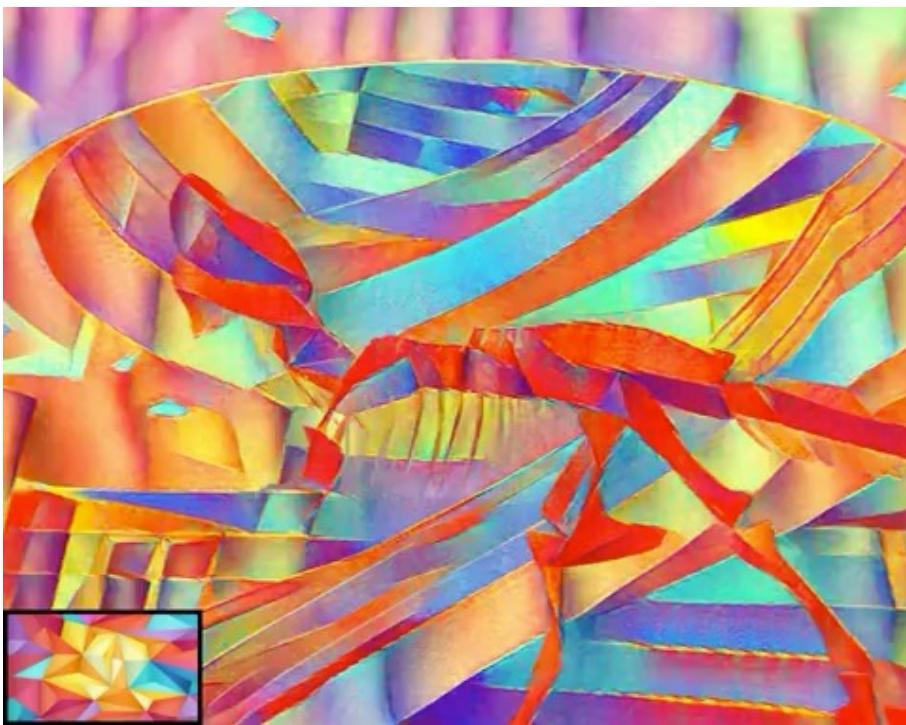
<https://imagic-editing.github.io/>

Kawar and Zada et al. Arxiv 2022

# Semantic Editing



# Manipulating captured scenes





Matthew Tancik\*, Ethan Weber\*, Evonne Ng\*, Ruilong Li, Brent Yi,  
Terrance Wang, Alexander Kristoffersen, Jake Austin, Kamyar Salahi,  
Abhik Ahuja, David McAllister, Angjoo Kanazawa

+14 additional Github collaborators



# Step 1: Capture Data



# Step 1: Capture Data

- Maximize view coverage
  - Try to get at least 5 views per point in scene.
  - Wide angle / Fisheye lenses work well
- Minimize motion blur
  - Noise is an OK tradeoff
- Minimize dynamic objects

## Step 2: Recover Camera Poses

```
(nerfstudio) :~/nerfstudio$ []
```

COLMAP scripts

## Step 2: Recover Camera Poses



COLMAP Alternative  
Record3D

# Step 3: Optimize NeRF!

```
(nerfstudio) :~/nerfstudio$ ll data/nerfstudio/desolation/
total 93541
drwxr-xr-x 7 evonneng users      9 Oct  5 00:58 .
drwxr-xr-x 9 evonneng users      9 Oct  5 00:58 ..
drwxr-xr-x 2 evonneng users     3 Oct  5 00:58 camera_paths/
drwxr-xr-x 8 evonneng users     9 Oct  5 01:00 colmap/
drwxr-xr-x 2 evonneng users    209 Oct  5 00:58 images/
drwxr-xr-x 2 evonneng users    209 Oct  5 00:58 images_2/
drwxr-xr-x 2 evonneng users    209 Oct  5 00:58 images_4/
-rw-r--r-- 1 evonneng users 95638275 Oct  5 00:58 IMG_8981.MOV
-rw-r--r-- 1 evonneng users 177295 Oct  5 00:58 transforms.json
(nerfstudio) :~/nerfstudio$ ns-train nerfacto --data data/nerfstudio/desolation/ --viewer.websocket-port 7
```

[GETTING STARTED](#)[GITHUB](#)[DOCUMENTATION](#)

VIEWPORT RENDER VIEW



▶ RESUME TRAINING

[Show Scene](#)[Show Images](#)

Refresh Page

Resolution: 640x1024px

Time Allocation: 100% spent on viewer

Server Connected | Render Connected



CONTROLS



RENDER



SCENE

[LOAD PATH](#)[EXPORT PATH](#)

Height

1080

Width

1920

FOV

50

Seconds

4

FPS

24

[+ ADD CAMERA](#)

Smoothness



0.00

0

1

2

3



CAMERA 0



## Step 4: Render

```
(nerfstudio) :~/nerfstudio$ █
```



# Goals of nerfstudio

- Modular Framework
- Open, Evolving Framework
- Reference Source

# Modularity

## Encoders

- Positional Encoding
- Fourier Features
- Hash Encoding
- Spherical Harmonics
- Matrix Decomposition

## Samplers

- Uniform
- Occupancy
- PDF
- Proposal
- Spacing Fn

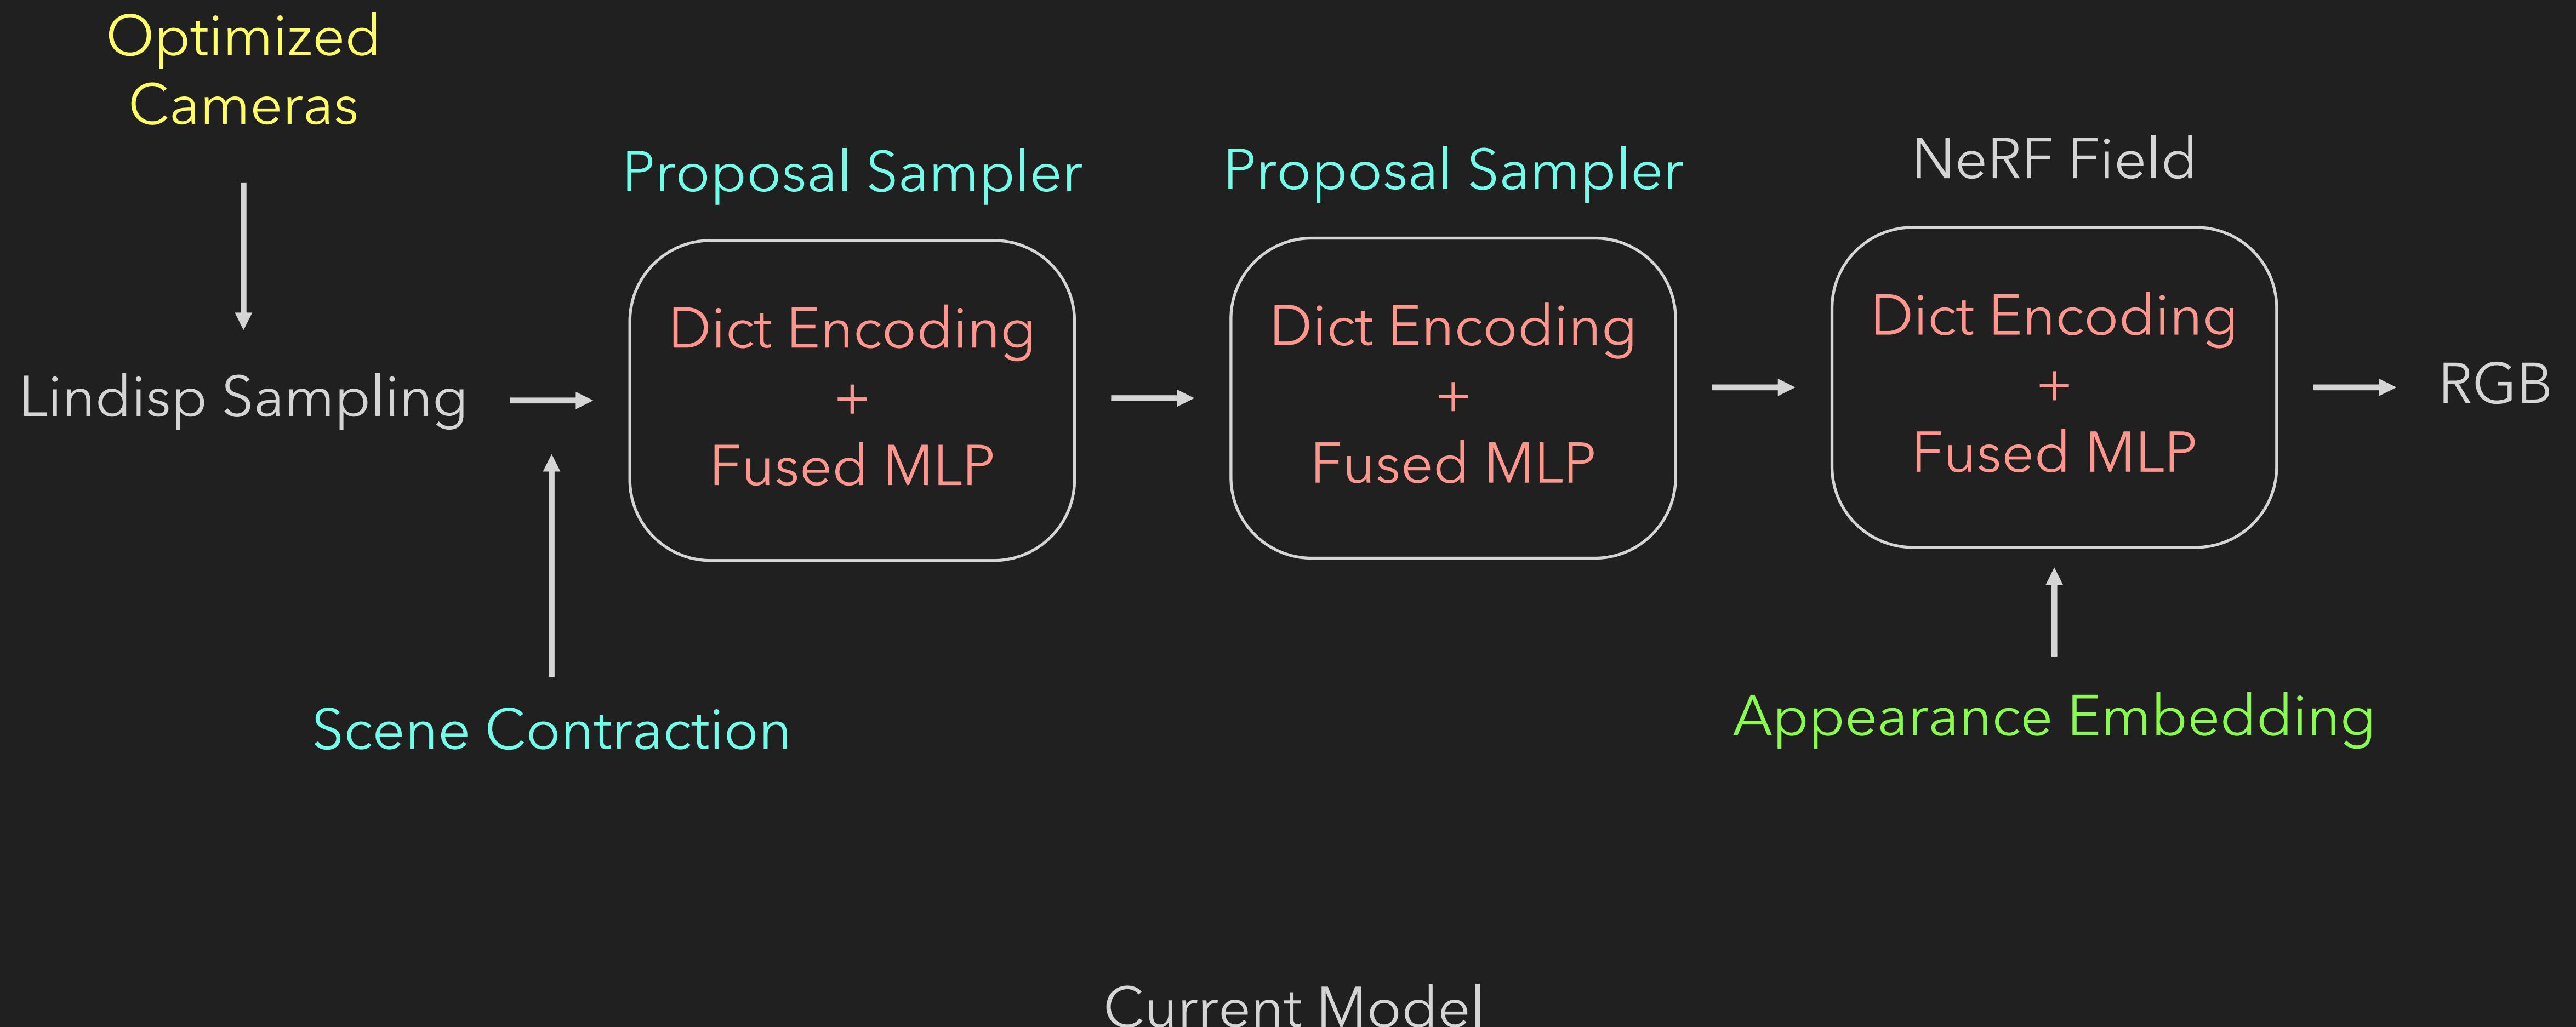
## Fields

- Fused MLP
- Voxel Grid

## Renderers

- RGB
- RGB-SH
- Depth
- Accumulation

# Case study: Nerfacto Model



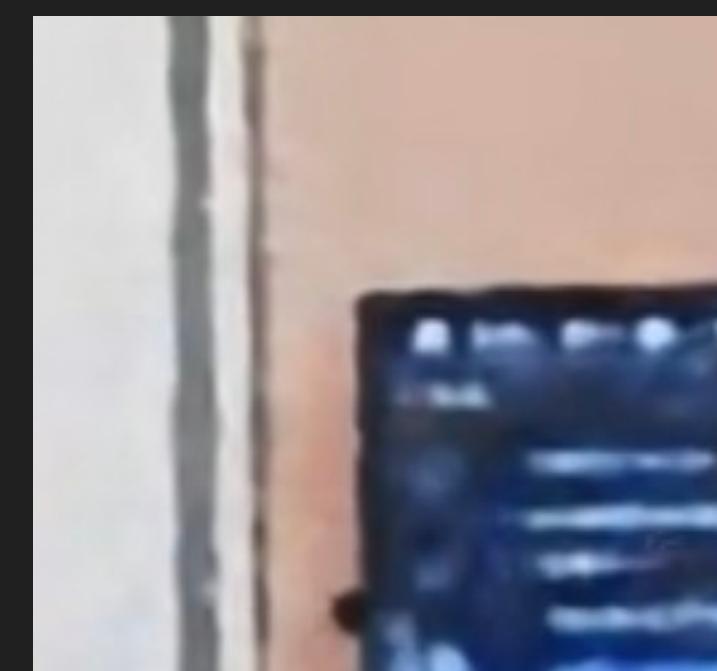
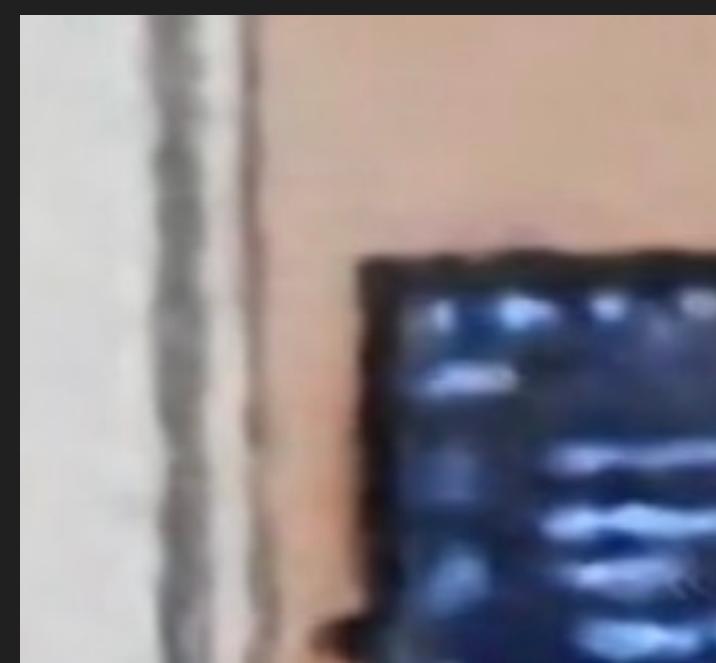
## Case study: Nerfacto Model

Since Release:

1.5x Faster training

3x Less Memory

Improved Quality



Before

After

## GETTING STARTED

[Installation](#)[Training your first model](#)[Using custom data](#)[Using the viewer](#)[Google Colab](#)[Contributing](#)

## NERFOLOGY

[Methods](#)[Model components](#)[Cameras models](#)[Sample representation](#)[Ray samplers](#)**Spatial distortions**[Encoders](#)

## DEVELOPER GUIDES



# Spatial Distortions

If you are trying to reconstruct an object floating in an empty void, you can stop reading. However if you are trying to reconstruct a scene or object from images, you may wish to consider adding a spatial distortion.

When rendering a target view of a scene, the camera will emit a camera ray for each pixel and query the scene at points along this ray. We can choose where to query these points using different [samplers](#). These samplers have some notion of *bounds* that define where the ray should start and terminate. If you know that everything in your scenes exists within some predefined bounds (ie. a cube that a room fits in) then the sampler will properly sample the entire space. If however the scene is unbounded (ie. an outdoor scene) defining where to stop sampling is challenging. One option to increase the far sampling distance to a large value (ie. 1km). Alternatively we can warp the space into a fixed volume. Below are supported distortions.

## Scene Contraction

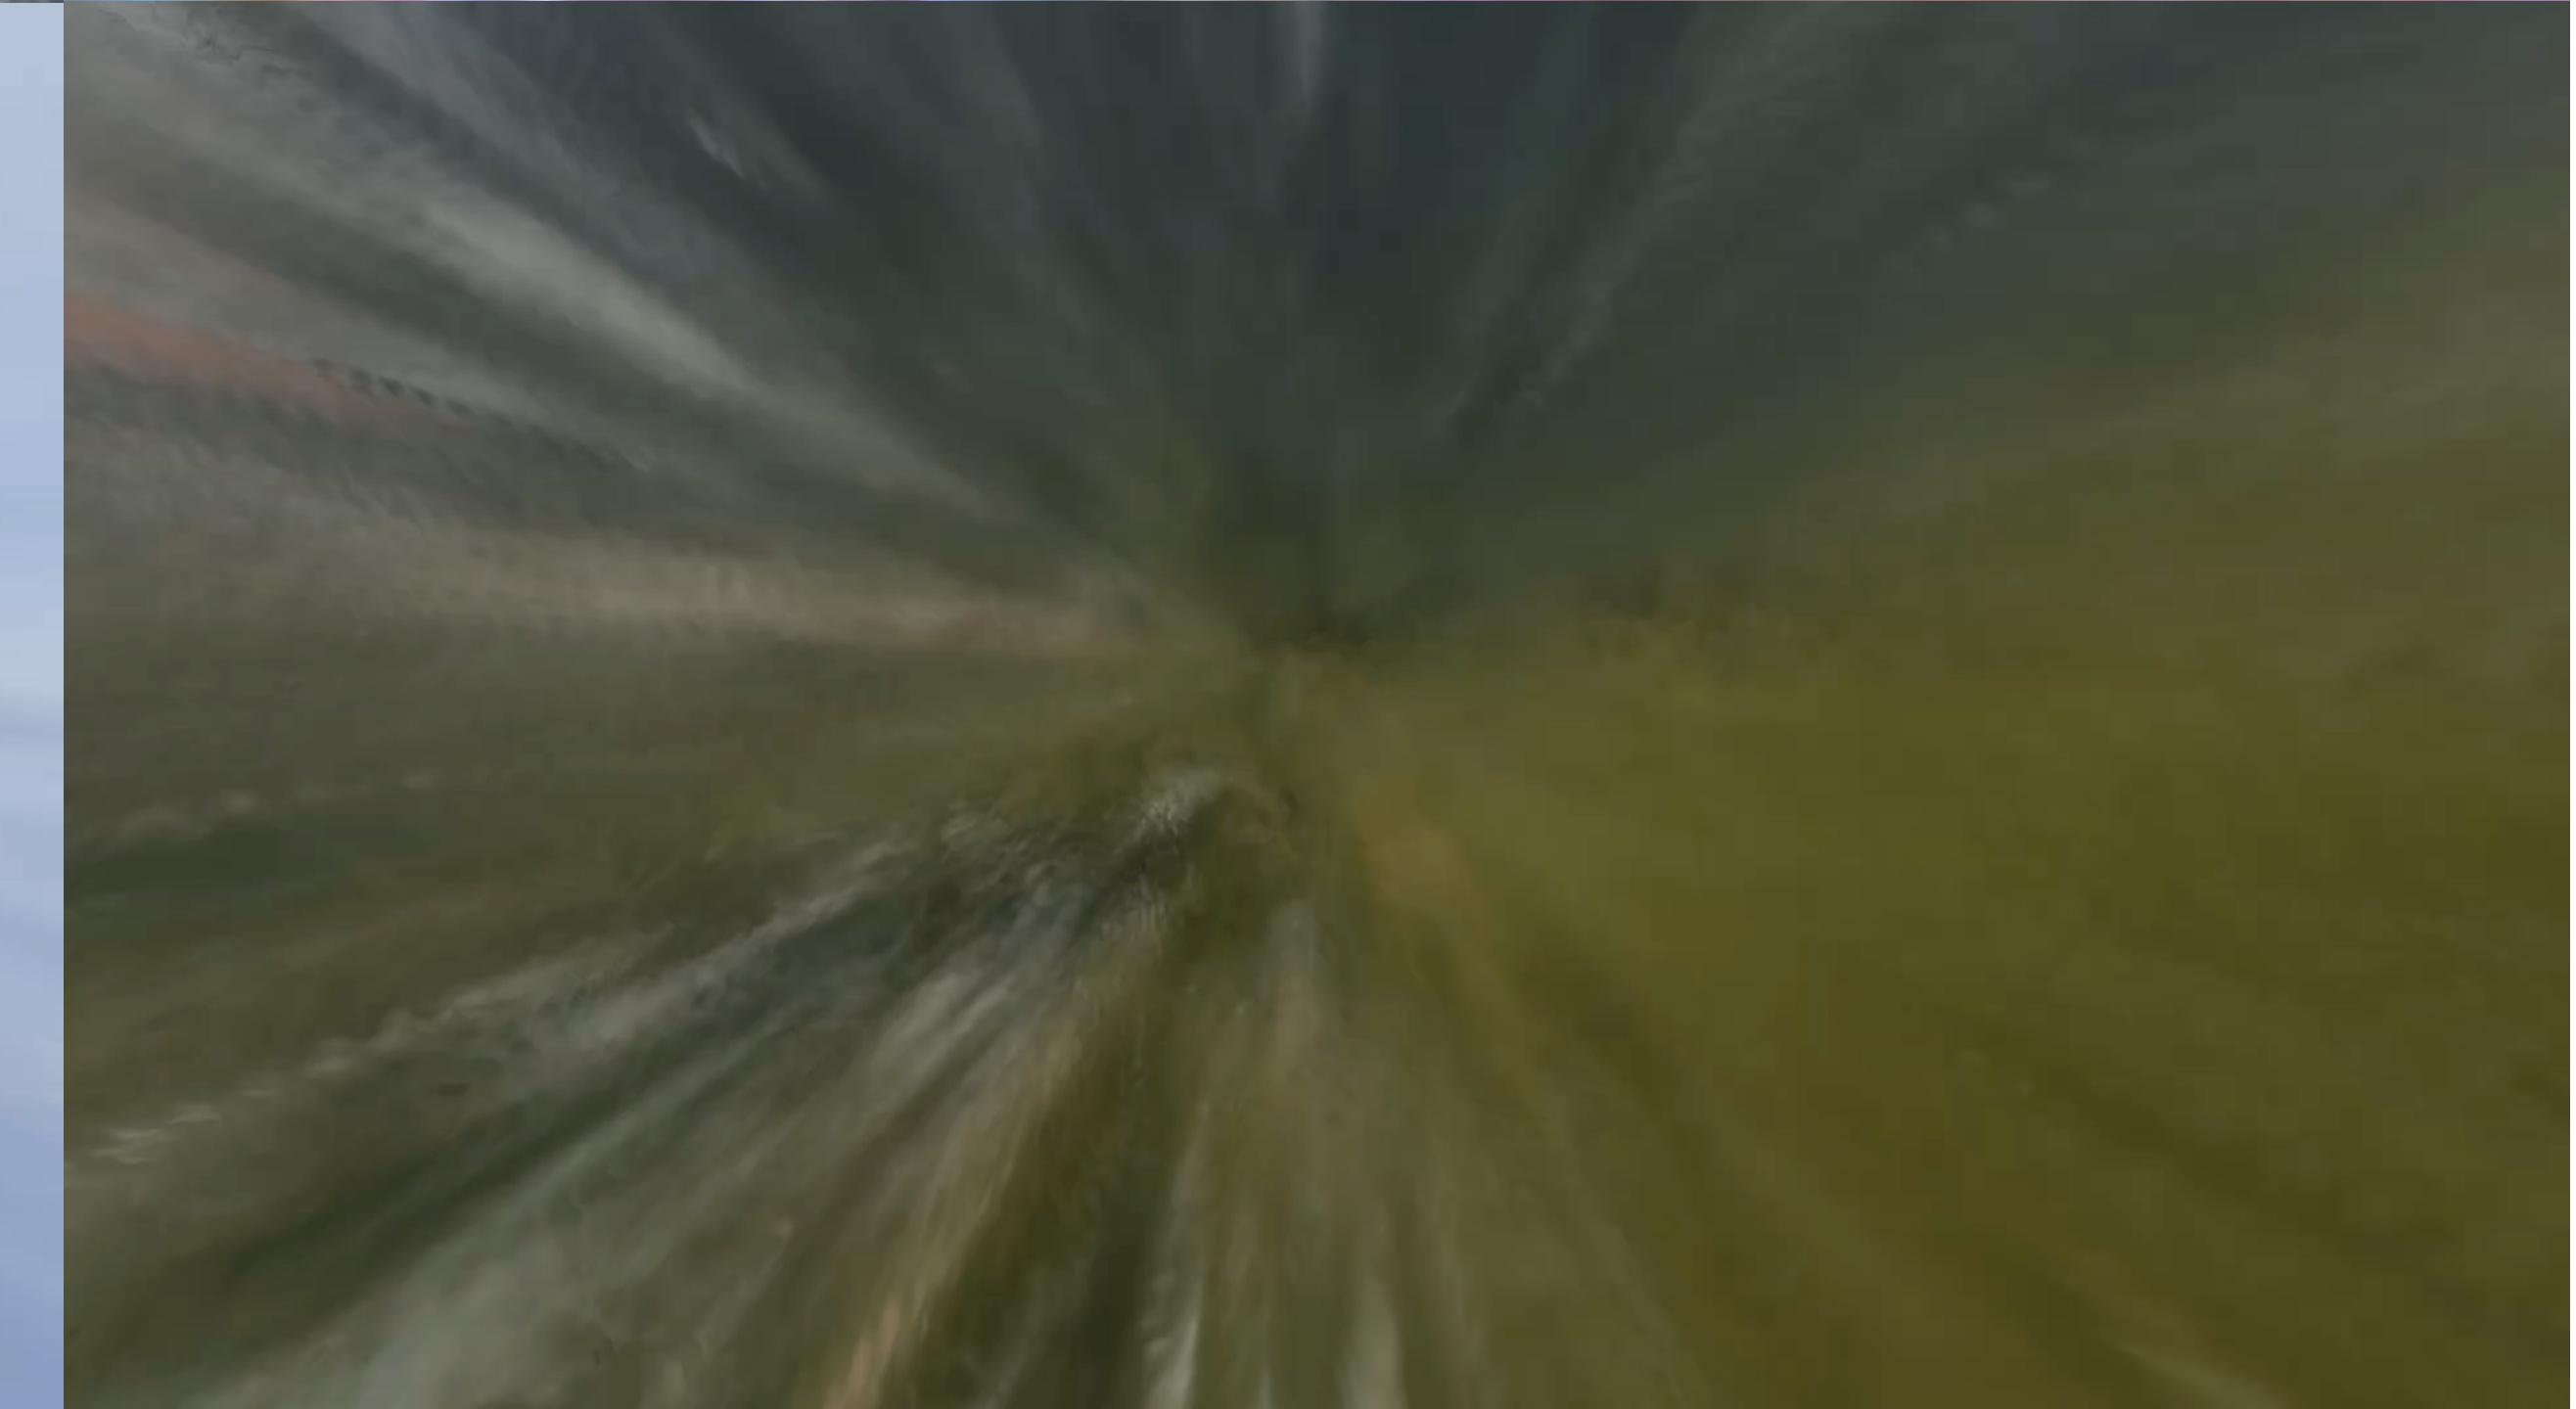
Contract unbounded space into a ball of radius 2. This contraction was proposed in [MipNeRF-360](#).

Samples within the unit ball are not modified, whereas sample outside the unit ball are contracted to fit within the ball of radius 2.

We use the following contraction equation:

$$f(x) = \begin{cases} x & \|x\| \leq 1 \\ (2 - \frac{1}{\|x\|})(\frac{x}{\|x\|}) & \|x\| > 1 \end{cases}$$

Below we visualize a ray before and after scene contraction. Visualized are 95% confidence intervals for the multivariate Gaussians for each sample location ([this guide](#) explains why the samples are represented by Gaussians). We are also visualizing both a unit sphere and a radius 2 sphere.





nerfstudio



[docs.nerf.studio](https://docs.nerf.studio)



Discord